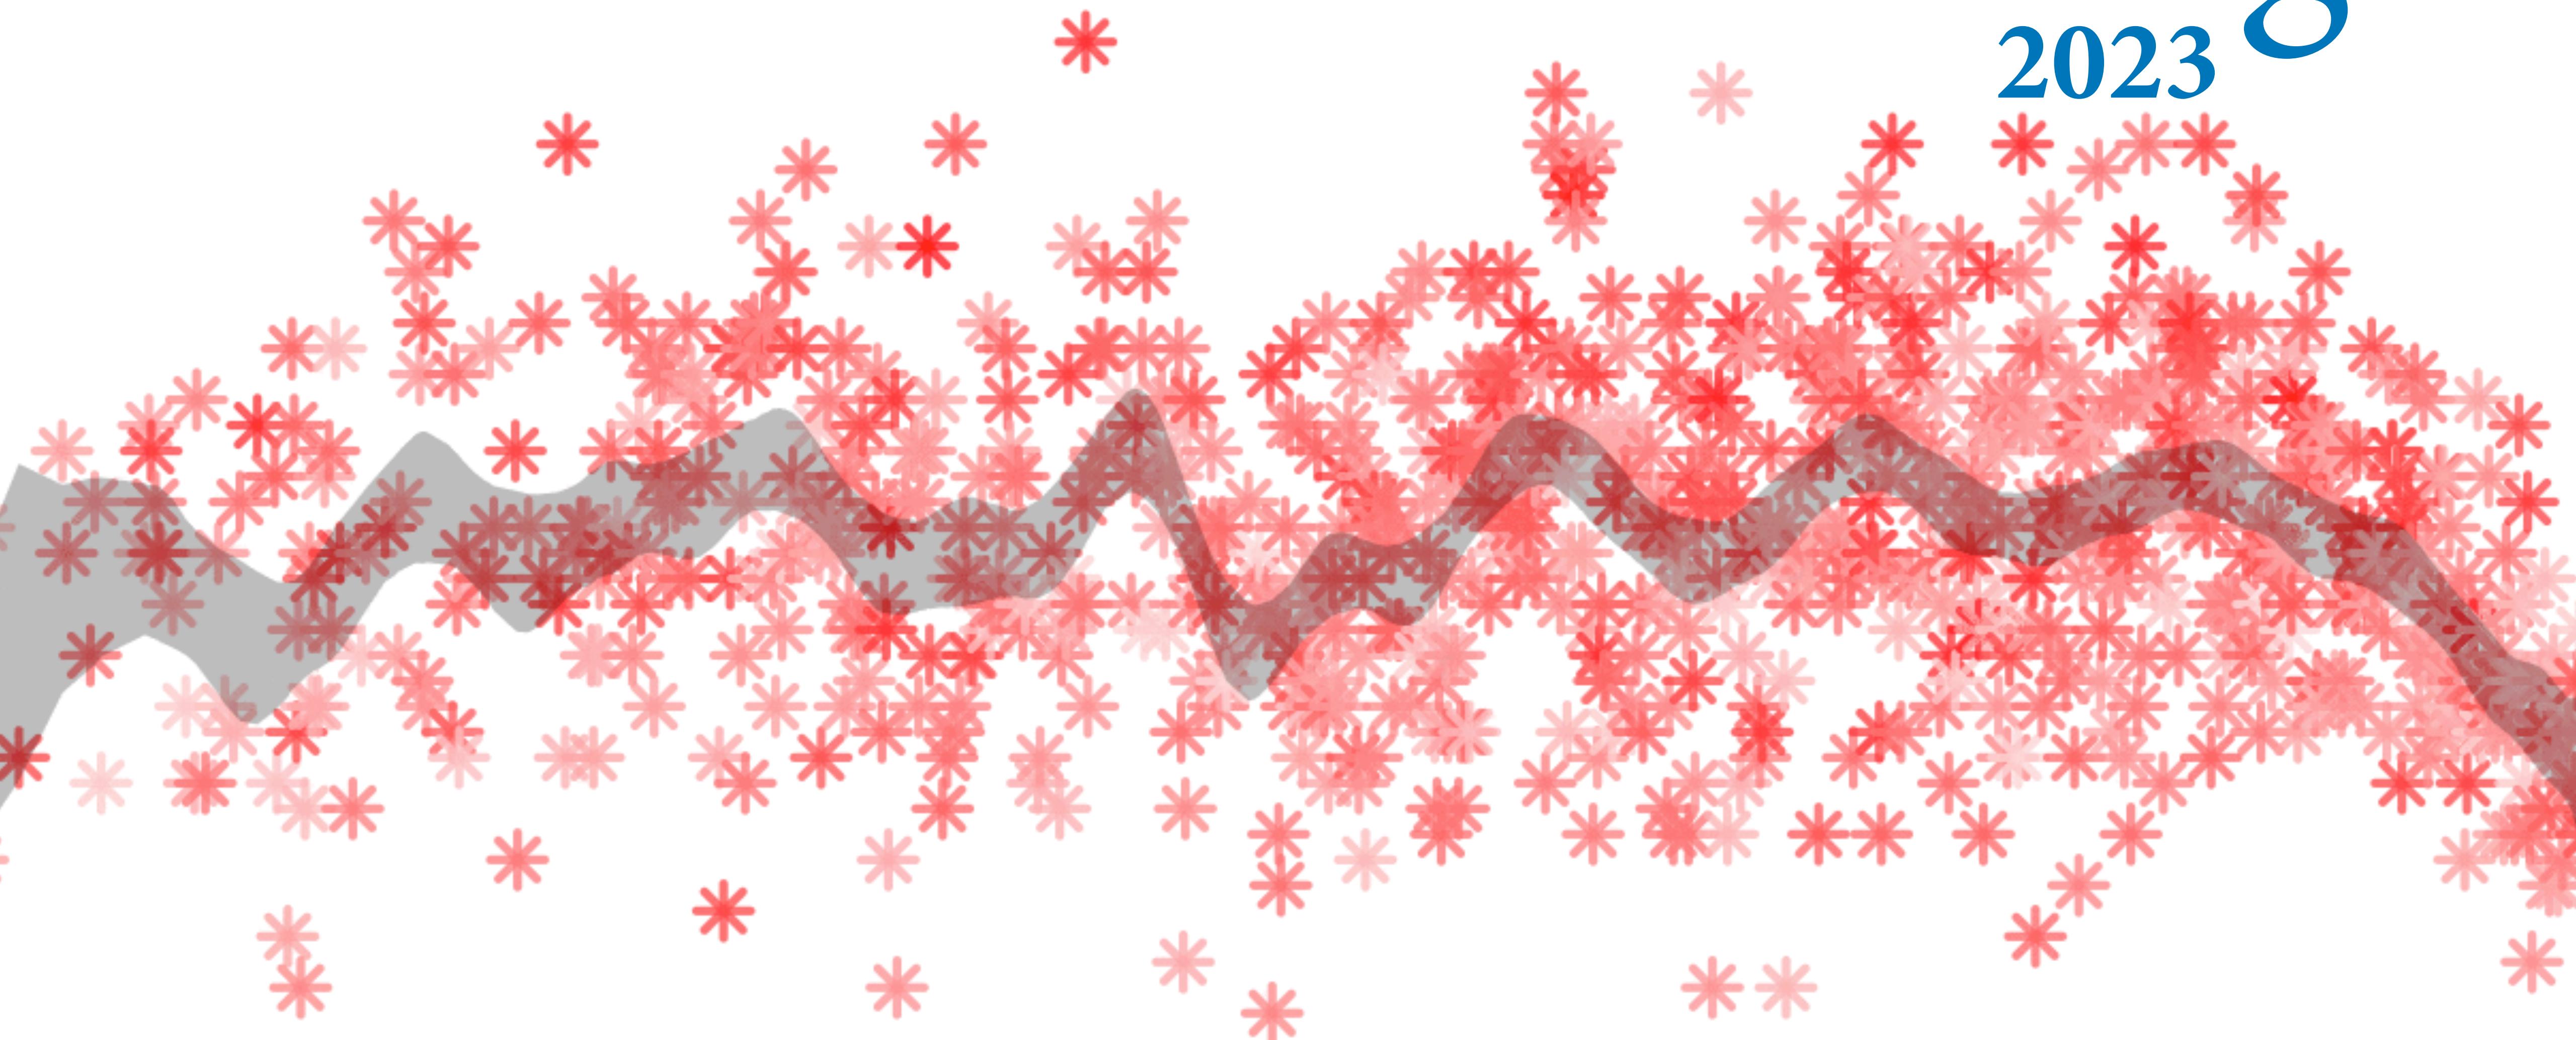


Statistical Rethinking

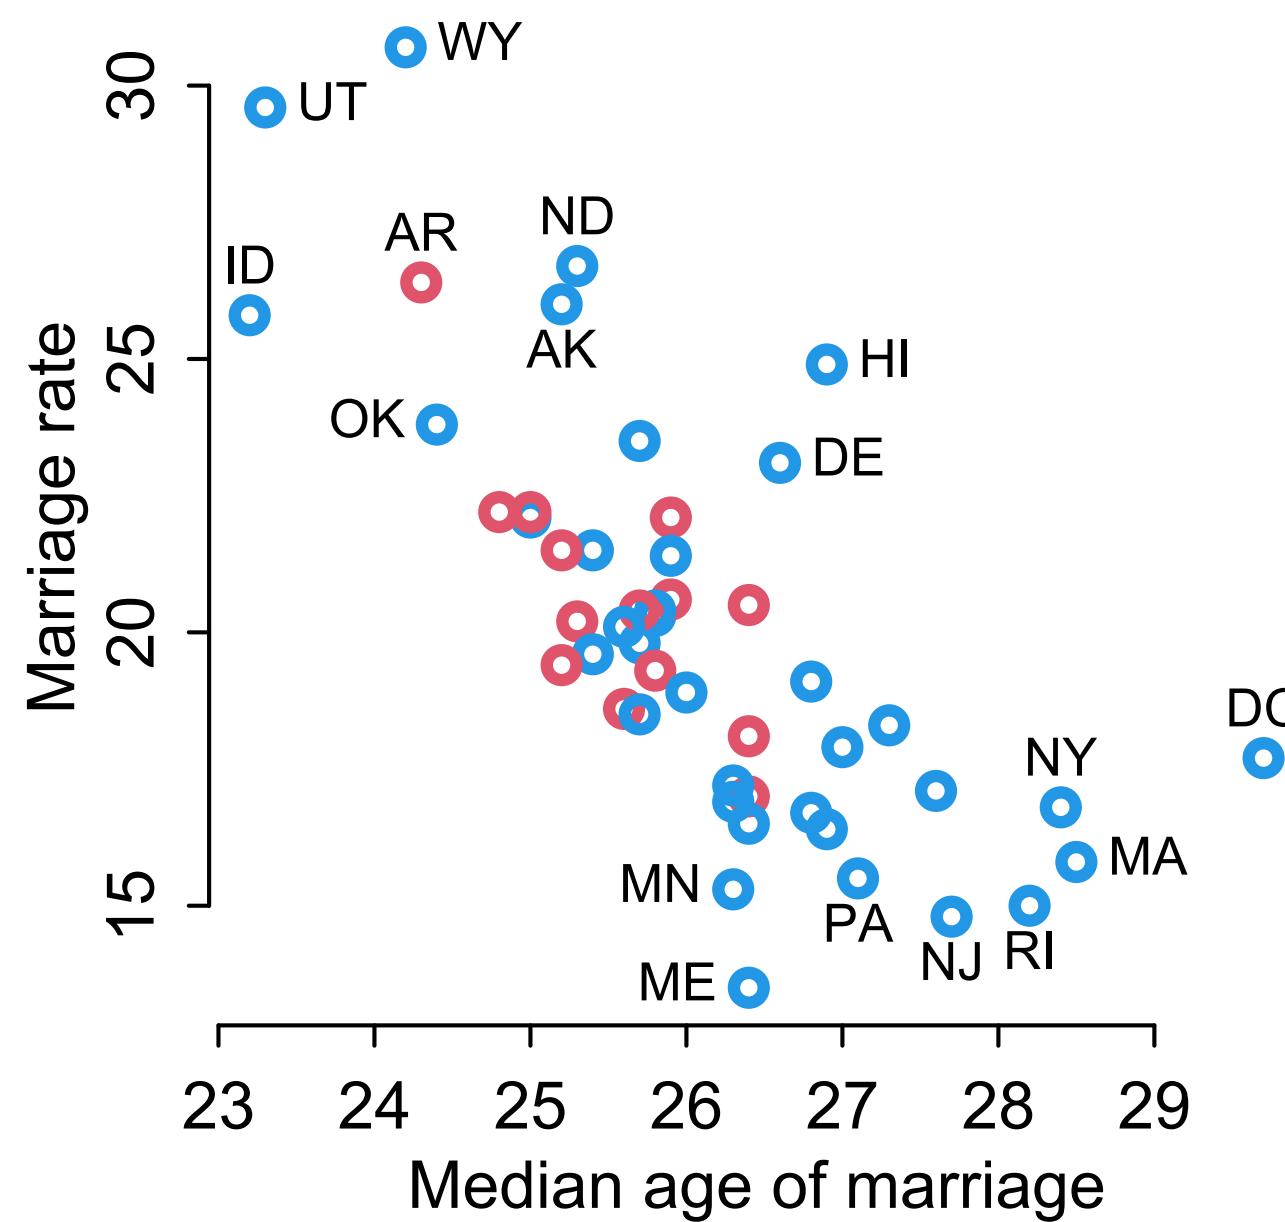
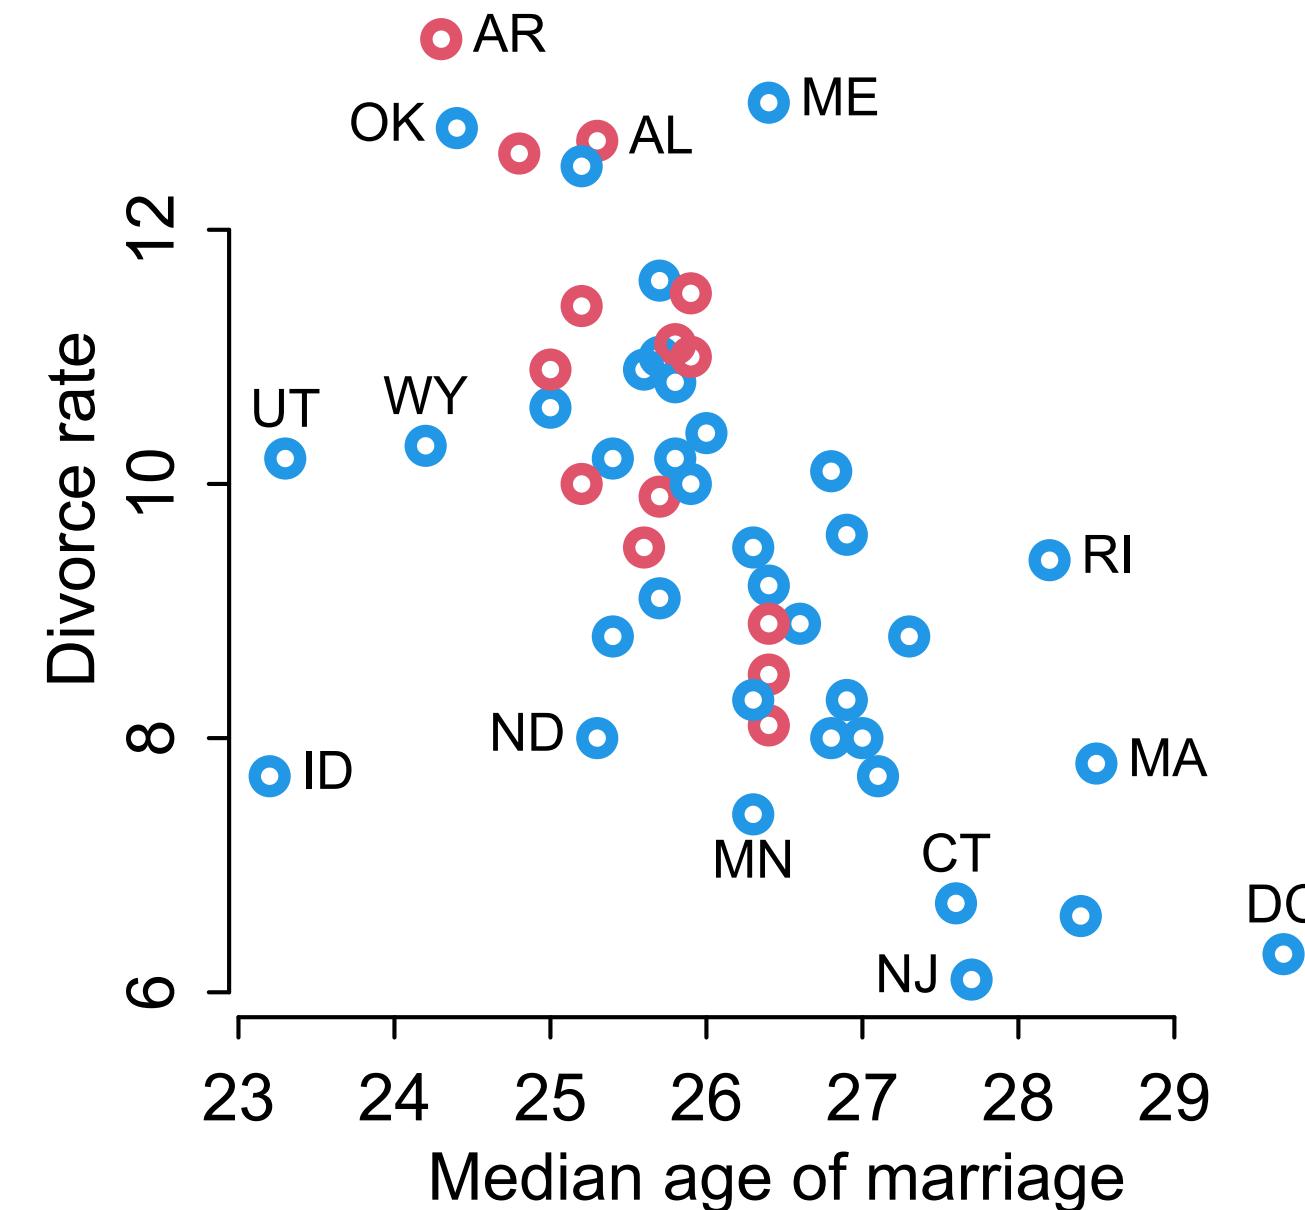
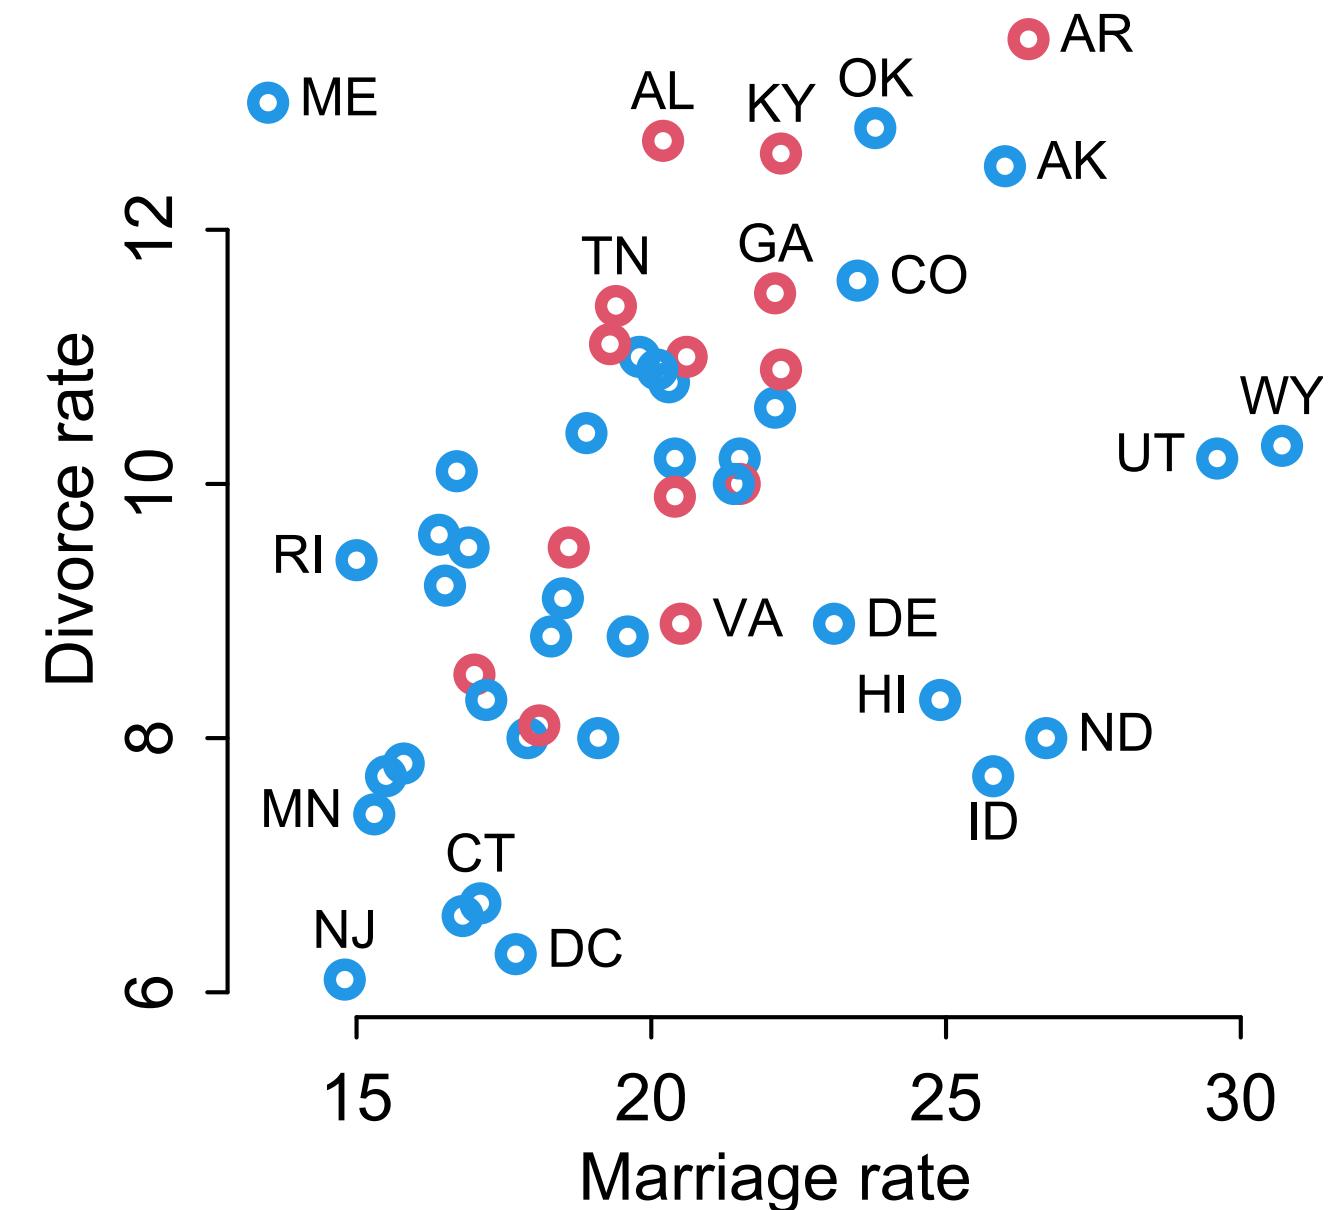
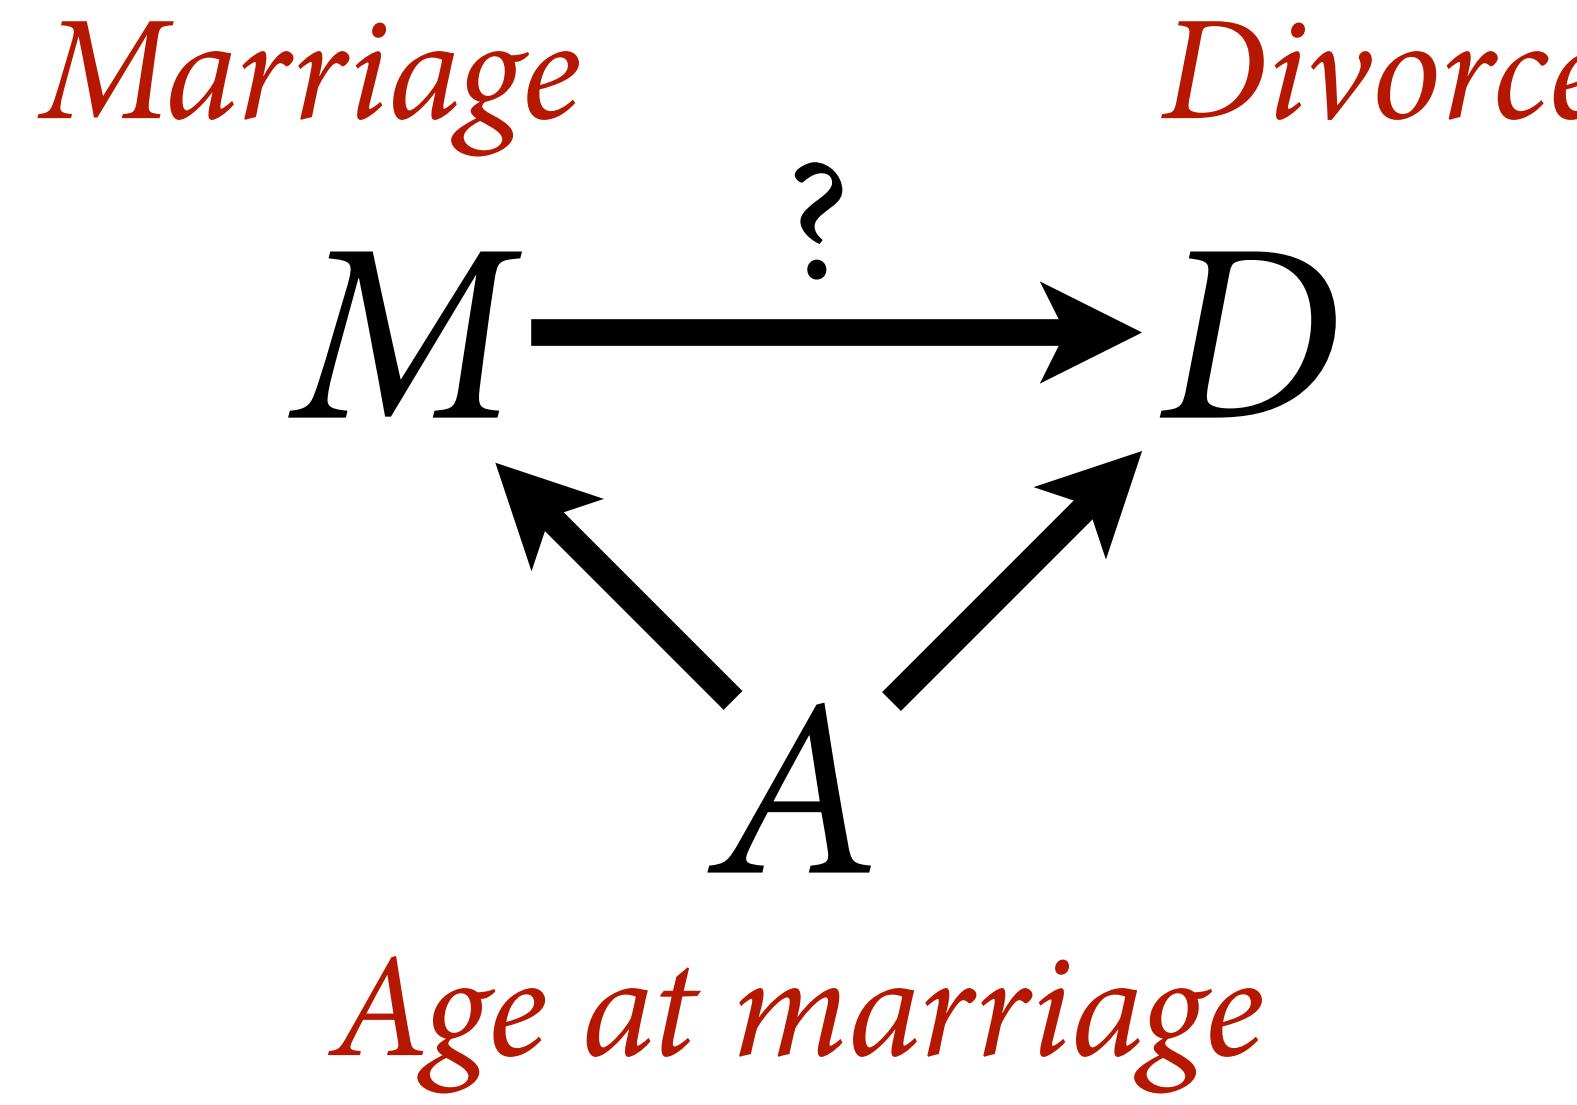
2023



8. Markov chain Monte Carlo

Table of Random Numbers

36518	36777	89116	05542	29705	83775	21564	81639	27973	62413	85652	62817	57881
46132	81380	75635	19428	88048	08747	20092	12615	35046	67753	69630	10883	13683
31841	77367	40791	97402	27569	90184	02338	39318	54936	34641	95525	86316	87384
84180	93793	64953	51472	65358	23701	75230	47200	78176	85248	90589	74567	22633
78435	37586	07015	98729	76703	16224	97661	79907	06611	26501	93389	92725	68158
41859	94198	37182	61345	88857	53204	86721	59613	67494	17292	94457	89520	77771
13019	07274	51068	93129	40386	51731	44254	66685	72835	01270	42523	45323	63481
82448	72430	29041	59208	95266	33978	70958	60017	39723	00606	17956	19024	15819
25432	96593	83112	96997	55340	80312	78839	09815	16887	22228	06206	54272	83516
69226	38655	03811	08342	47863	02743	11547	38250	58140	98470	24364	99797	73498
25837	68821	66426	20496	84843	18360	91252	99134	48931	99538	21160	09411	44659
38914	82707	24769	72026	56813	49336	71767	04474	32909	74162	50404	68562	14088
04070	60681	64290	26905	65617	76039	91657	71362	32246	49595	50663	47459	57072
01674	14751	28637	86980	11951	10479	41454	48527	53868	37846	85912	15156	00865
70294	35450	39982	79503	34382	43186	69890	63222	30110	56004	04879	05138	57476
73903	98066	52136	89925	50000	96334	30773	80571	31178	52799	41050	76298	43995
87789	56408	77107	88452	80975	03406	36114	64549	79244	82044	00202	45727	35709
92320	95929	58545	70699	07679	23296	03002	63885	54677	55745	52540	62154	33314
46391	60276	92061	43591	42118	73094	53608	58949	42927	90993	46795	05947	01934
67090	45063	84584	66022	48268	74971	94861	61749	61085	81758	89640	39437	90044
11666	99916	35165	29420	73213	15275	62532	47319	39842	62273	94980	23415	64668
40910	59068	04594	94576	51187	54796	17411	56123	66545	82163	61868	22752	40101
41169	37965	47578	92180	05257	19143	77486	02457	00985	31960	39033	44374	28352



*Marriage
observed*

M^*

M

D

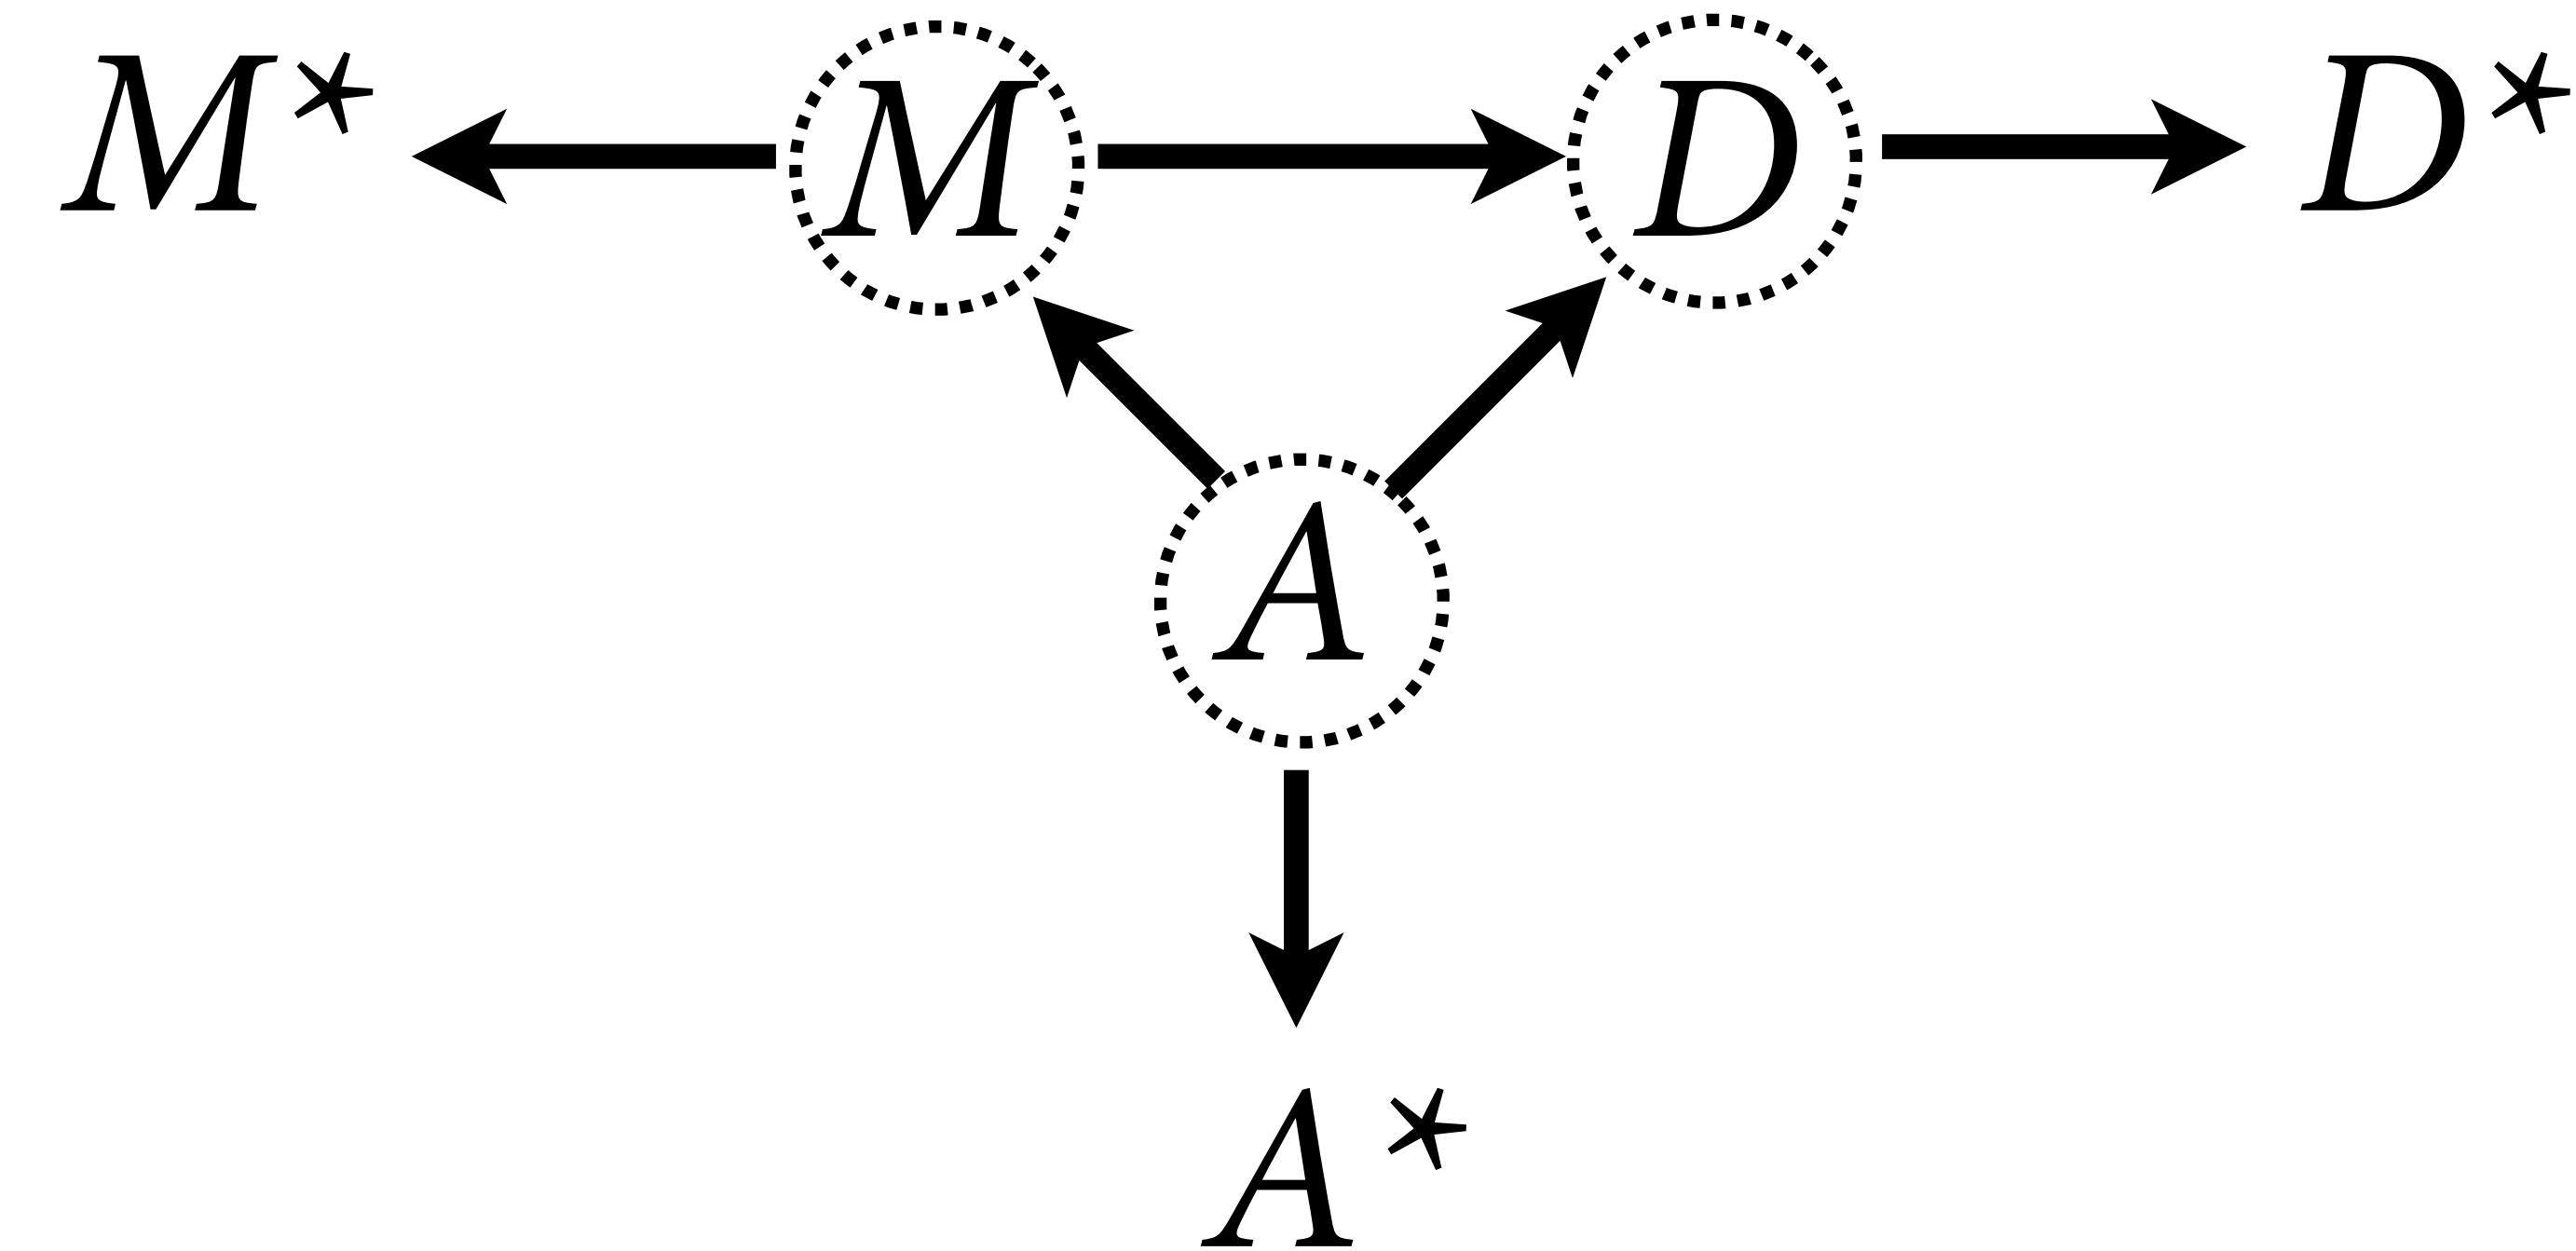
*Divorce
observed*

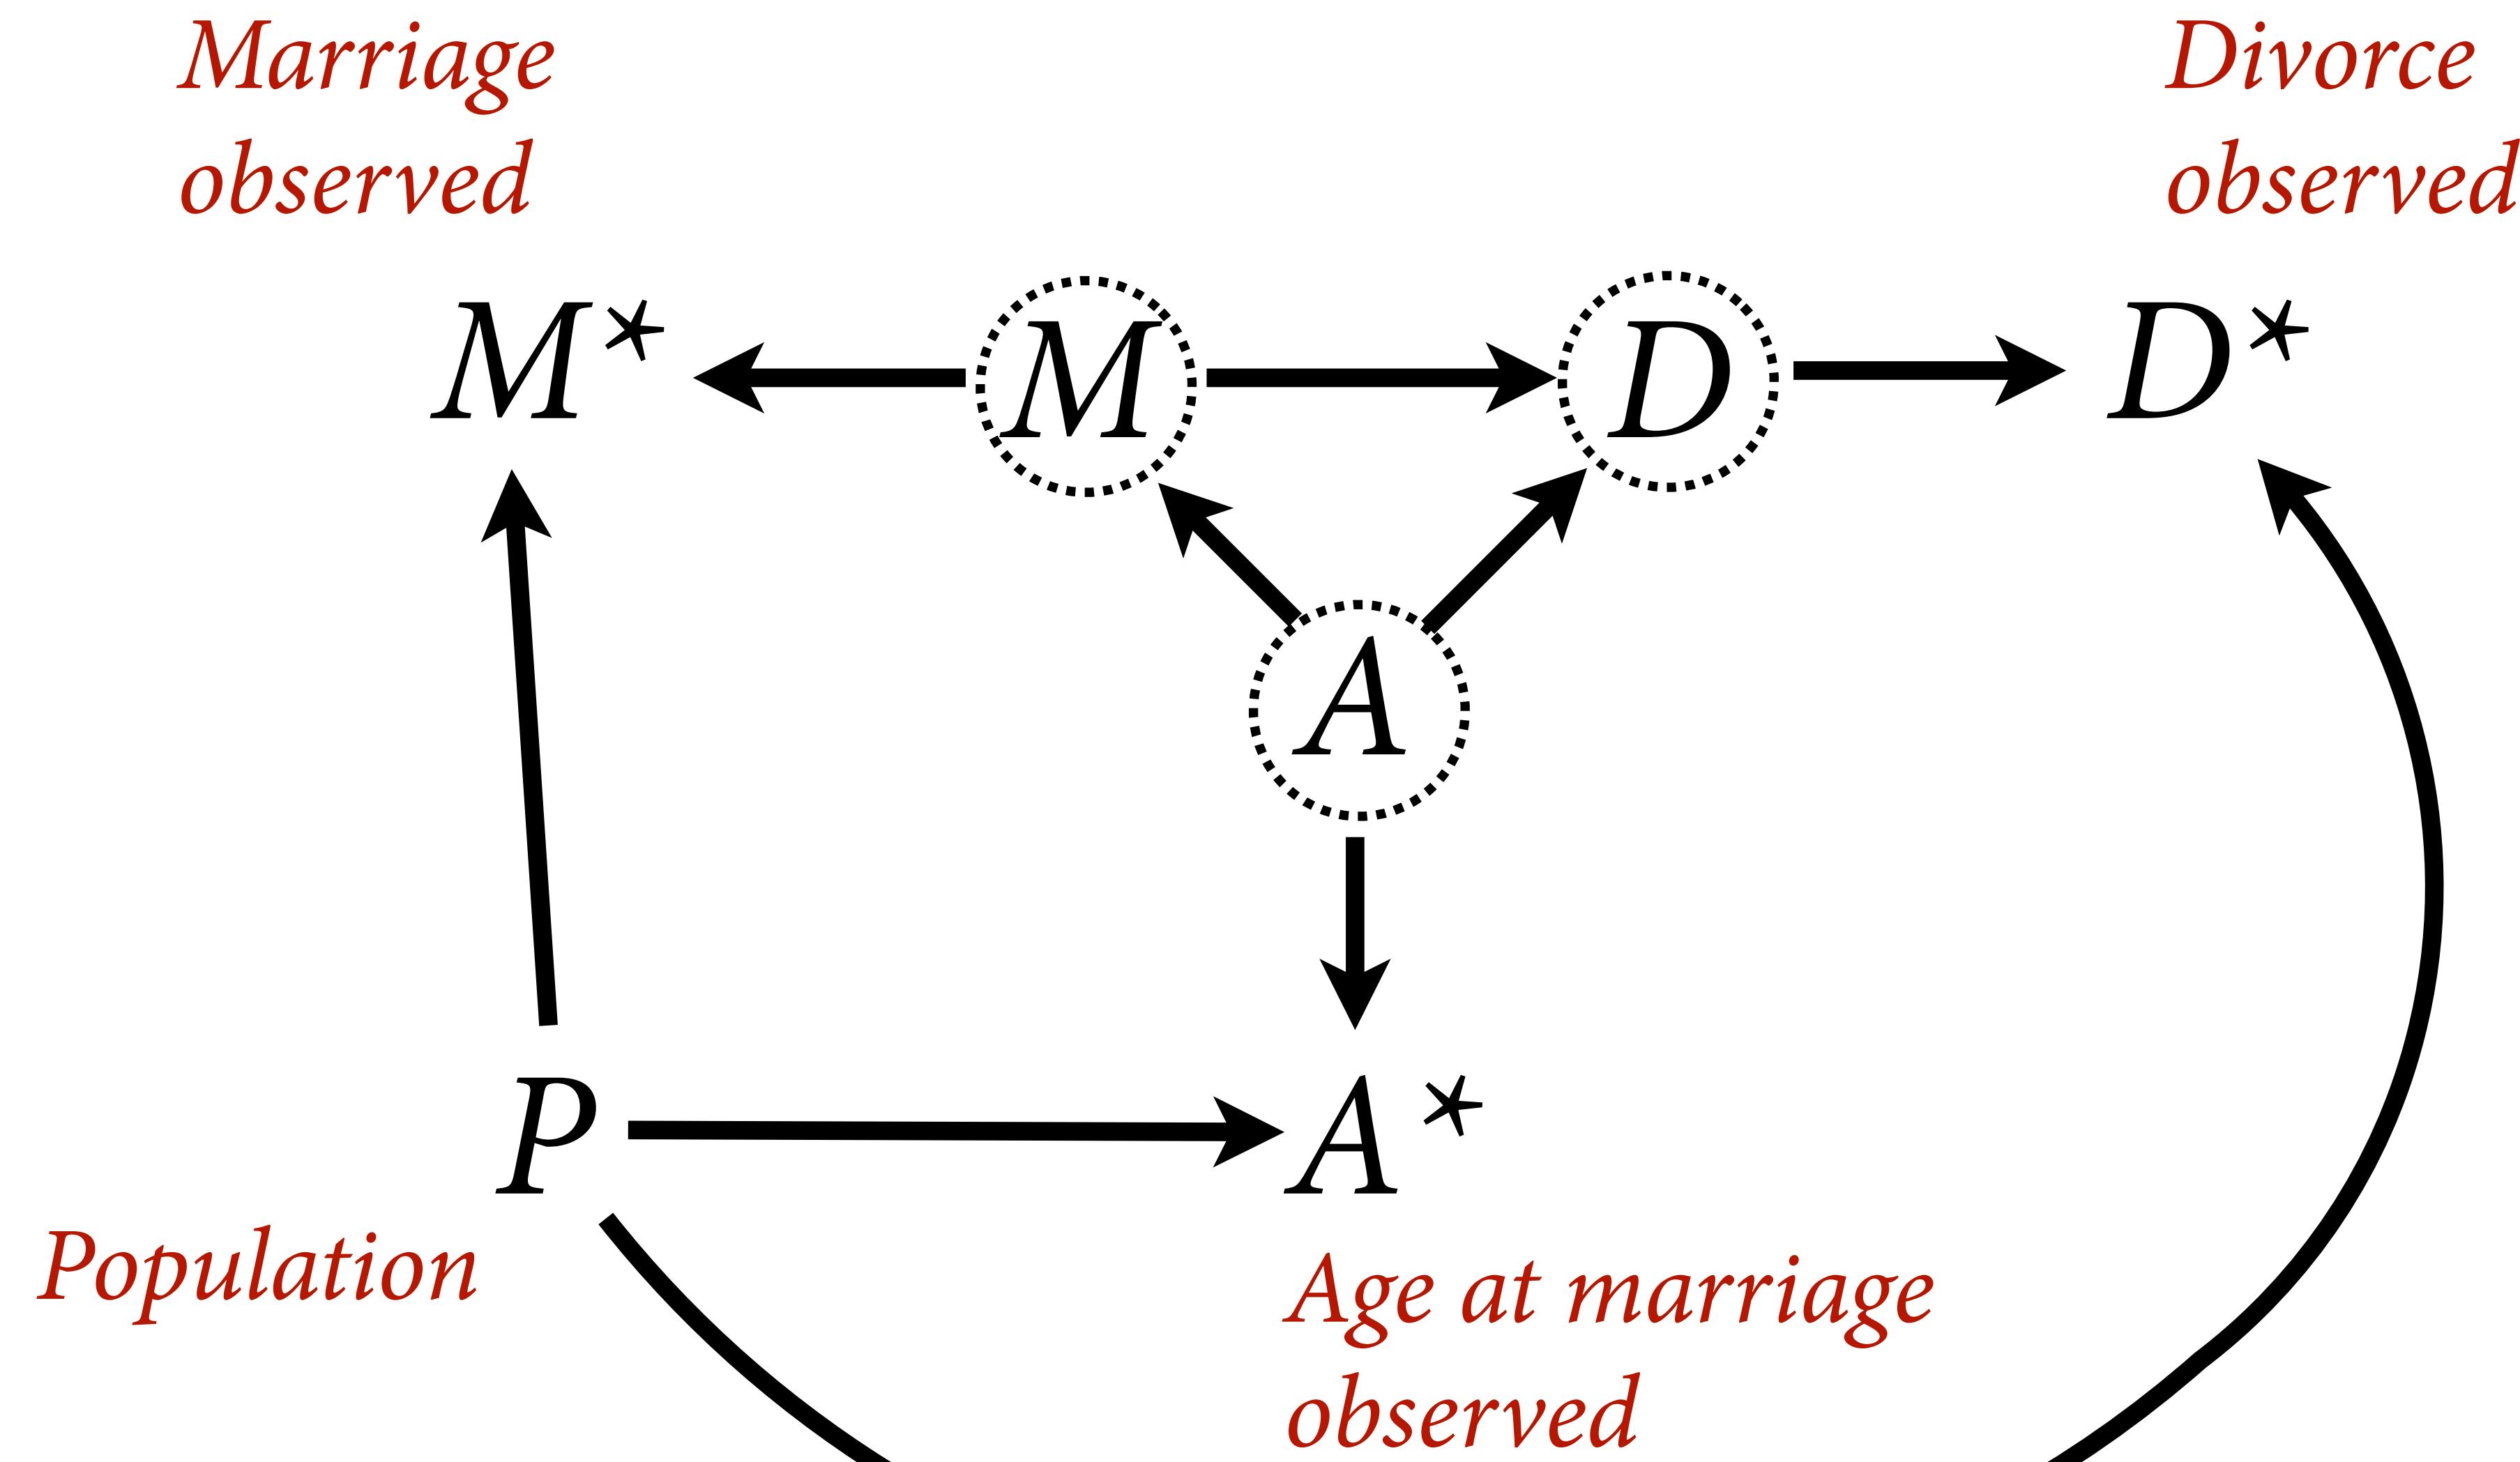
D^*

A

A^*

*Age at marriage
observed*



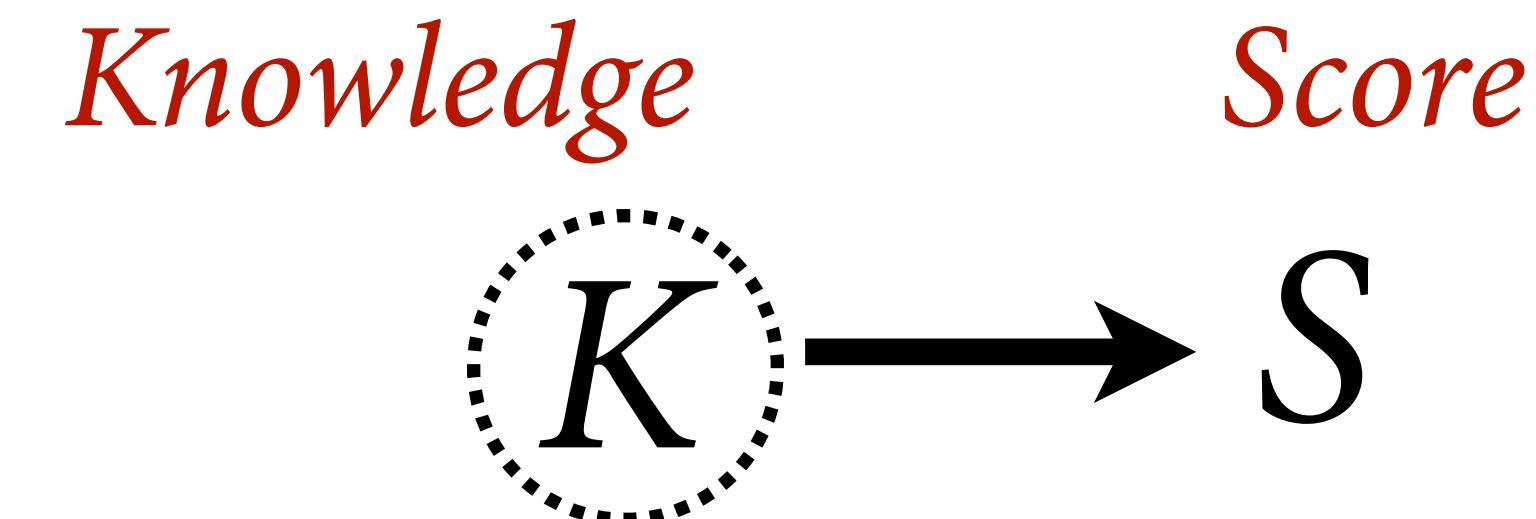


Real, latent modeling problems

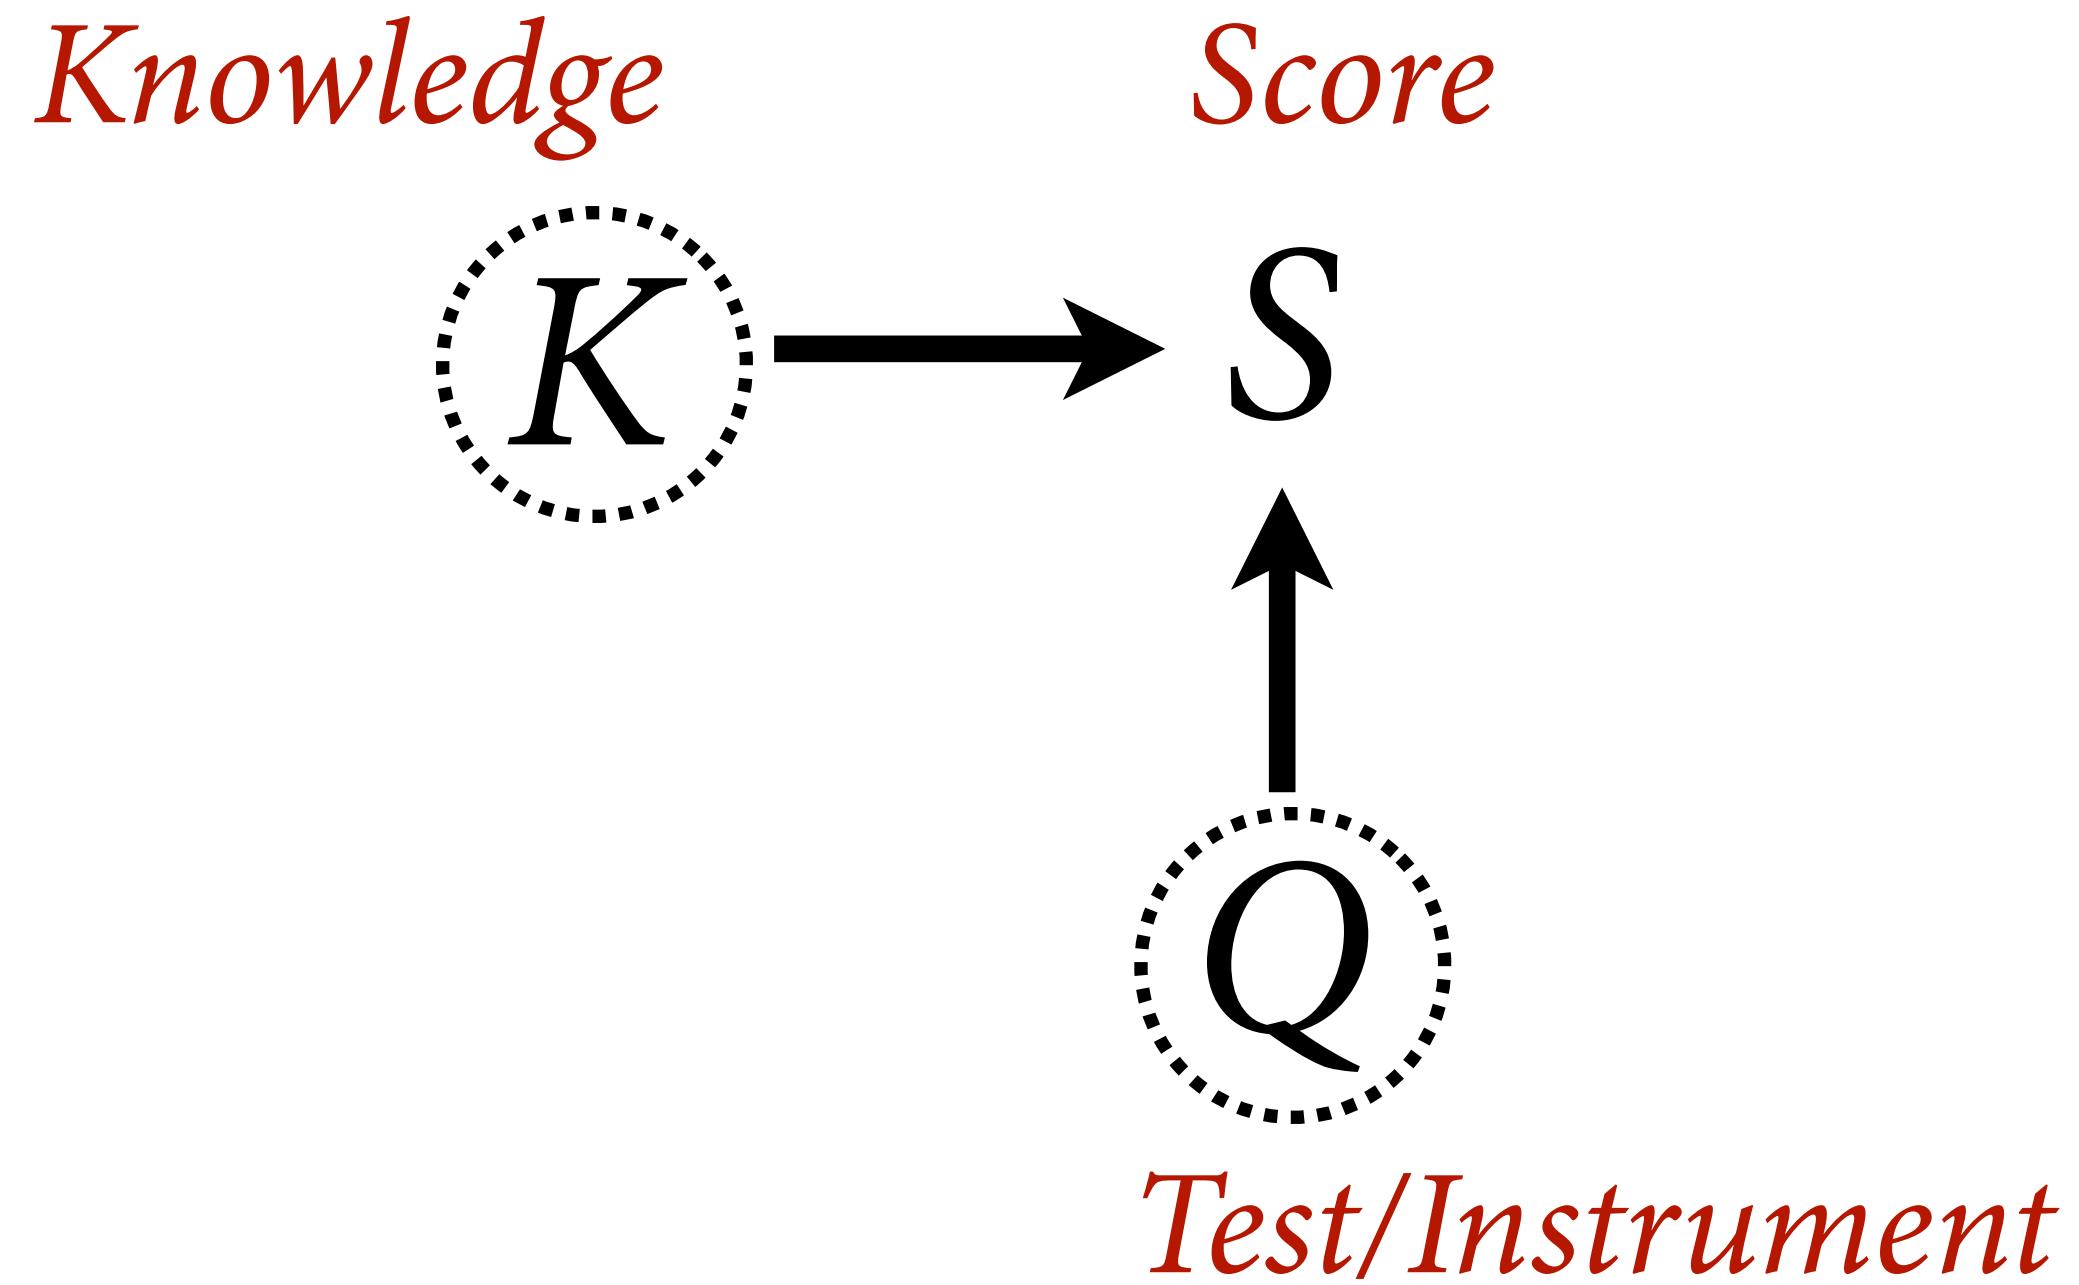
Knowledge

K

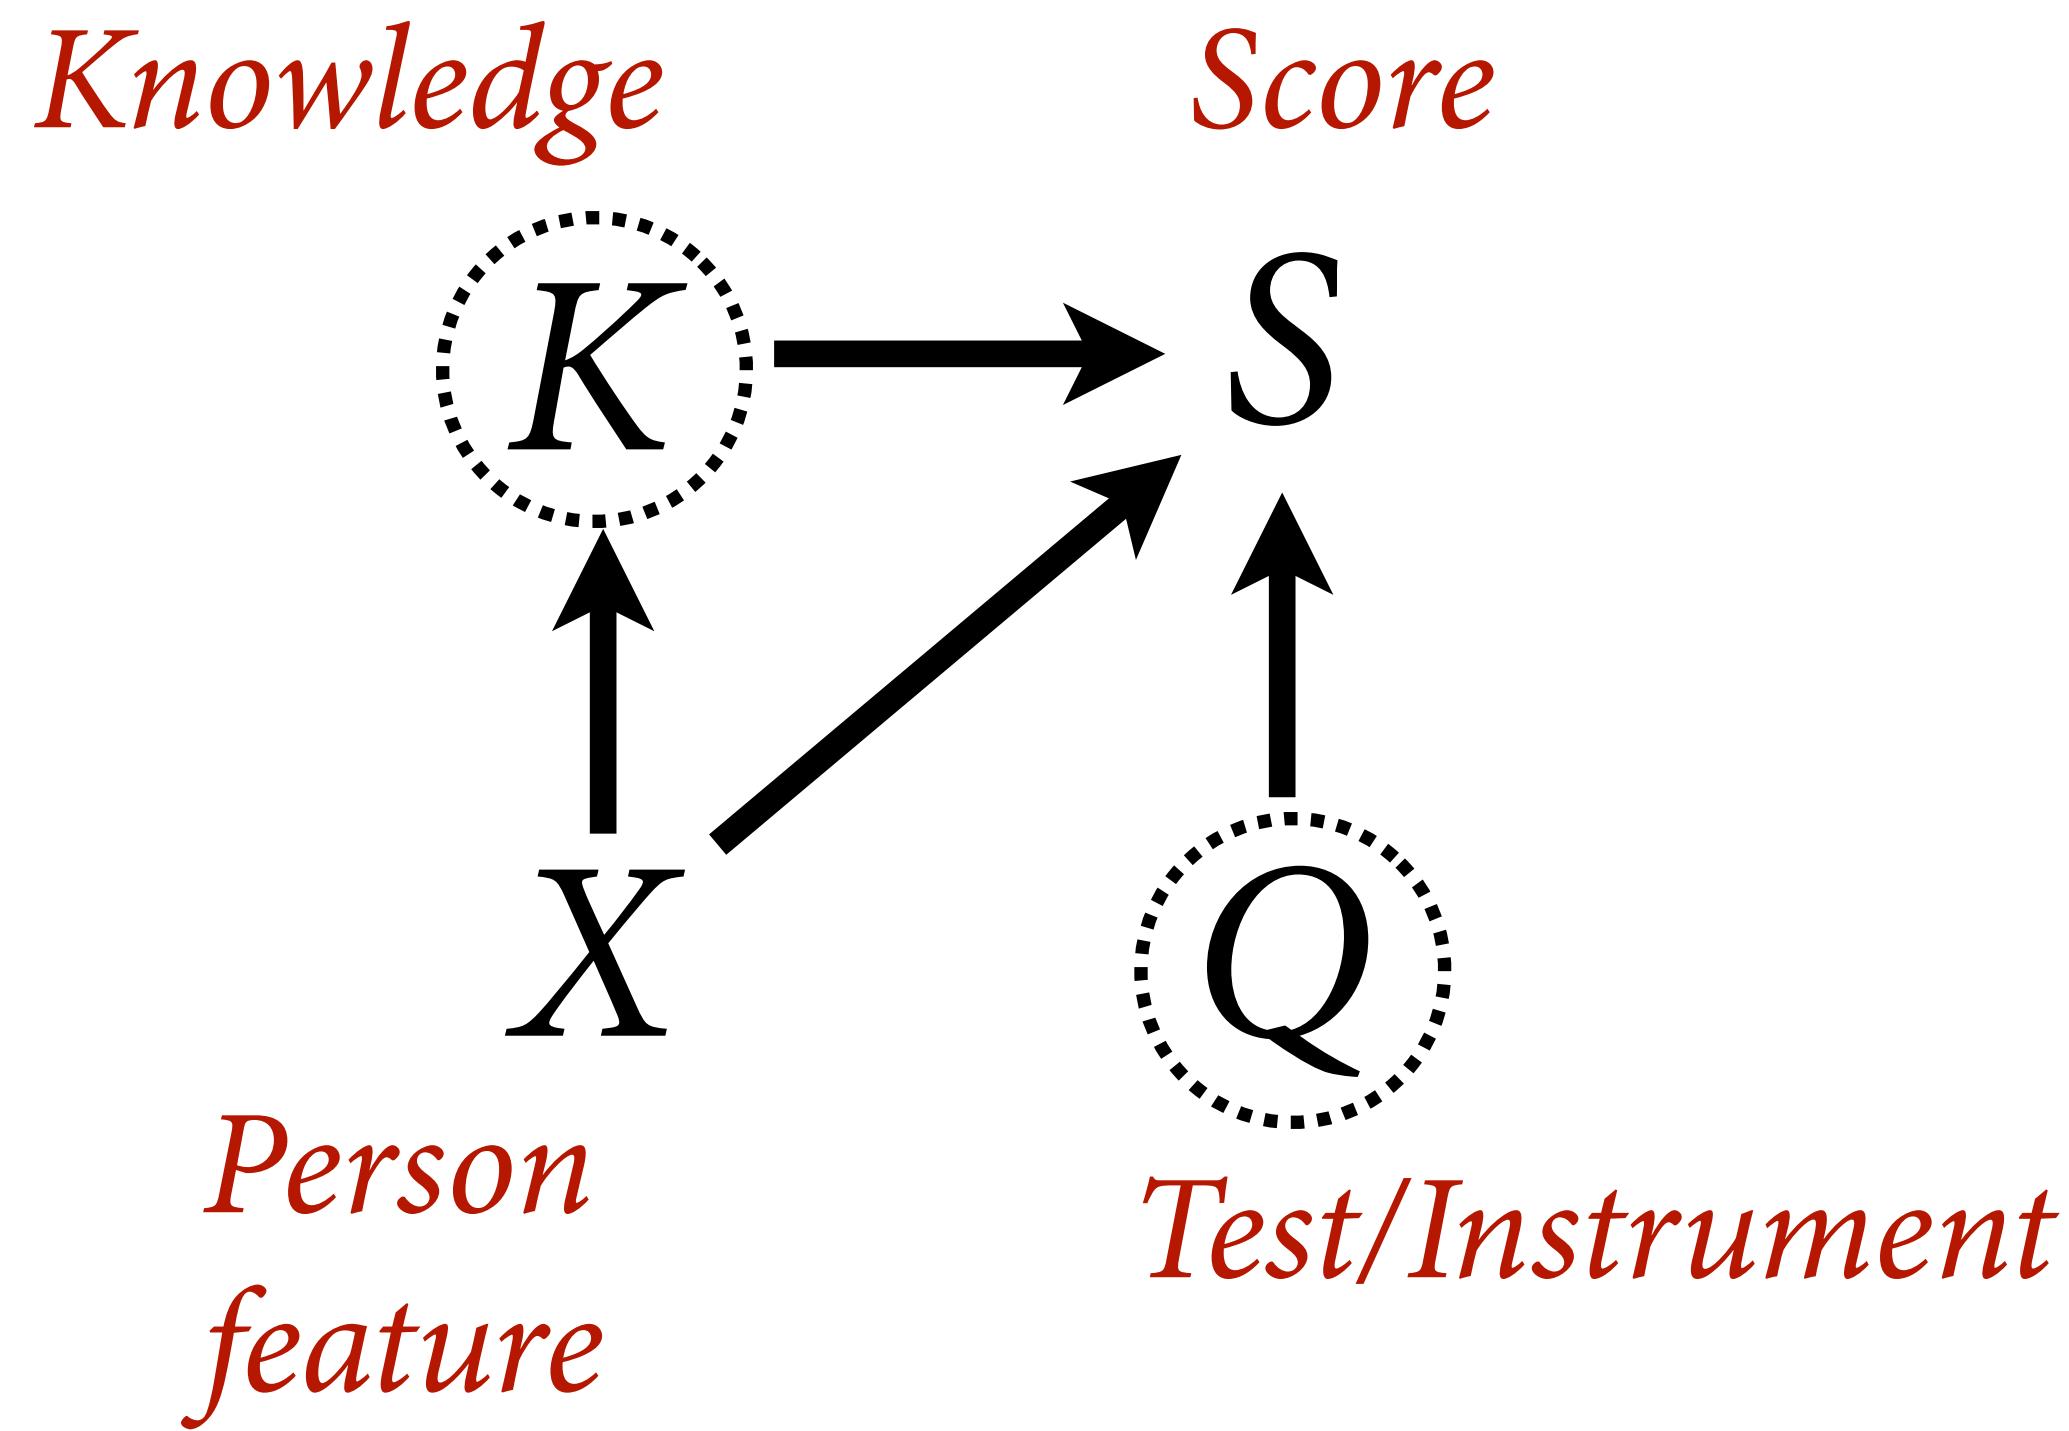
Real, latent modeling problems



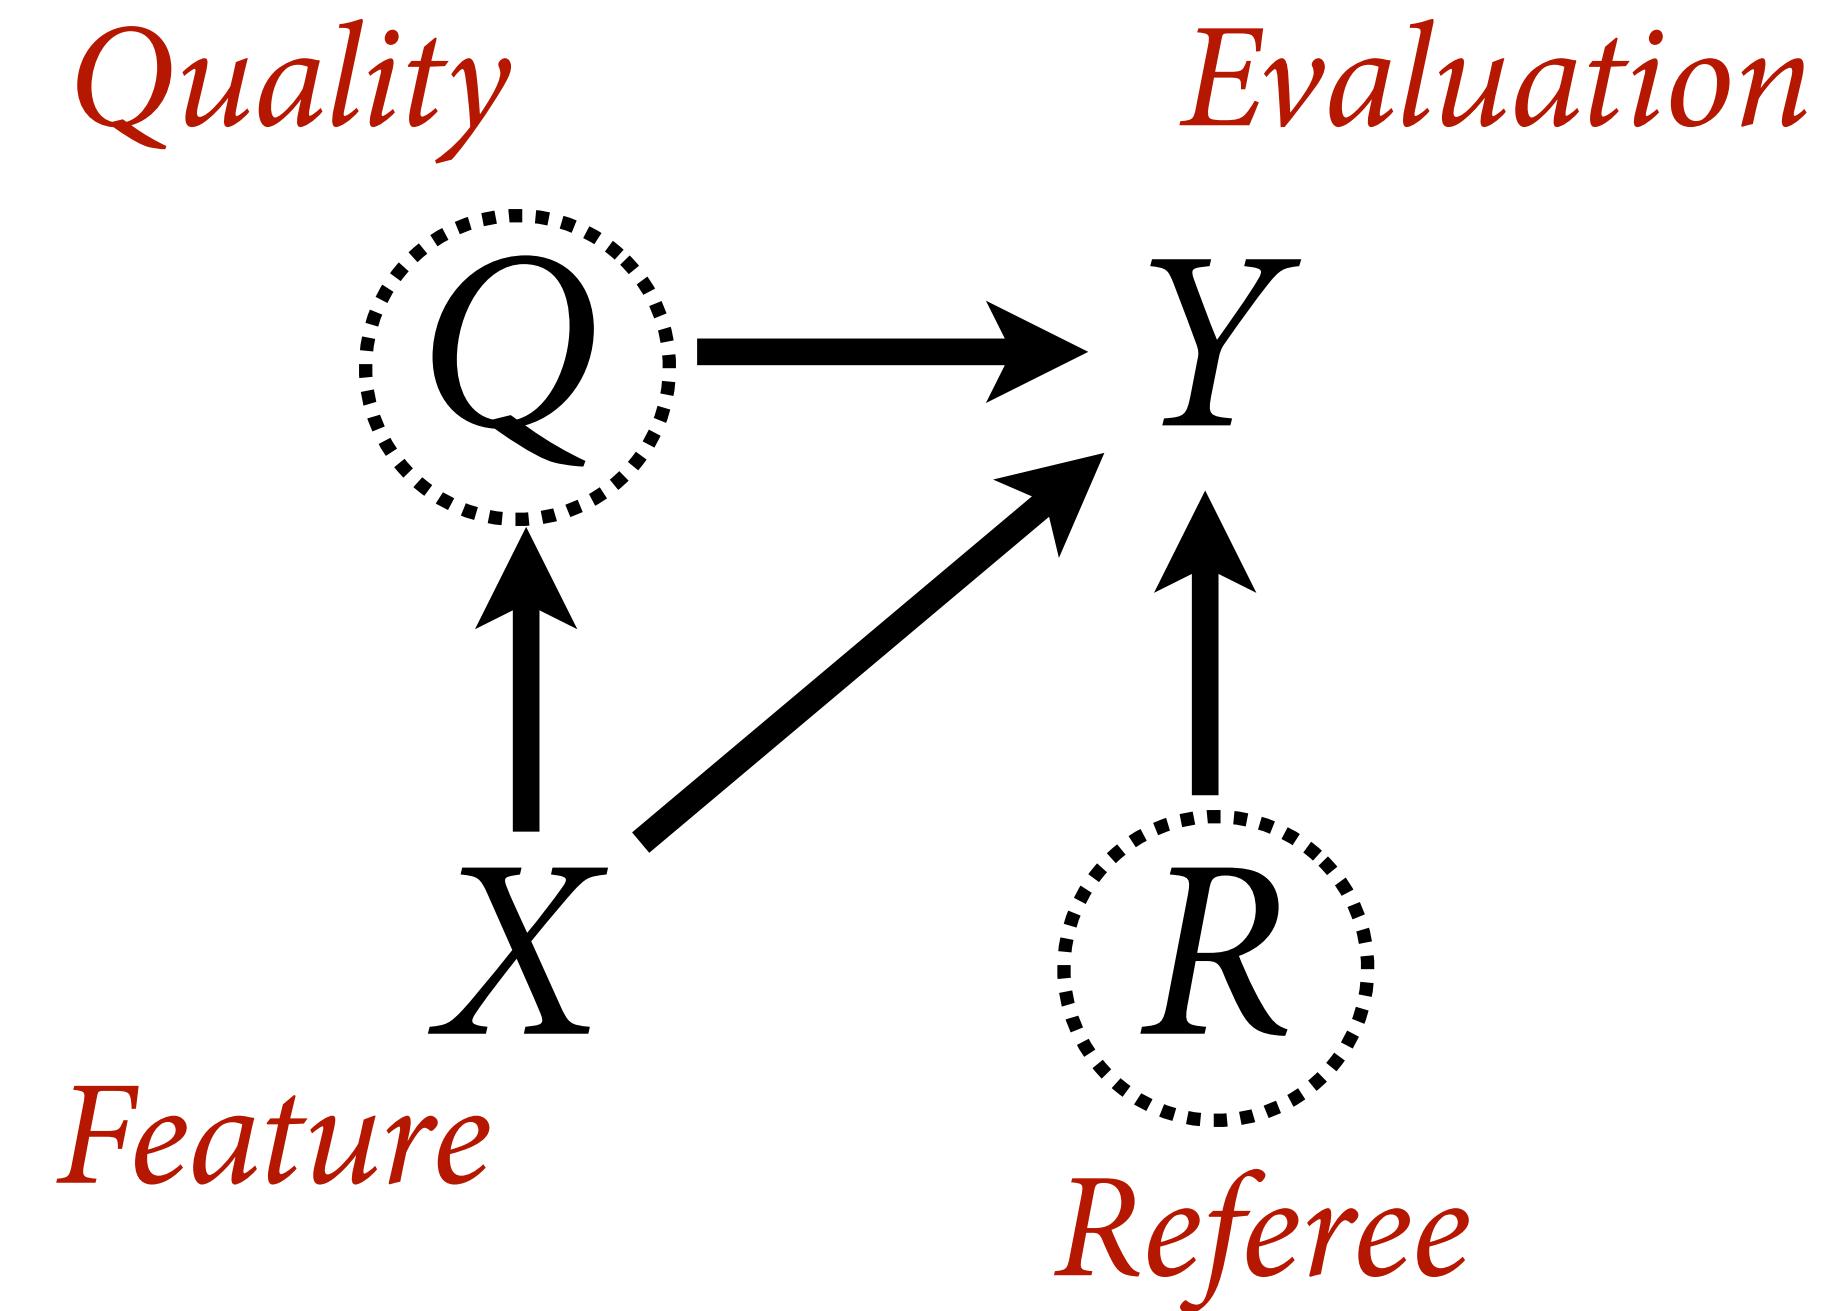
Real, latent modeling problems



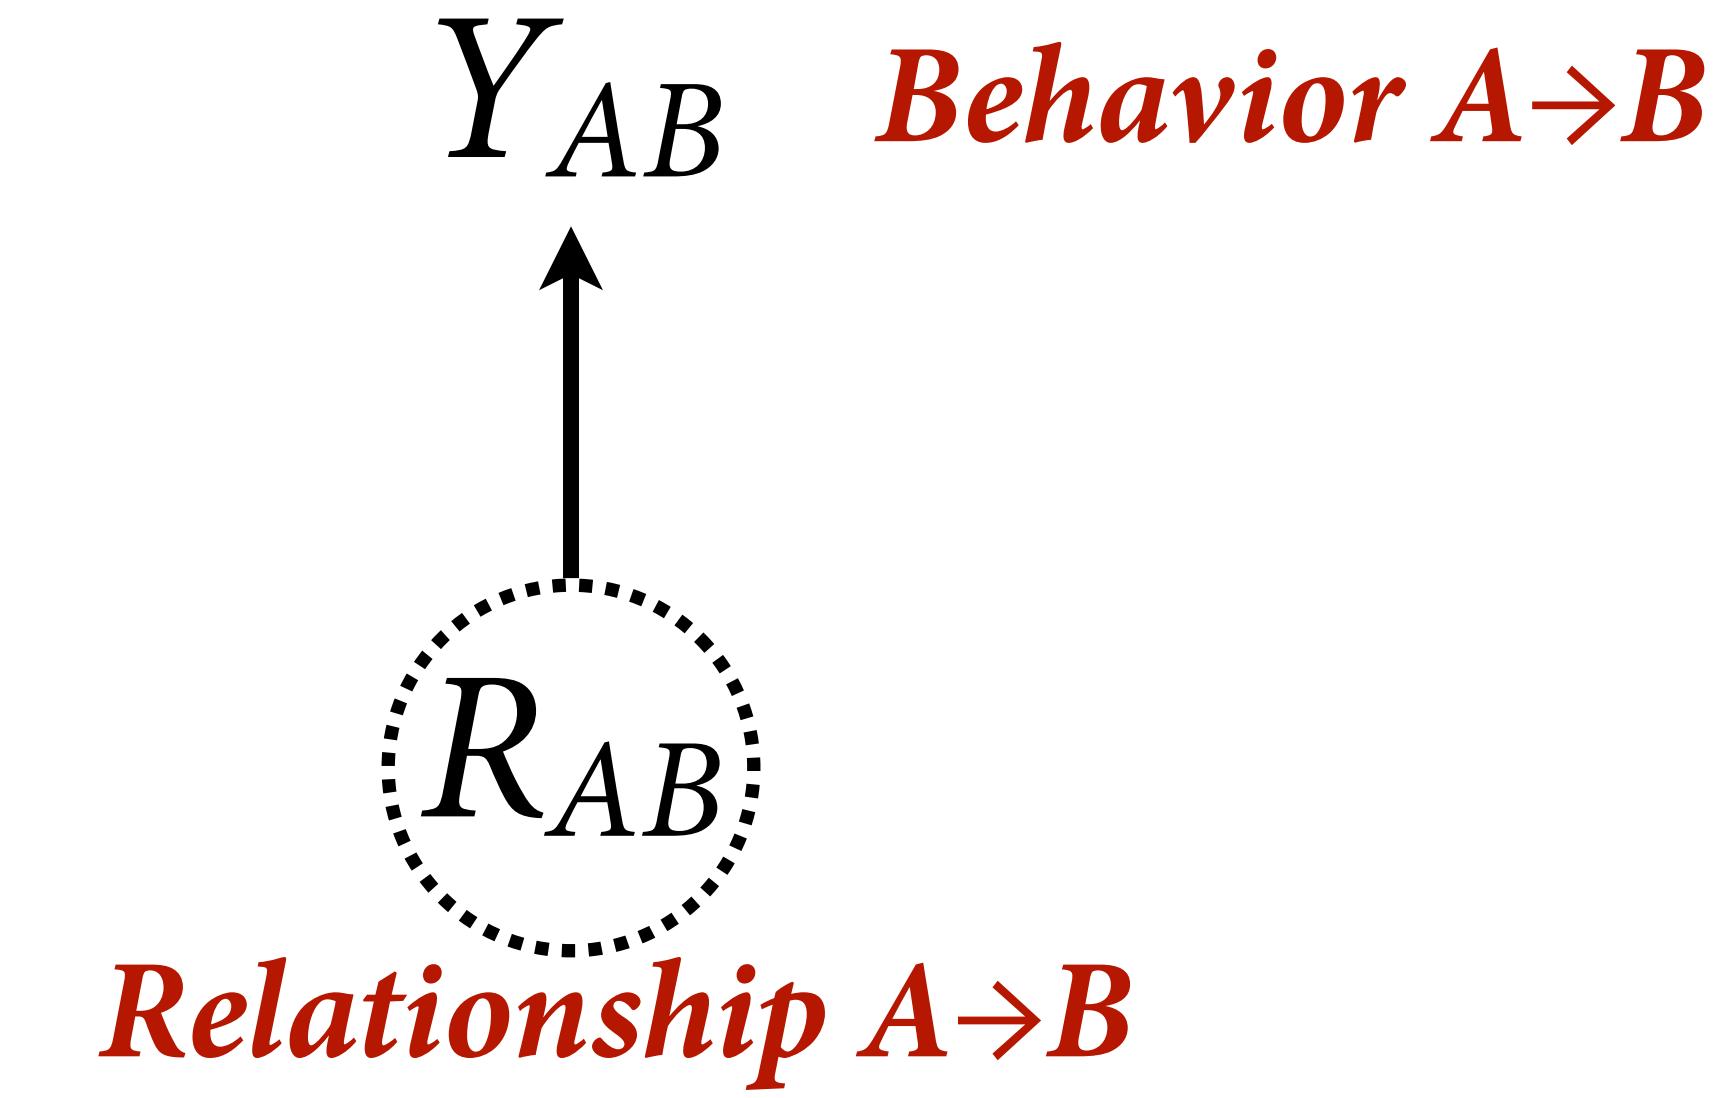
Real, latent modeling problems



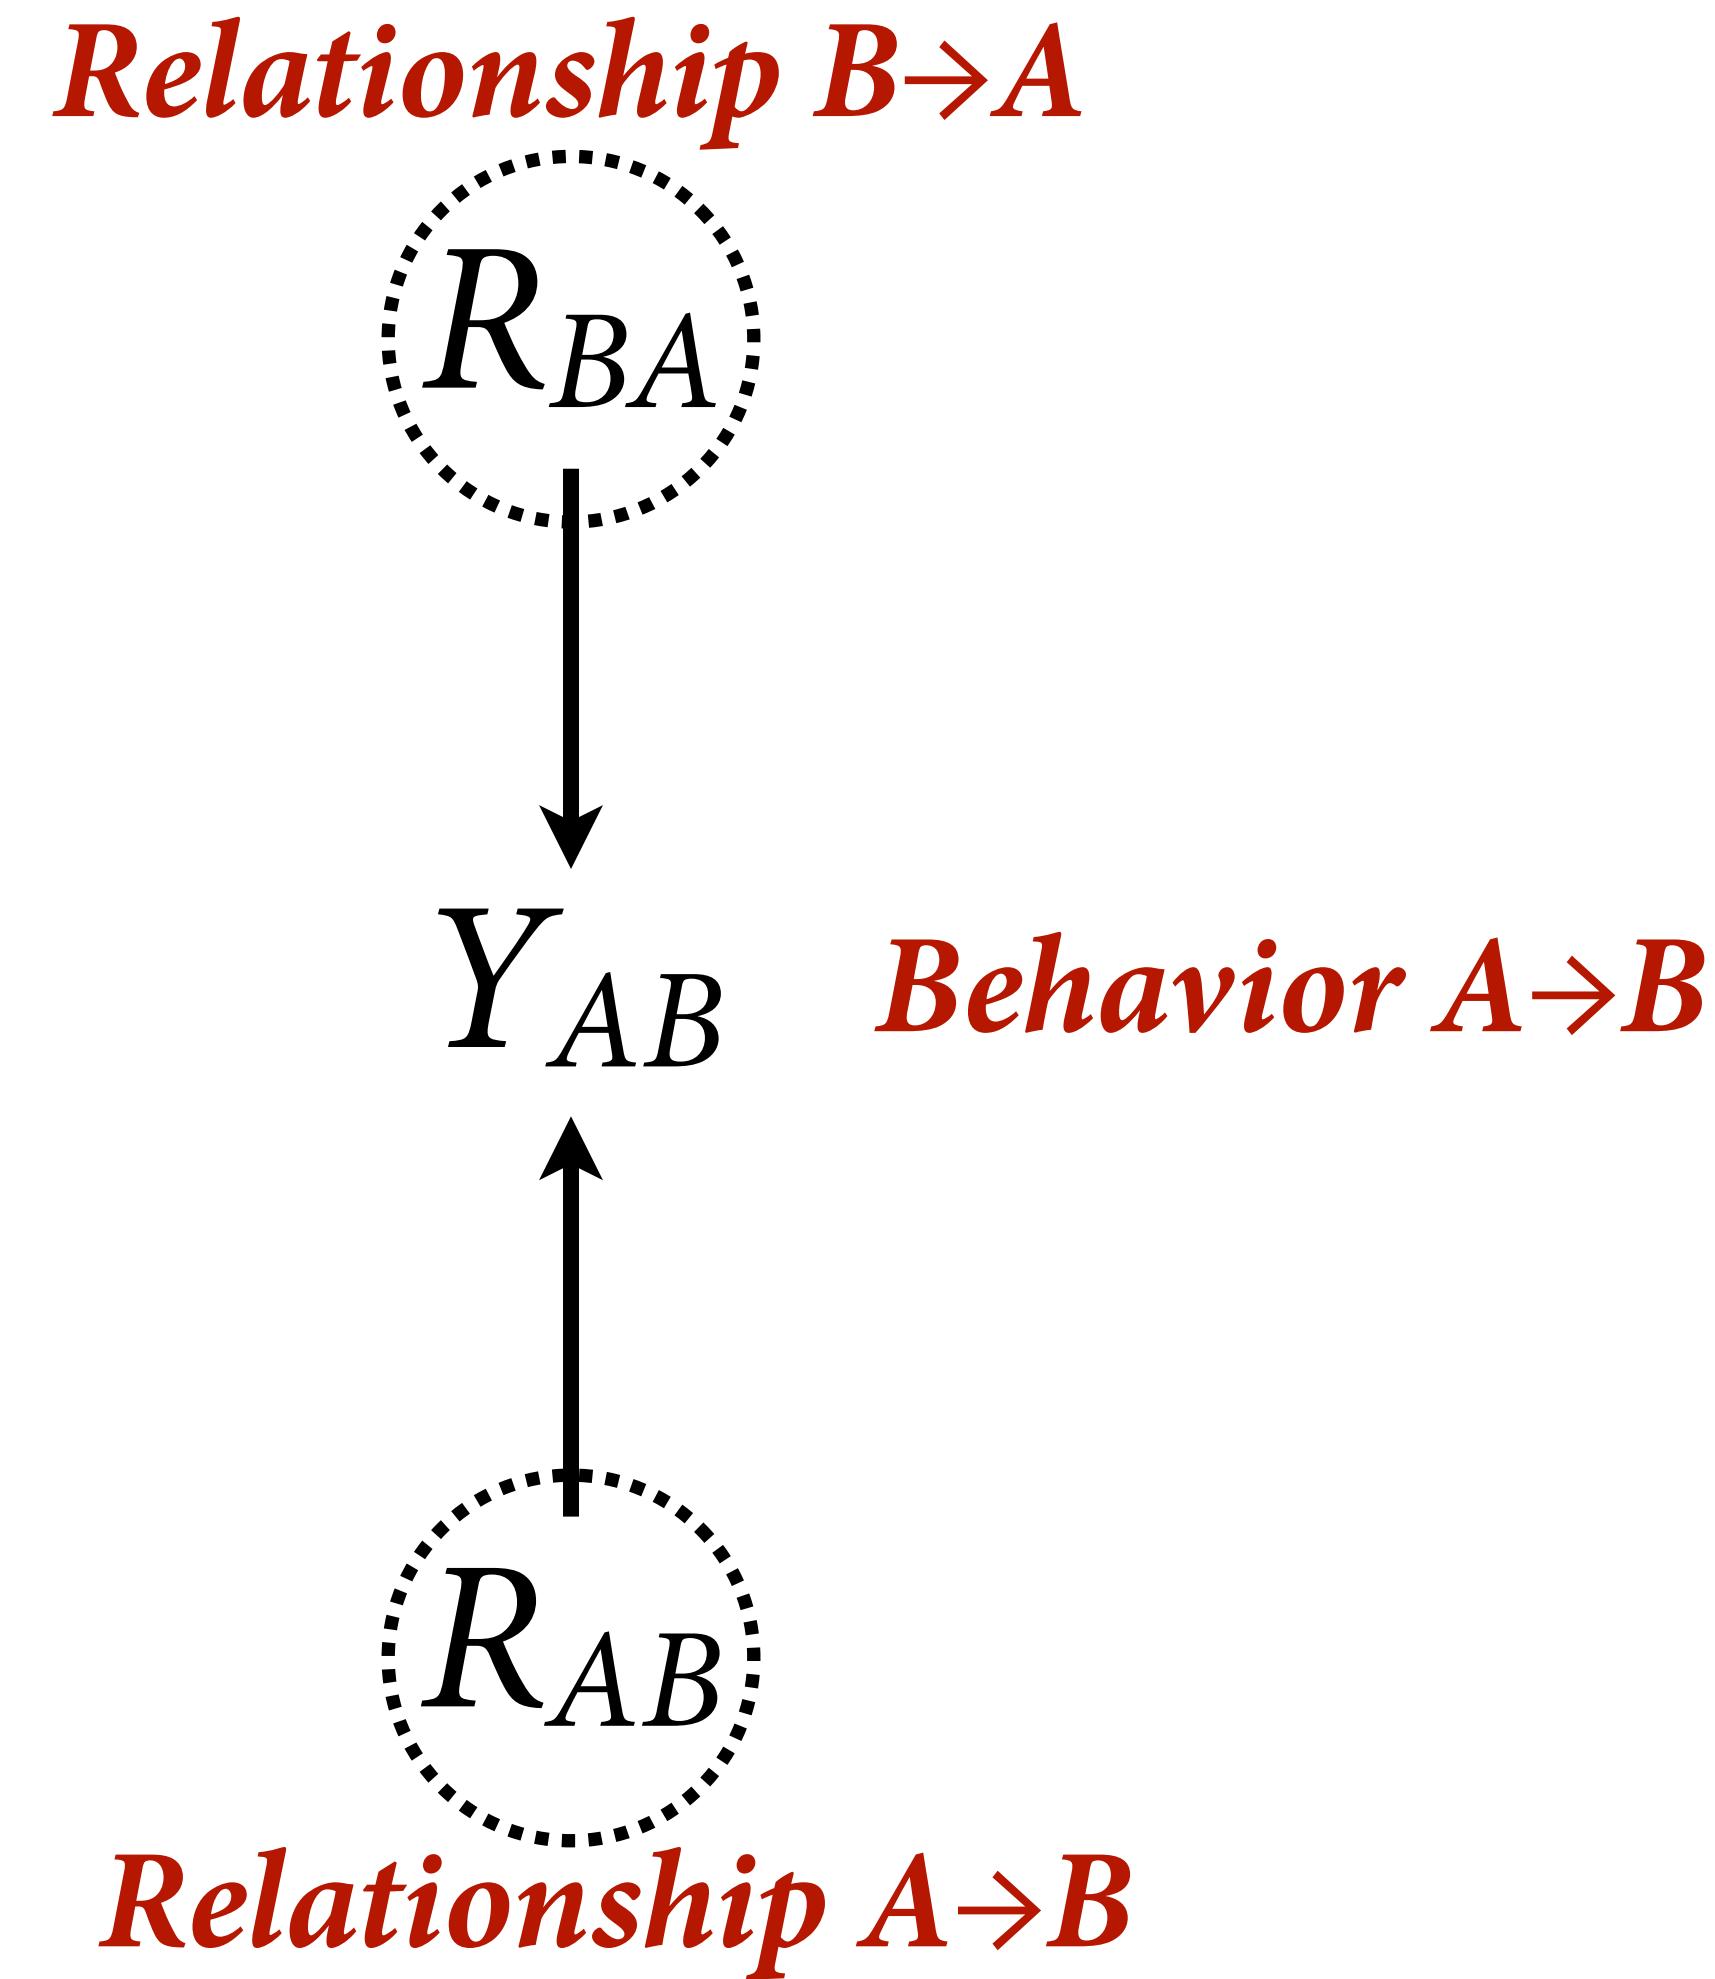
Real, latent modeling problems



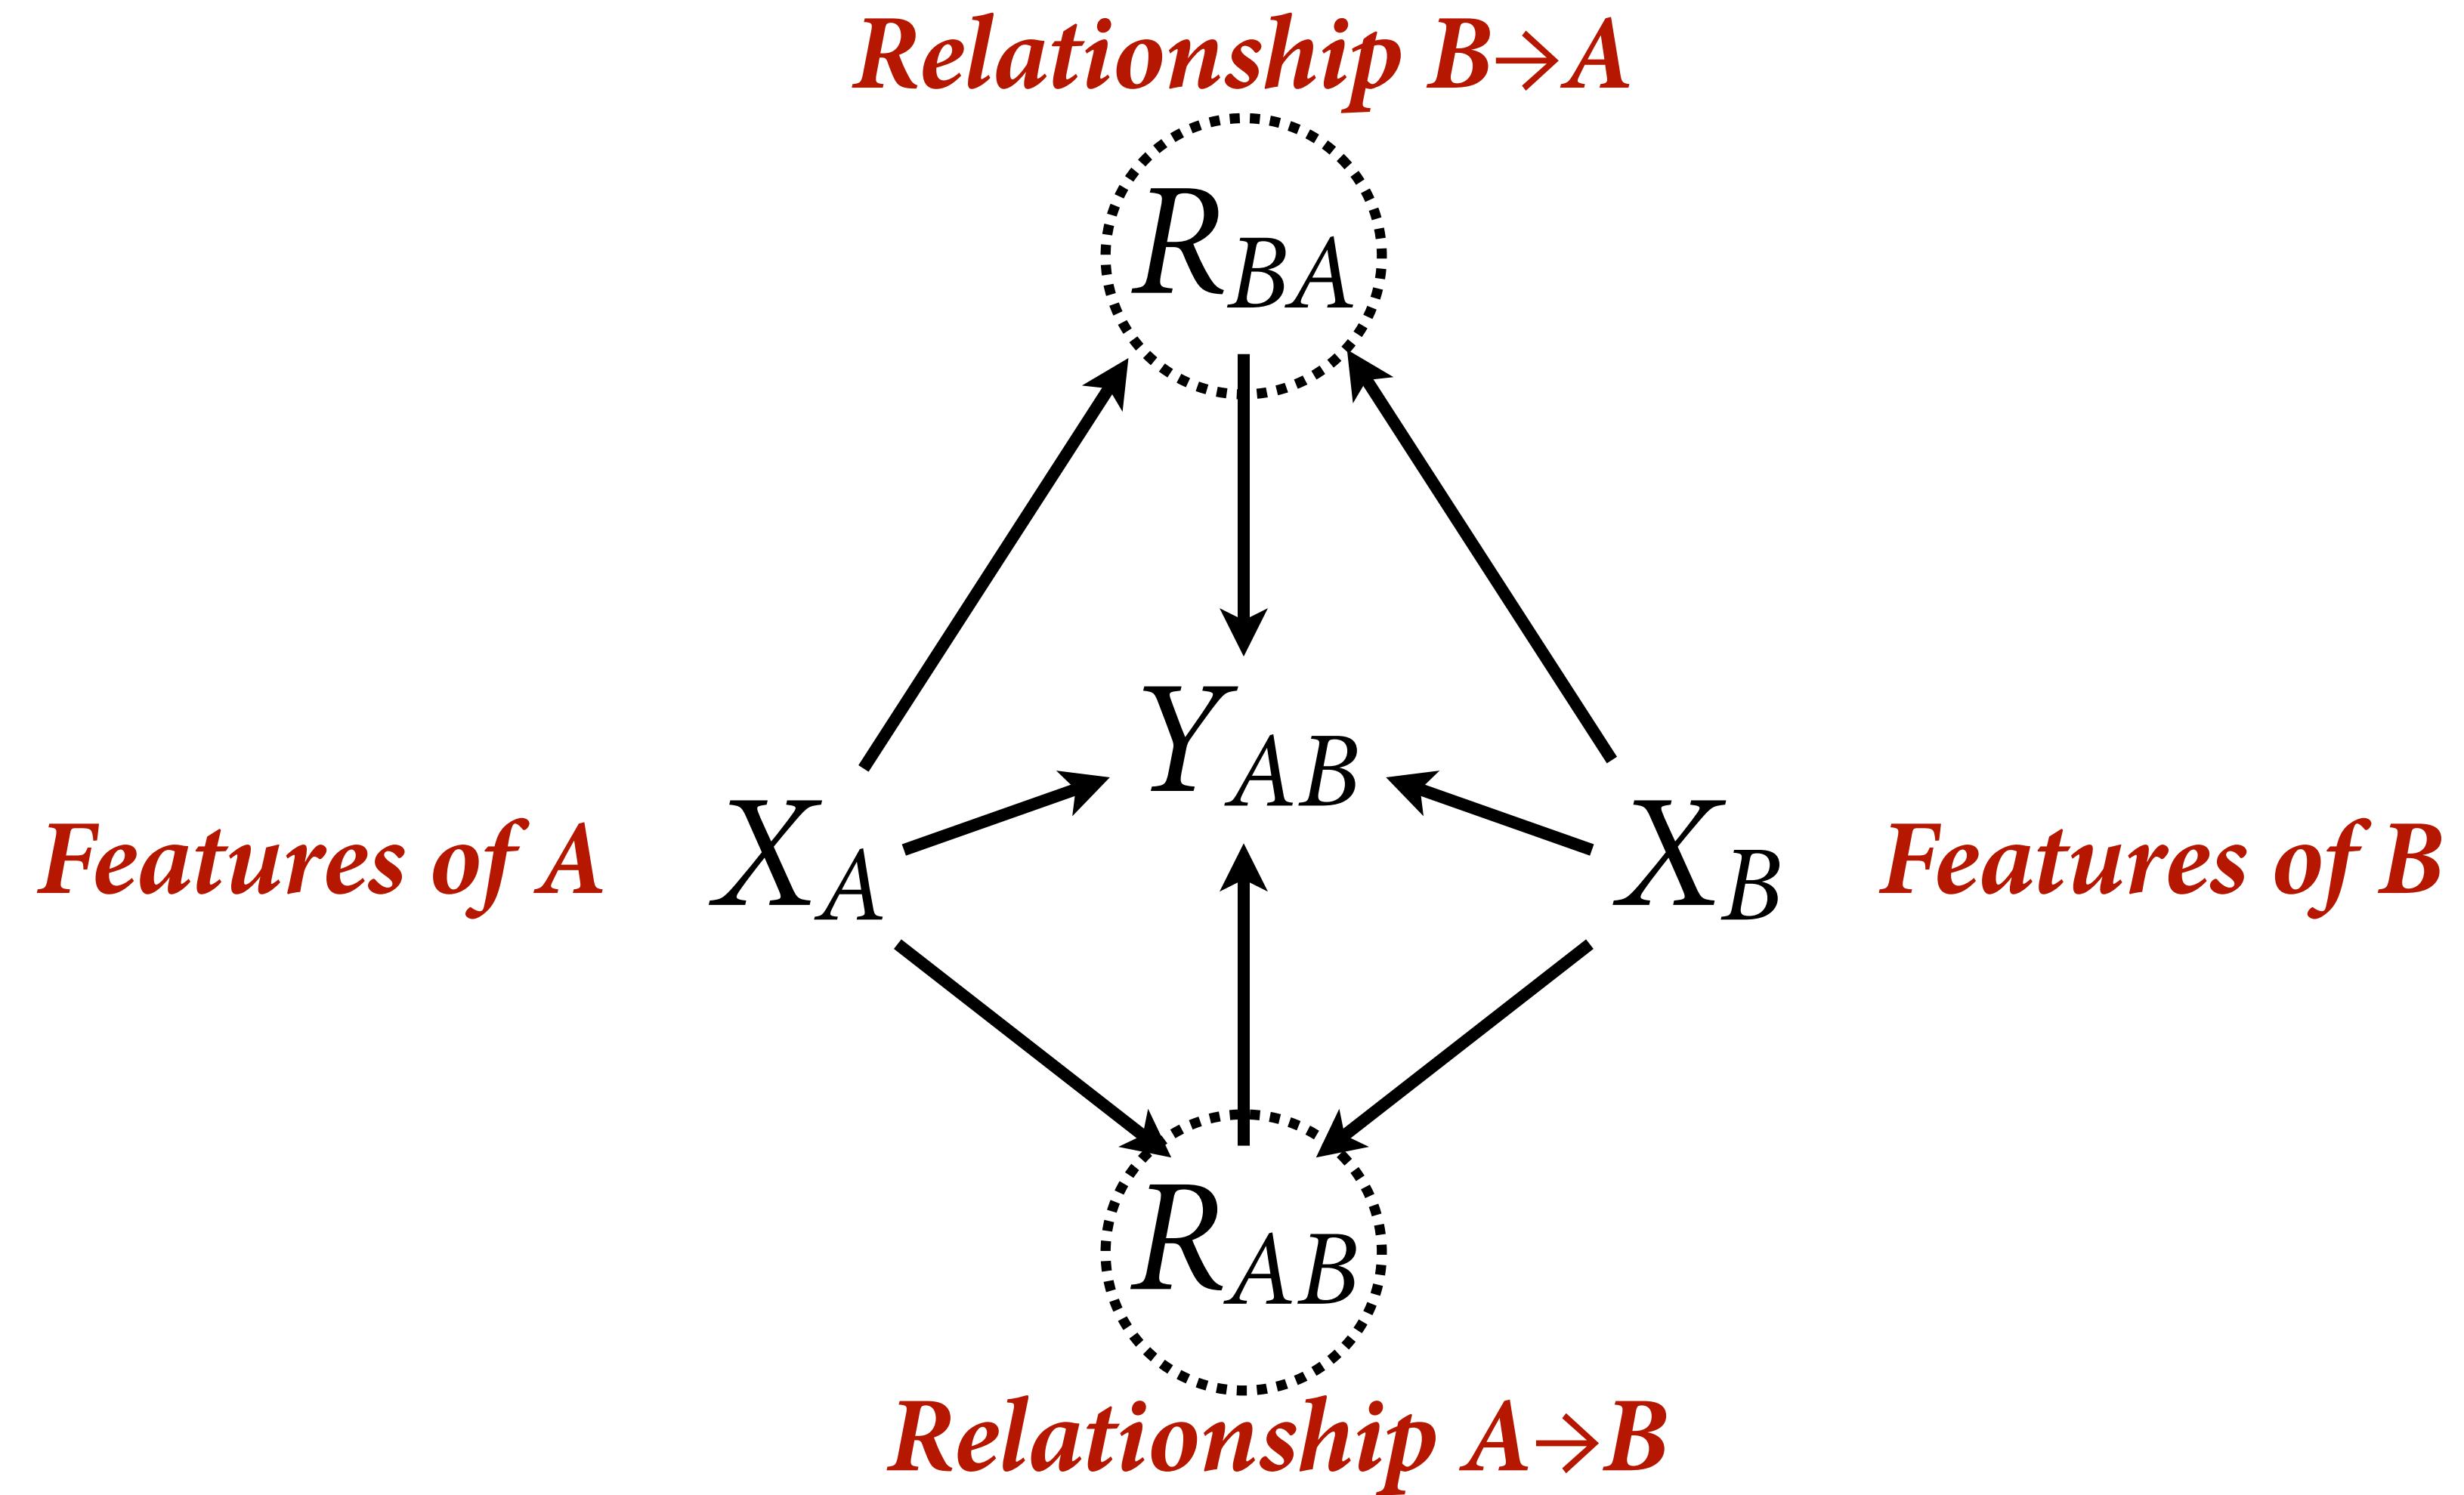
Social networks



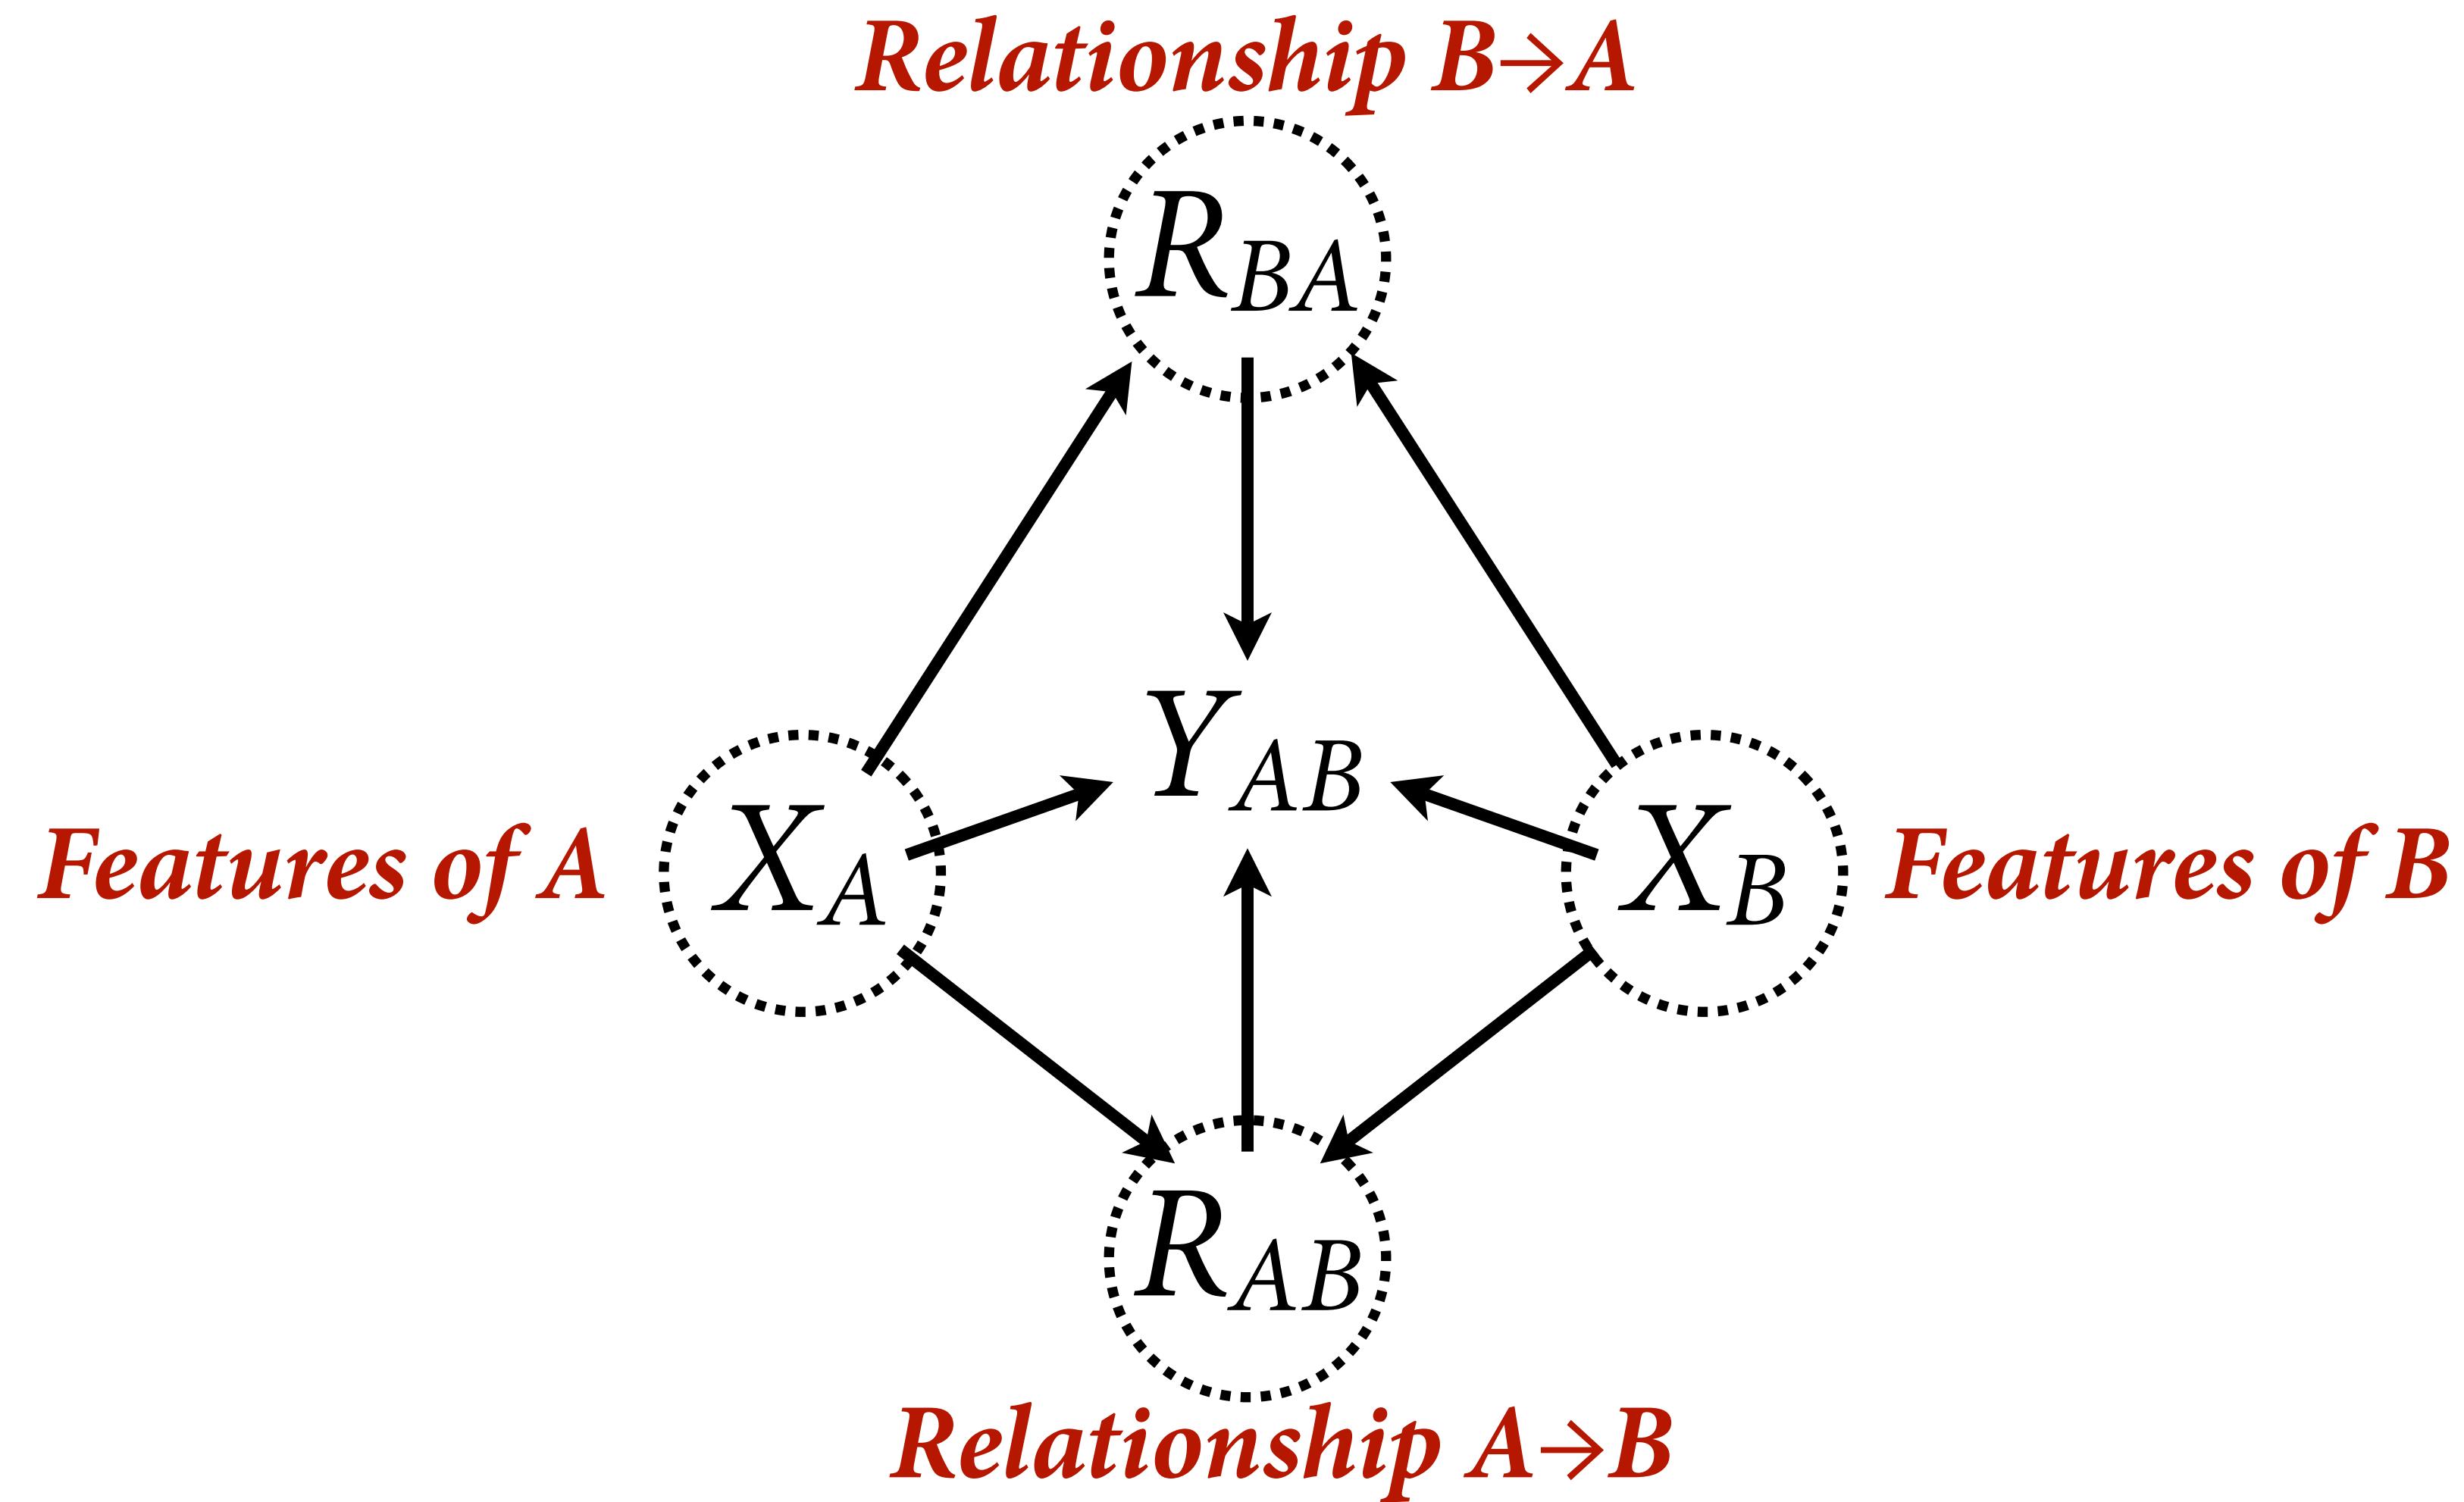
Social networks



Social networks



Social networks



Problems & solutions

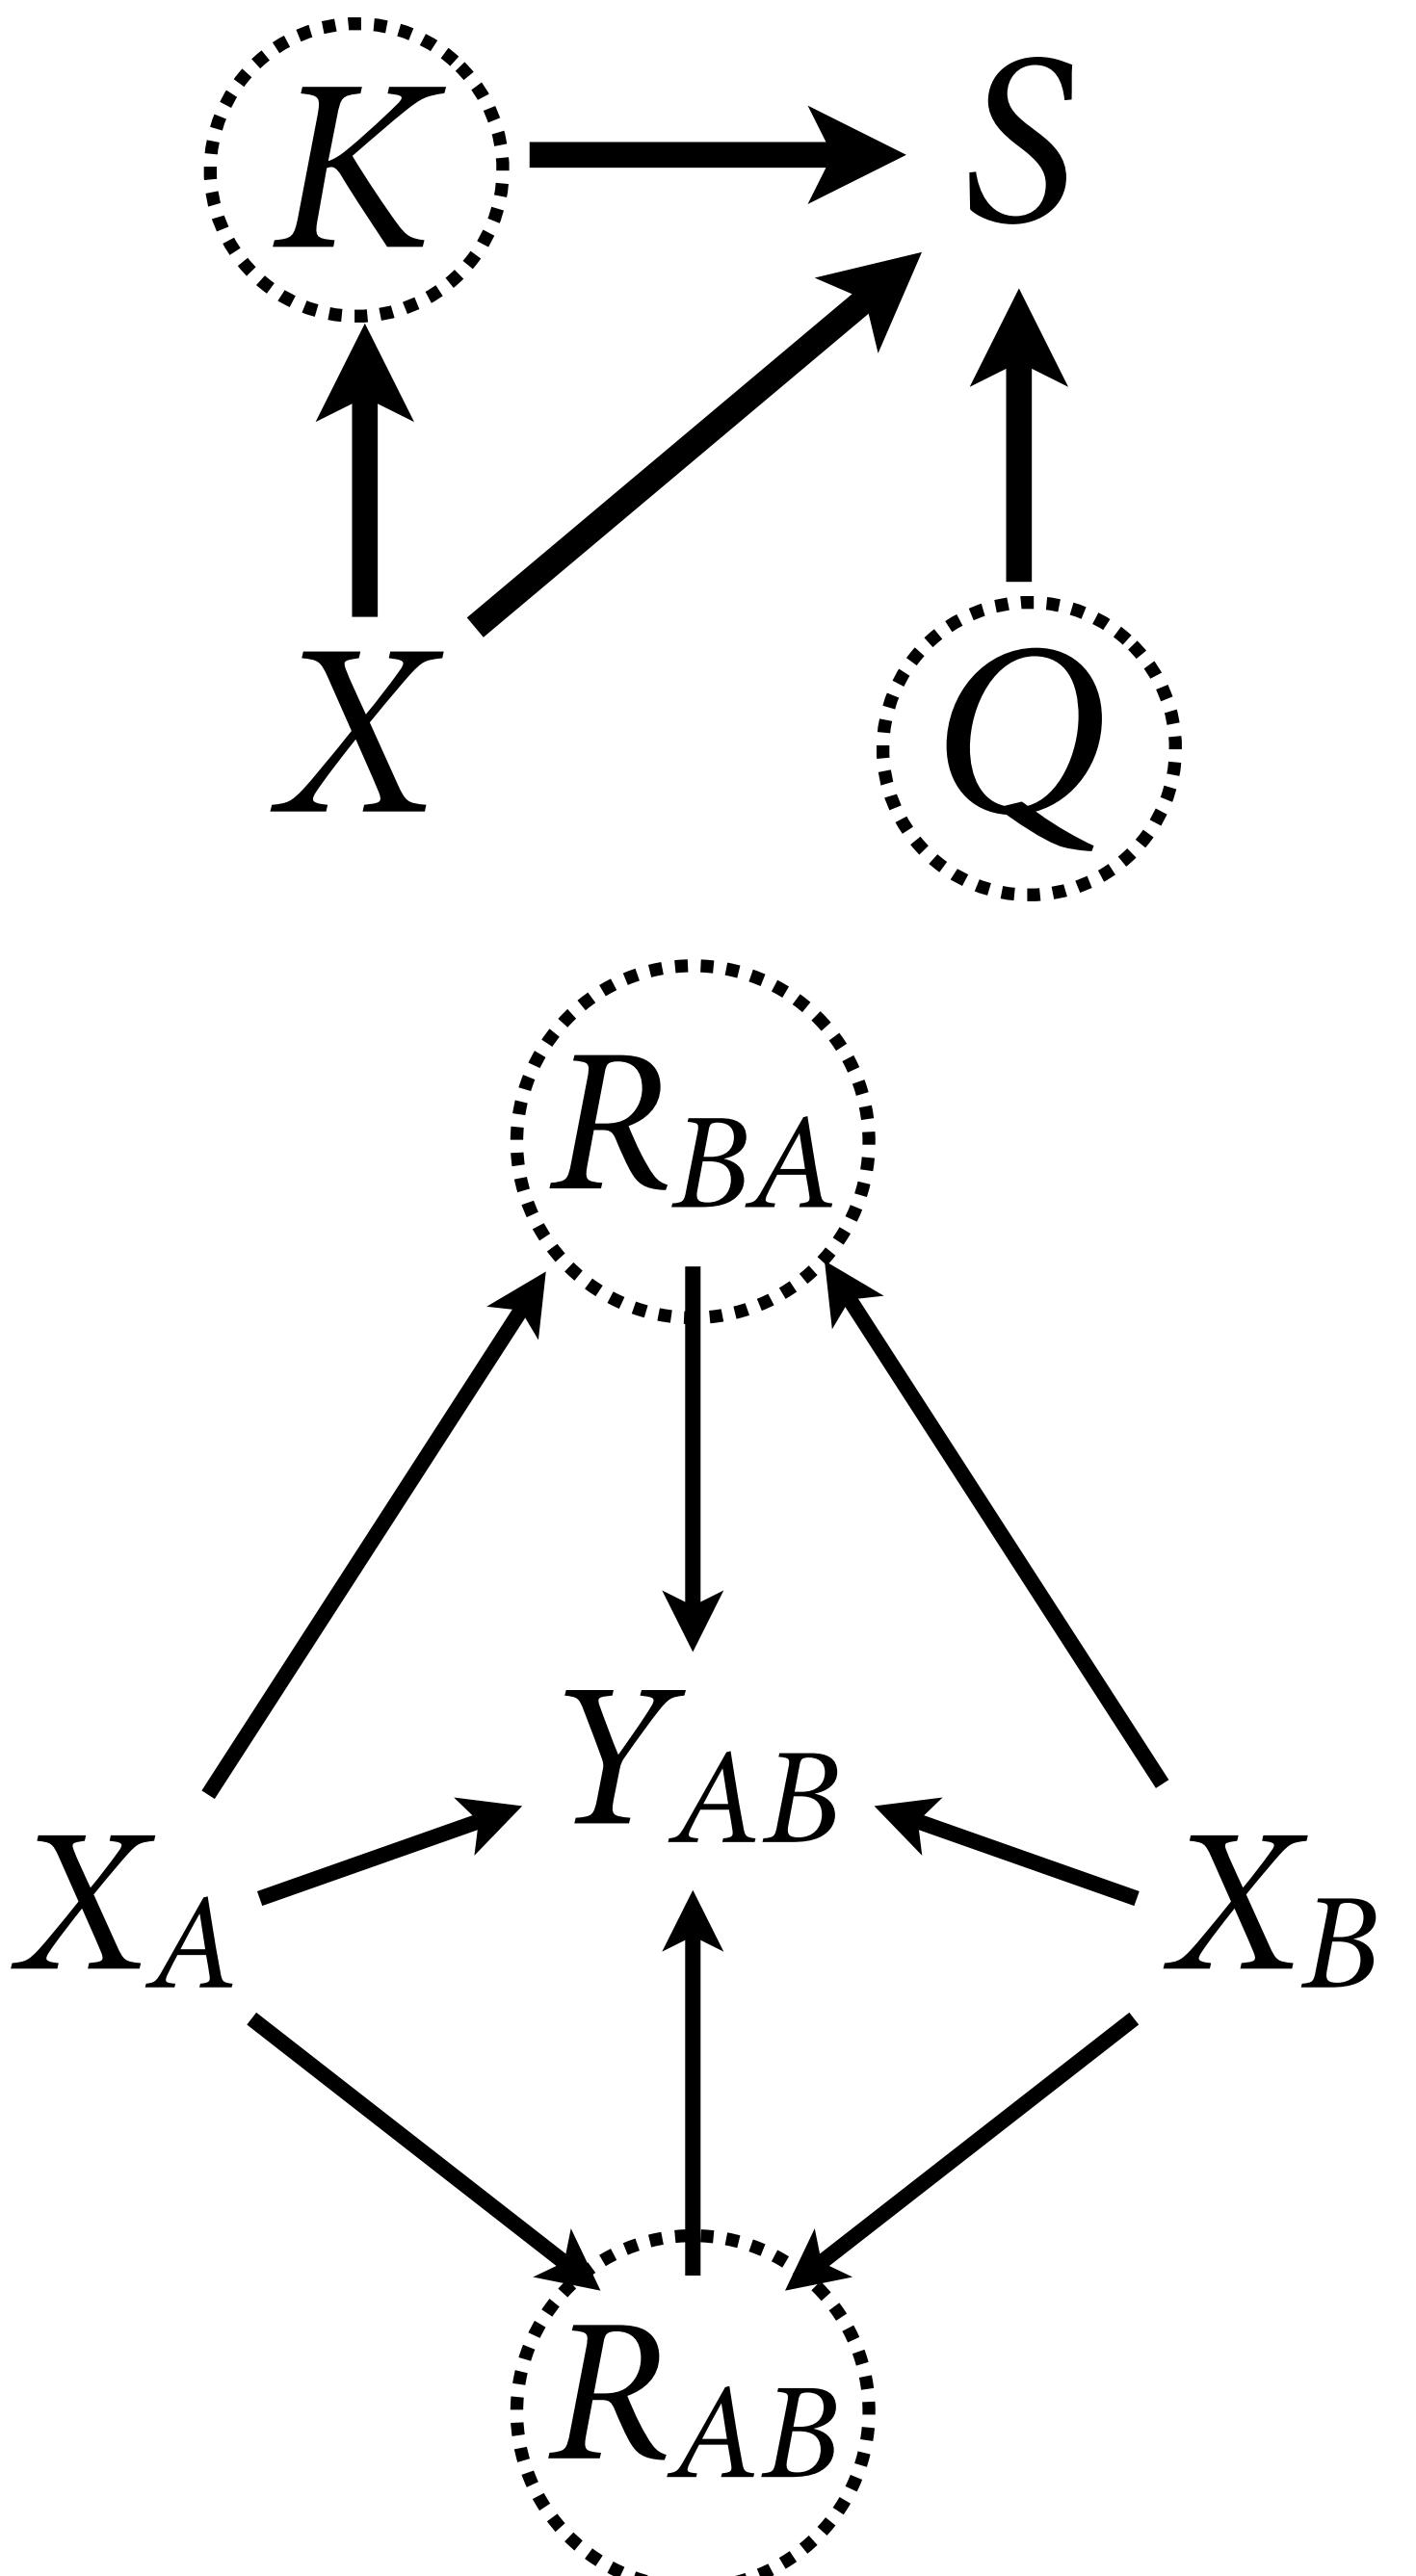
Real research problems:

Many unknowns

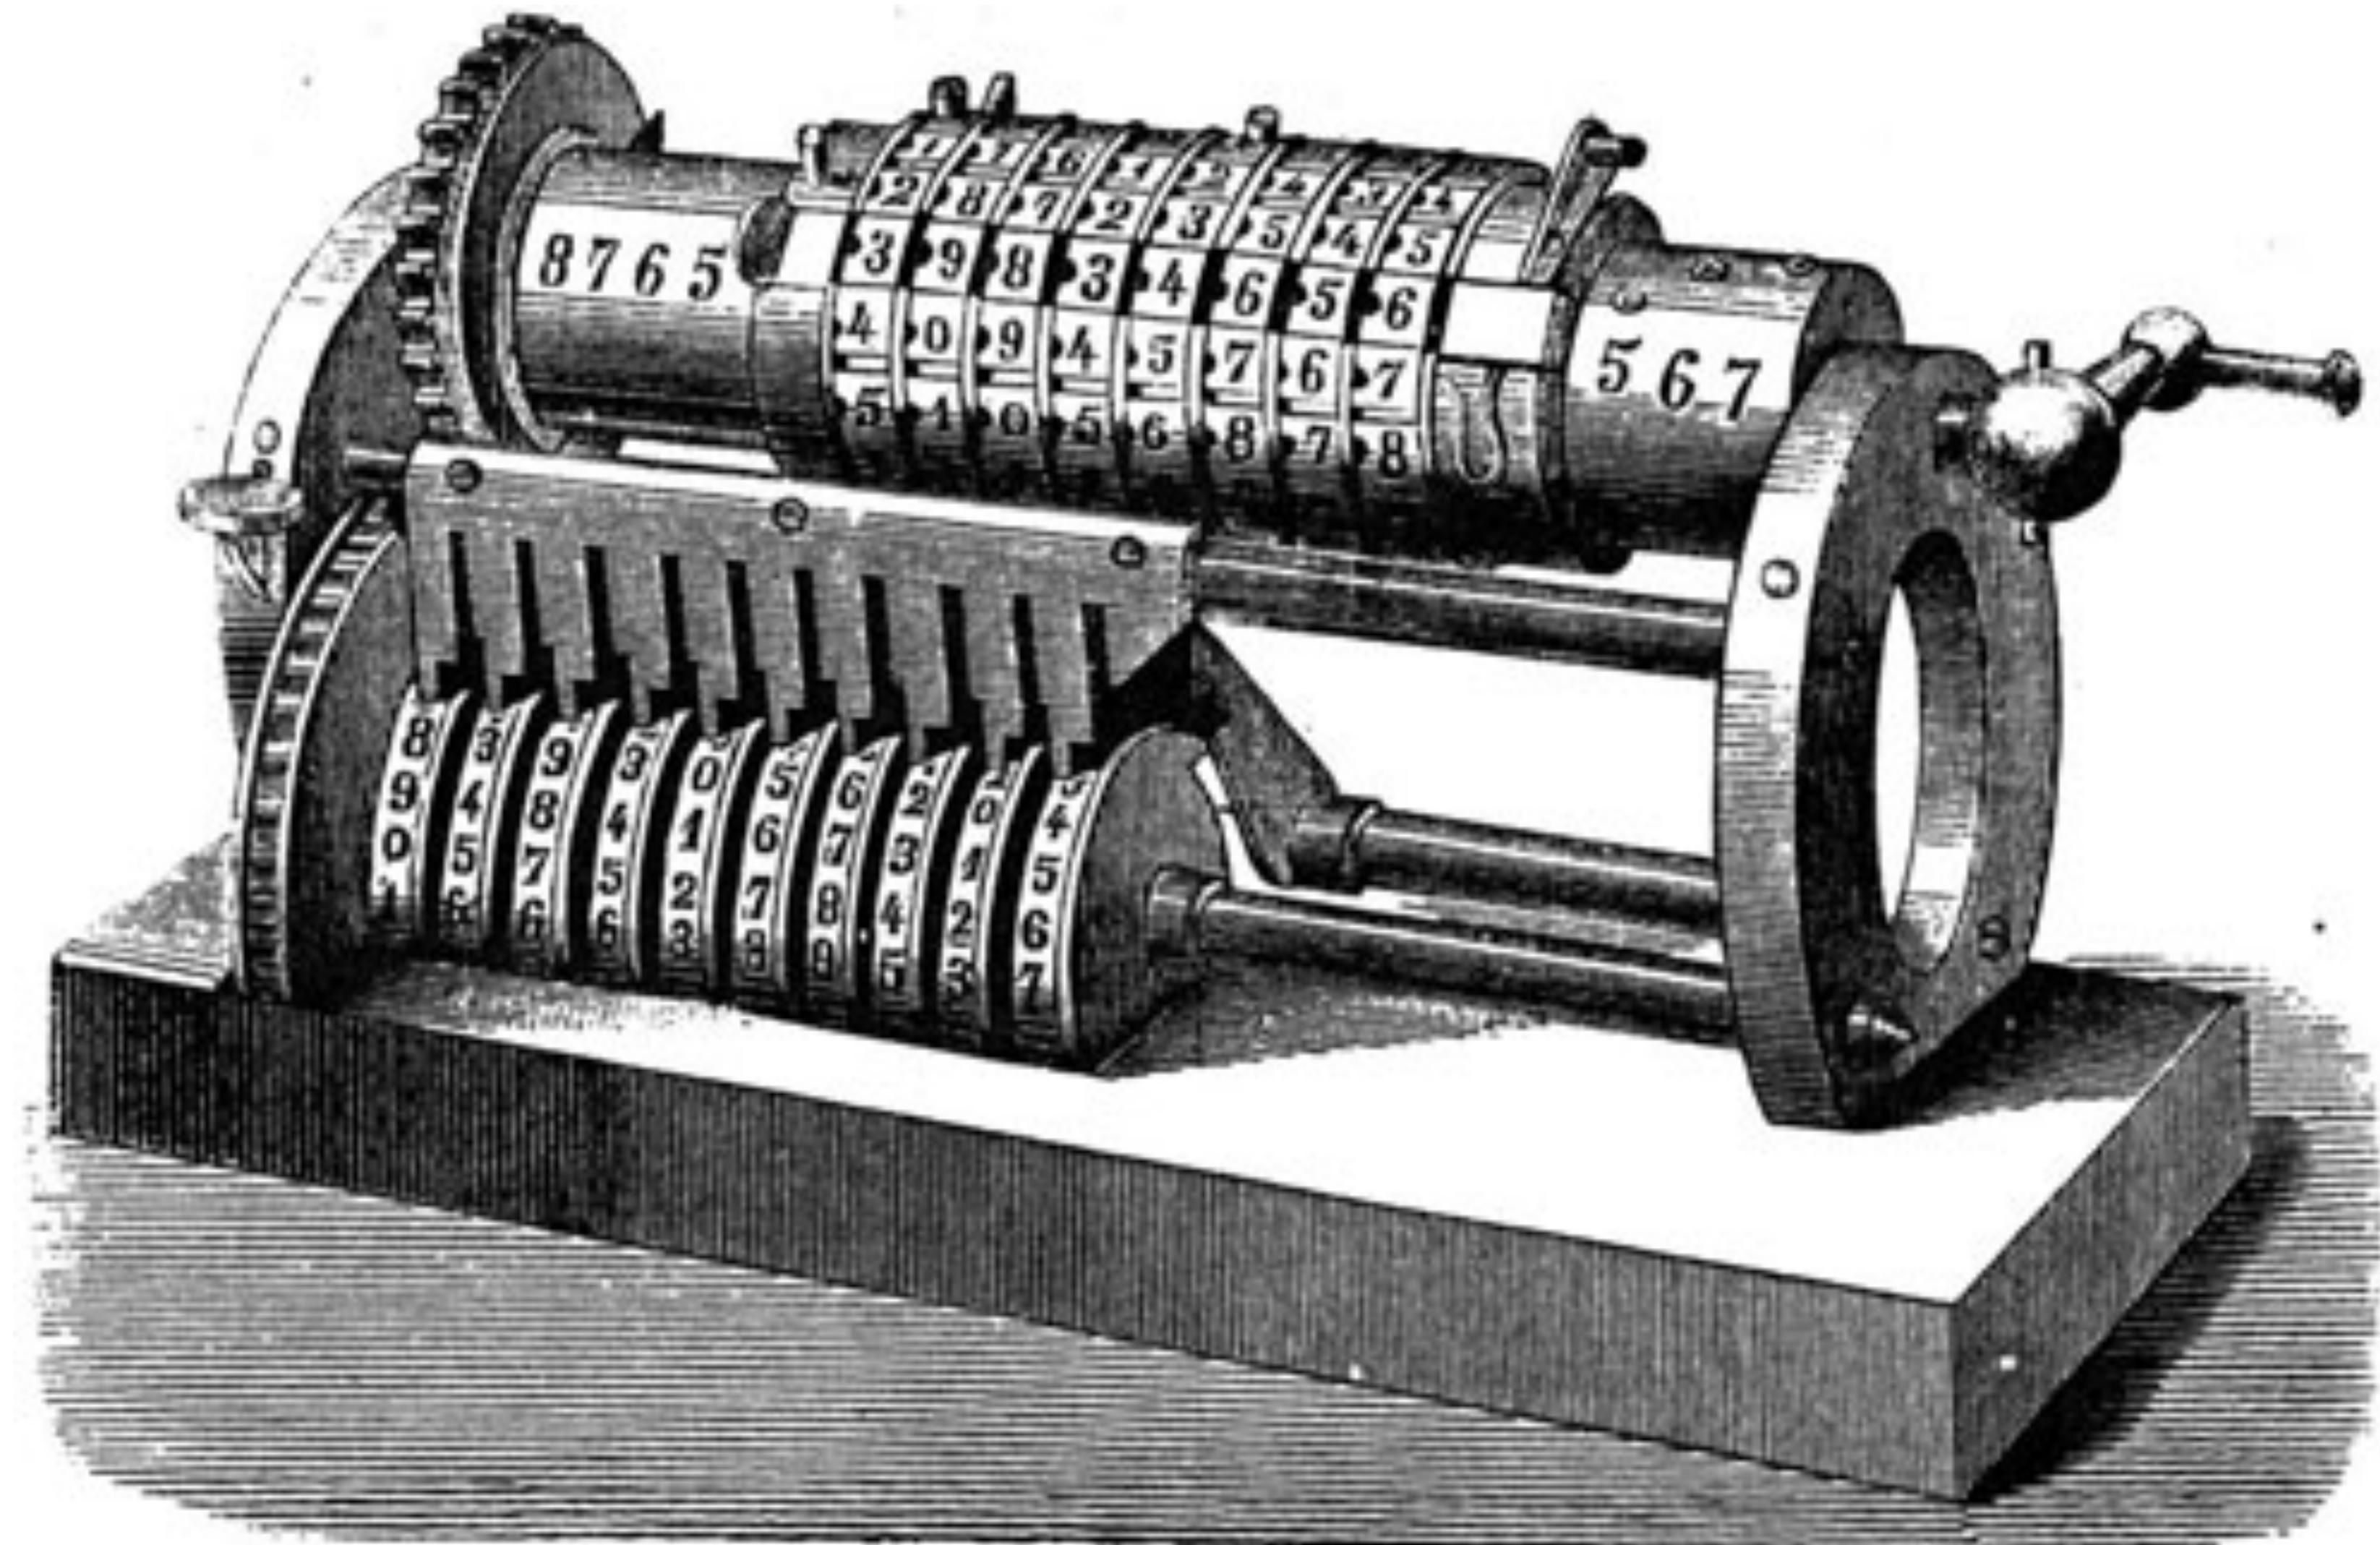
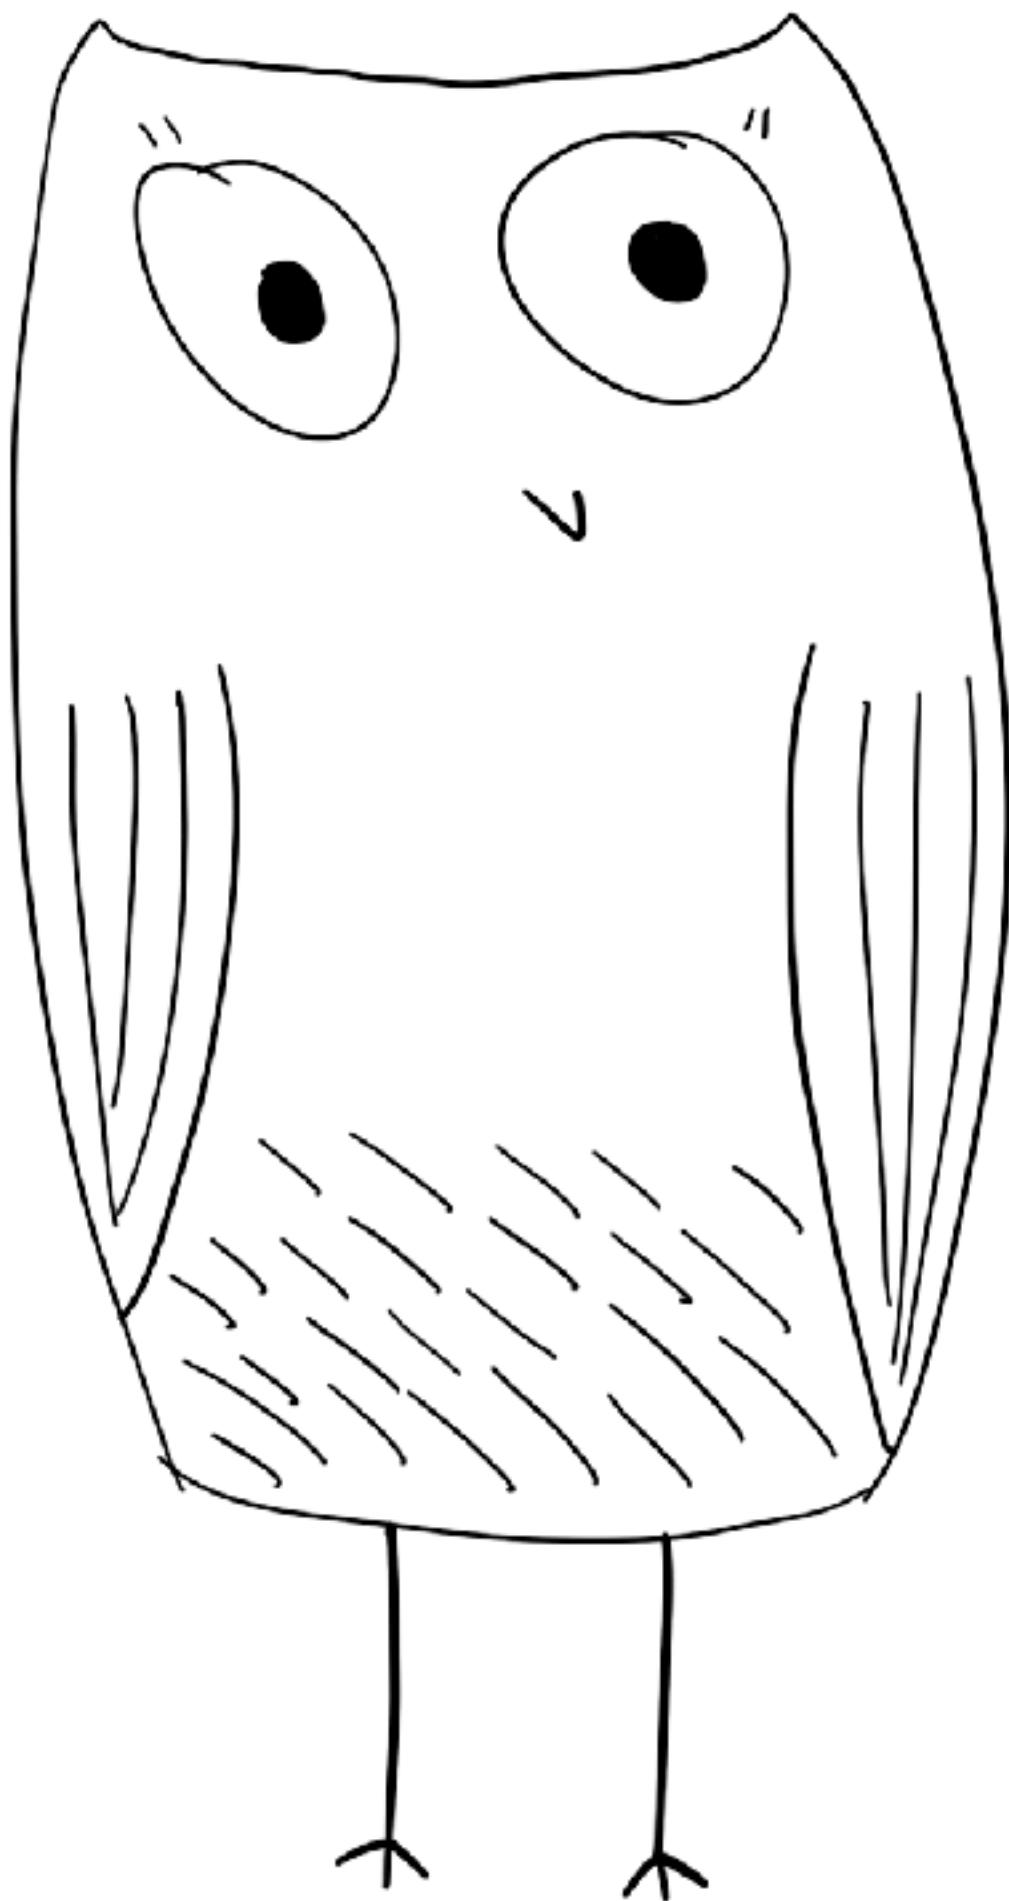
Nested relationships

Bounded outcomes

Difficult calculations

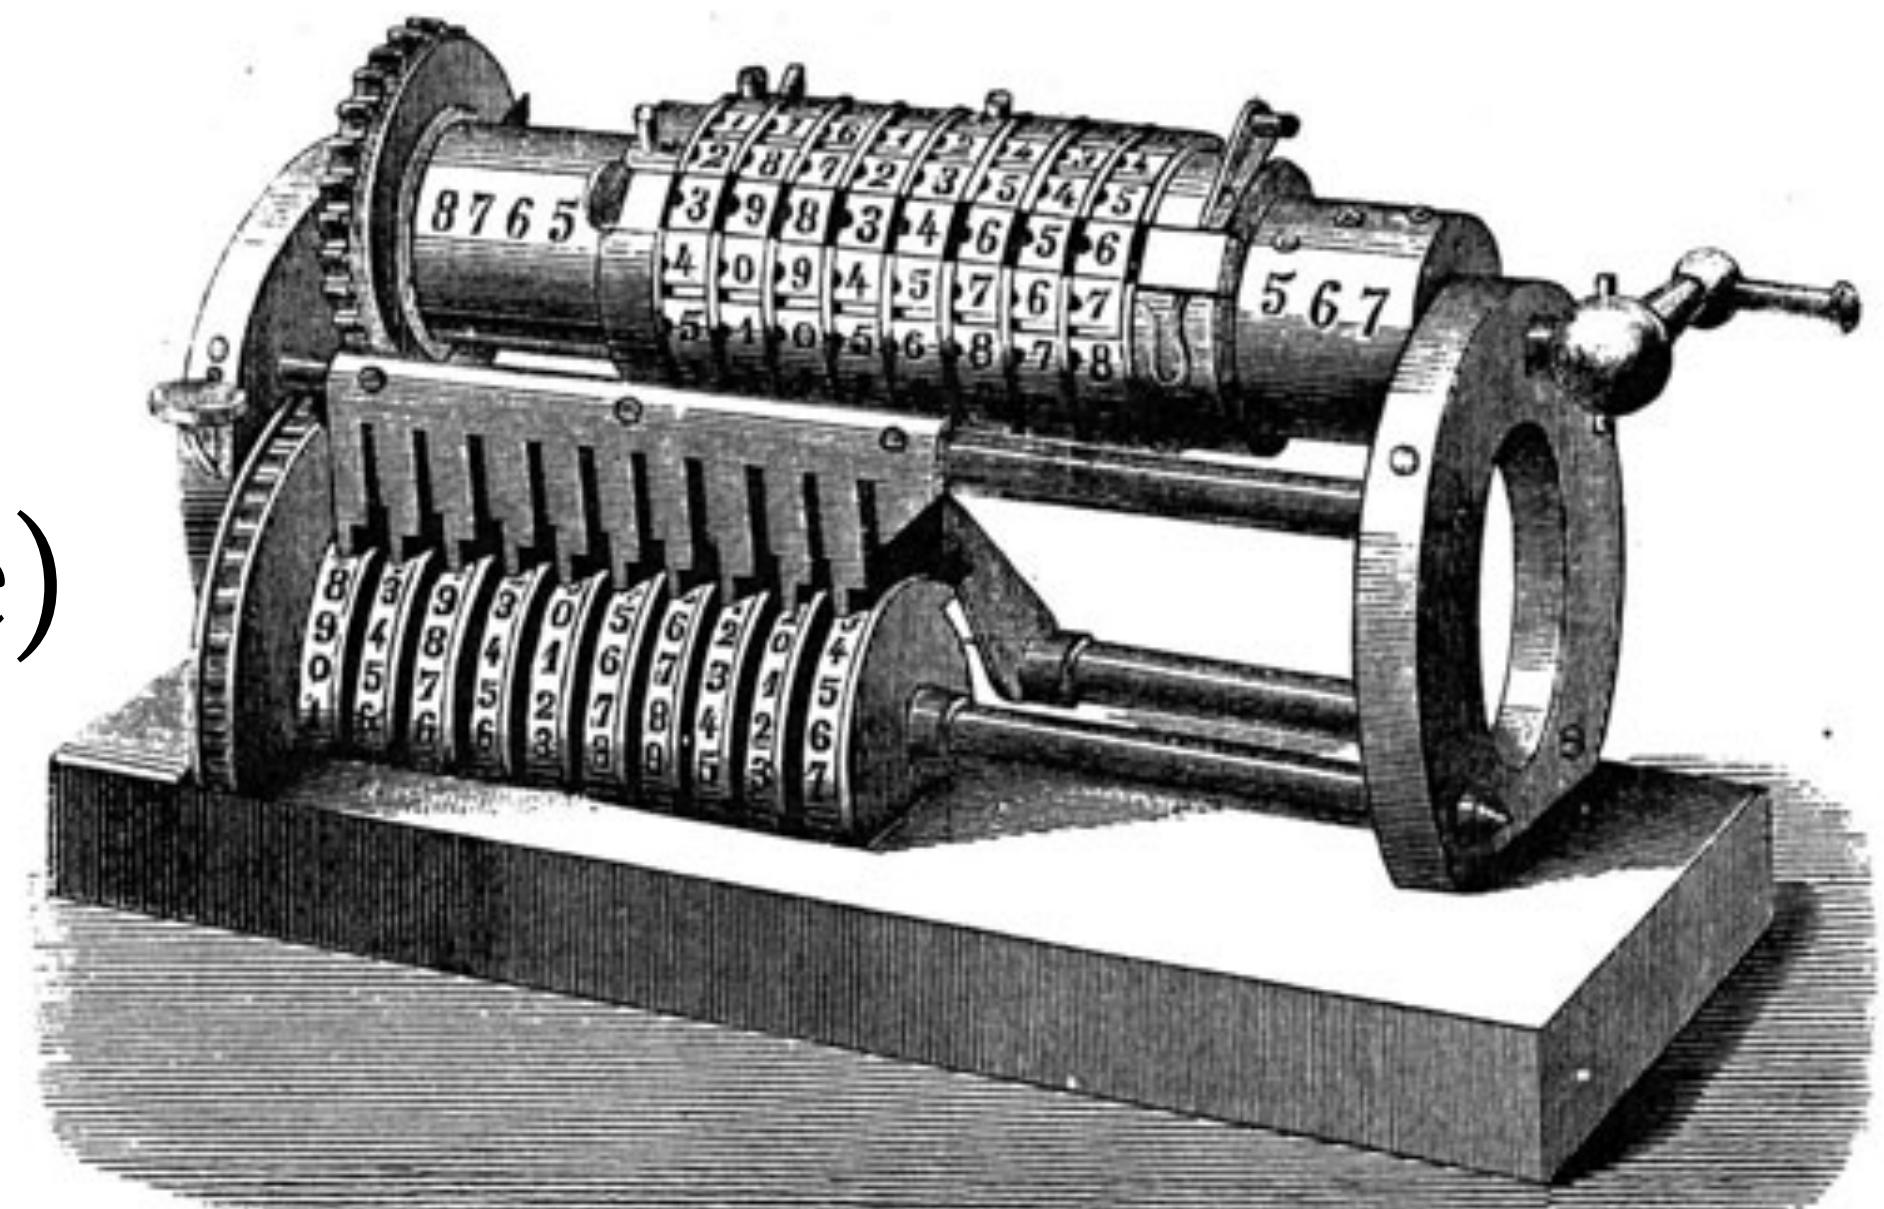


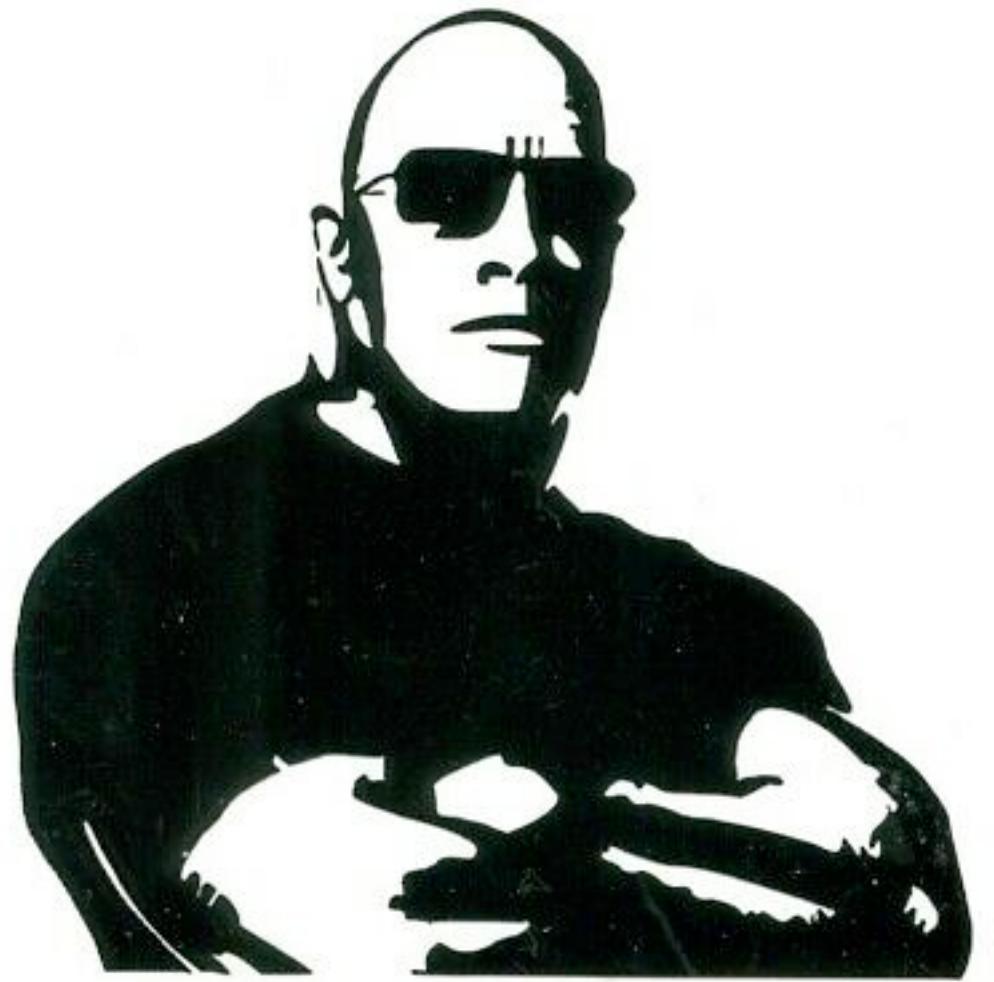
aNAlYZE tHe DatA



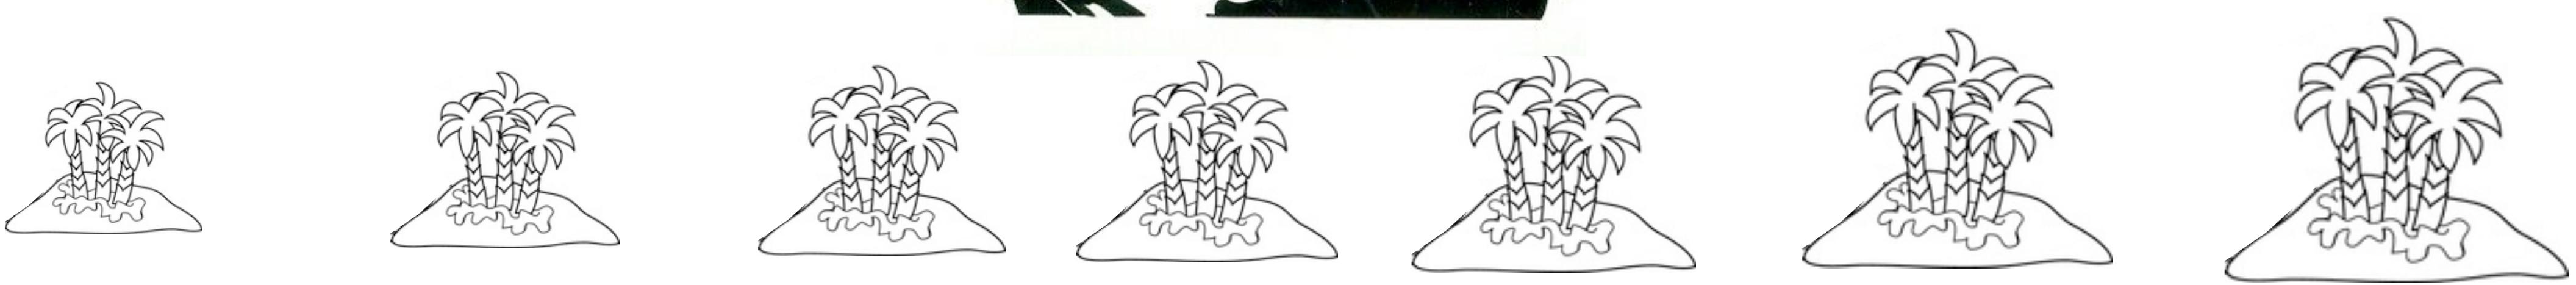
Computing the posterior

1. Analytical approach (often impossible)
2. Grid approximation (very intensive)
3. Quadratic approximation (limited)
4. Markov chain Monte Carlo (intensive)



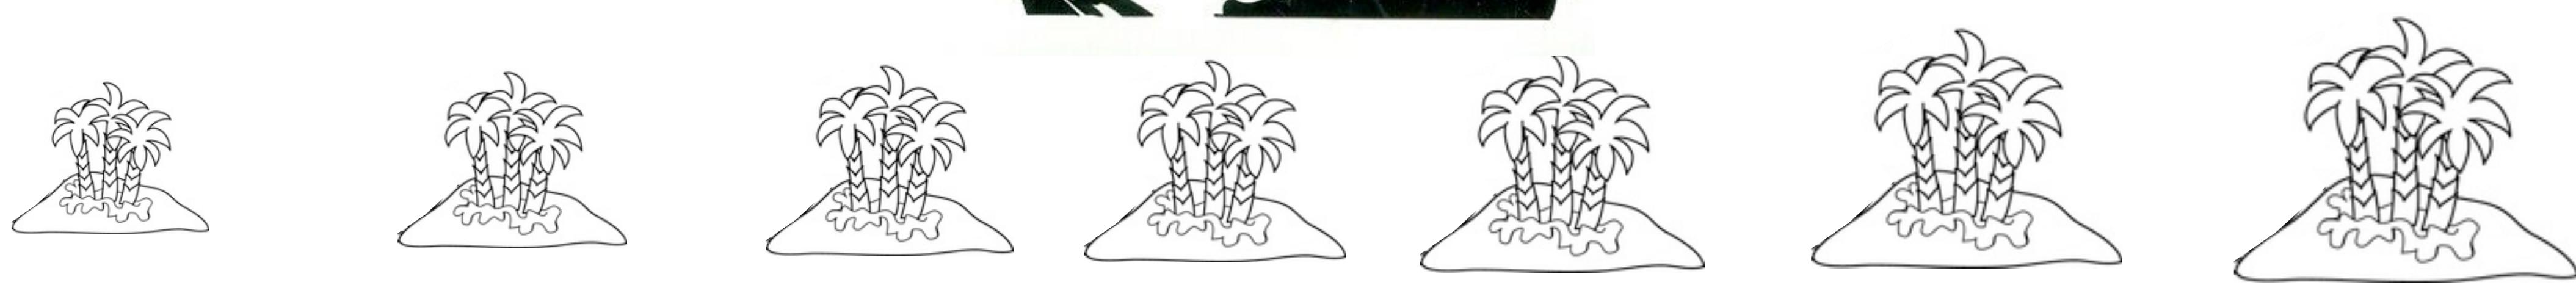


King Markov

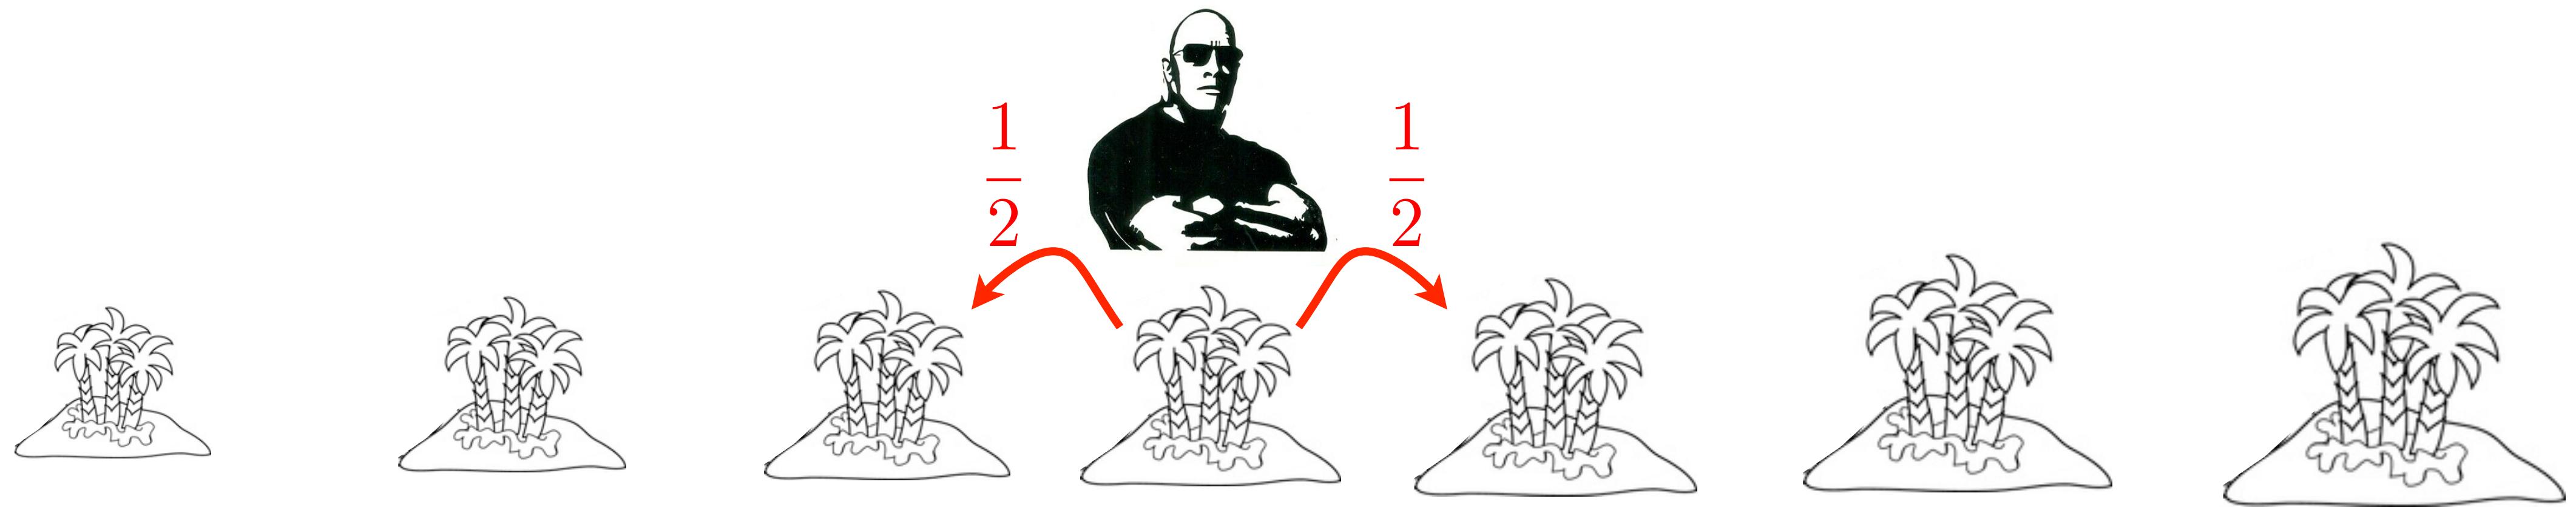


The Metropolis Archipelago

Contract: King Markov must visit each island in proportion to its population size

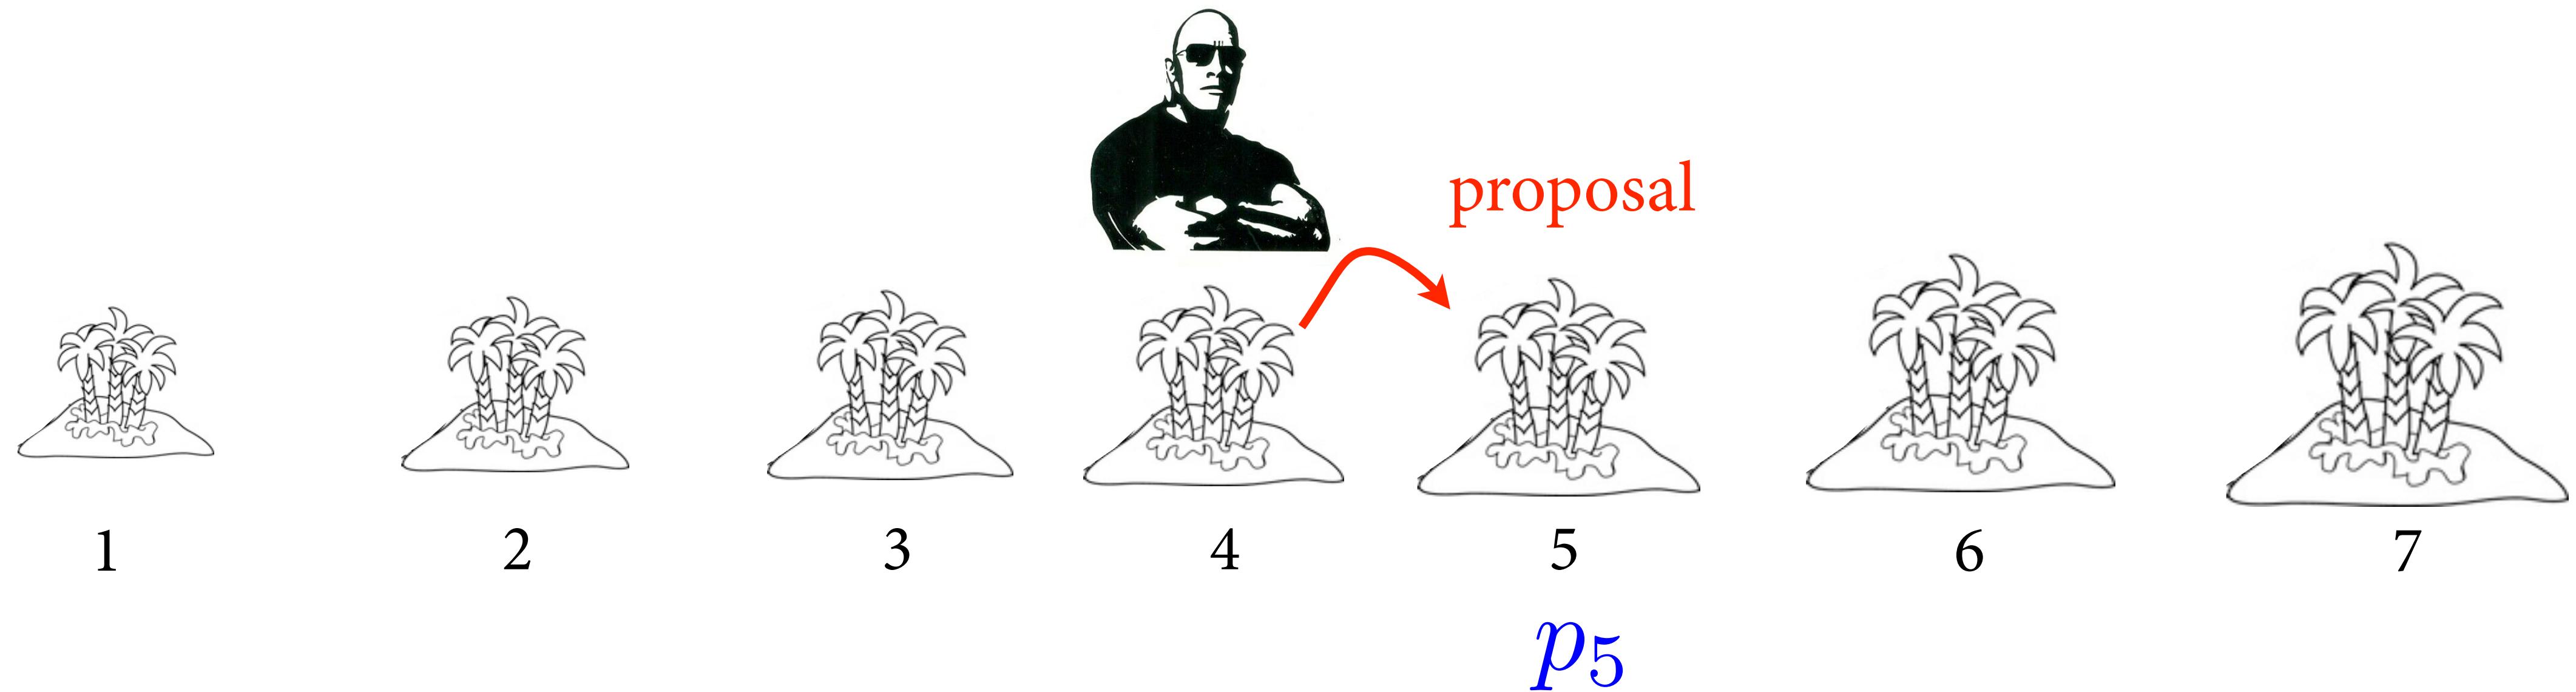


Here's how he does it...



(1) Flip a coin to choose island on left or right.
Call it the “proposal” island.

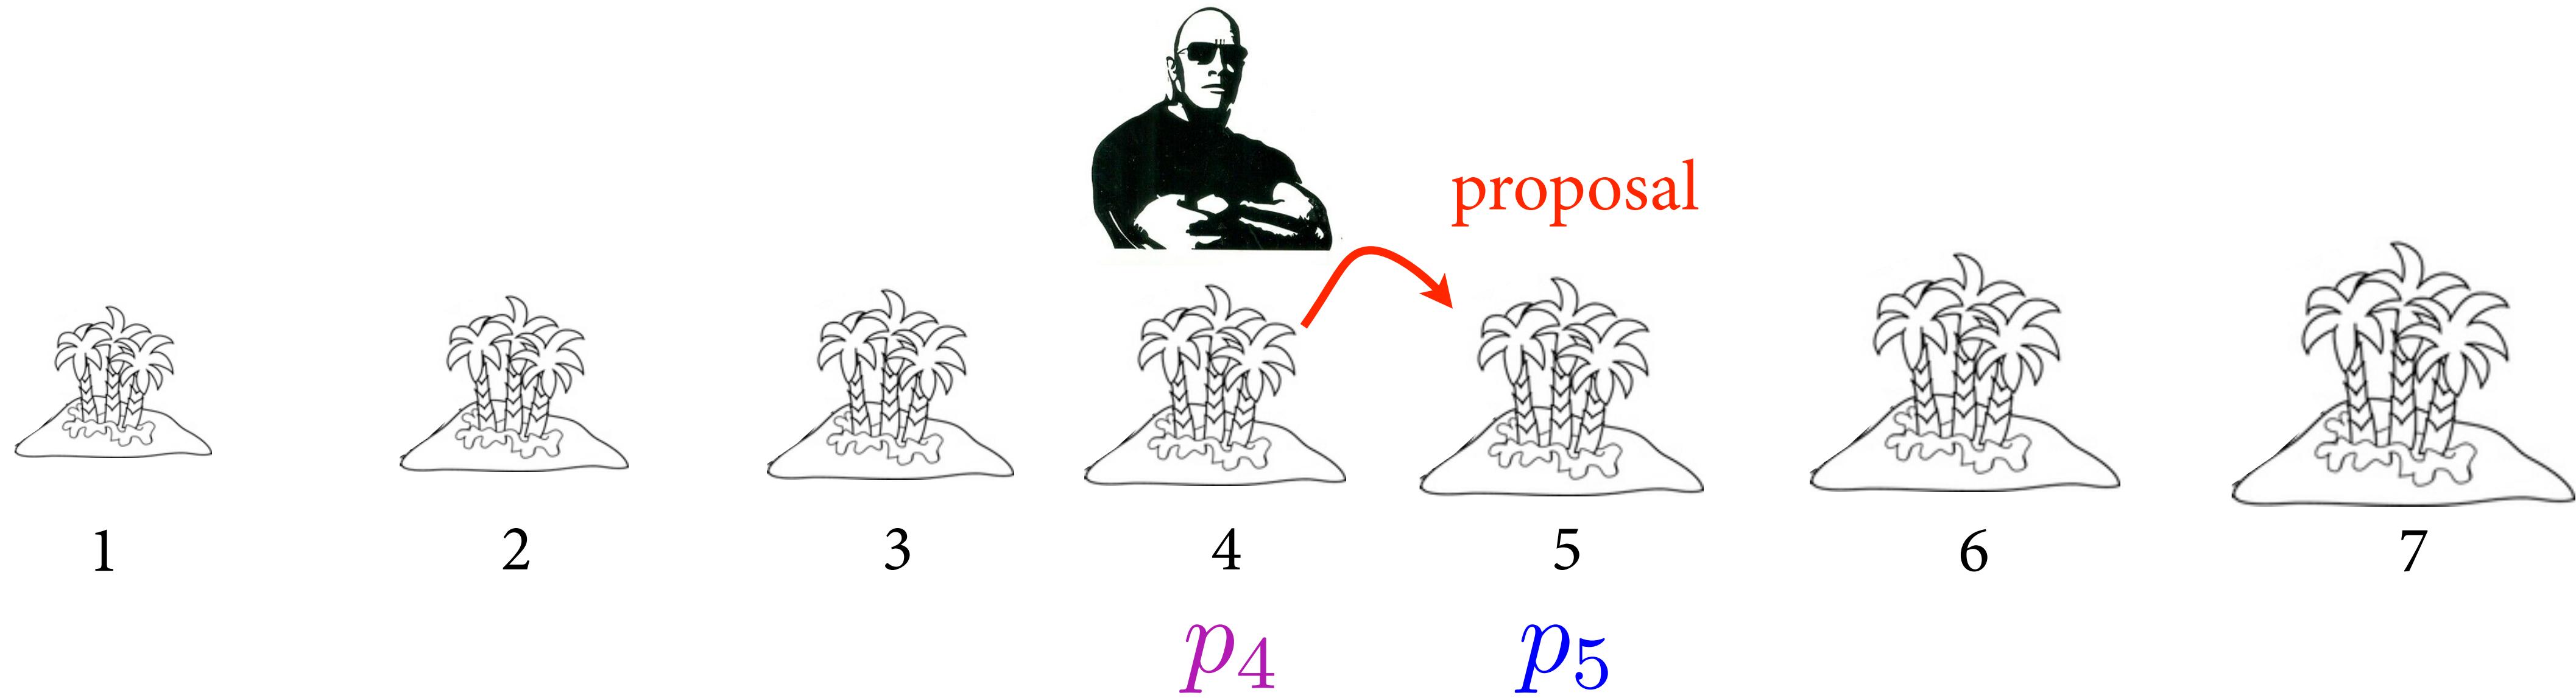
(1) Flip a coin to choose island on left or right.
Call it the “proposal” island.



(2) Find population of proposal island.

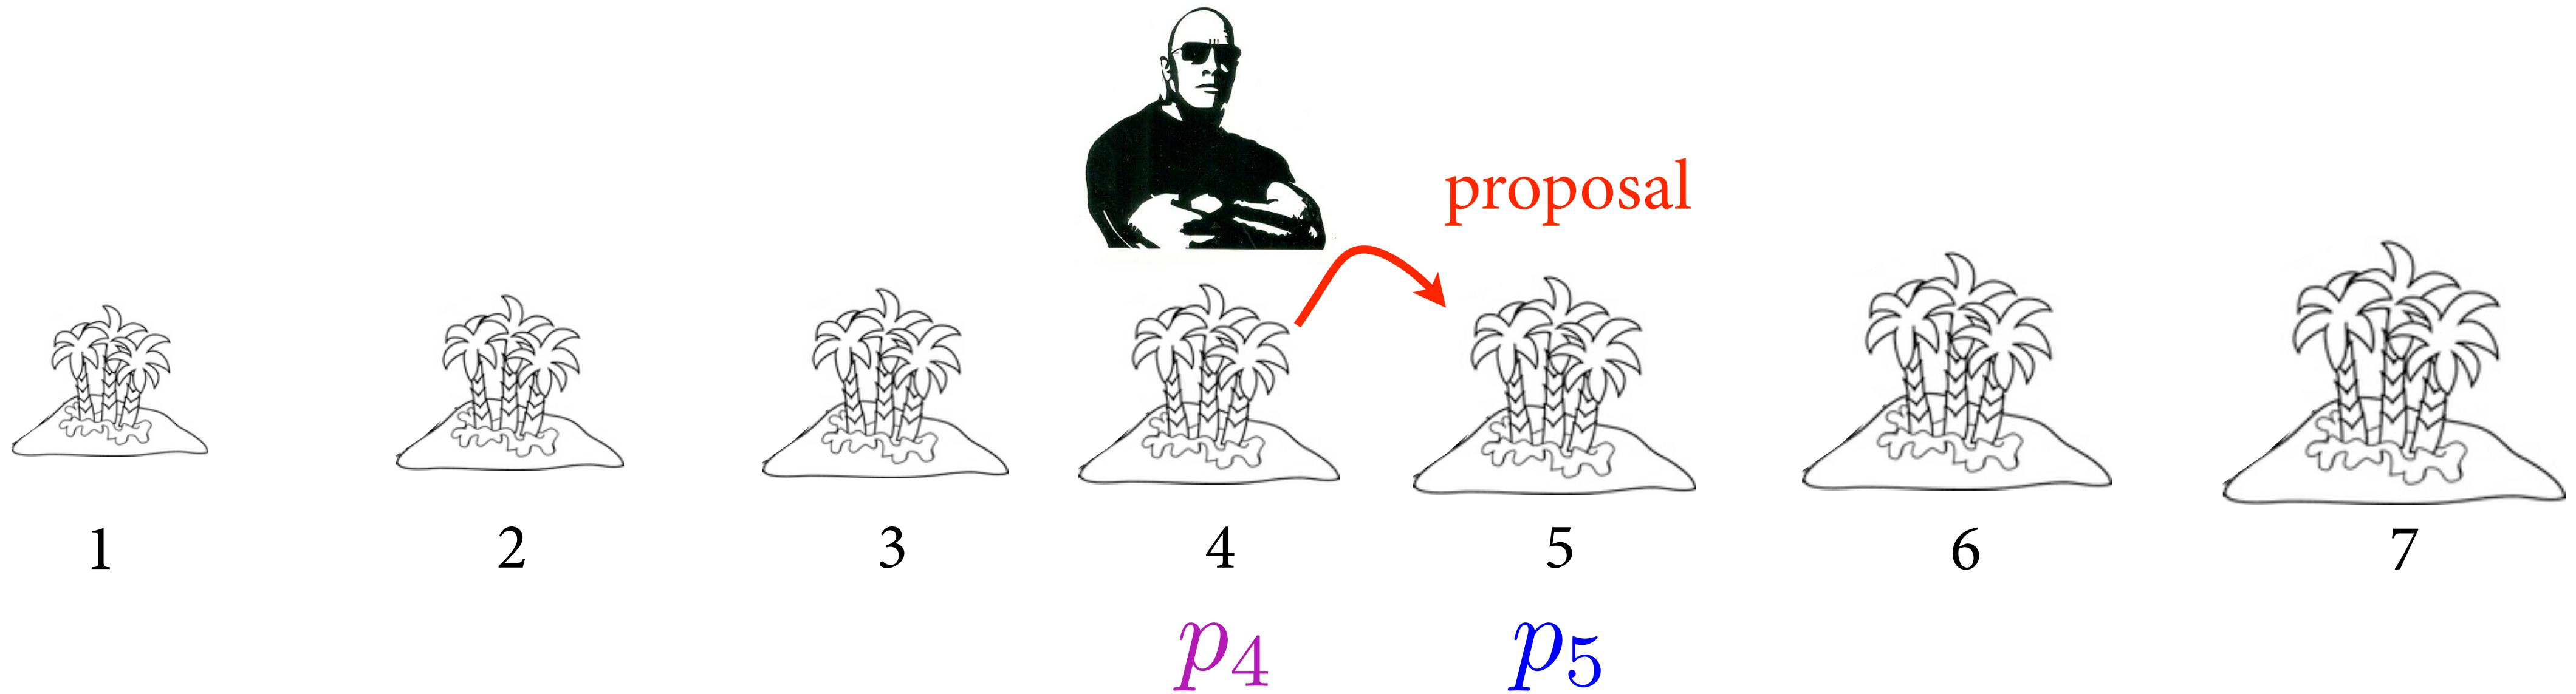
(1) Flip a coin to choose island on left or right.
Call it the “proposal” island.

(2) Find population of proposal island.



(3) Find population of current island.

- (1) Flip a coin to choose island on left or right.
Call it the “proposal” island.
- (2) Find population of proposal island.
- (3) Find population of current island.



(4) Move to proposal, with probability = $\frac{p_5}{p_4}$

(1) Flip a coin to choose island on left or right.

Call it the “proposal” island.

(2) Find population of proposal island.

(3) Find population of current island.

(4) Move to proposal, with probability = $\frac{p_5}{p_4}$



1



2



3



4



5



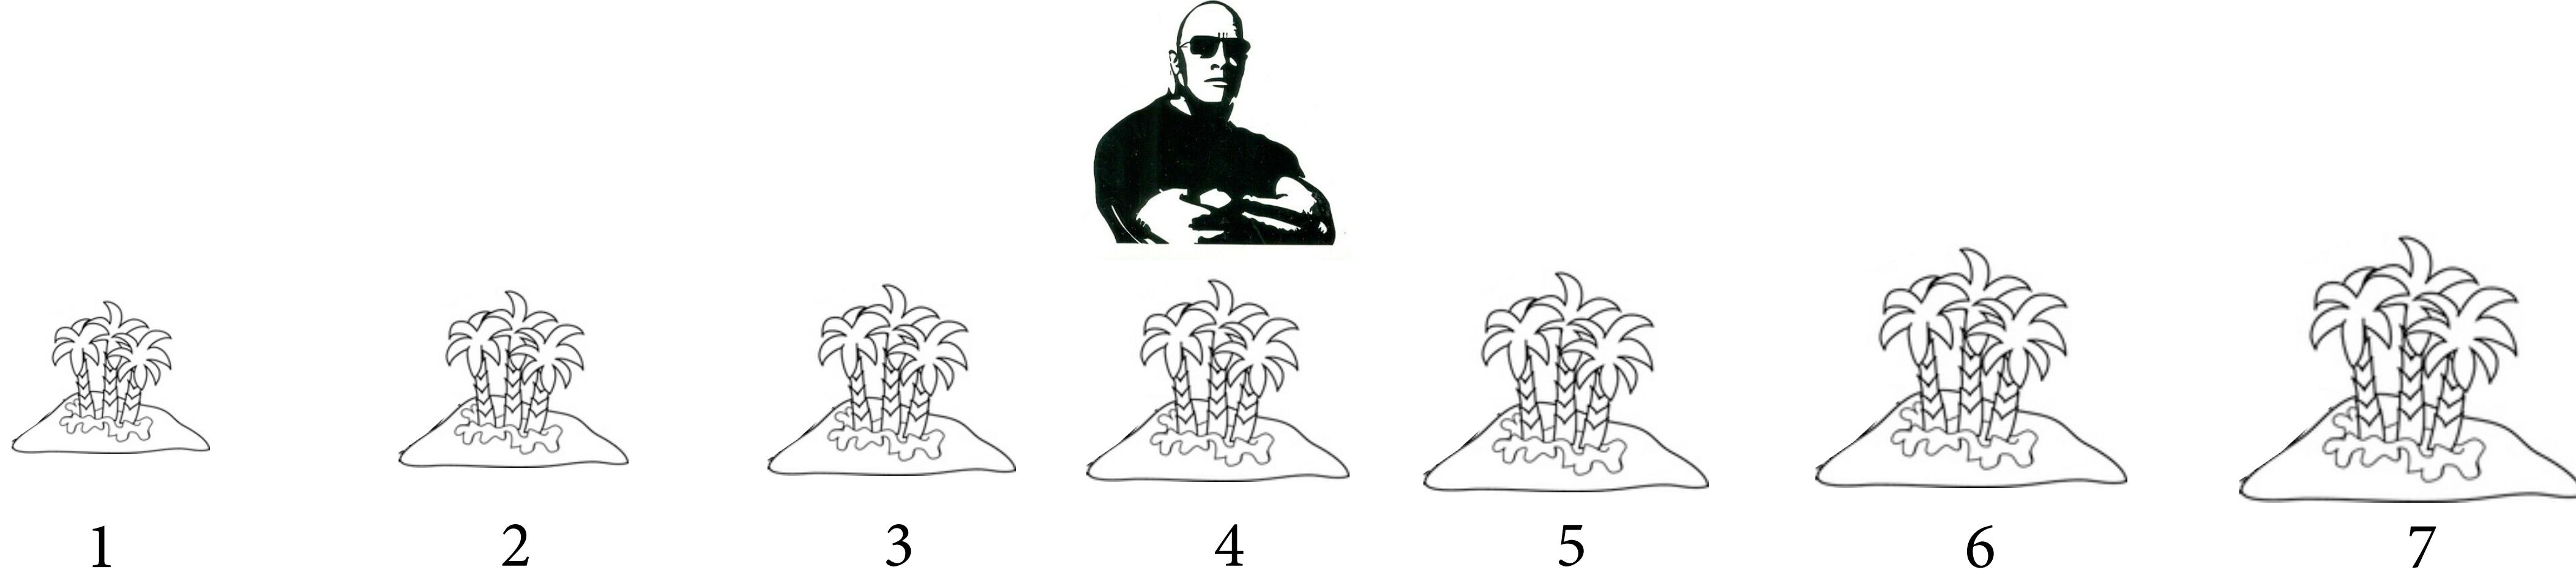
6



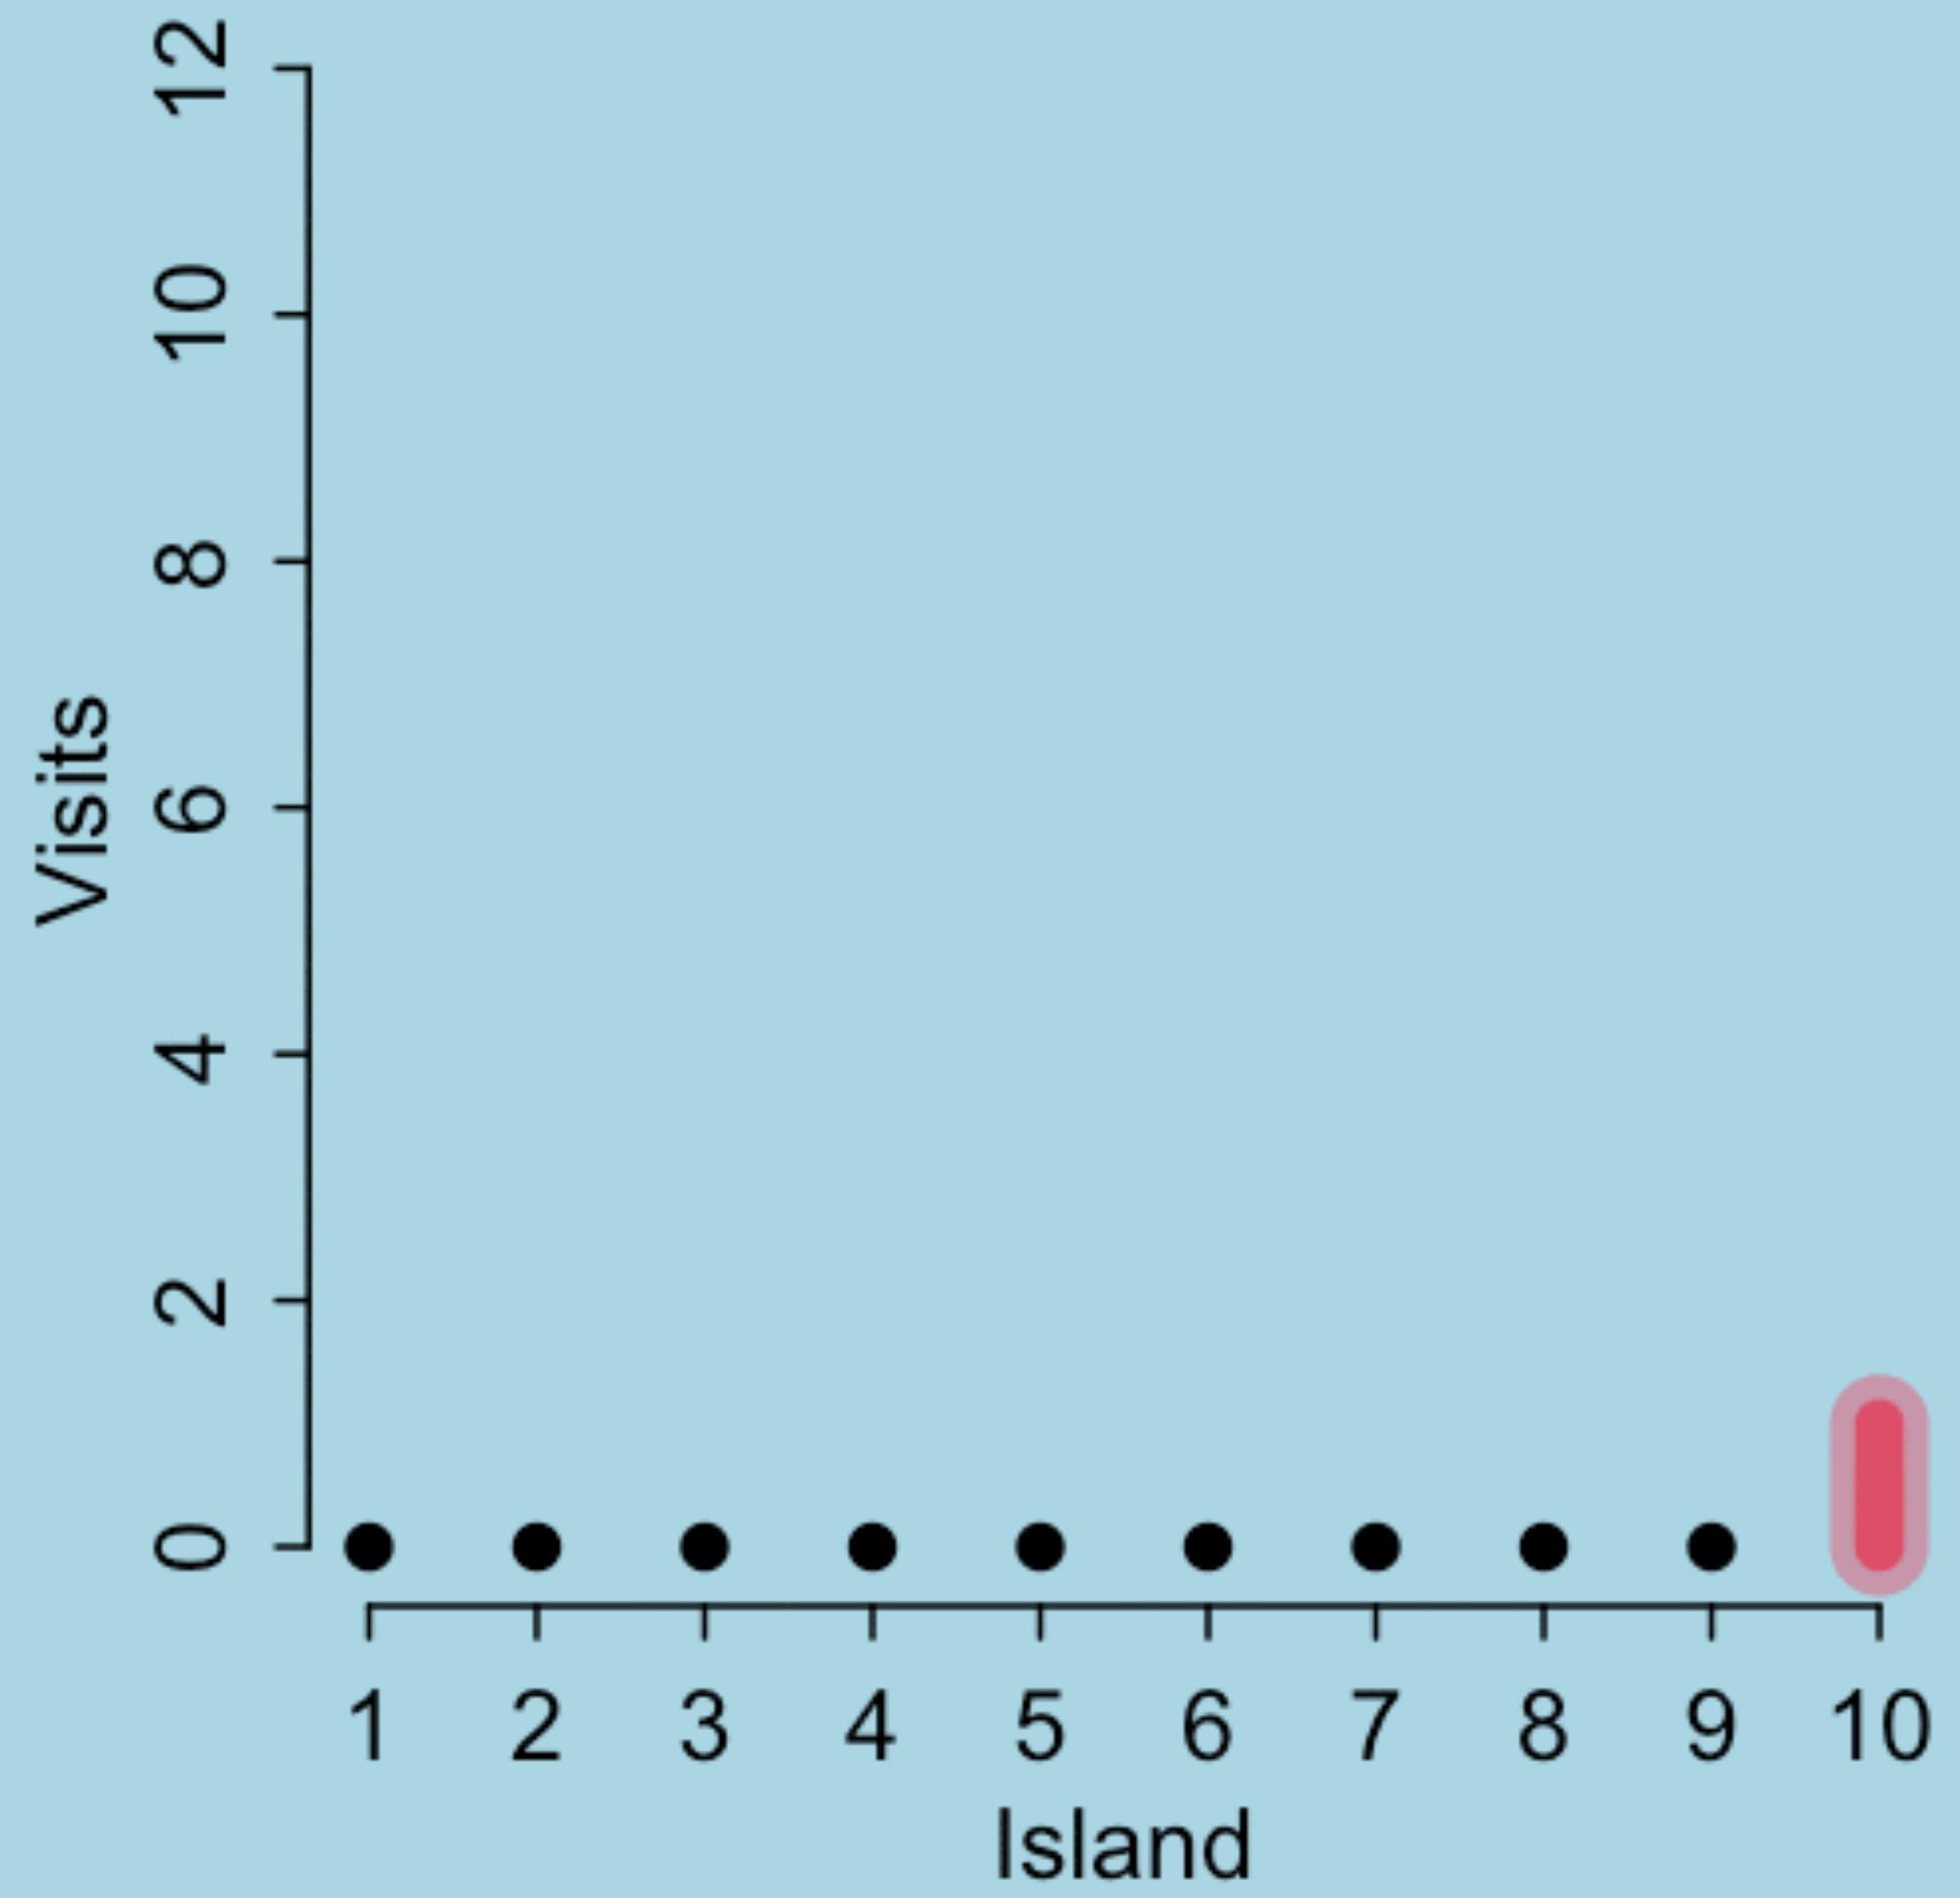
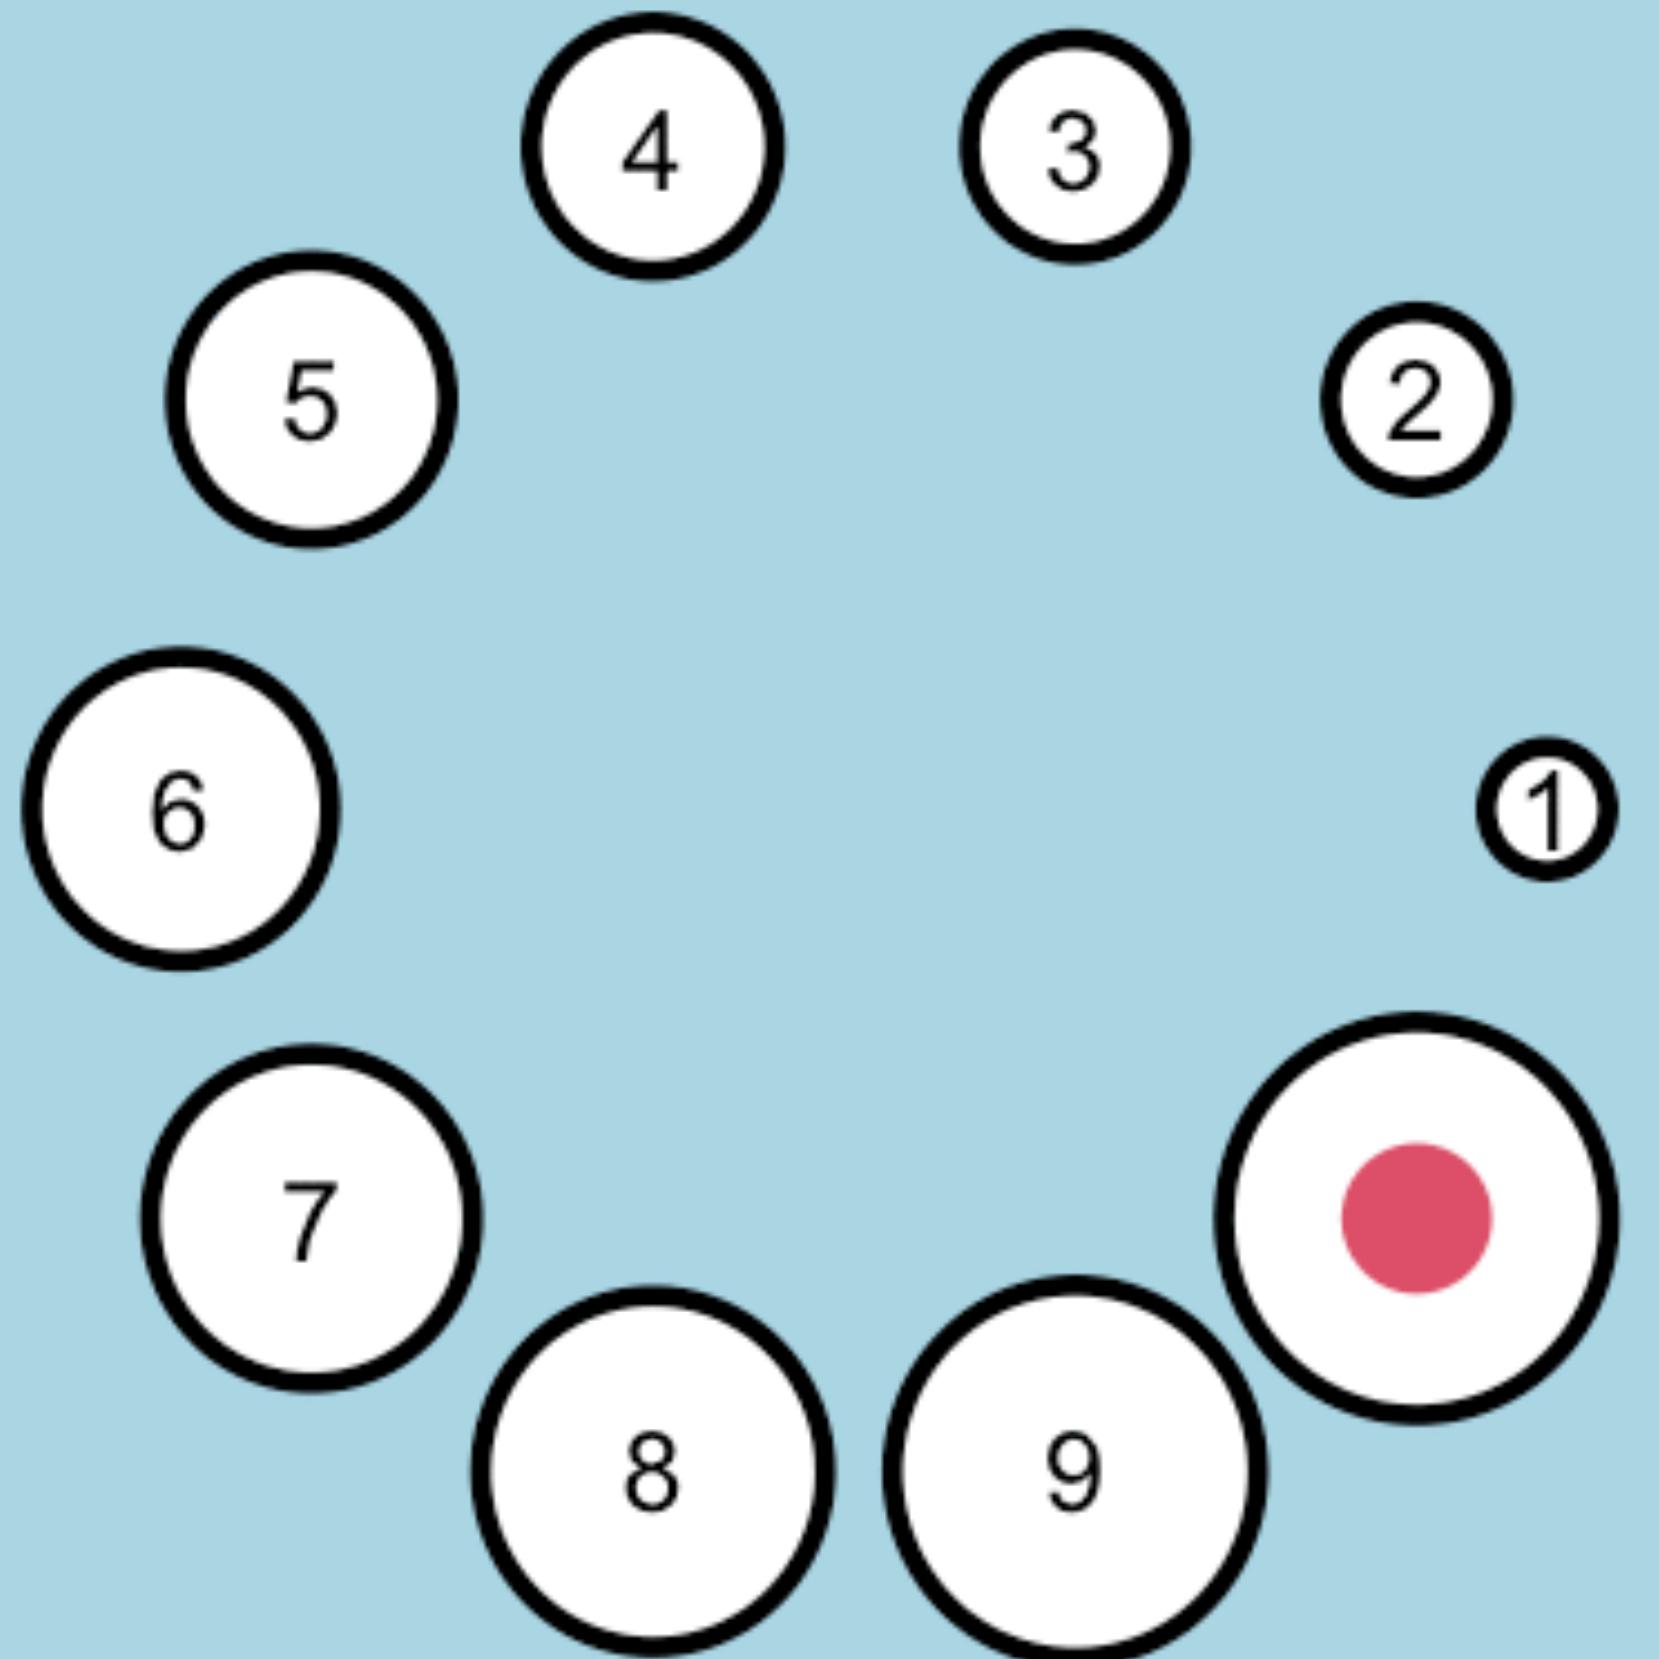
7

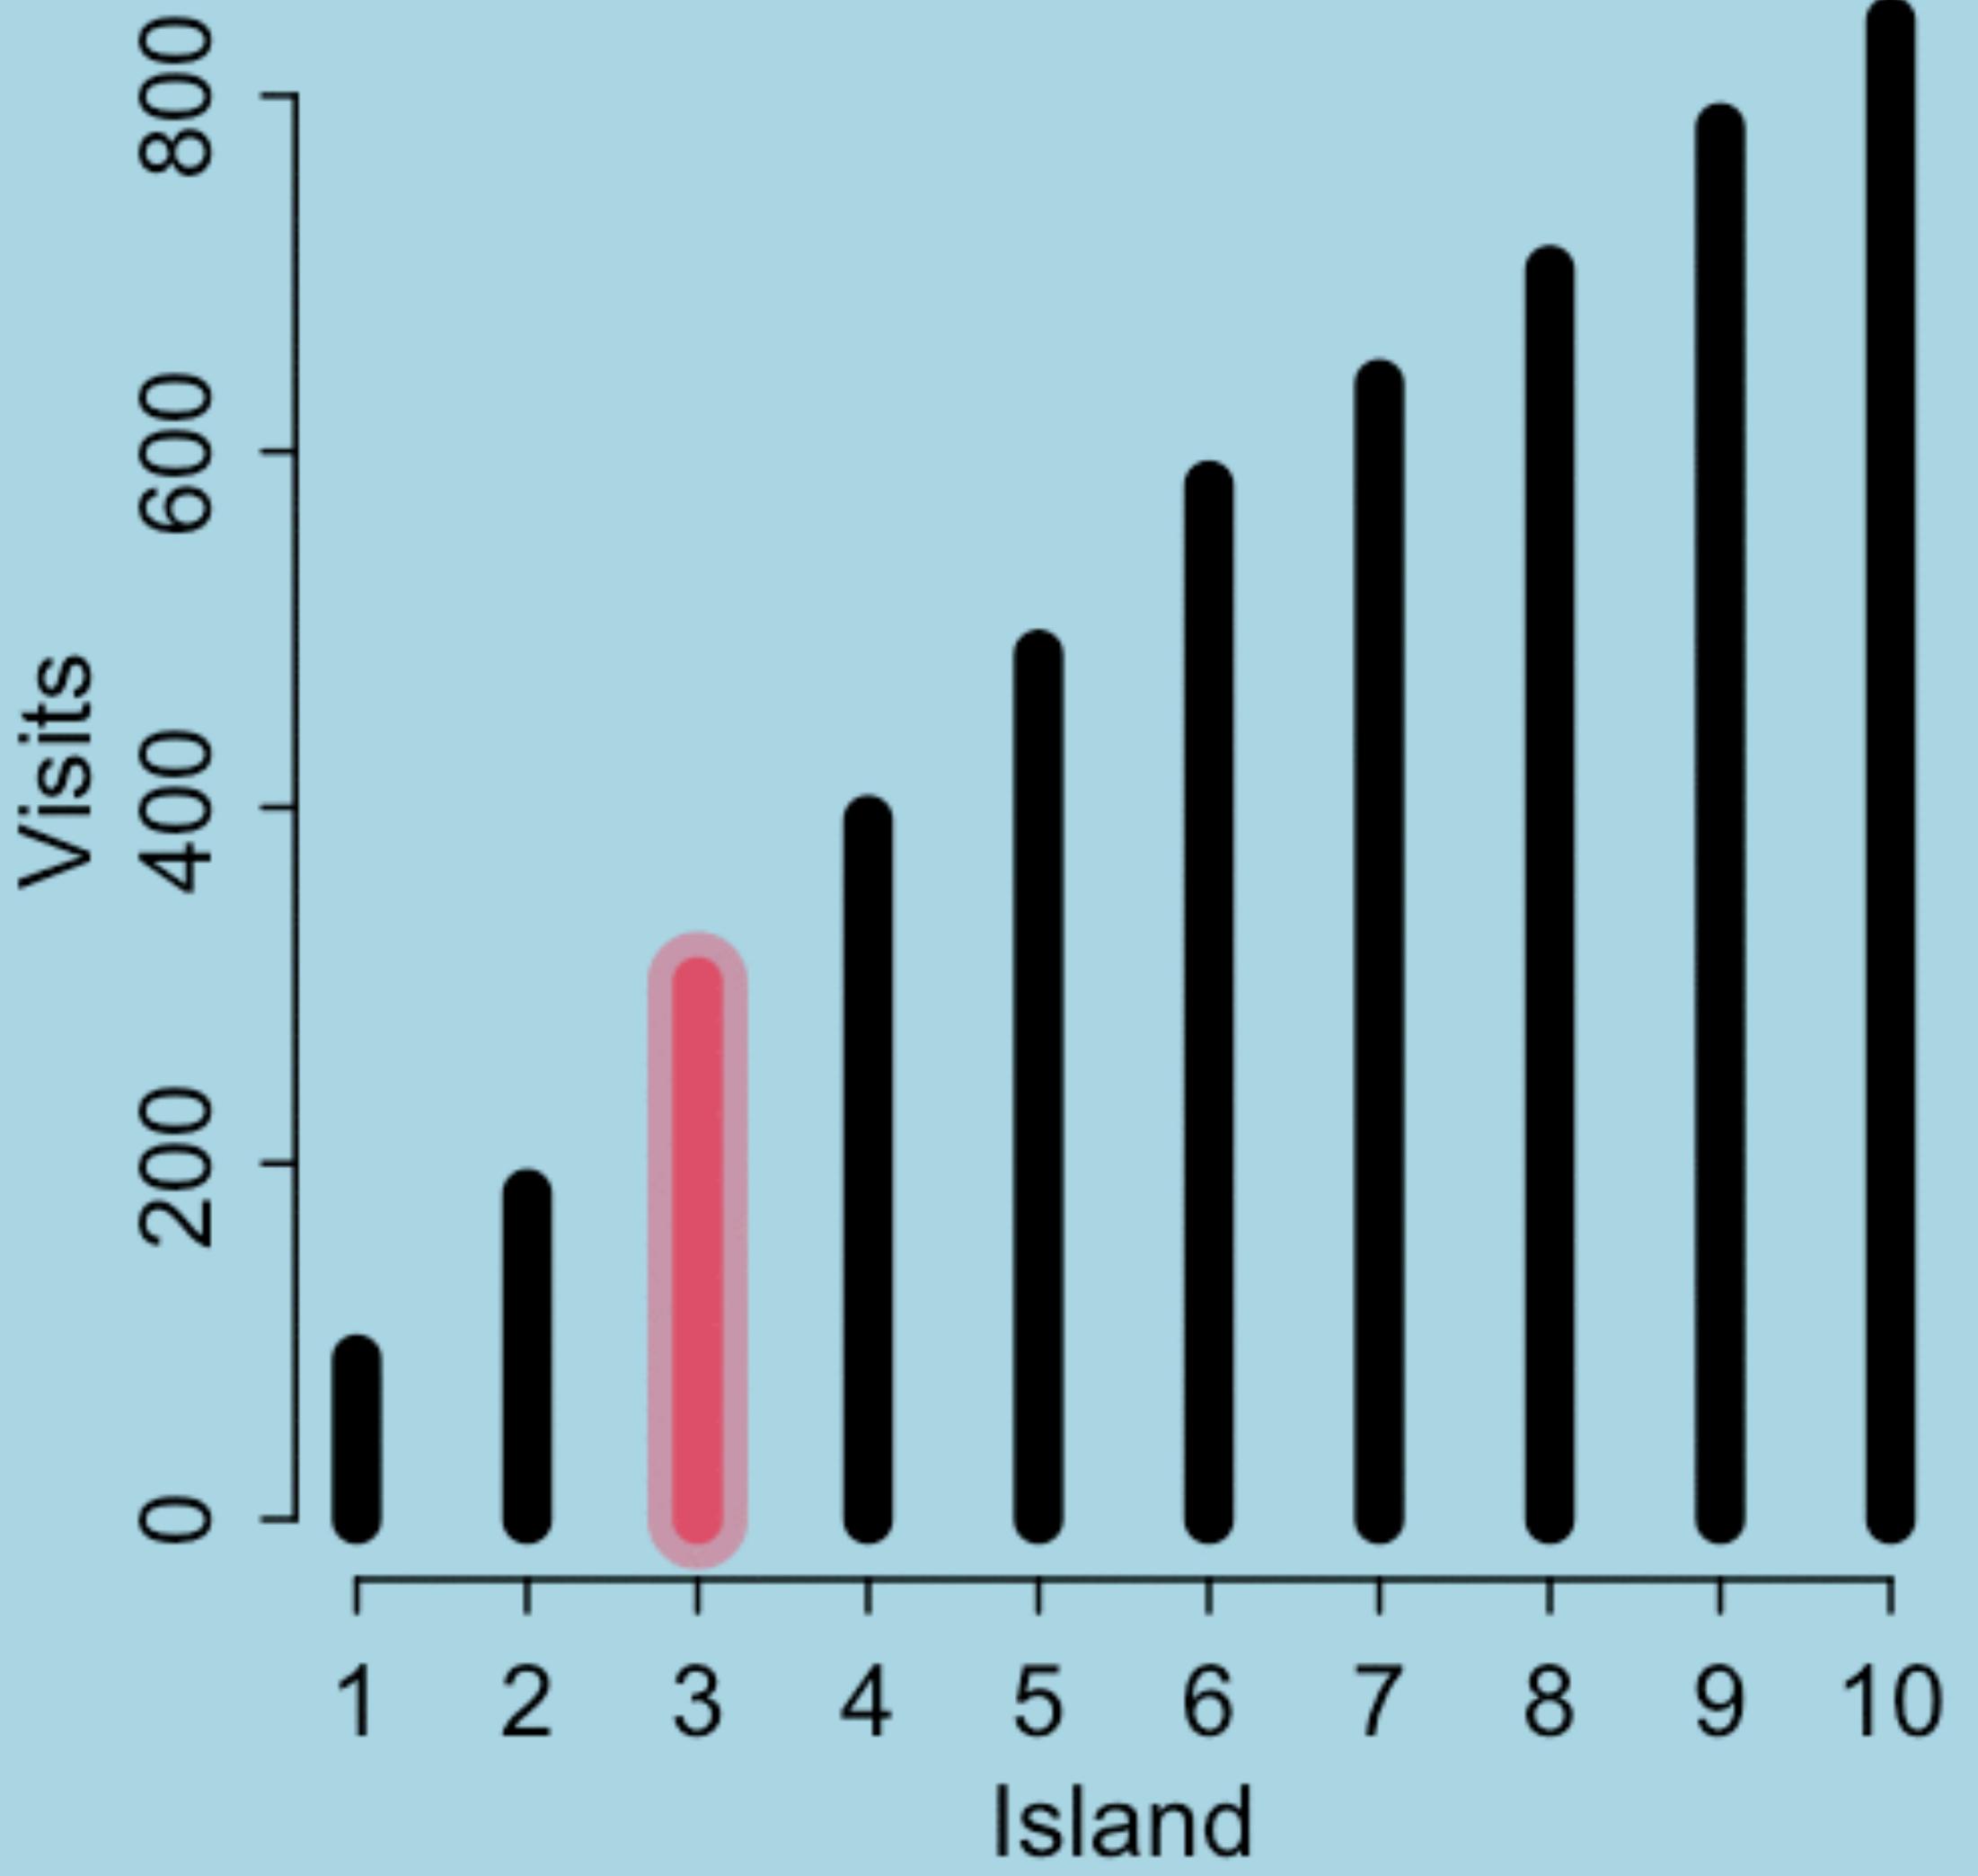
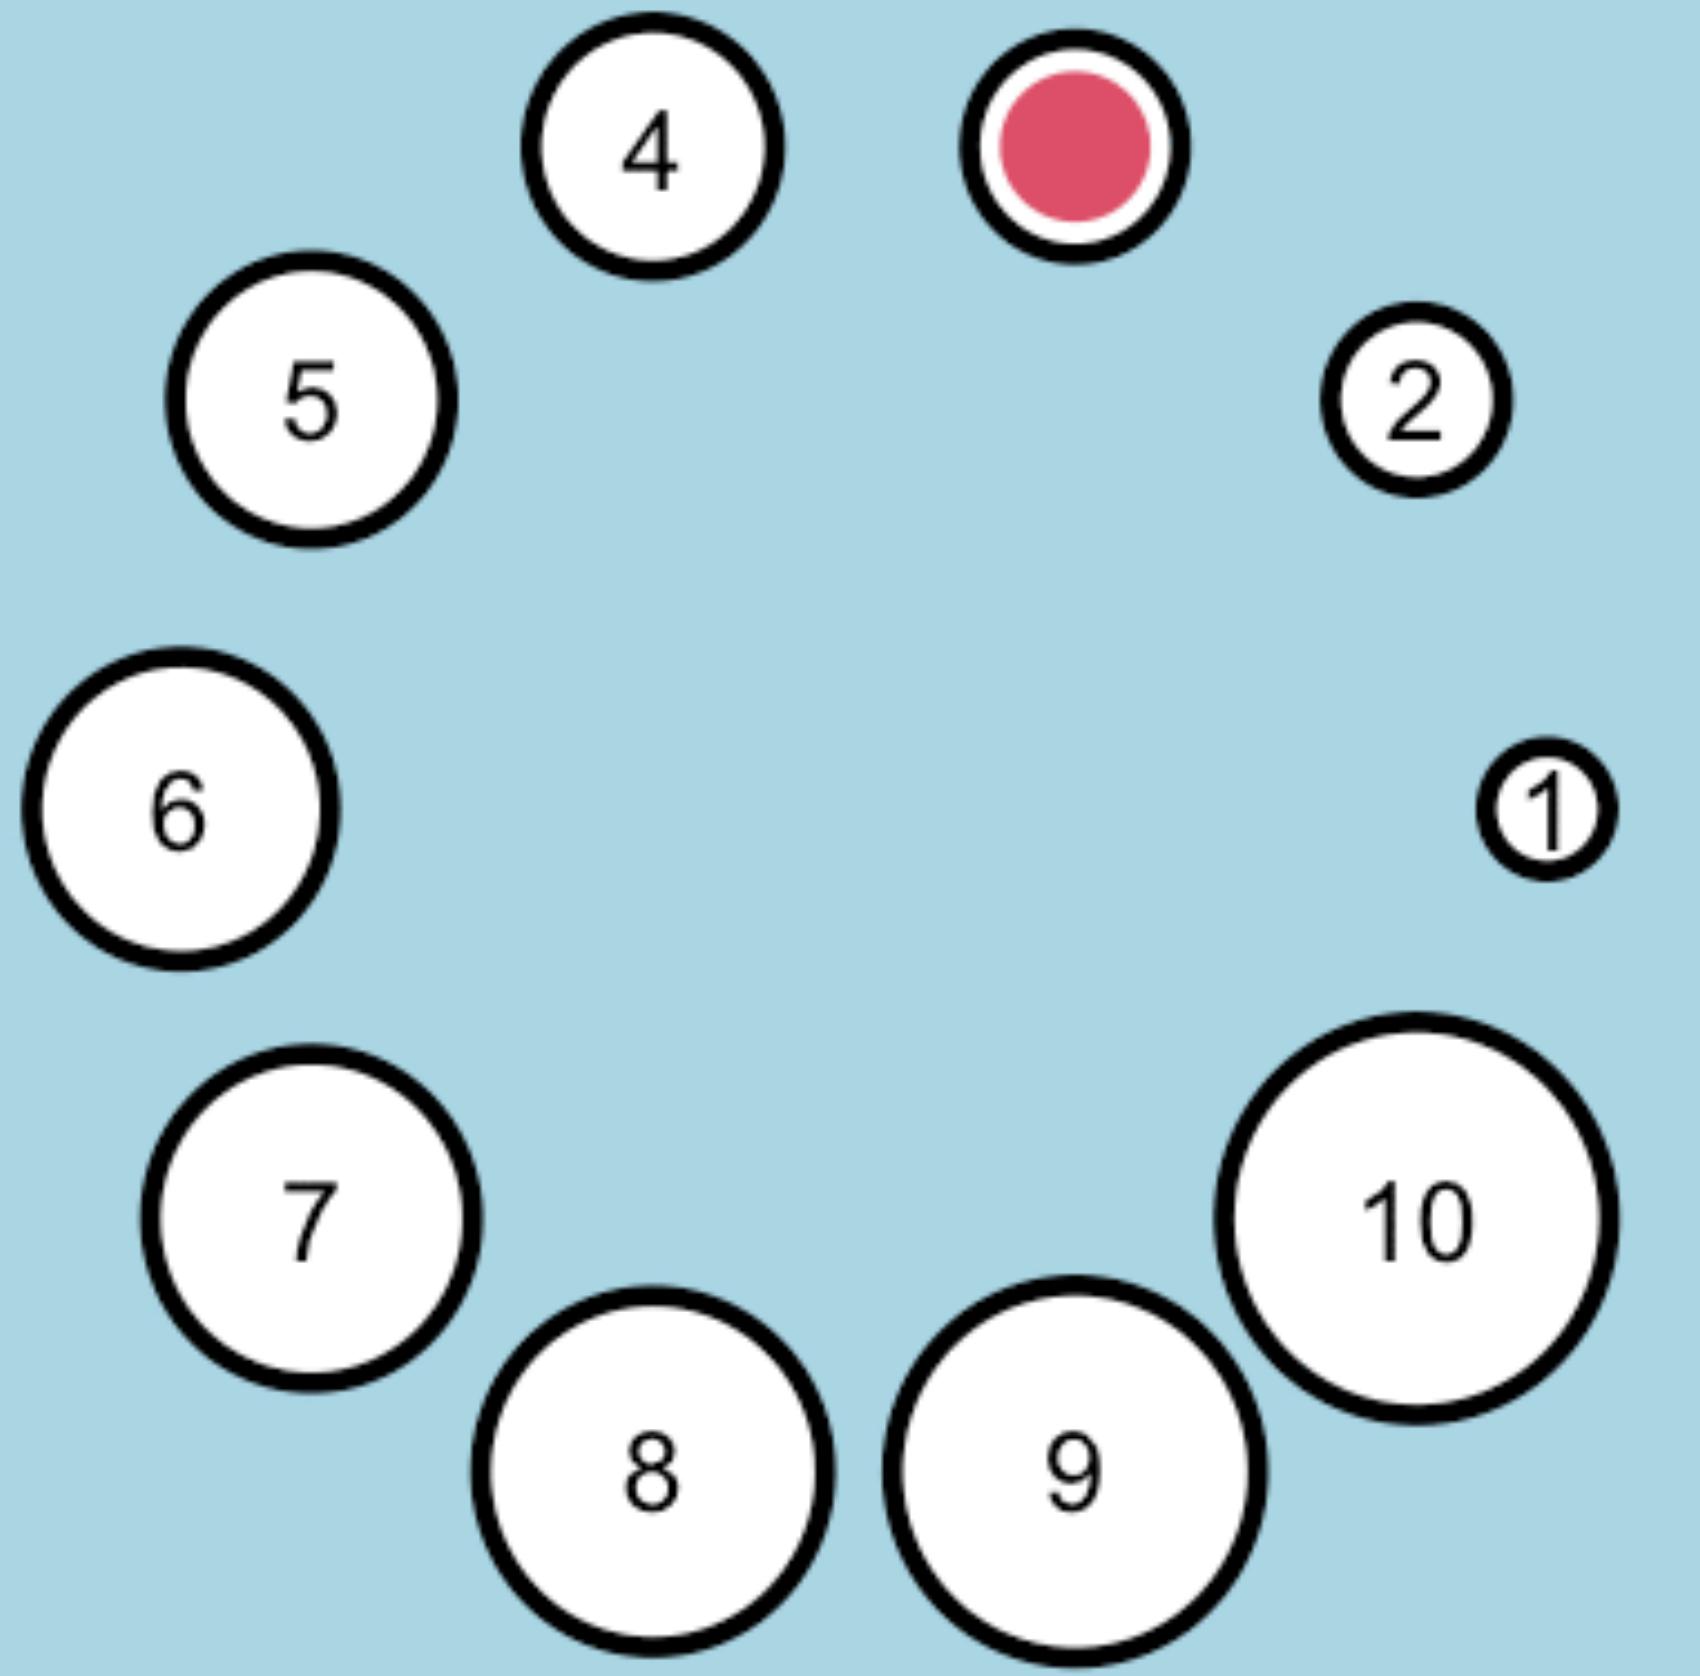
(5) Repeat from (1)

- (1) Flip a coin to choose island on left or right.
Call it the “proposal” island.
- (2) Find population of proposal island.
- (3) Find population of current island.
- (4) Move to proposal, with probability = $\frac{p_5}{p_4}$
- (5) Repeat from (1)



This procedure ensures visiting each island in proportion to its population, *in the long run*.





Markov chain Monte Carlo

Usual use: Draw samples from a posterior distribution



“Islands”: parameter values

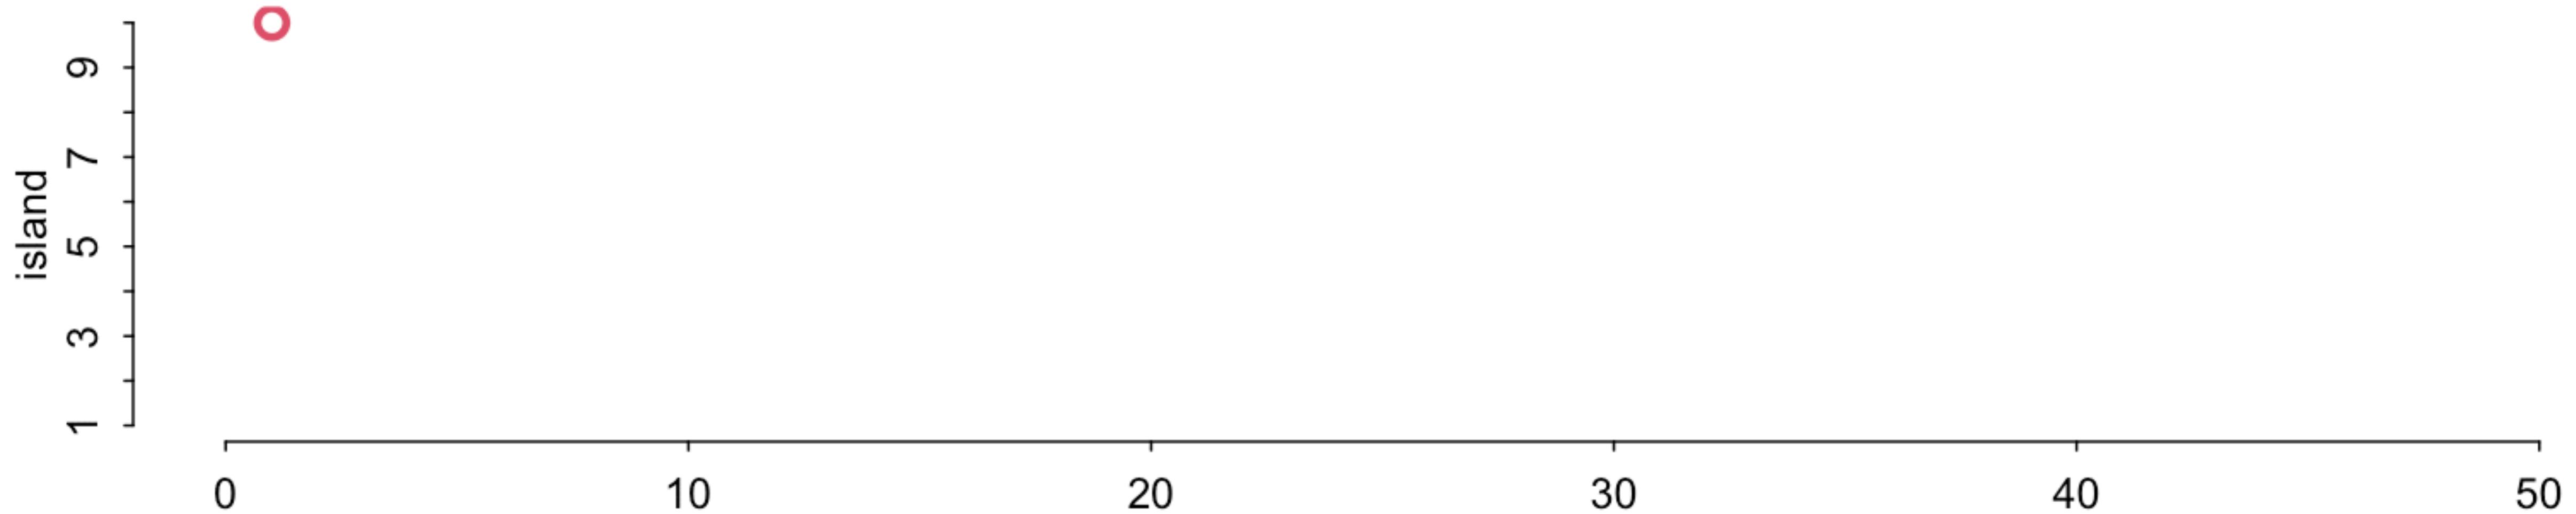
“Population size”: posterior probability

Visit each parameter value in proportion to its posterior probability



Any number of dimensions (parameters)

“Markov chain Monte Carlo”

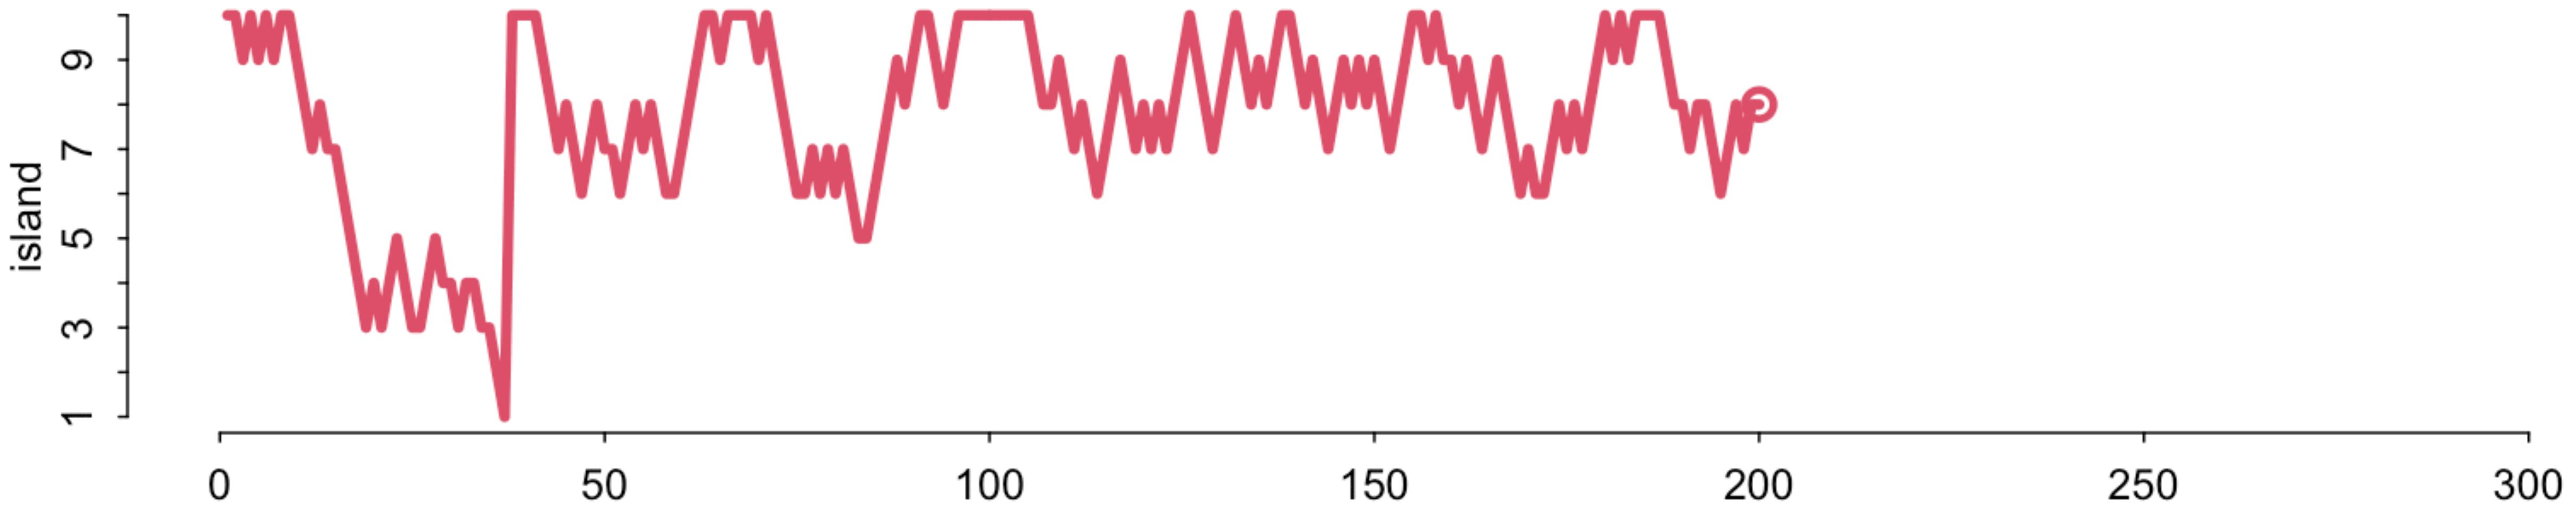


Chain: Sequence of draws from distribution

Markov chain: History doesn't matter, just where you are now

Monte Carlo: Random simulation

“Markov chain Monte Carlo”



Metropolis algorithm: Simple version of *Markov chain Monte Carlo* (MCMC)

Easy to write, very general, often inefficient

Metropolis, Rosenbluth, Rosenbluth, Teller and Teller (1953)

Equation of State Calculations by Fast Computing Machines

NICHOLAS METROPOLIS, ARIANNA W. ROSENBLUTH, MARSHALL N. ROSENBLUTH, AND AUGUSTA H. TELLER,
Los Alamos Scientific Laboratory, Los Alamos, New Mexico

AND

EDWARD TELLER,* *Department of Physics, University of Chicago, Chicago, Illinois*
(Received March 6, 1953)

A general method, suitable for fast computing machines, for investigating such properties as equations of state for substances consisting of interacting individual molecules is described. The method consists of a modified Monte Carlo integration over configuration space. Results for the two-dimensional rigid-sphere system have been obtained on the Los Alamos MANIAC and are presented here. These results are compared to the free volume equation of state and to a four-term virial coefficient expansion.

Metropolis, Rosenbluth, Rosenbluth, Teller and Teller (1953)

Equation of state calculations by fast computing machines

NICHOLAS METROPOLIS

Editor

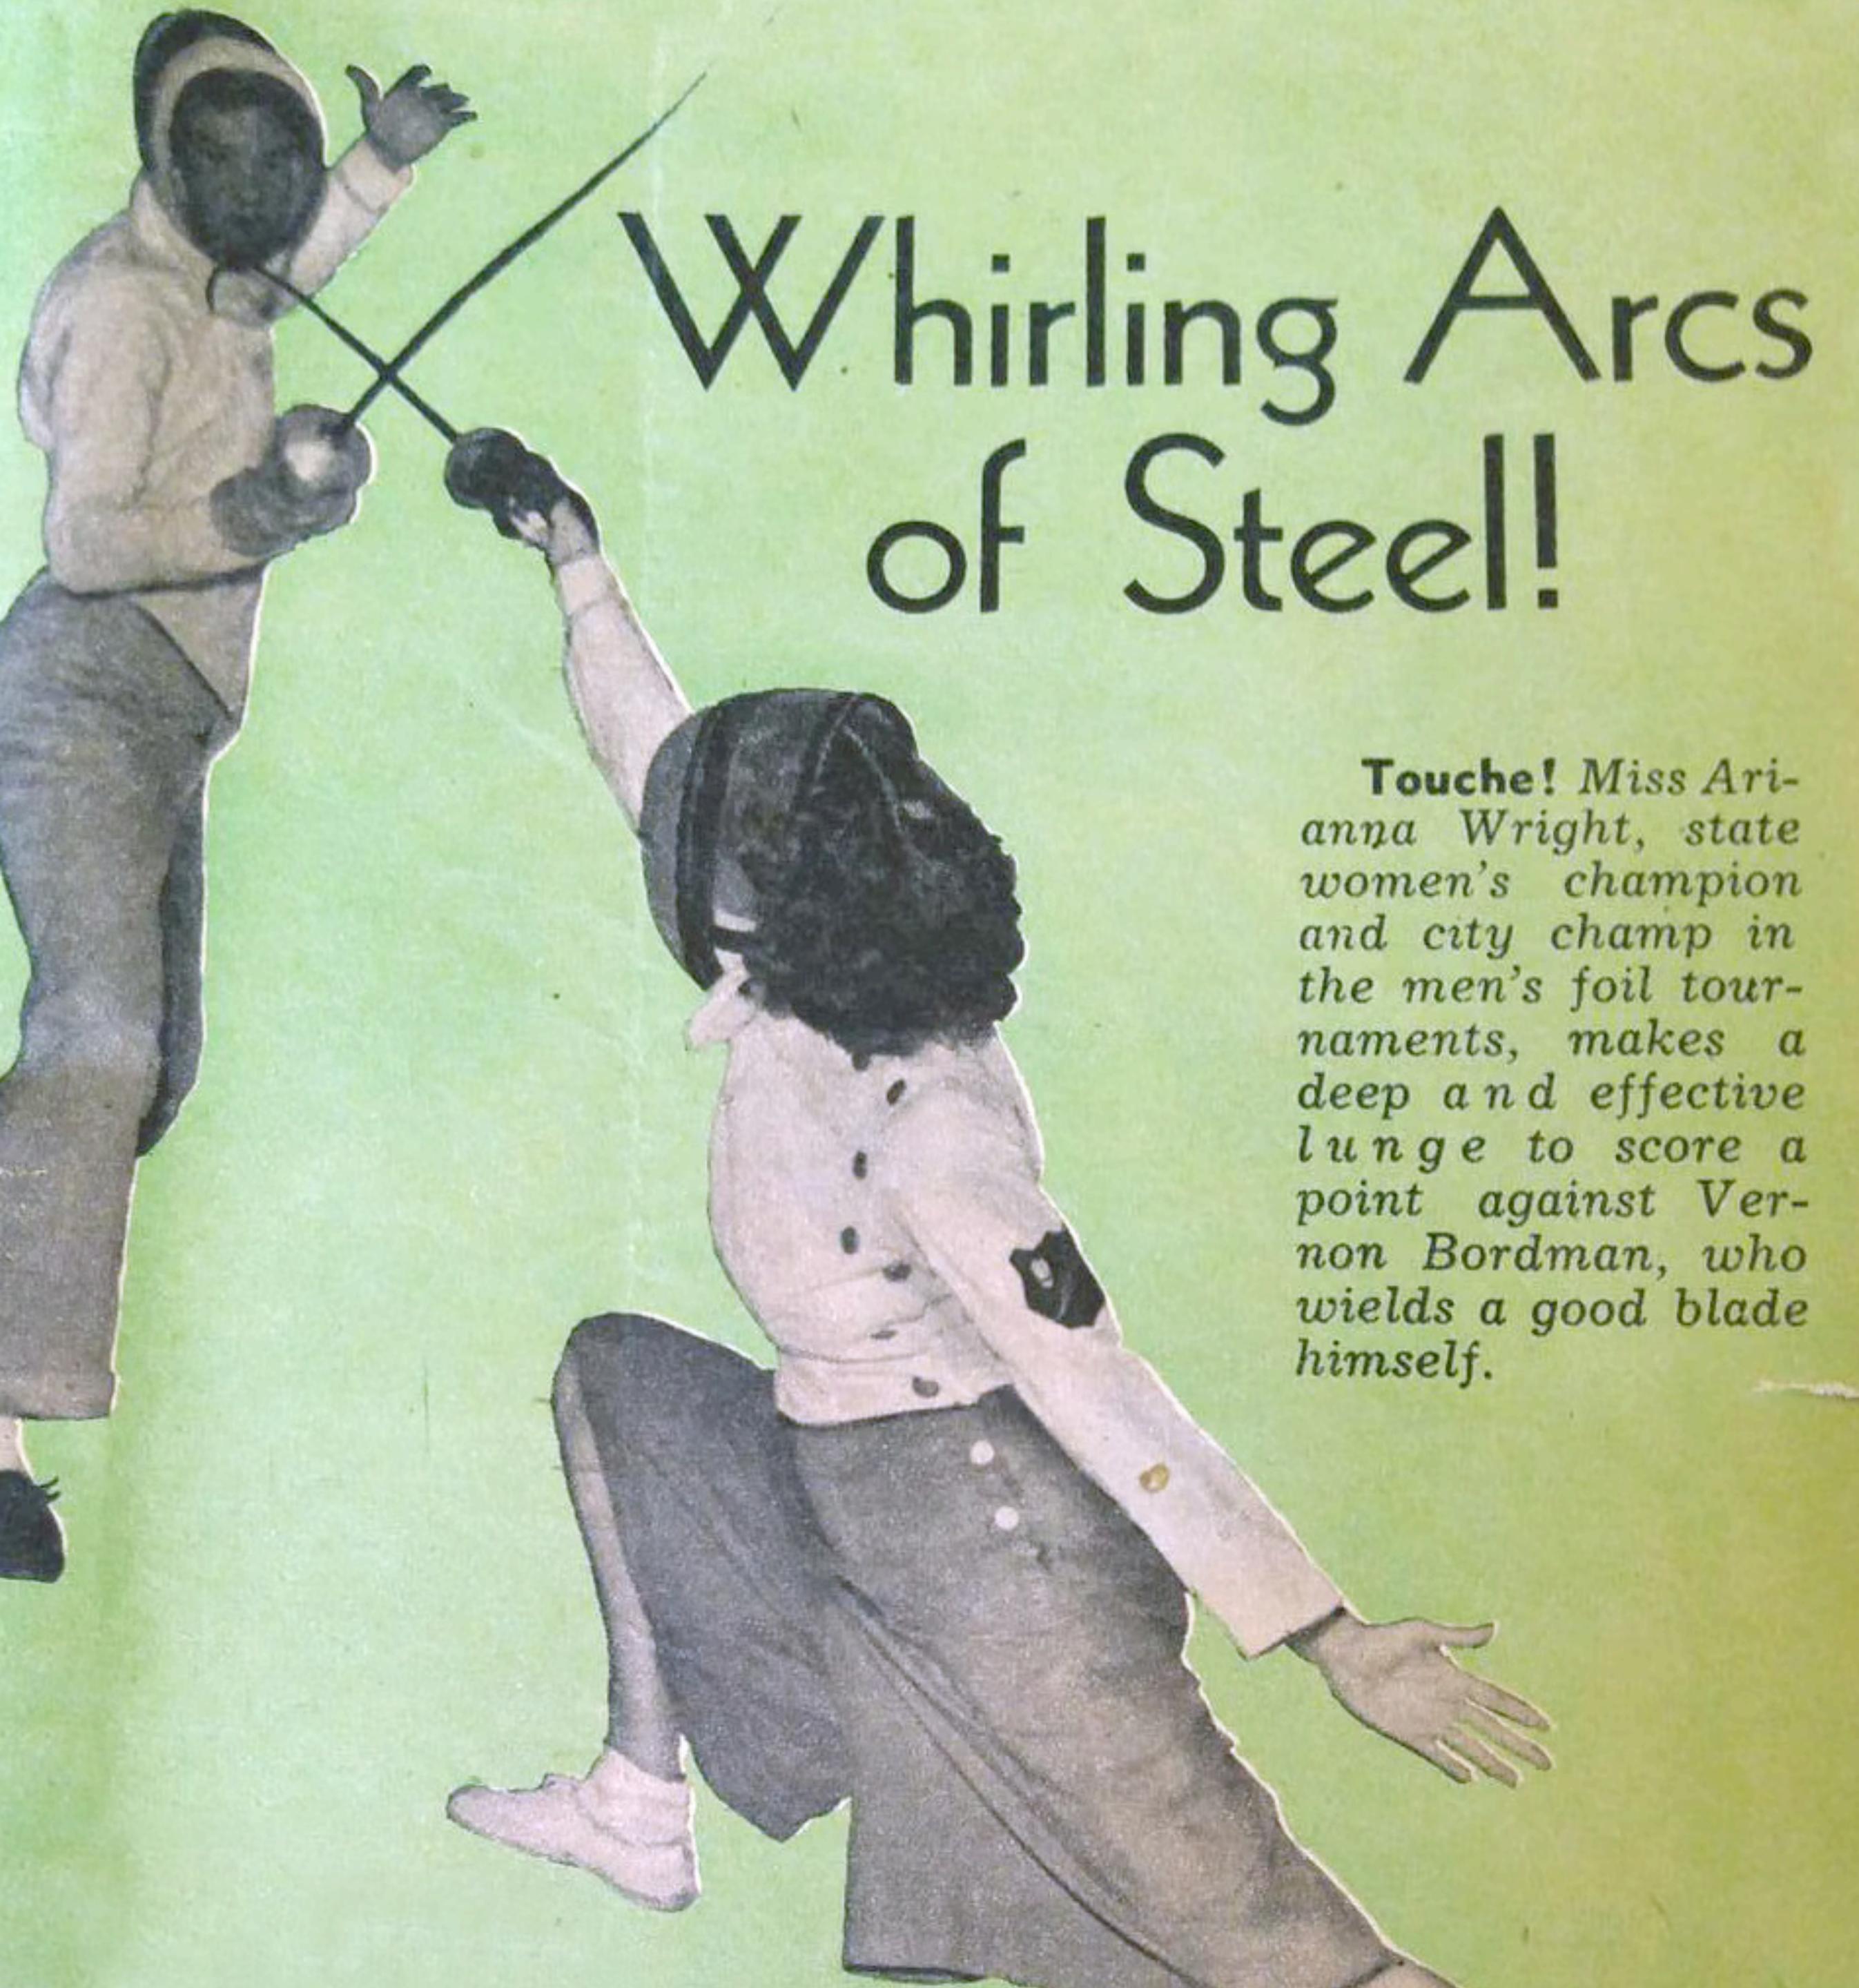
Equation of state calculations by fast computing machines

..., AW Rosenbluth, MN Rosenbluth... - The journal of ..., 1953 - aip.scitation.org

A general method, suitable for fast computing machines, for investigating such properties as equations of state for substances consisting of interacting individual molecules is described. The method consists of a modified Monte Carlo integration over configuration space. Results for the two-dimensional rigid-sphere system have been obtained on the Los Alamos MANIAC and are presented here. These results are compared to the free volume equation of state and to a four-term virial coefficient expansion.

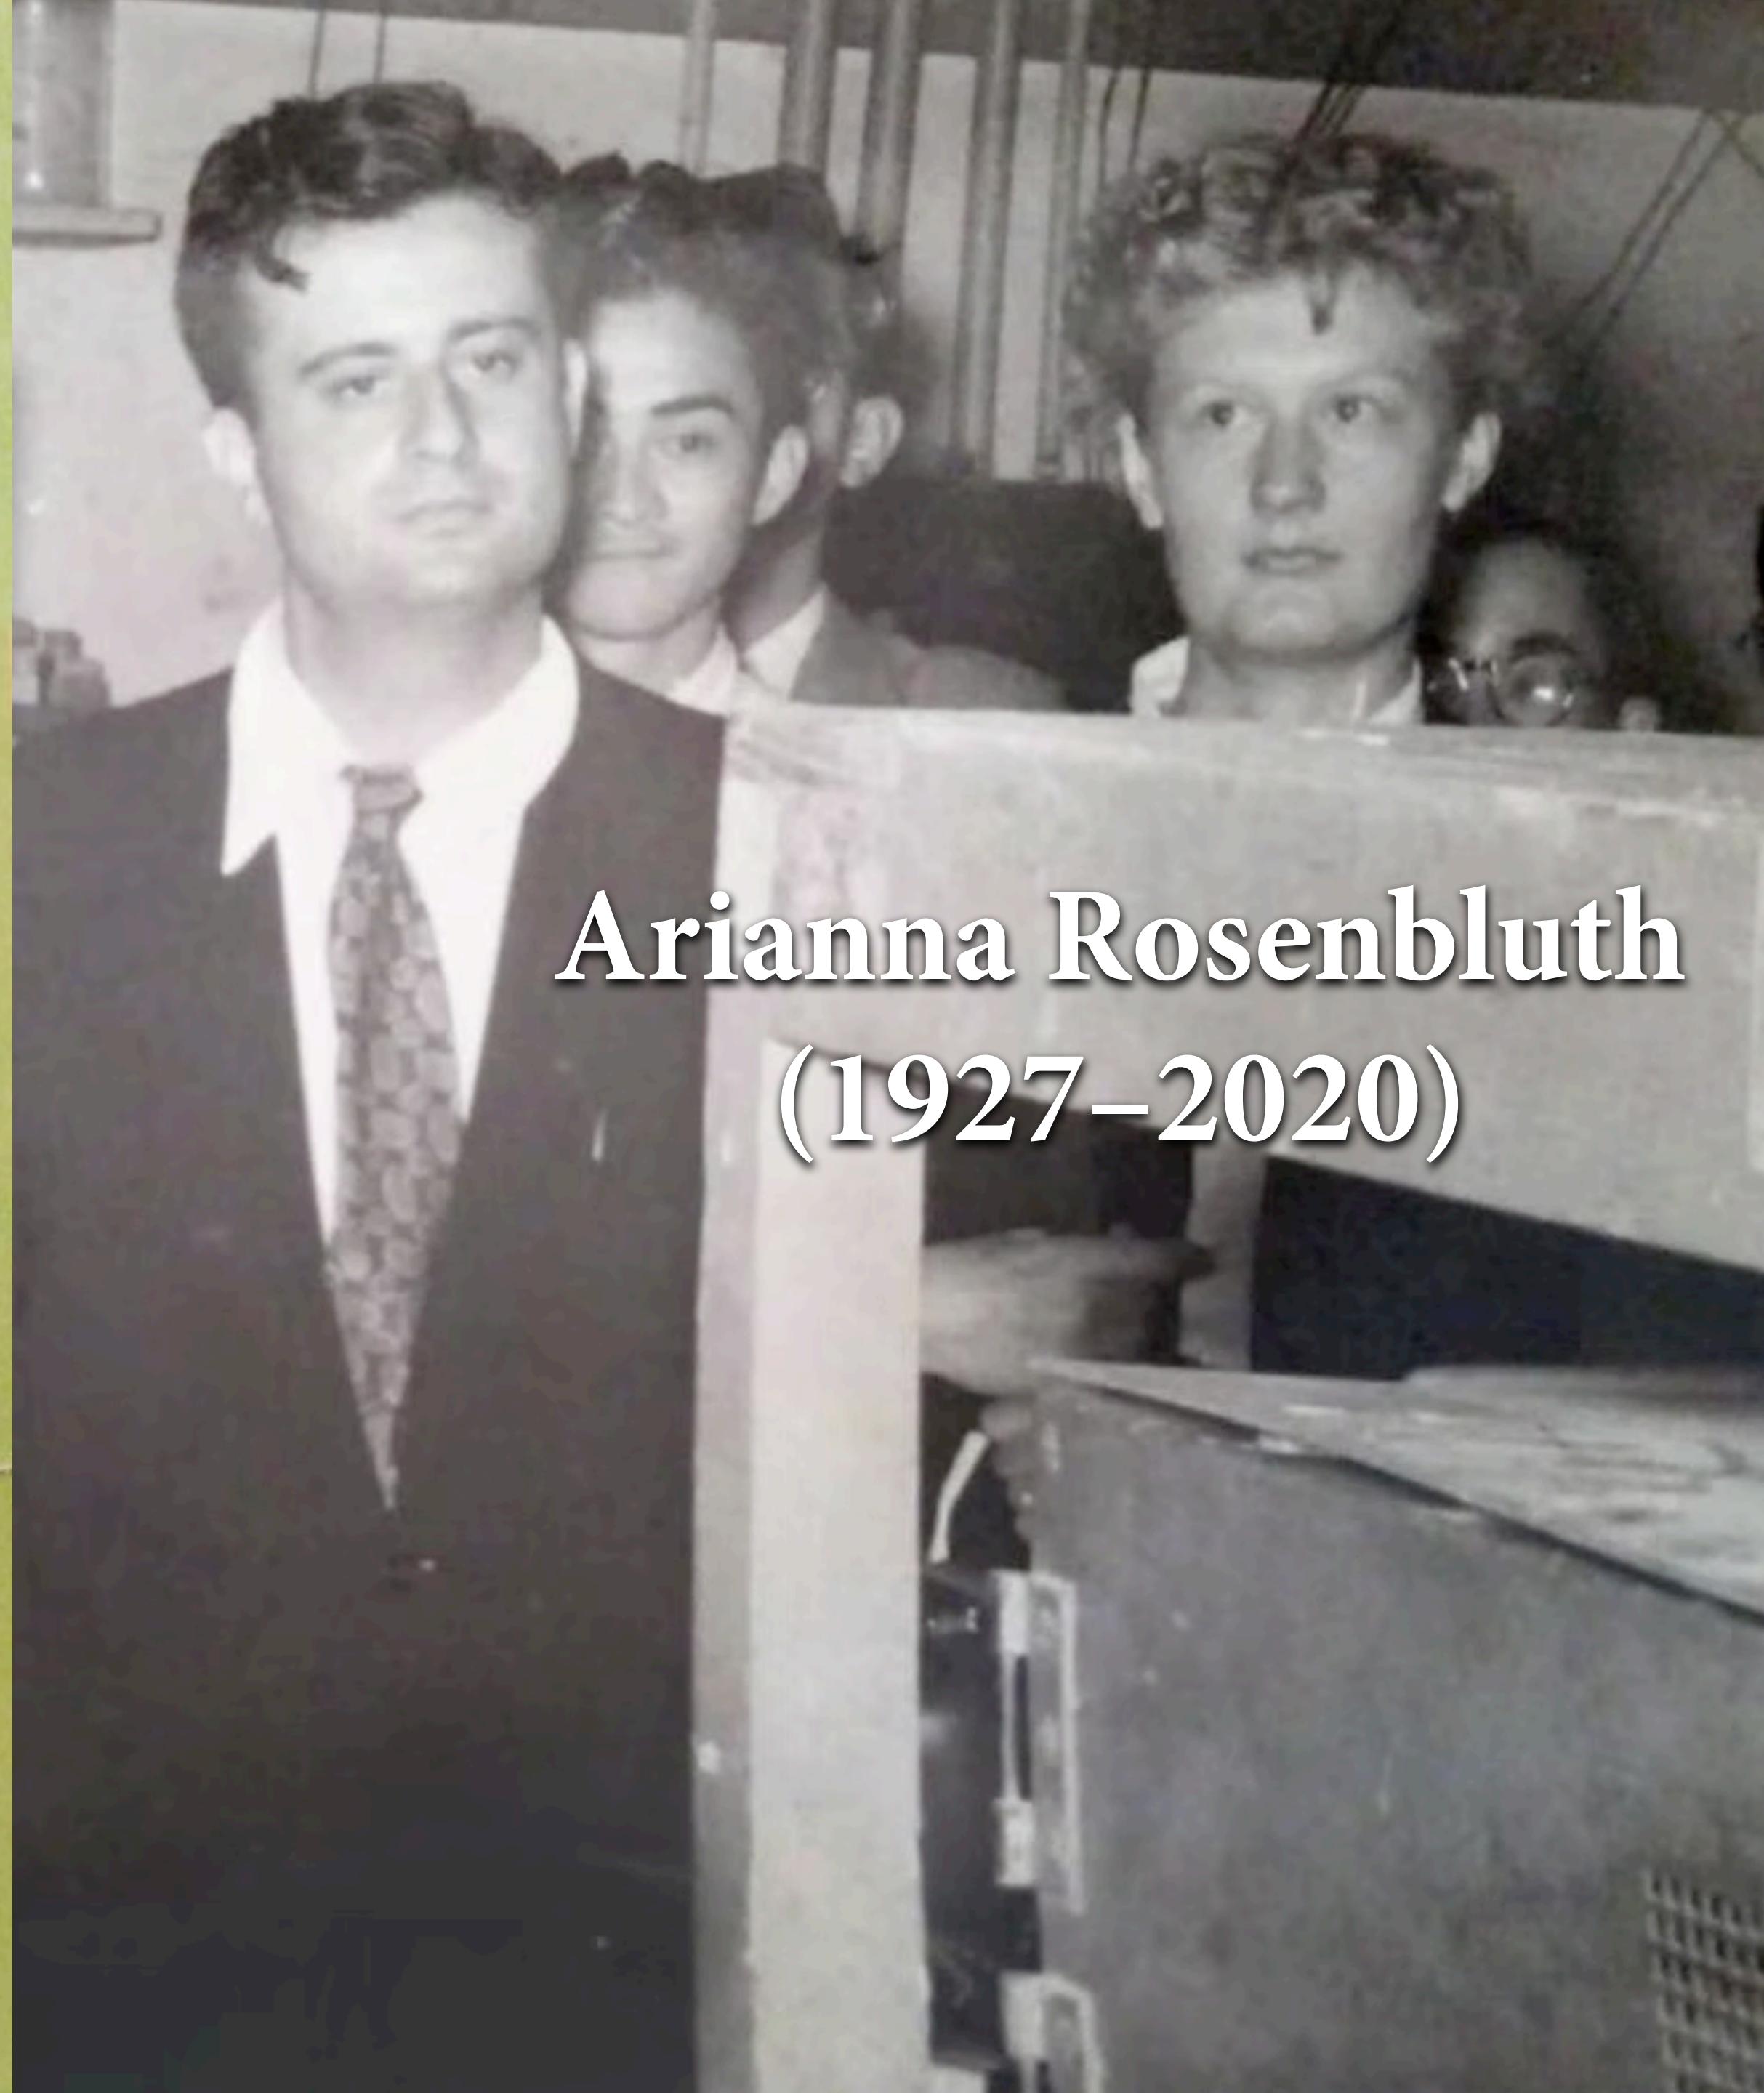
☆ Save 99 Cite Cited by 47206 Related articles All 29 versions

A general method, suitable for fast computing machines, for investigating such properties as equations of state for substances consisting of interacting individual molecules is described. The method consists of a modified Monte Carlo integration over configuration space. Results for the two-dimensional rigid-sphere system have been obtained on the Los Alamos MANIAC and are presented here. These results are compared to the free volume equation of state and to a four-term virial coefficient expansion.



Whirling Arcs of Steel!

Touche! Miss Arianna Wright, state women's champion and city champ in the men's foil tournaments, makes a deep and effective lunge to score a point against Vernon Bordman, who wields a good blade himself.



Arianna Rosenbluth
(1927–2020)

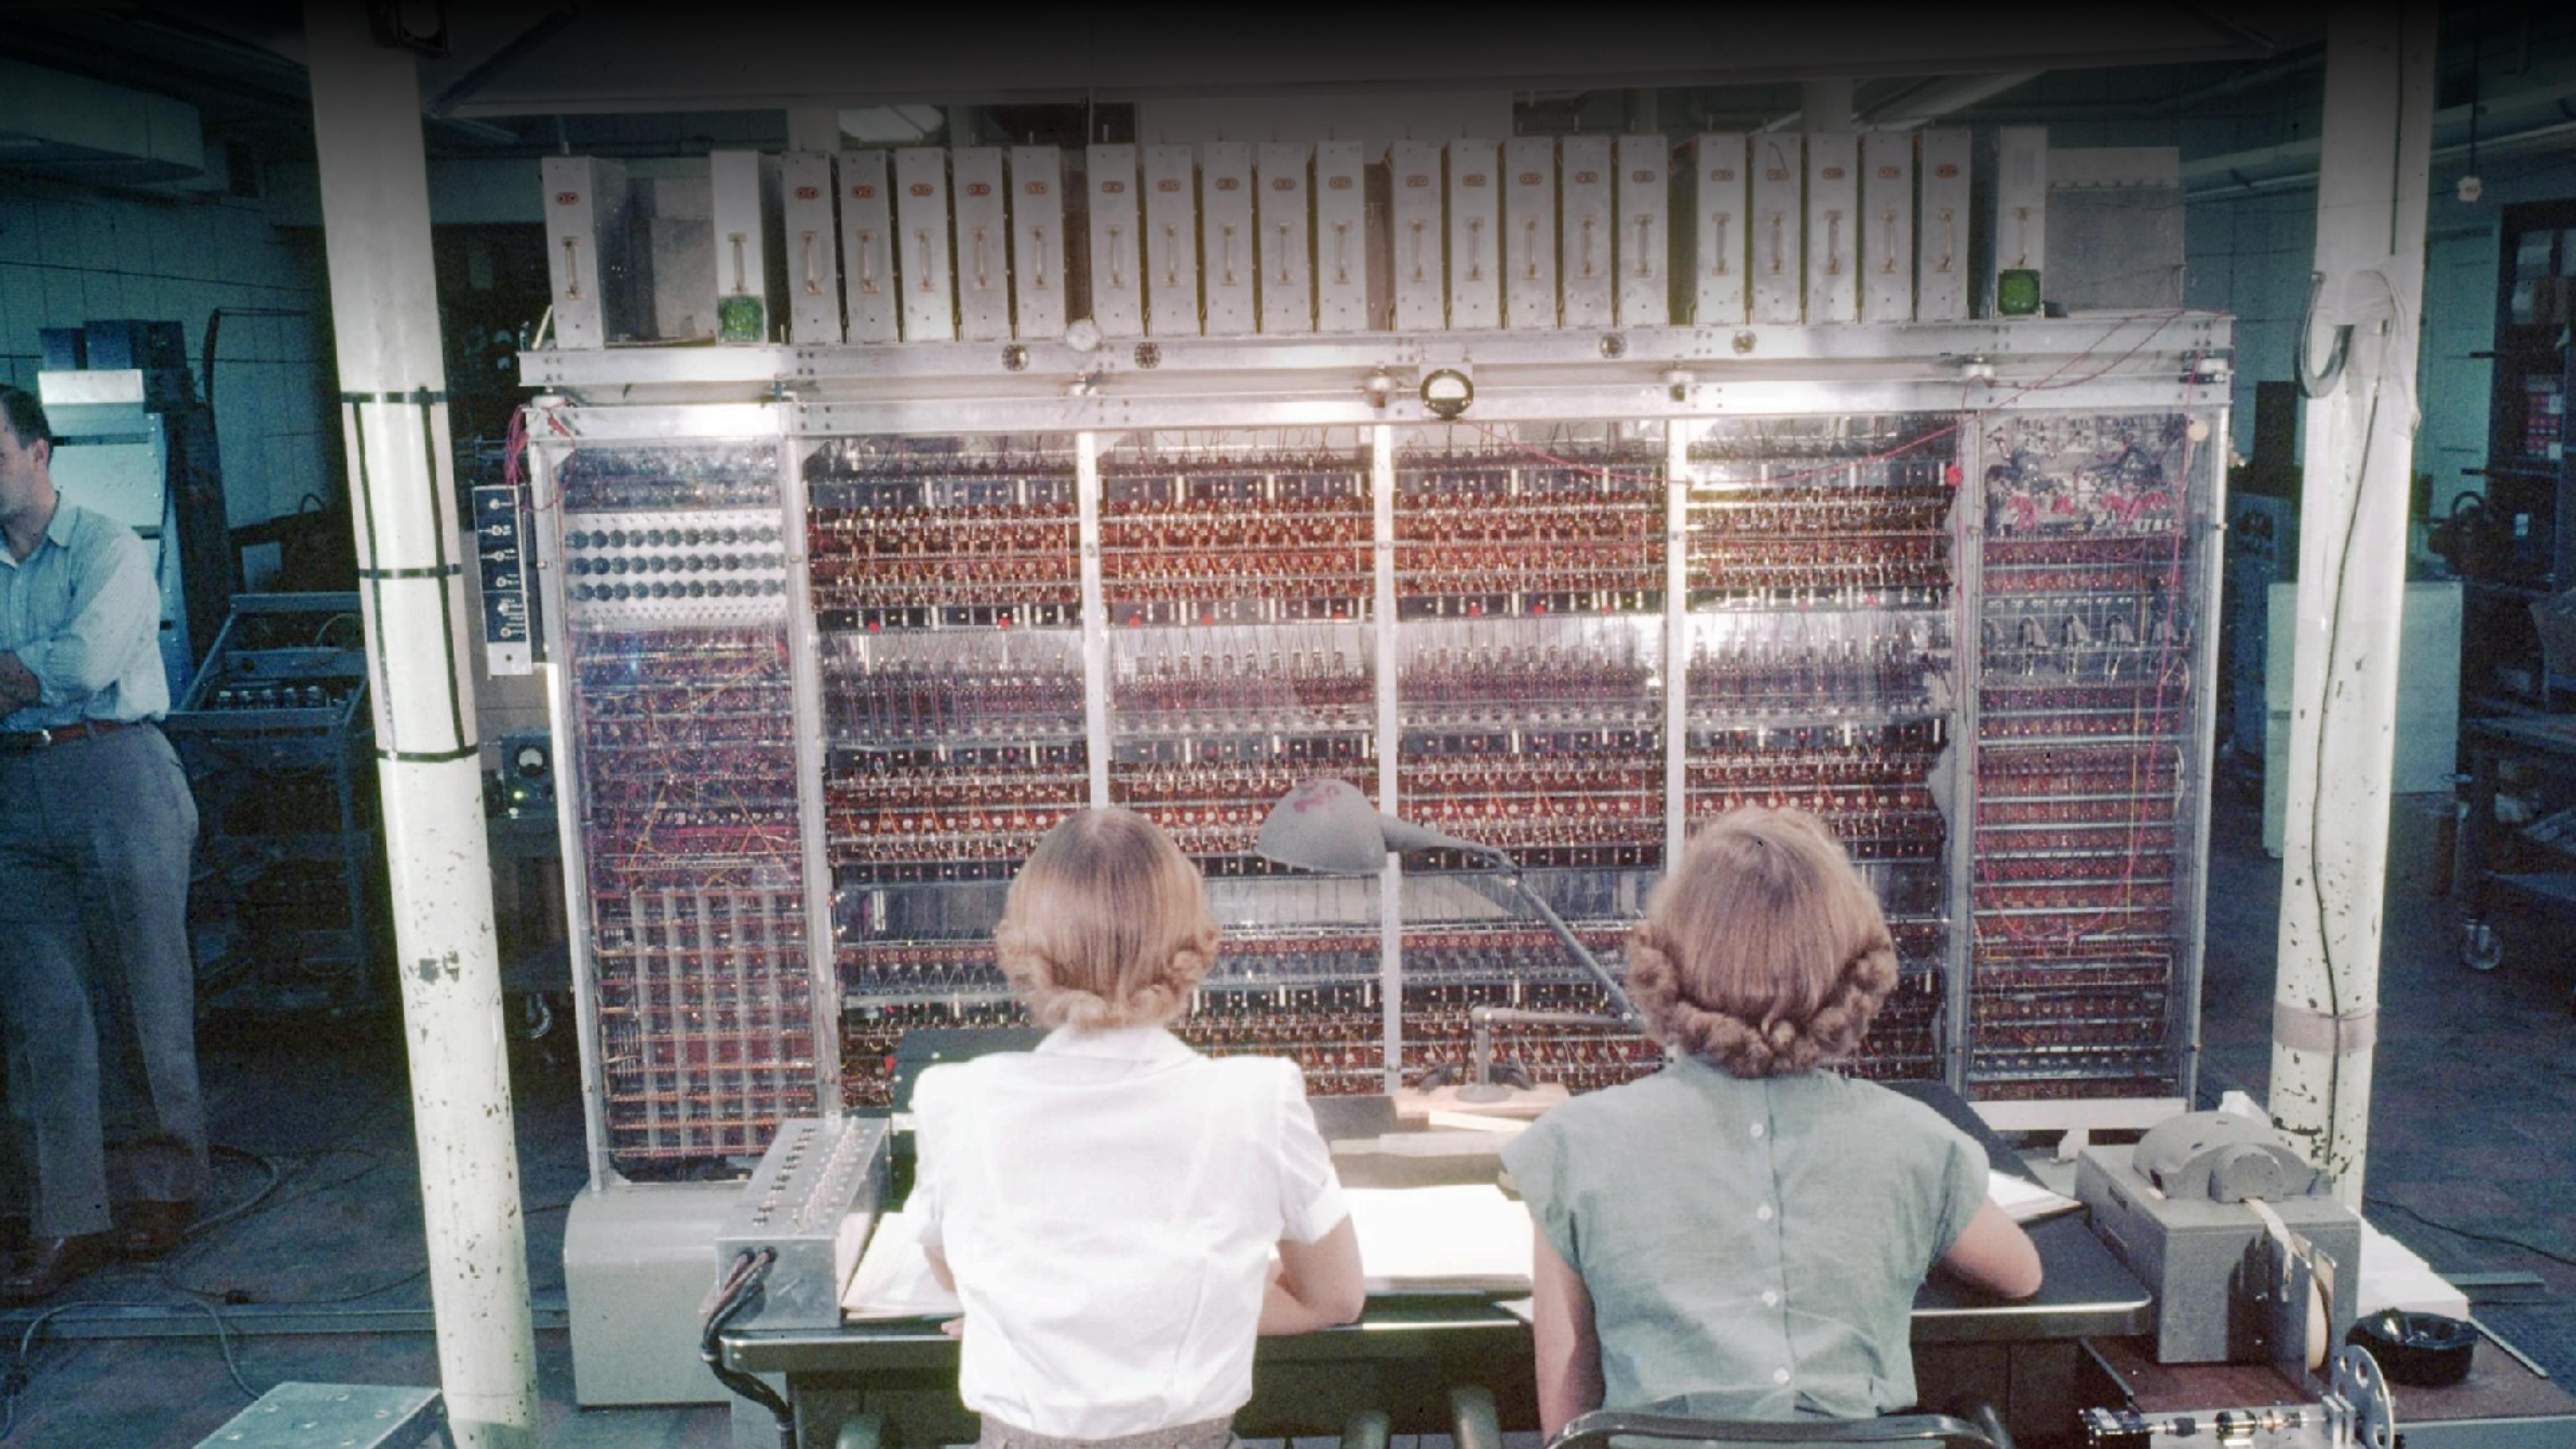


Whirling Arcs of Steel!

Touche! Miss Arianna Wright, state women's champion and city champ in the men's foil tournaments, makes a deep and effective lunge to score a point against Vernon Bordman, who wields a good blade himself.



Arianna Rosenbluth
(1927-2020)



Mathematical Analyzer, Numerical Integrator, and Computer

MANIAC:
1000 pounds
5 kilobytes of memory
70k multiplications/sec

Your laptop:
4-7 pounds
8+ million kilobytes memory
Billions of multiplications/sec

MCMC is diverse

Metropolis has yielded to newer, more efficient algorithms

Many innovations in the last decades

Best methods use *gradients*

We'll use Hamiltonian Monte Carlo

Chapman & Hall/CRC
Handbooks of Modern
Statistical Methods

Handbook of Markov Chain Monte Carlo

Edited by

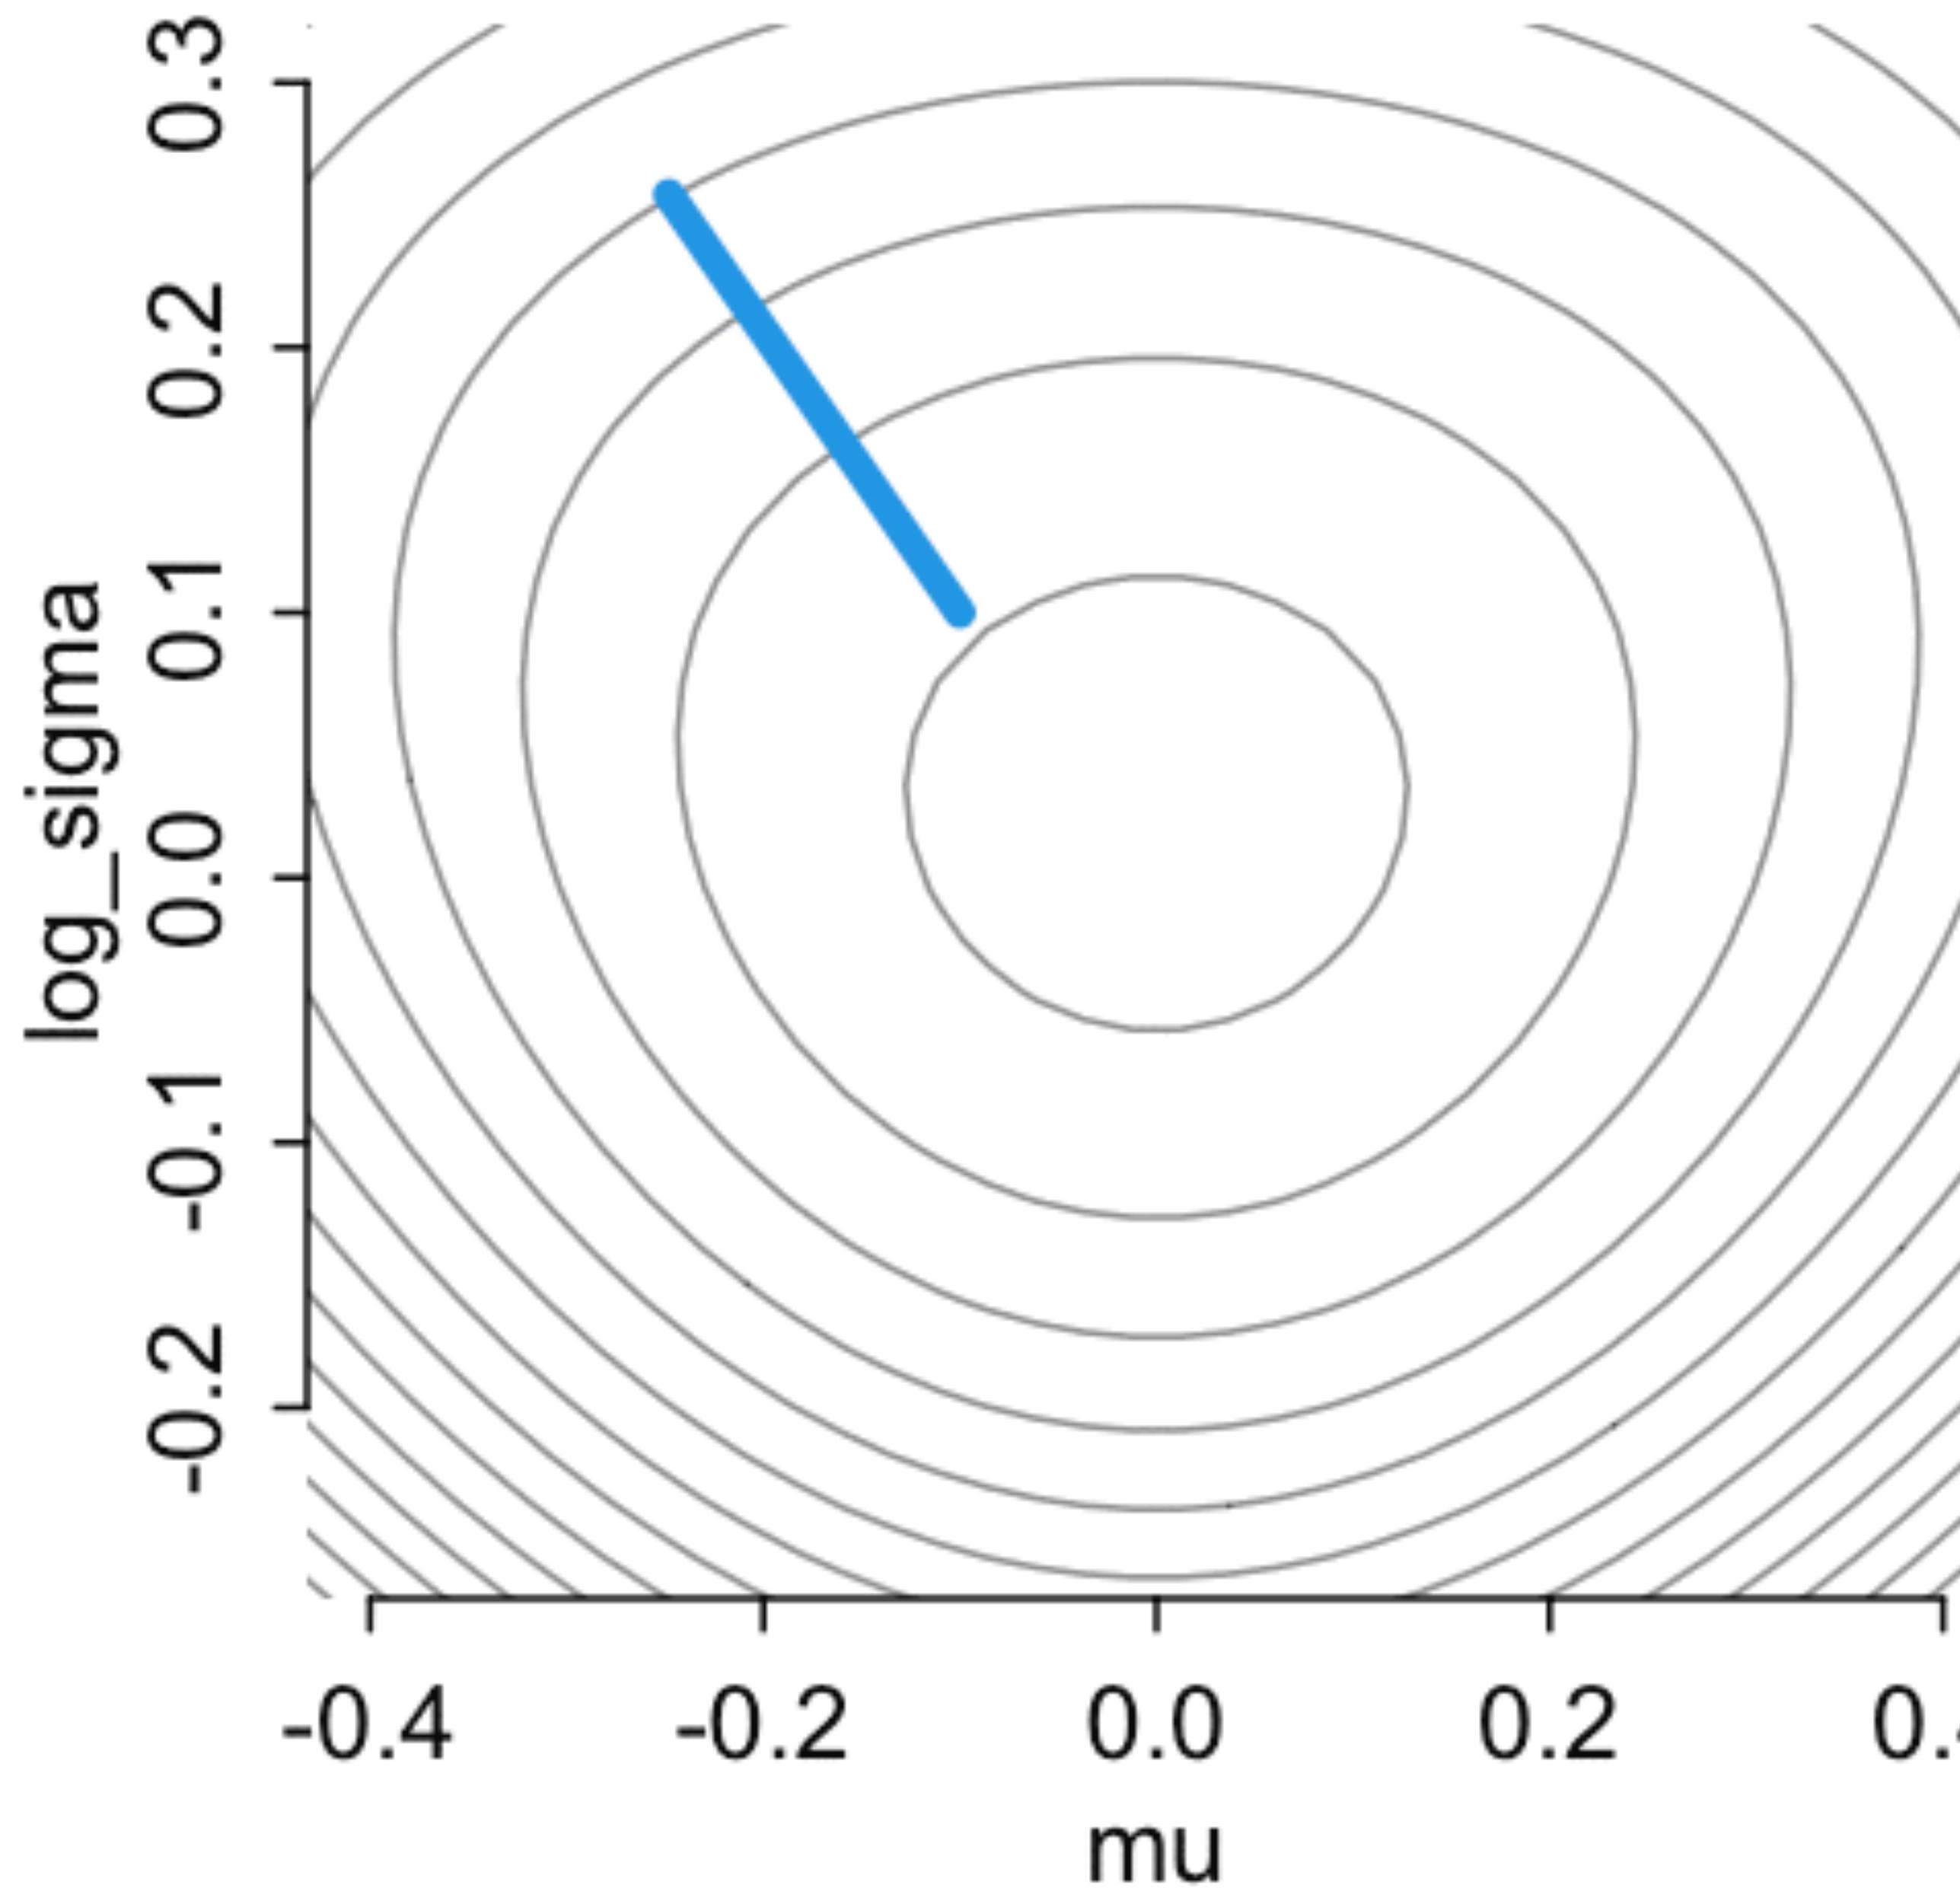
Steve Brooks

Andrew Gelman

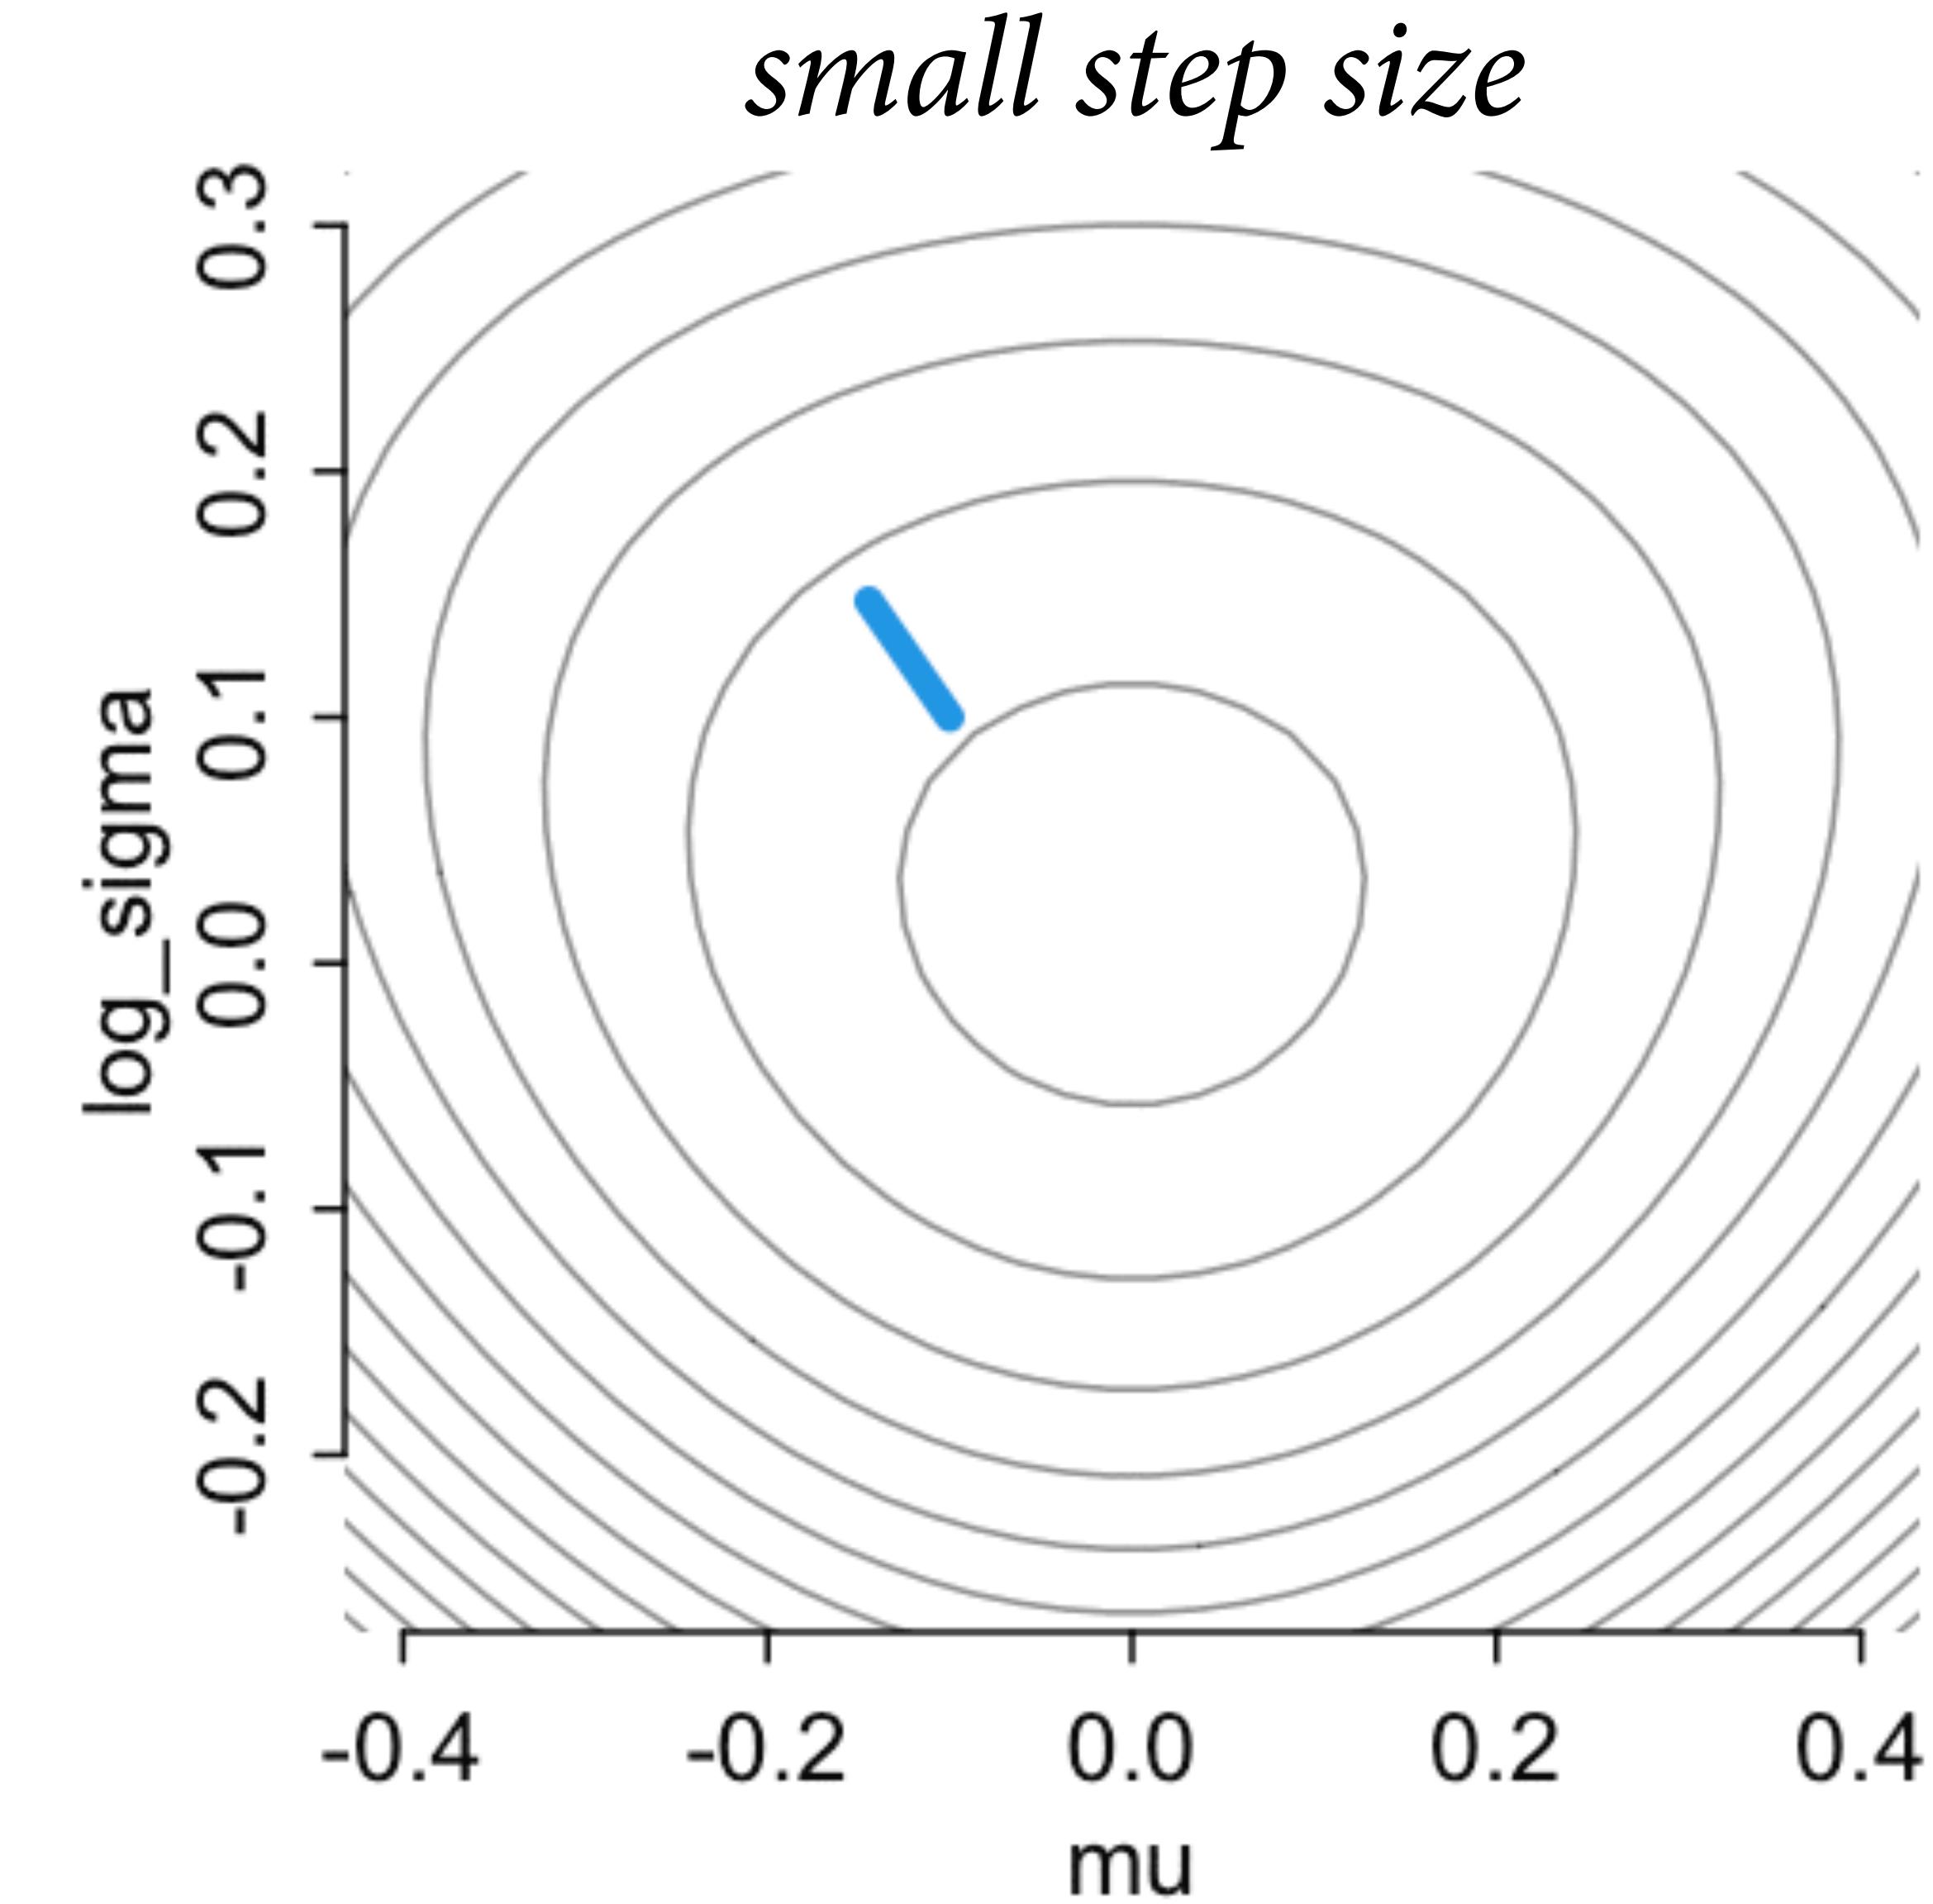
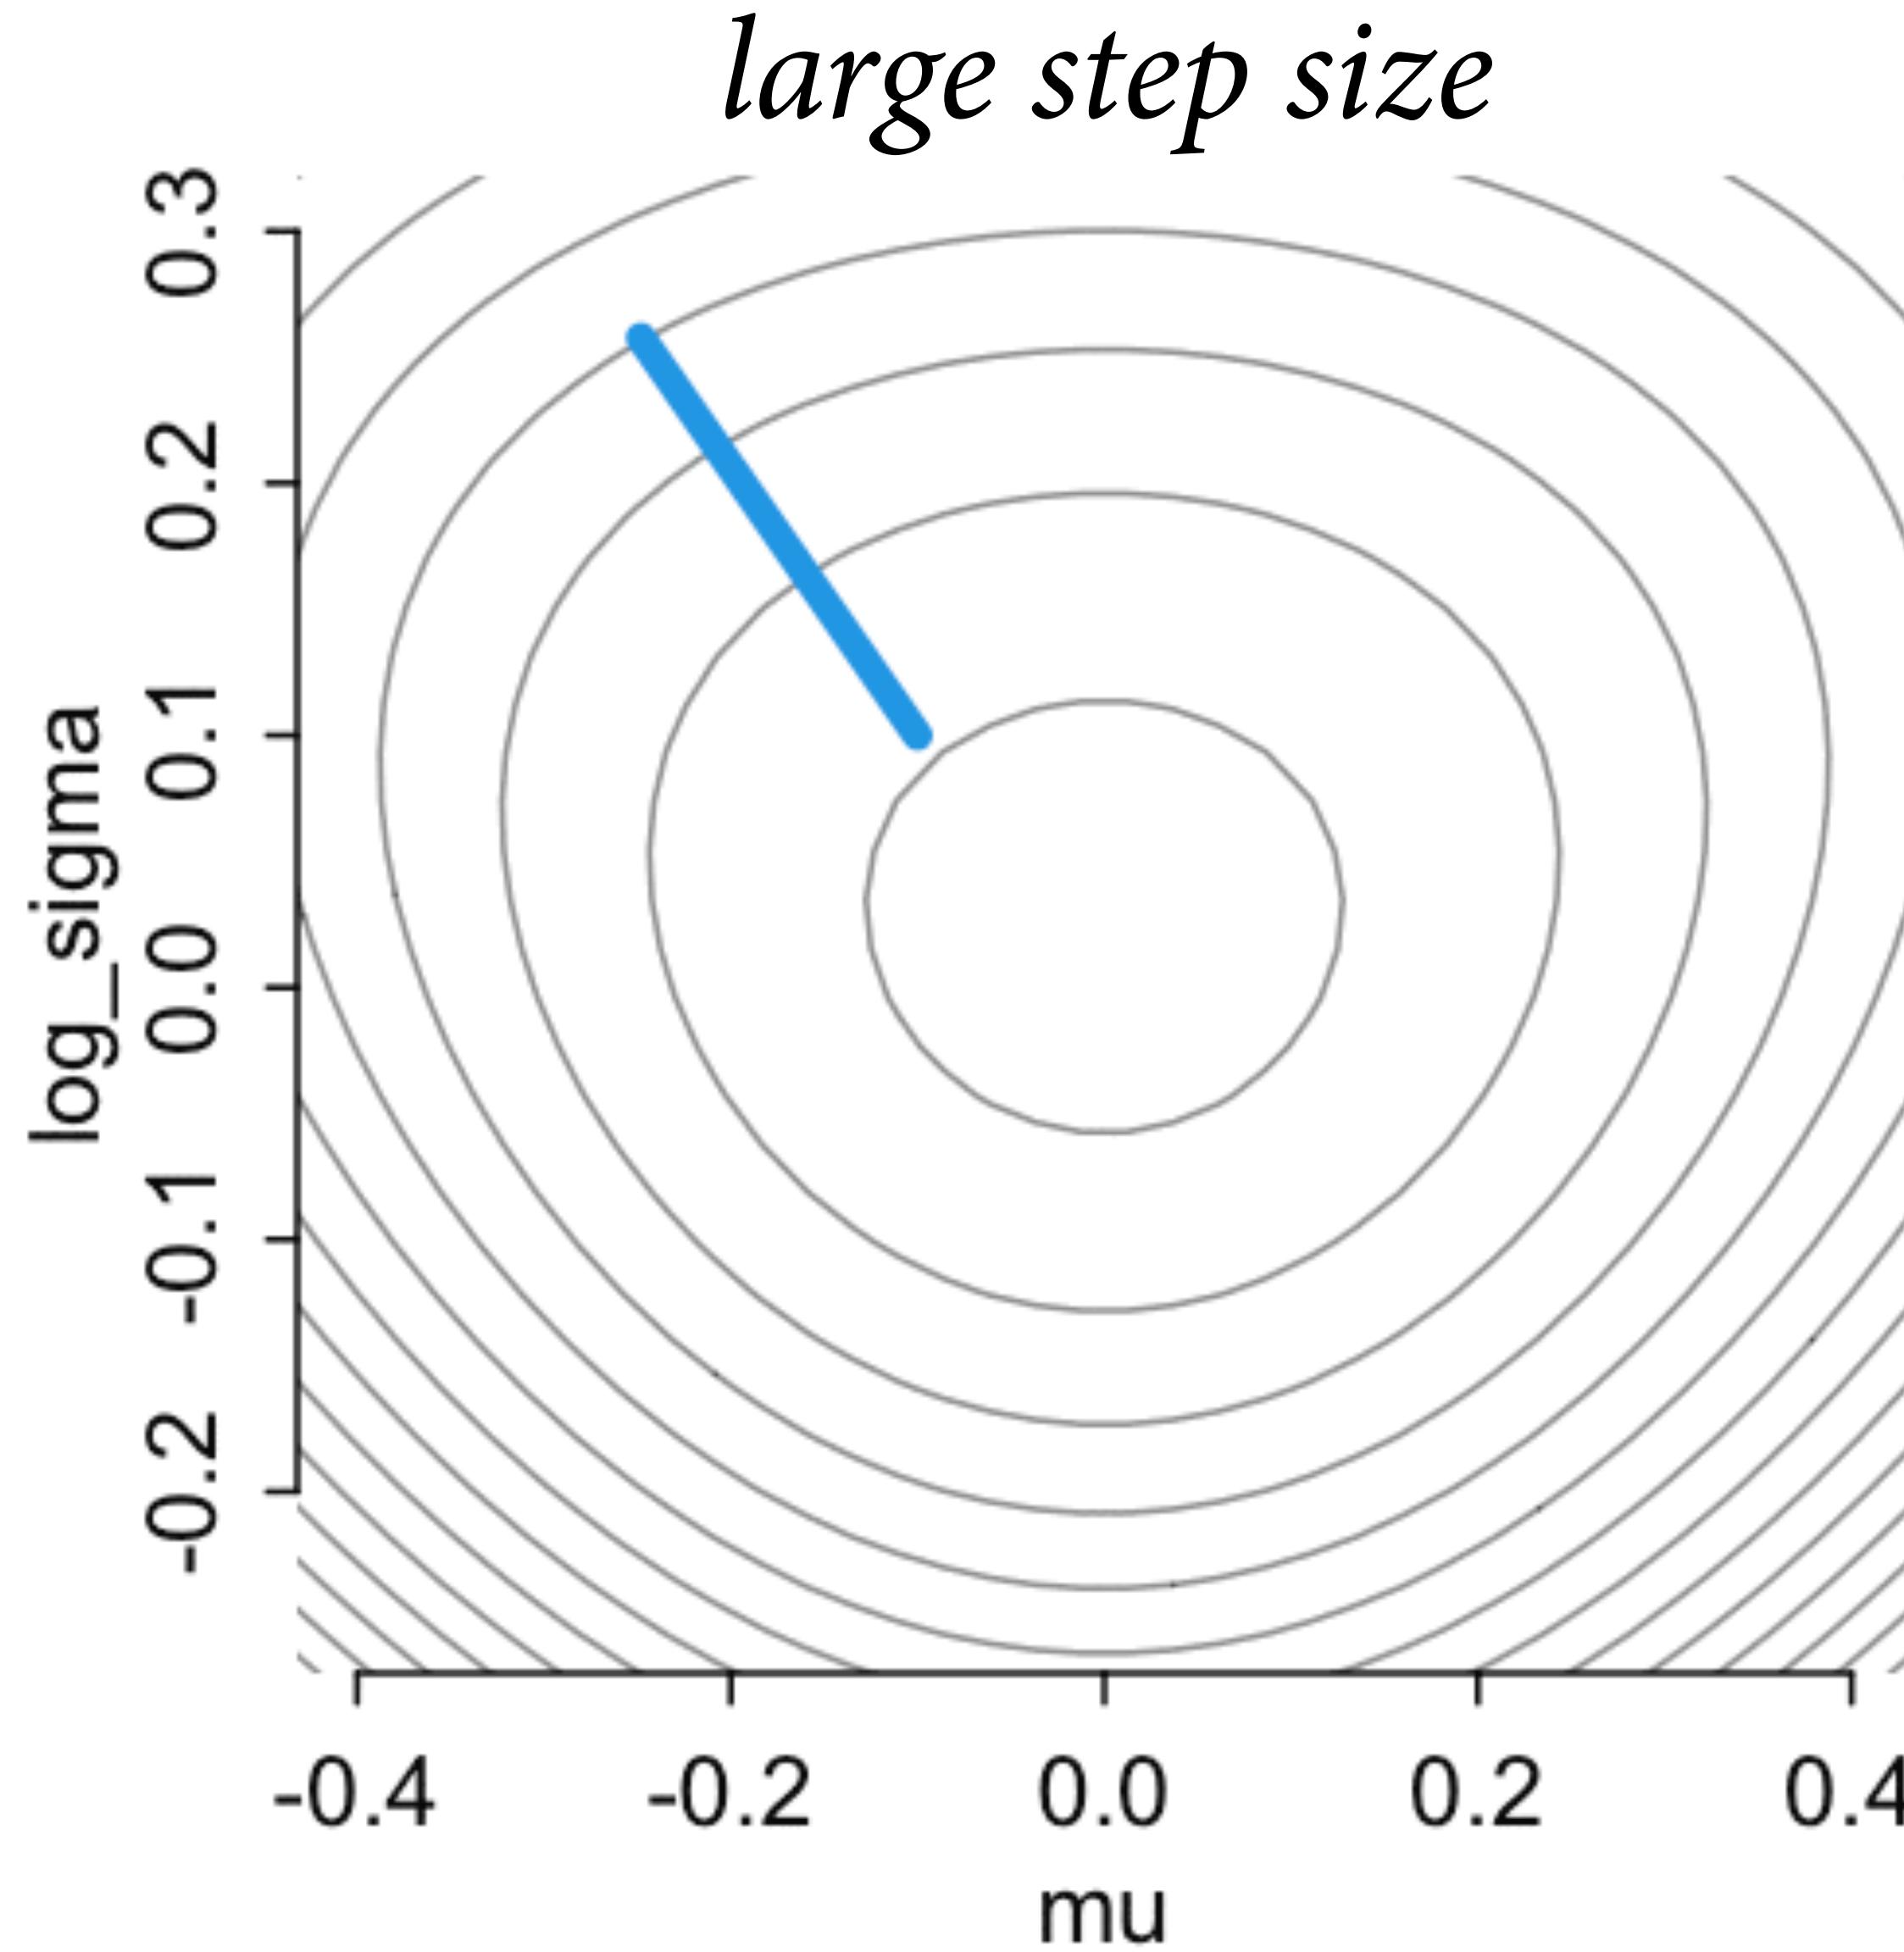
Galin L. Jones

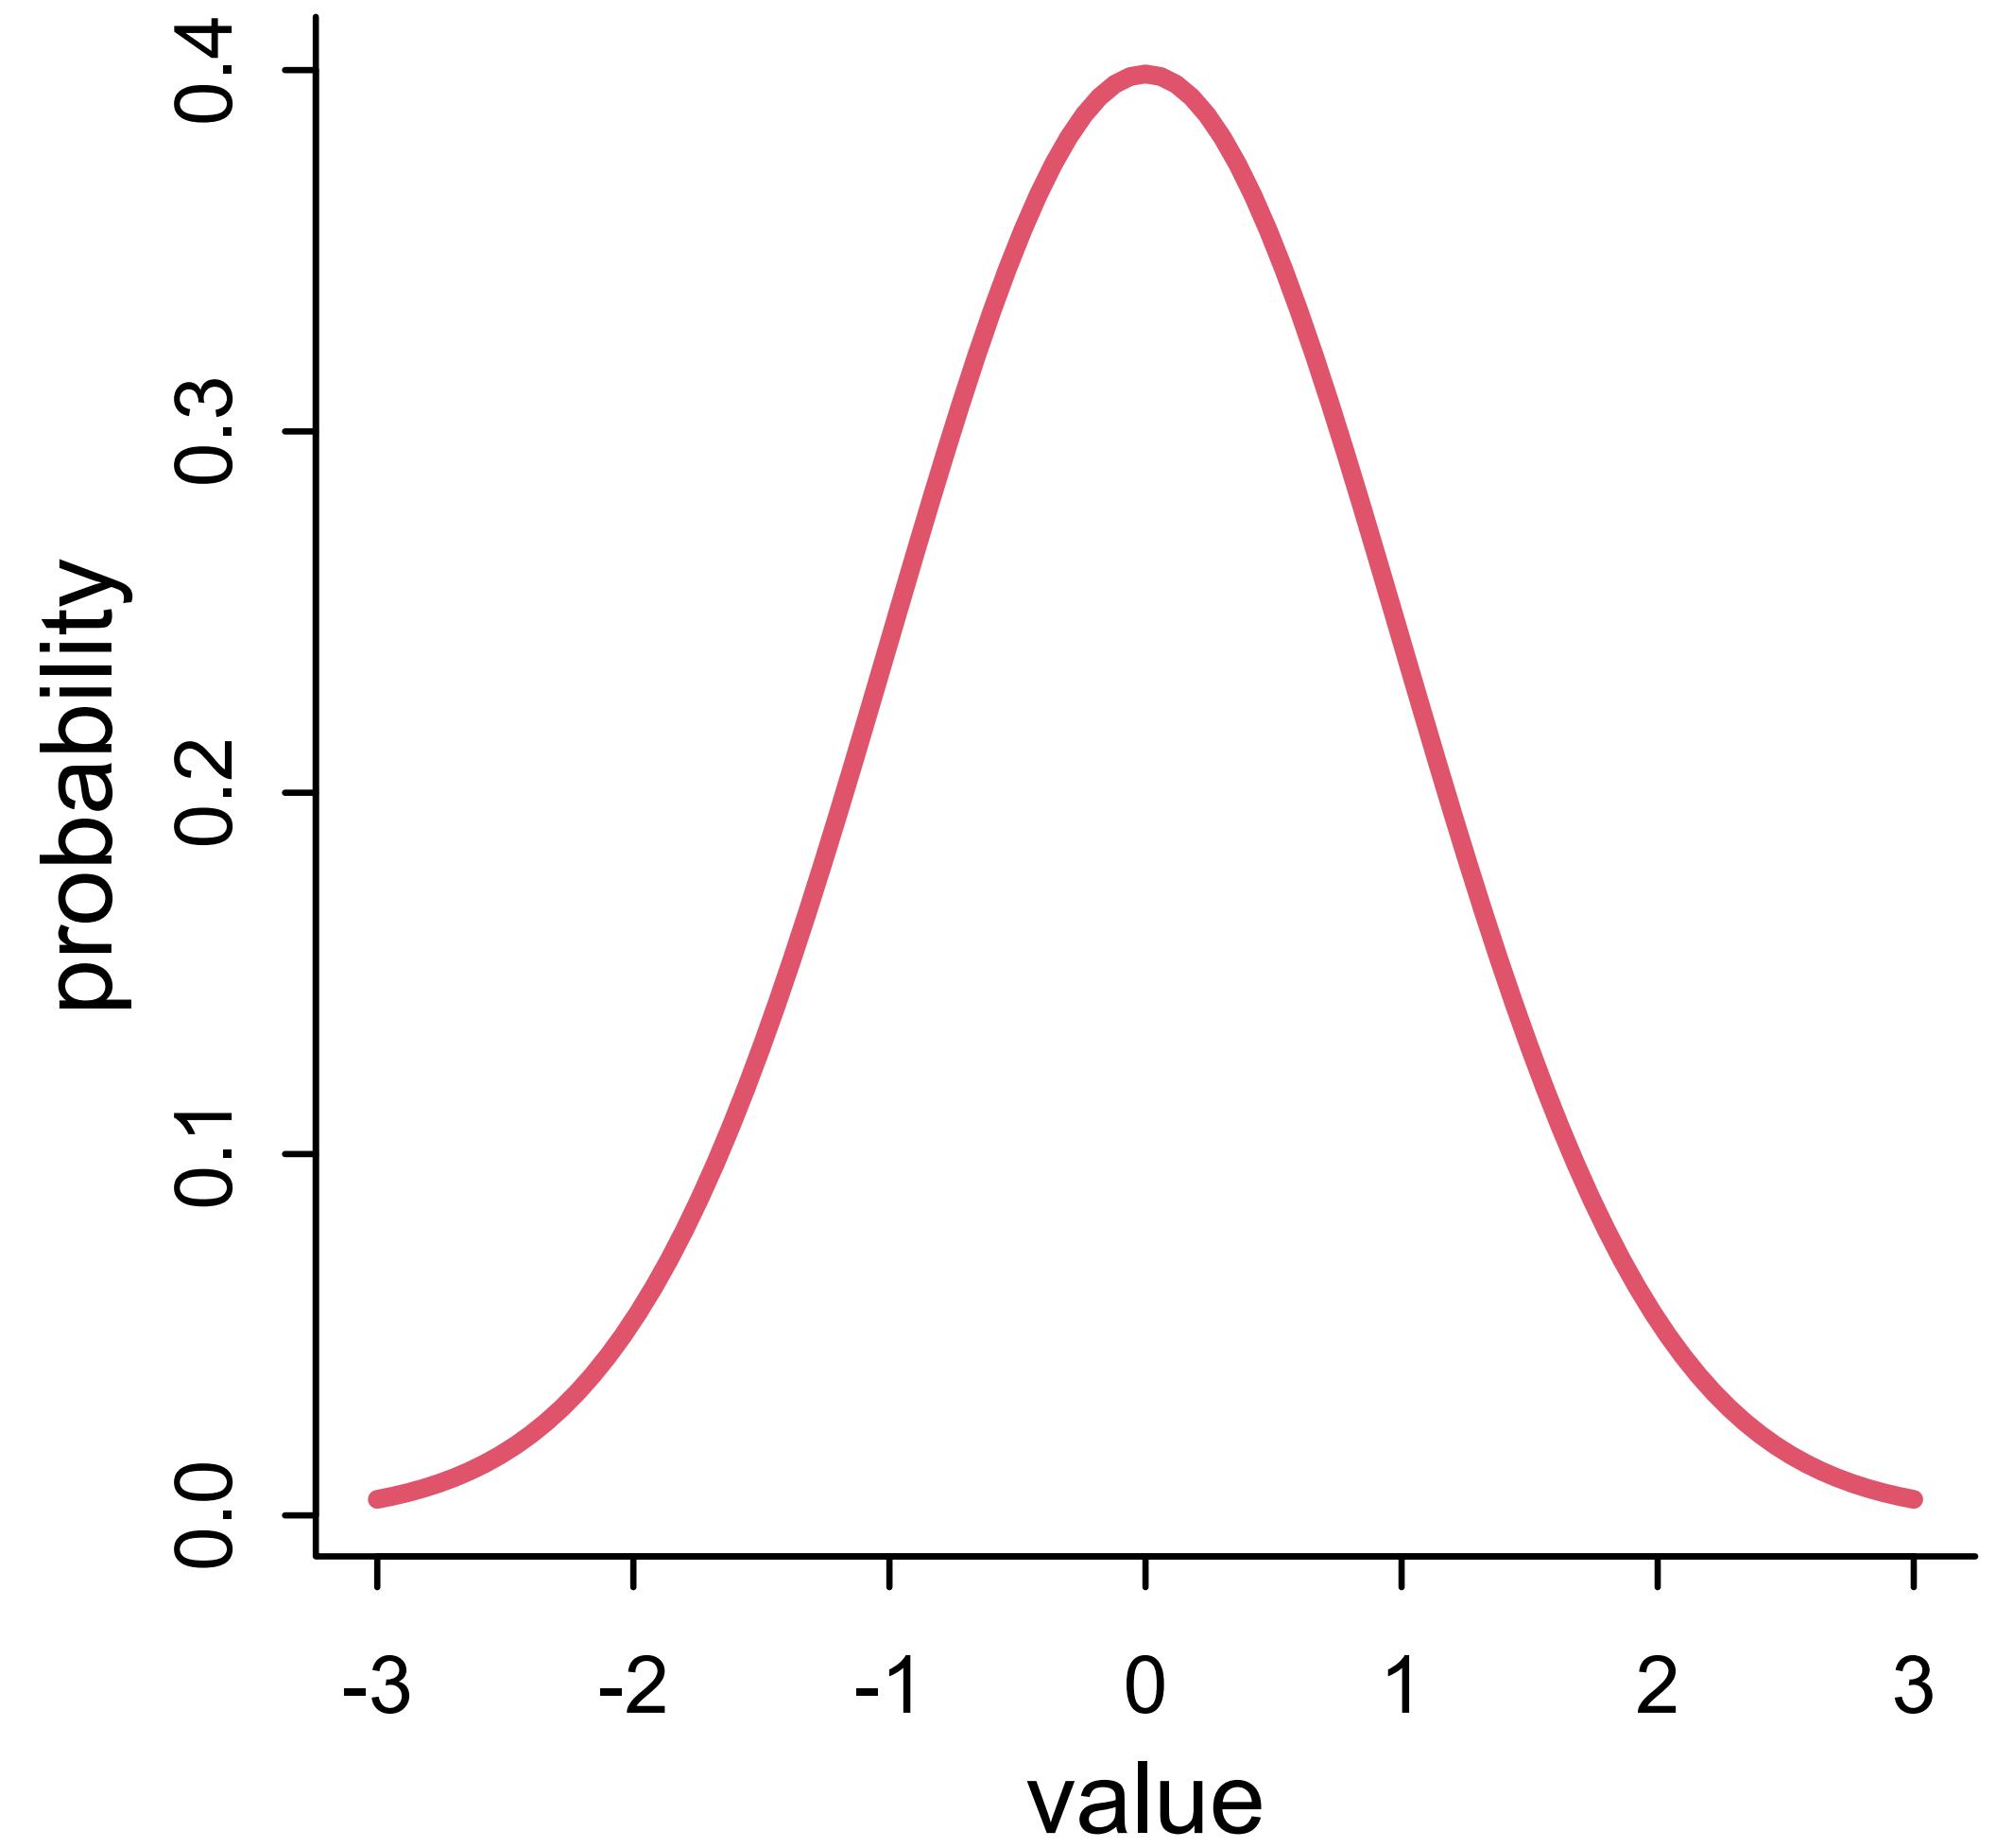
Xiao-Li Meng

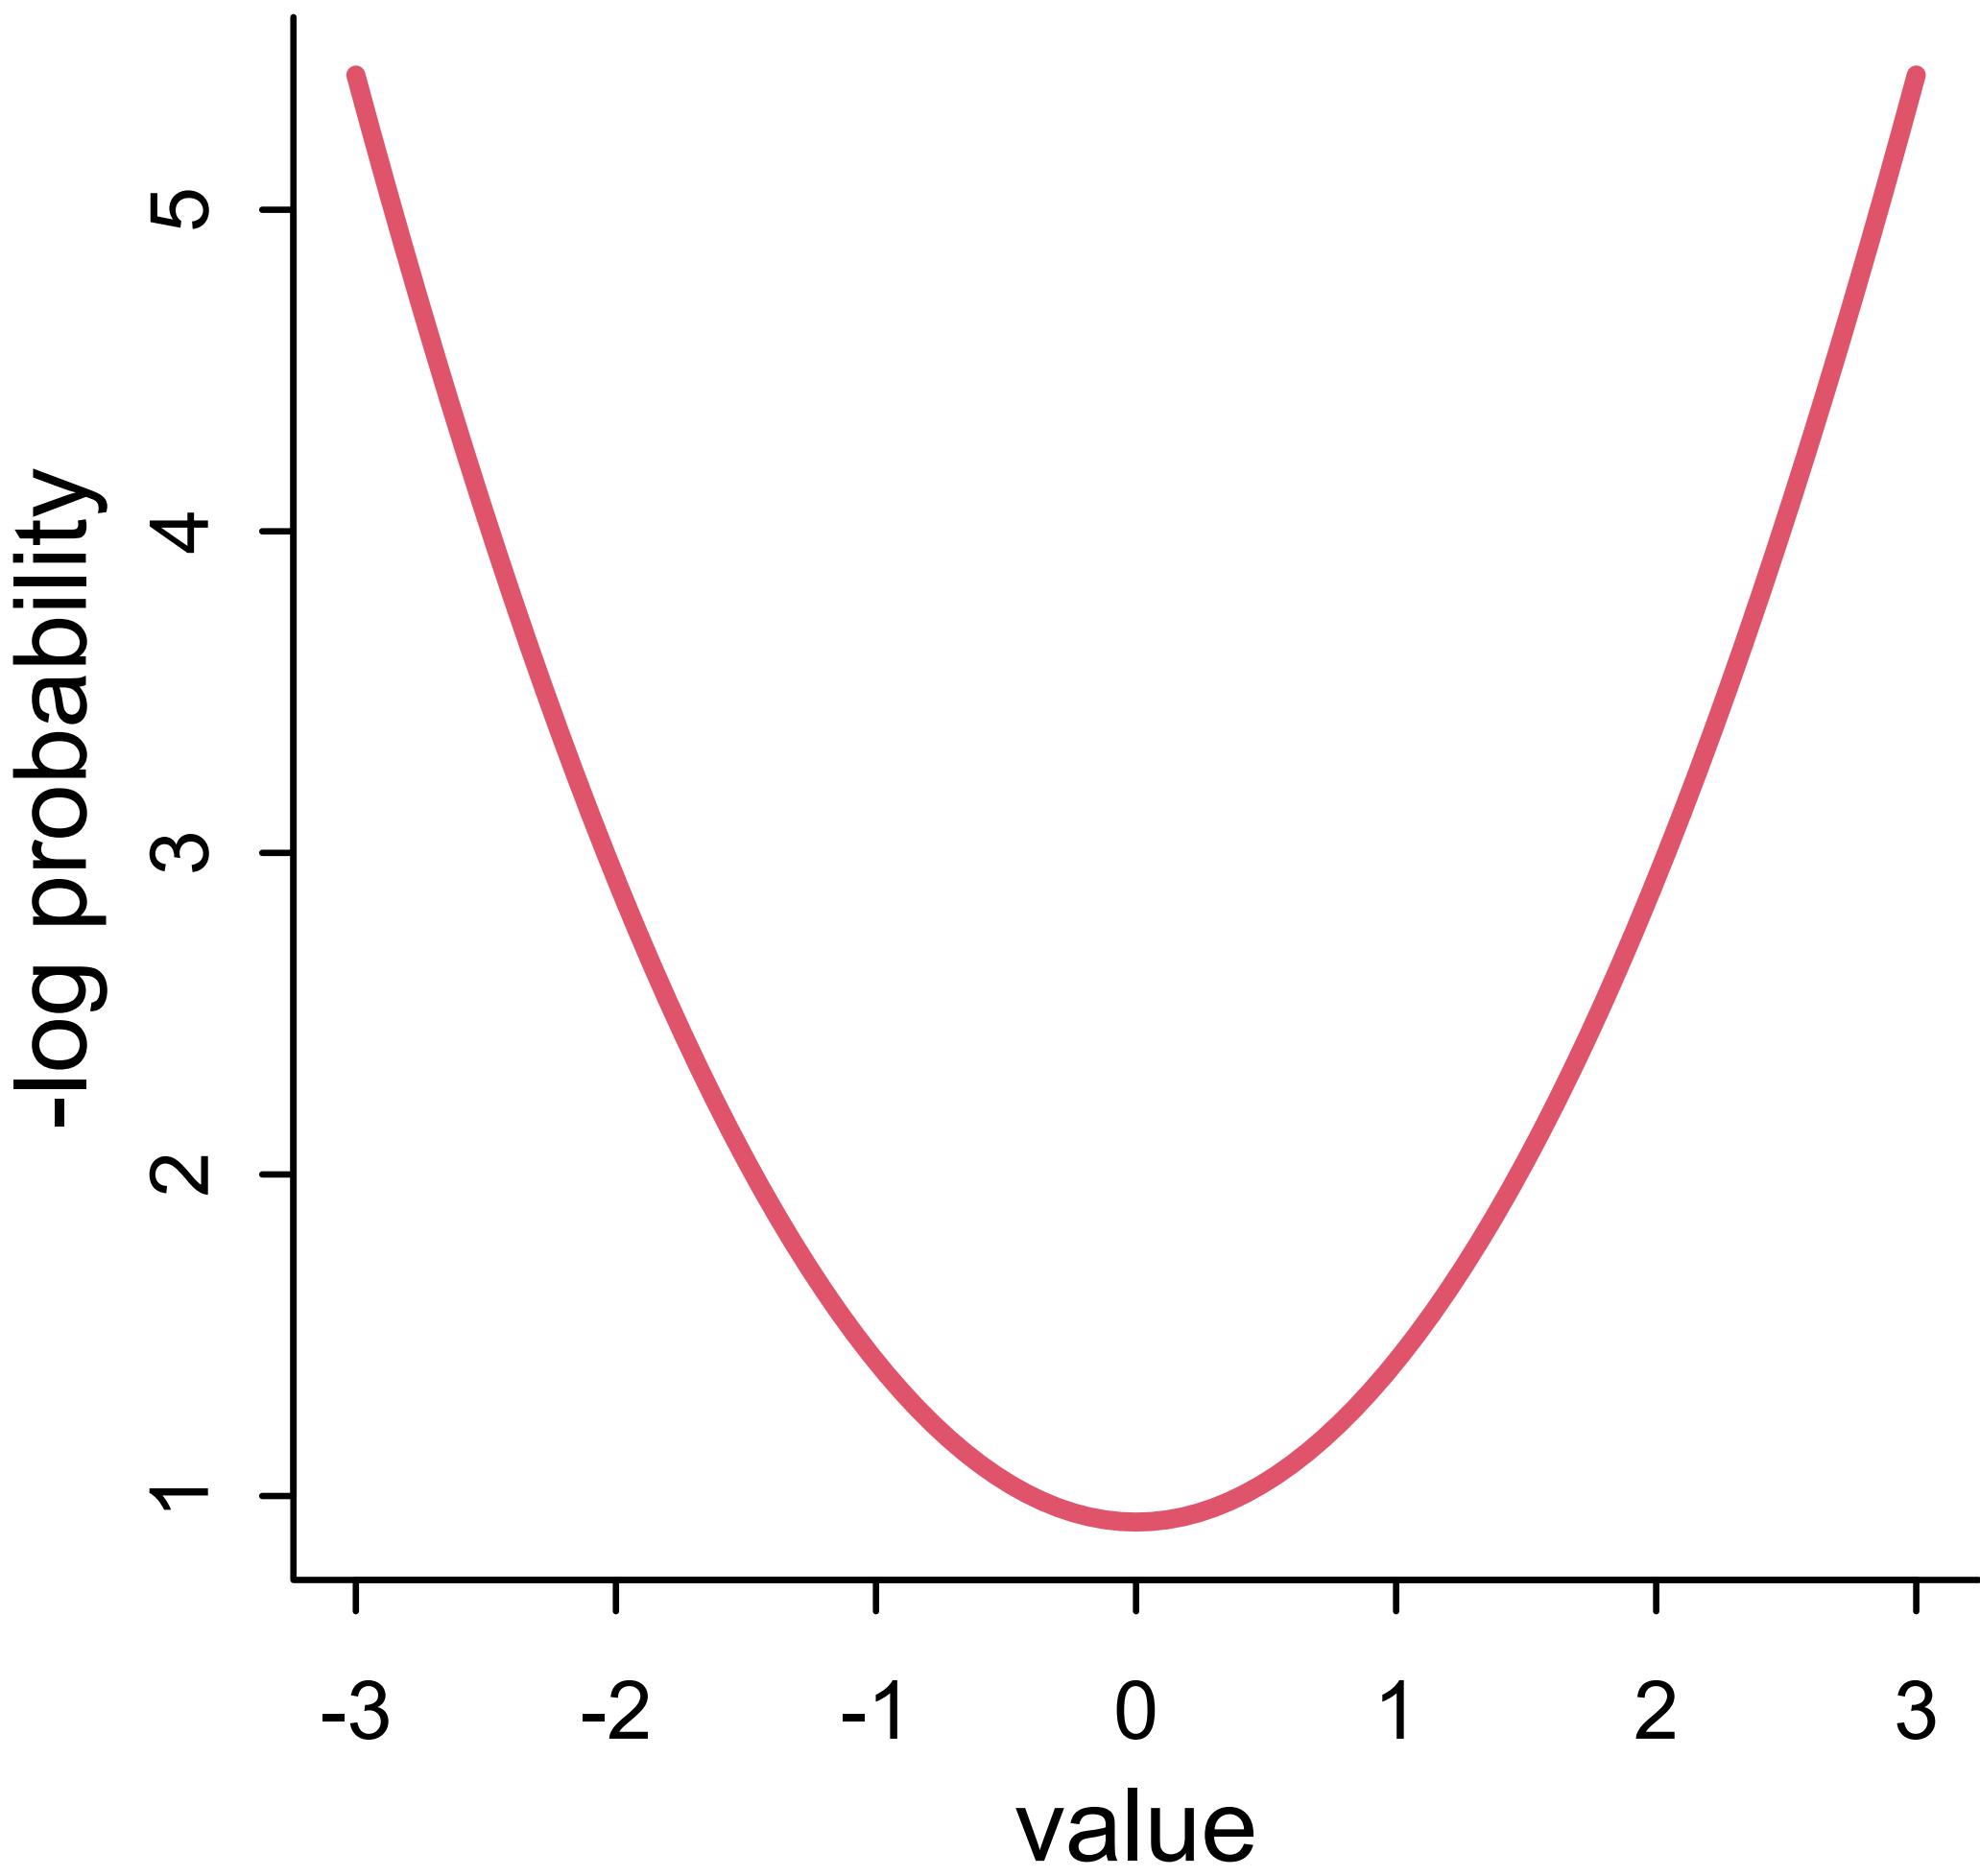
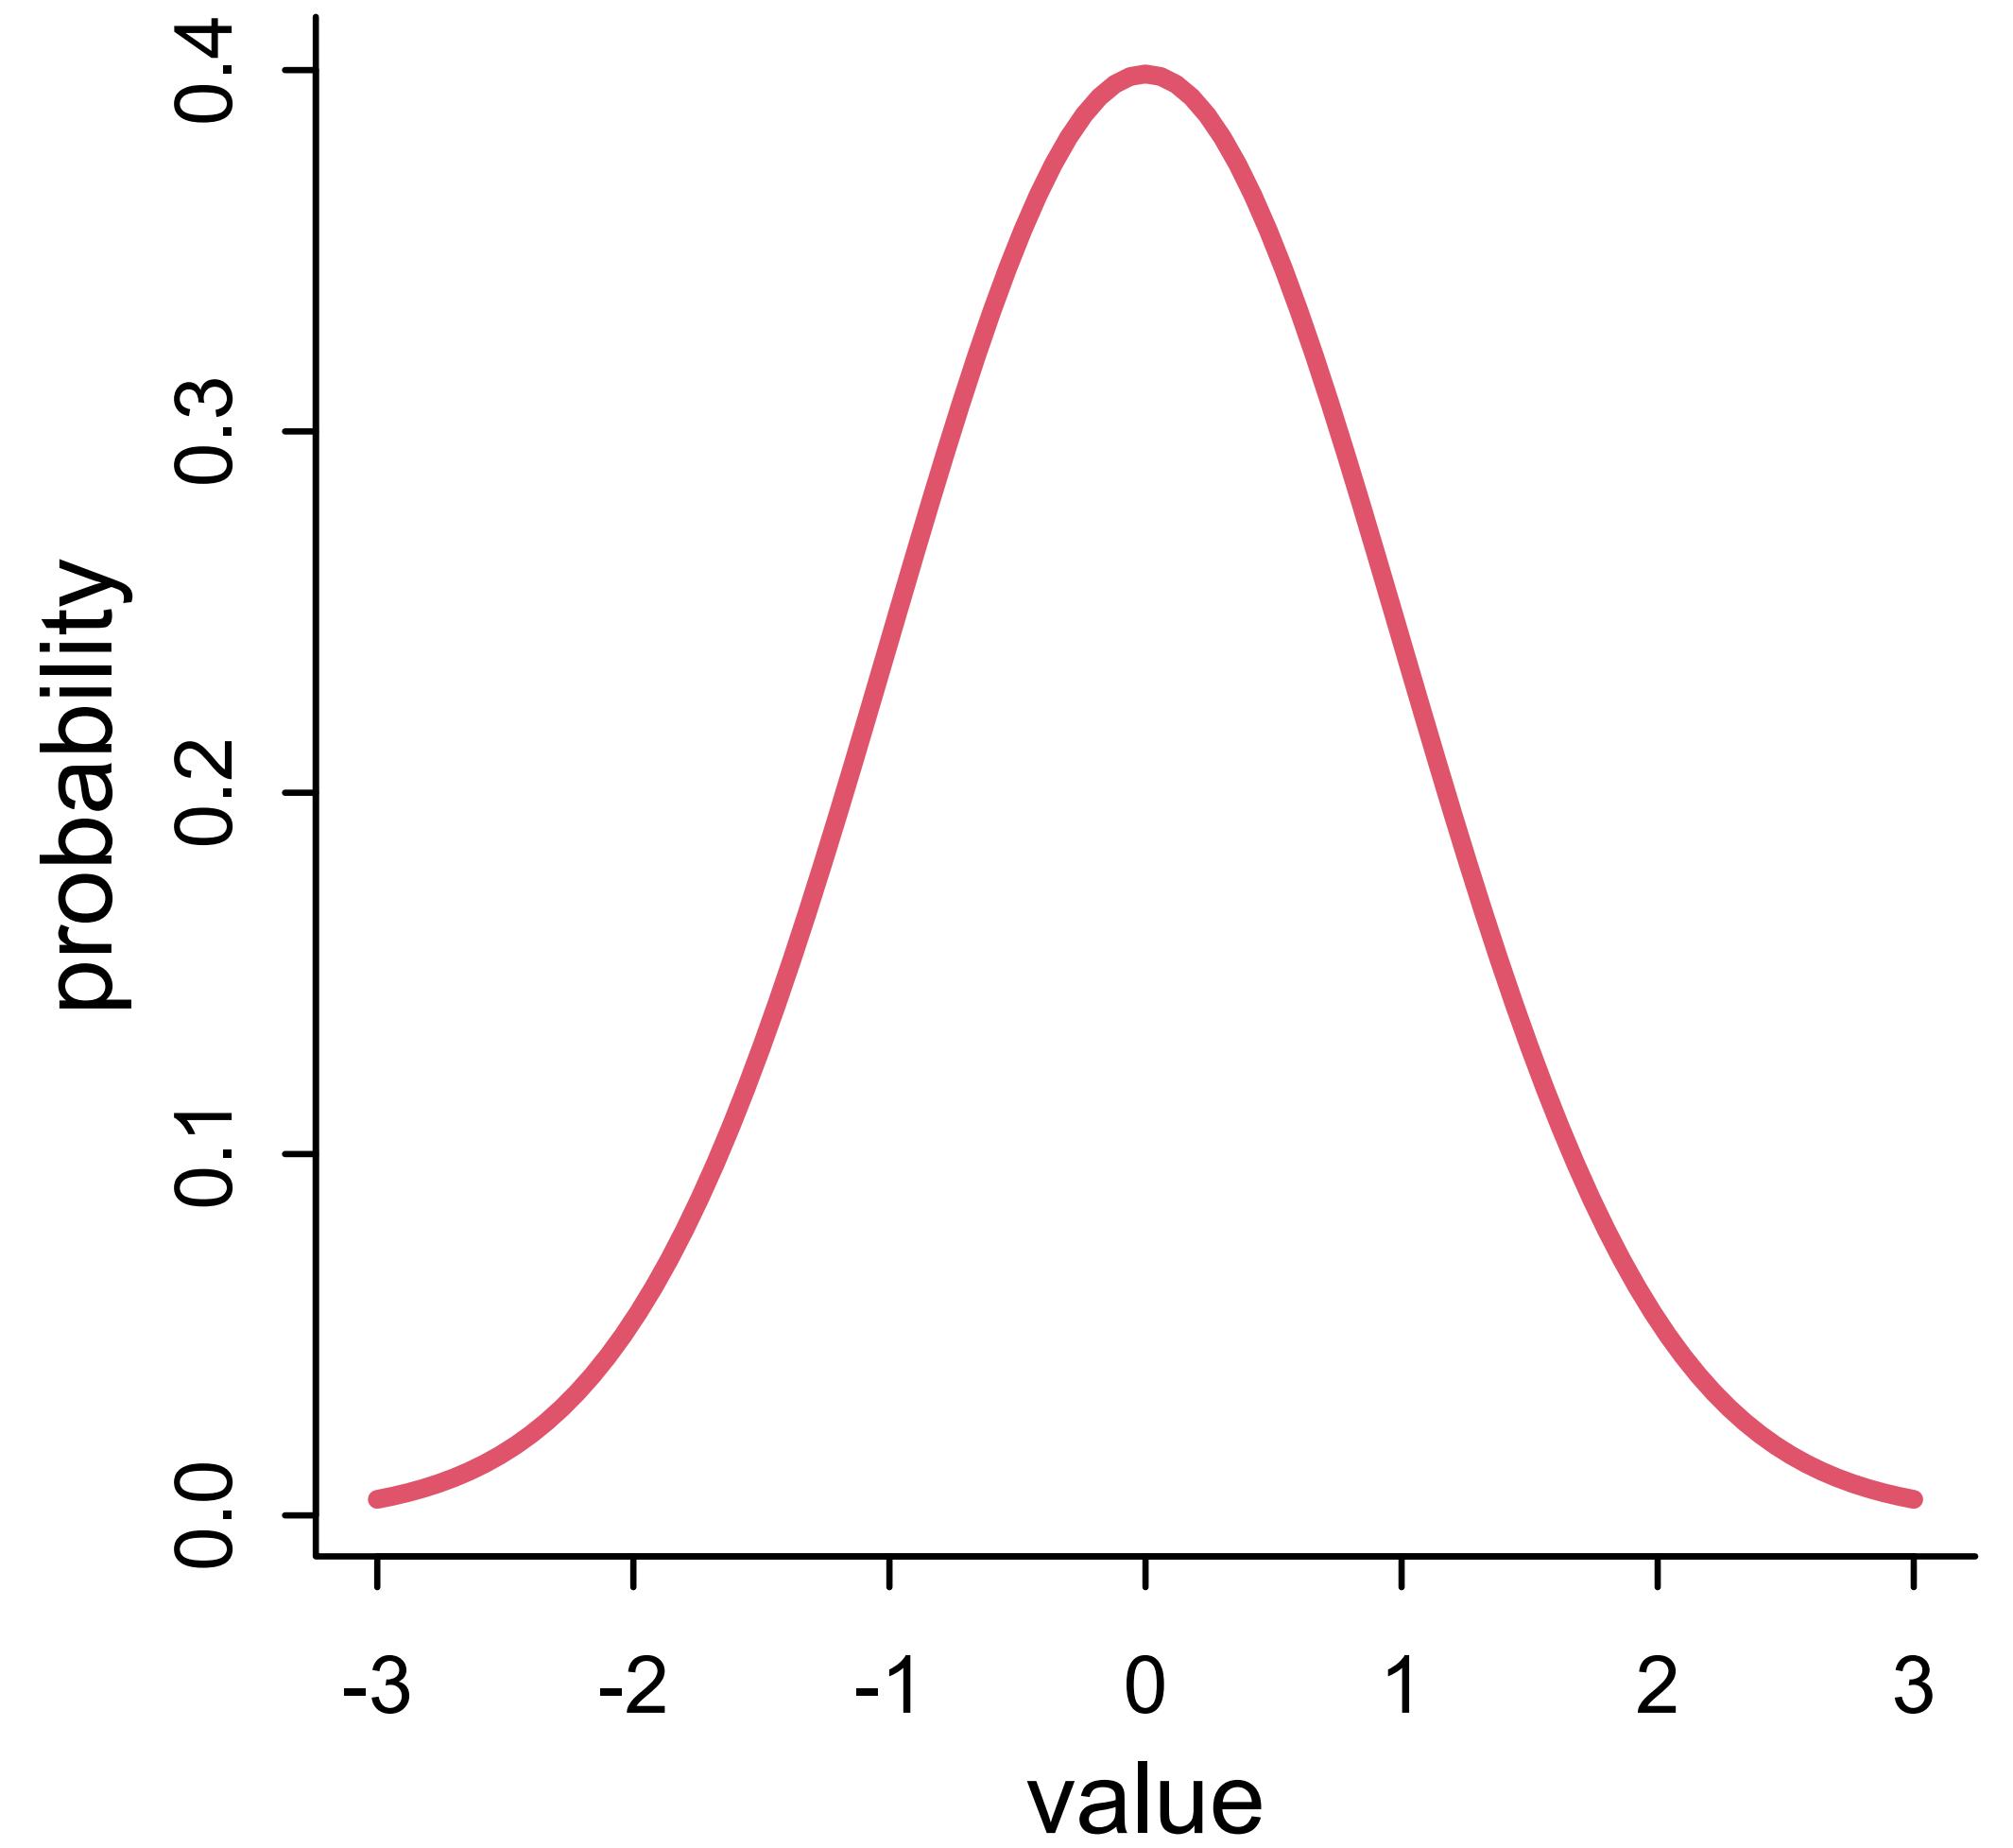
Basic Rosenbluth (aka Metropolis) algorithm

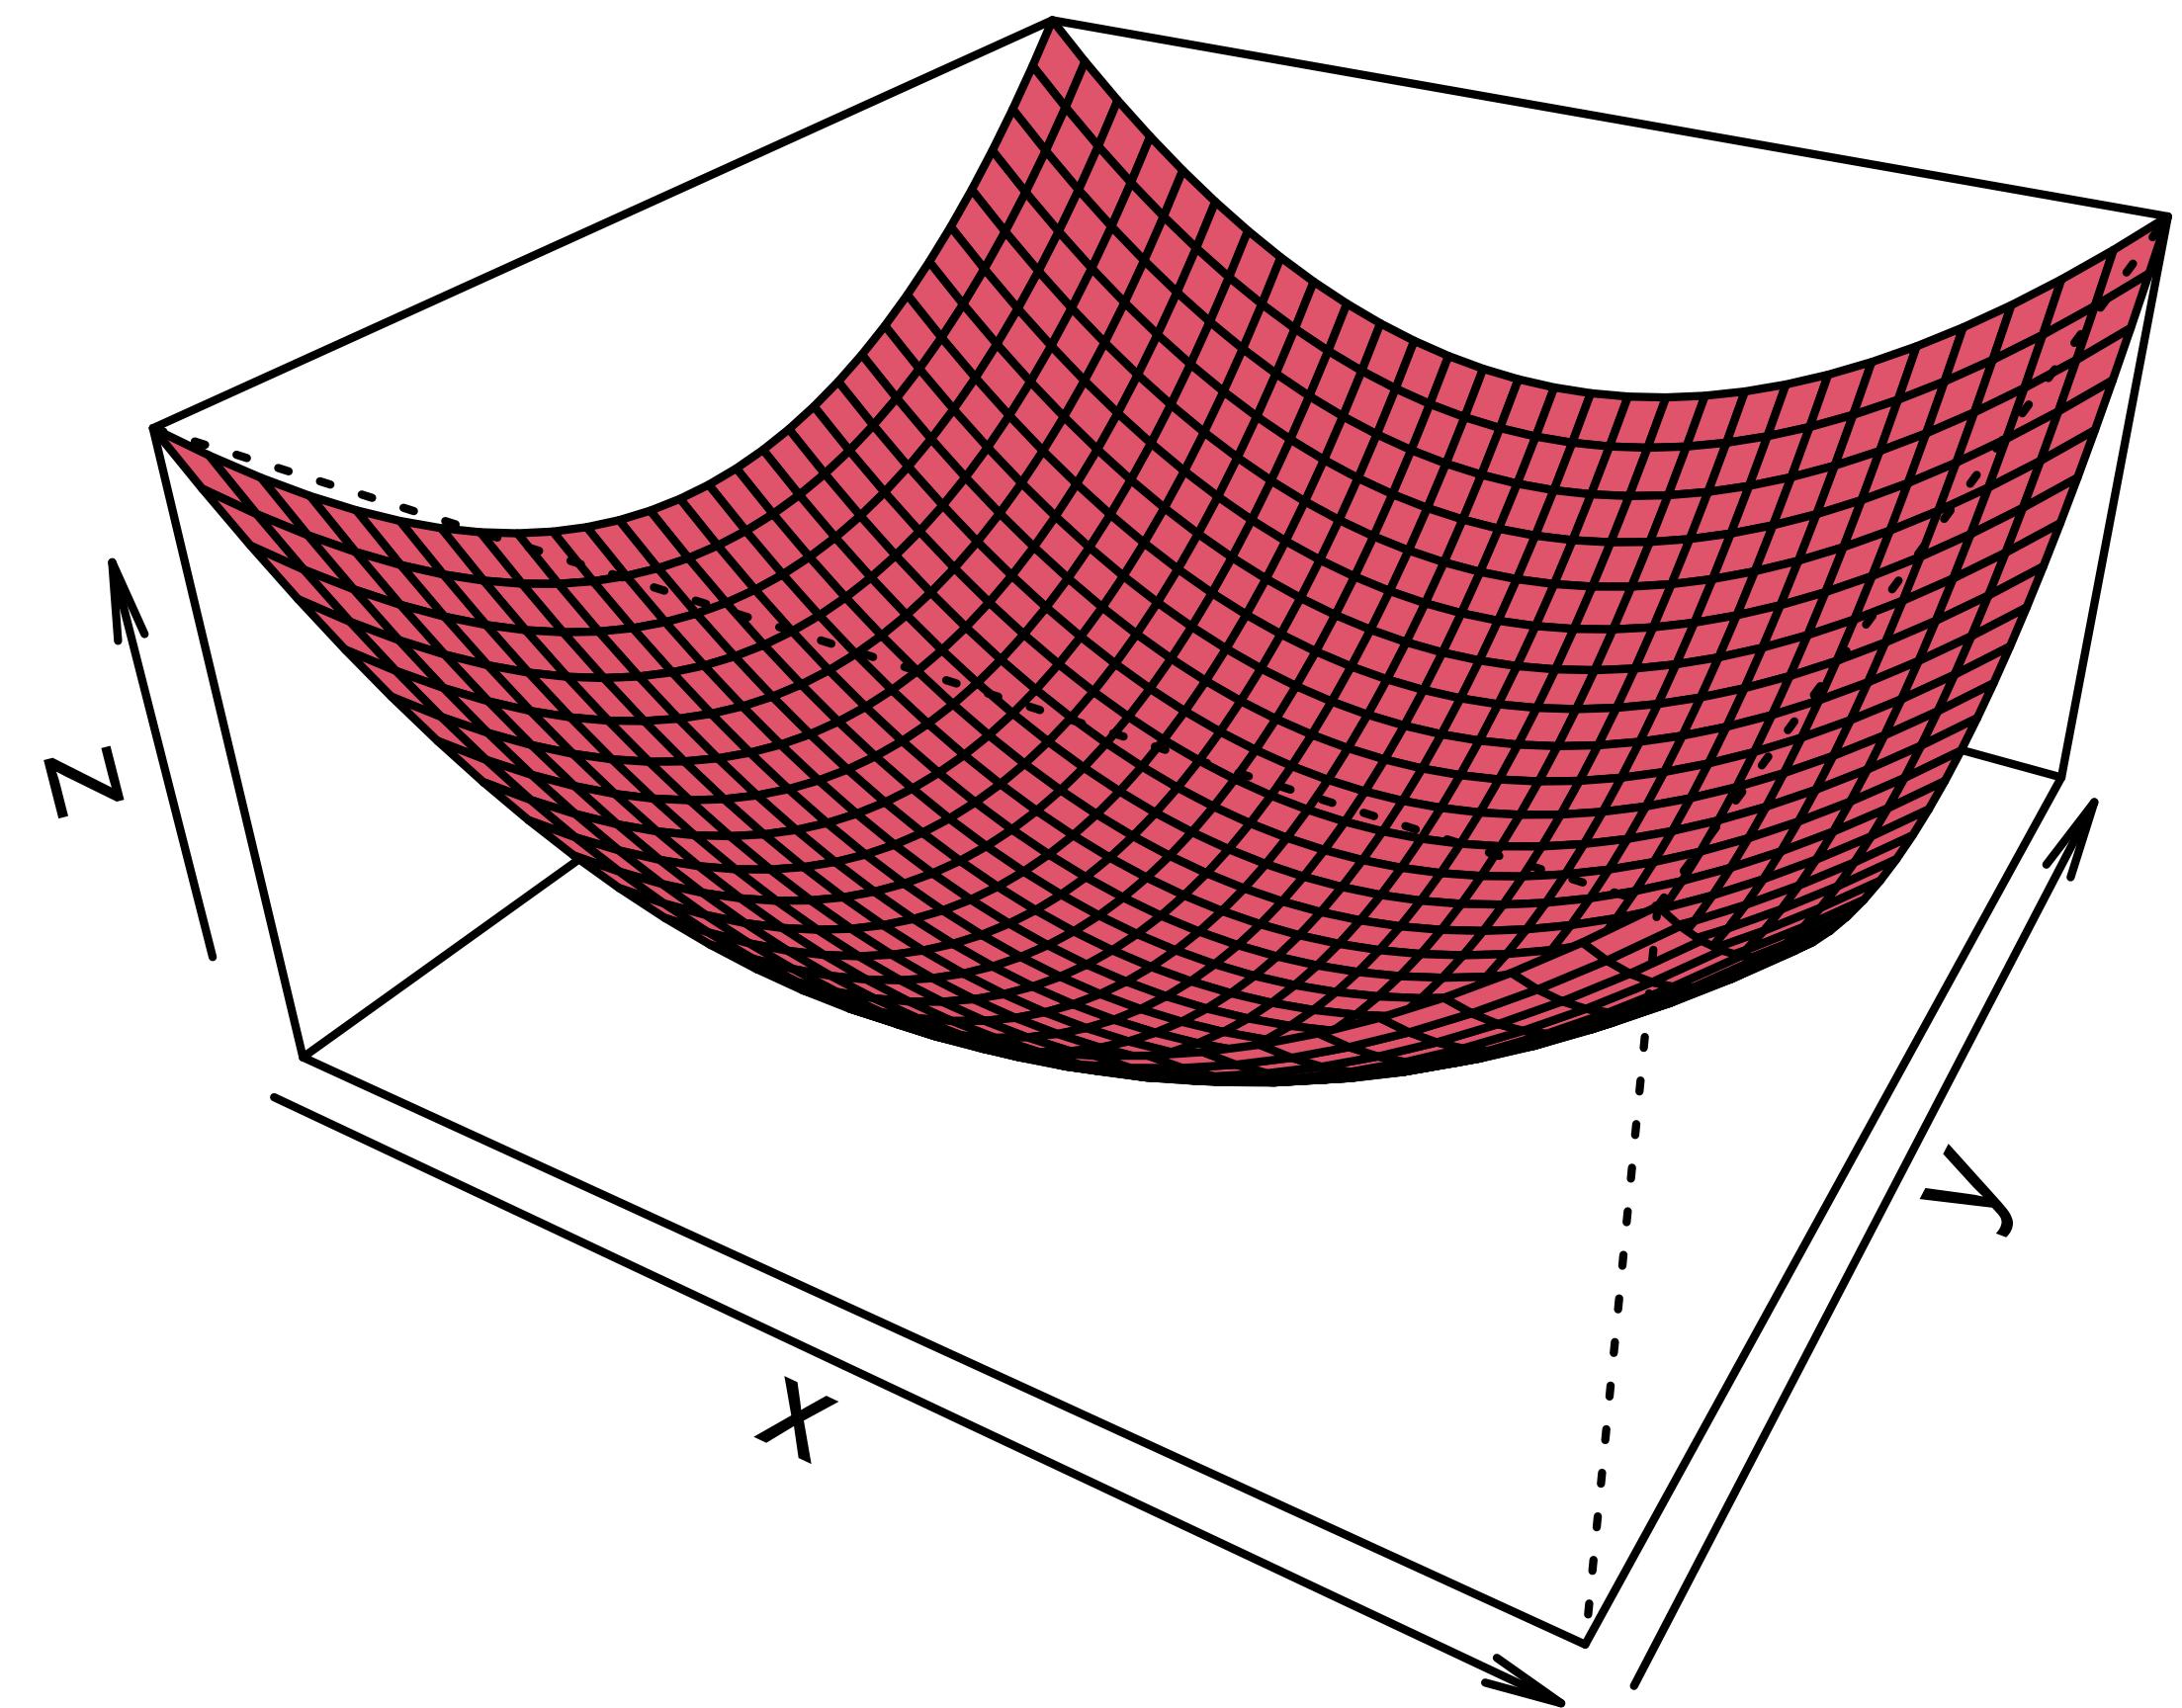
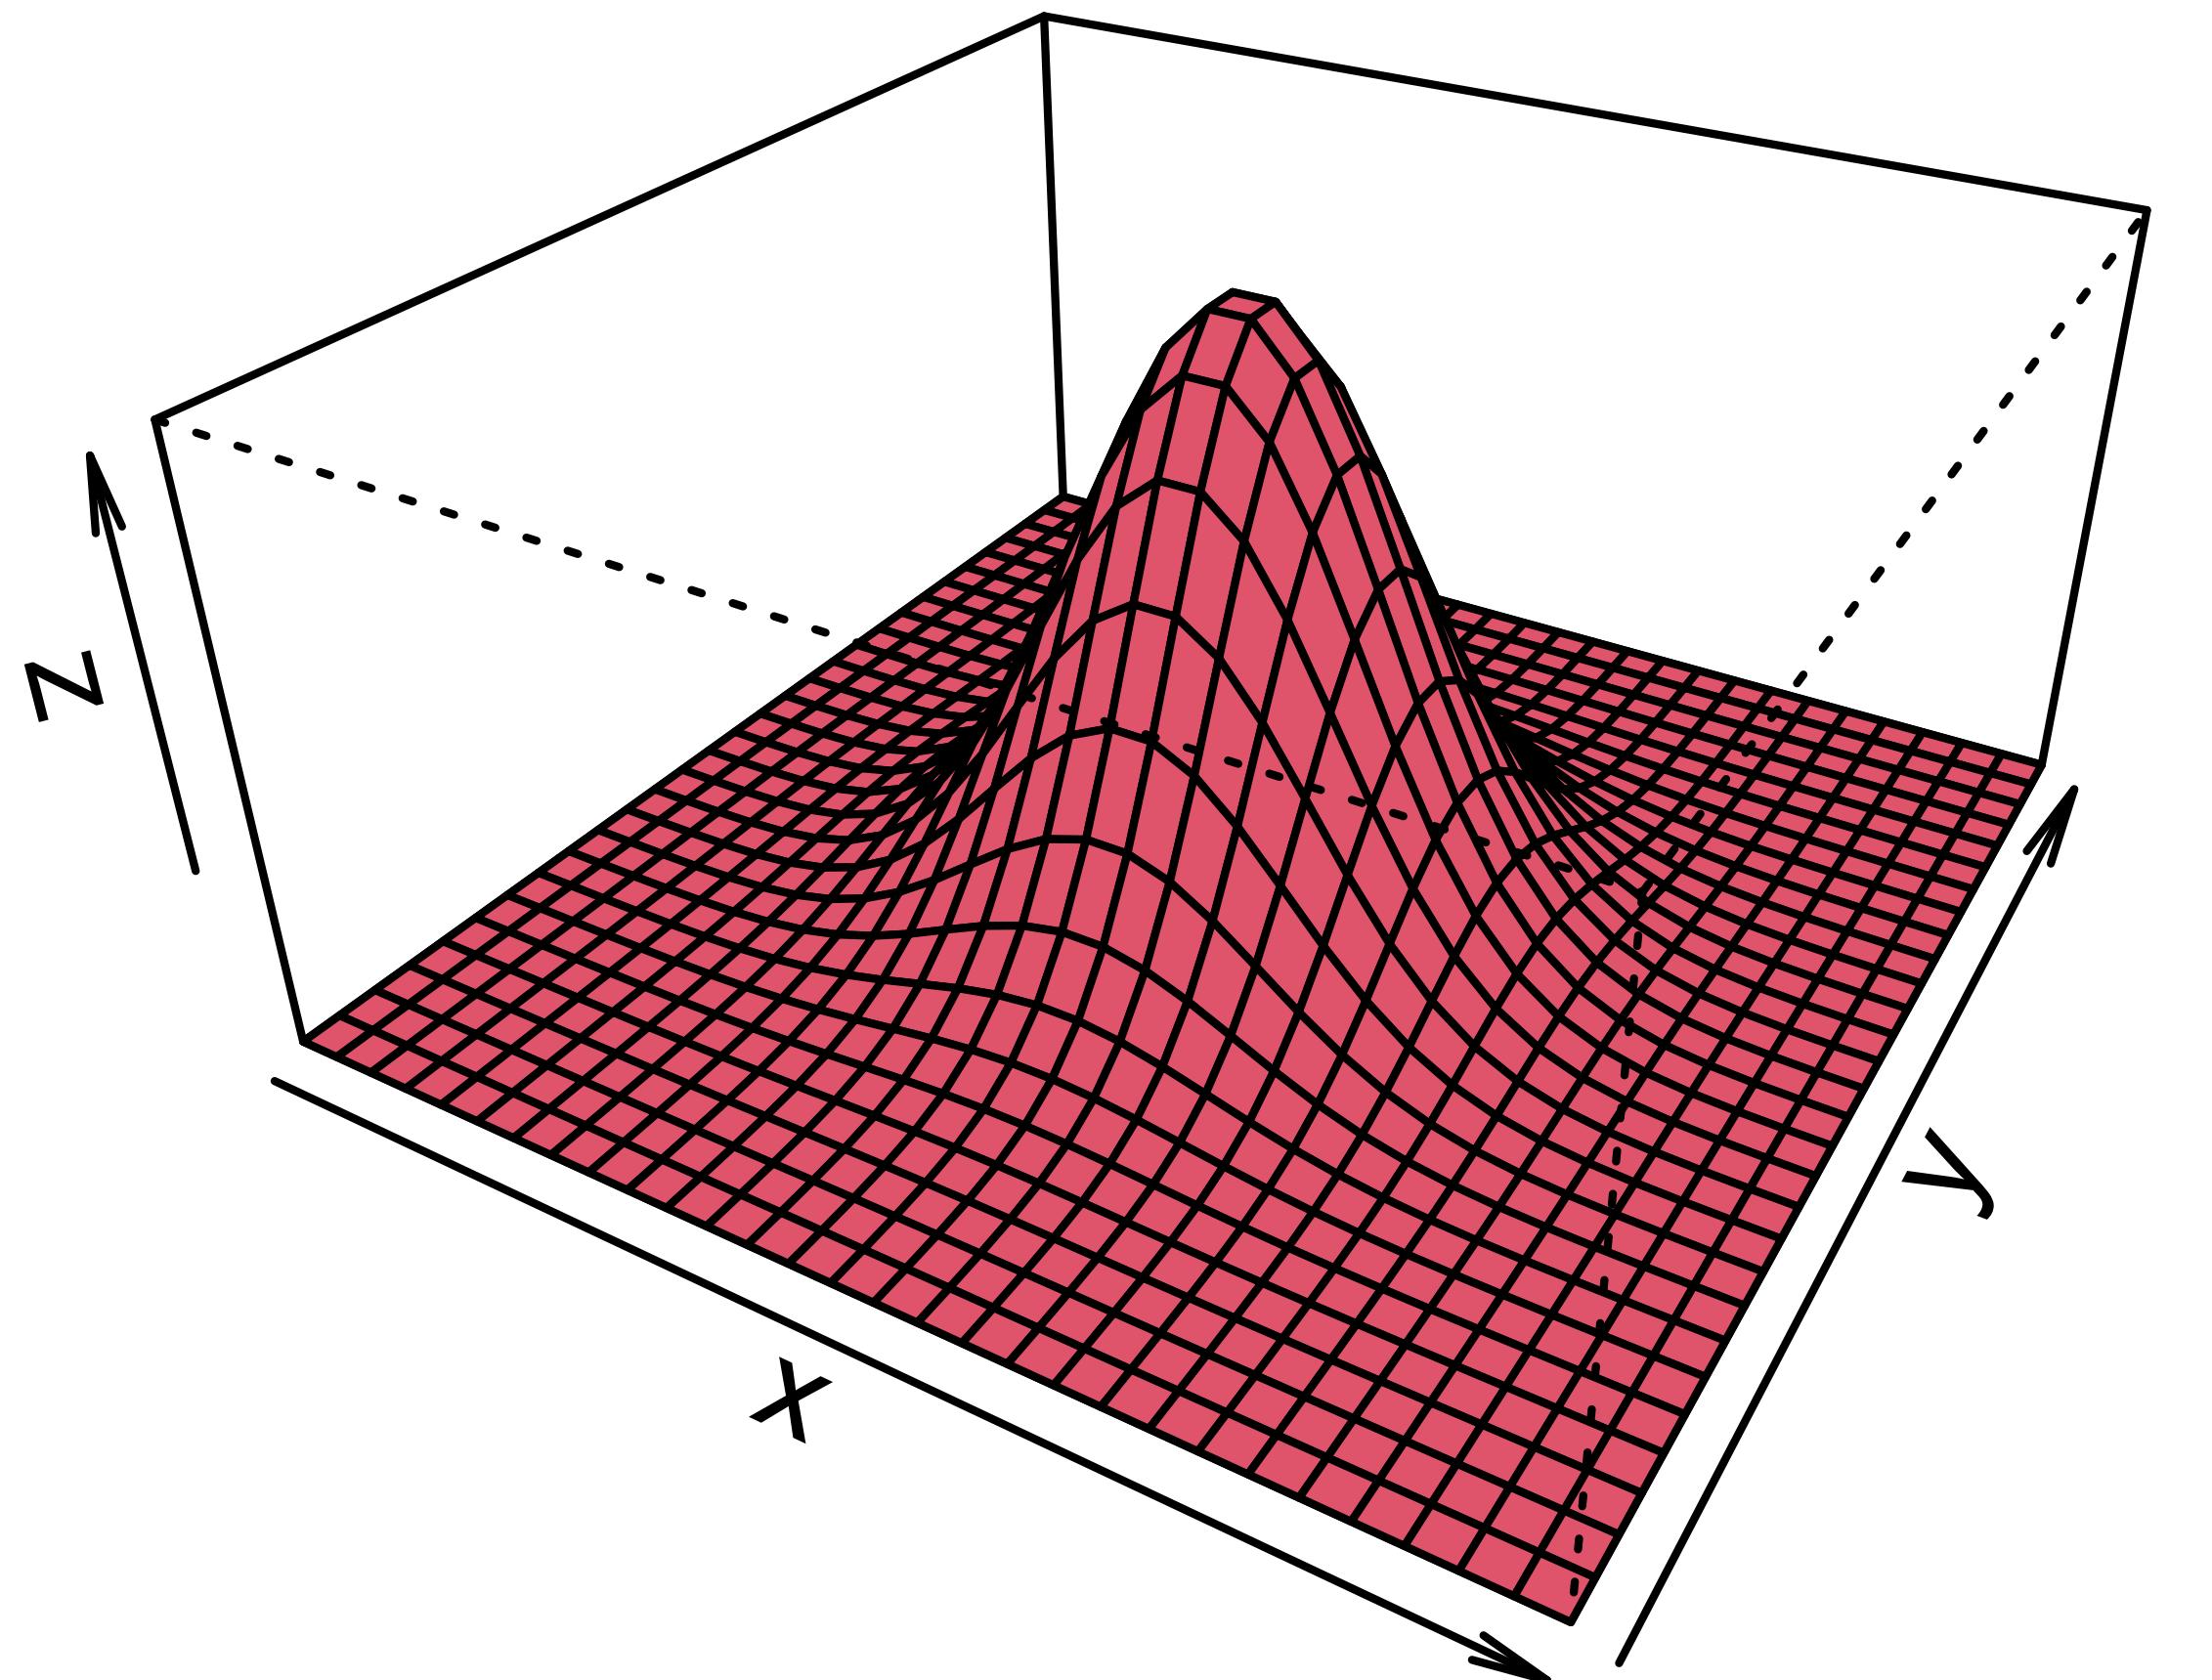


Basic Rosenbluth (aka Metropolis) algorithm









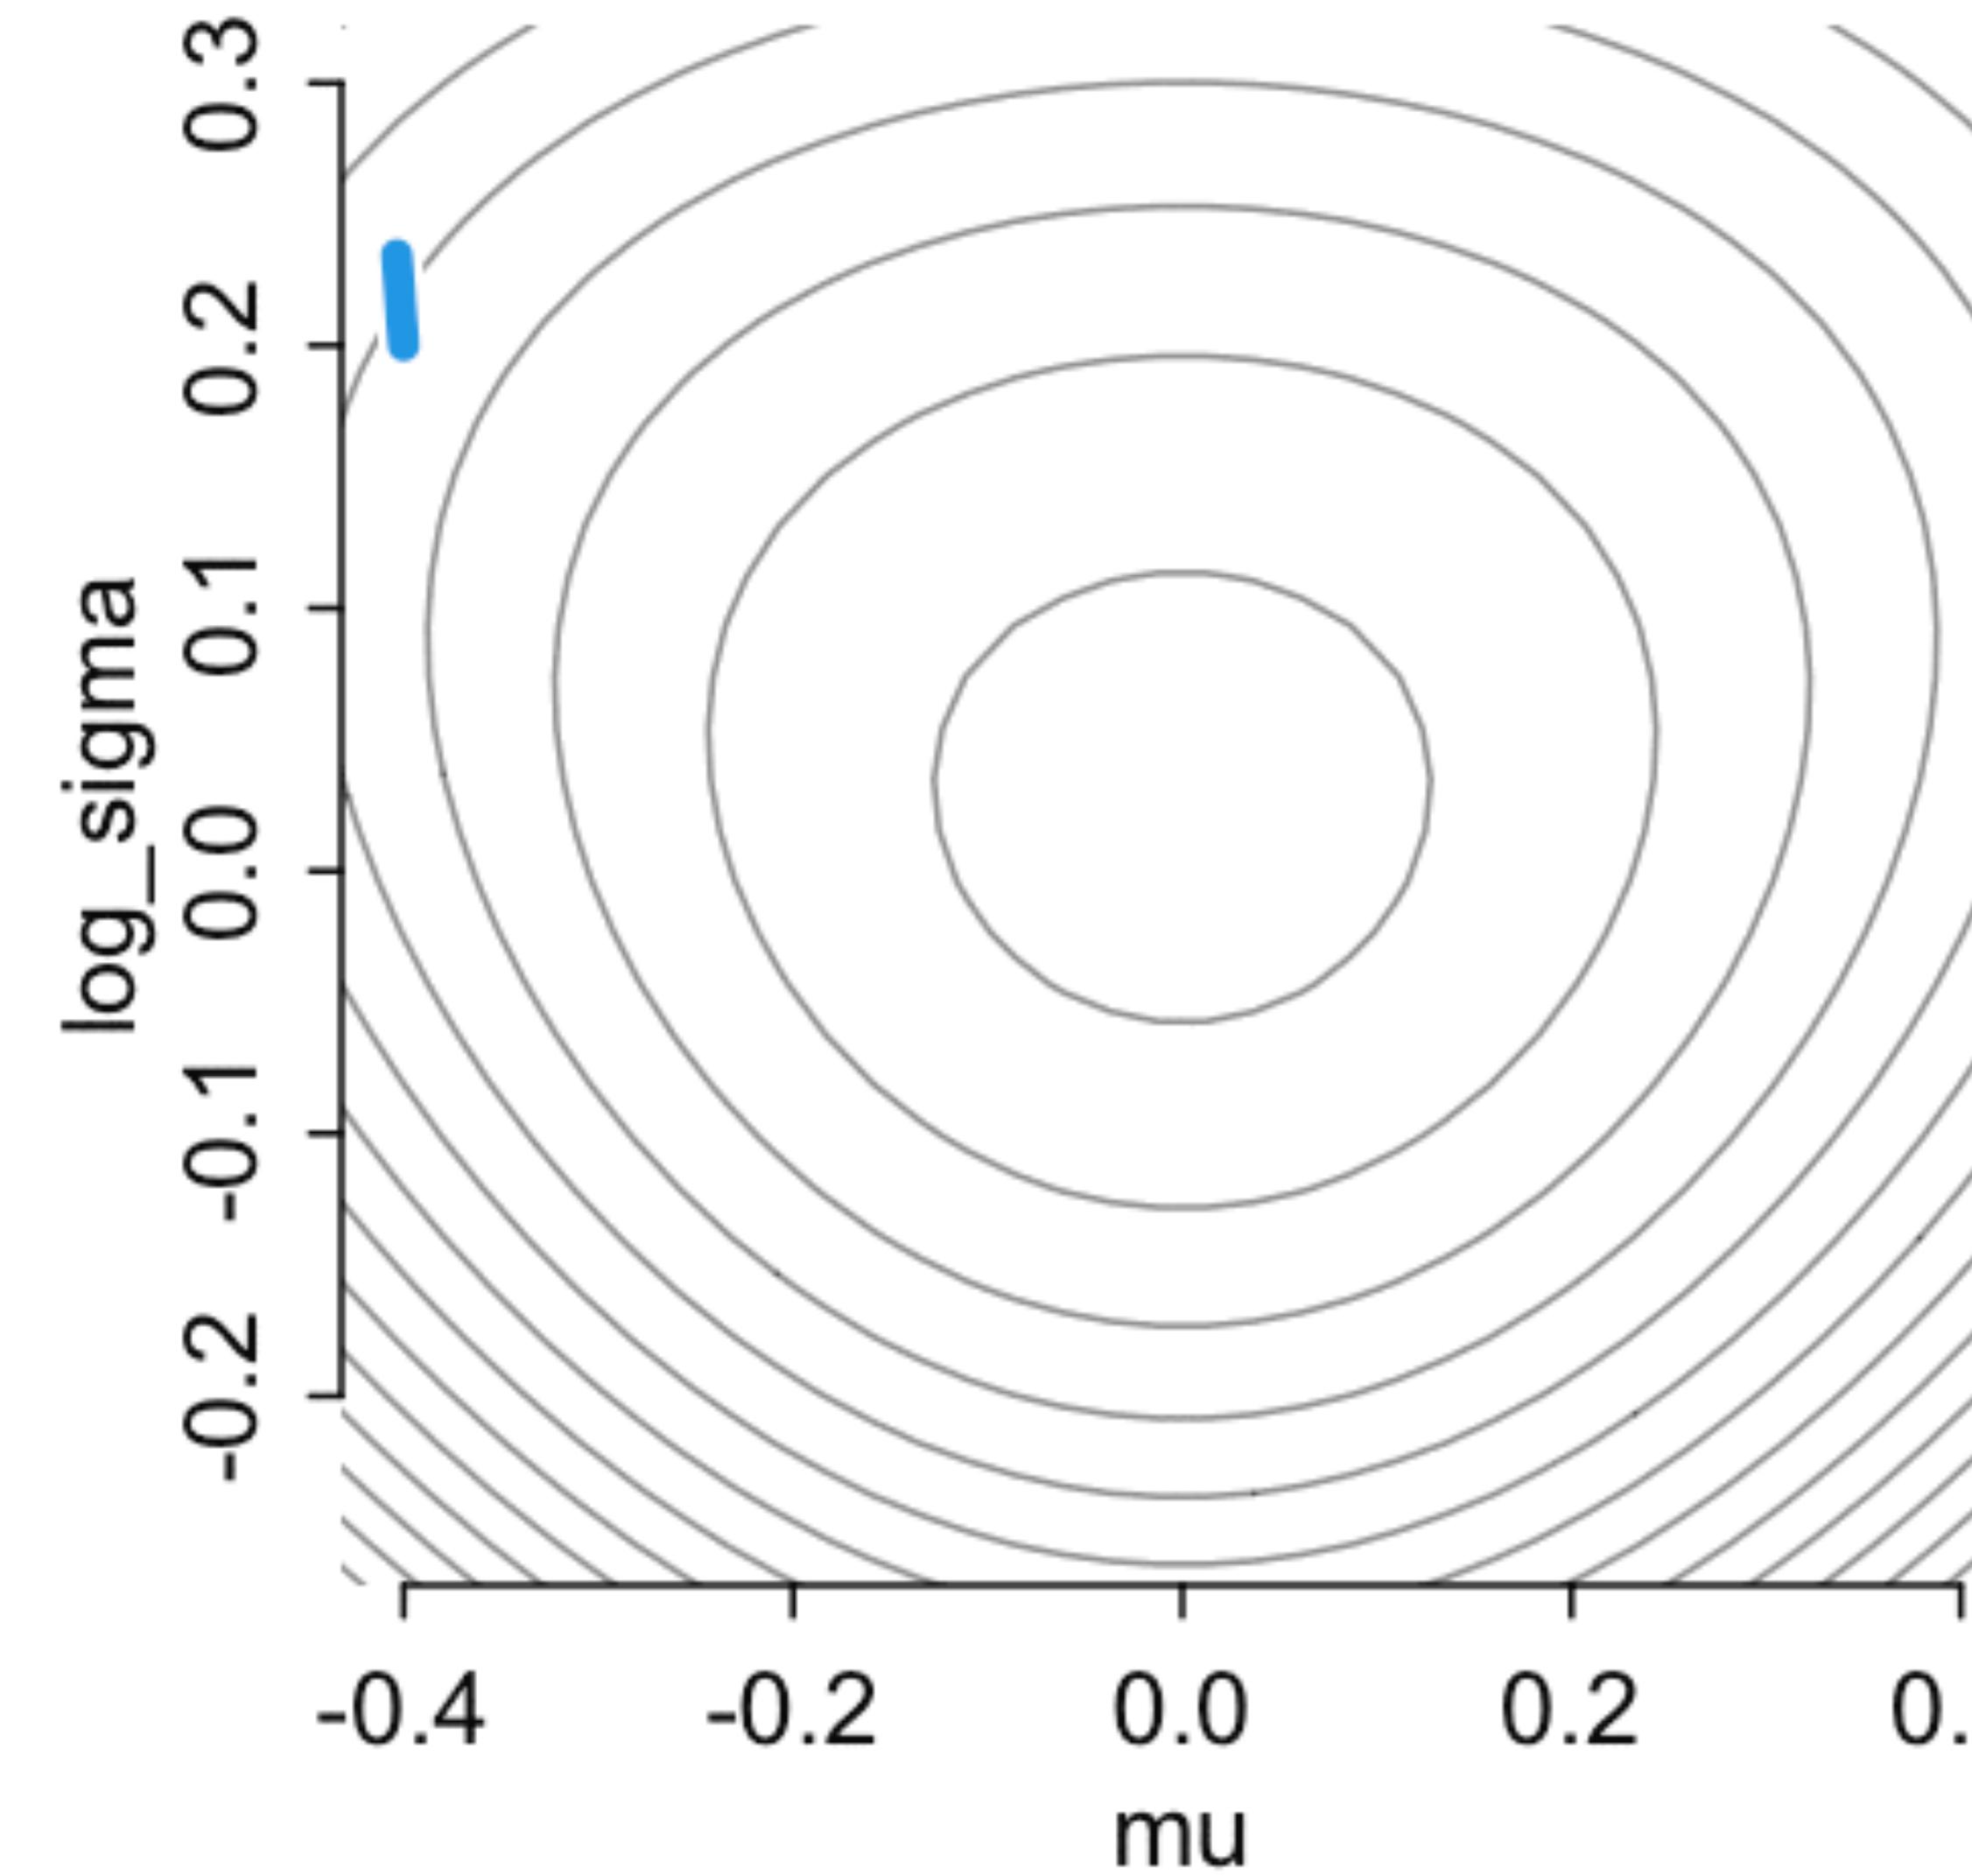


A wide-angle photograph of a large, shallow concrete bowl at a public park. A skateboarder in a white t-shirt and dark pants is performing a trick in the center of the bowl. The bowl is surrounded by a metal railing where several spectators are watching. The background shows trees and streetlights under a clear sky.

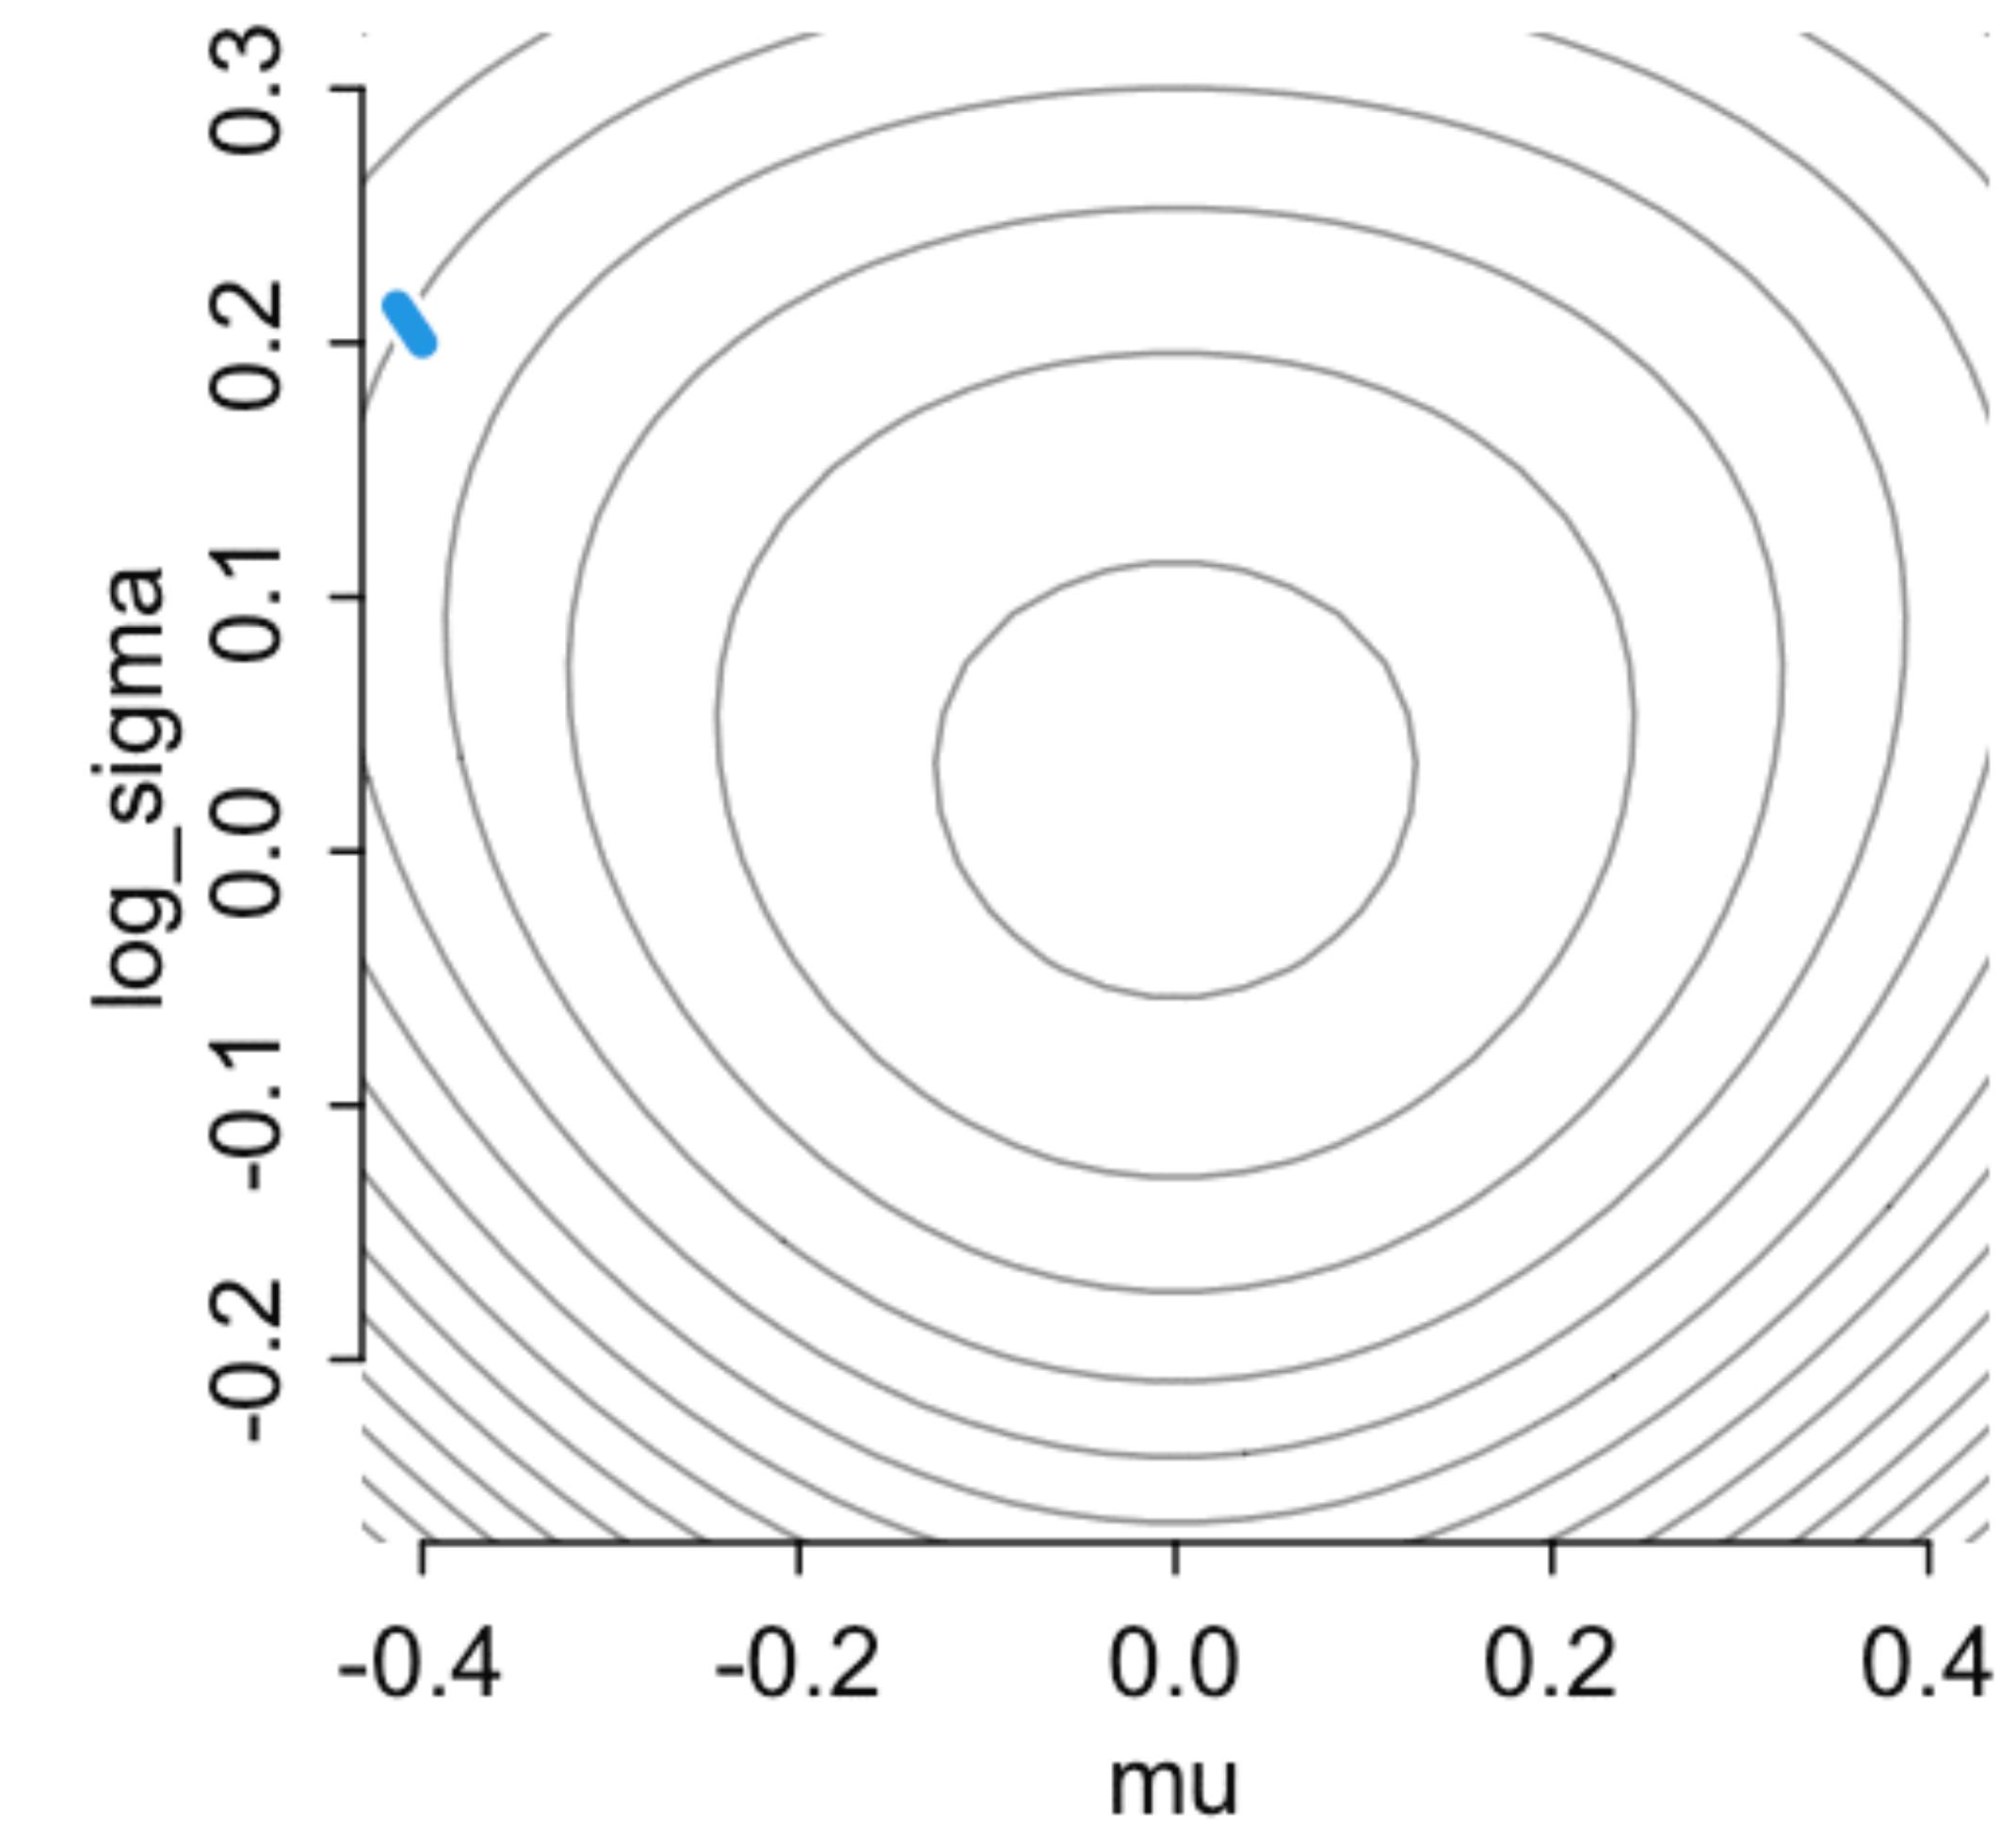
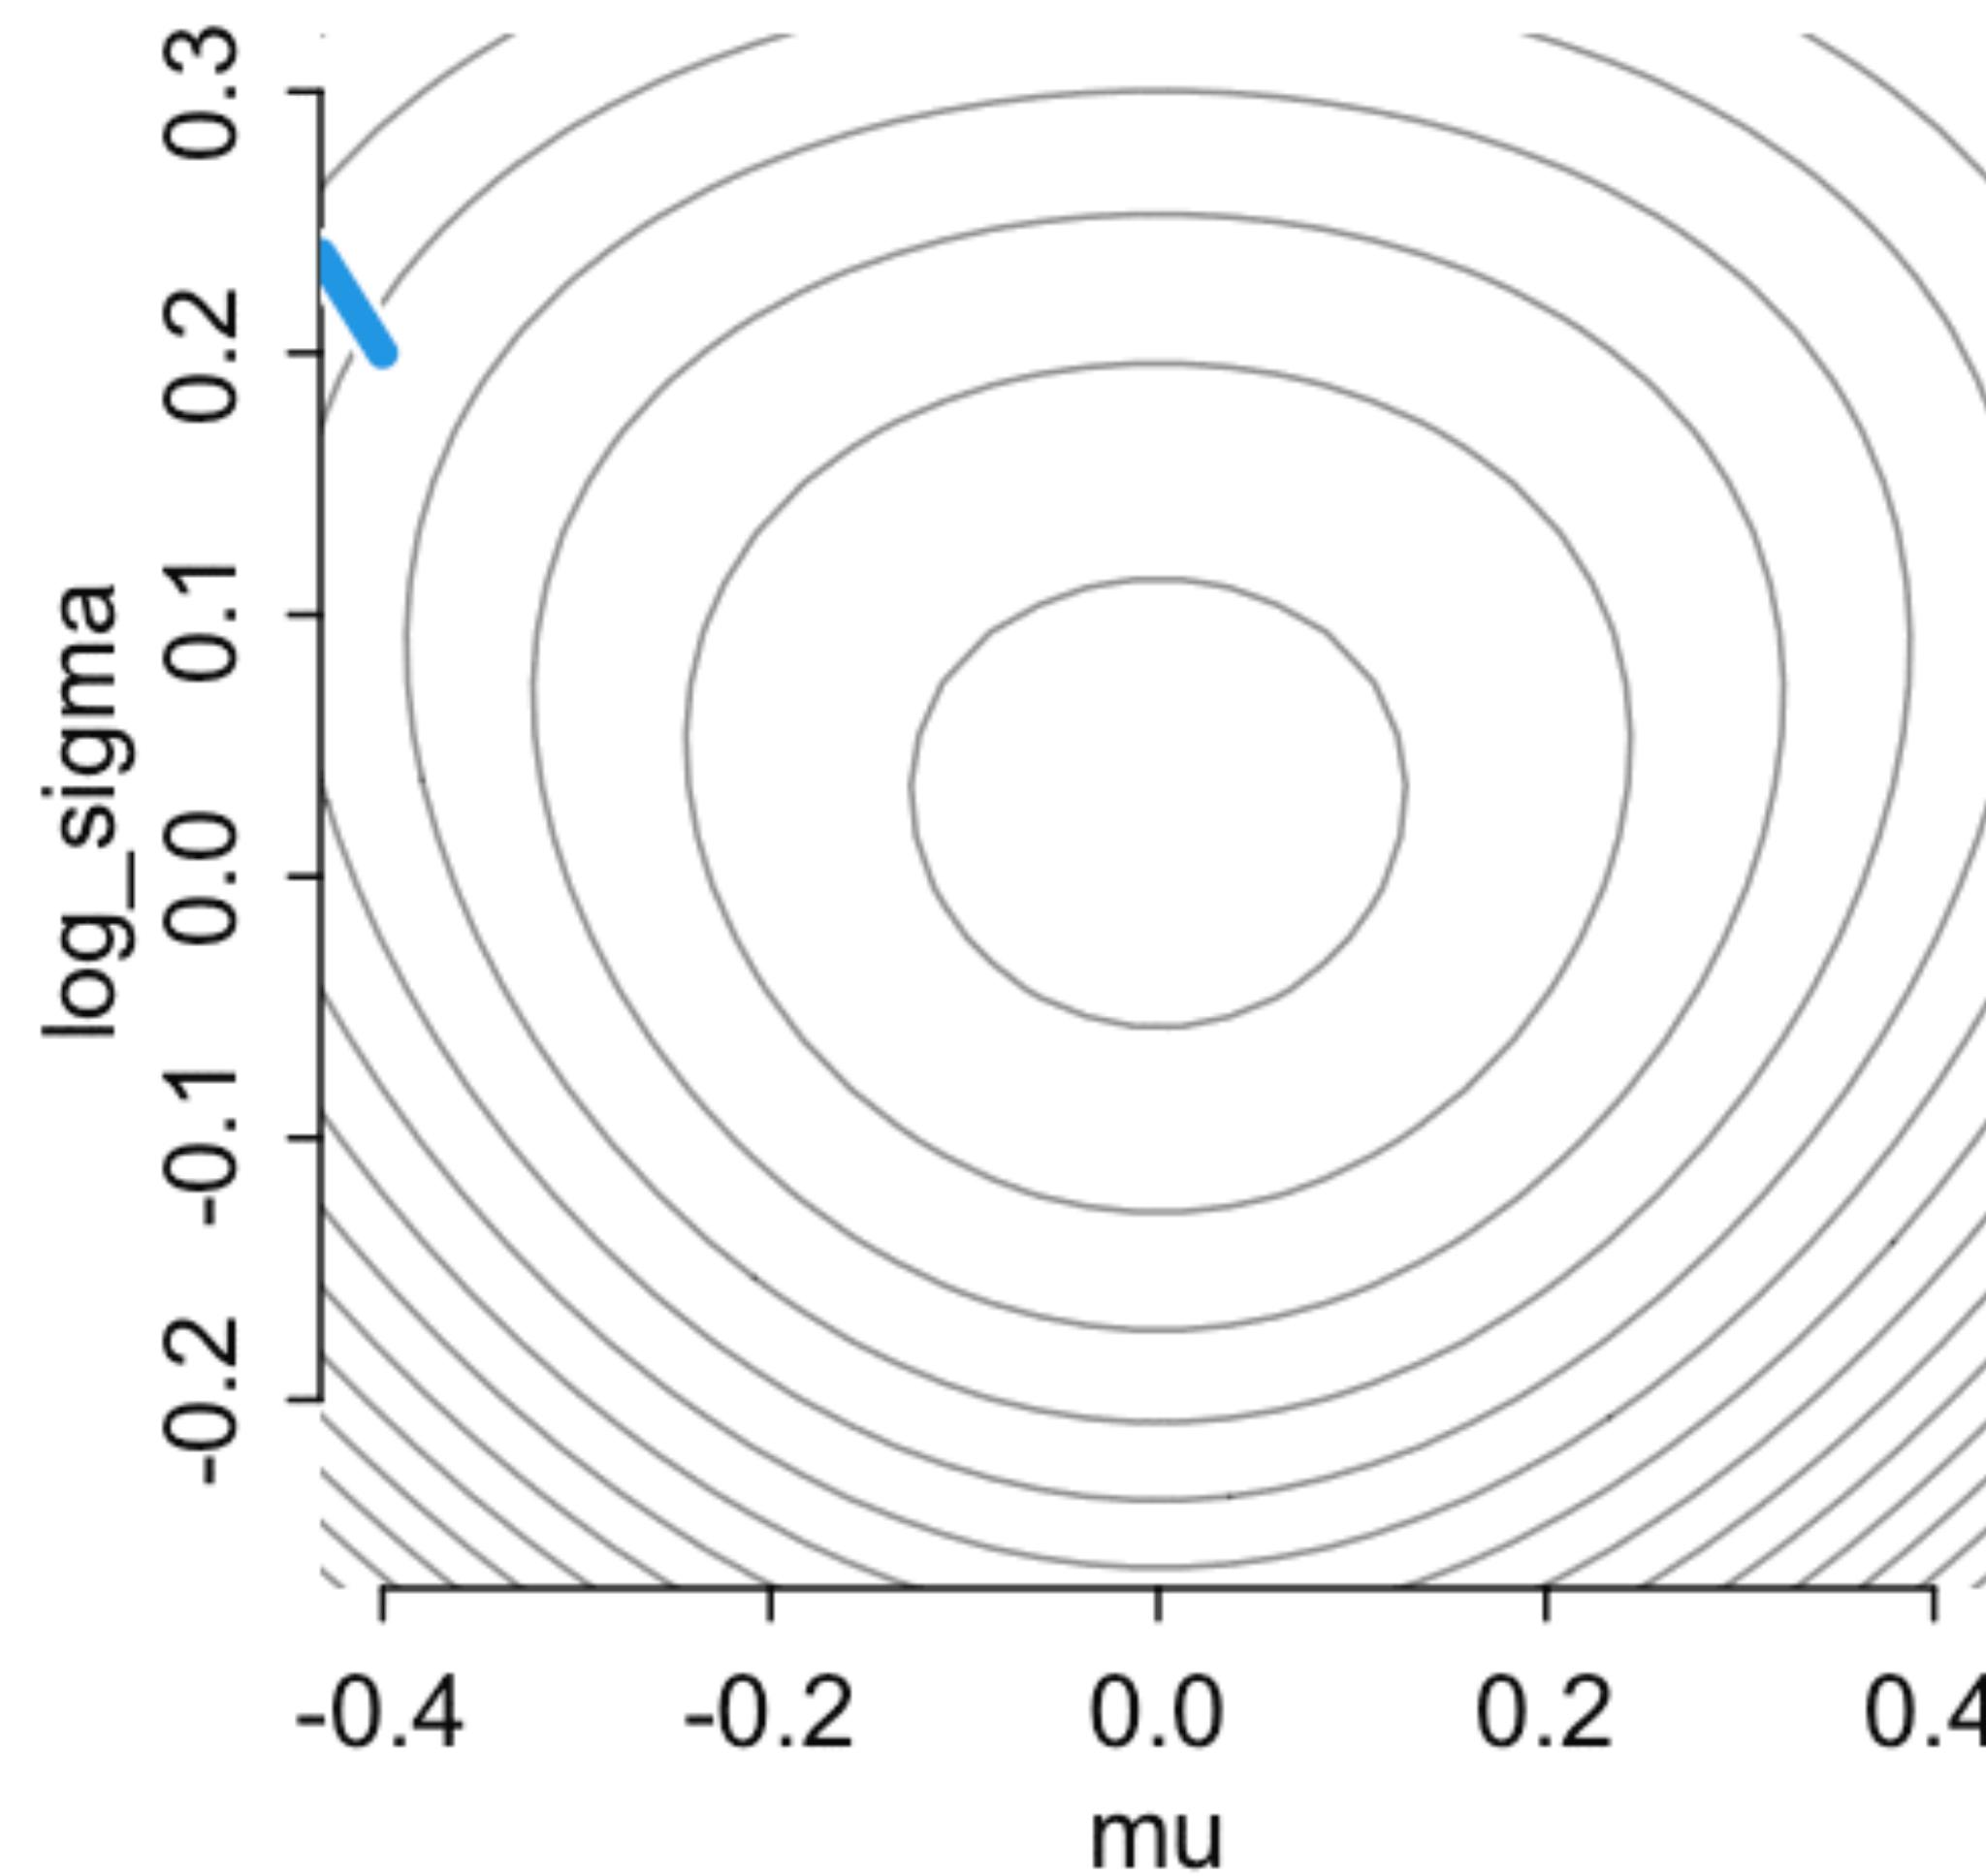
Low probability

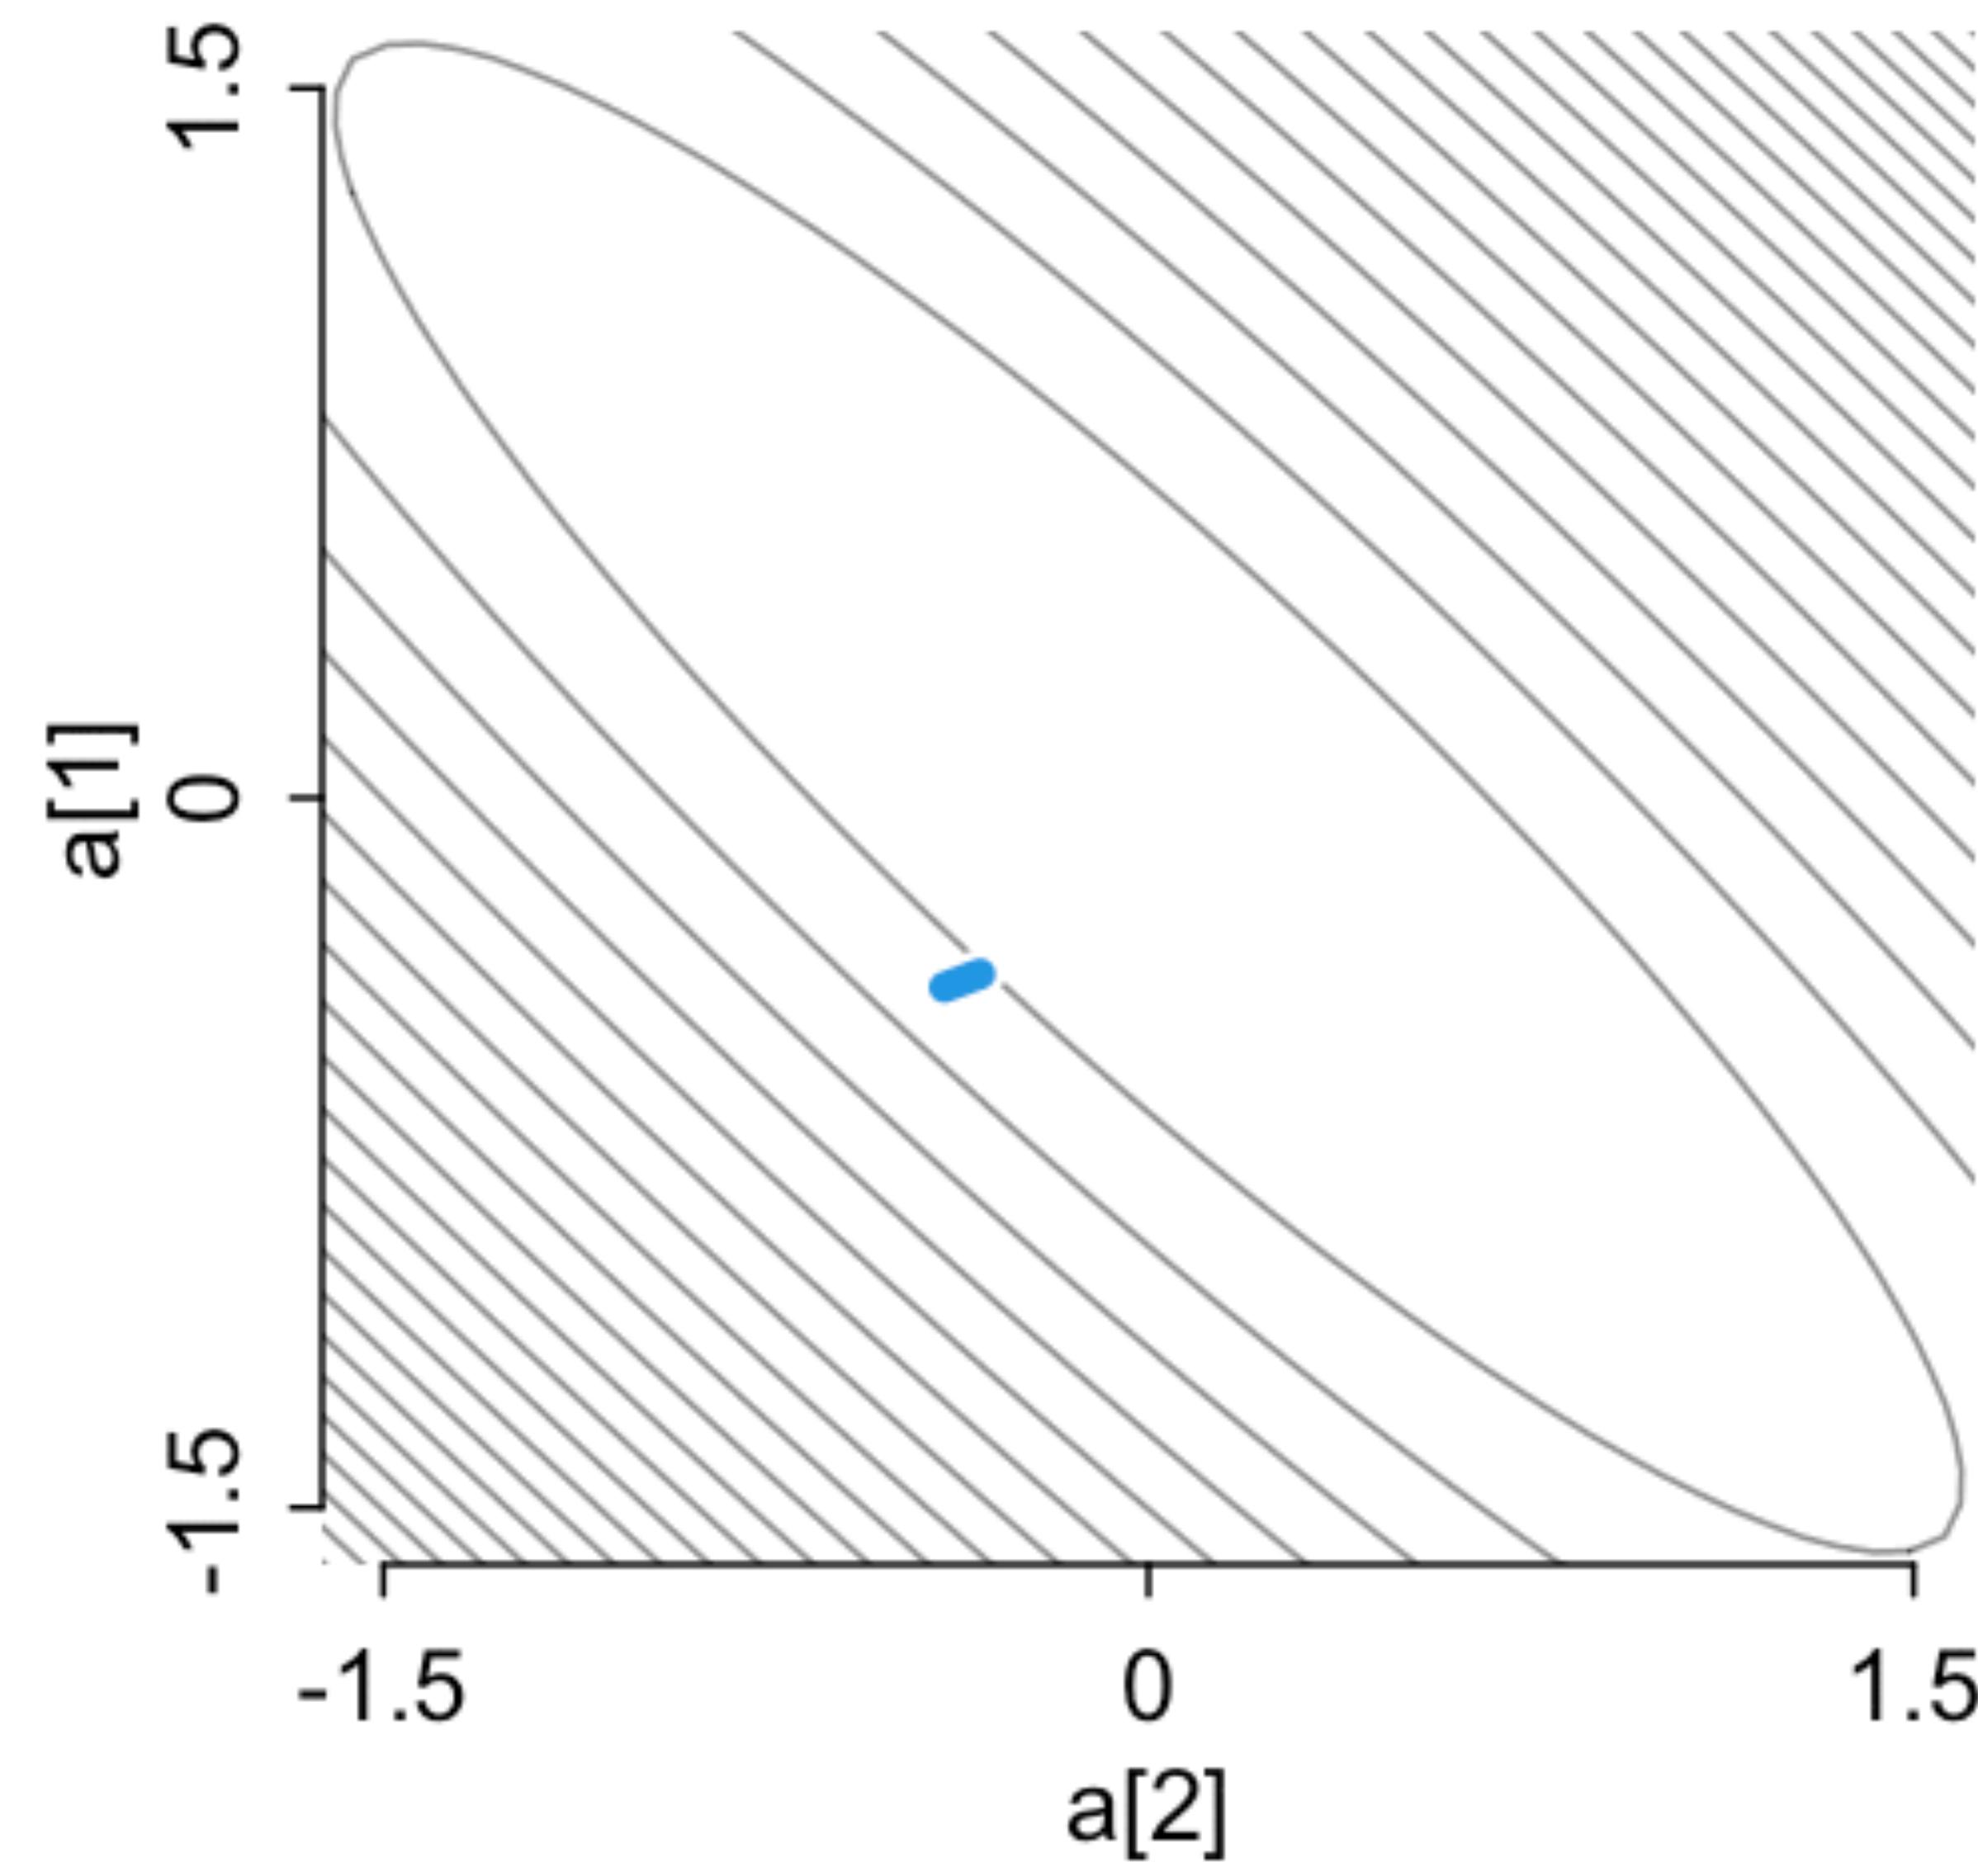
High probability

Hamiltonian Monte Carlo



Hamiltonian Monte Carlo





Overthinking: Hamiltonian Monte Carlo in the raw. The HMC algorithm needs five things to go: (1) a function U that returns the negative log-probability of the data at the current position (parameter values), (2) a function grad_U that returns the *gradient* of the negative log-probability at the current position, (3) a step size epsilon , (4) a count of leapfrog steps L , and (5) a starting position current_q . Keep in mind that the position is a vector of parameter values and that the gradient also needs to return a vector of the same length. So that these U and grad_U functions make more sense, let's present them first, built custom for the 2D Gaussian example. The U function just expresses the log-posterior, as stated before in the main text:

$$\sum_i \log p(y_i|\mu_y, 1) + \sum_i \log p(x_i|\mu_x, 1) + \log p(\mu_y|0, 0.5) + \log p(\mu_x|0, 0.5)$$

So it's just four calls to `dnorm` really:

R code
9.5

```
# U needs to return neg-log-probability
U <- function( q , a=0 , b=1 , k=0 , d=1 ) {
  muy <- q[1]
  mux <- q[2]
  U <- sum( dnorm(y,muy,1,log=TRUE) ) + sum( dnorm(x,mux,1,log=TRUE) ) +
    dnorm(muy,a,b,log=TRUE) + dnorm(mux,k,d,log=TRUE)
  return( -U )
}
```

Now the gradient function requires two partial derivatives. Luckily, Gaussian derivatives are very clean. The derivative of the logarithm of any univariate Gaussian with mean a and standard deviation b with respect to a is:

$$\frac{\partial \log N(y|a, b)}{\partial a} = \frac{y - a}{b^2}$$

And since the derivative of a sum is a sum of derivatives, this is all we need to write the gradients:

$$\frac{\partial U}{\partial \mu_x} = \frac{\partial \log N(x|\mu_x, 1)}{\partial \mu_x} + \frac{\partial \log N(\mu_x|0, 0.5)}{\partial \mu_x} = \sum_i \frac{x_i - \mu_x}{1^2} + \frac{0 - \mu_x}{0.5^2}$$

Calculus is a superpower

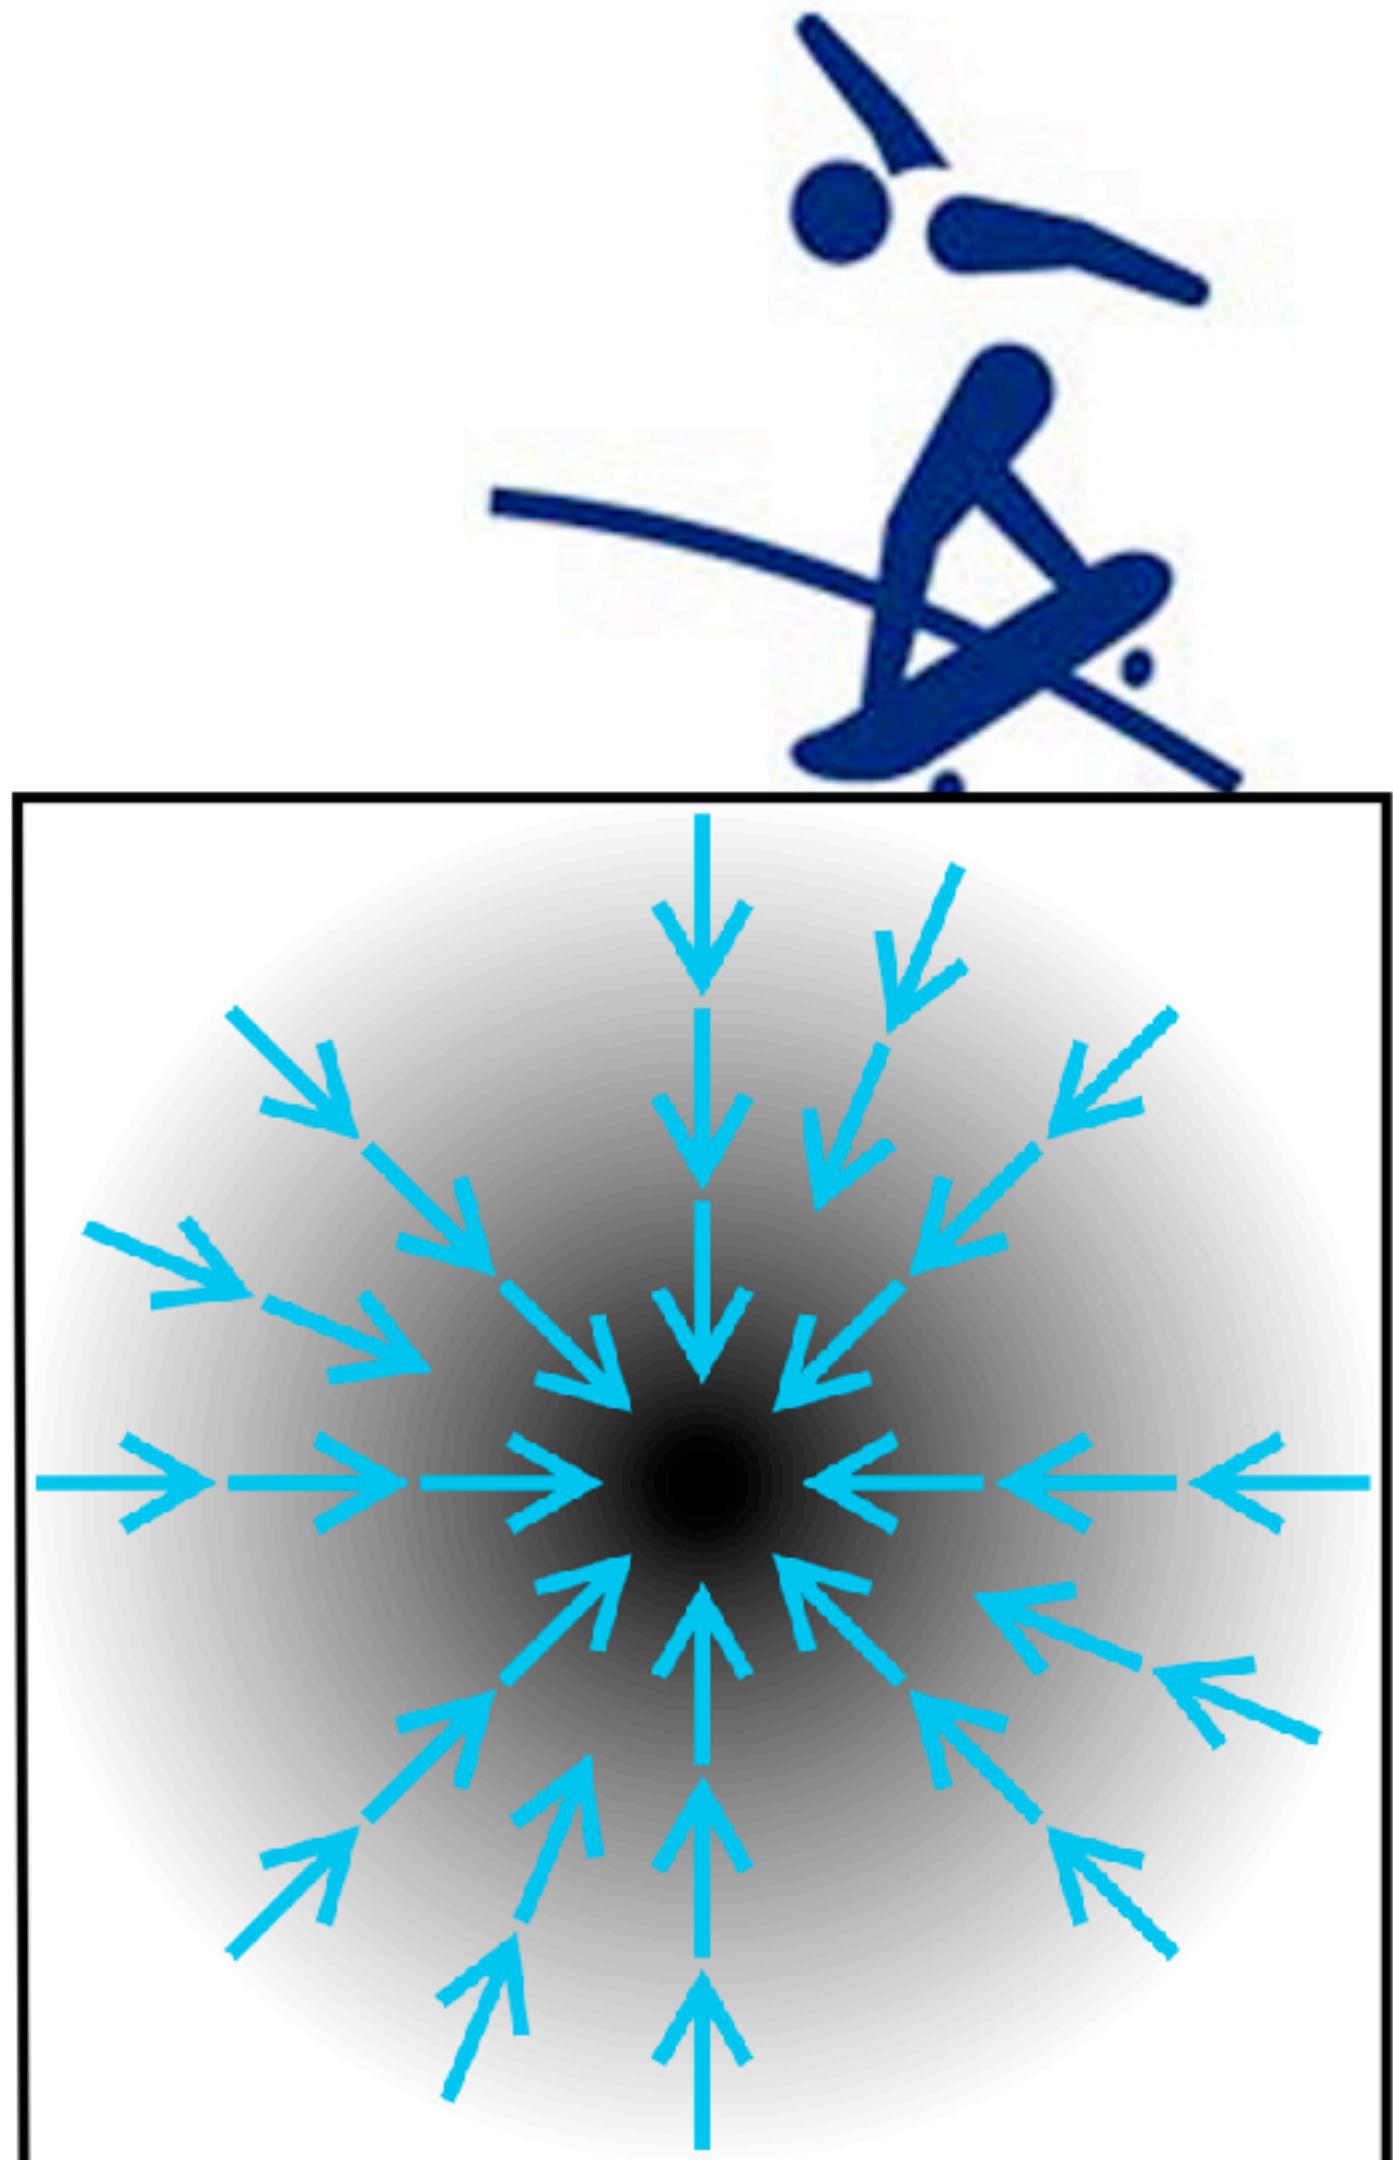
Hamiltonian Monte Carlo needs **gradients**

How does it get them? Write them yourself or...

Auto-diff: Automatic differentiation

Symbolic derivatives of your model code

Used in many machine learning approaches;
“Backpropagation” is special case



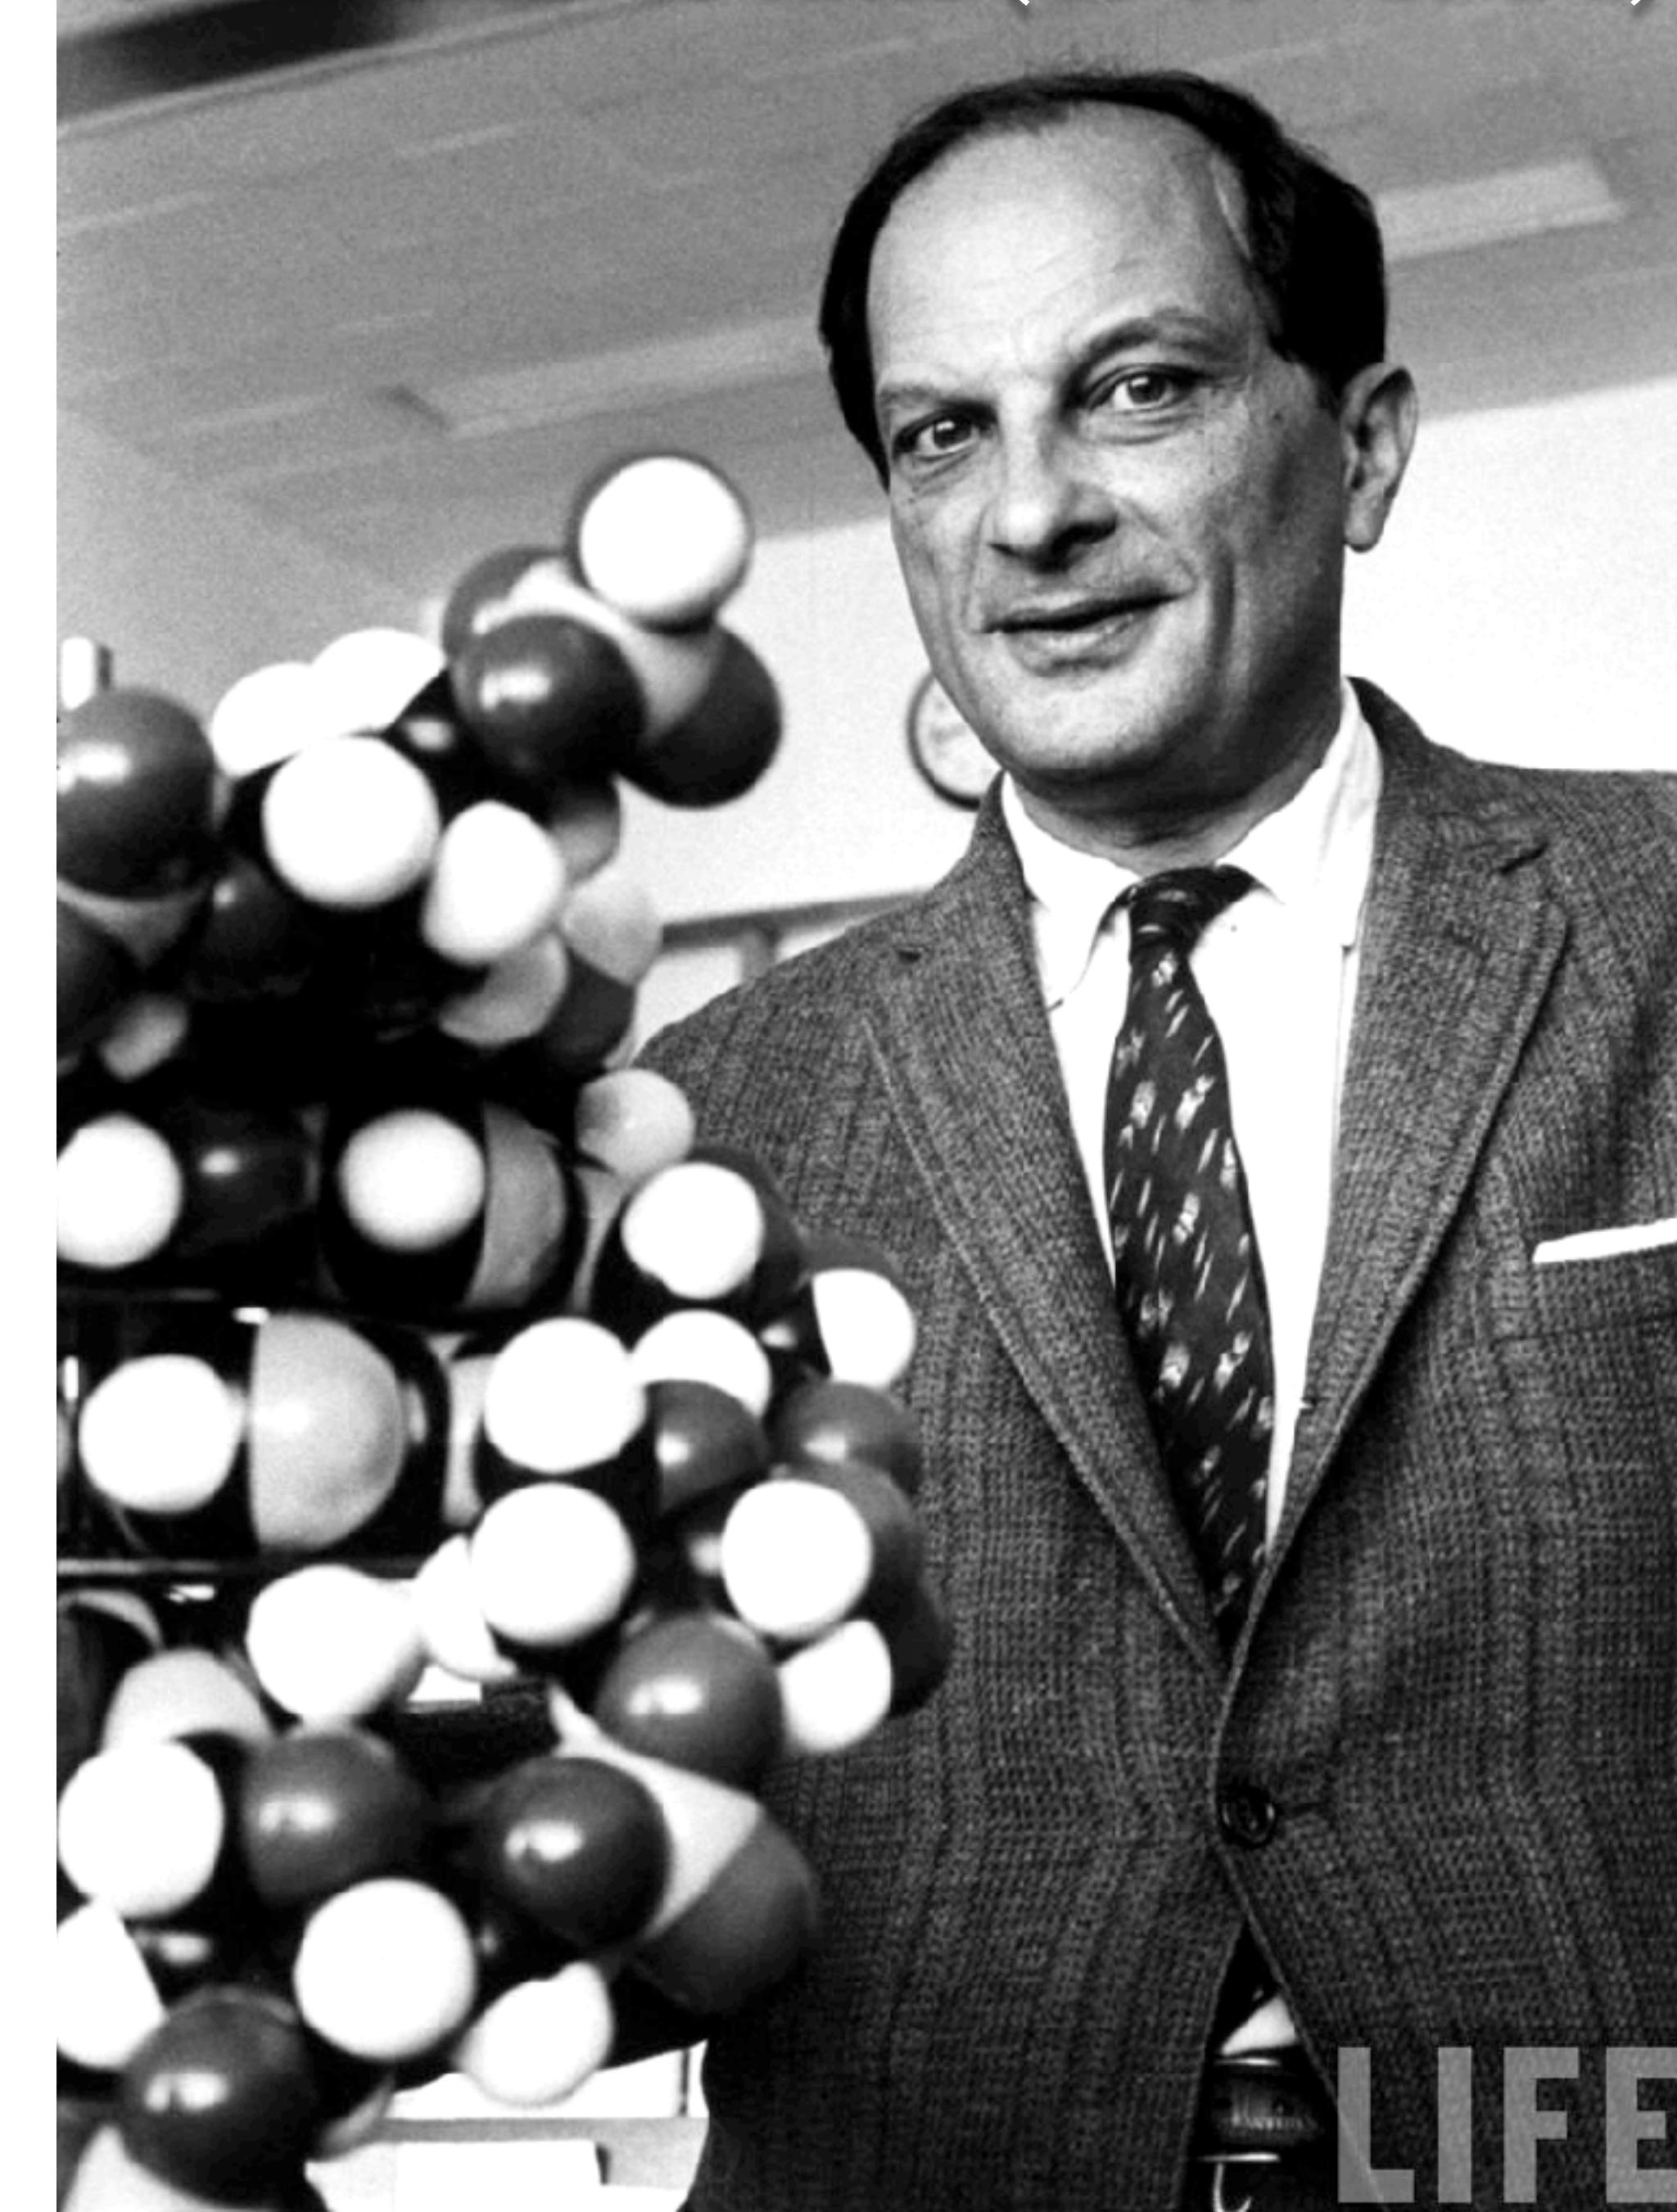


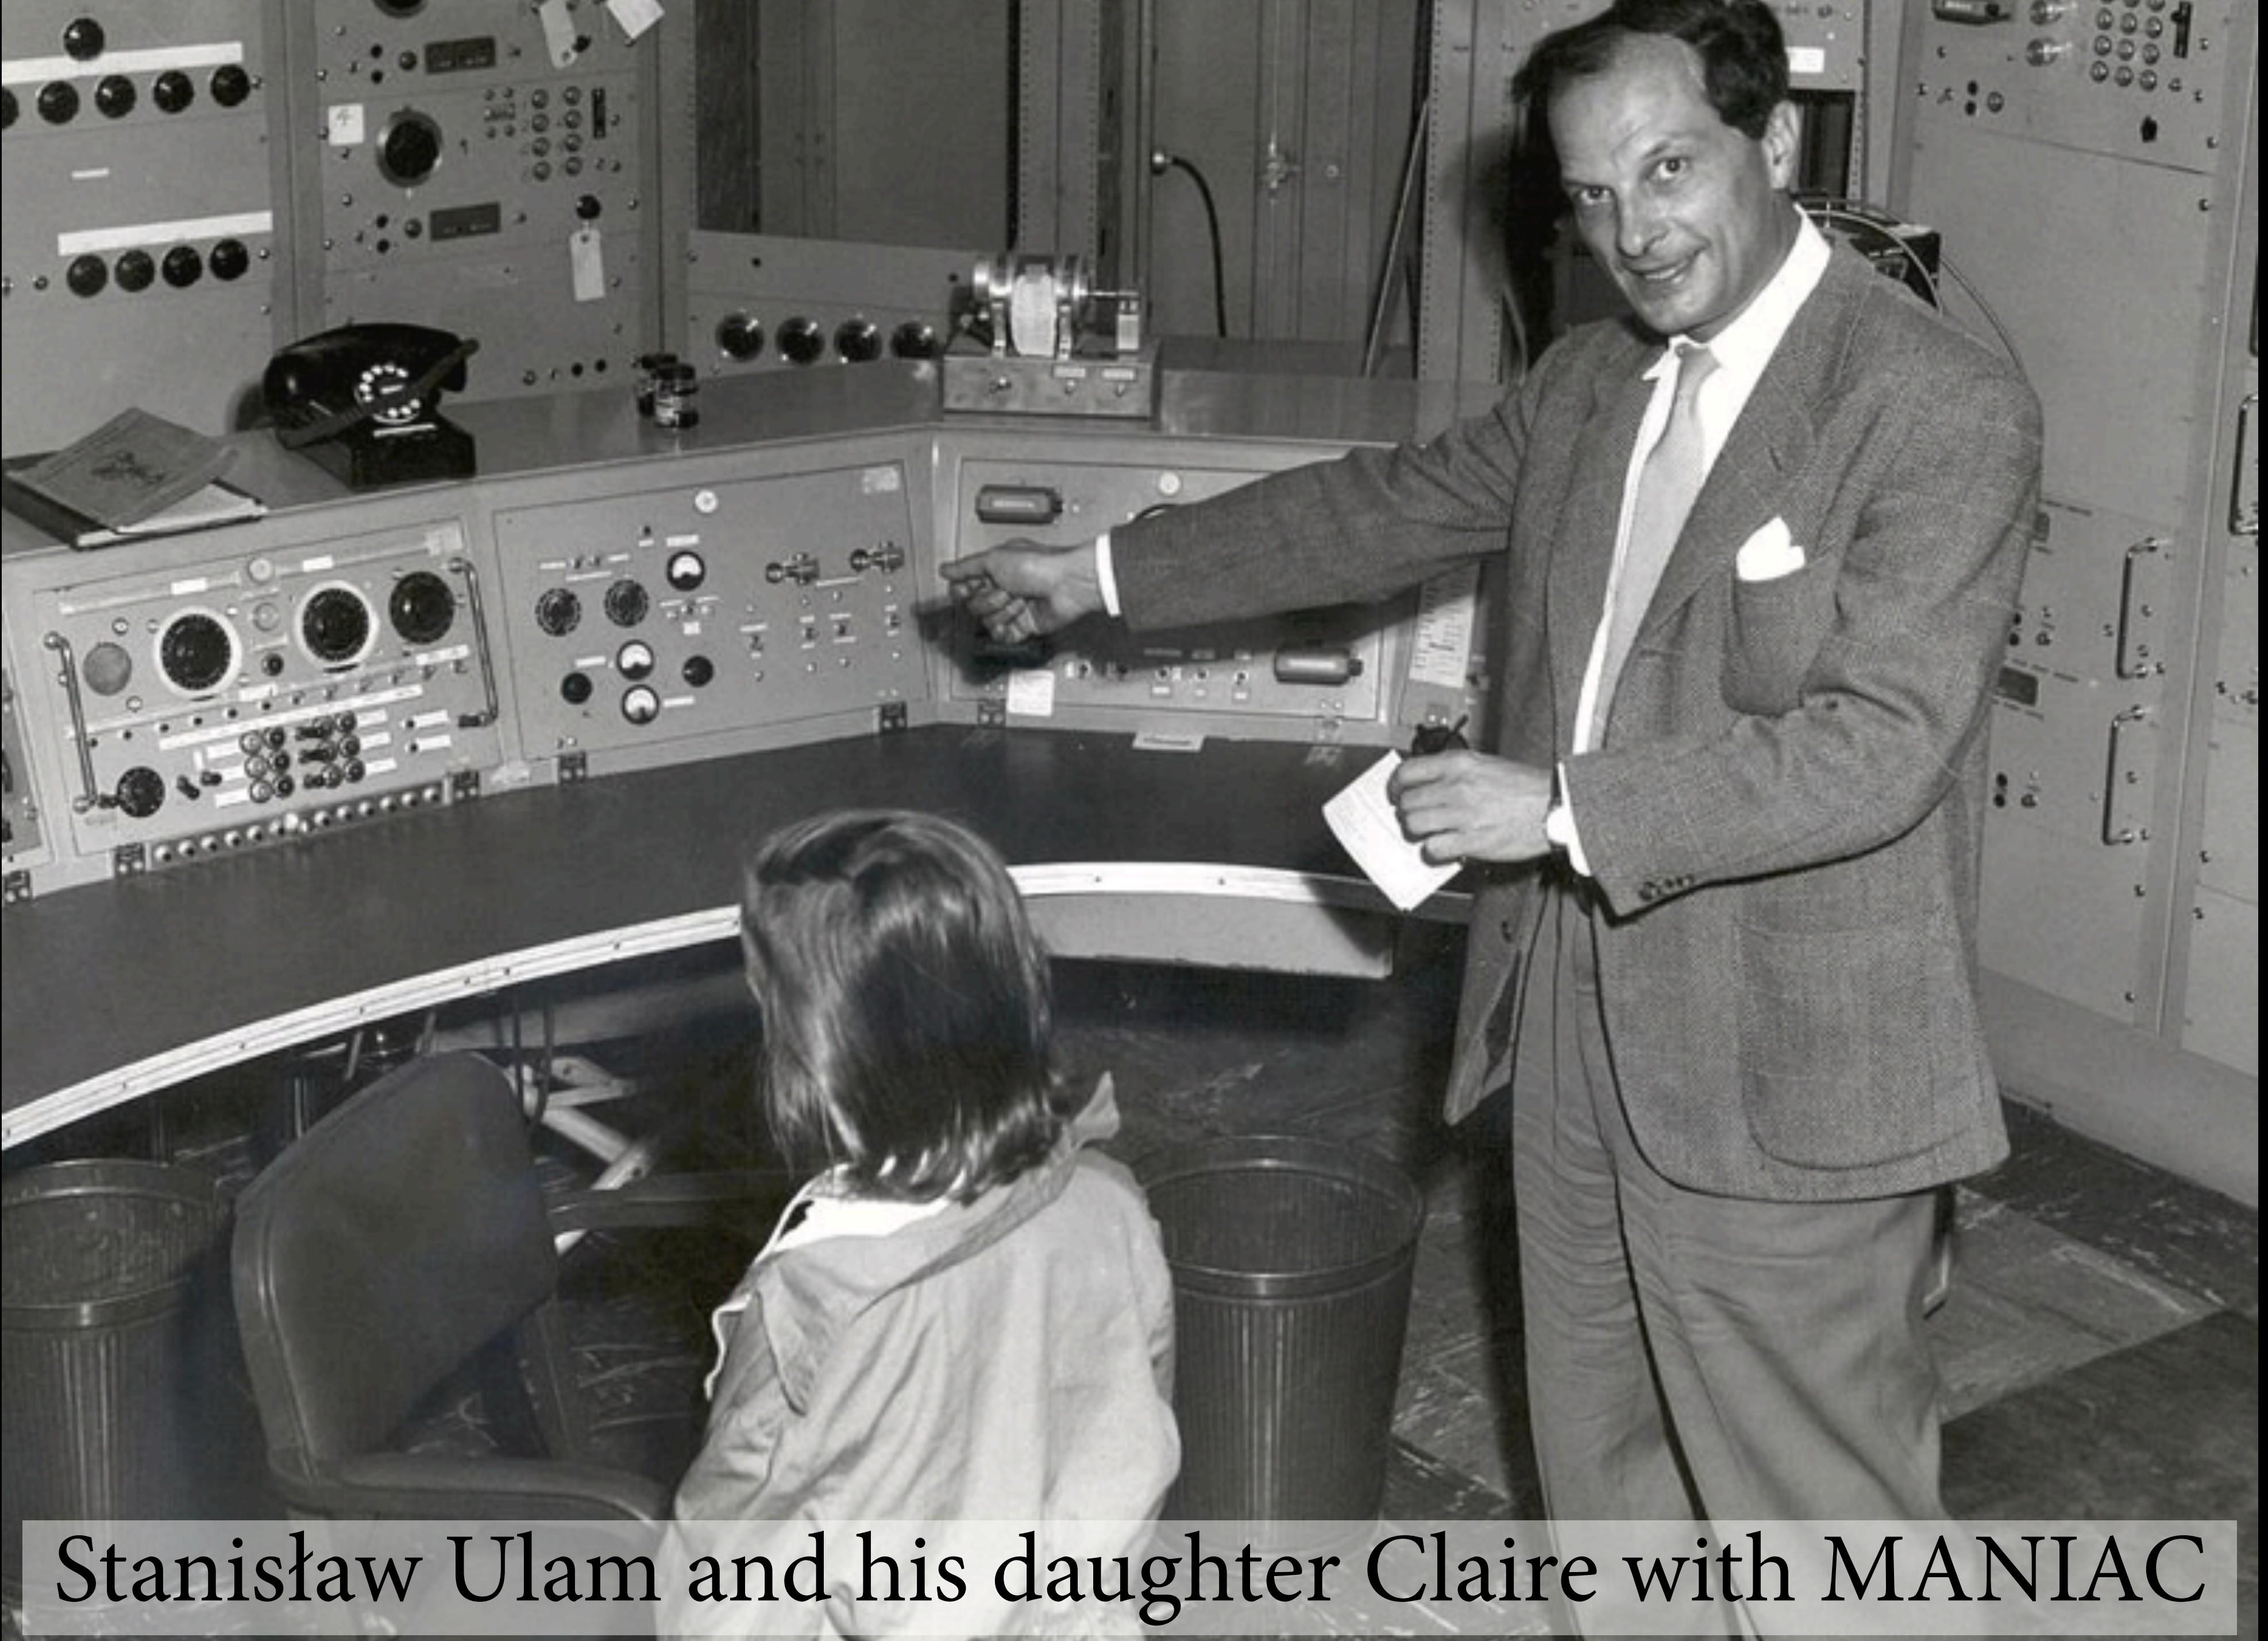
Stan

About Stan

Stan is a state-of-the-art platform for statistical modeling and high-performance statistical computation. Thousands of users rely on Stan for statistical modeling, data analysis, and prediction in the social, biological, and physical sciences, engineering, and business.

Stanisław Ulam (1909–1984)





Stanisław Ulam and his daughter Claire with MANIAC

PAUSE

Example: The Judgment at Princeton



Example: The Judgment at Princeton

```
library(rethinking)  
data(Wines2012)
```

A CENTURY OF NEW JERSEY WINE:

1920

Prohibition
Begins

1933

Prohibition
Ends

1980

7 NJ Wineries
in Operation

1981

NJ Farm
Winery Act

1984

Central Delaware
Valley AVA
Designated

1987

Garden State
Wine Growers
Association Forms

1988

Warren
Hills AVA
Designated

1990

17 NJ Wineries
in Operation

2006

Outer
Coastal AVA
Designated

2008

38 NJ Wineries
in Operation

2012

Judgment at
Princeton

2016

50+ NJ
Wineries in
Operation

2018

NJ Wines Become Exclusive Wine Served
at Governor's Mansion, Drumthwacket;
Cape May Peninsula AVA Designated

2022

60 NJ Wineries
in Operation

Example: The Judgment at Princeton

```
library(rethinking)  
data(Wines2012)
```

A CENTURY OF NEW JERSEY WINE:

1920
Prohibition
Begins

1933
Prohibition
Ends

1980
7 NJ Wineries
in Operation

1988
NJ First
Wine

2006
Outer
Coastal AVA
Designated

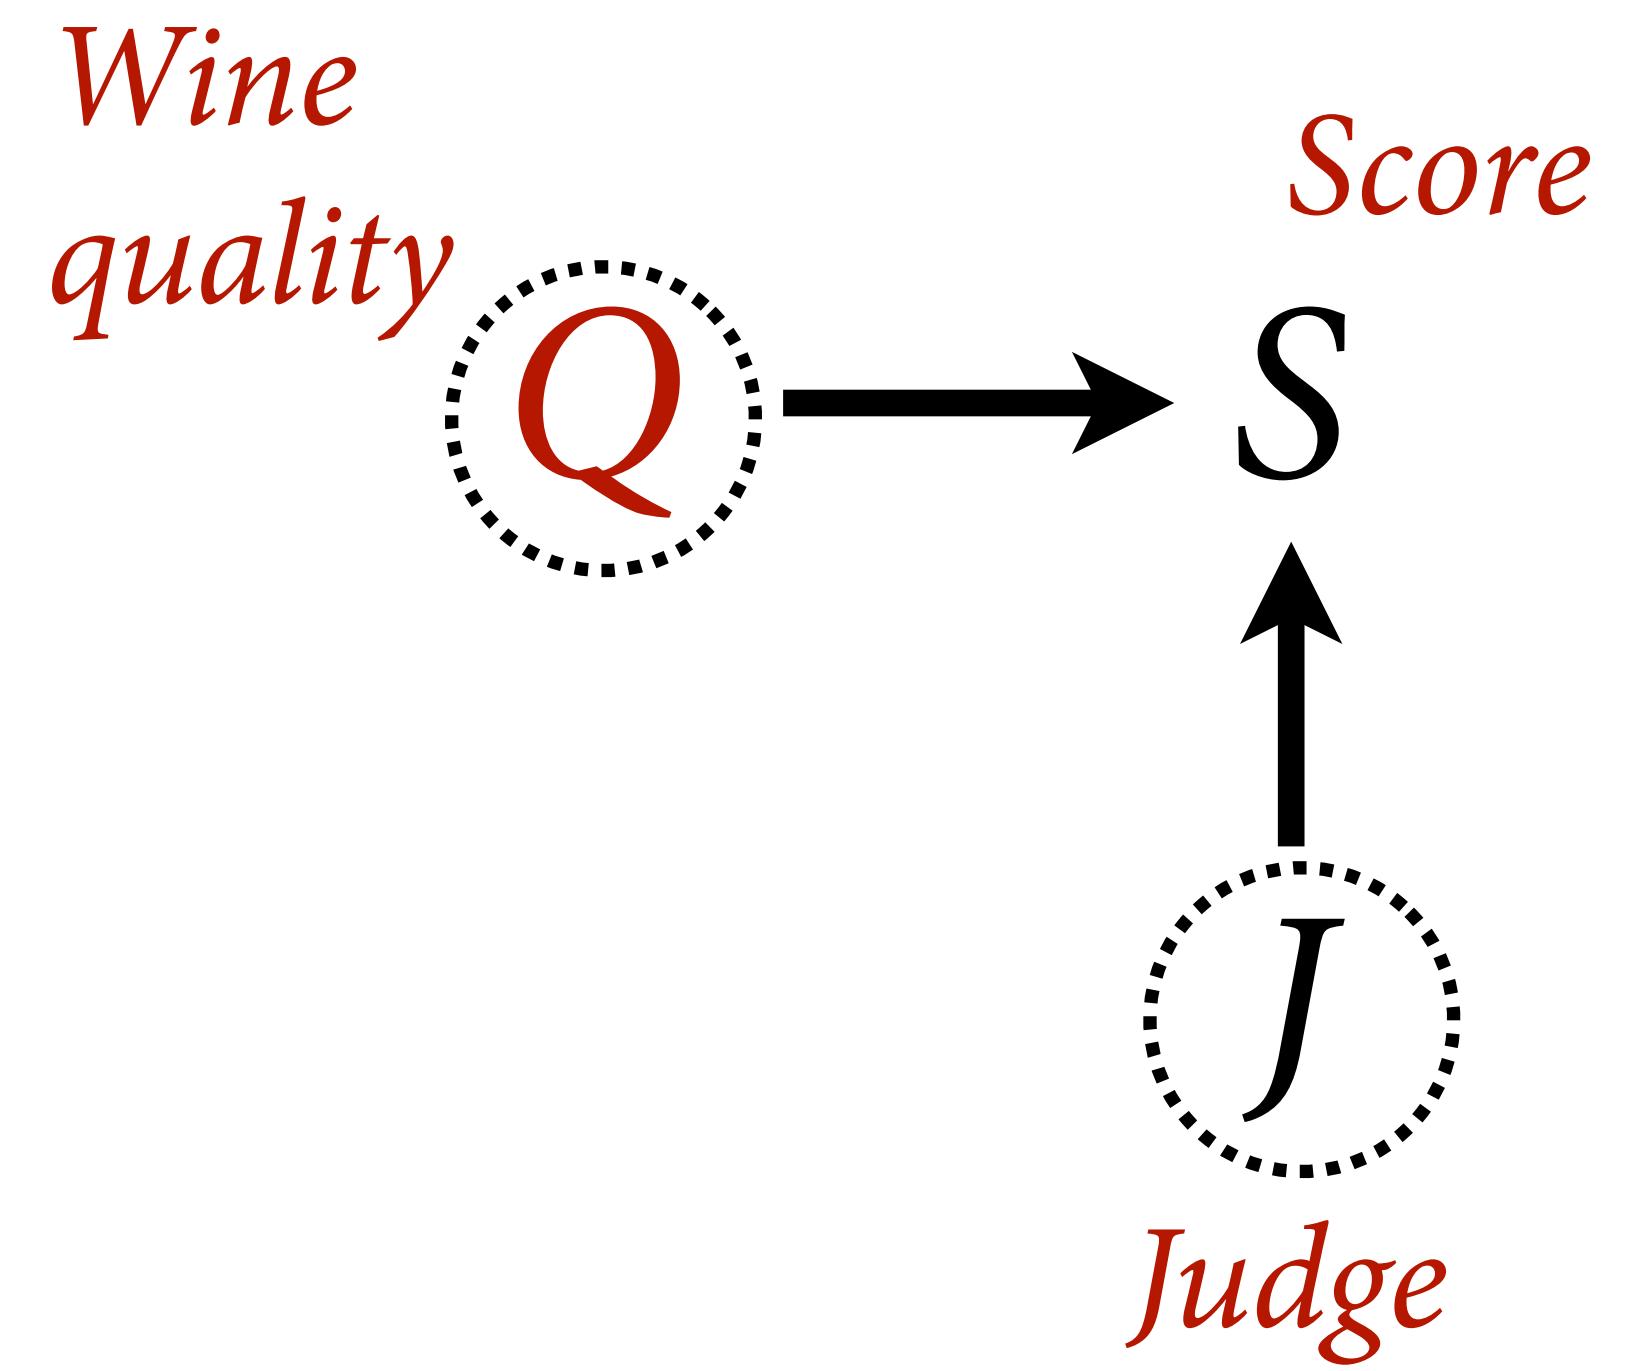
2008
38 NJ Wineries
in Operation

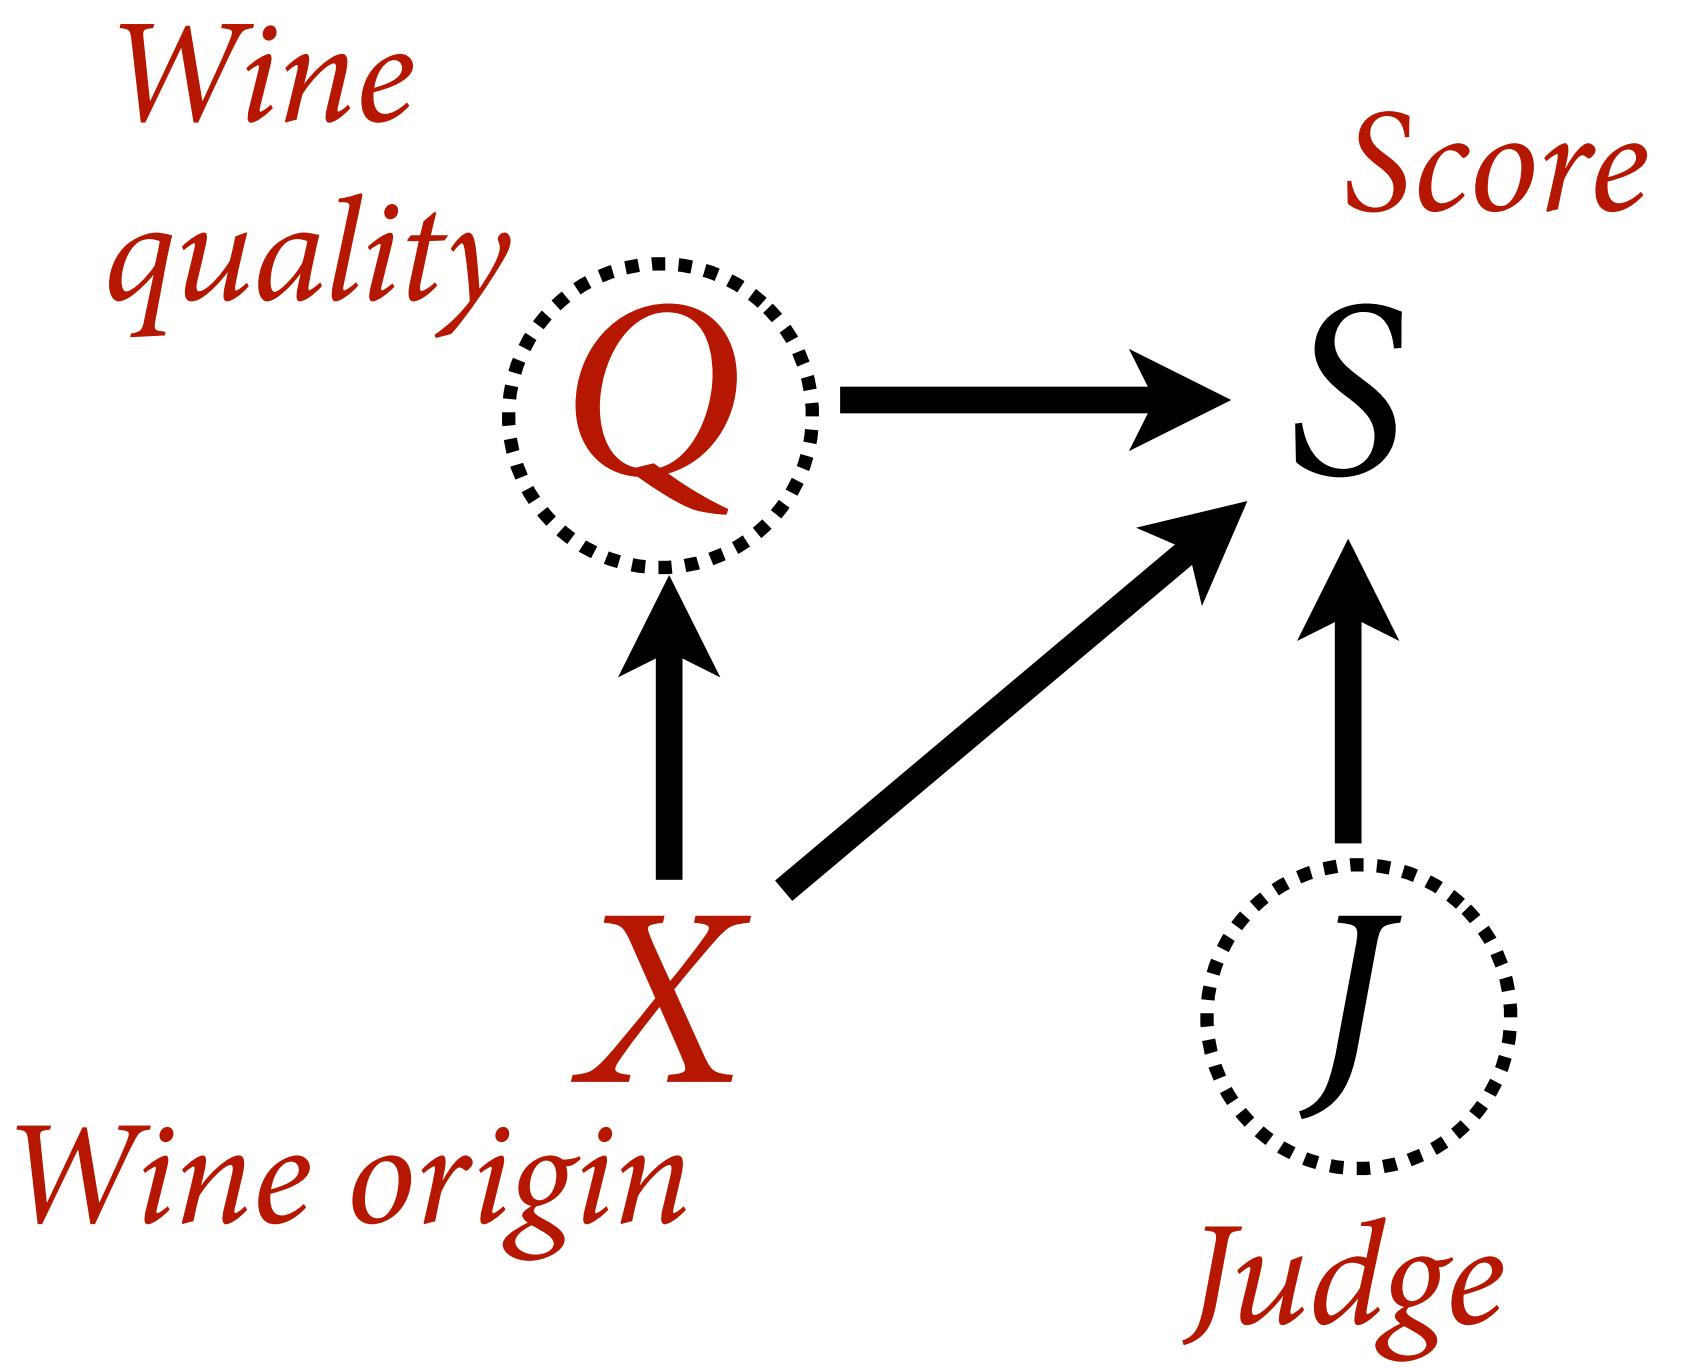
2012
Judgment at
Princeton

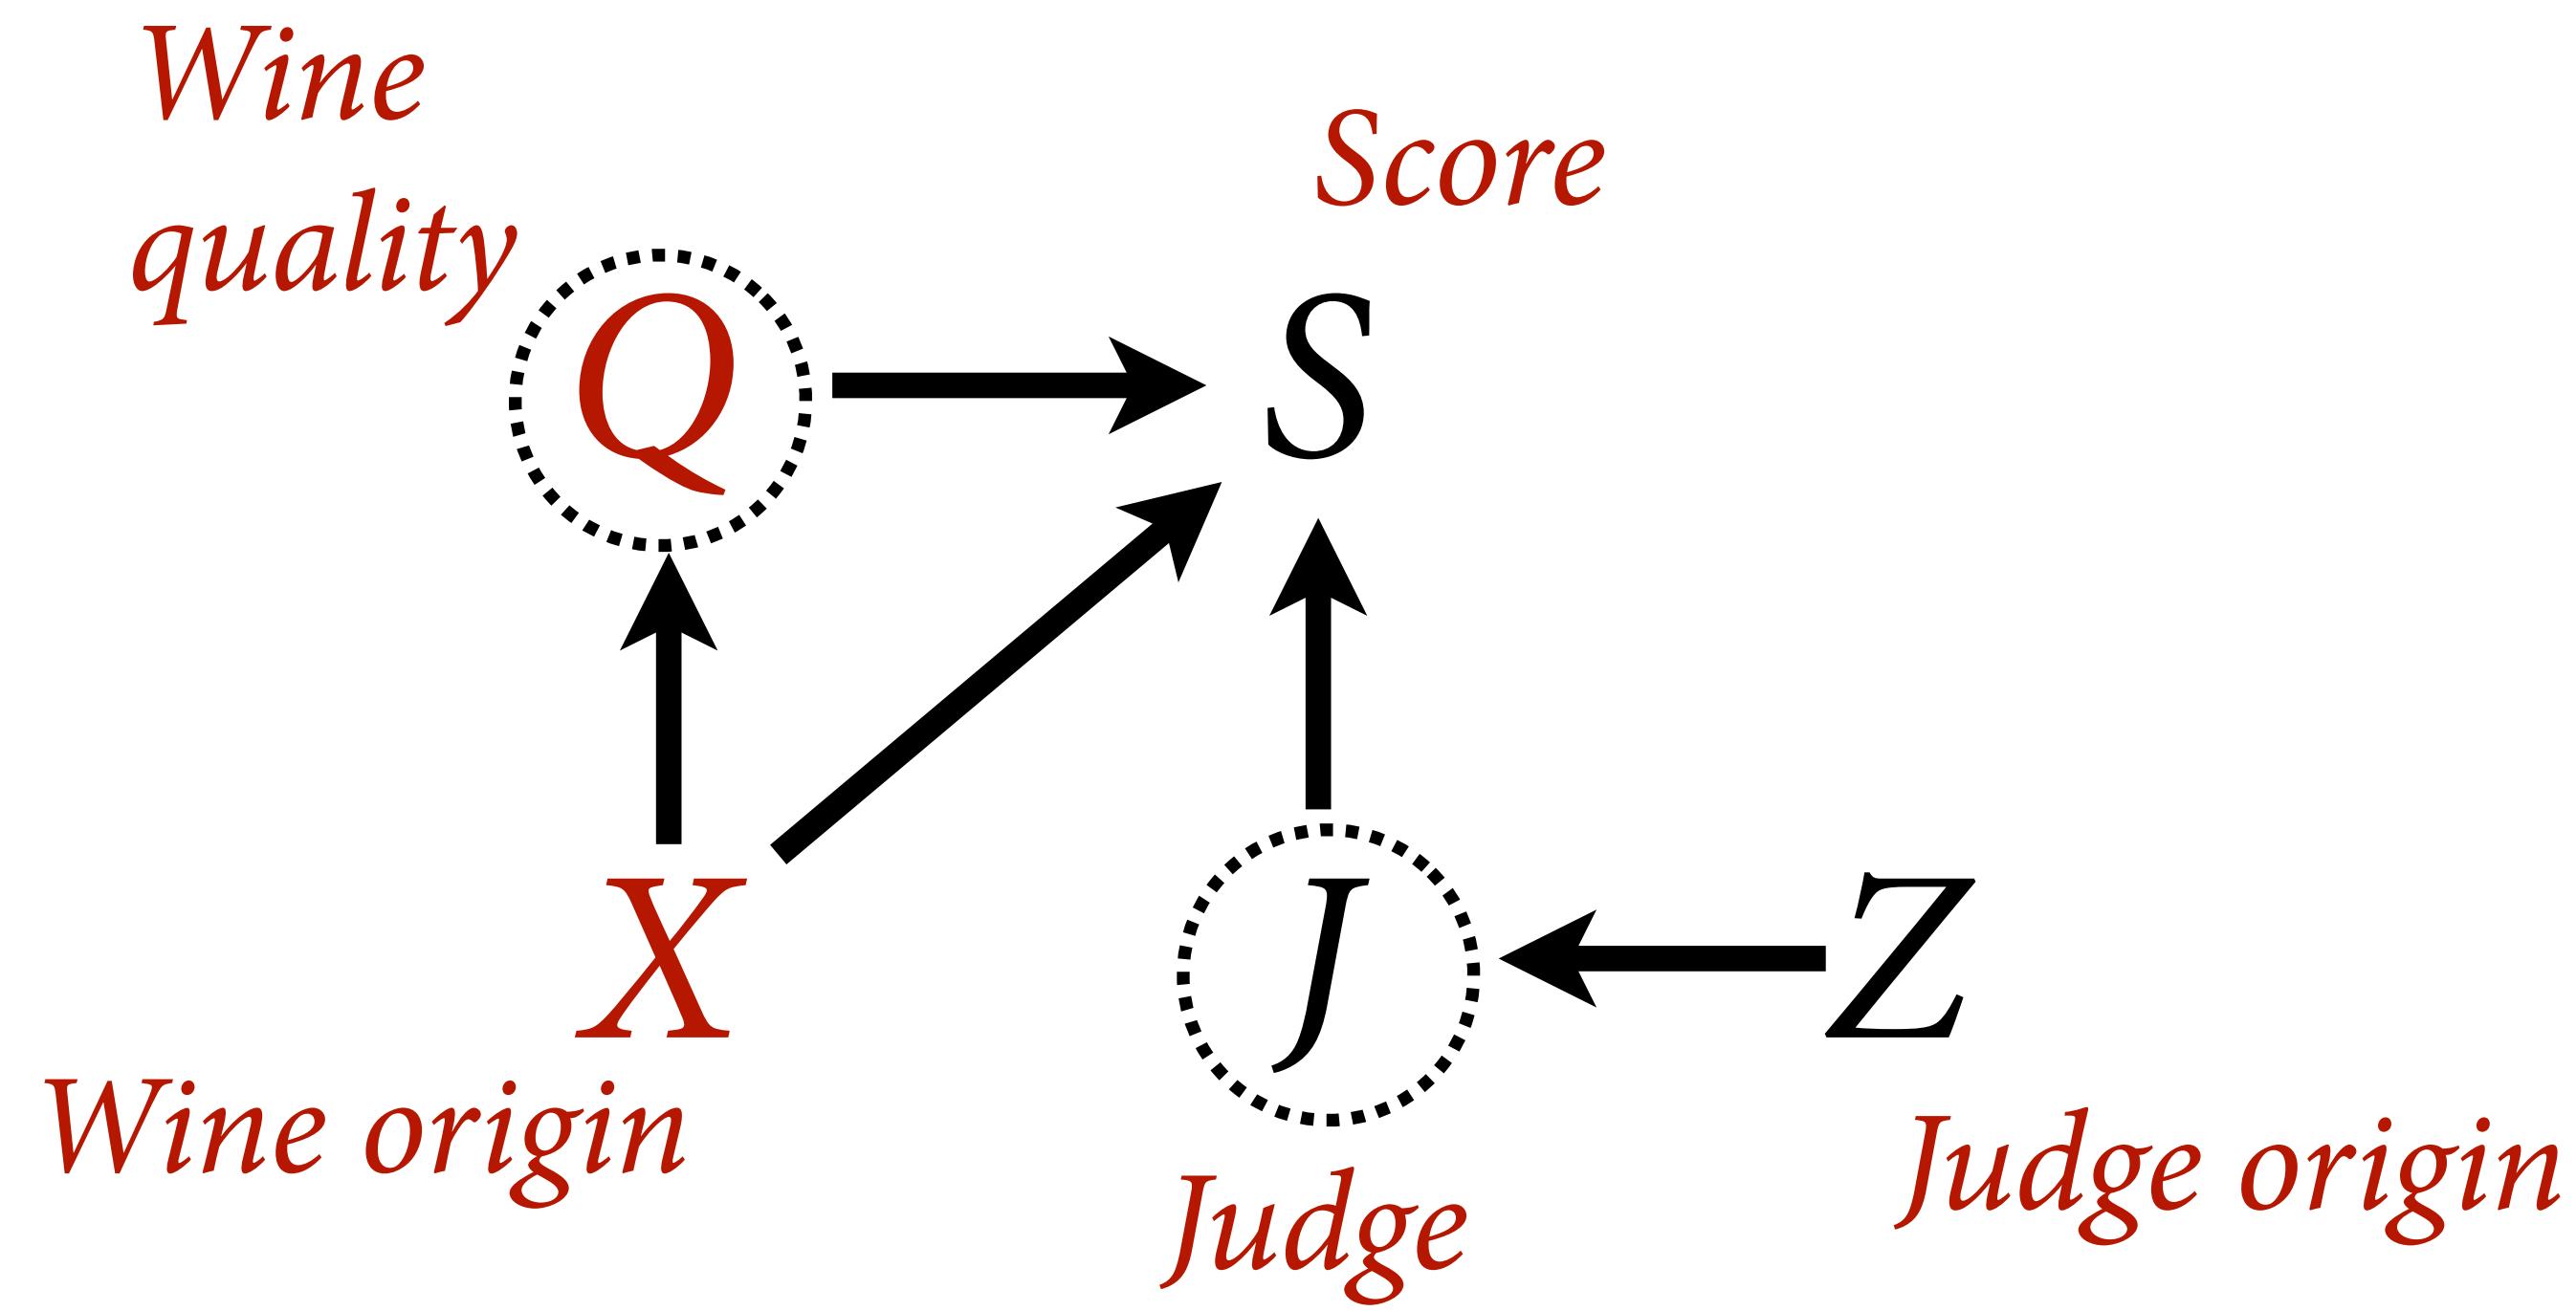
20 wines (10 French, 10 NJ)

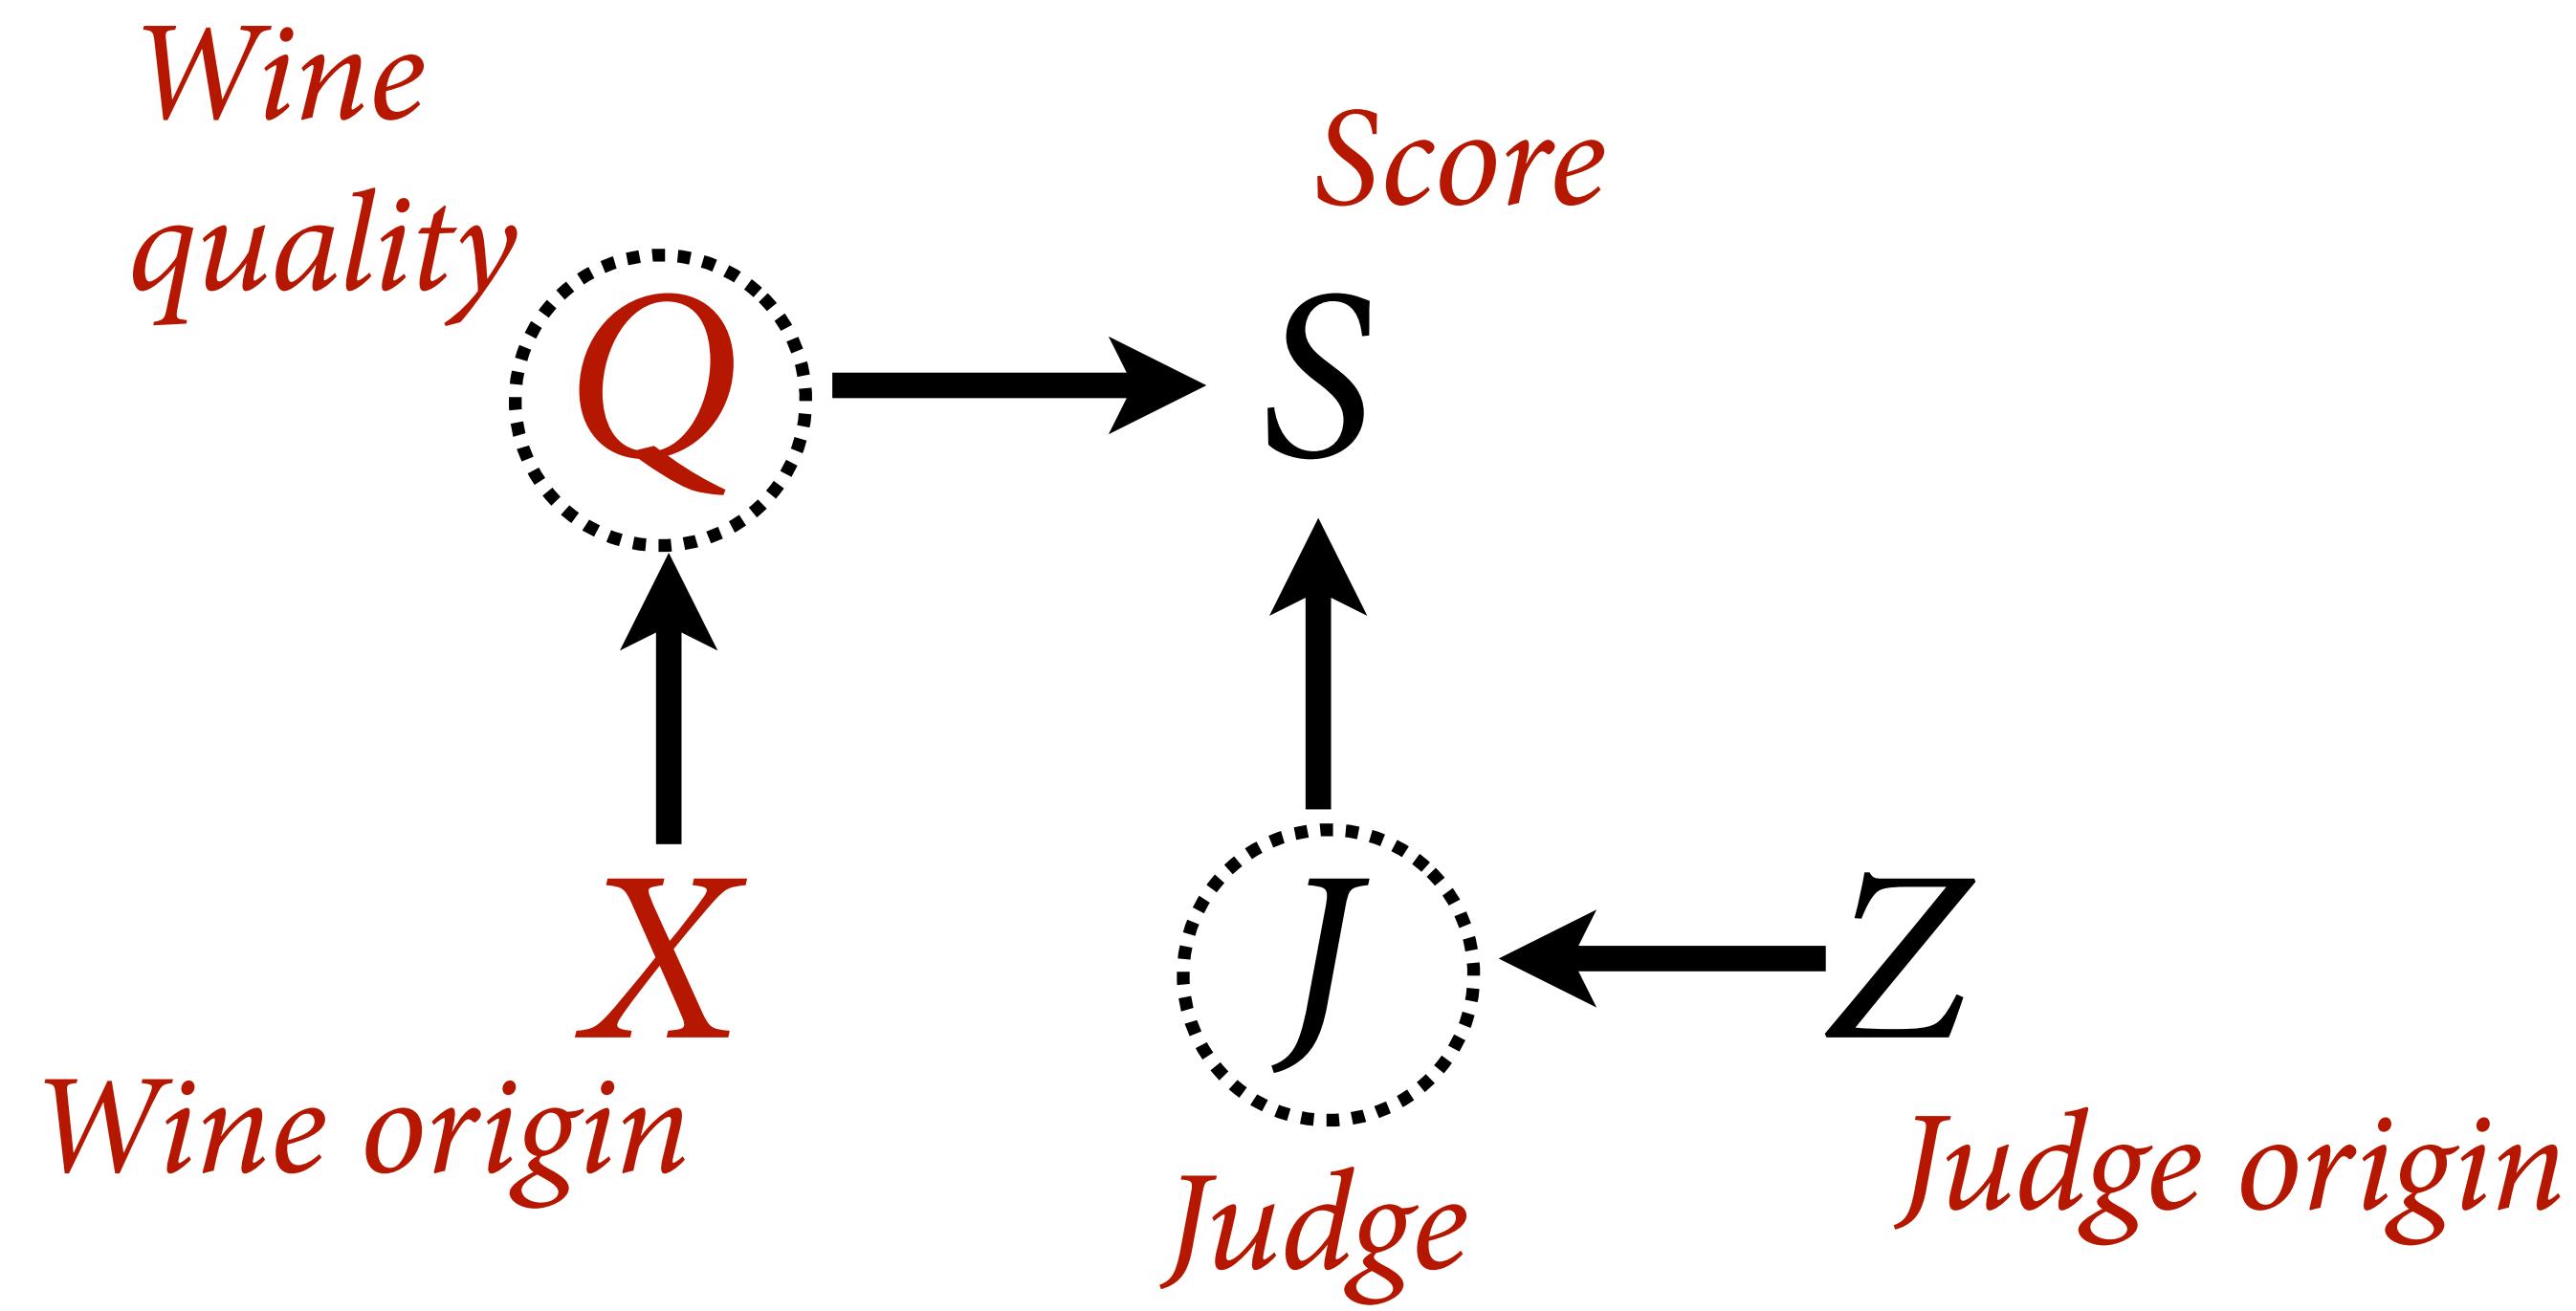
9 French & American judges

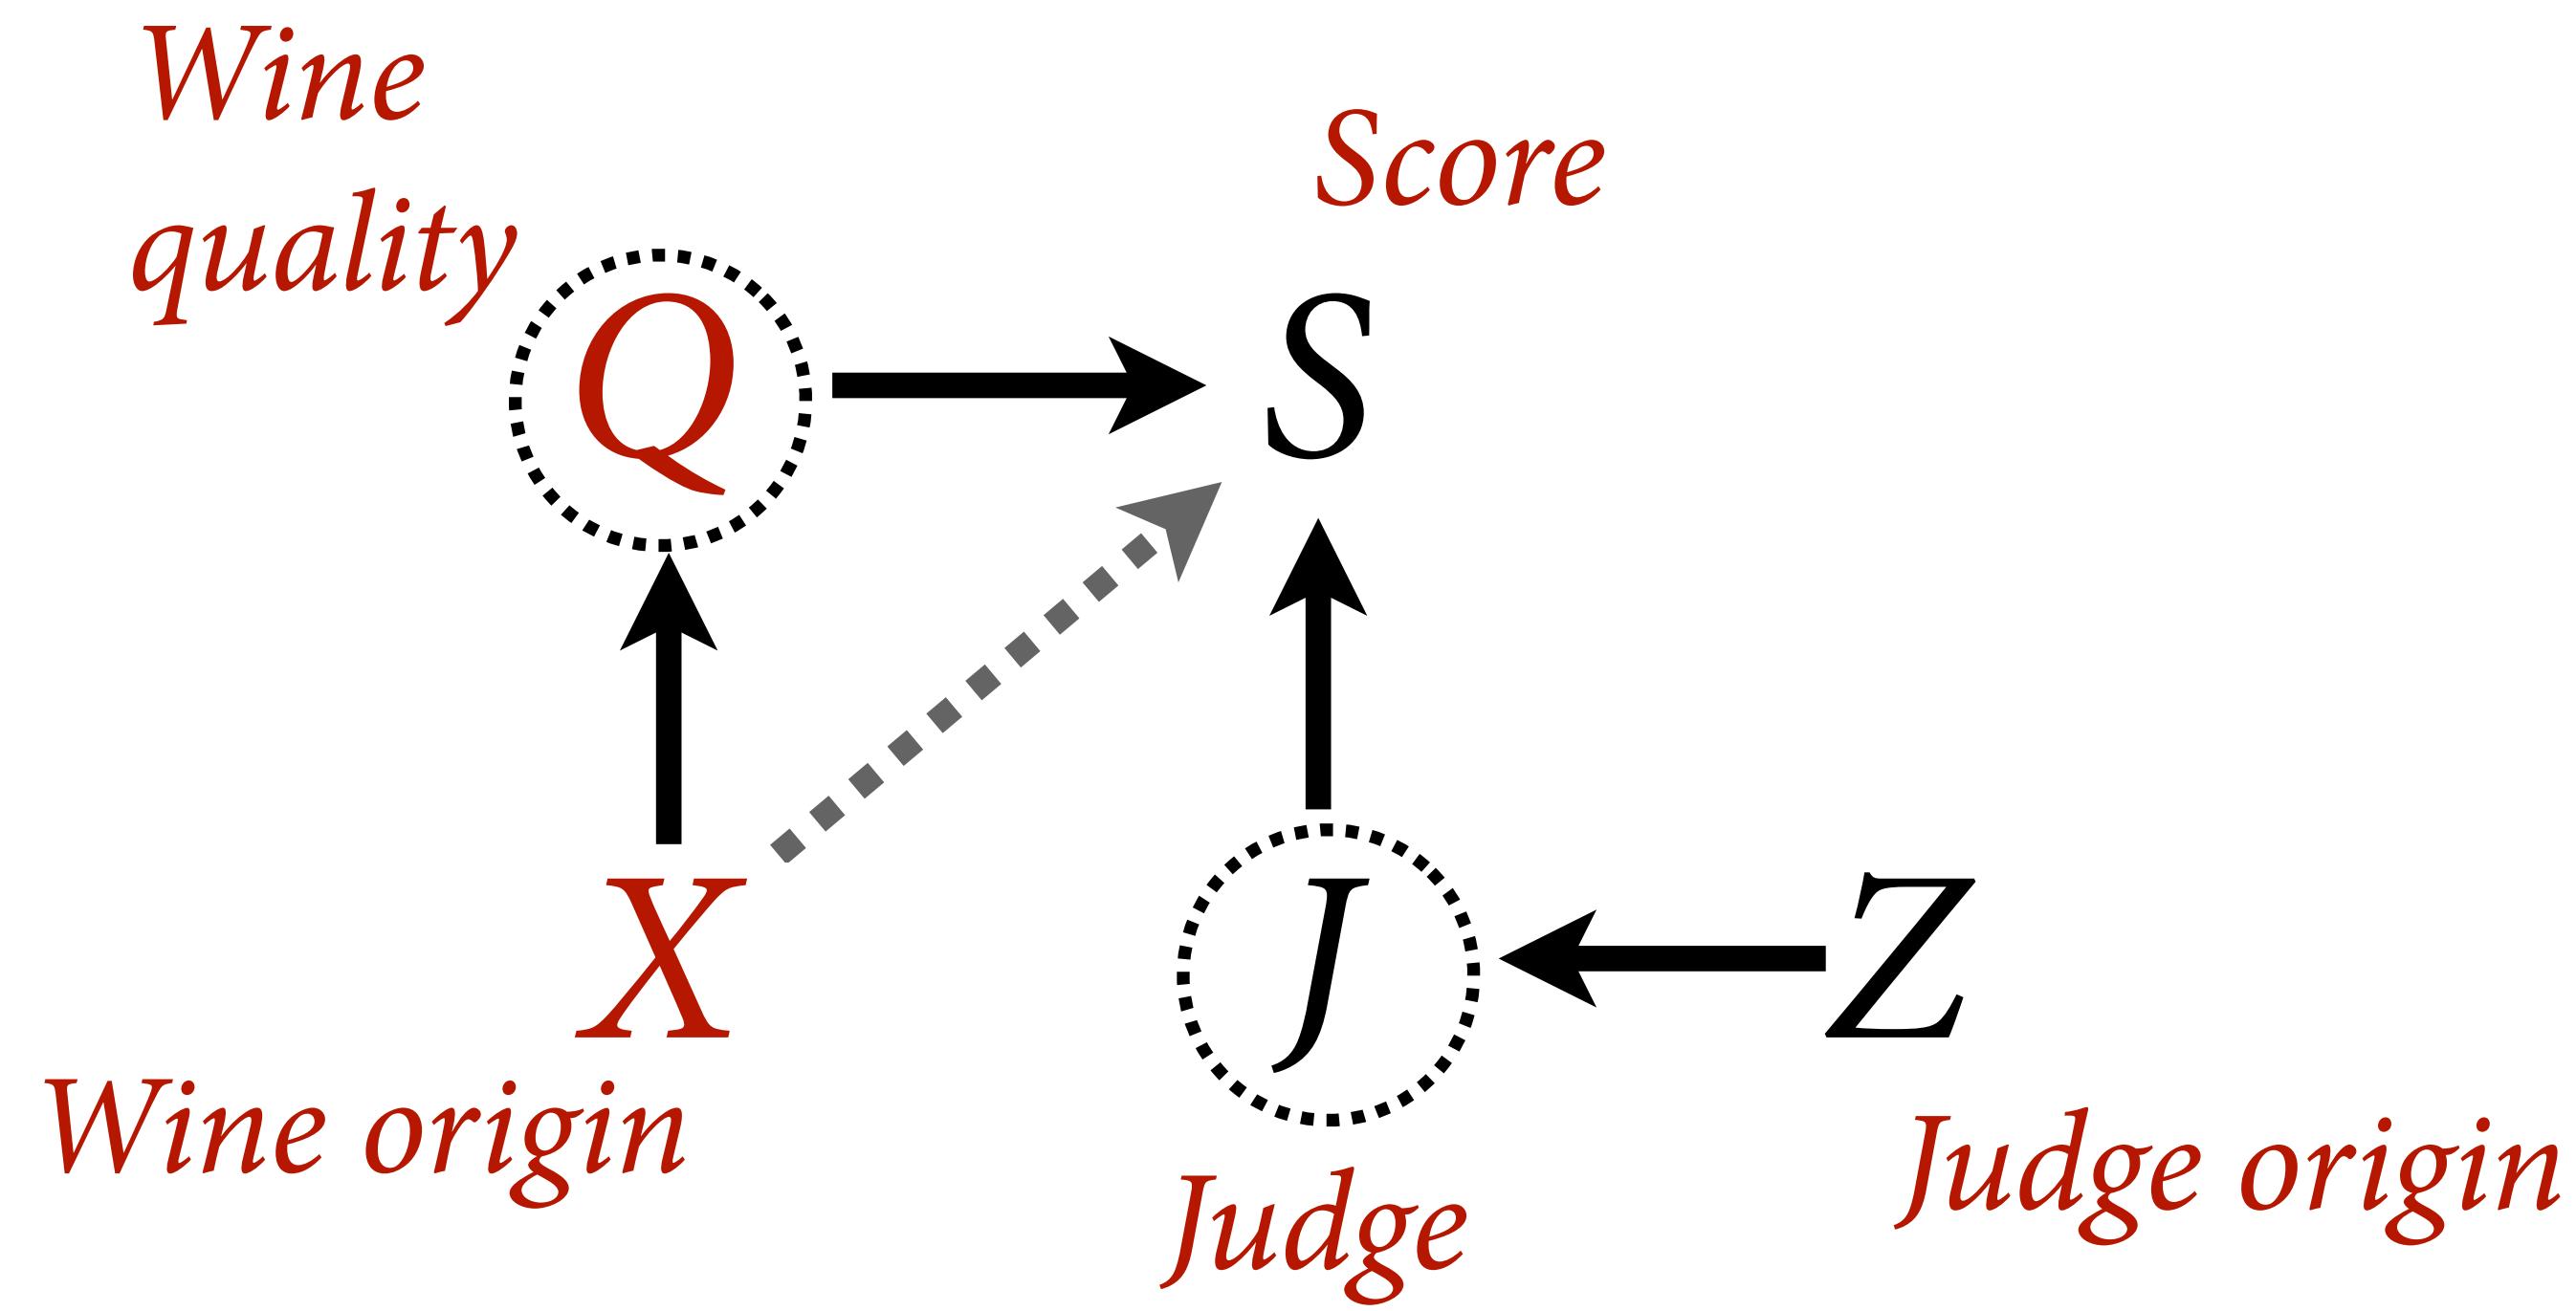
180 scores

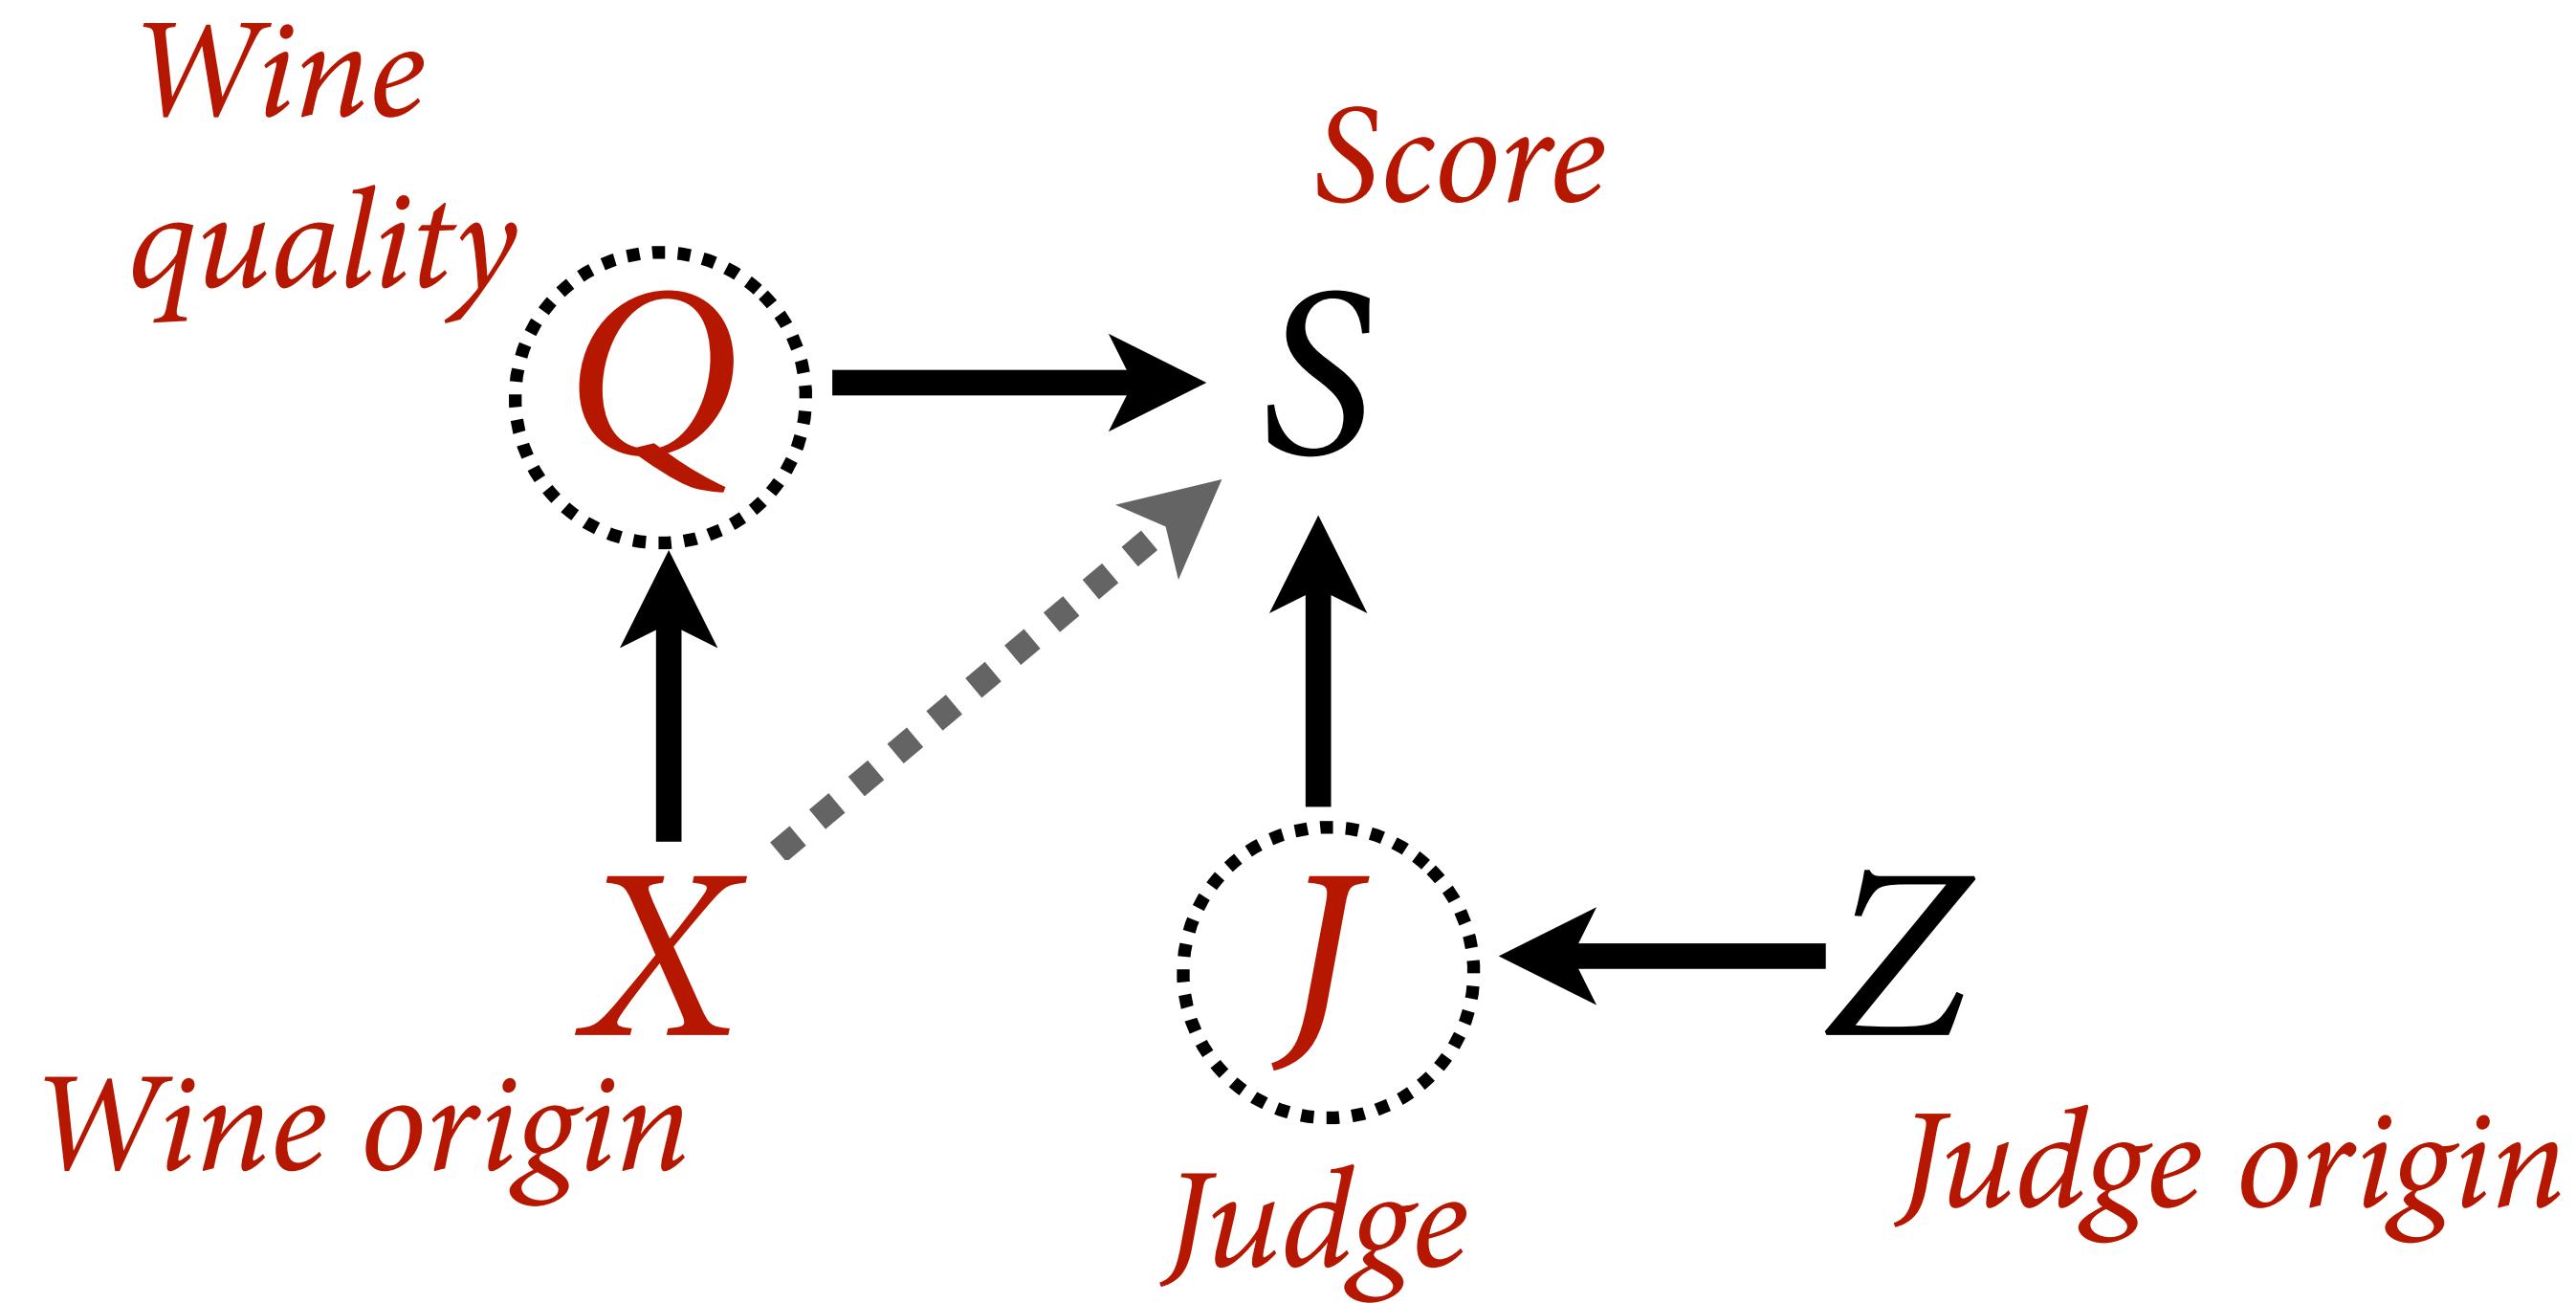












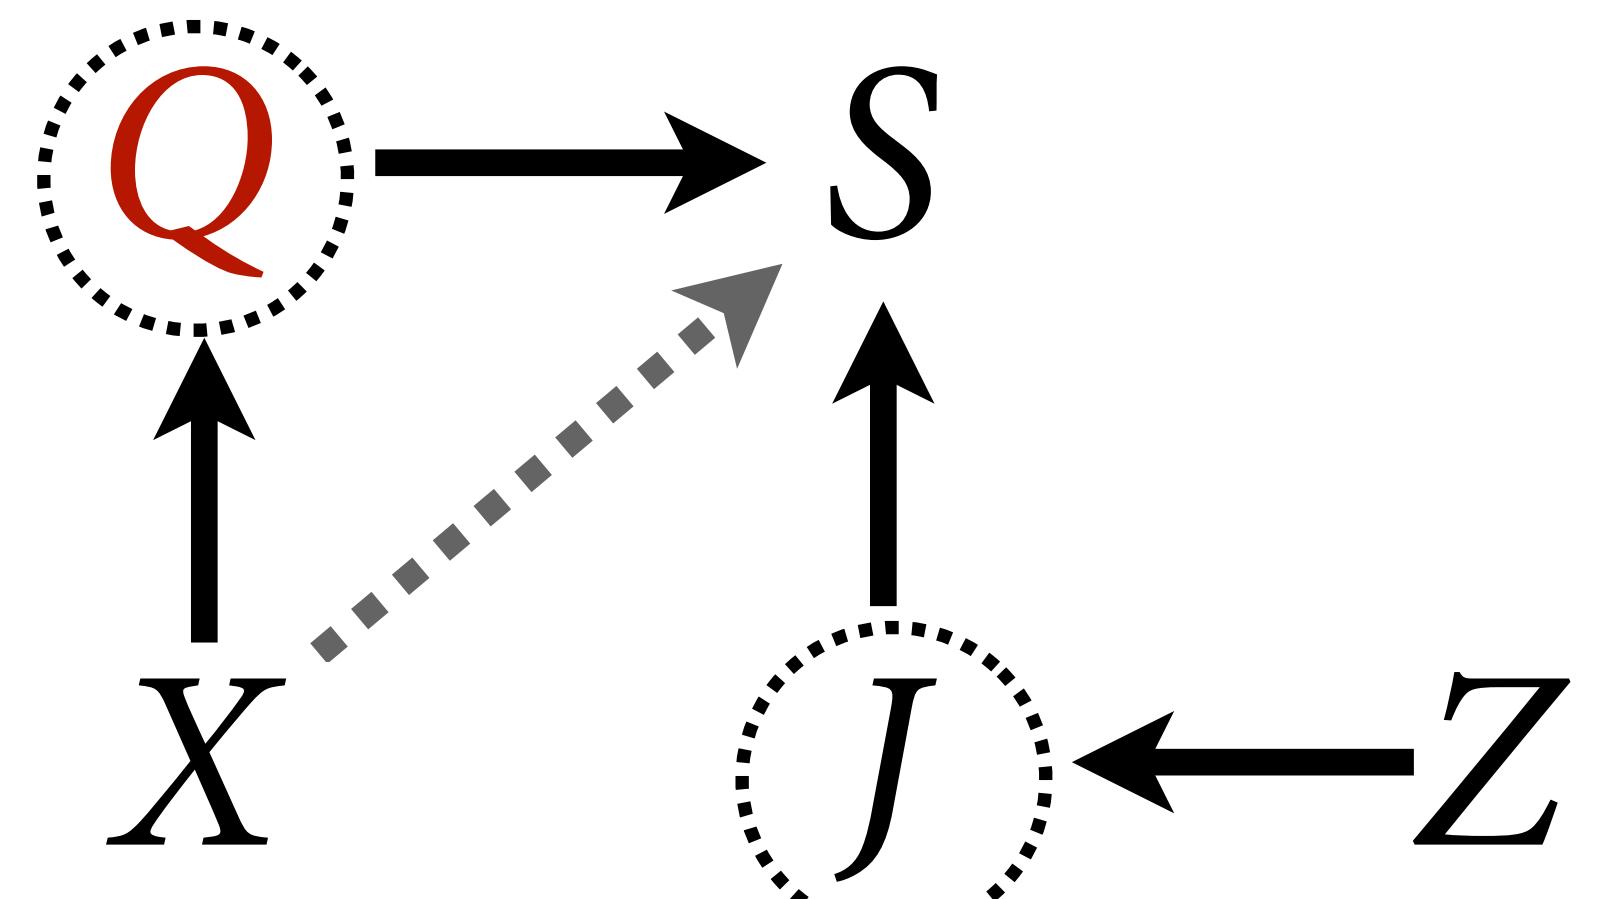
Estimand: Association between wine **quality** and wine **origin**. Stratify by **judge** for efficiency.

```

library(rethinking)
data(Wines2012)
d <- Wines2012

dat <- list(
  S=standardize(d$sscore),
  J=as.numeric(d$judge),
  W=as.numeric(d$wine),
  X=ifelse(d$wine.amer==1,1,2),
  Z=ifelse(d$judge.amer==1,1,2)
)
mQ <- ulam(
  alist(
    S ~ dnorm(mu,sigma),
    mu <- Q[W],
    Q[W] ~ dnorm(0,1),
    sigma ~ dexp(1)
  ) , data=dat , chains=4 , cores=4
)
precis(mQ,2)

```



$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = Q_{W[i]}$$

$$Q_j \sim \text{Normal}(0, 1)$$

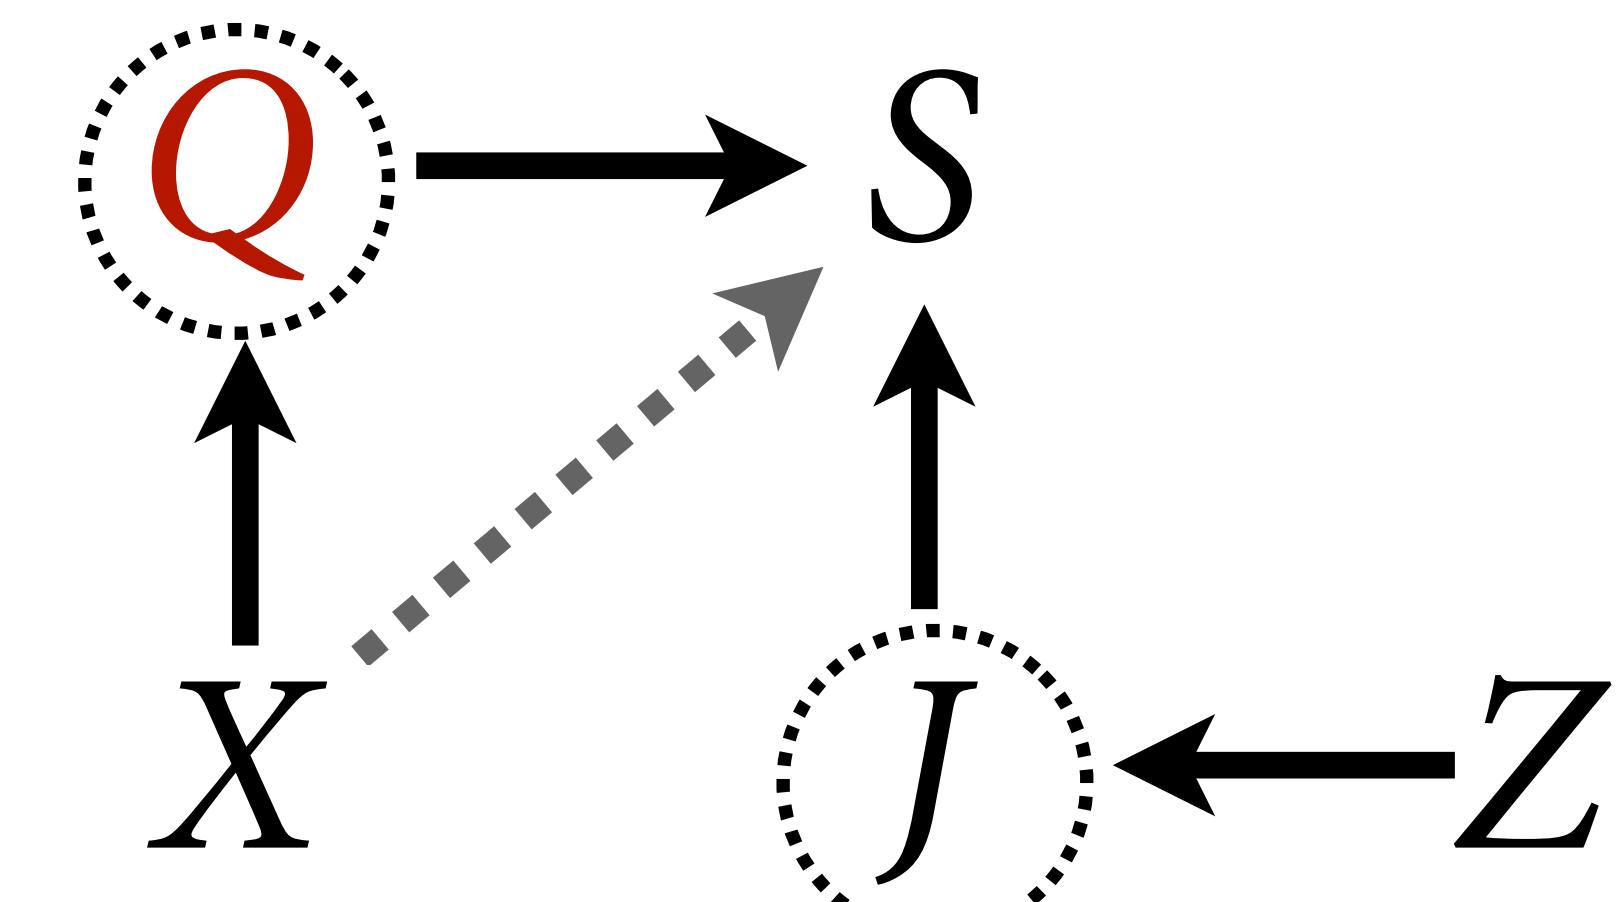
$$\sigma \sim \text{Exponential}(1)$$

```

library(rethinking)
data(Wines2
d <- Wines2
dat <- list(
  S=stand(d),
  J=as.numeric(as.factor(d$Type)),
  W=as.numeric(as.factor(d$Wine))
)
X=ifelse(J==1, 1, 0)
Z=ifelse(J==2, 1, 0)
)
mQ <- ulam(
  alist(
    S ~ mu + sig * Z,
    mu ~ Q[W],
    sig ~ Q[w]
  ) , dat
)
precis(mQ, 2)

```

	mean	sd	5.5%	94.5%	n_eff	Rhat4
Q[1]	0.14	0.30	-0.34	0.64	2962	1
Q[2]	0.11	0.32	-0.40	0.61	3033	1
Q[3]	0.27	0.31	-0.21	0.75	2608	1
Q[4]	0.55	0.33	0.04	1.09	2598	1
Q[5]	-0.13	0.32	-0.63	0.37	2726	1
Q[6]	-0.37	0.32	-0.88	0.12	2710	1
Q[7]	0.29	0.32	-0.22	0.81	2679	1
Q[8]	0.27	0.32	-0.24	0.79	3248	1
Q[9]	0.08	0.30	-0.40	0.56	3410	1
Q[10]	0.12	0.32	-0.40	0.63	2697	1
Q[11]	-0.02	0.31	-0.52	0.50	2740	1
Q[12]	-0.03	0.32	-0.53	0.48	2809	1
Q[13]	-0.11	0.32	-0.62	0.39	3225	1
Q[14]	0.01	0.33	-0.51	0.52	2690	1
Q[15]	-0.22	0.31	-0.71	0.28	2754	1
Q[16]	-0.21	0.32	-0.72	0.30	2120	1
Q[17]	-0.14	0.31	-0.64	0.37	3831	1
Q[18]	-0.86	0.32	-1.37	-0.35	3093	1
Q[19]	-0.16	0.30	-0.65	0.31	2783	1
Q[20]	0.38	0.32	-0.12	0.88	2873	1
sigma	1.00	0.06	0.91	1.09	2759	1



$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = Q_{W[i]}$$

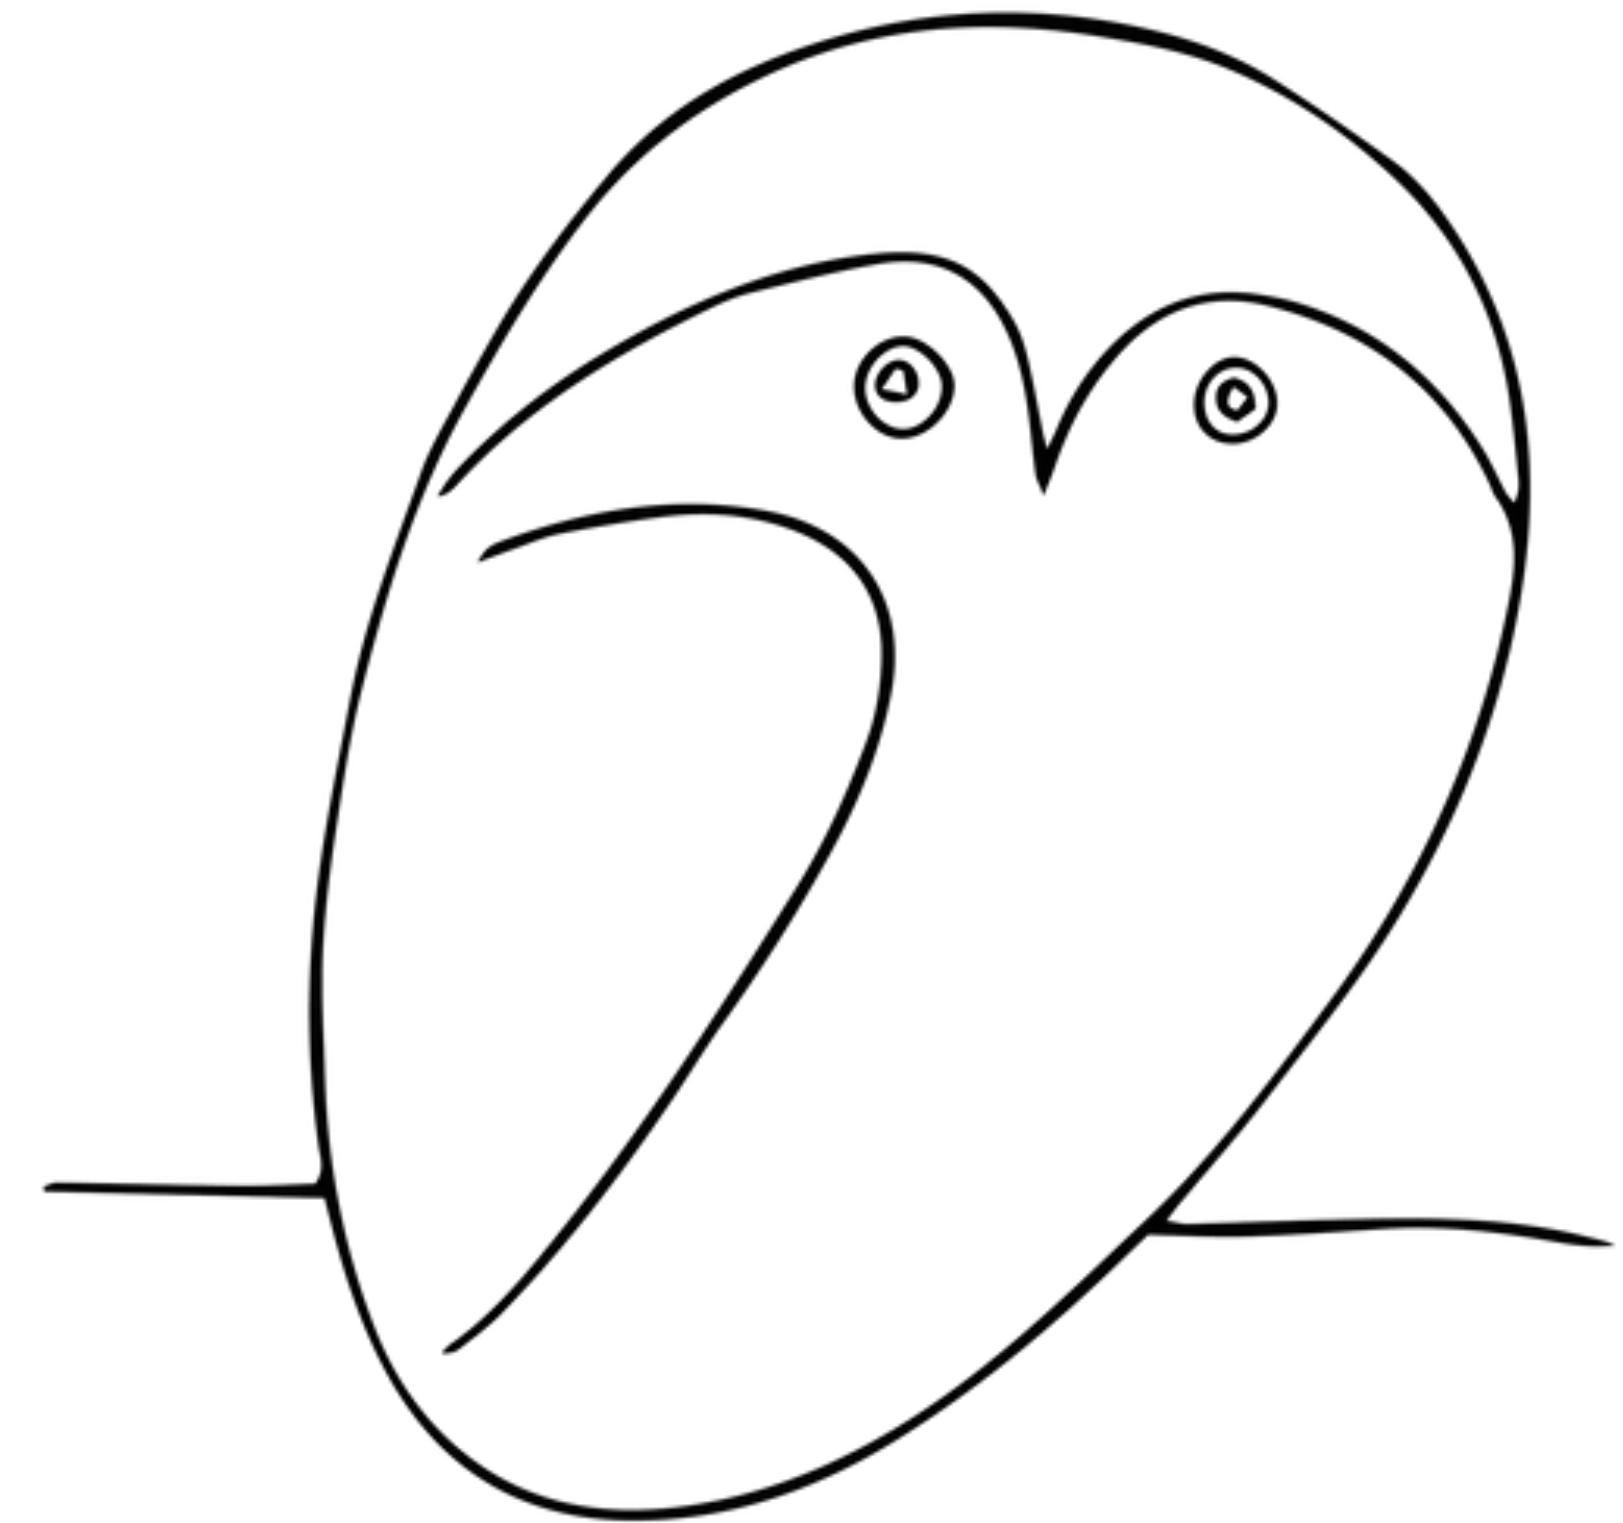
$$Q_j \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$

Drawing the Markov Owl

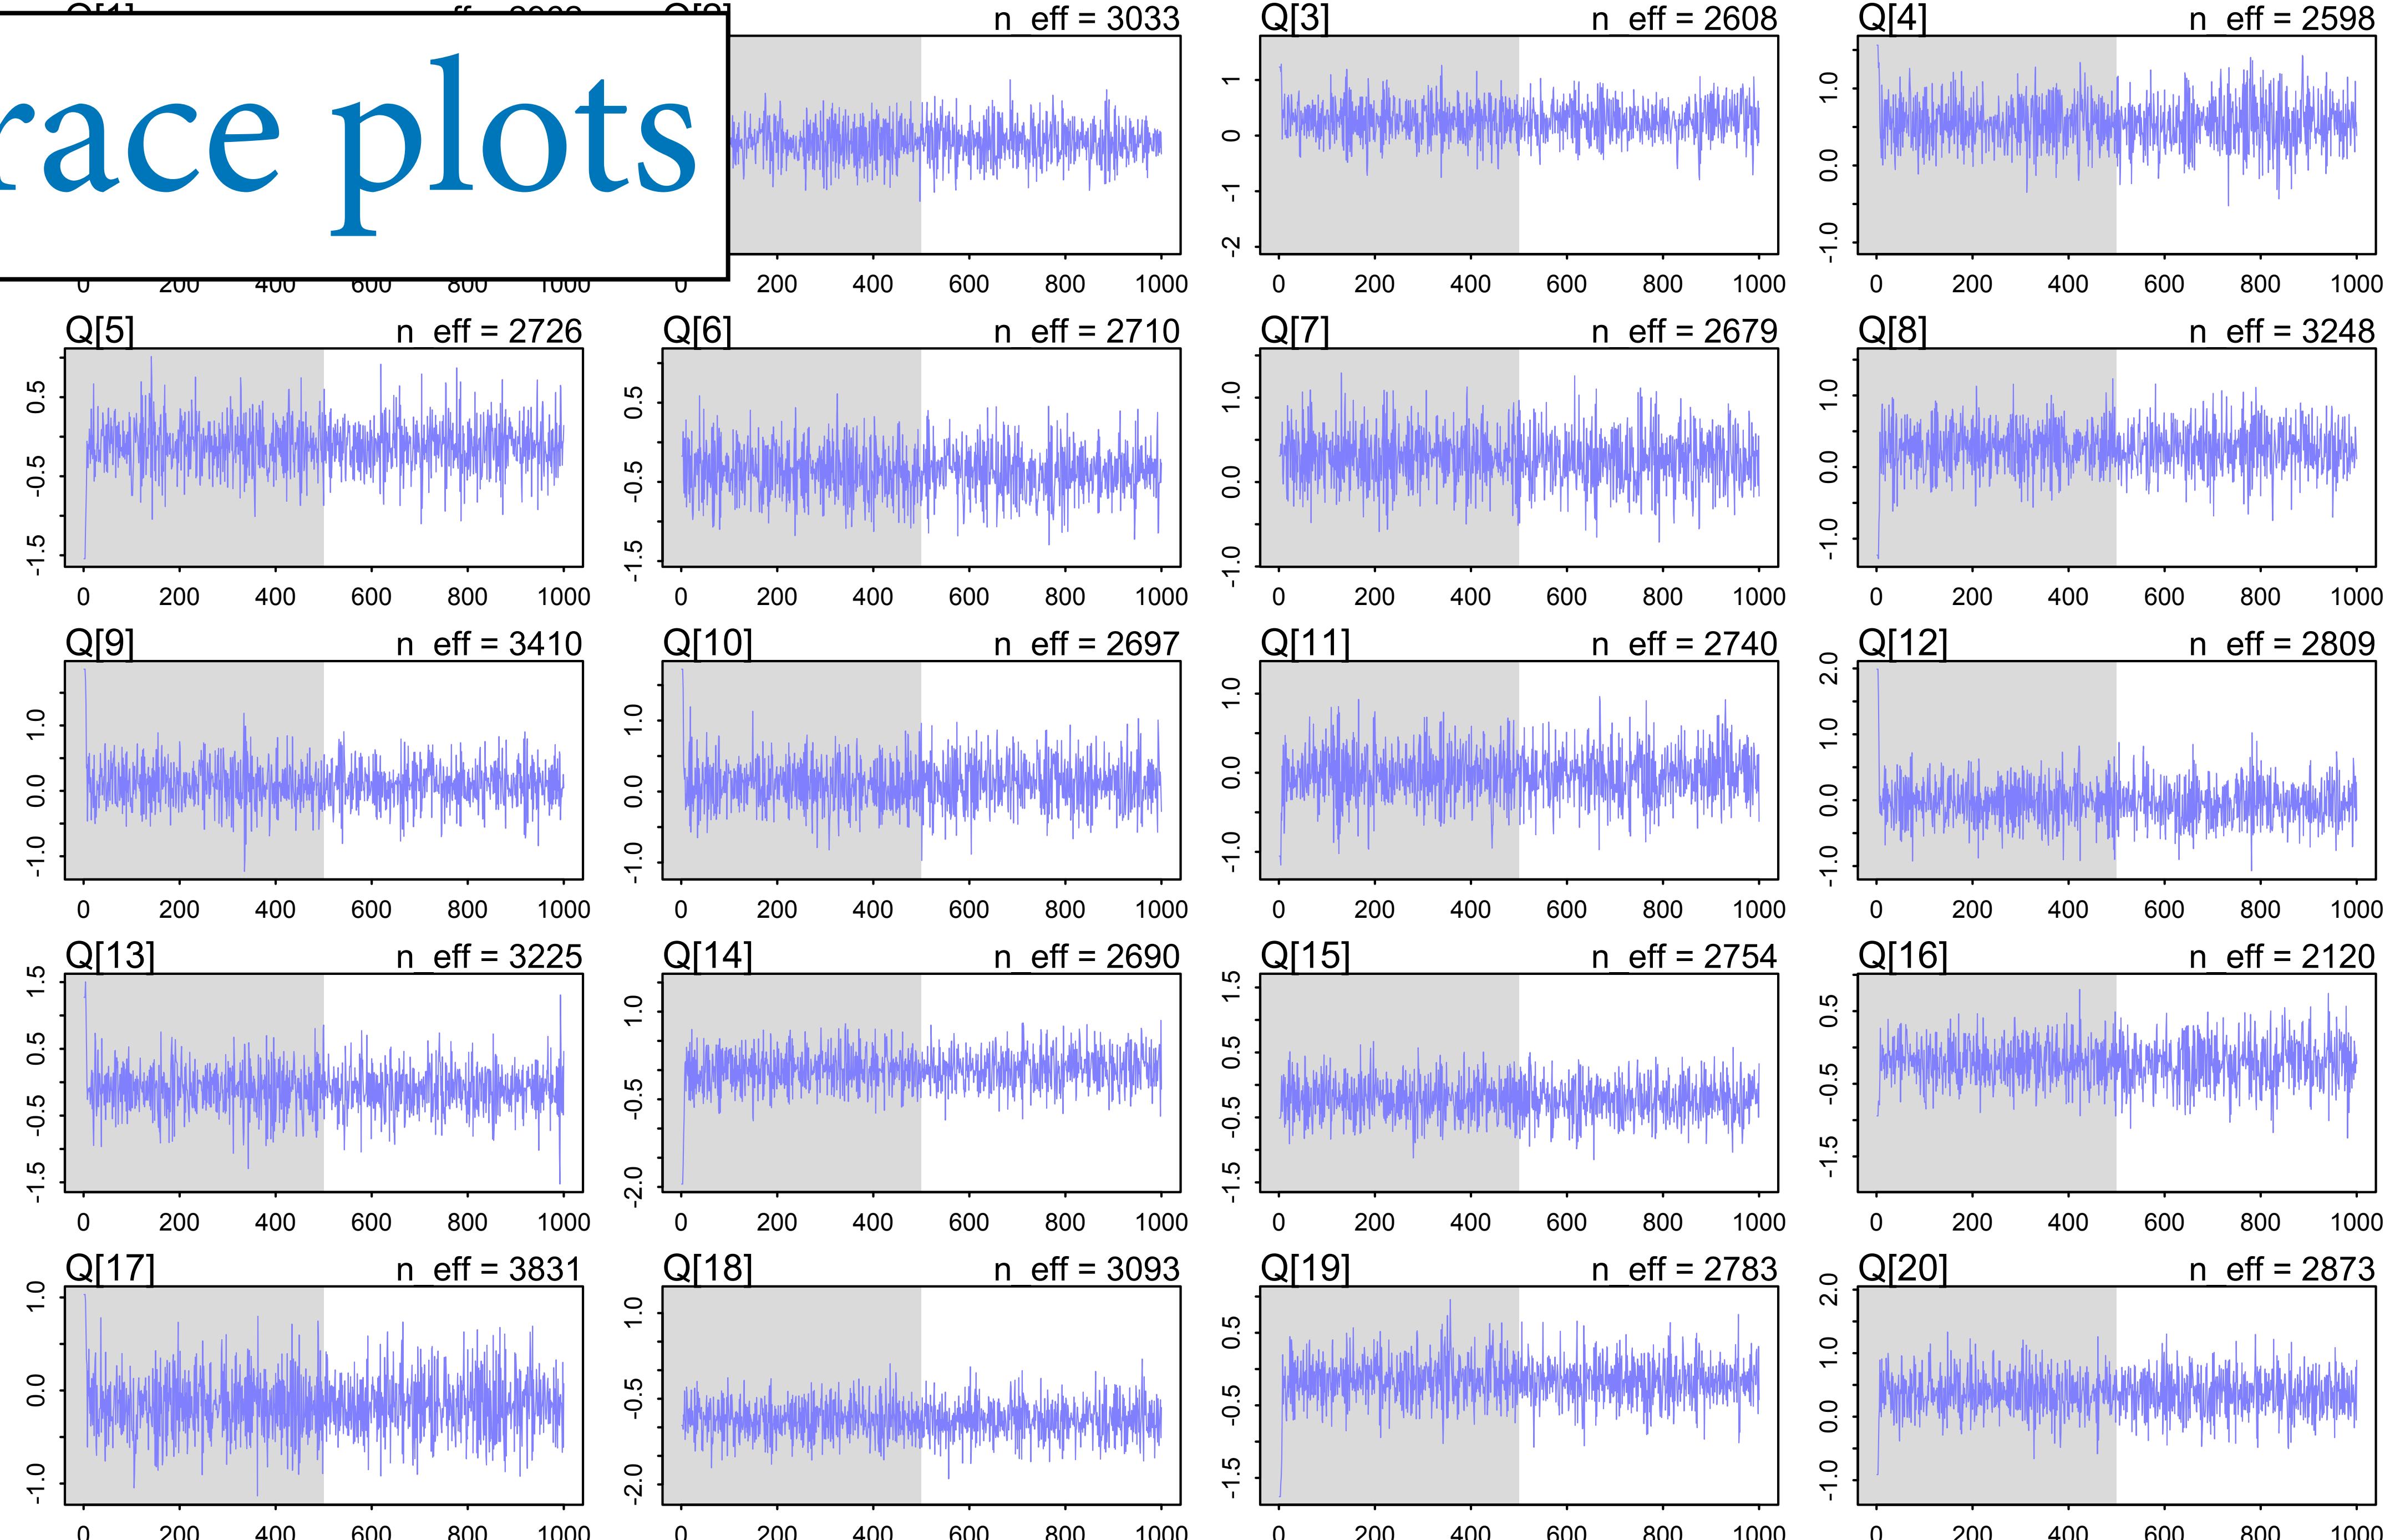
Complex machinery, but lots of diagnostics

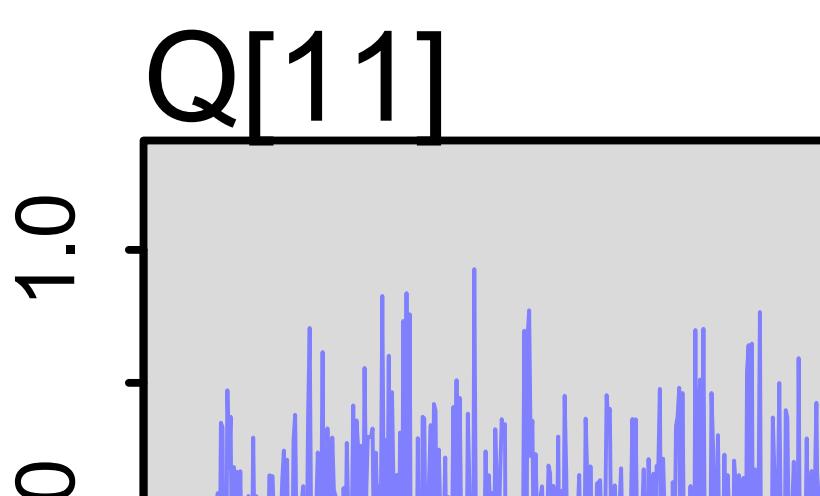
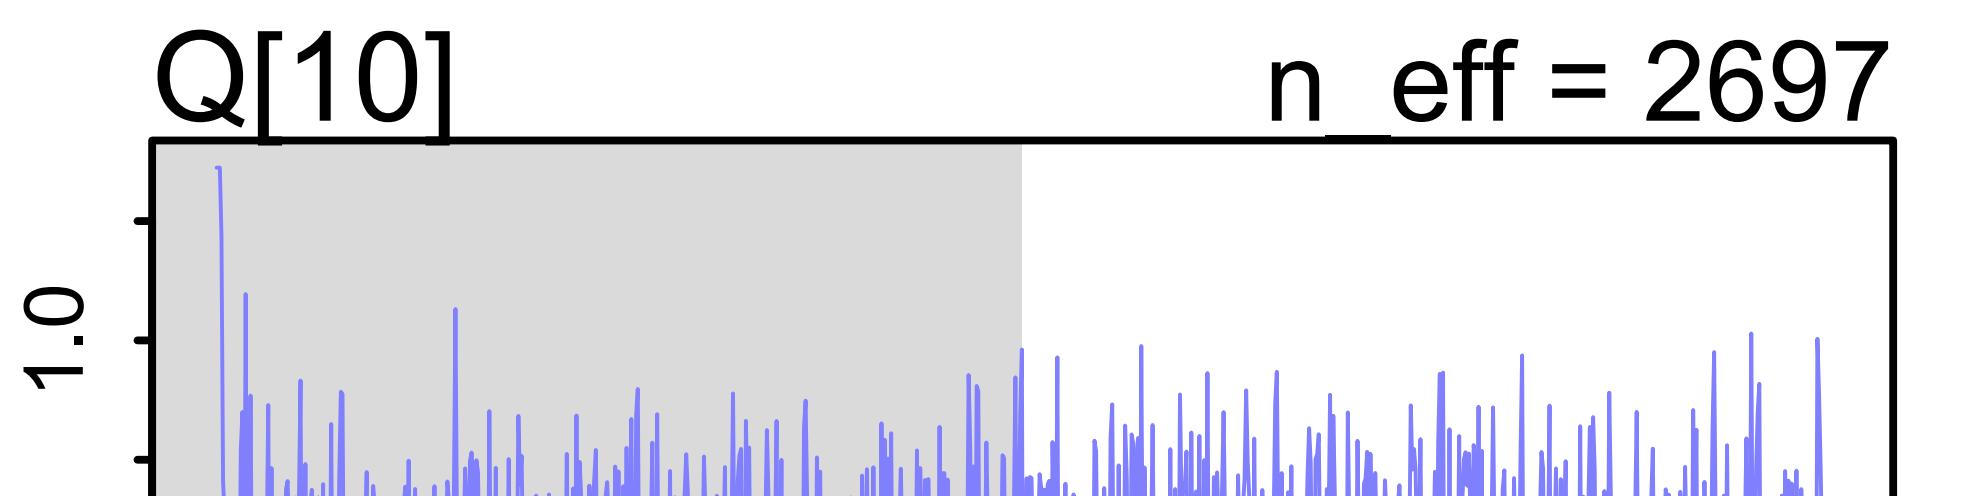
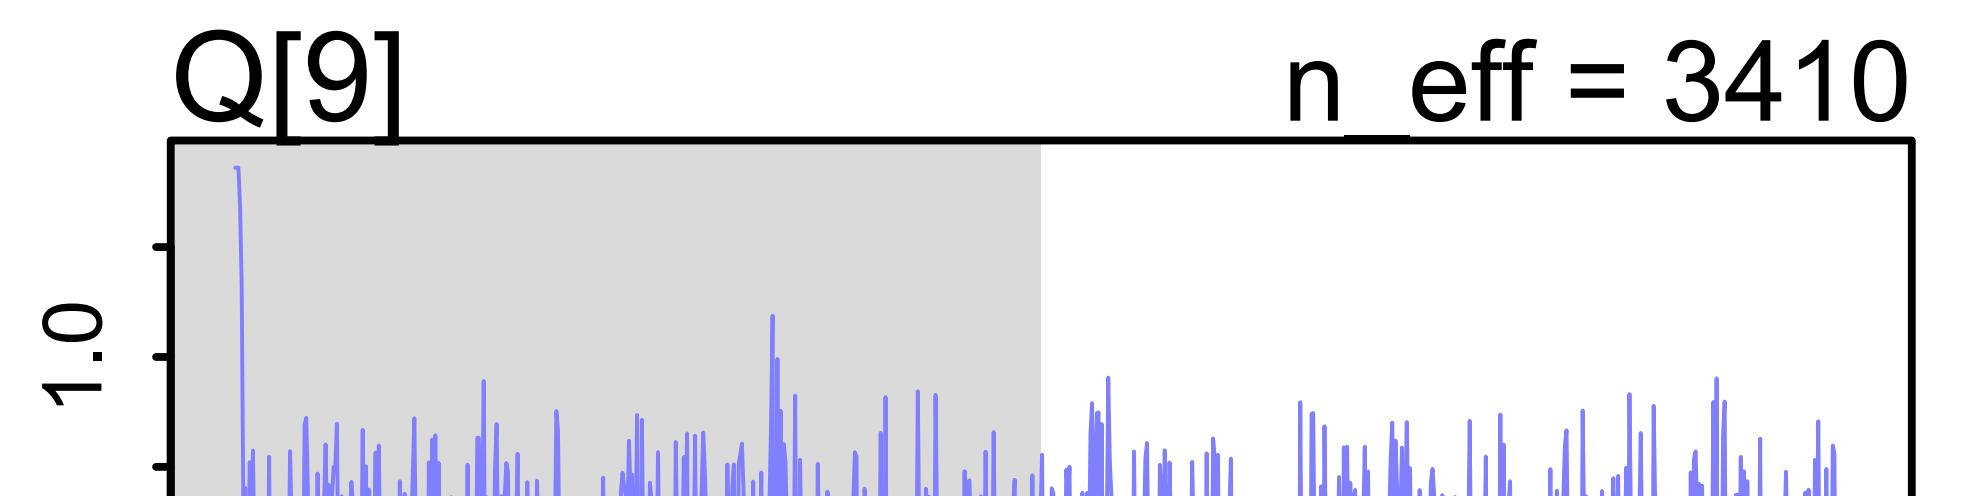
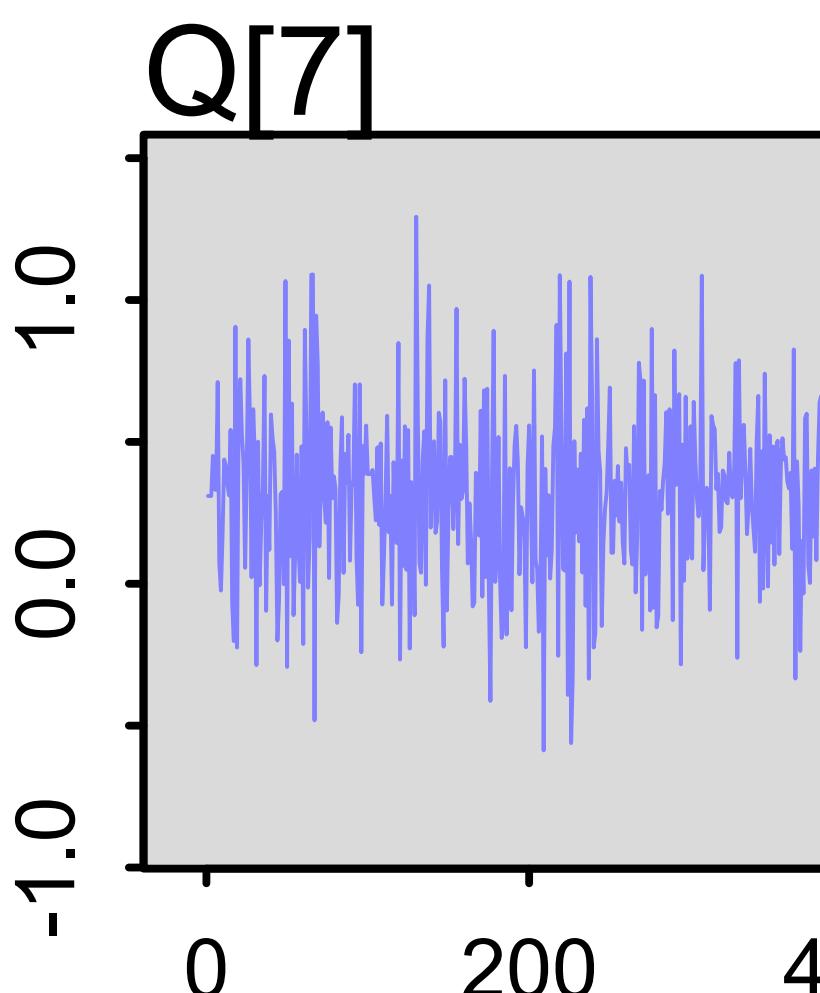
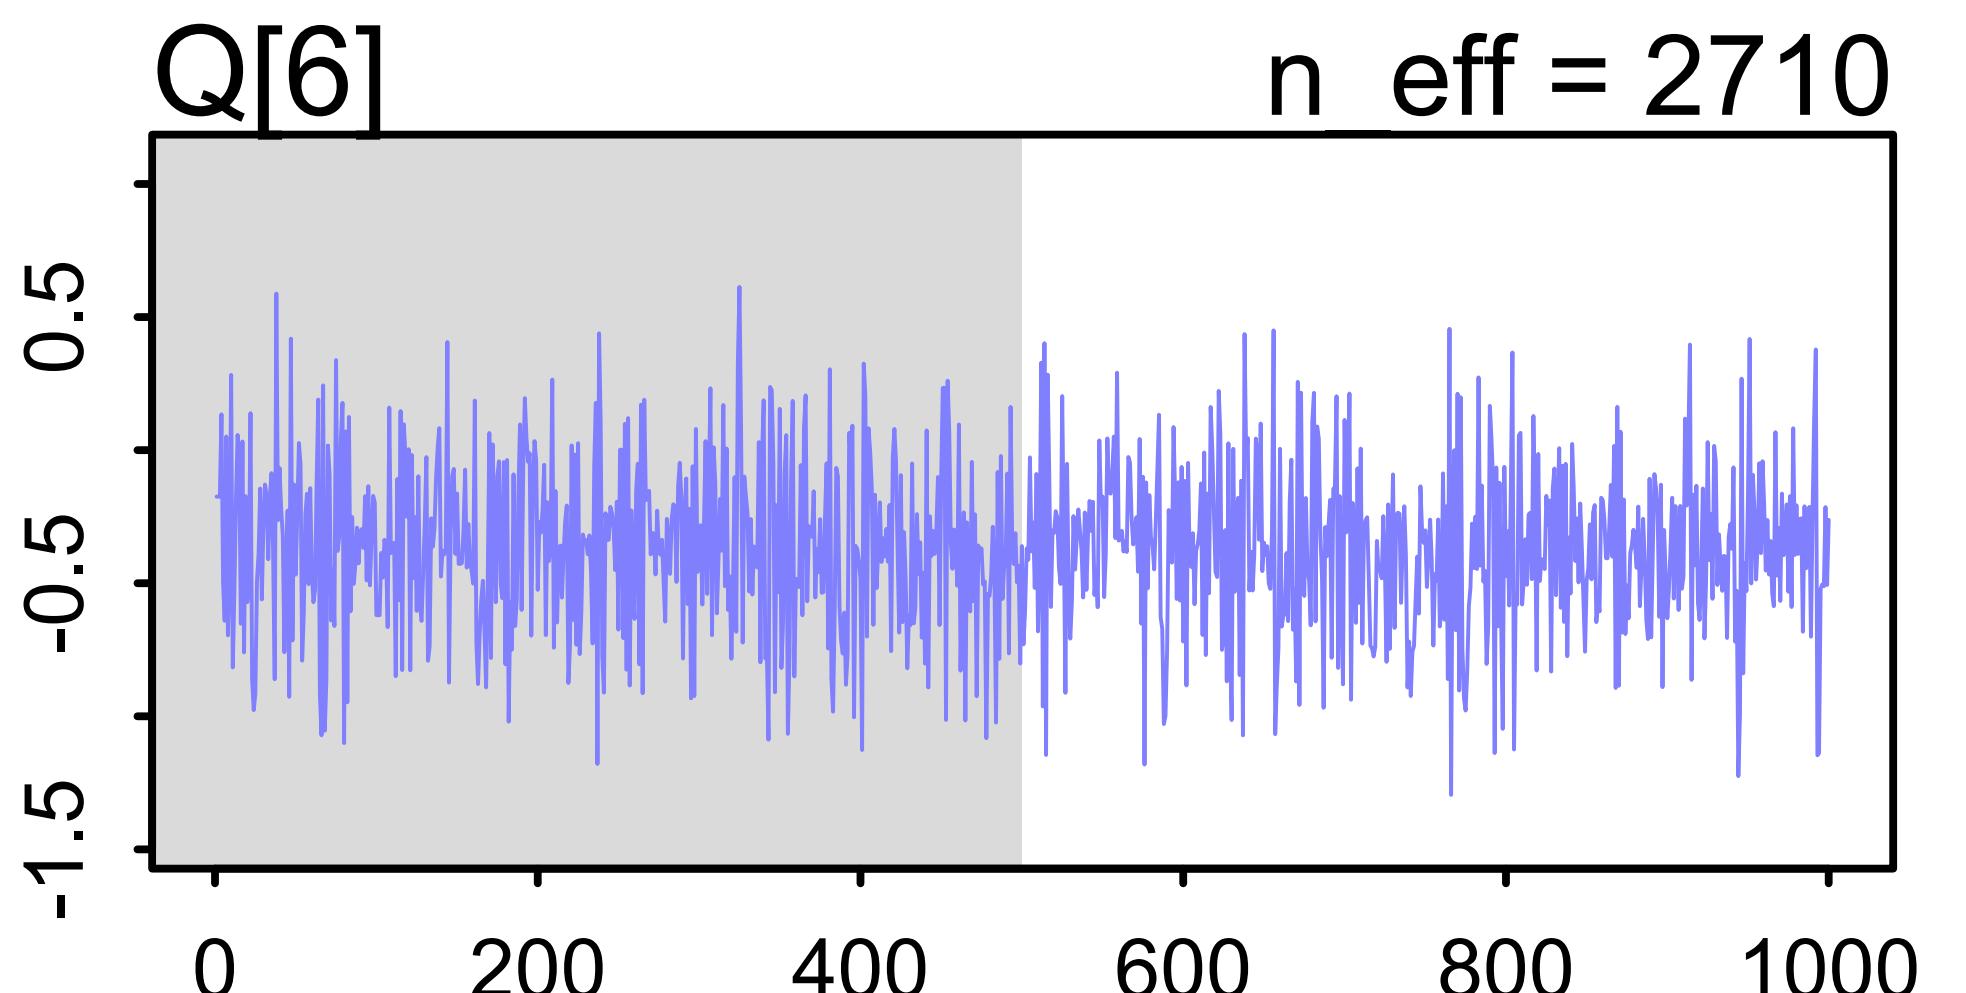
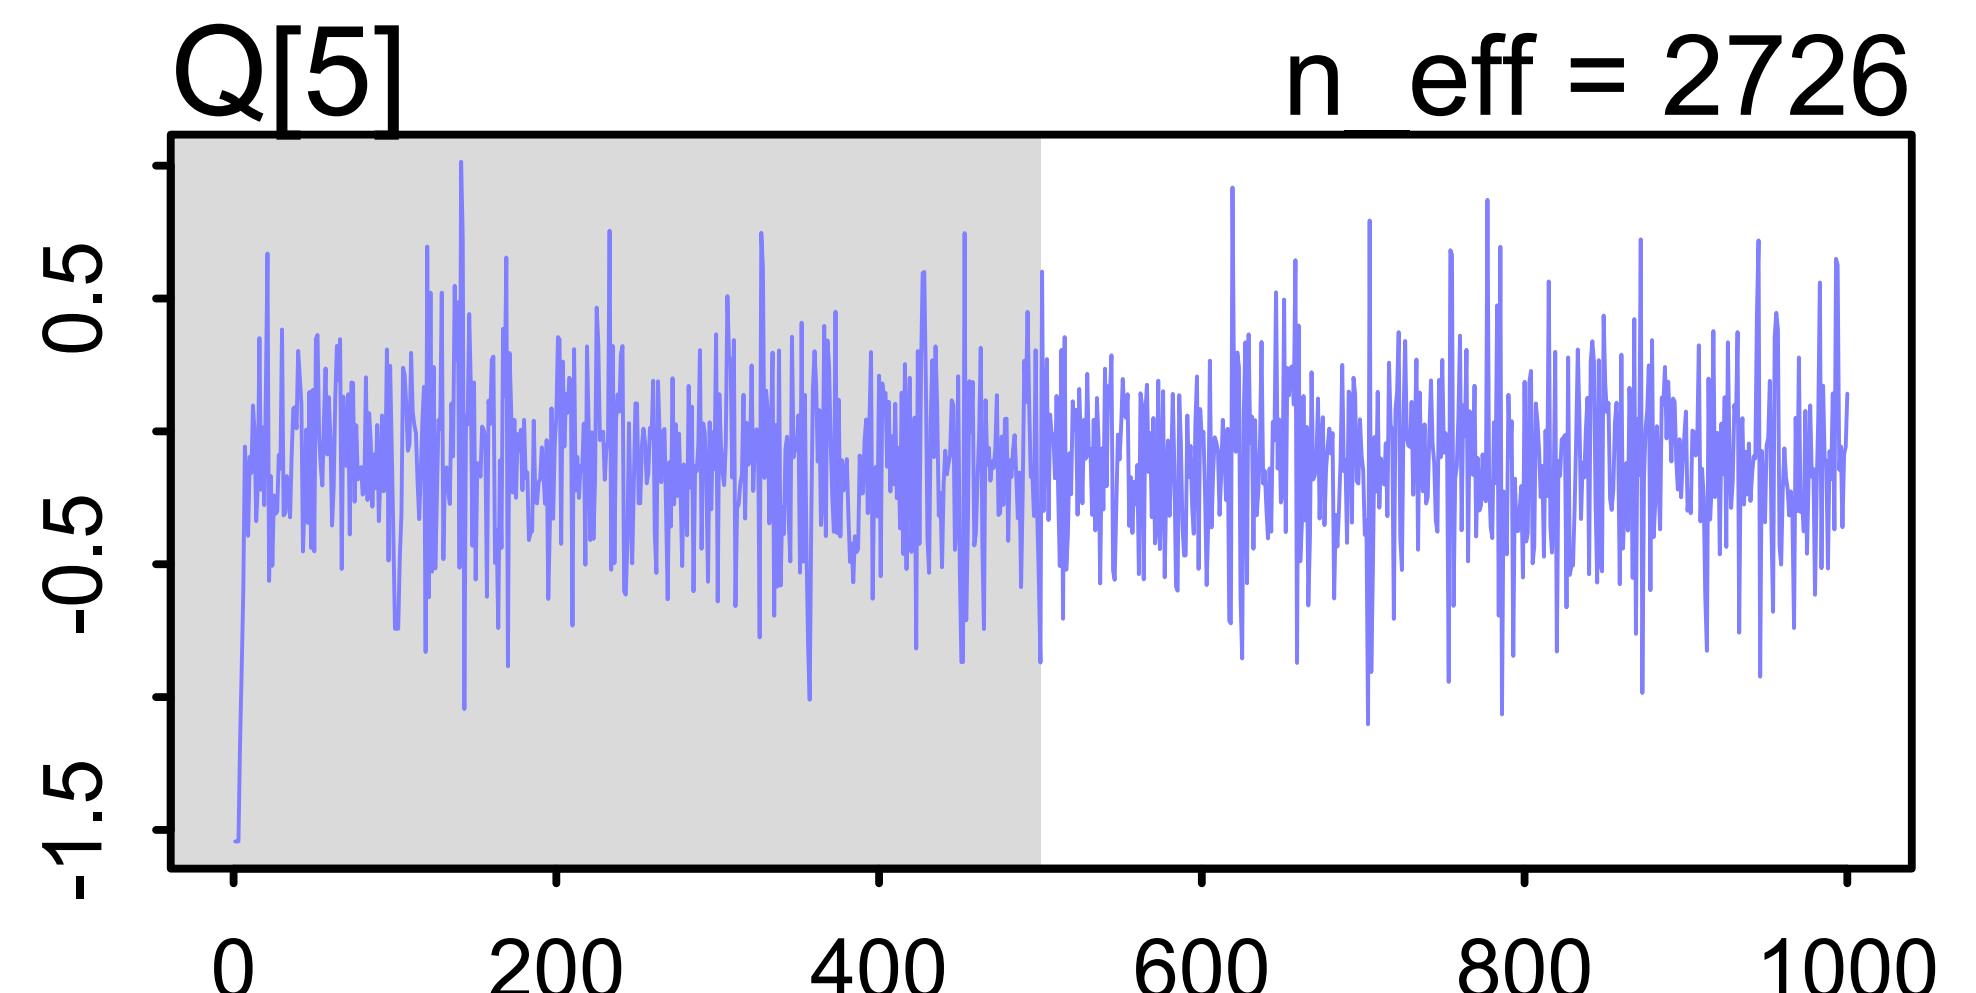
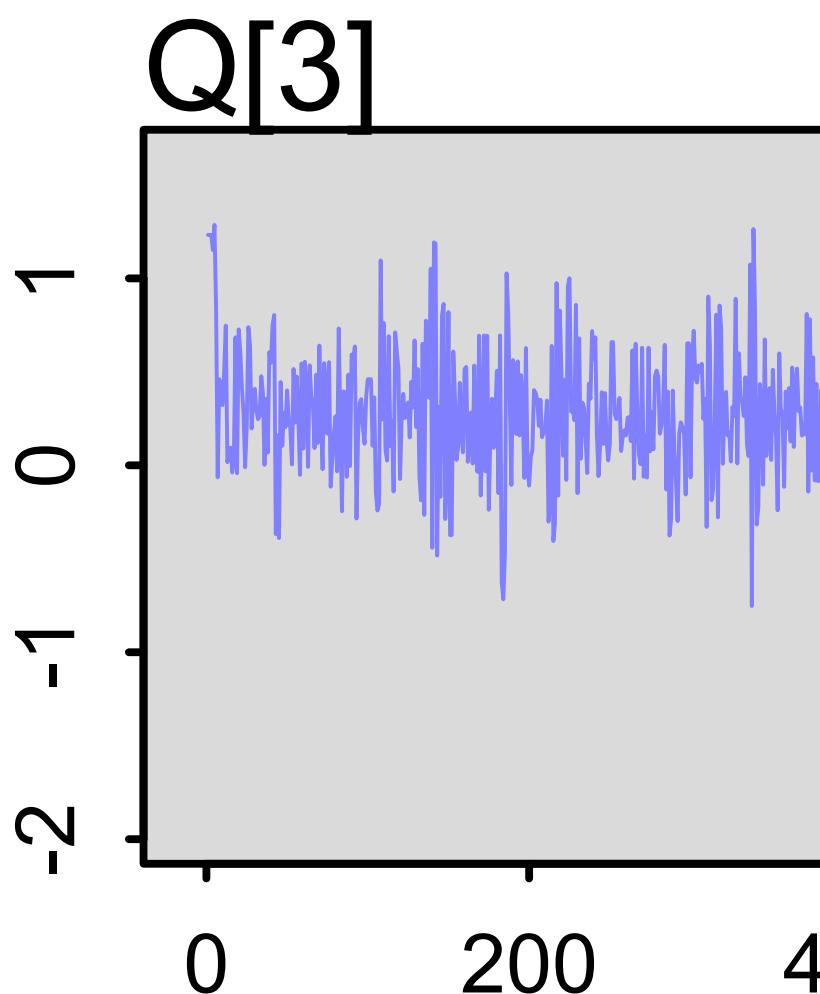
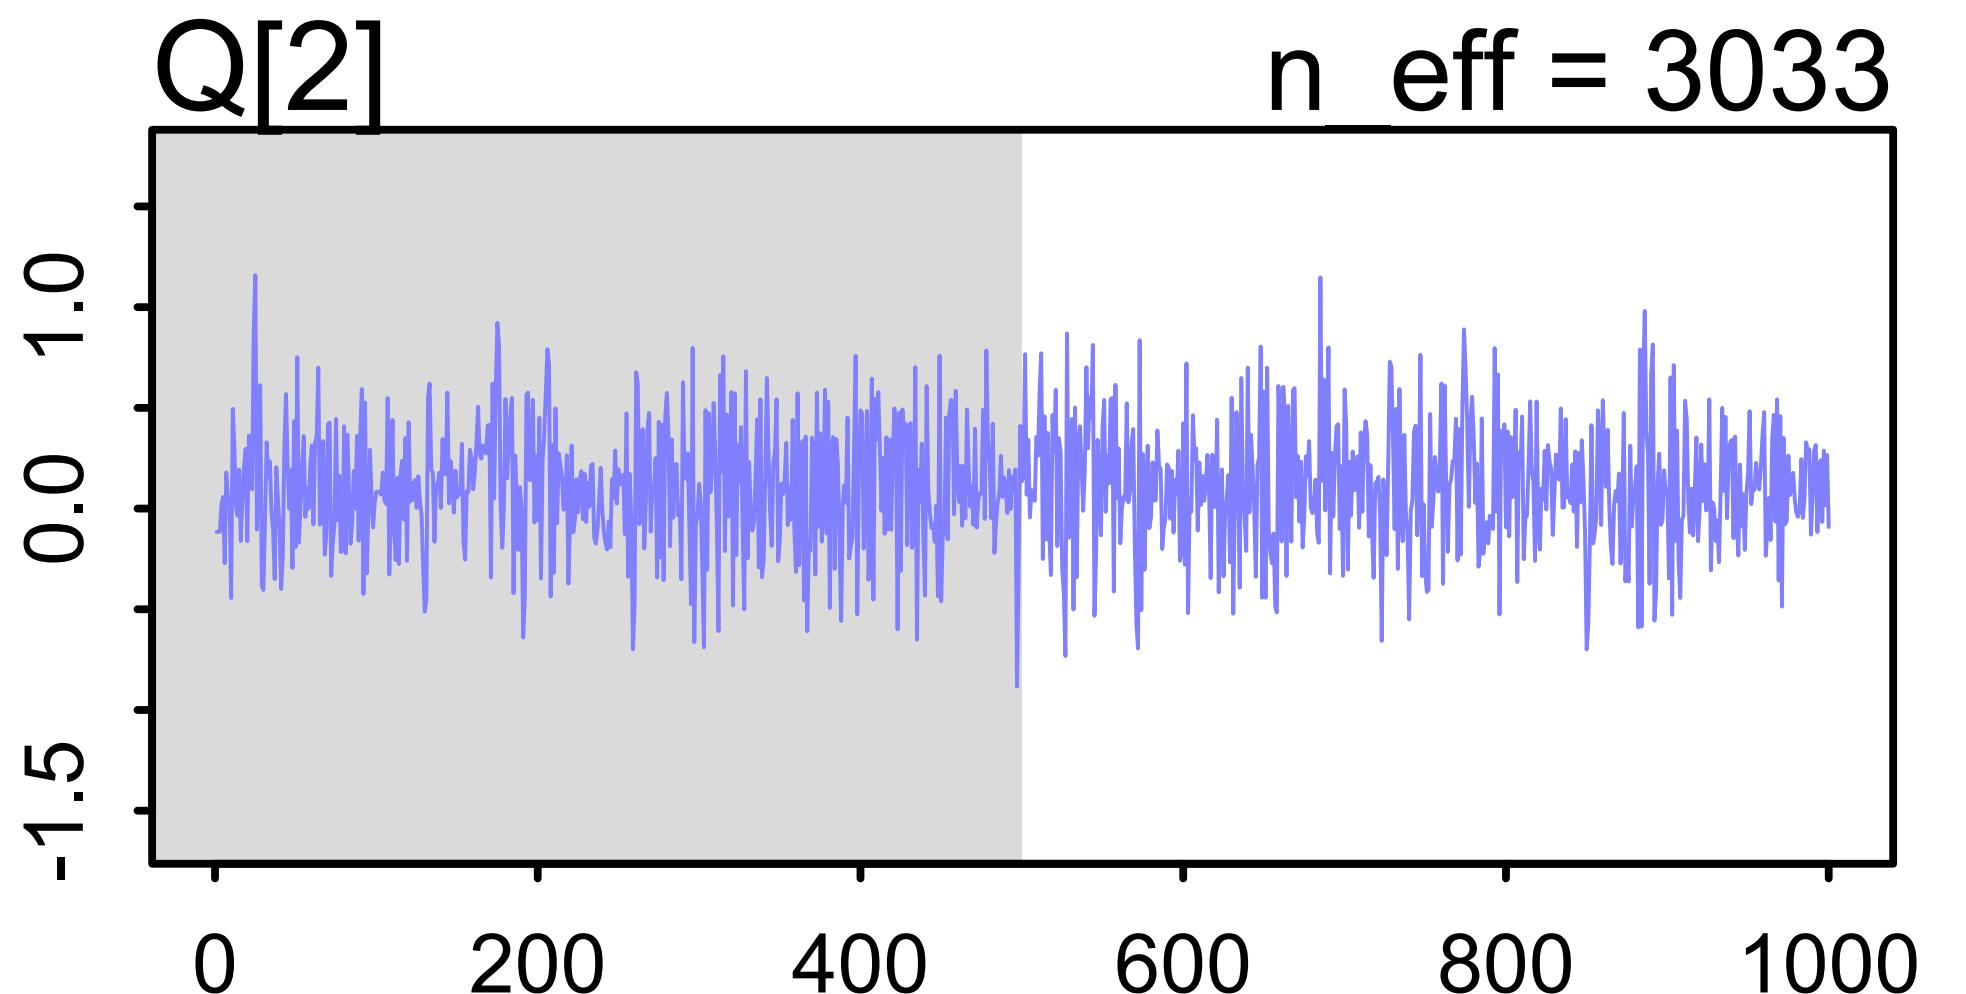
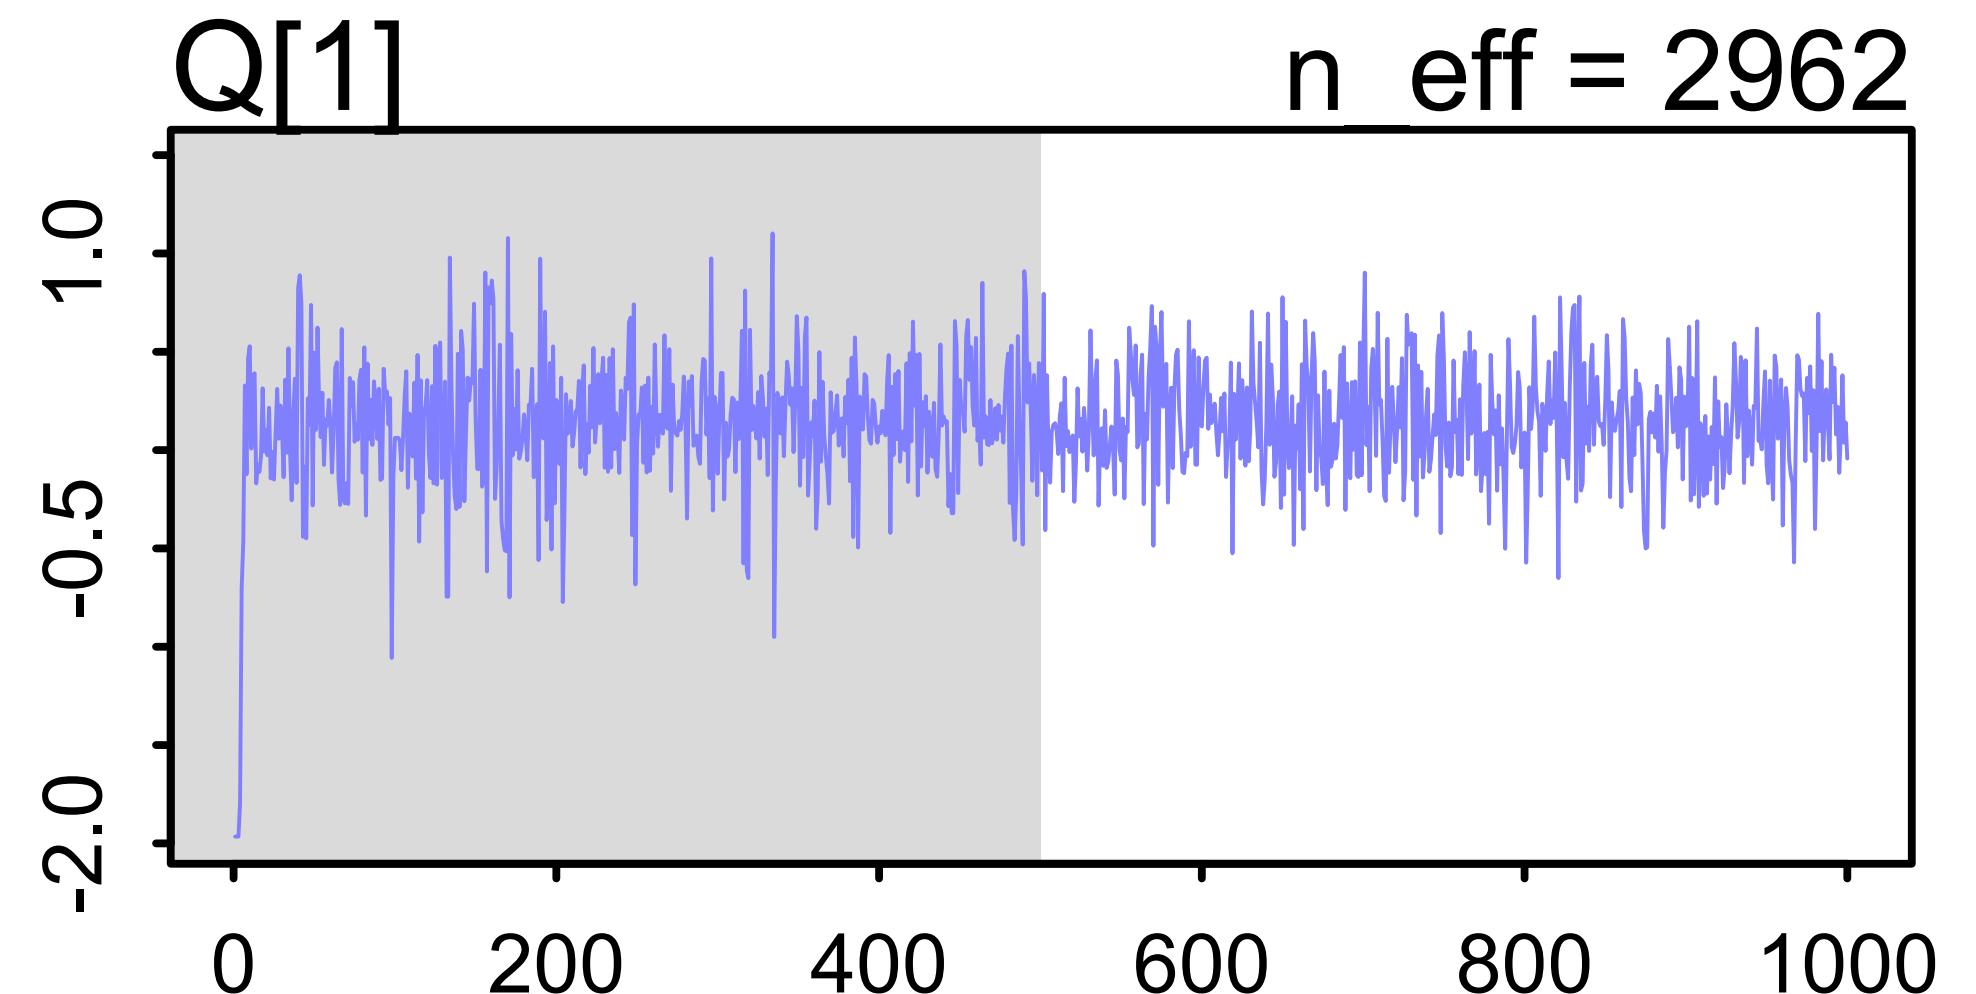
- (1) Trace plots
- (2) Trace rank plots
- (3) R-hat convergence measure
- (4) Number of effective samples
- (5) Divergent transitions

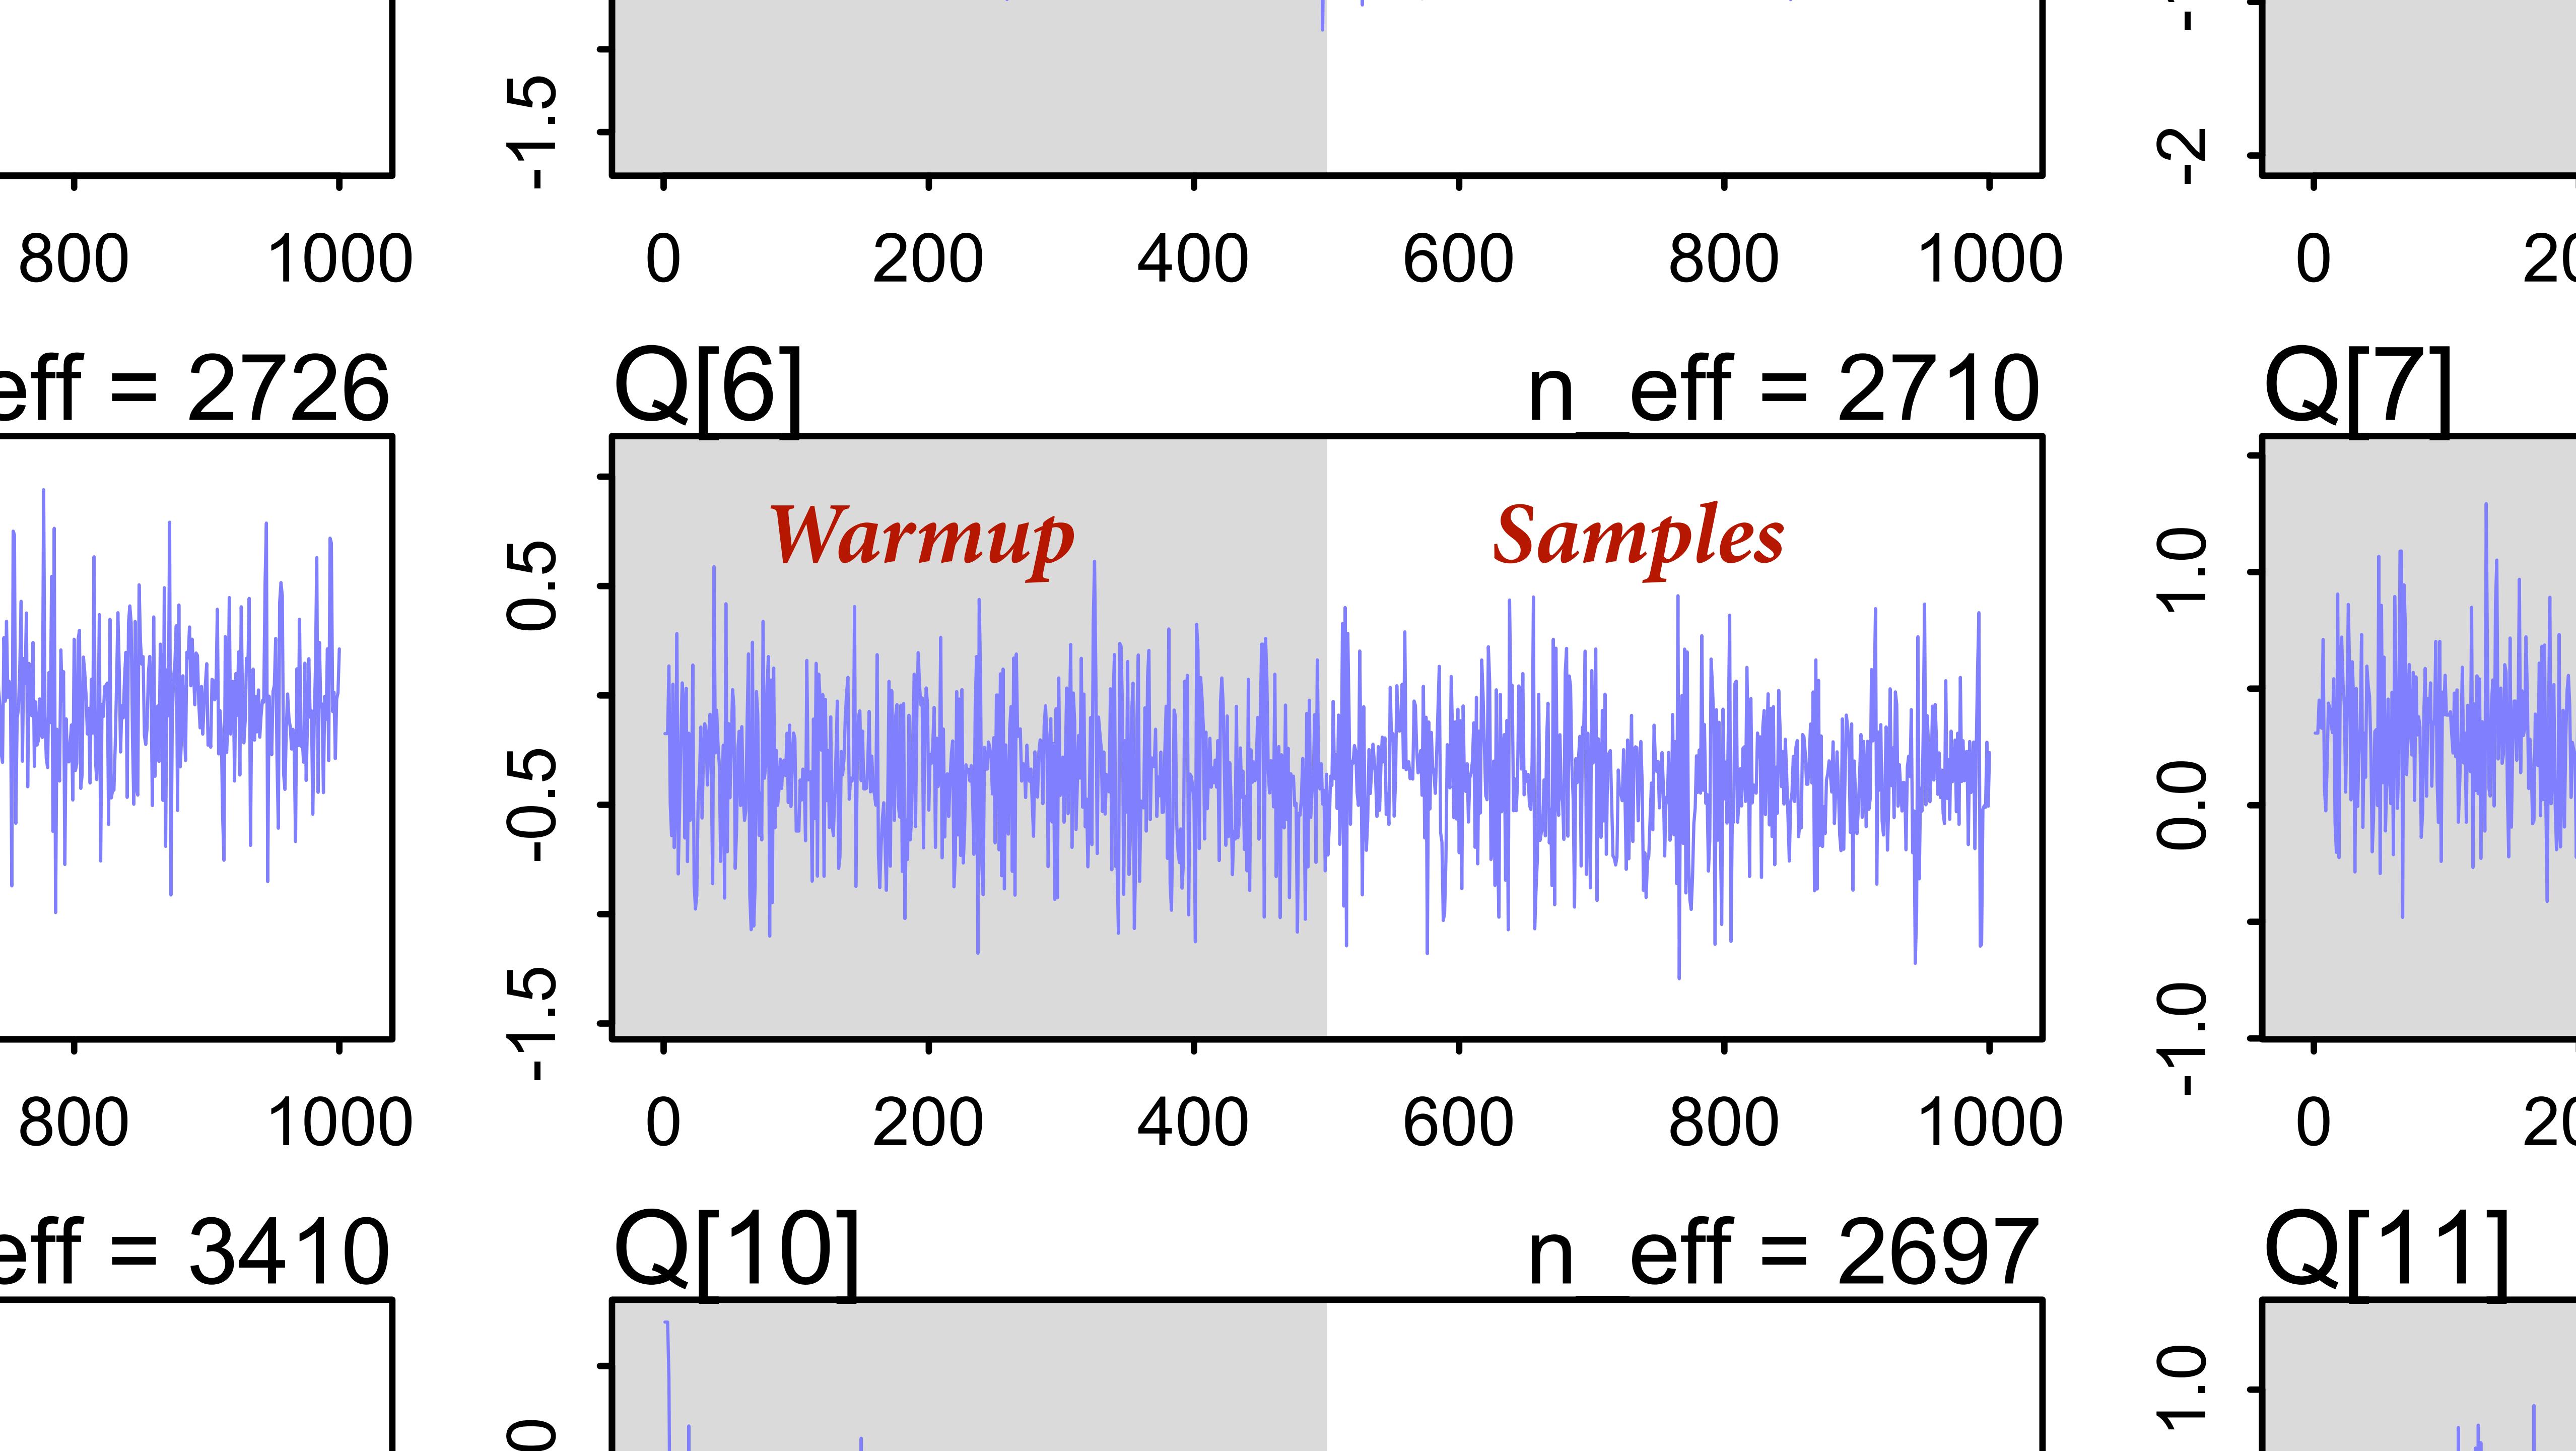


Picaso

Trace plots







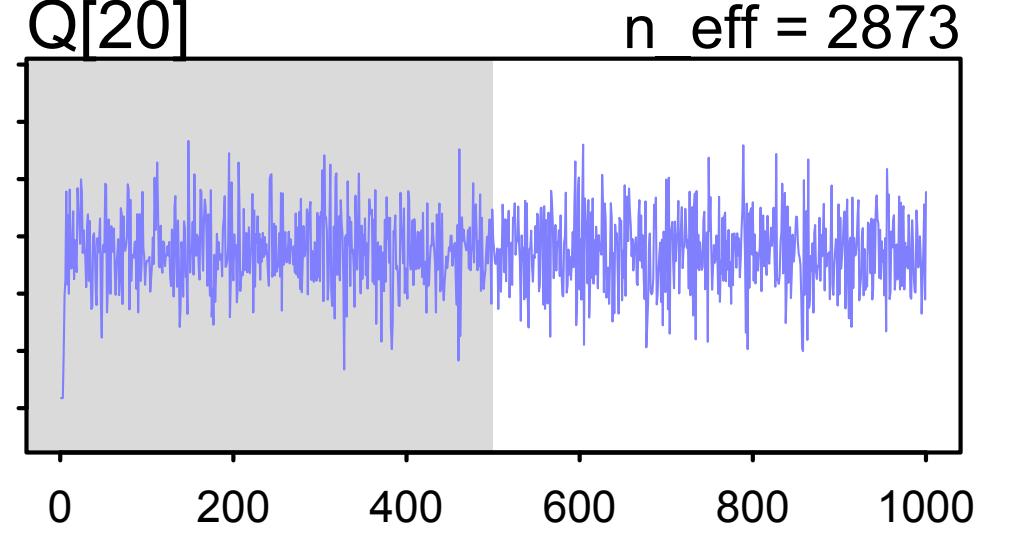
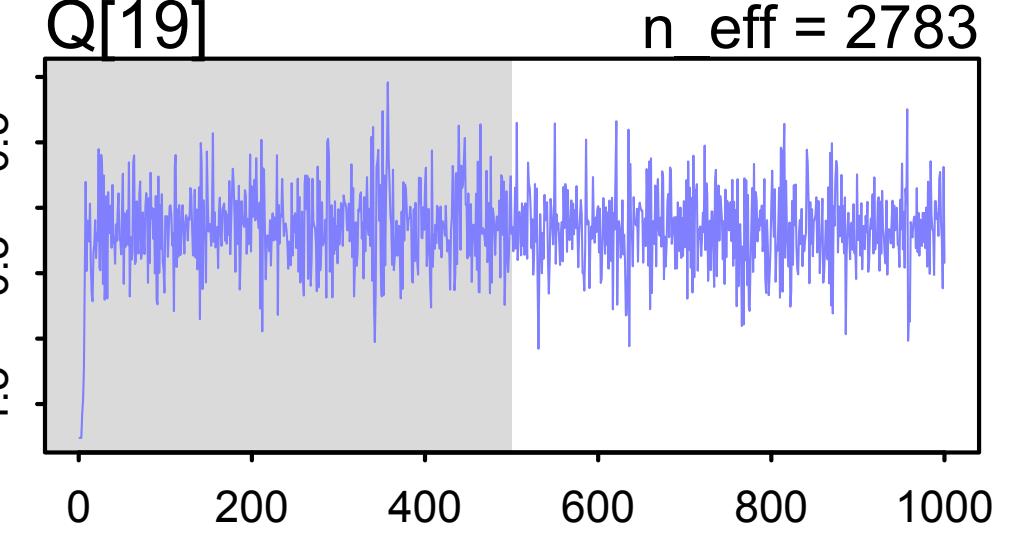
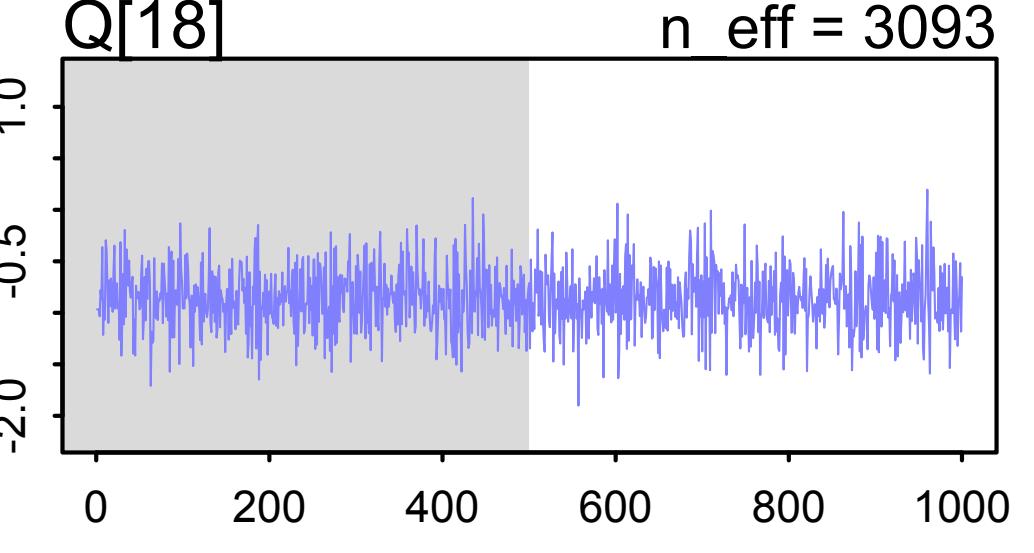
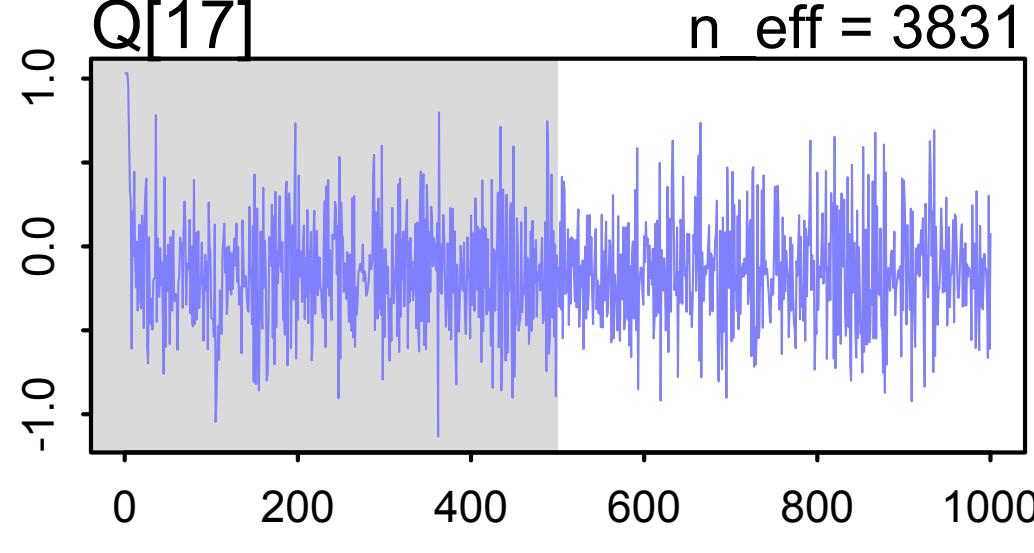
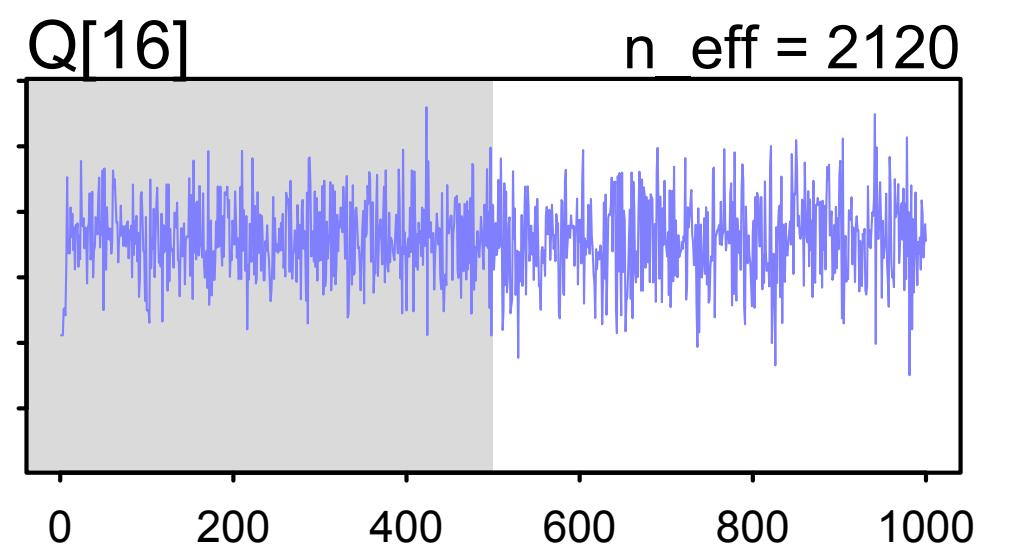
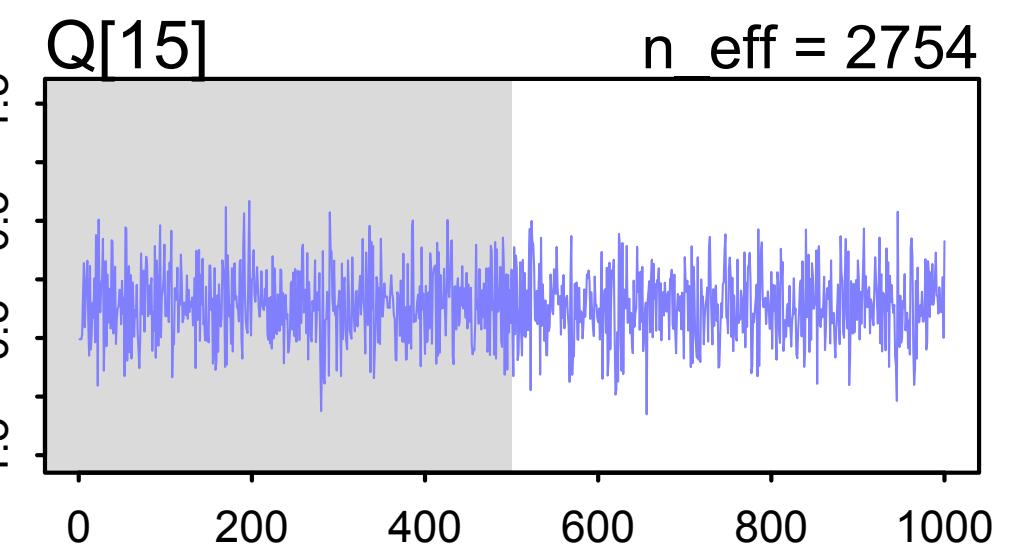
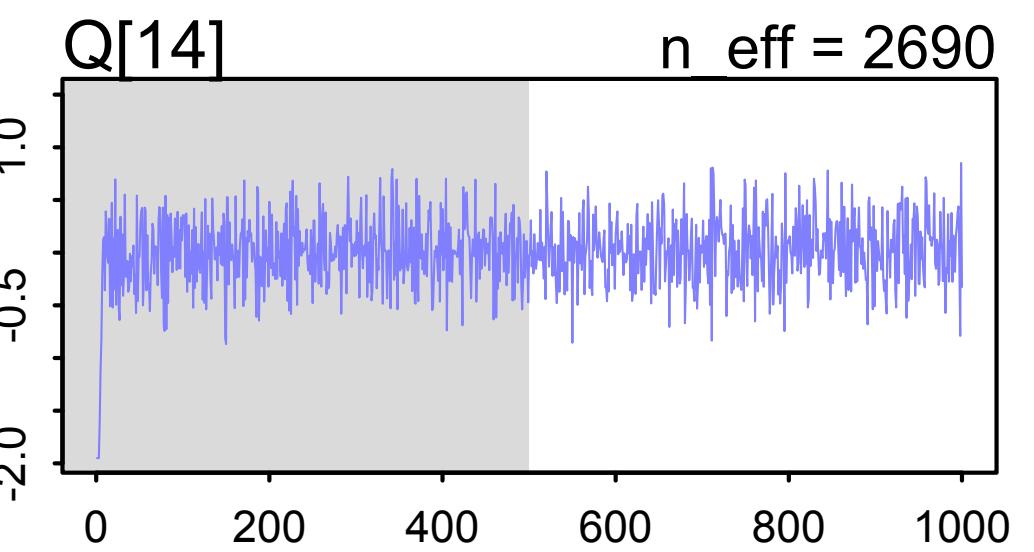
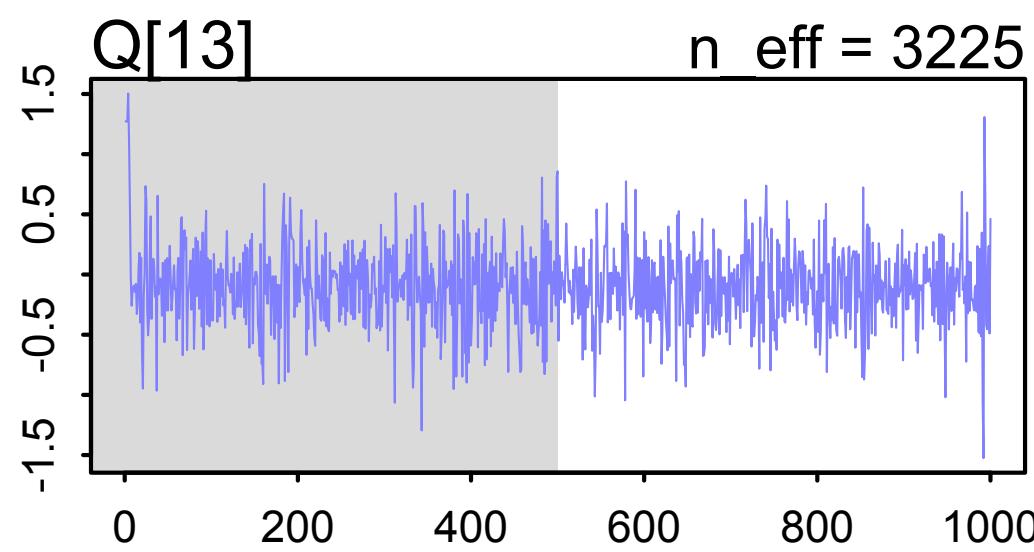
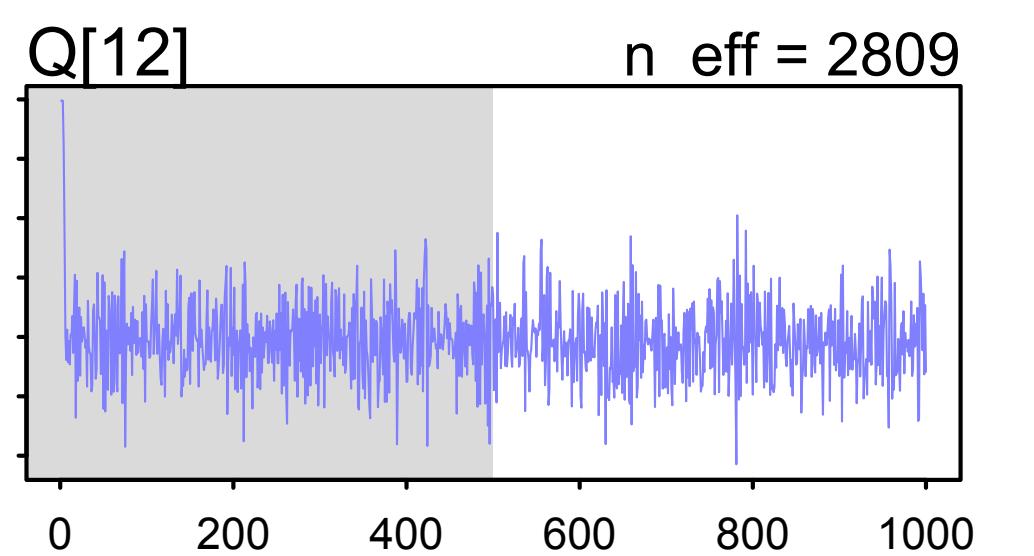
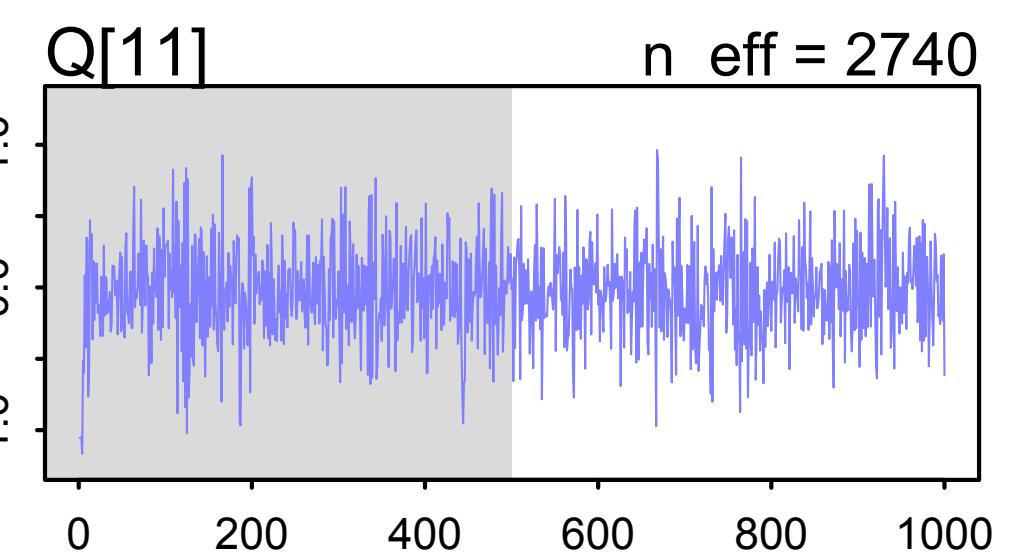
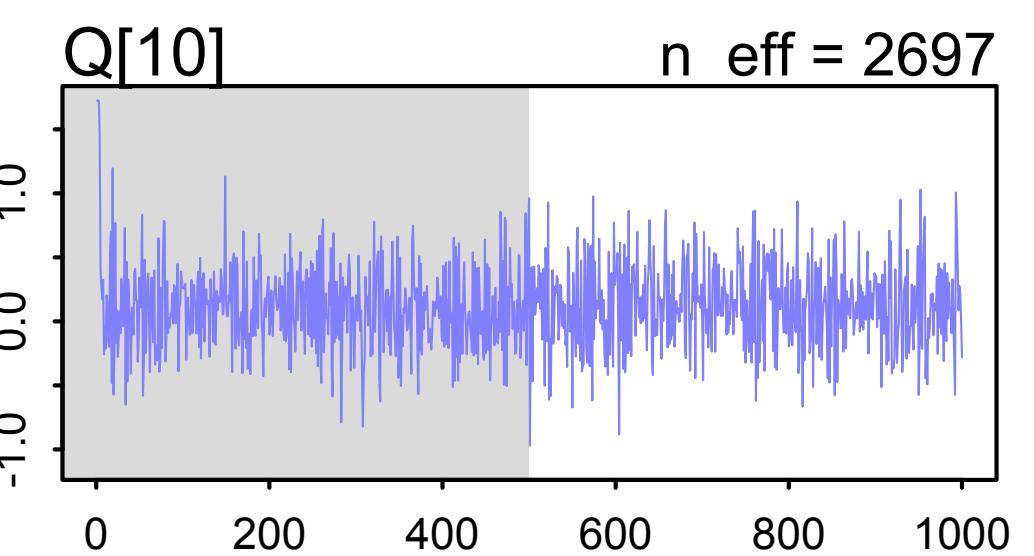
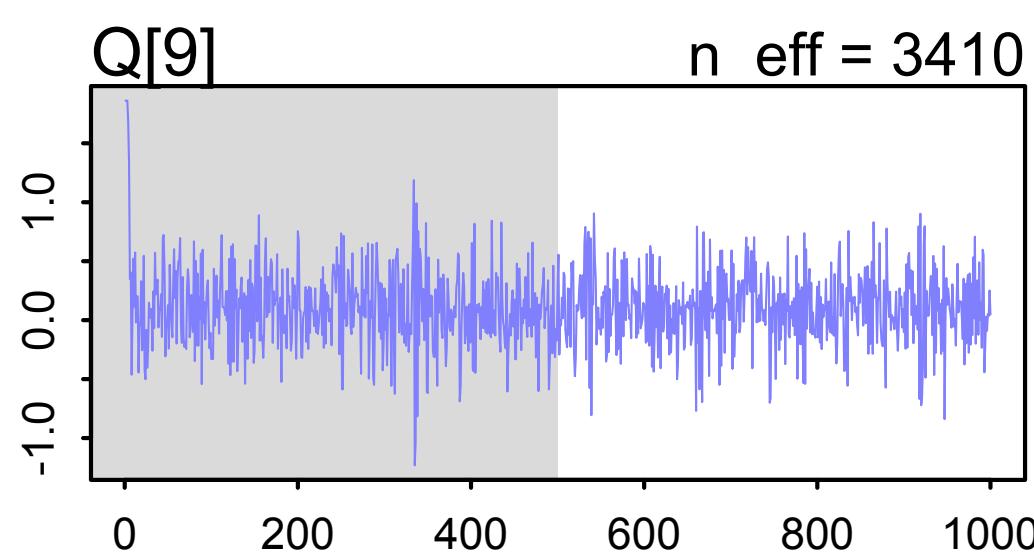
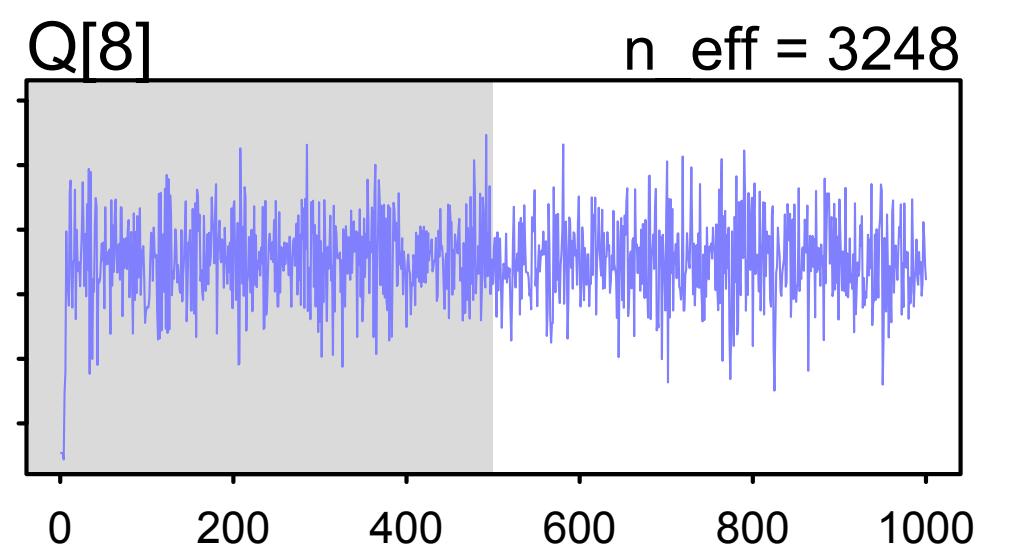
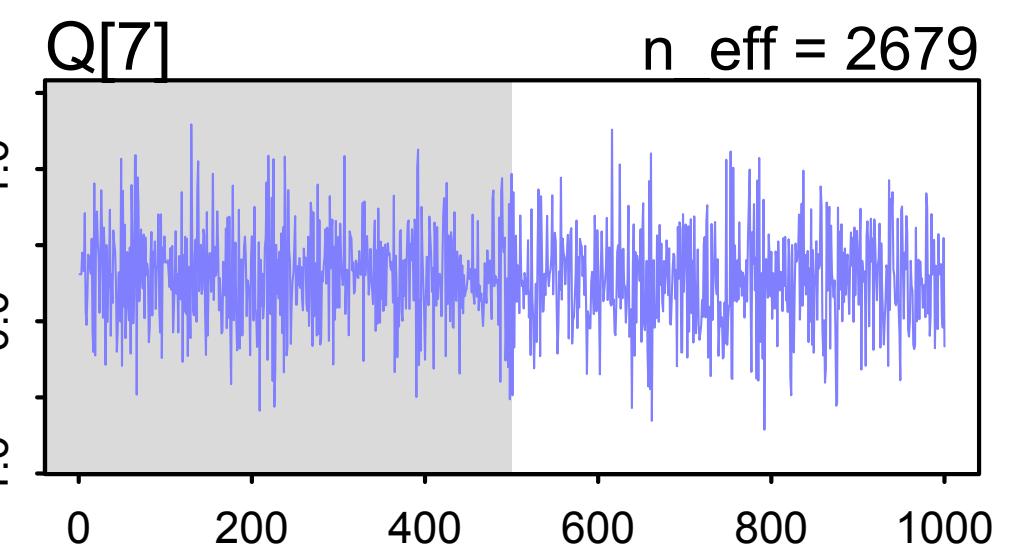
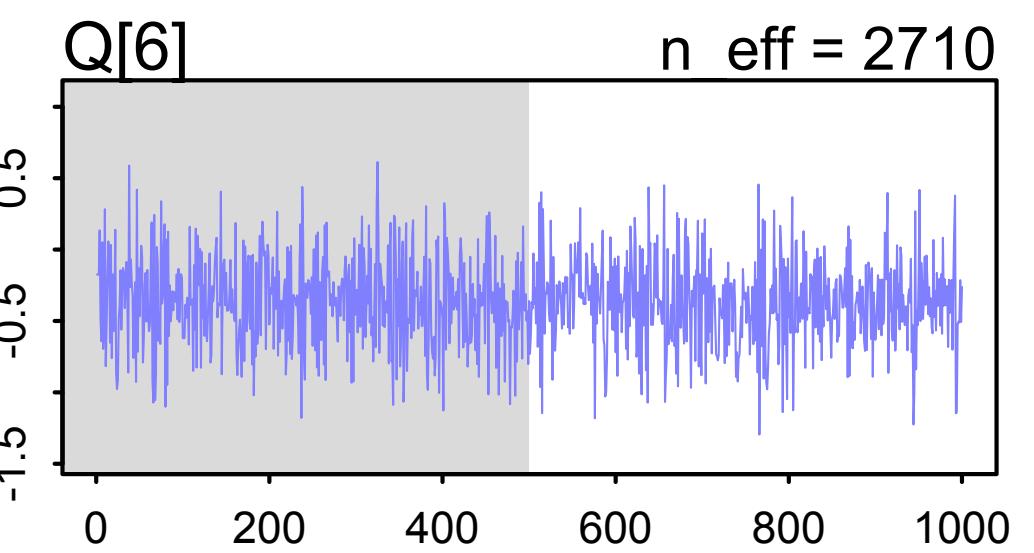
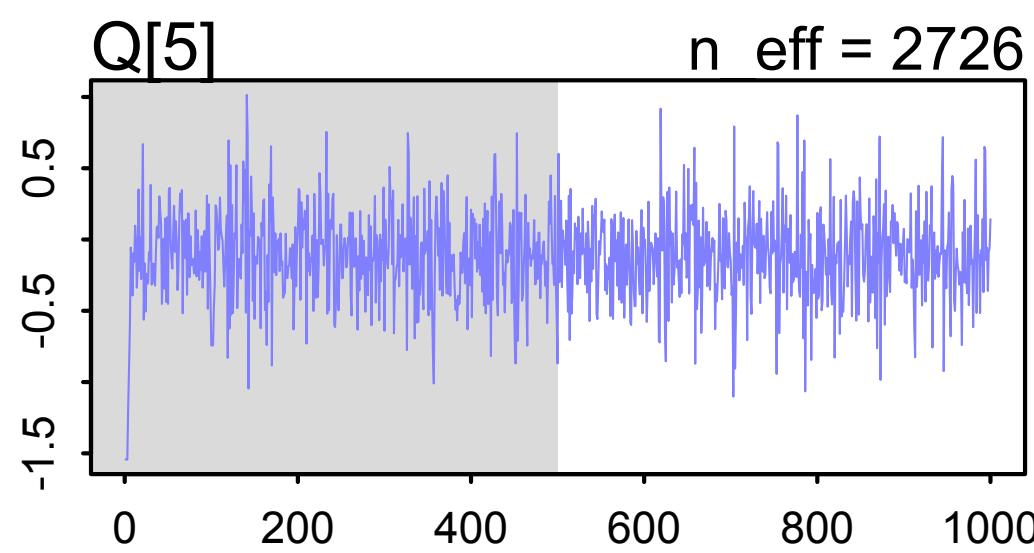
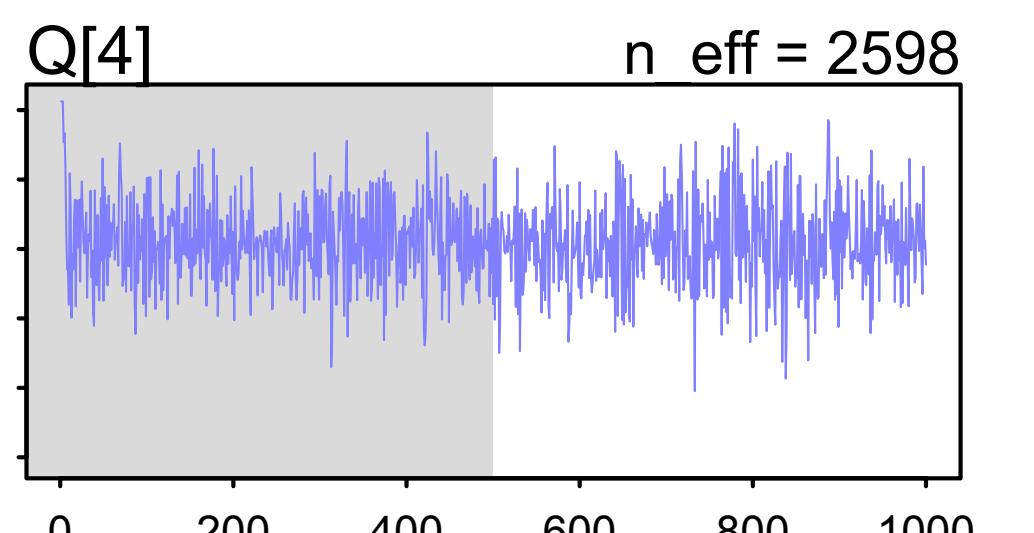
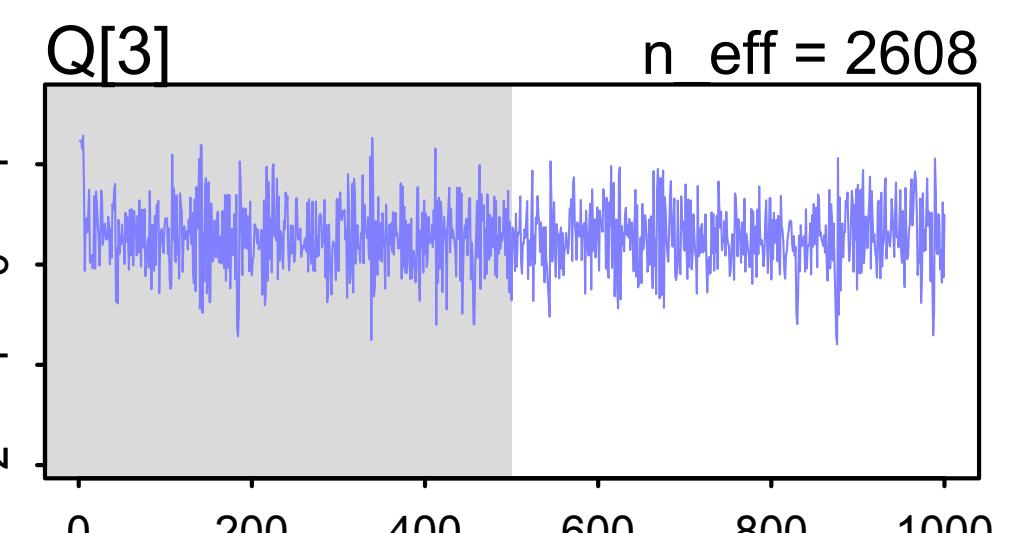
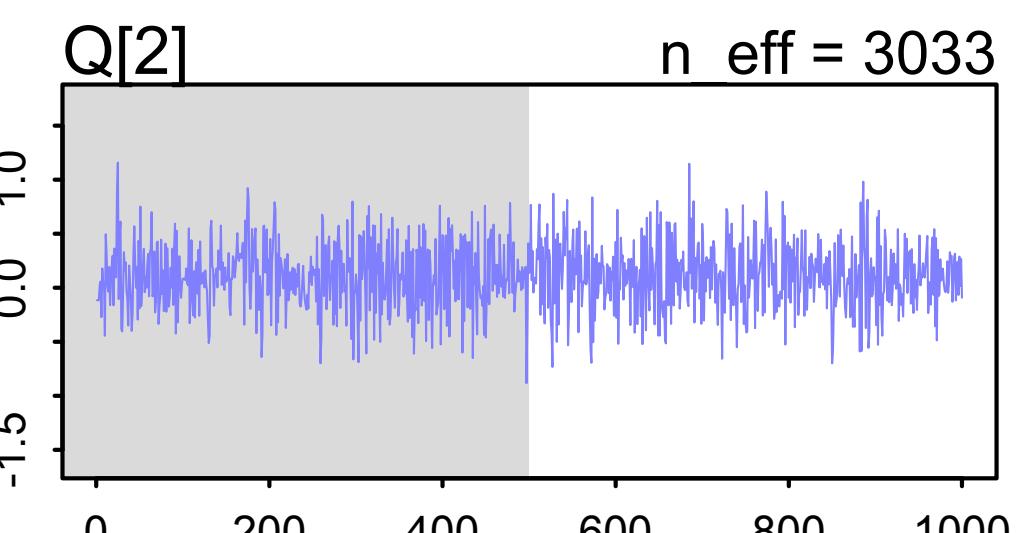
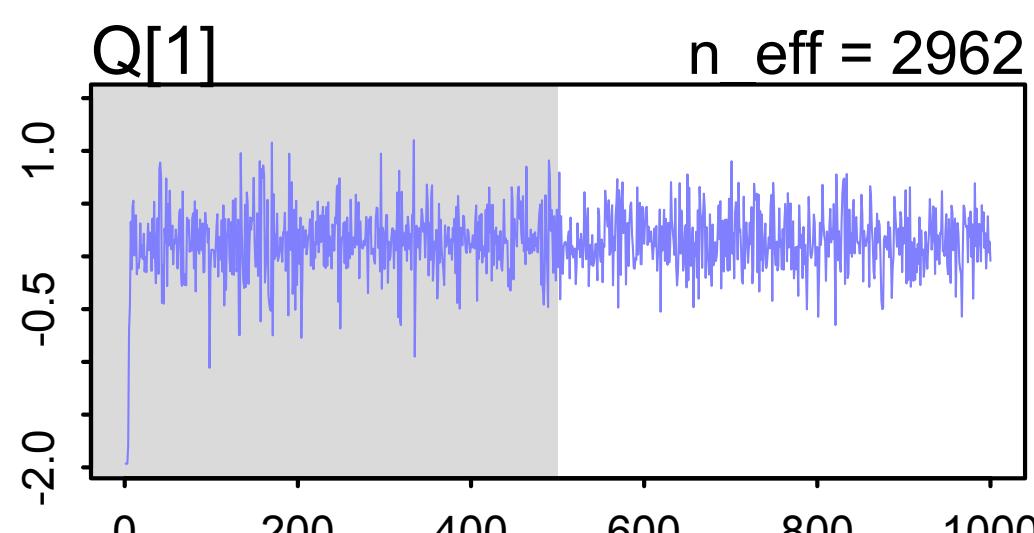
```
library(rethinking)
data(Wines2012)
d <- Wines2012

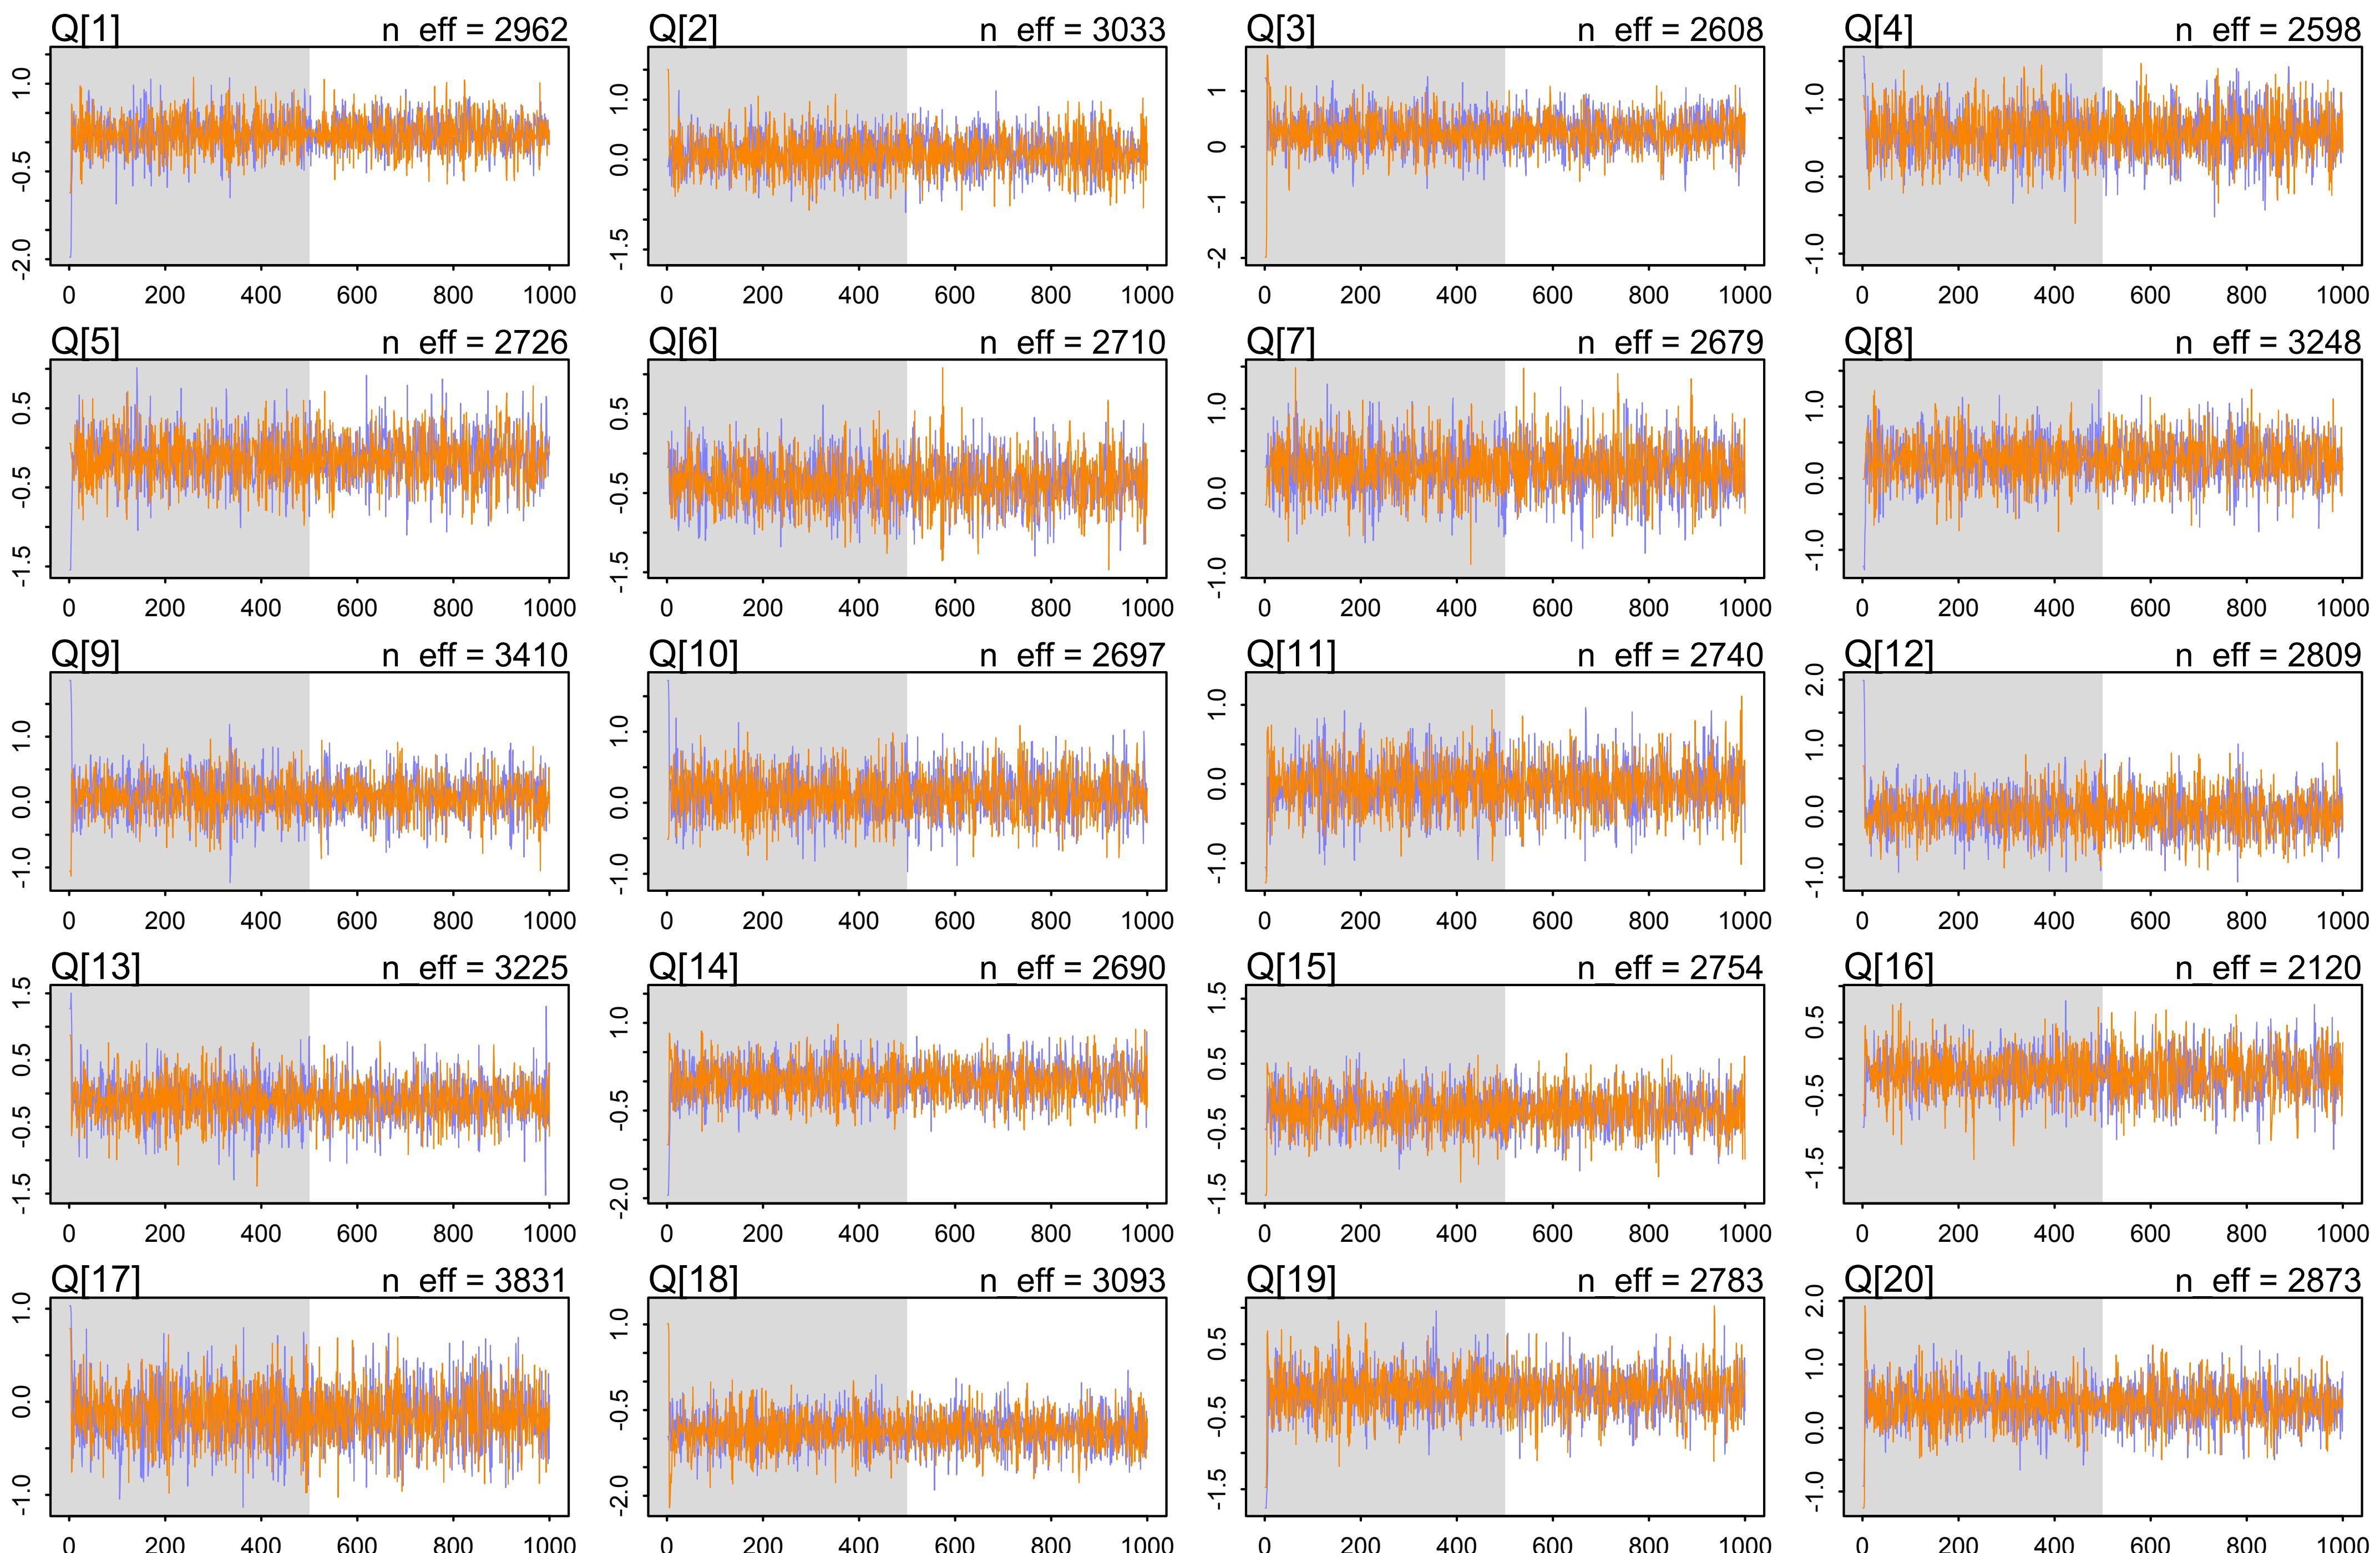
dat <- list(
  S=standardize(d$score),
  J=as.numeric(d$judge),
  W=as.numeric(d$wine),
  X=ifelse(d$wine.amer==1,1,2),
  Z=ifelse(d$judge.amer==1,1,2)
)
mQ <- ulam(
  alist(
    S ~ dnorm(mu,sigma),
    mu <- Q[W],
    Q[W] ~ dnorm(0,1),
    sigma ~ dexp(1)
  ) , data=dat , chains=4 , cores=4 )

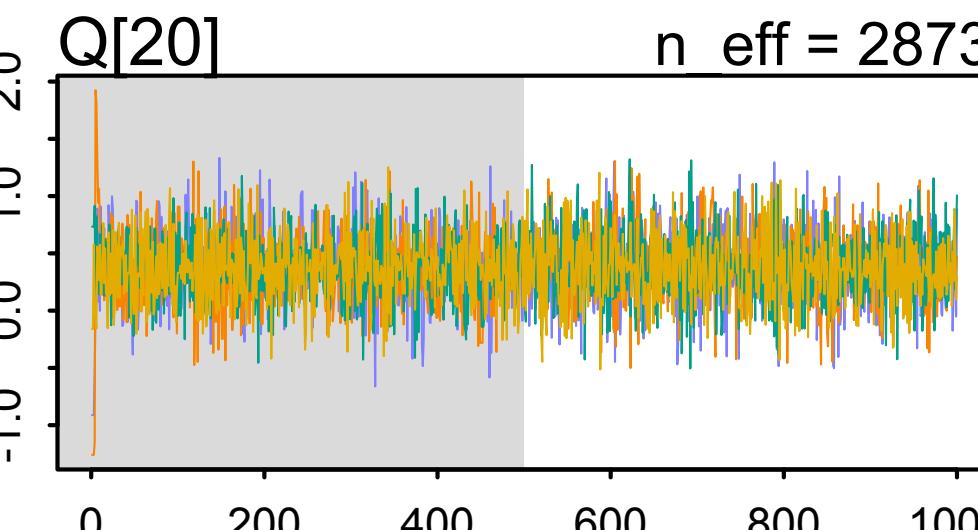
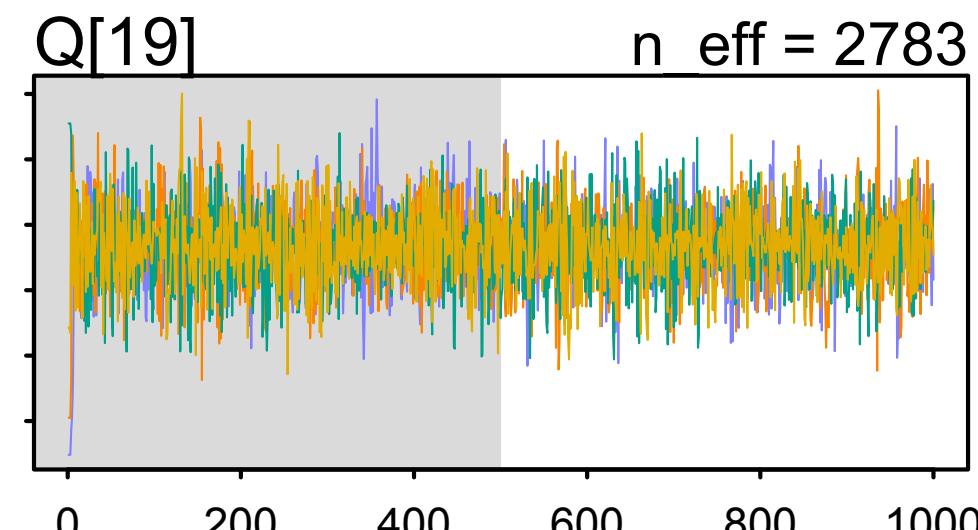
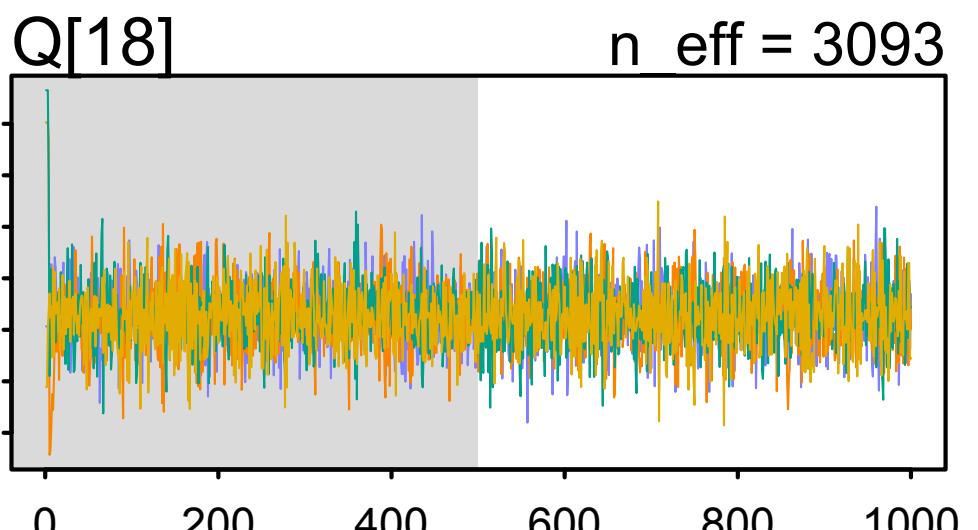
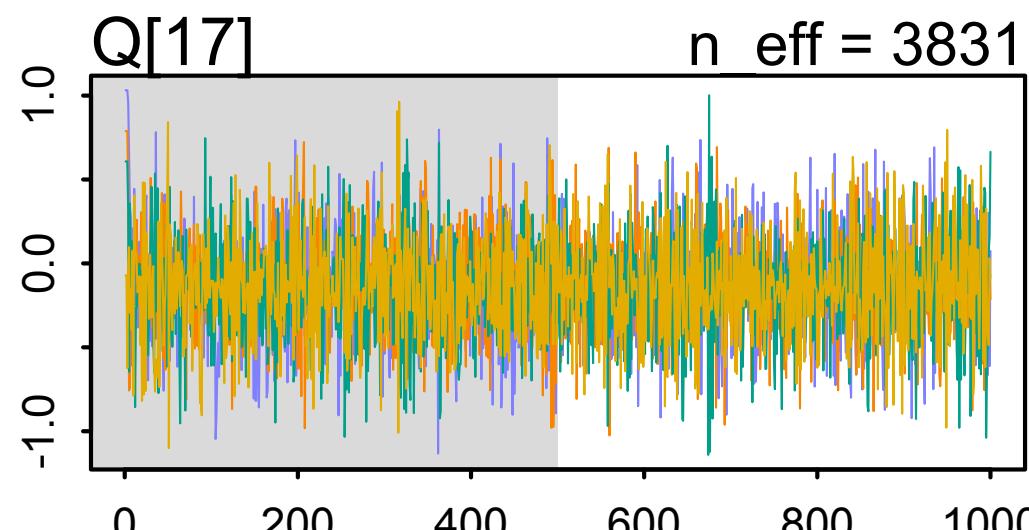
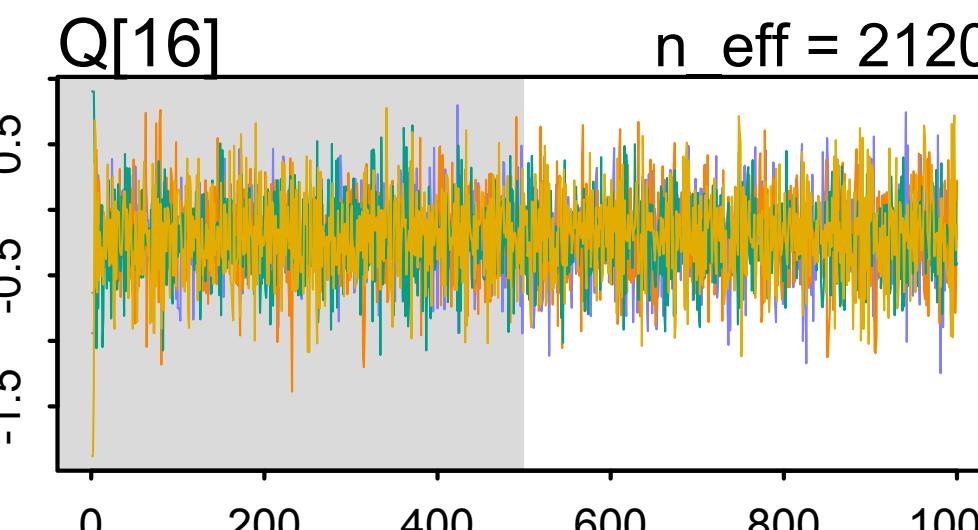
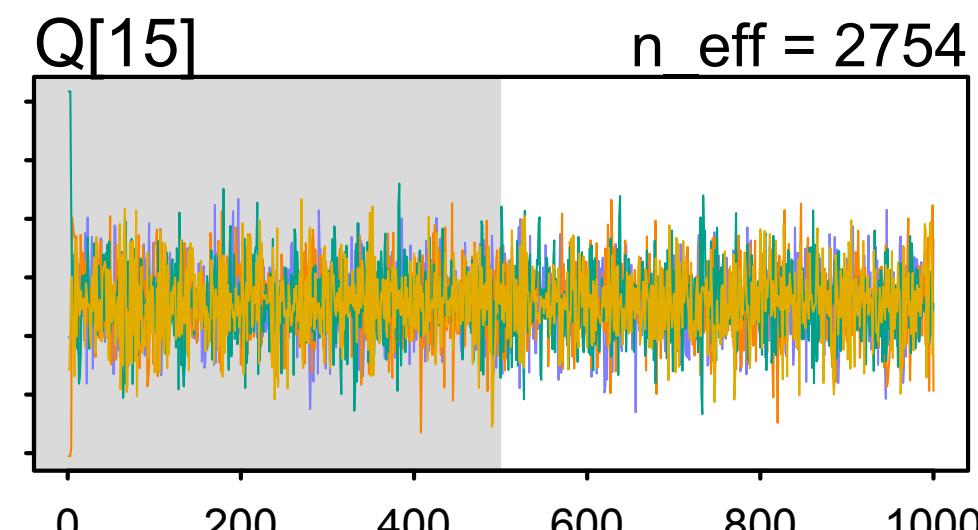
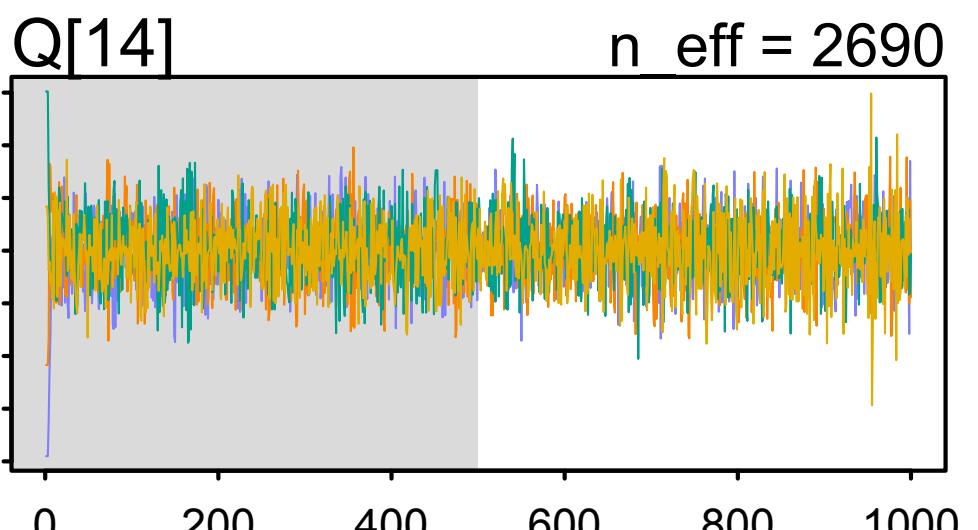
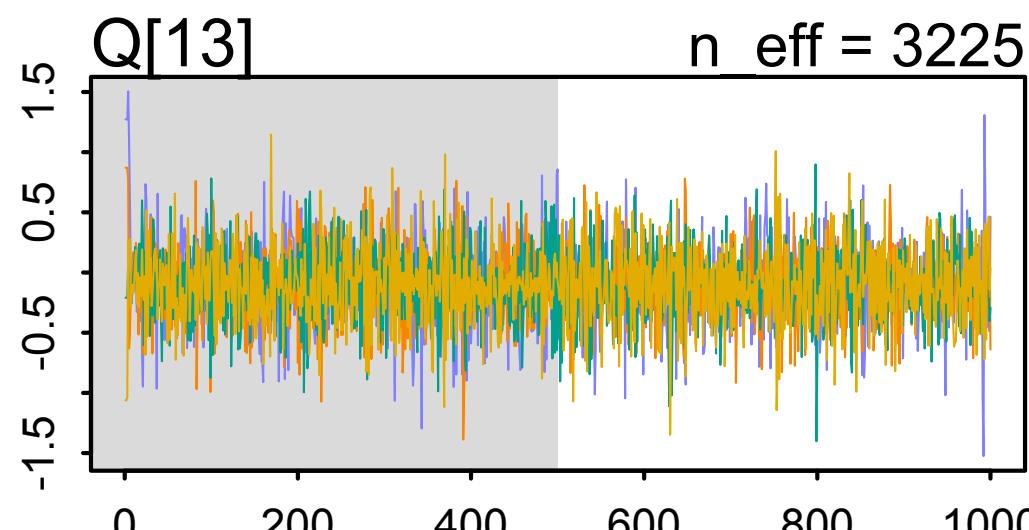
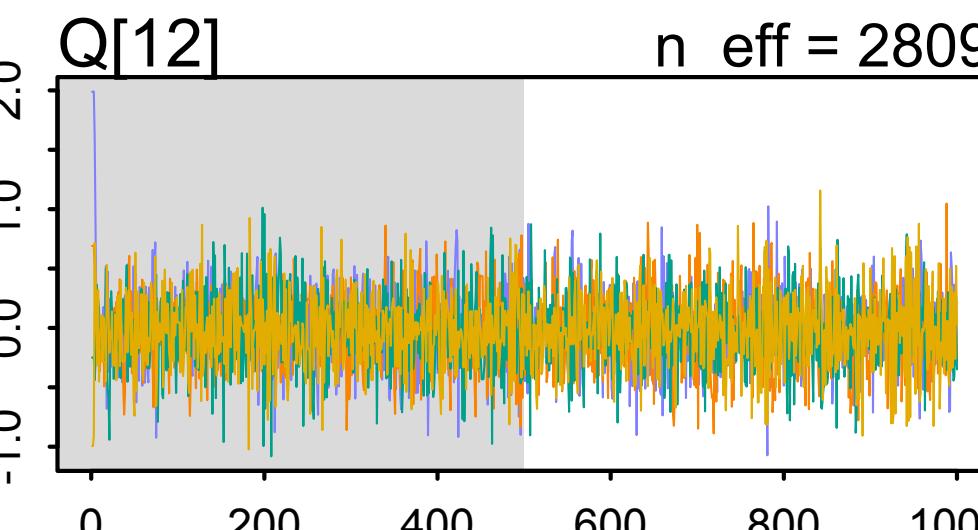
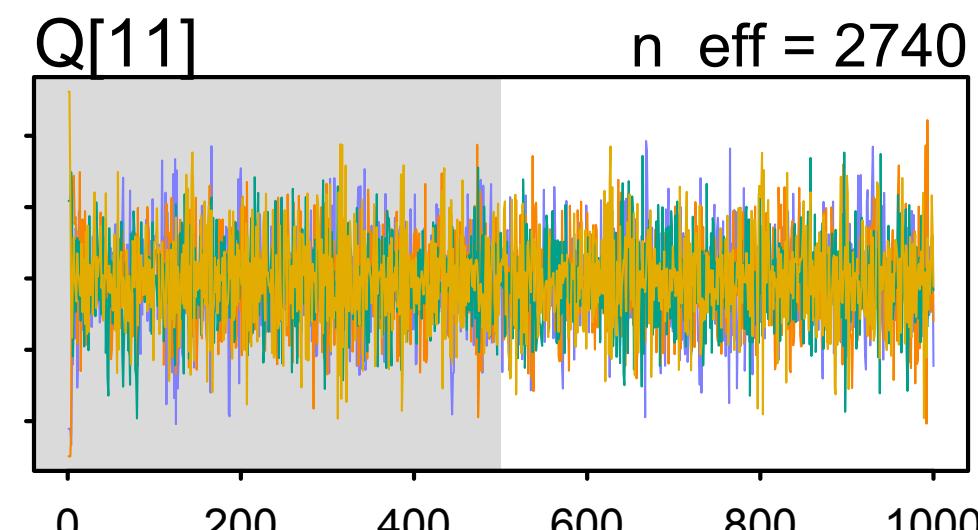
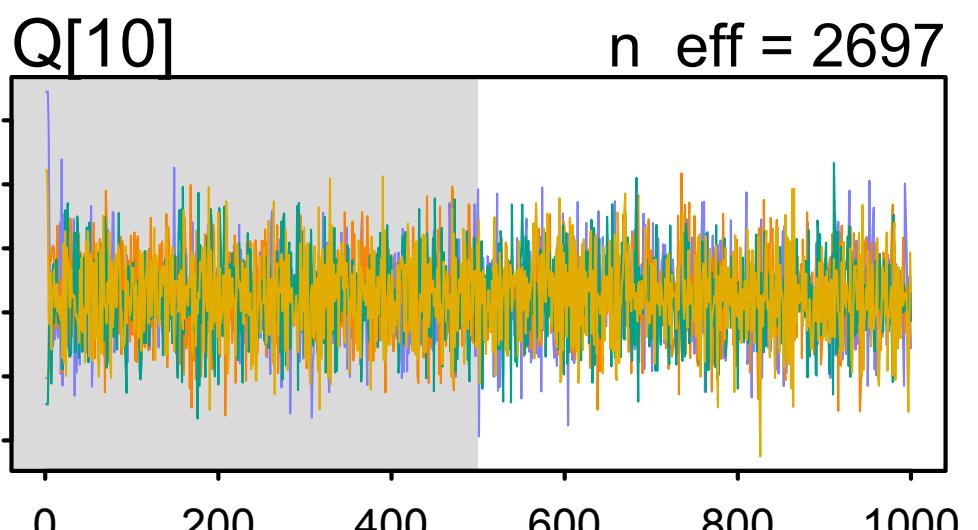
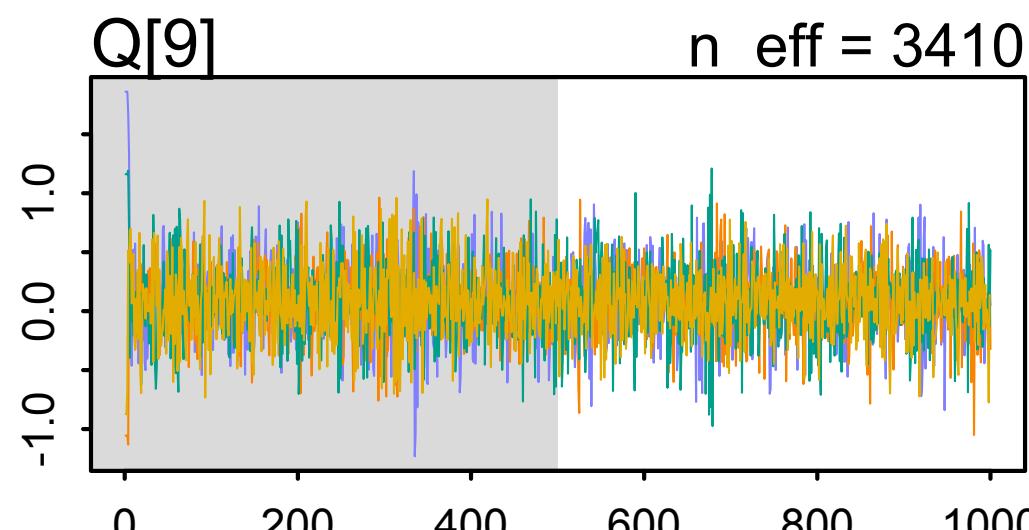
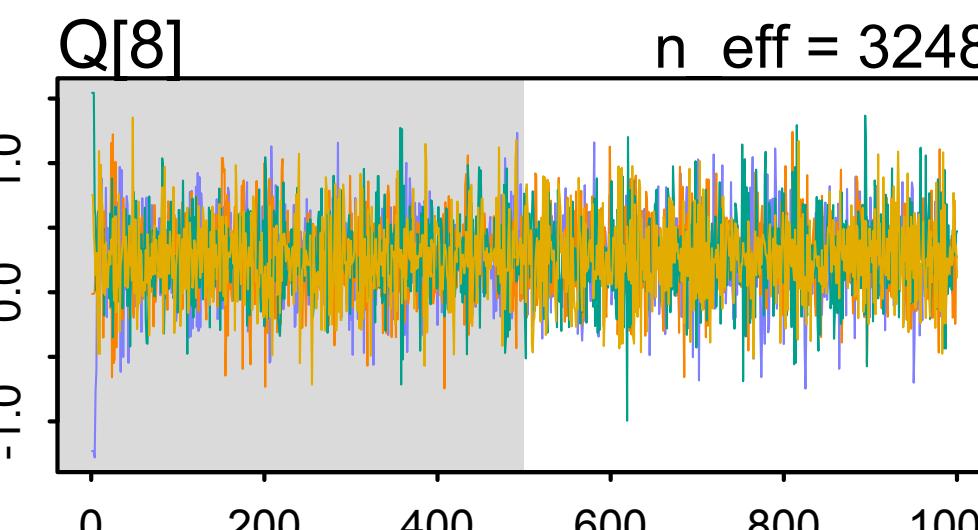
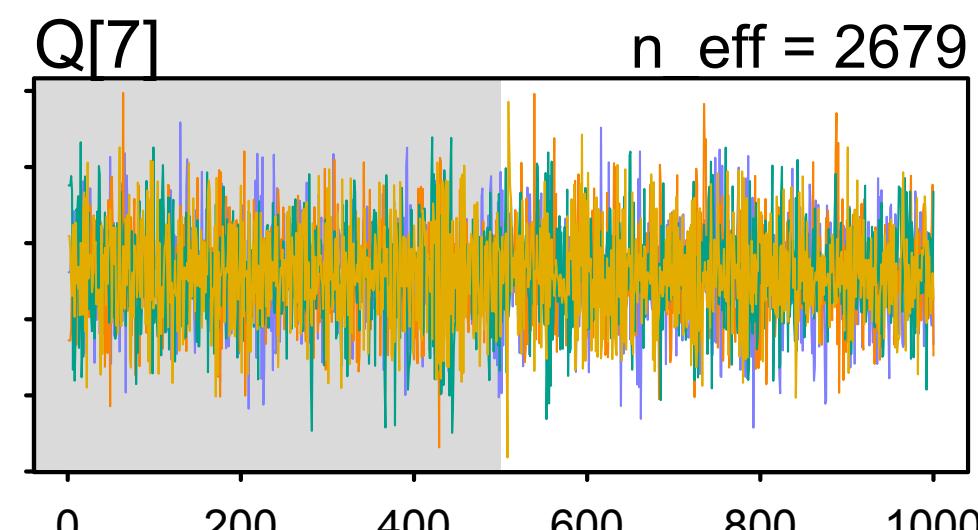
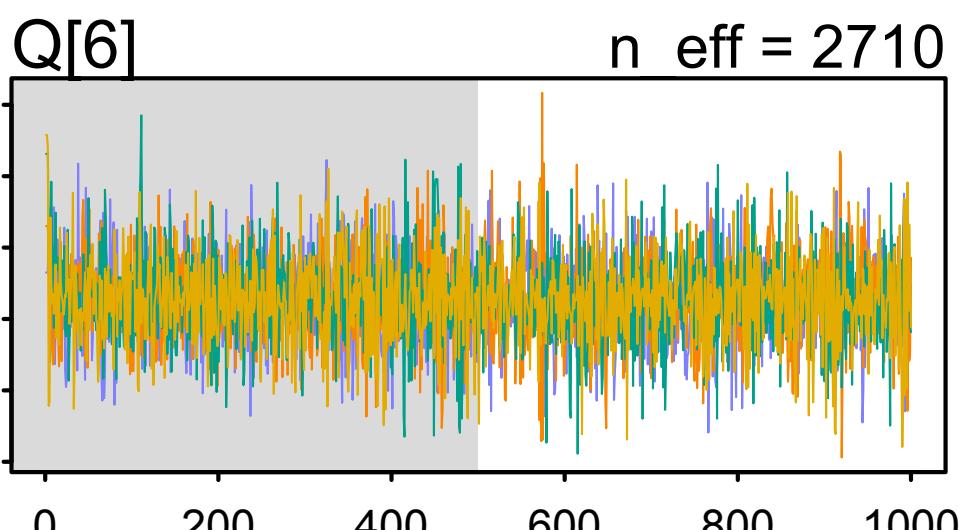
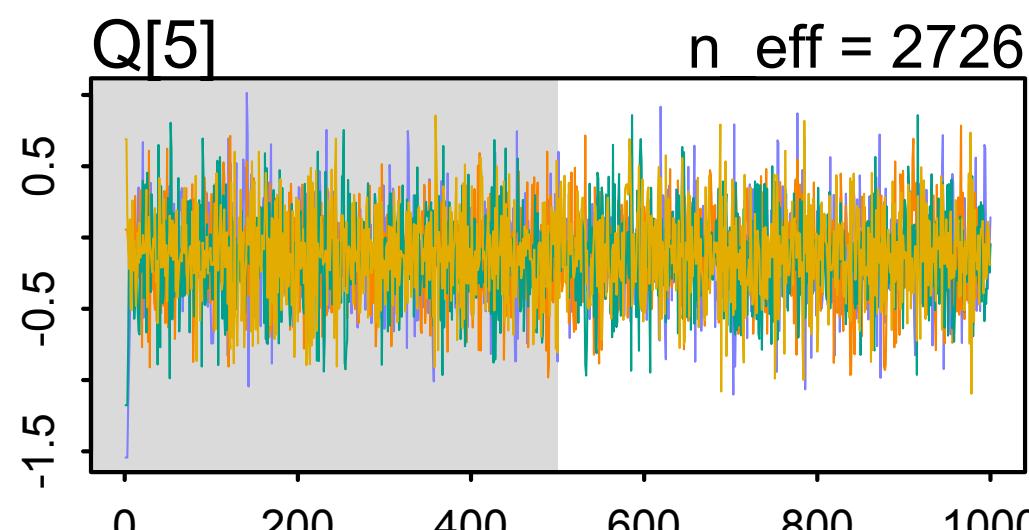
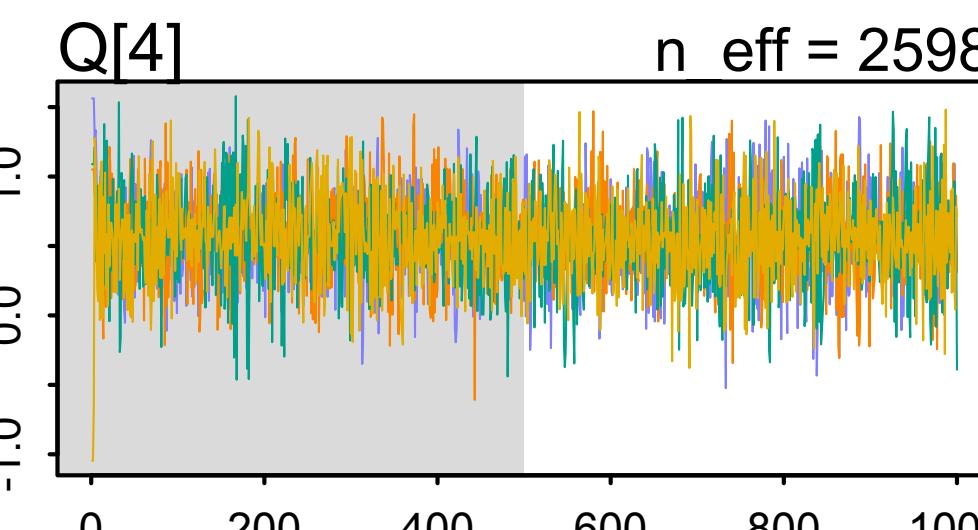
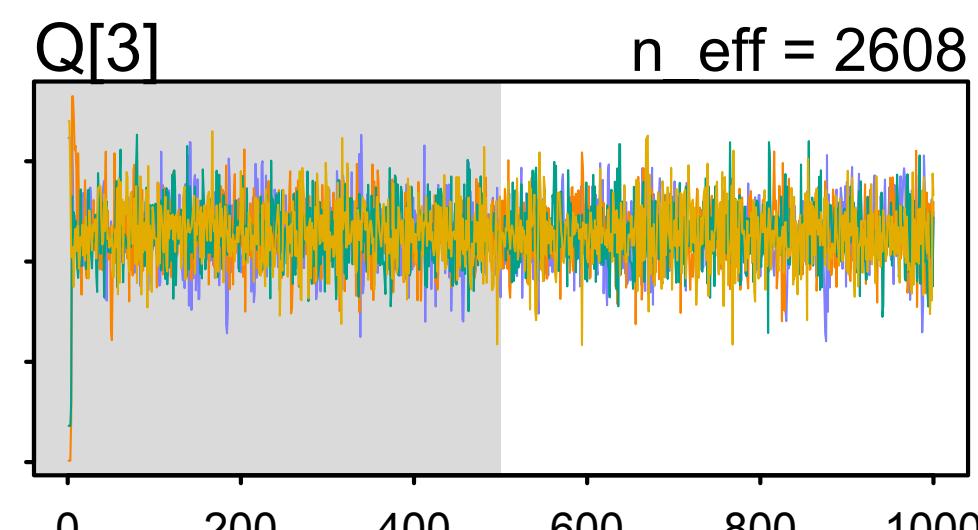
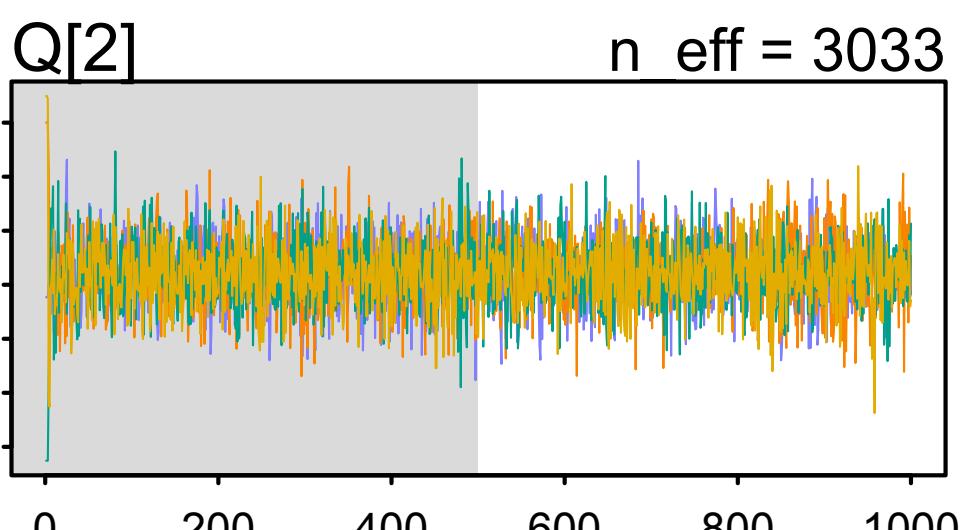
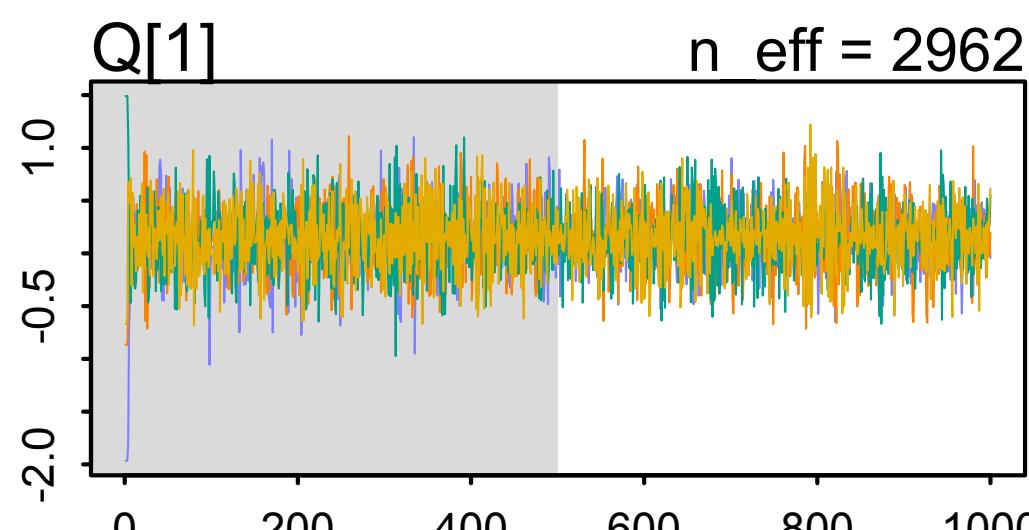
precis(mQ,2)
```

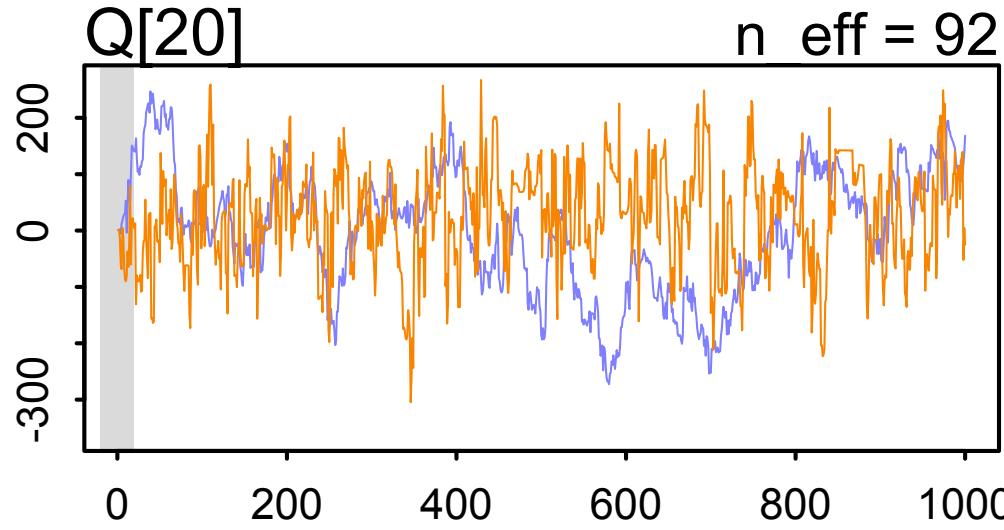
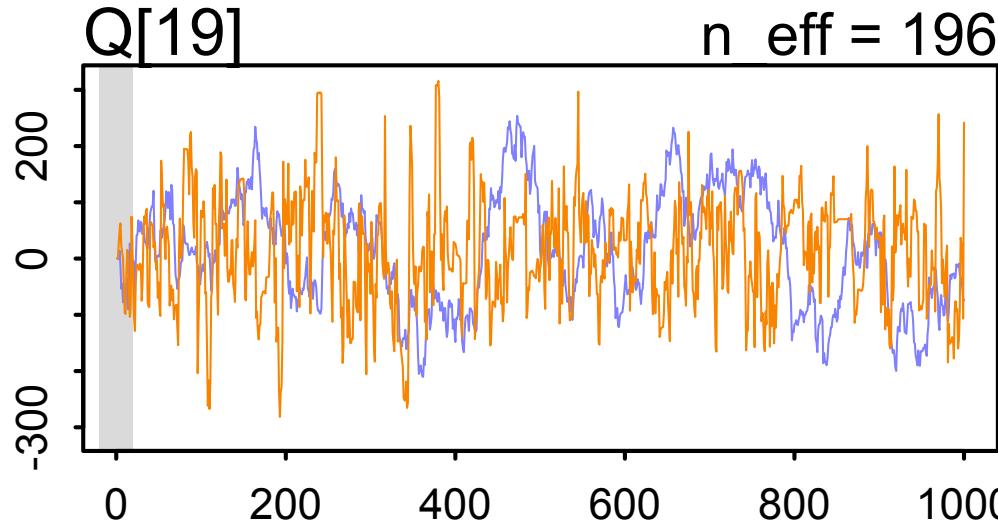
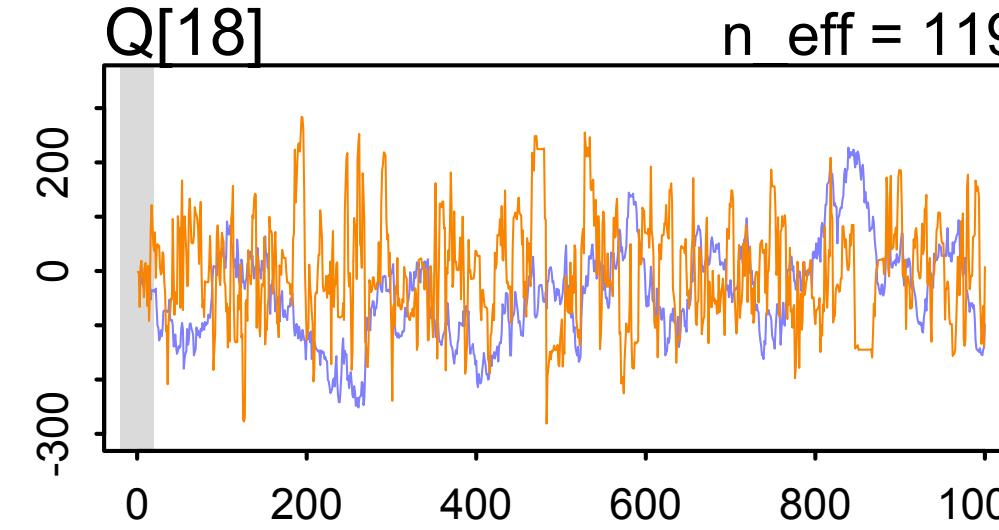
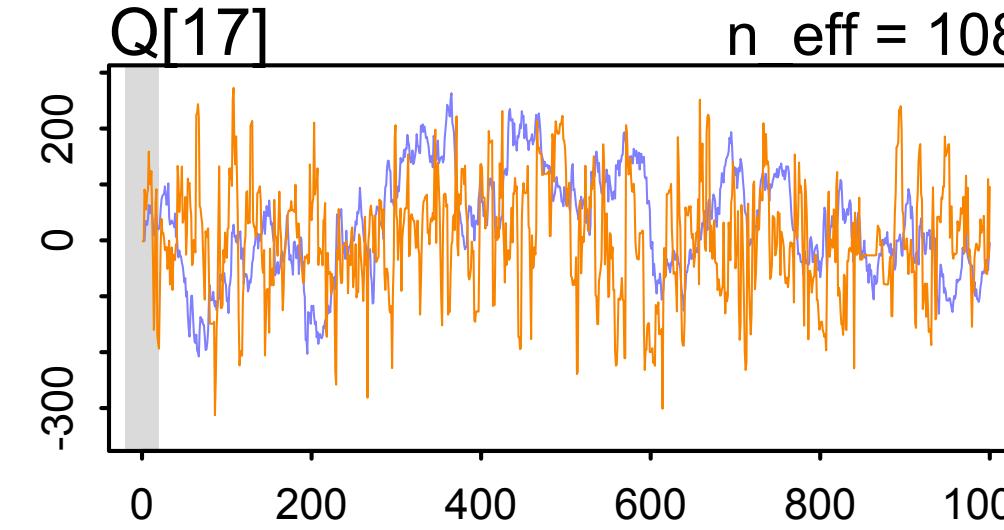
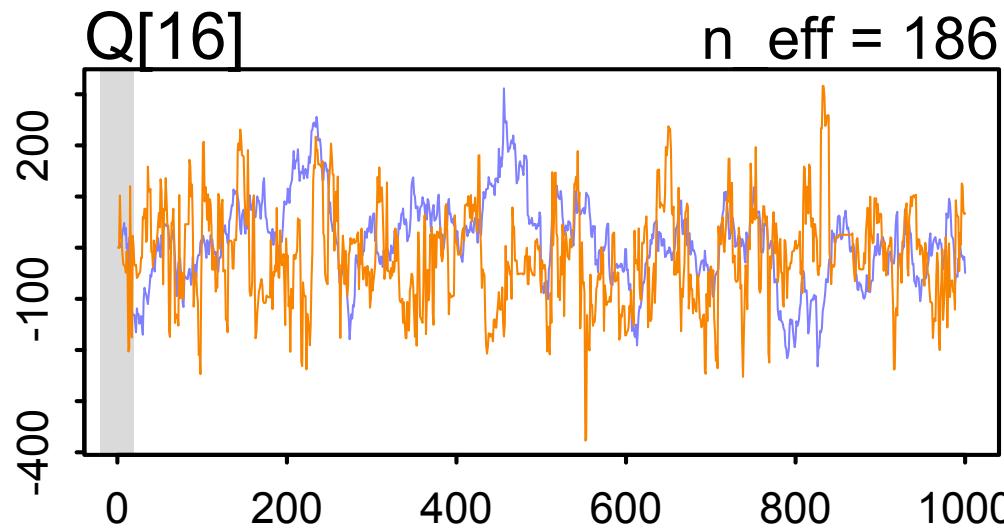
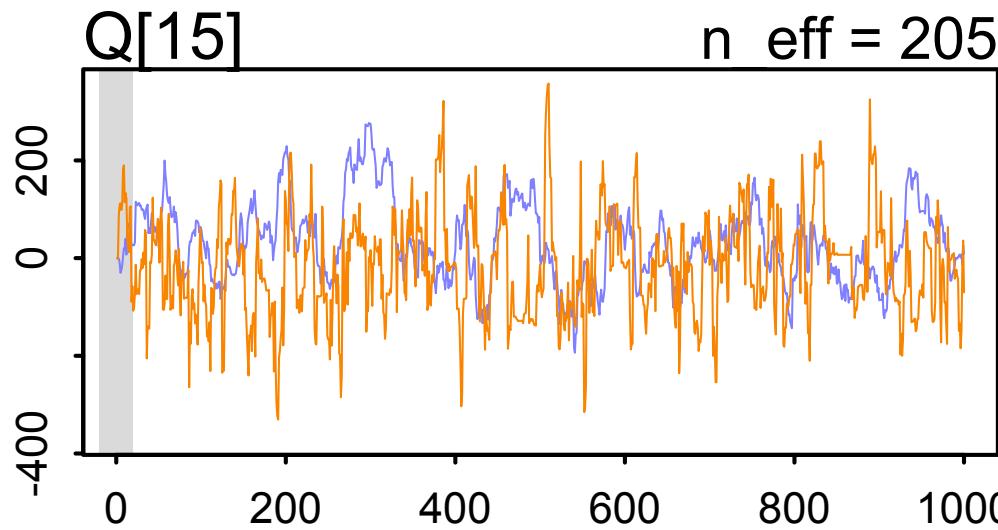
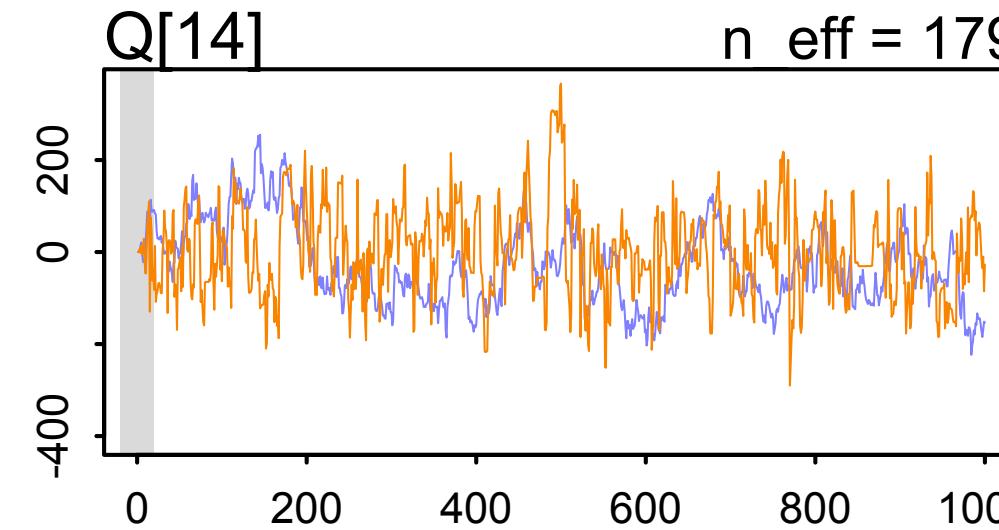
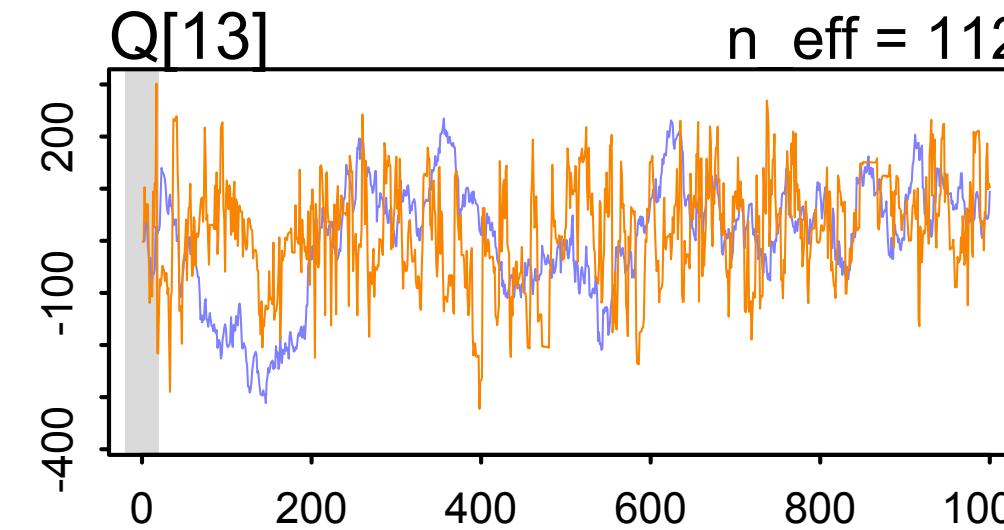
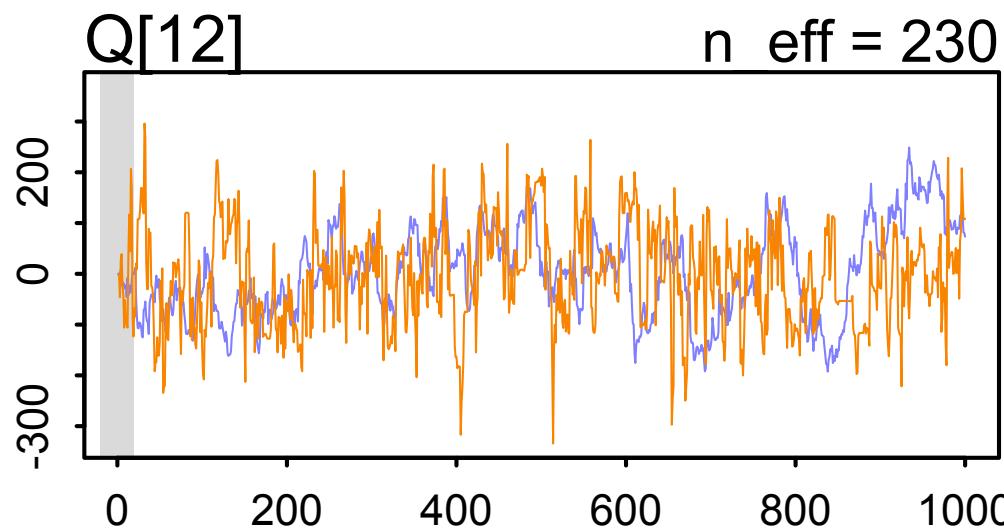
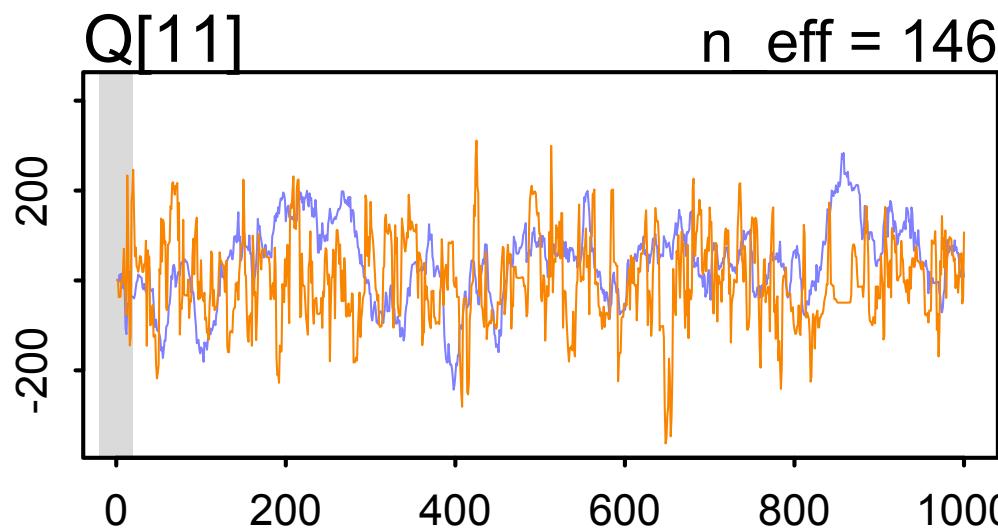
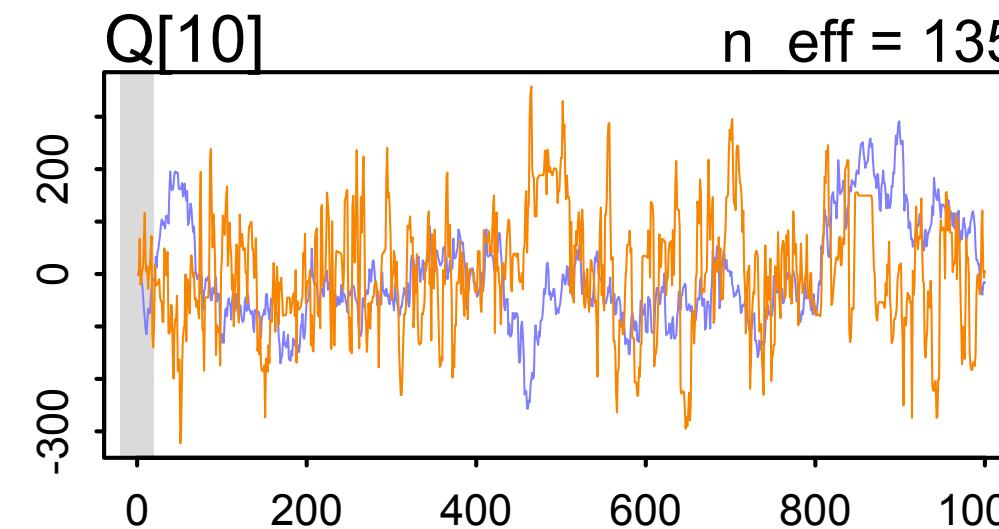
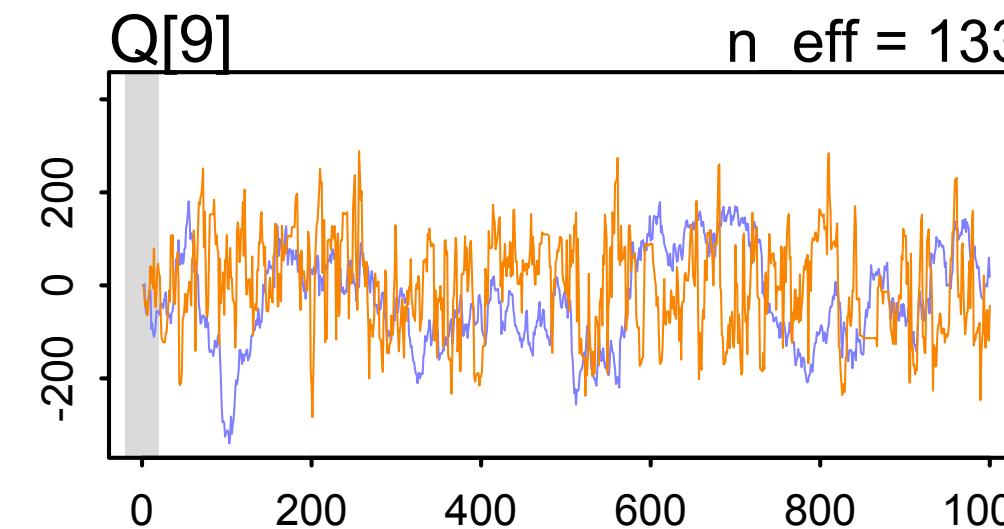
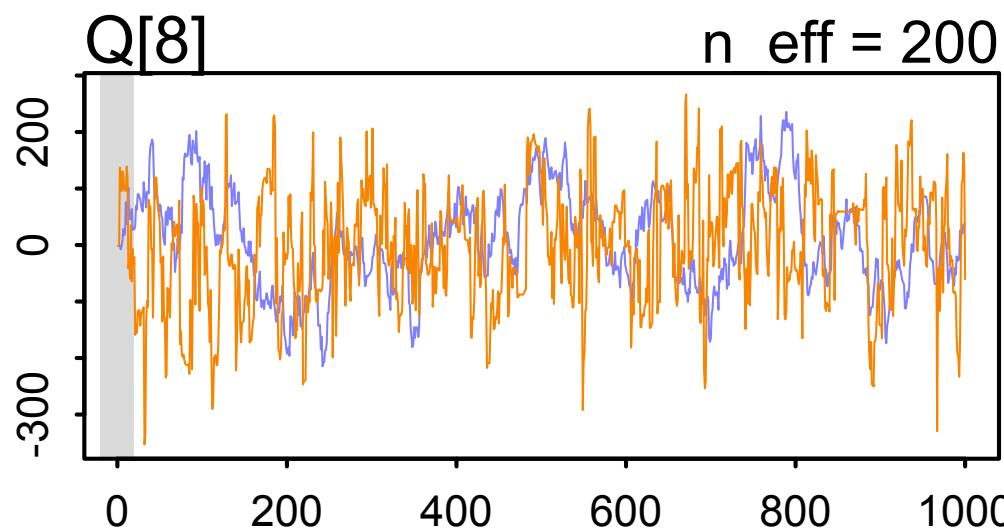
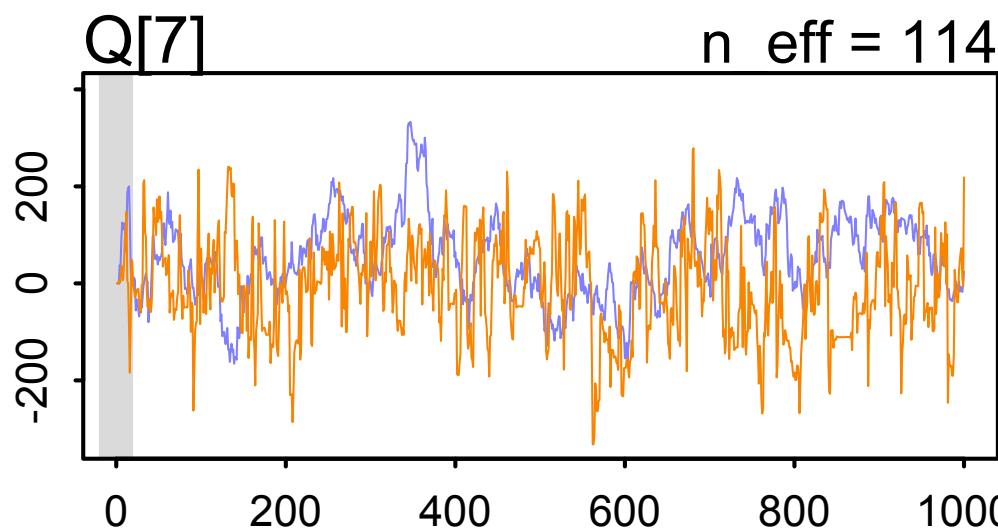
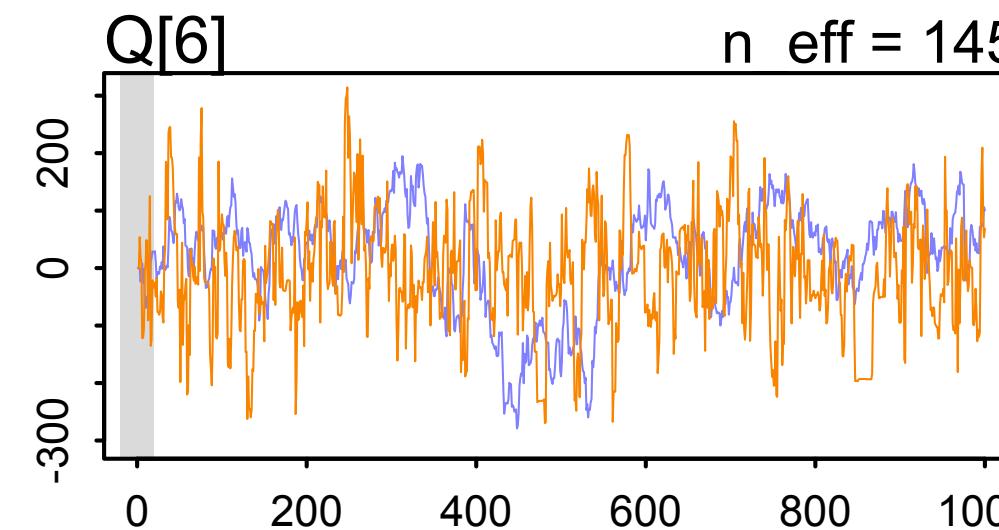
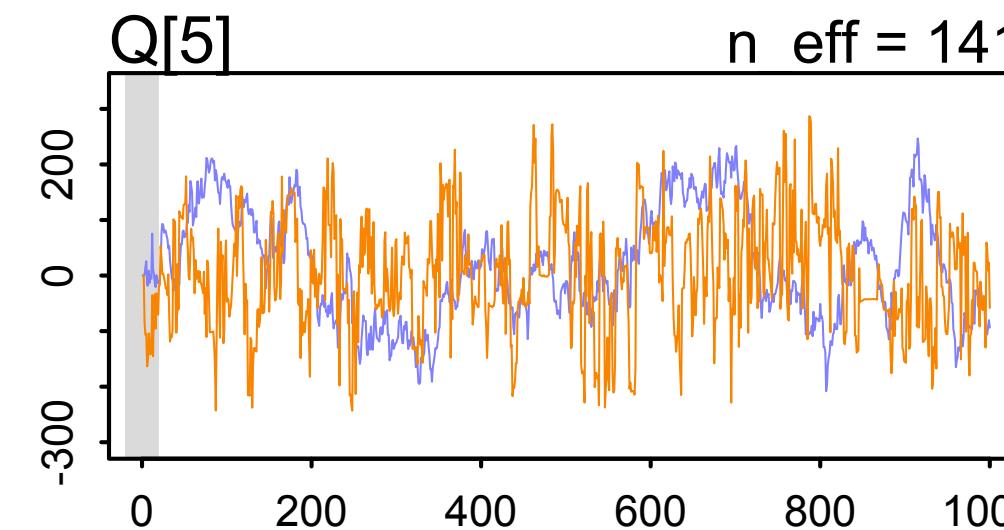
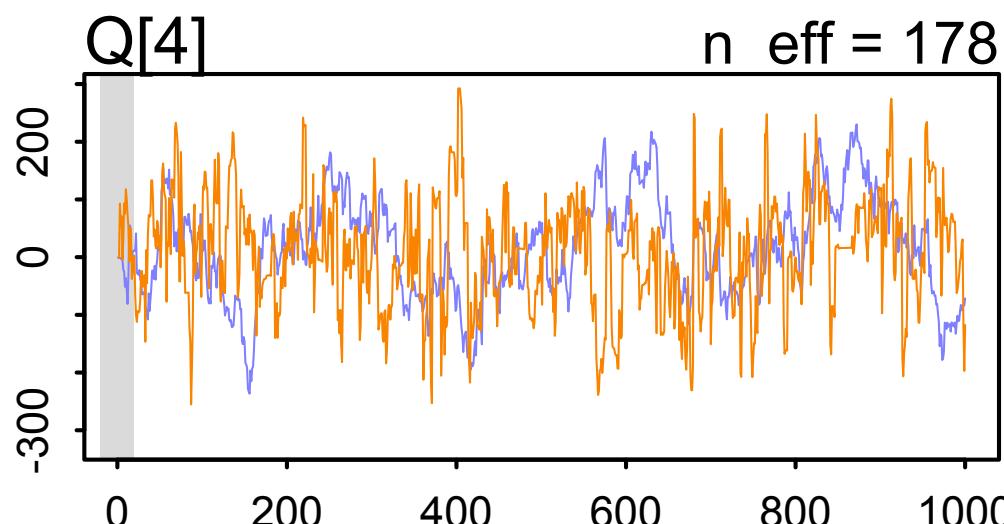
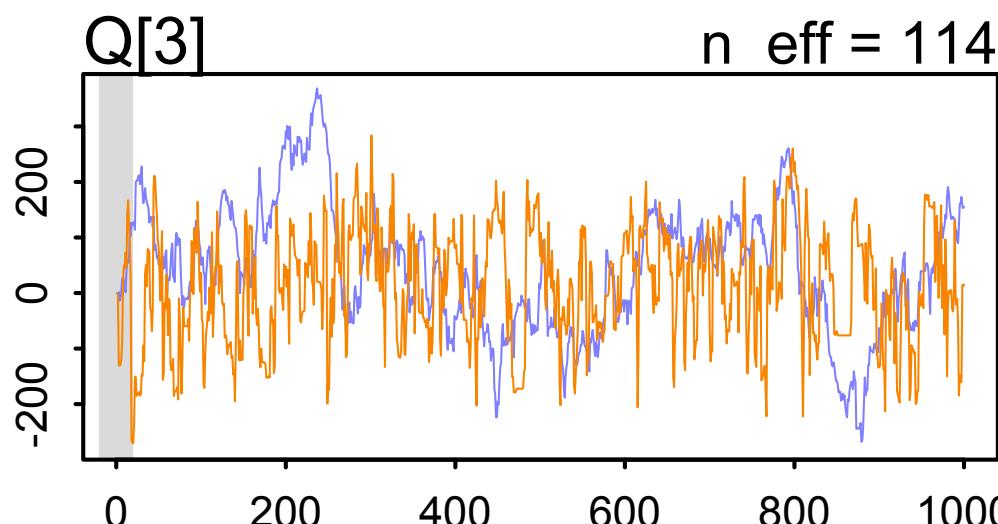
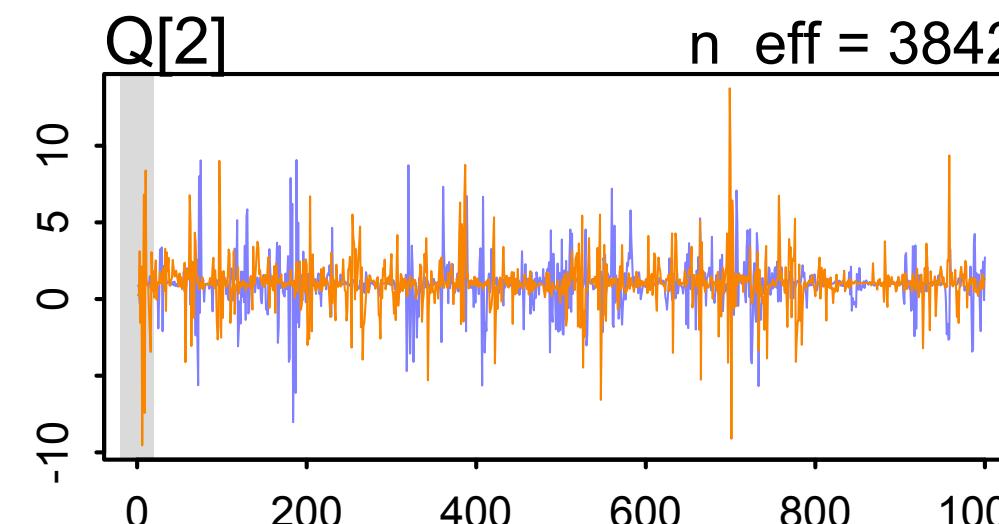
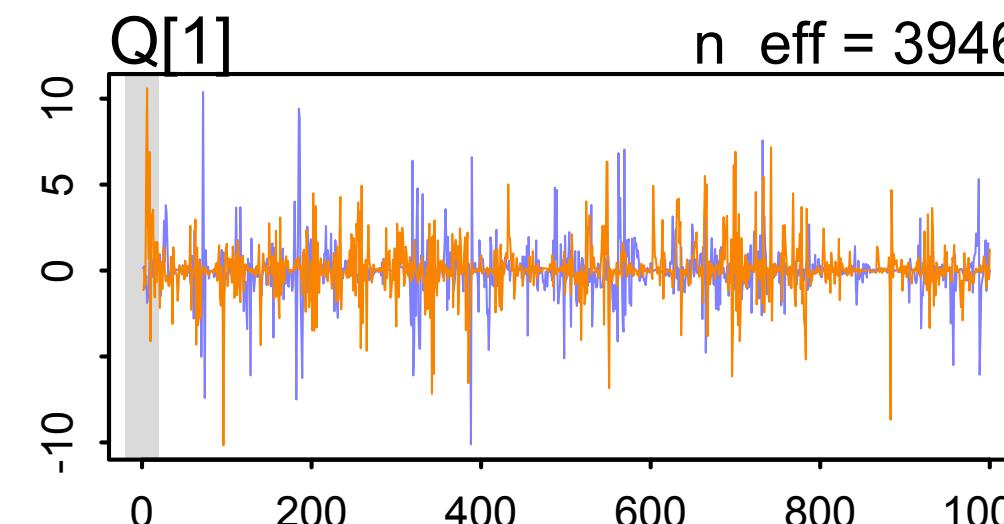
Need more than 1 chain to check convergence

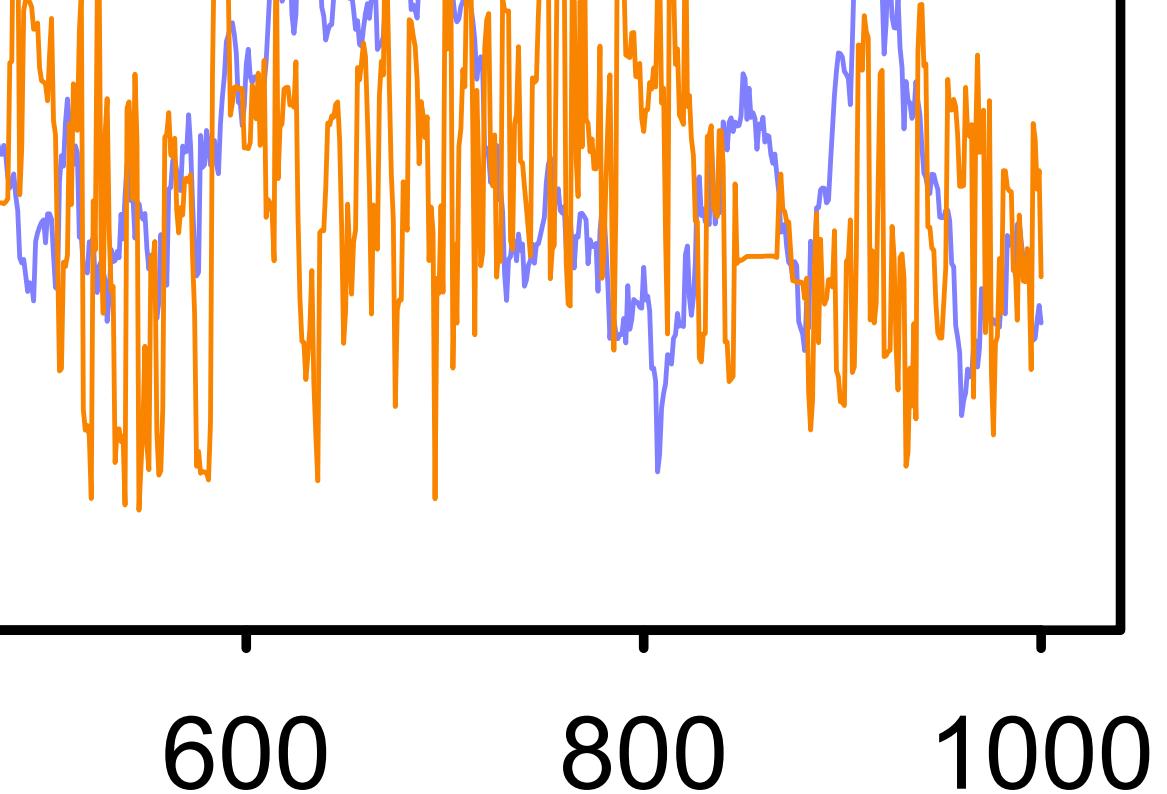
Convergence: Each chain explores the right distribution and every chain explores the same distribution



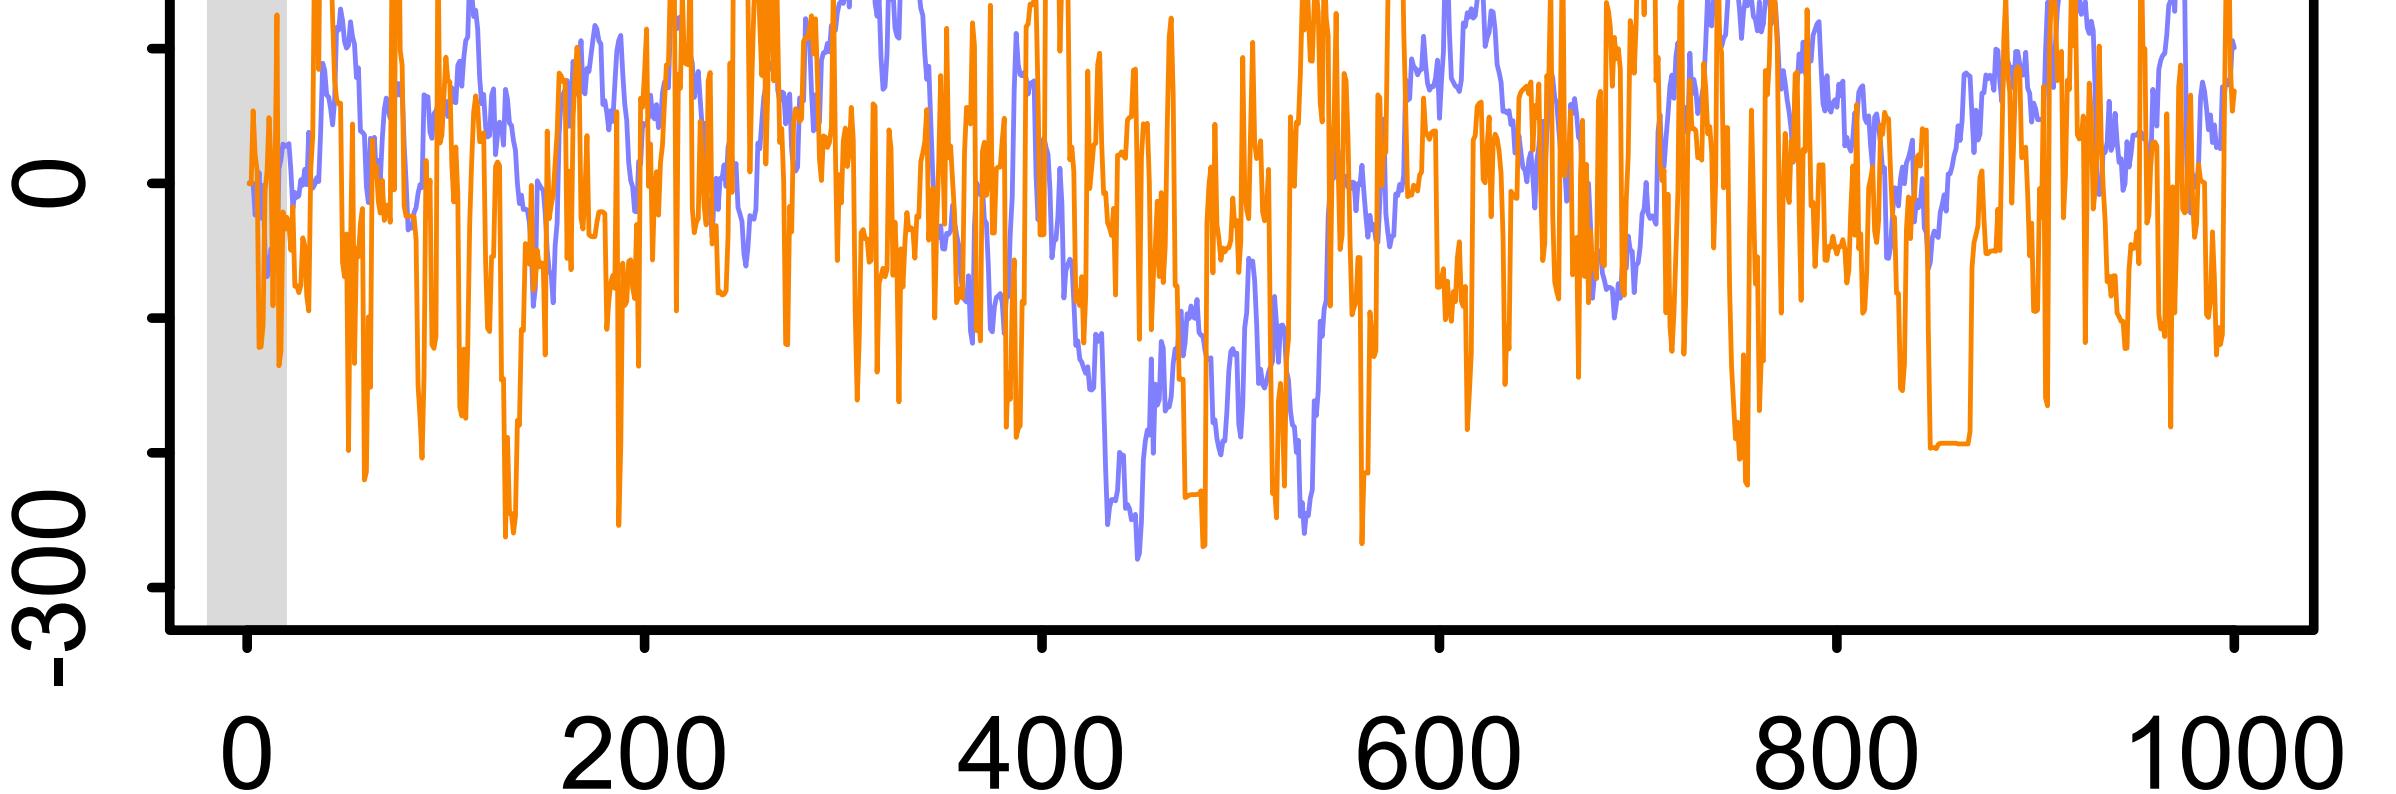






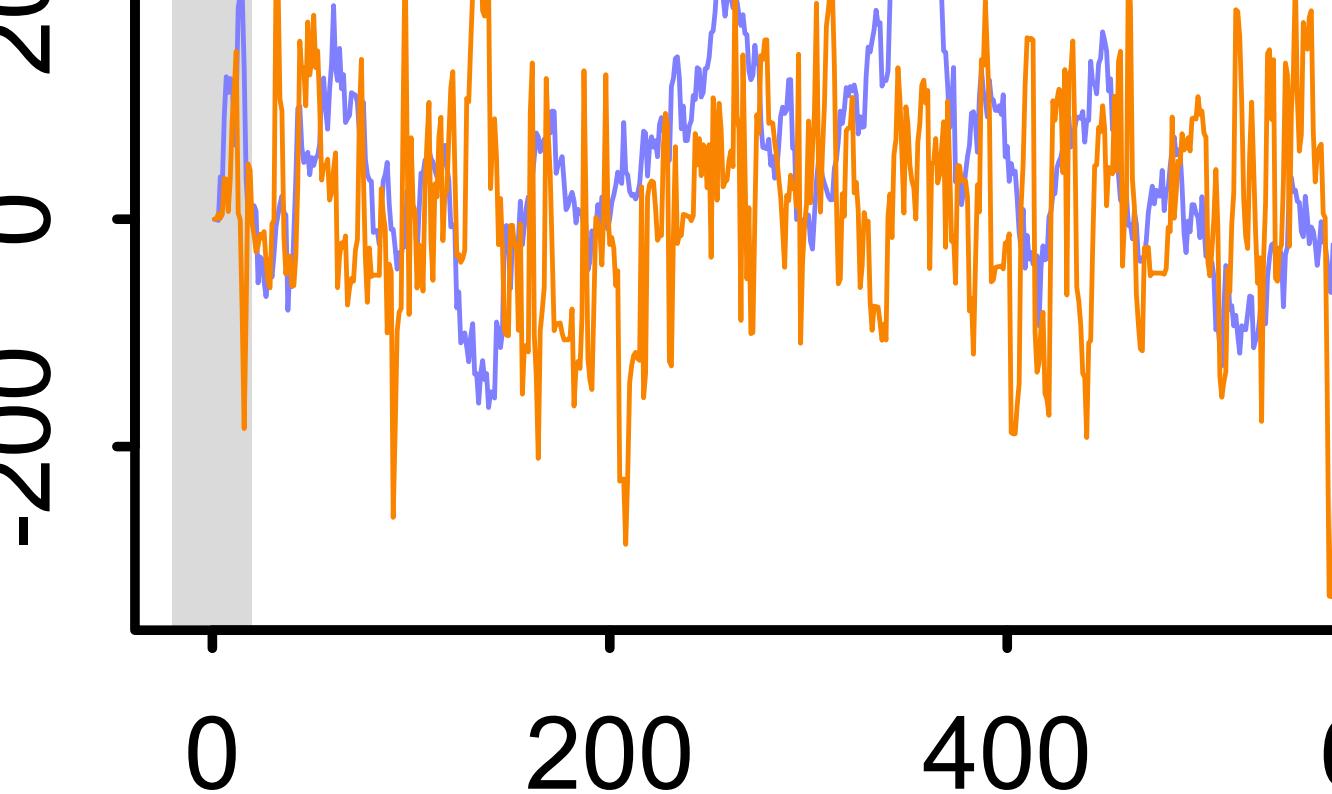


$n \text{ eff} = 133$

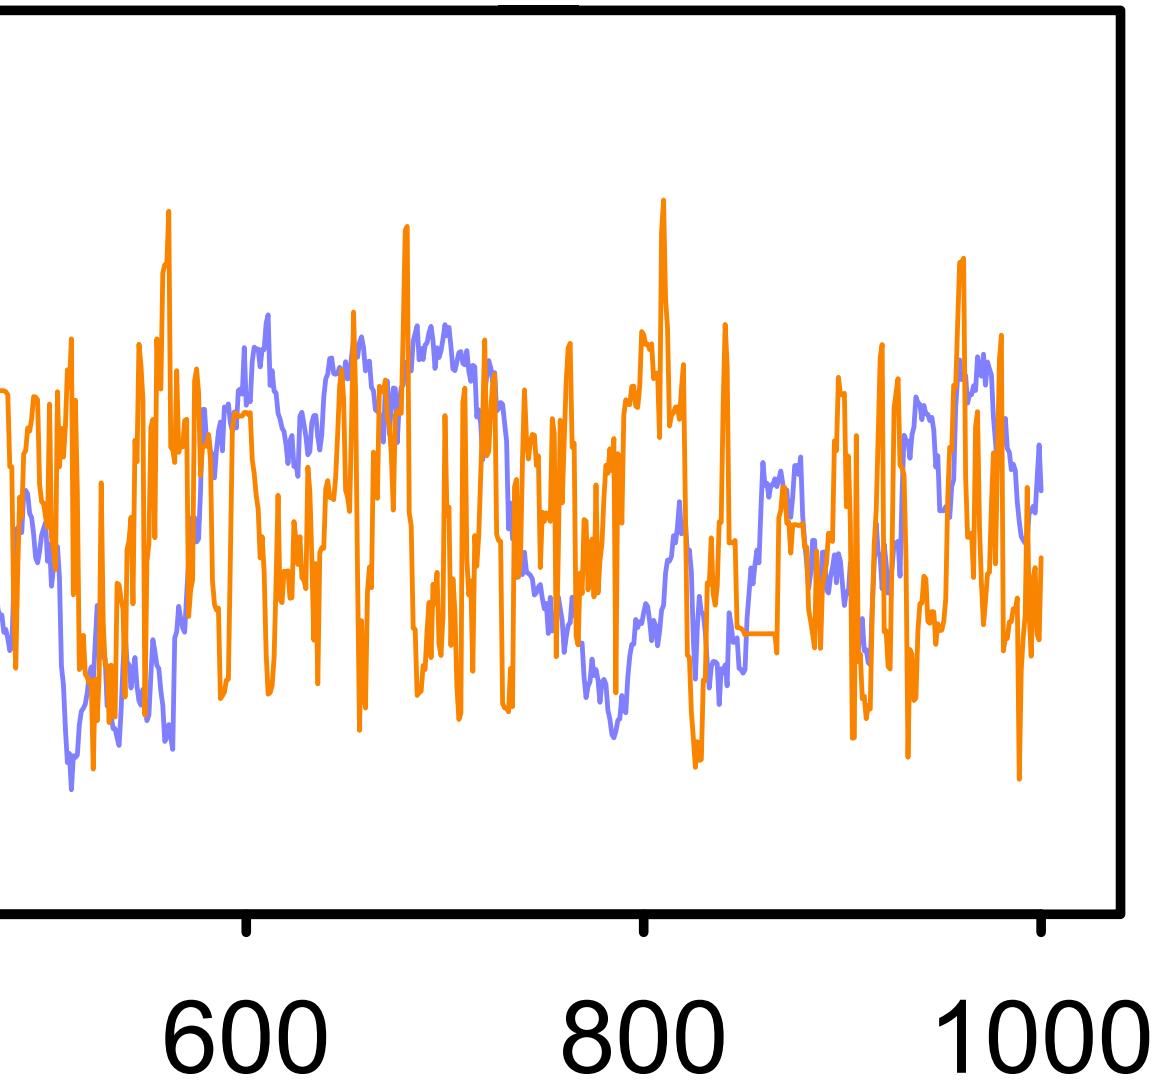


$Q[10]$

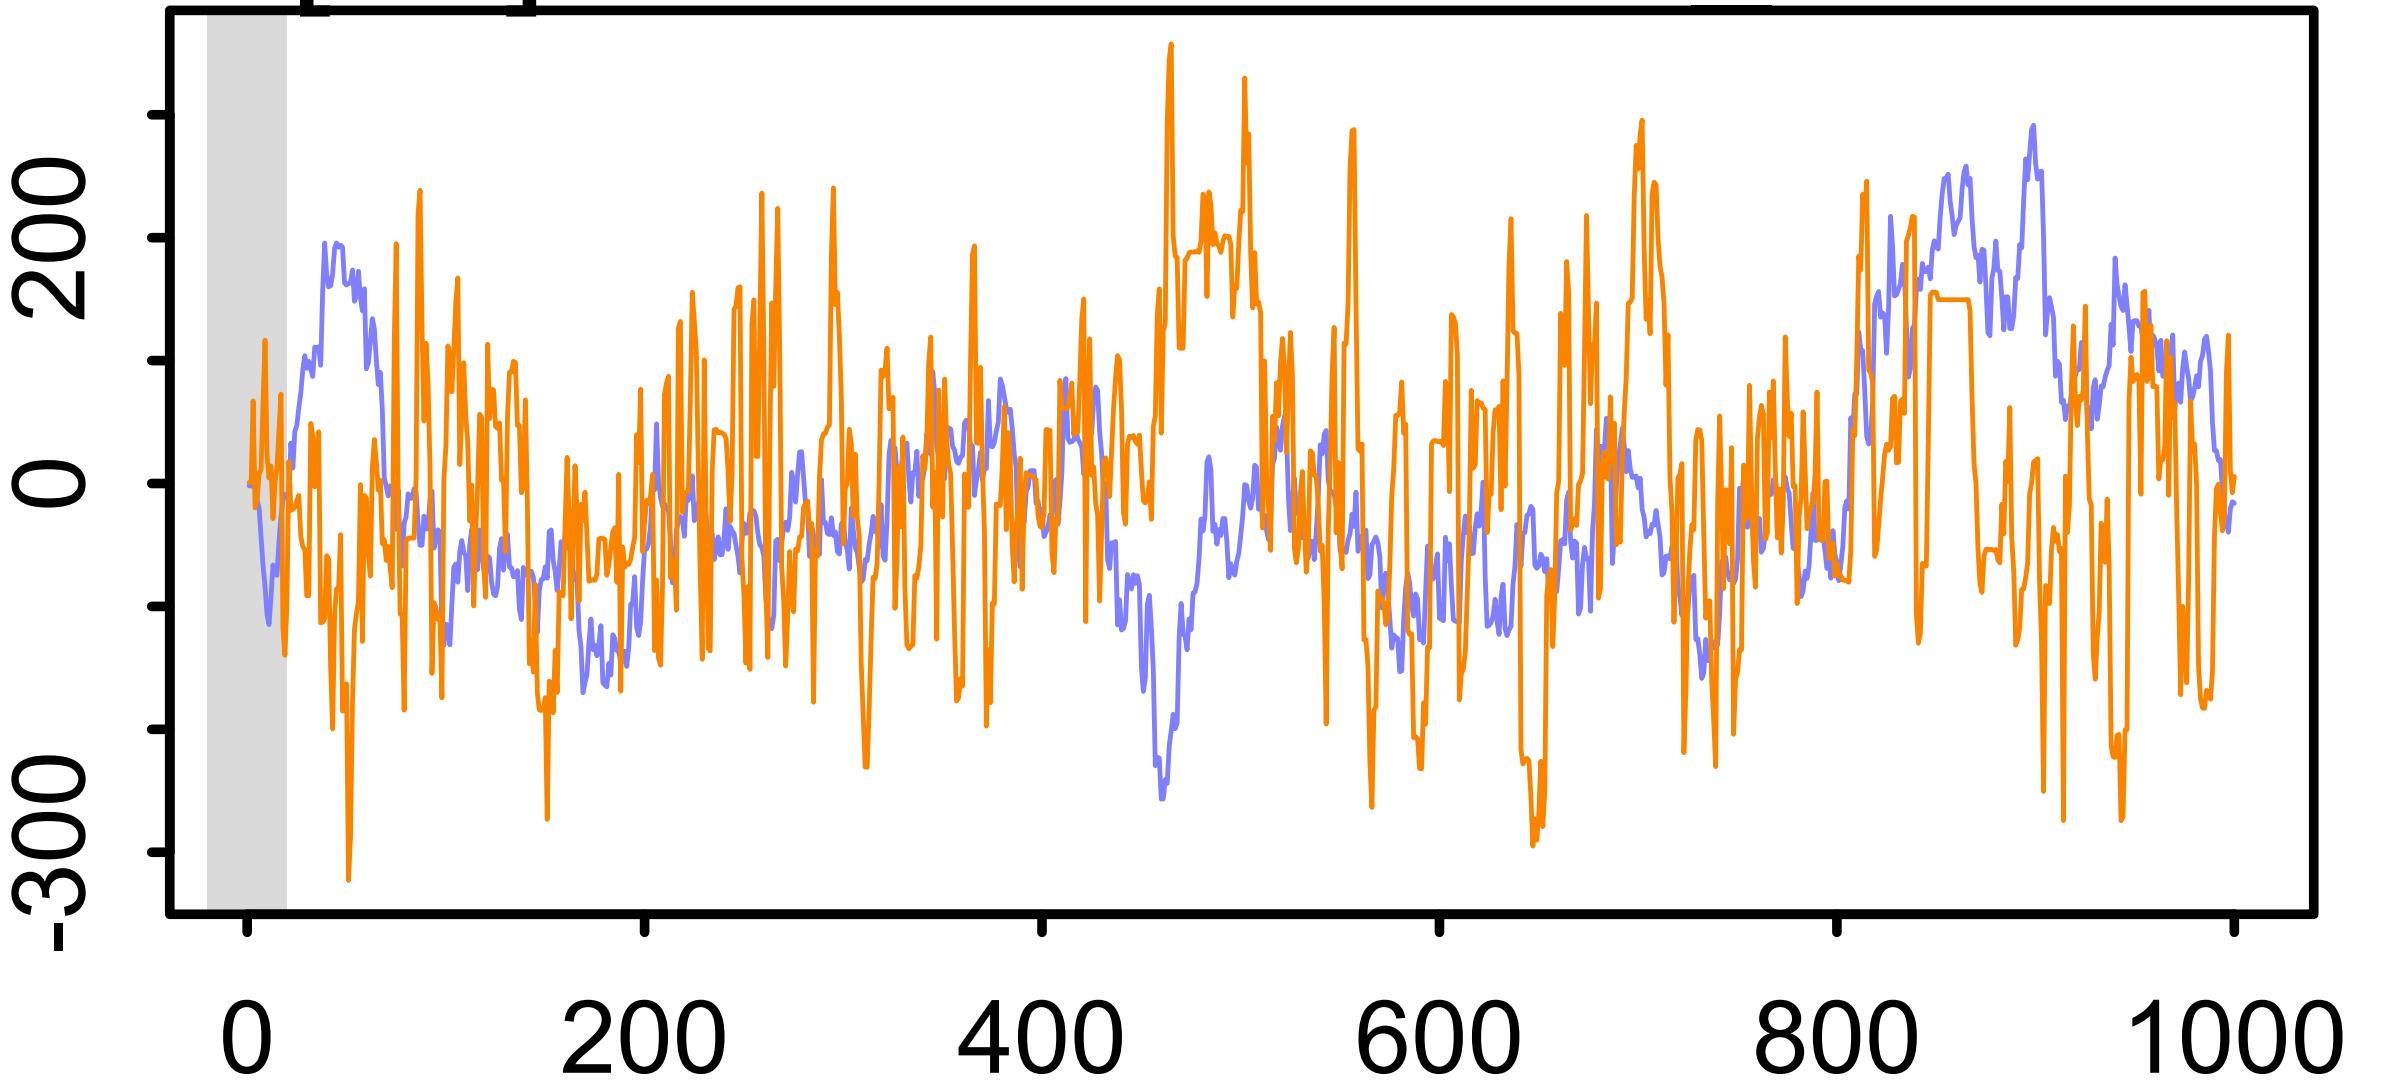
$n \text{ eff} = 135$



$Q[11]$

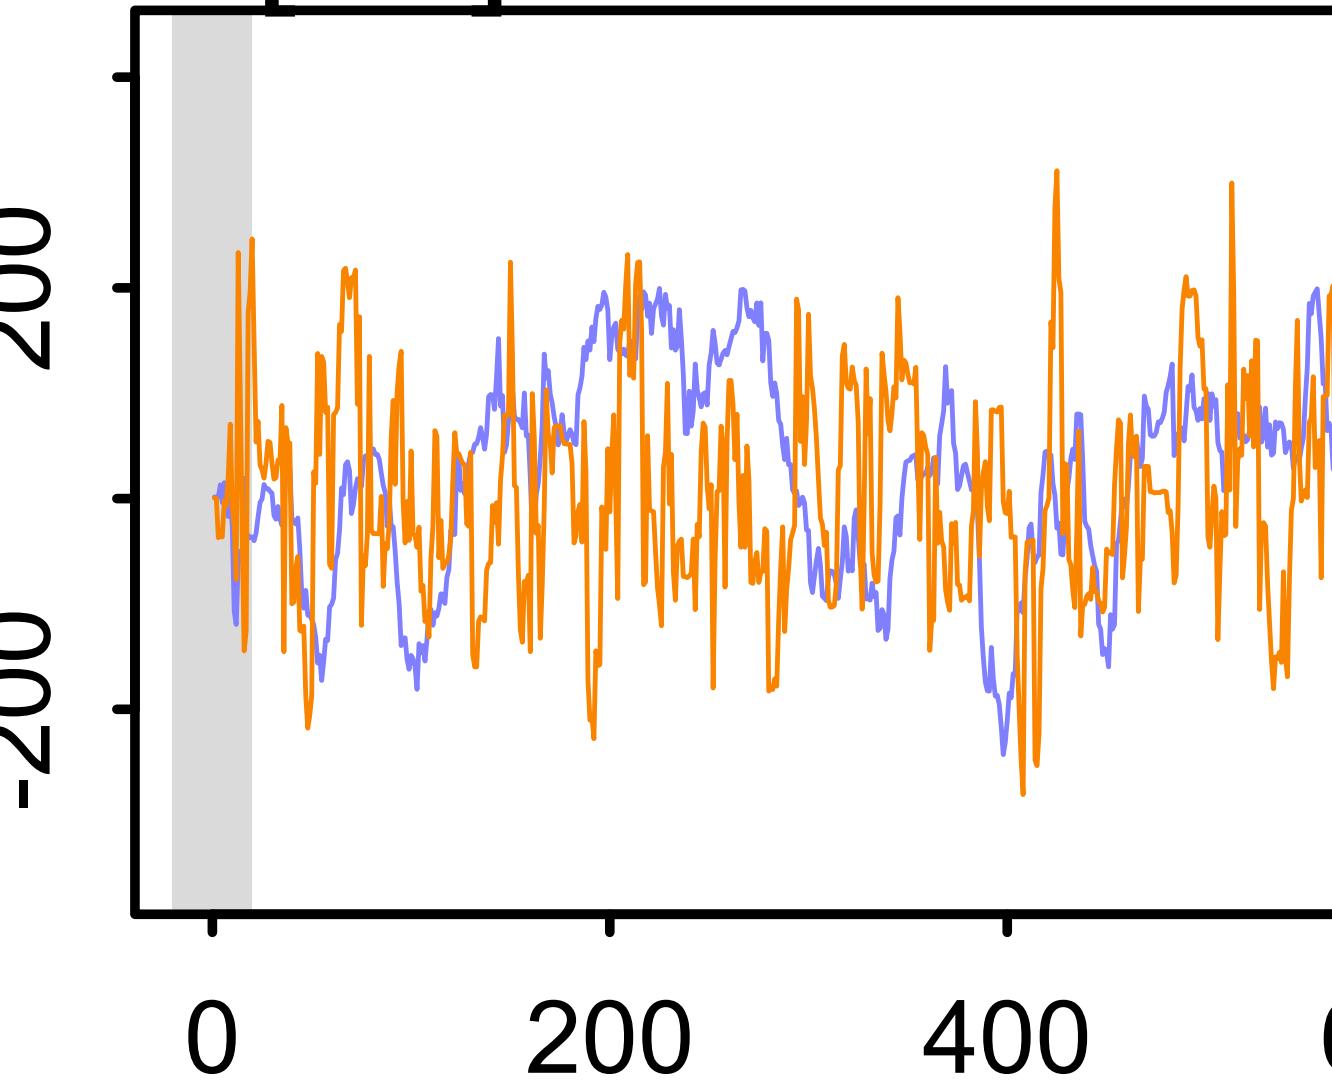


$n \text{ eff} = 112$

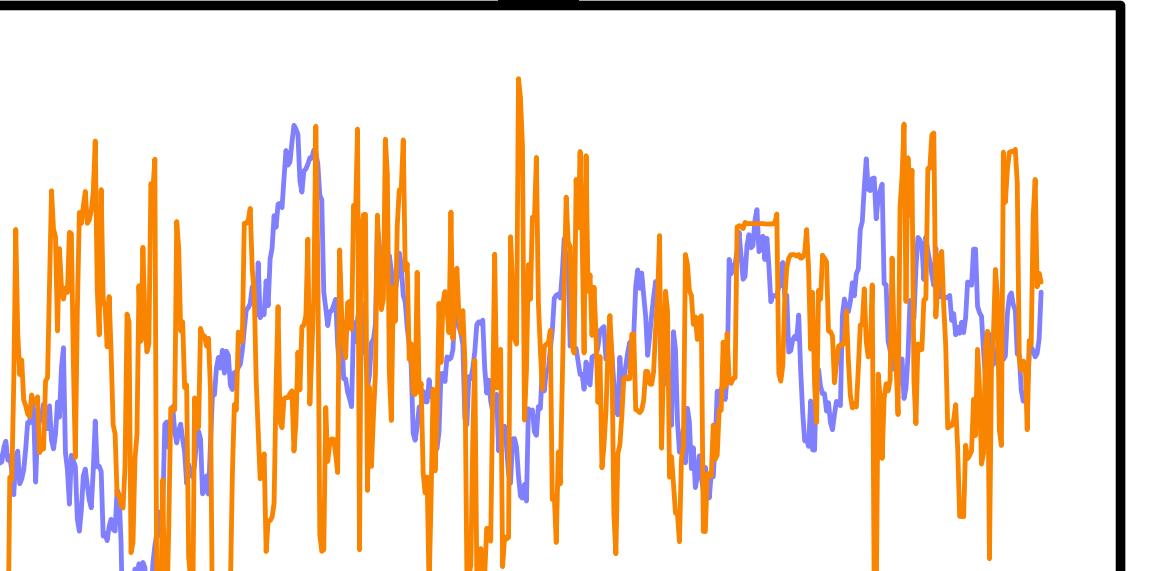


$Q[14]$

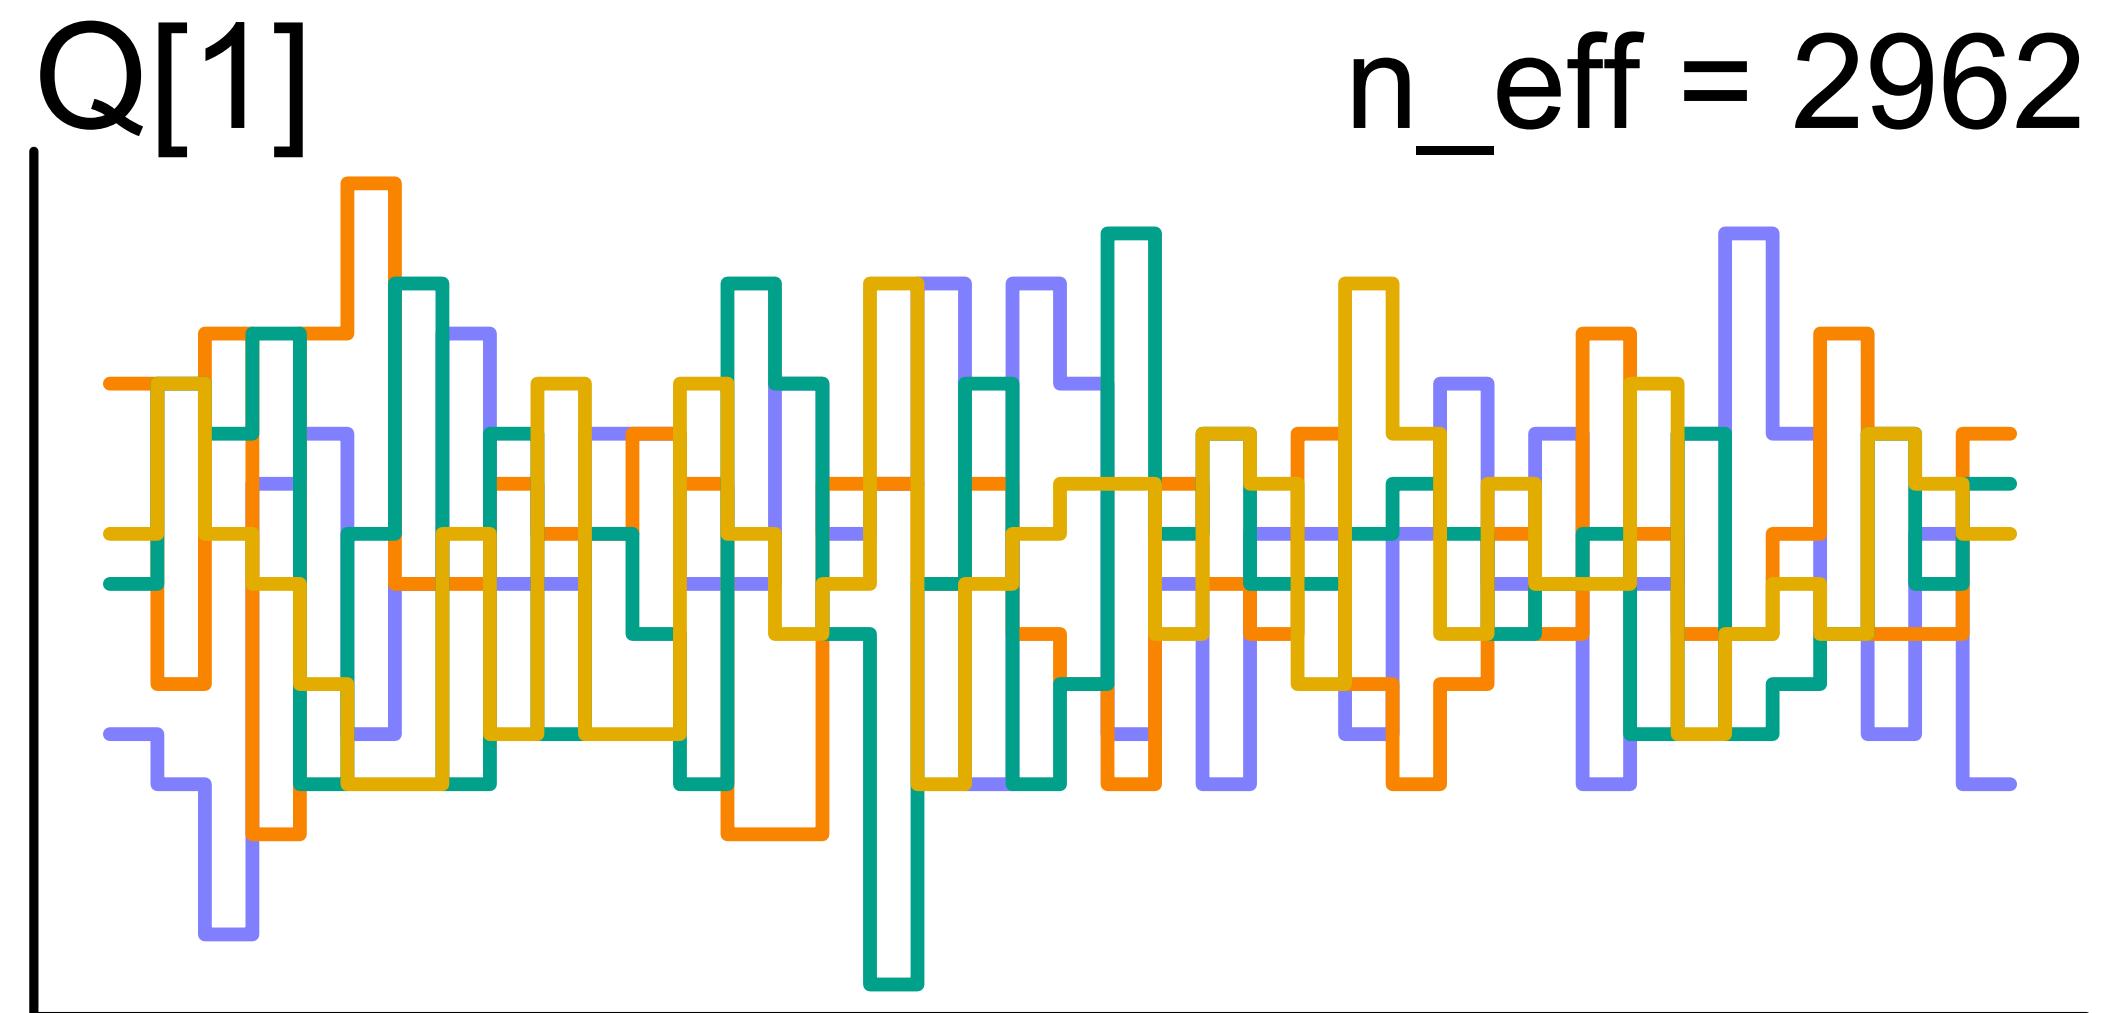
$n \text{ eff} = 179$



$Q[15]$

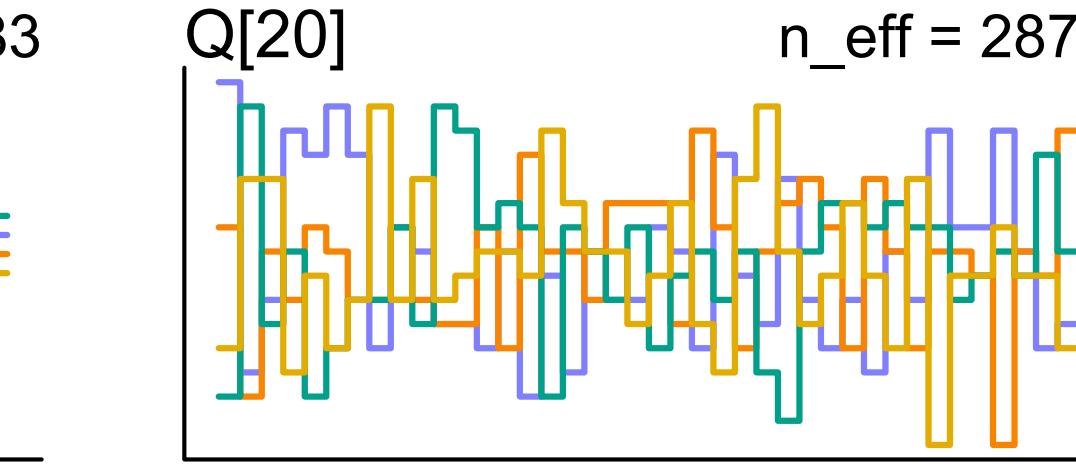
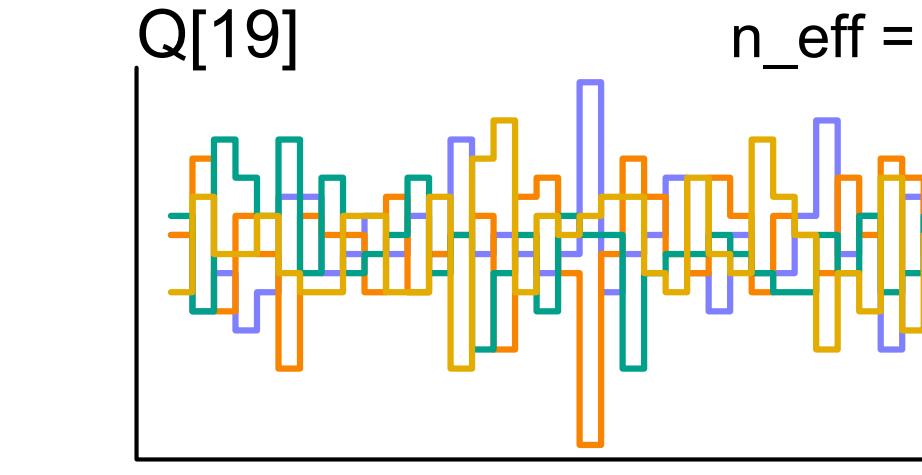
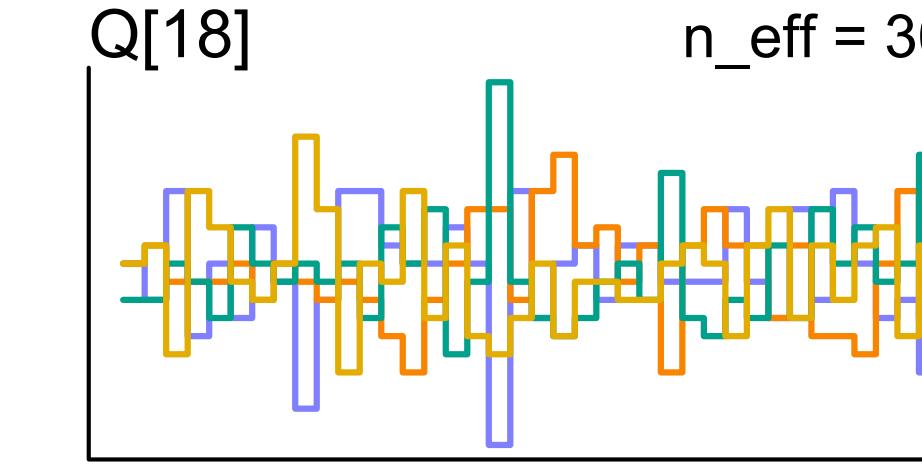
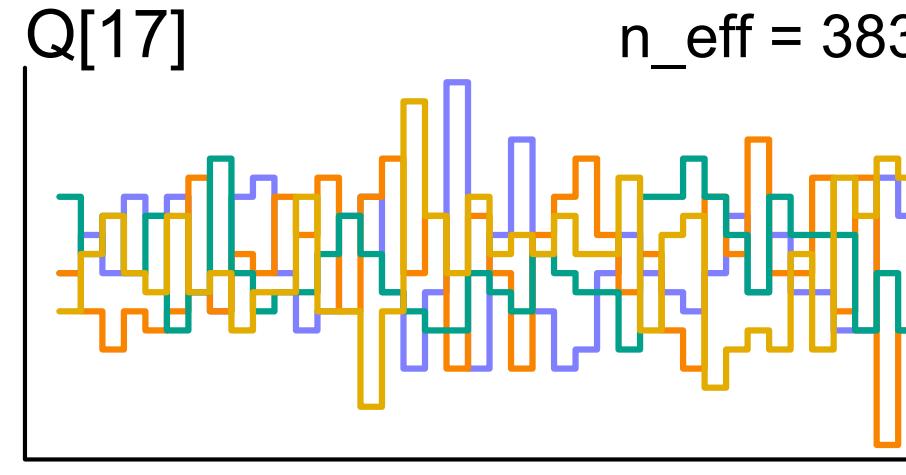
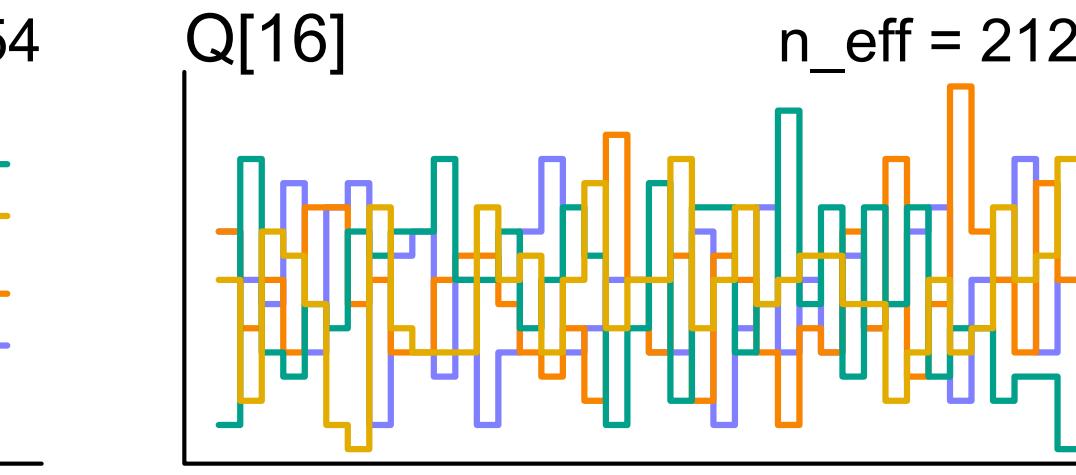
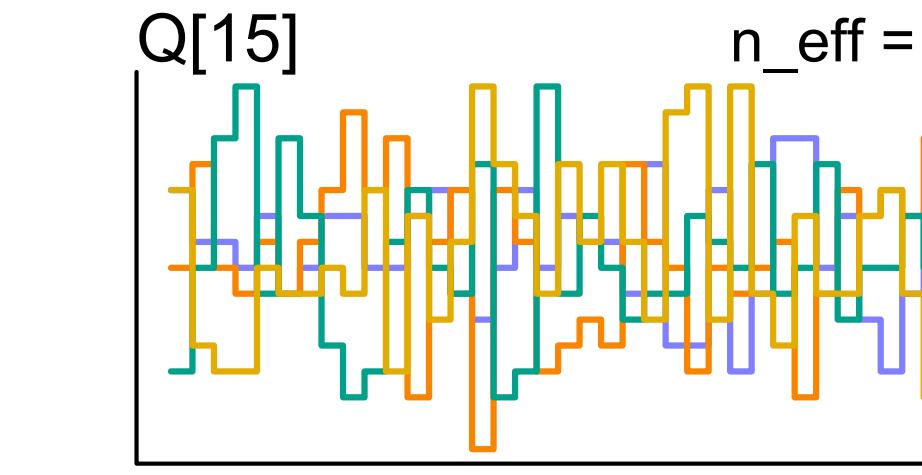
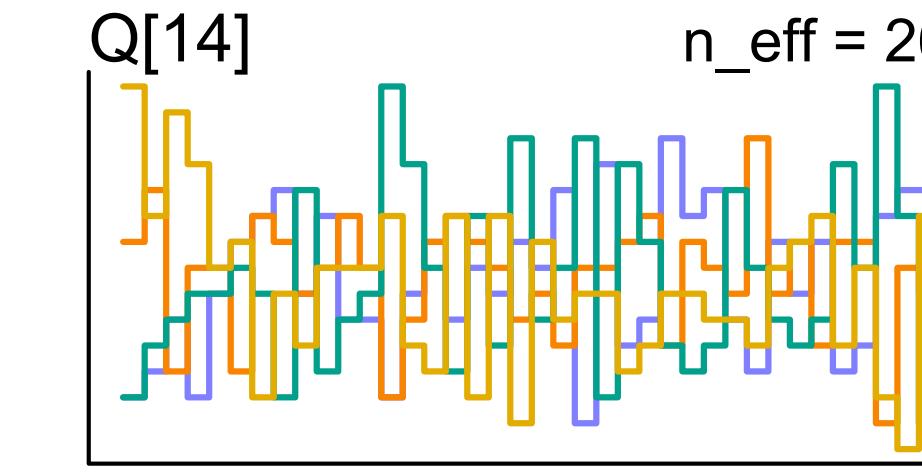
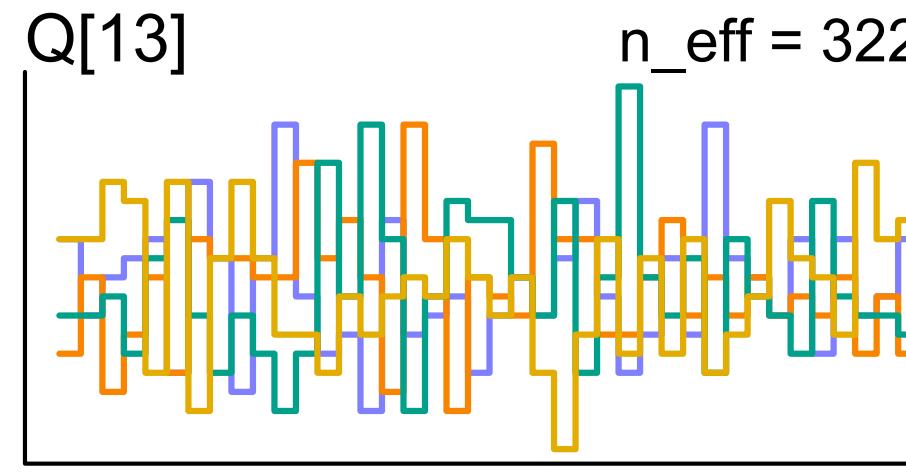
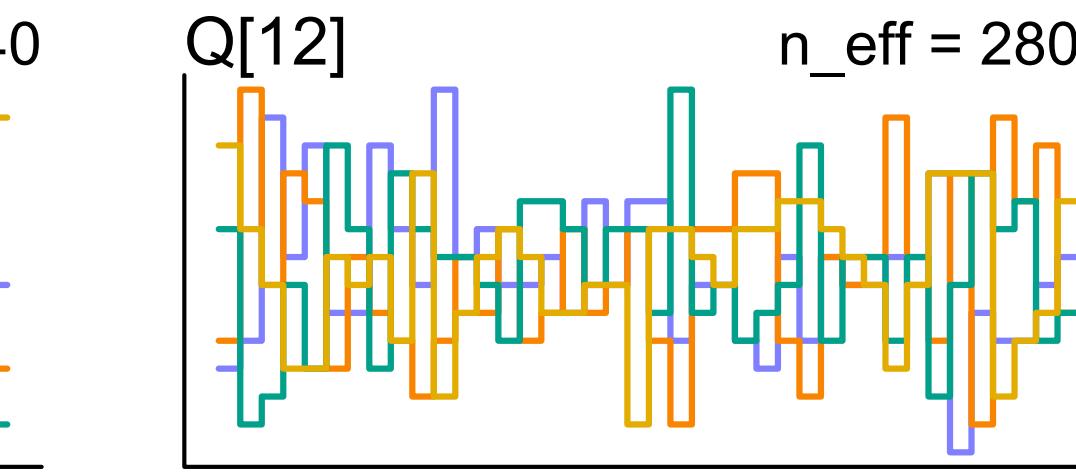
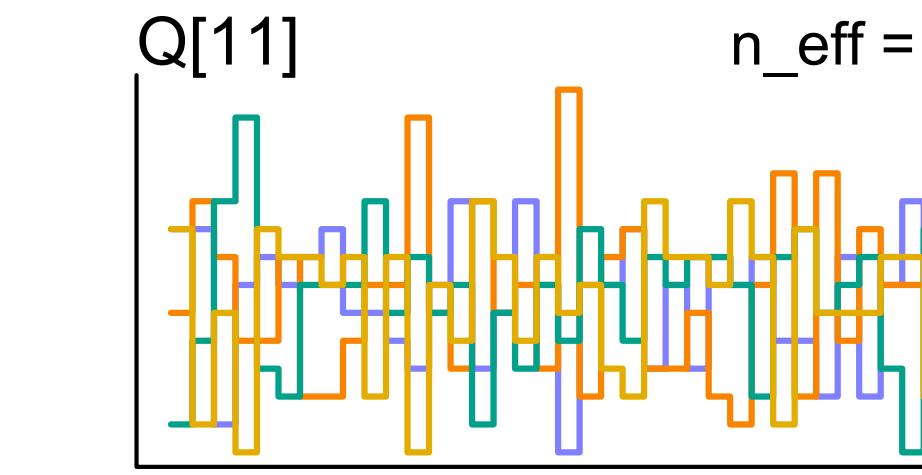
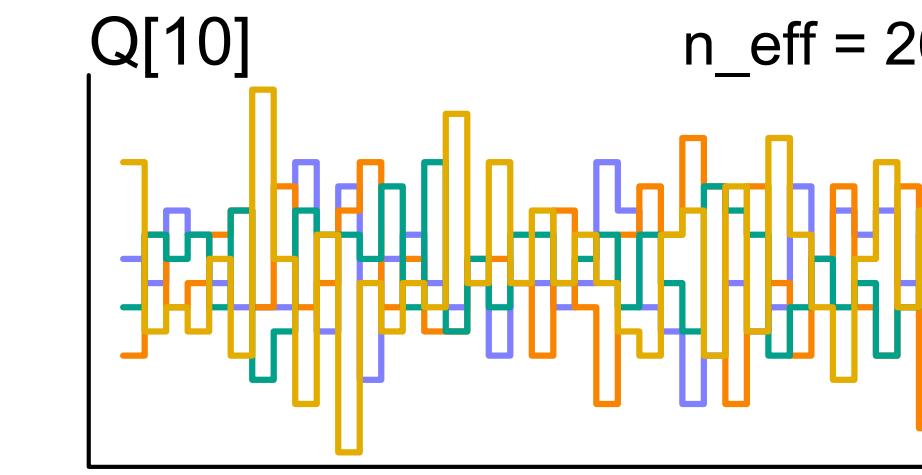
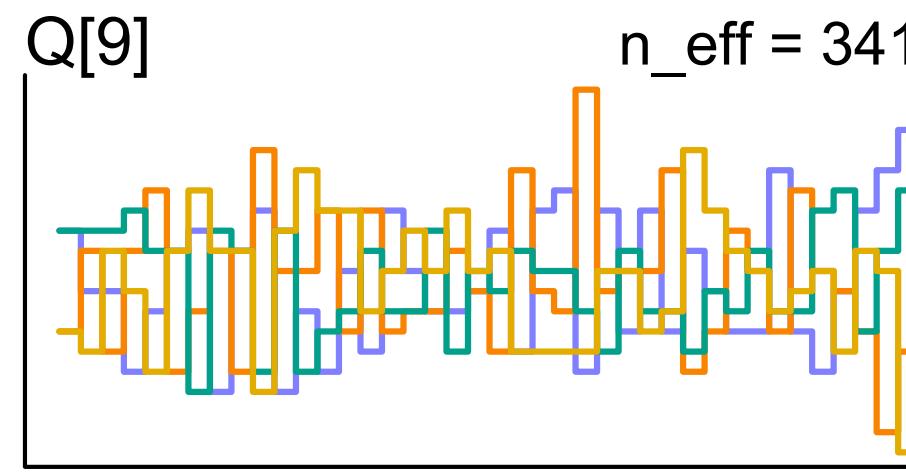
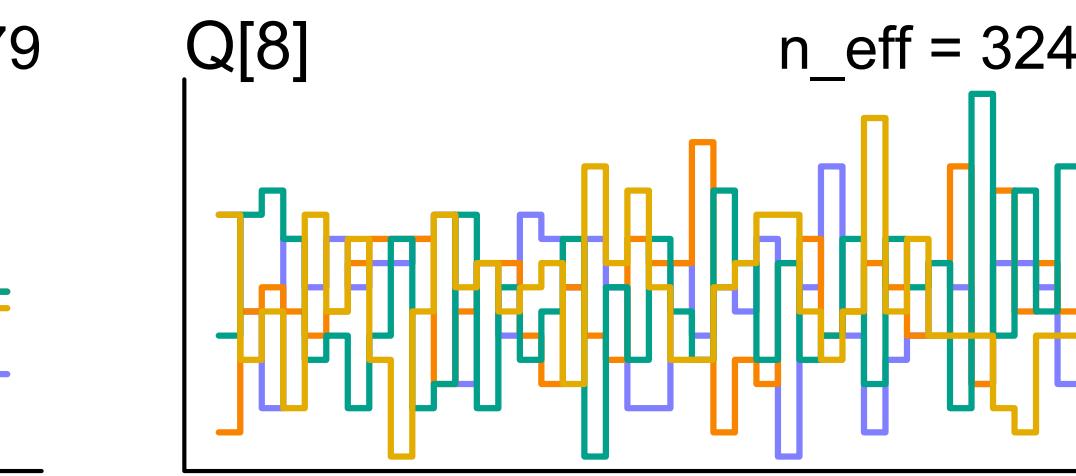
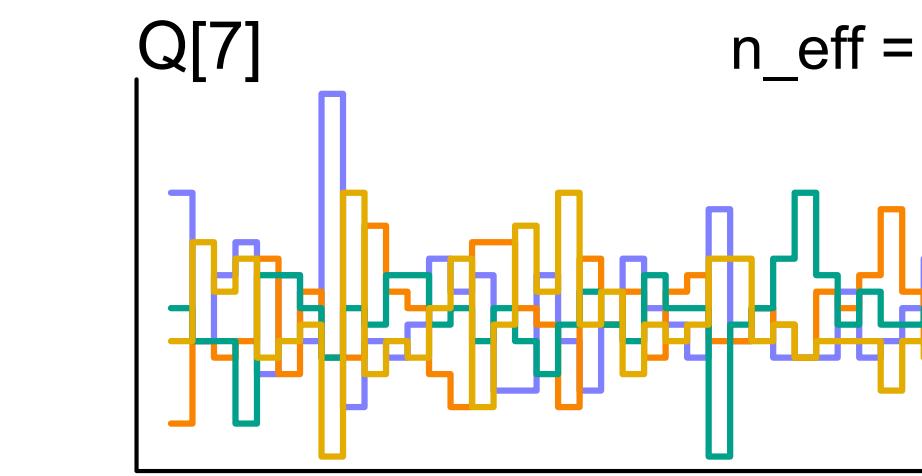
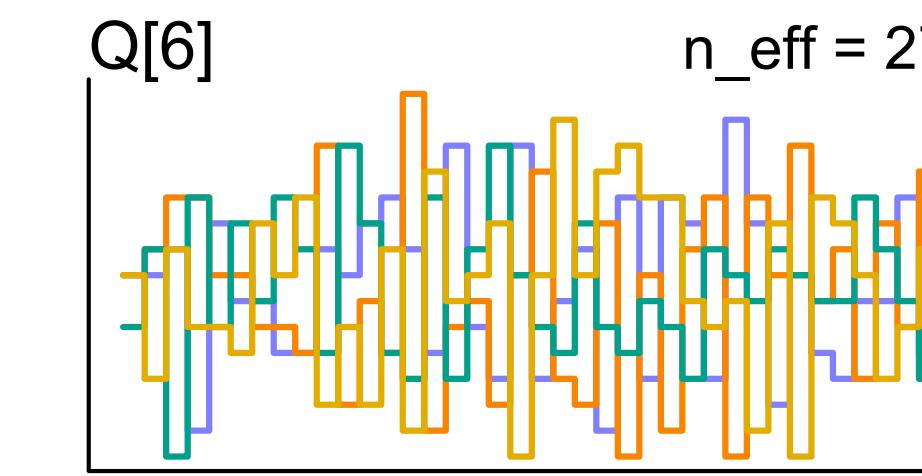
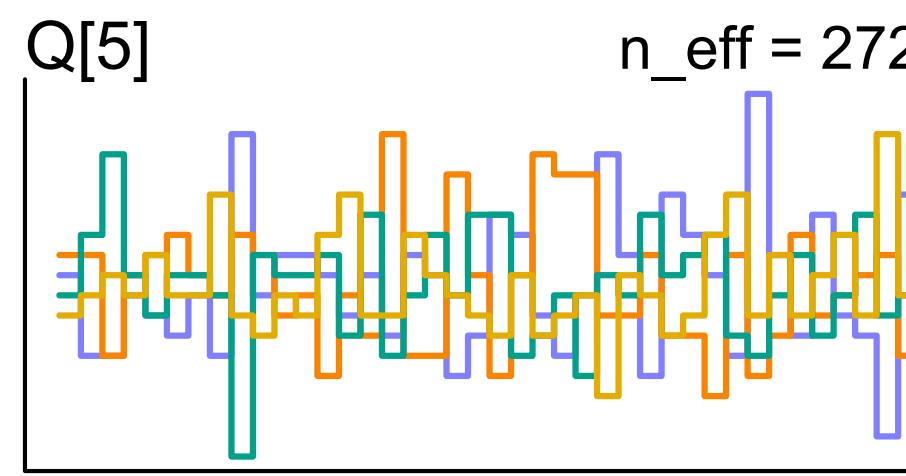
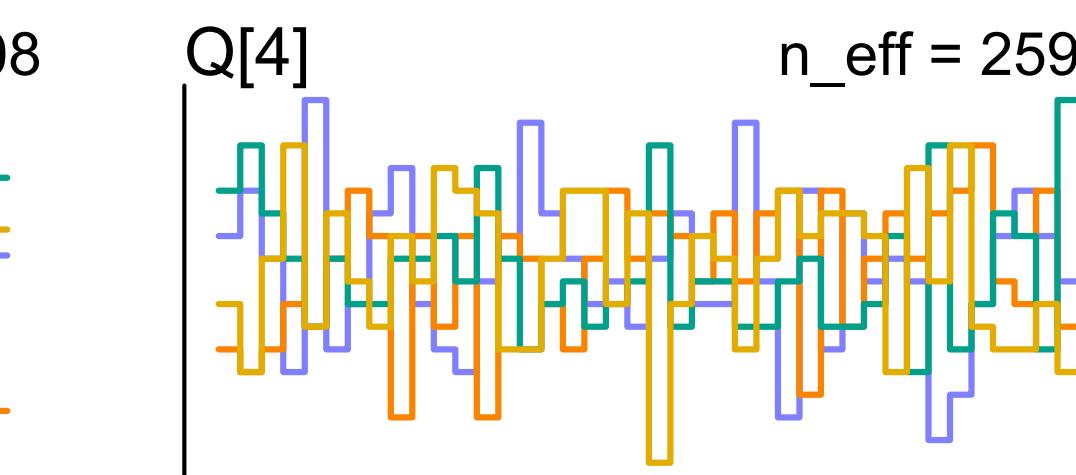
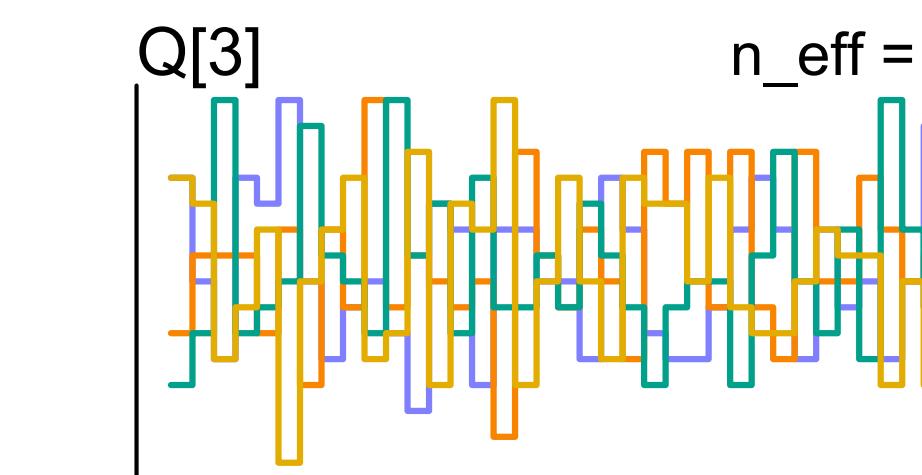
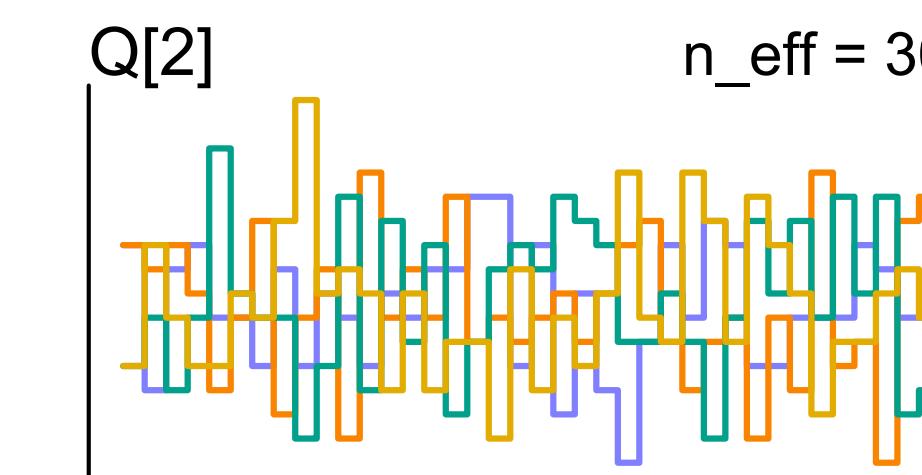
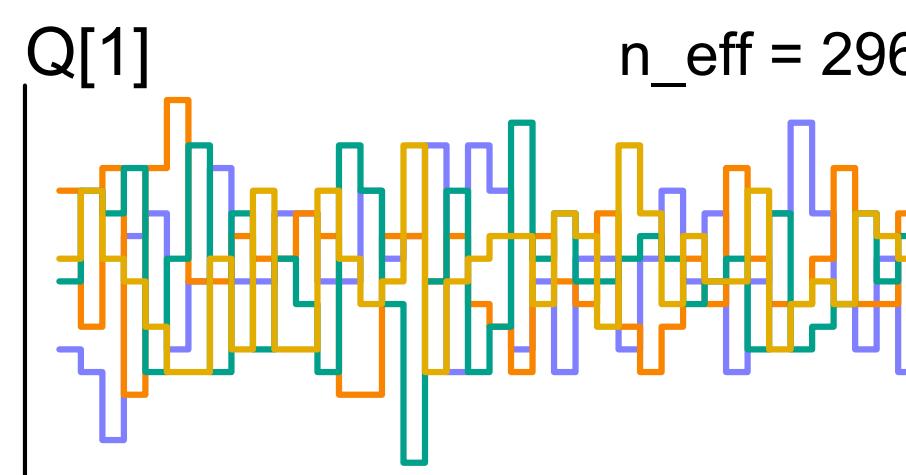


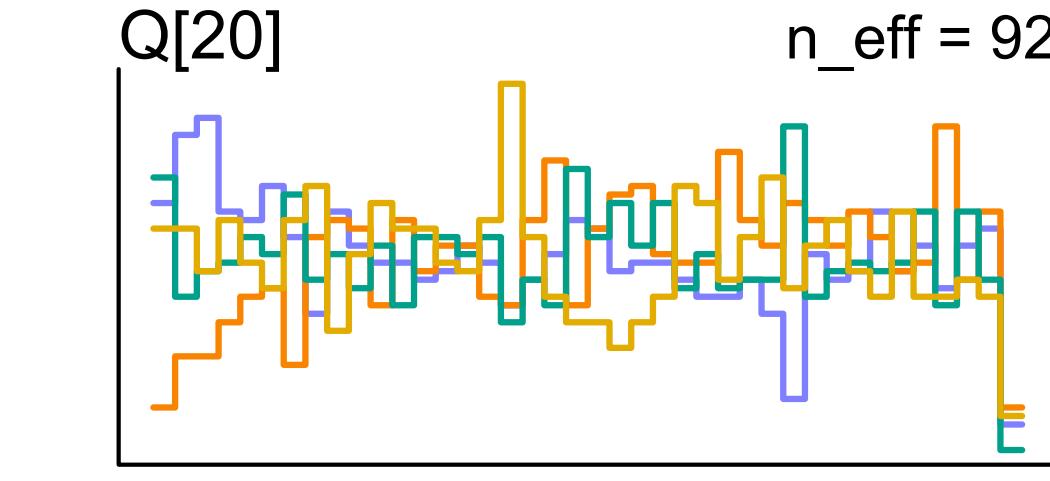
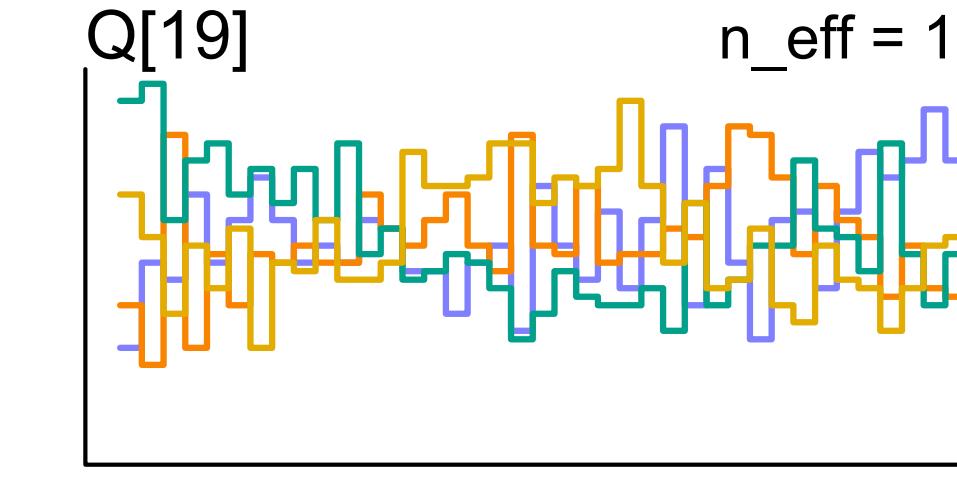
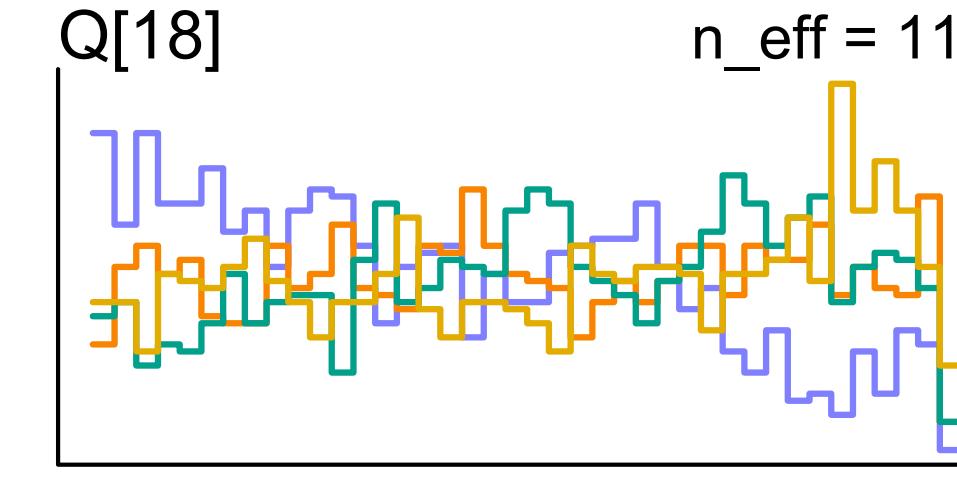
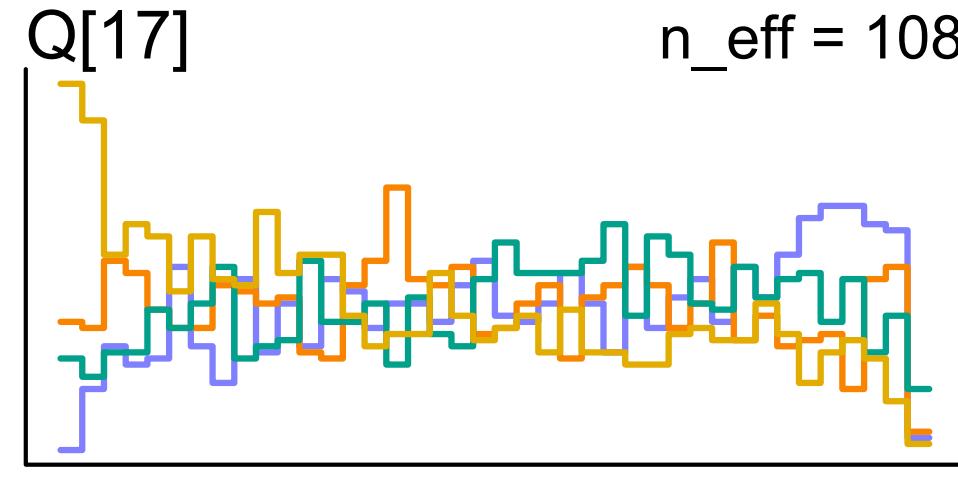
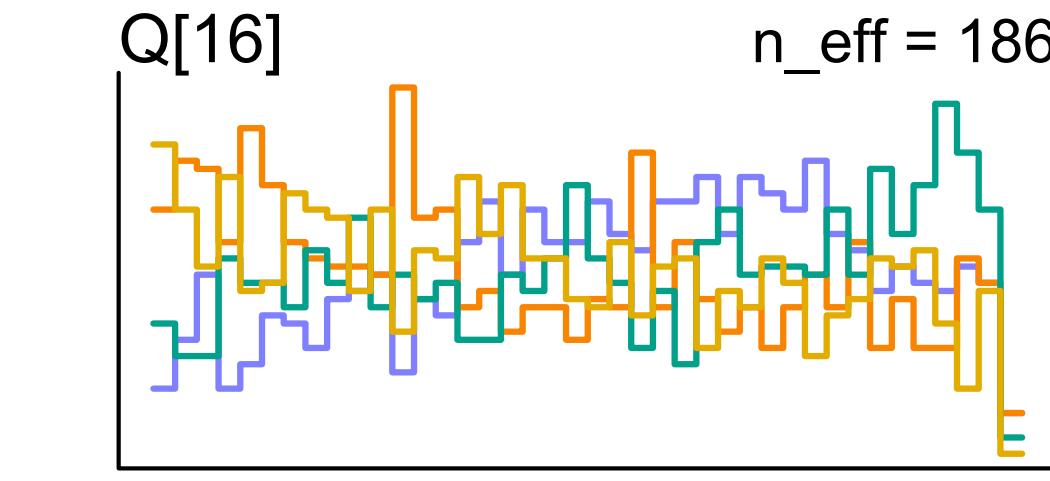
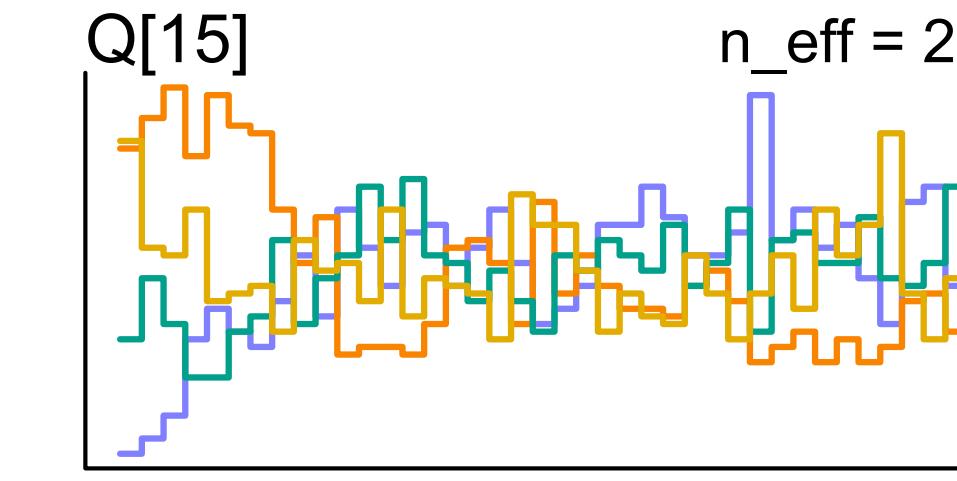
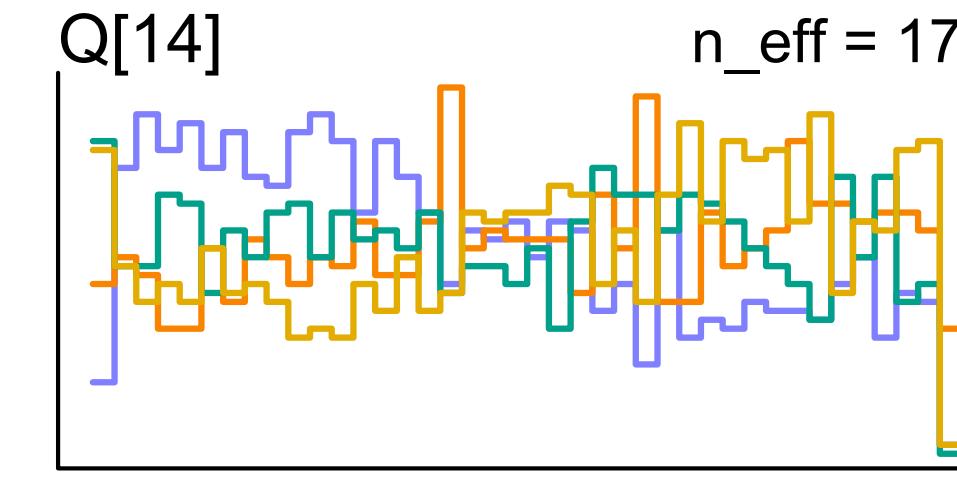
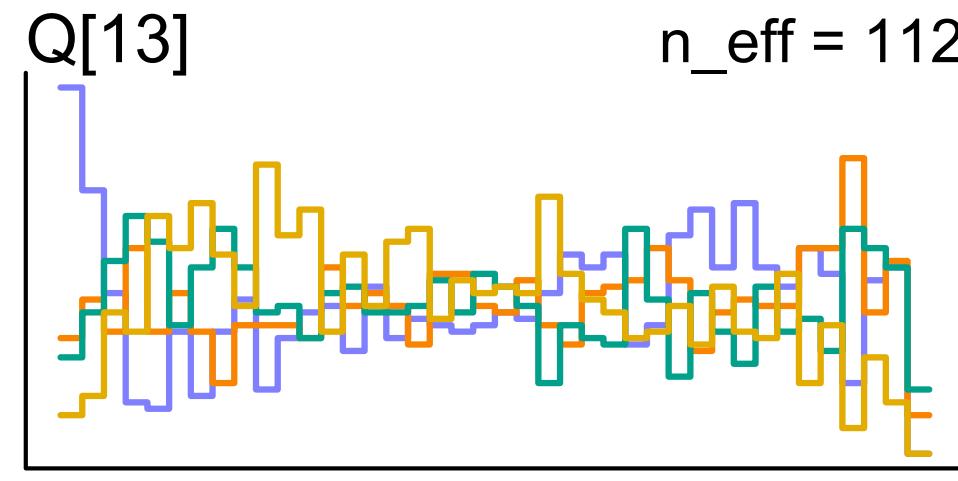
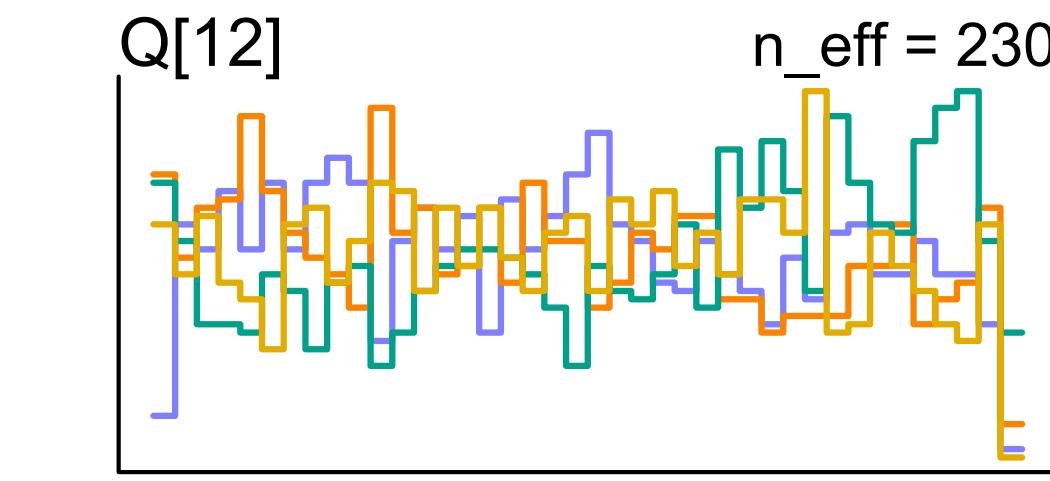
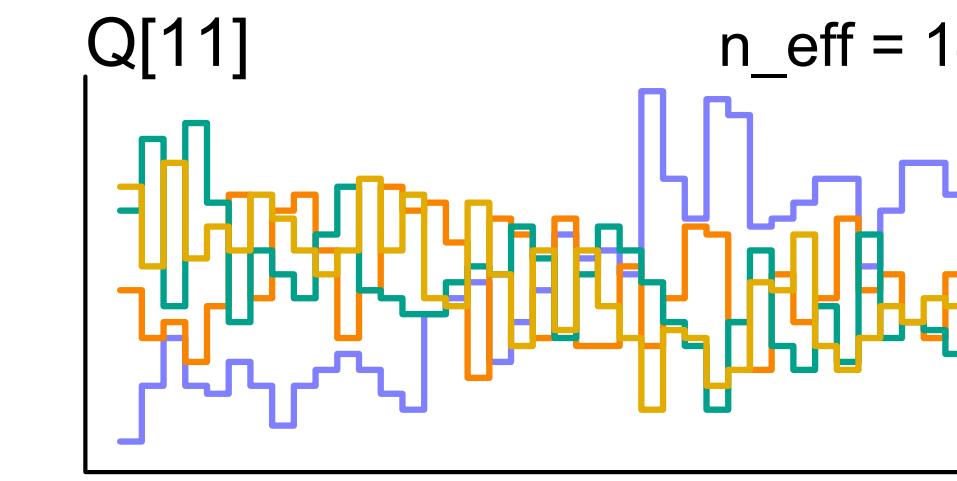
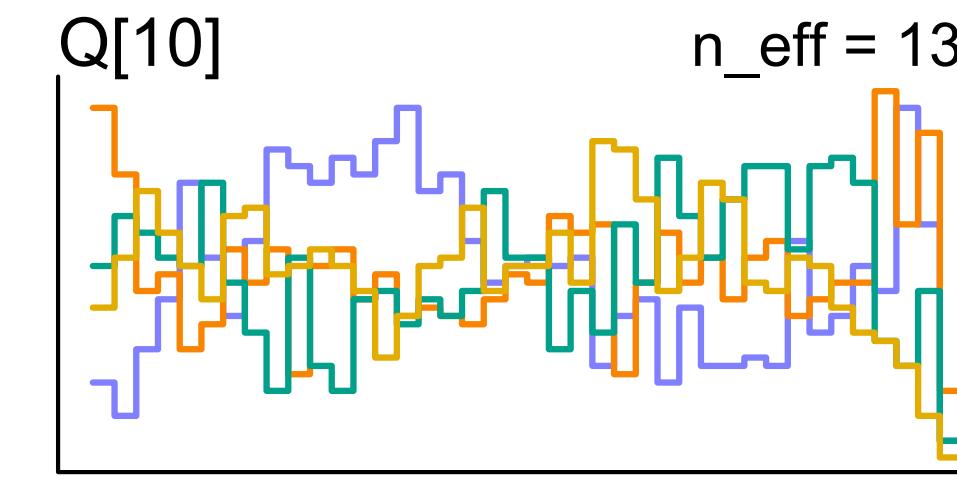
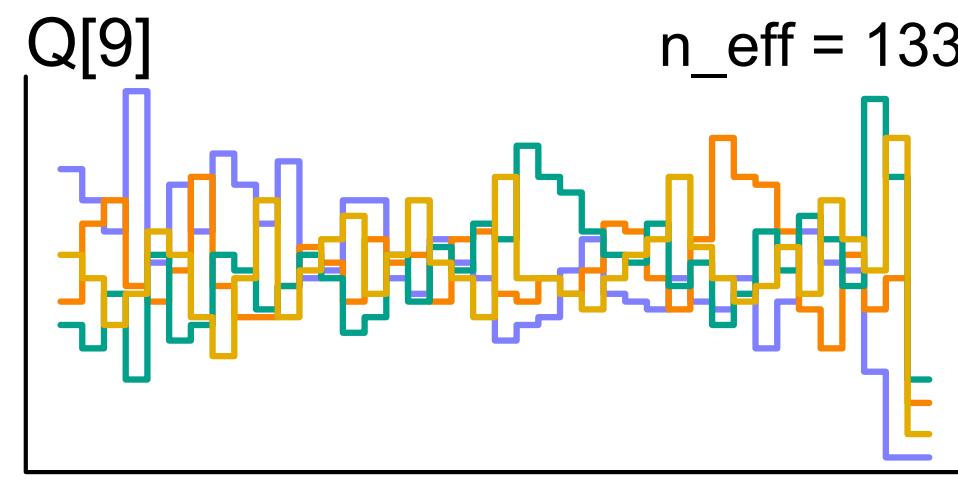
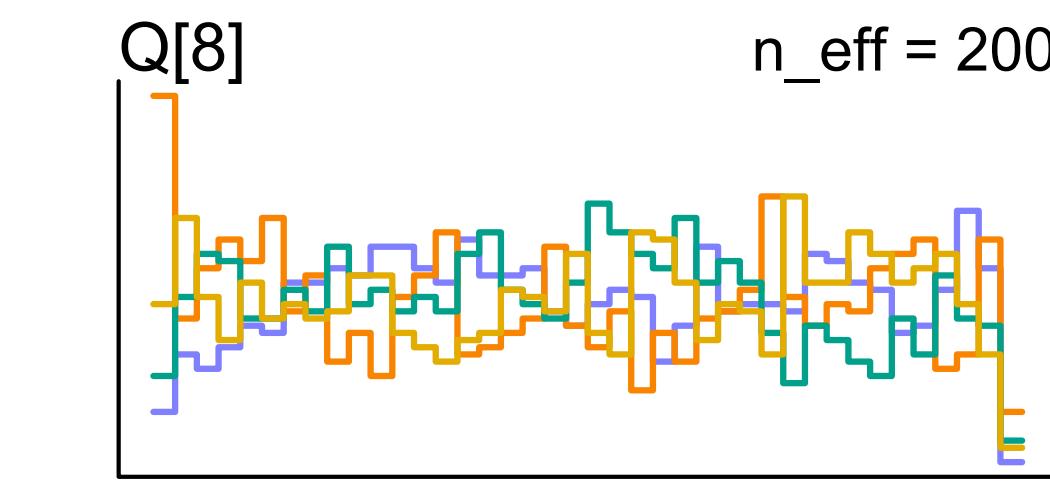
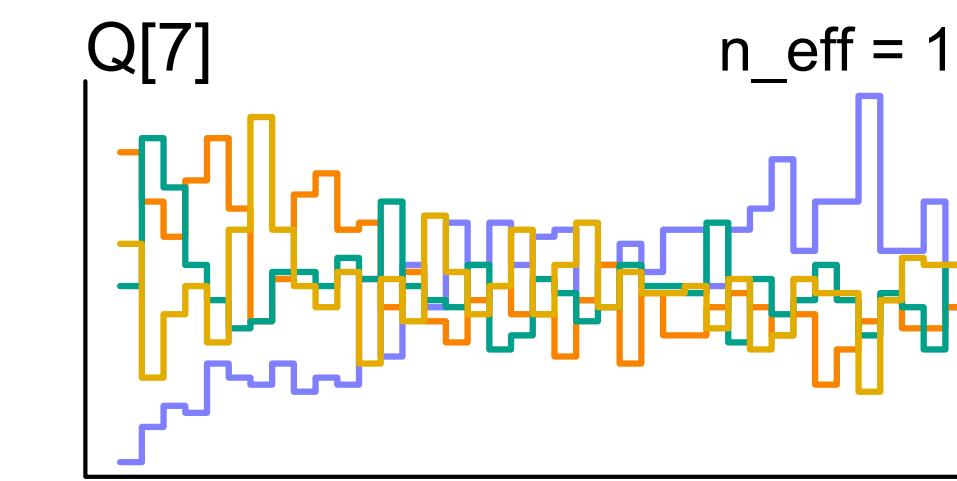
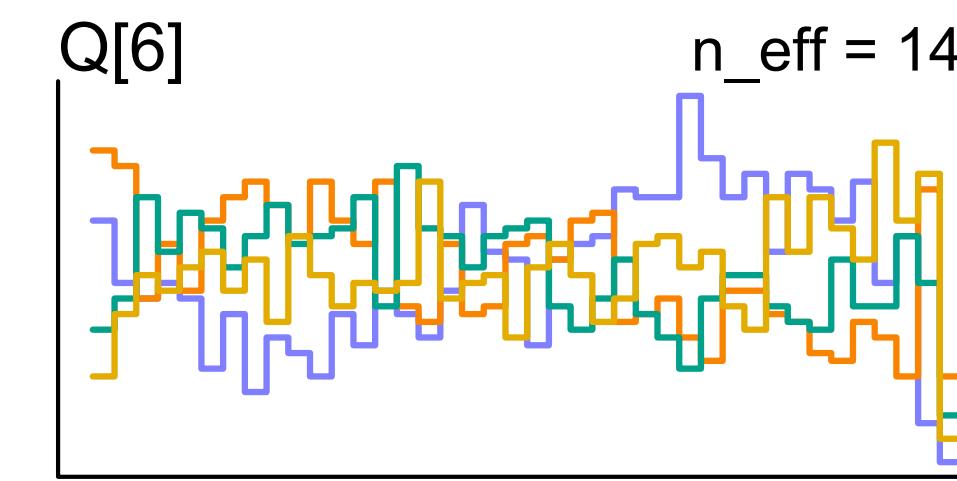
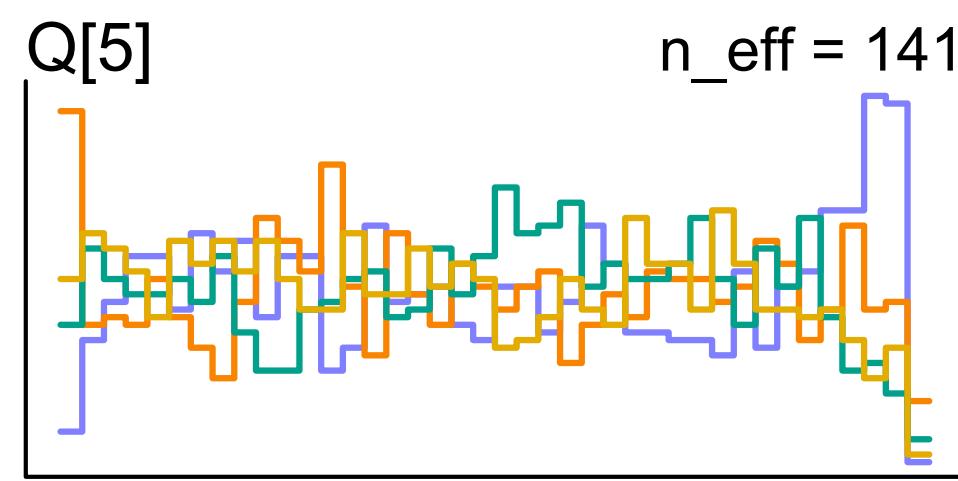
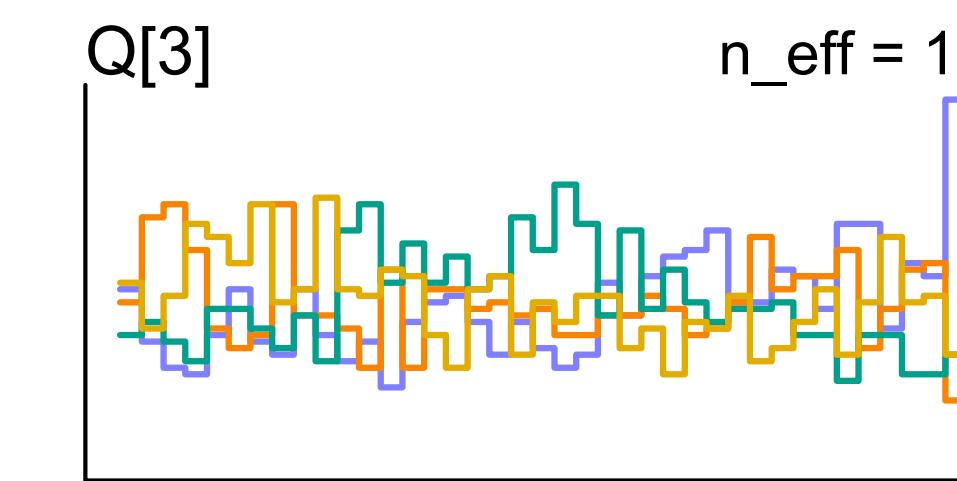
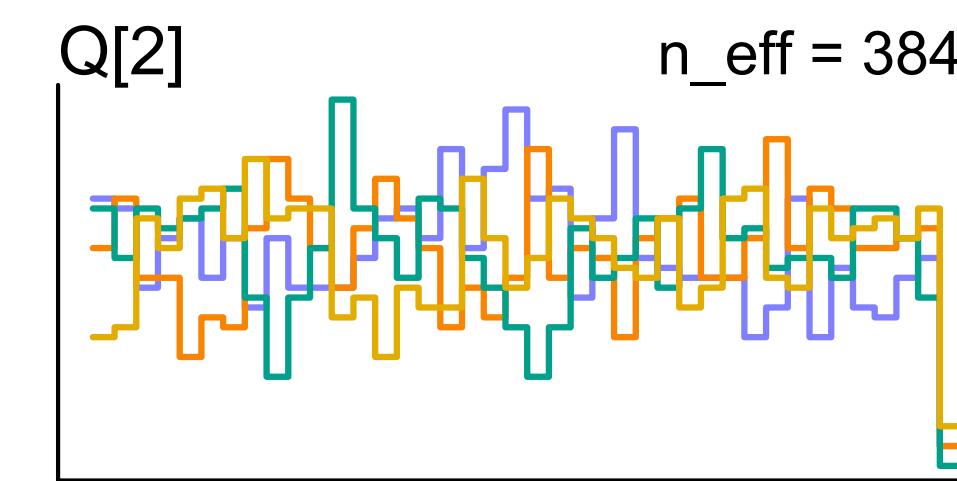
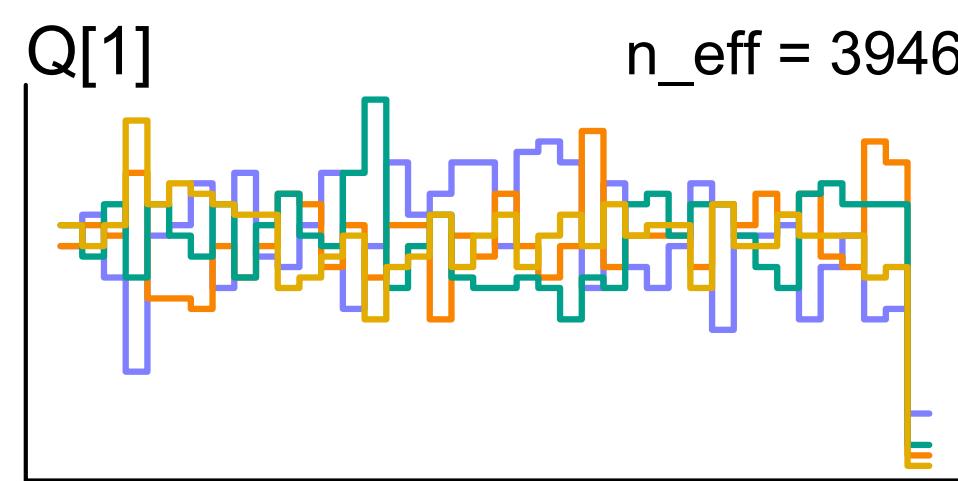
Trace rank (Trank) plots

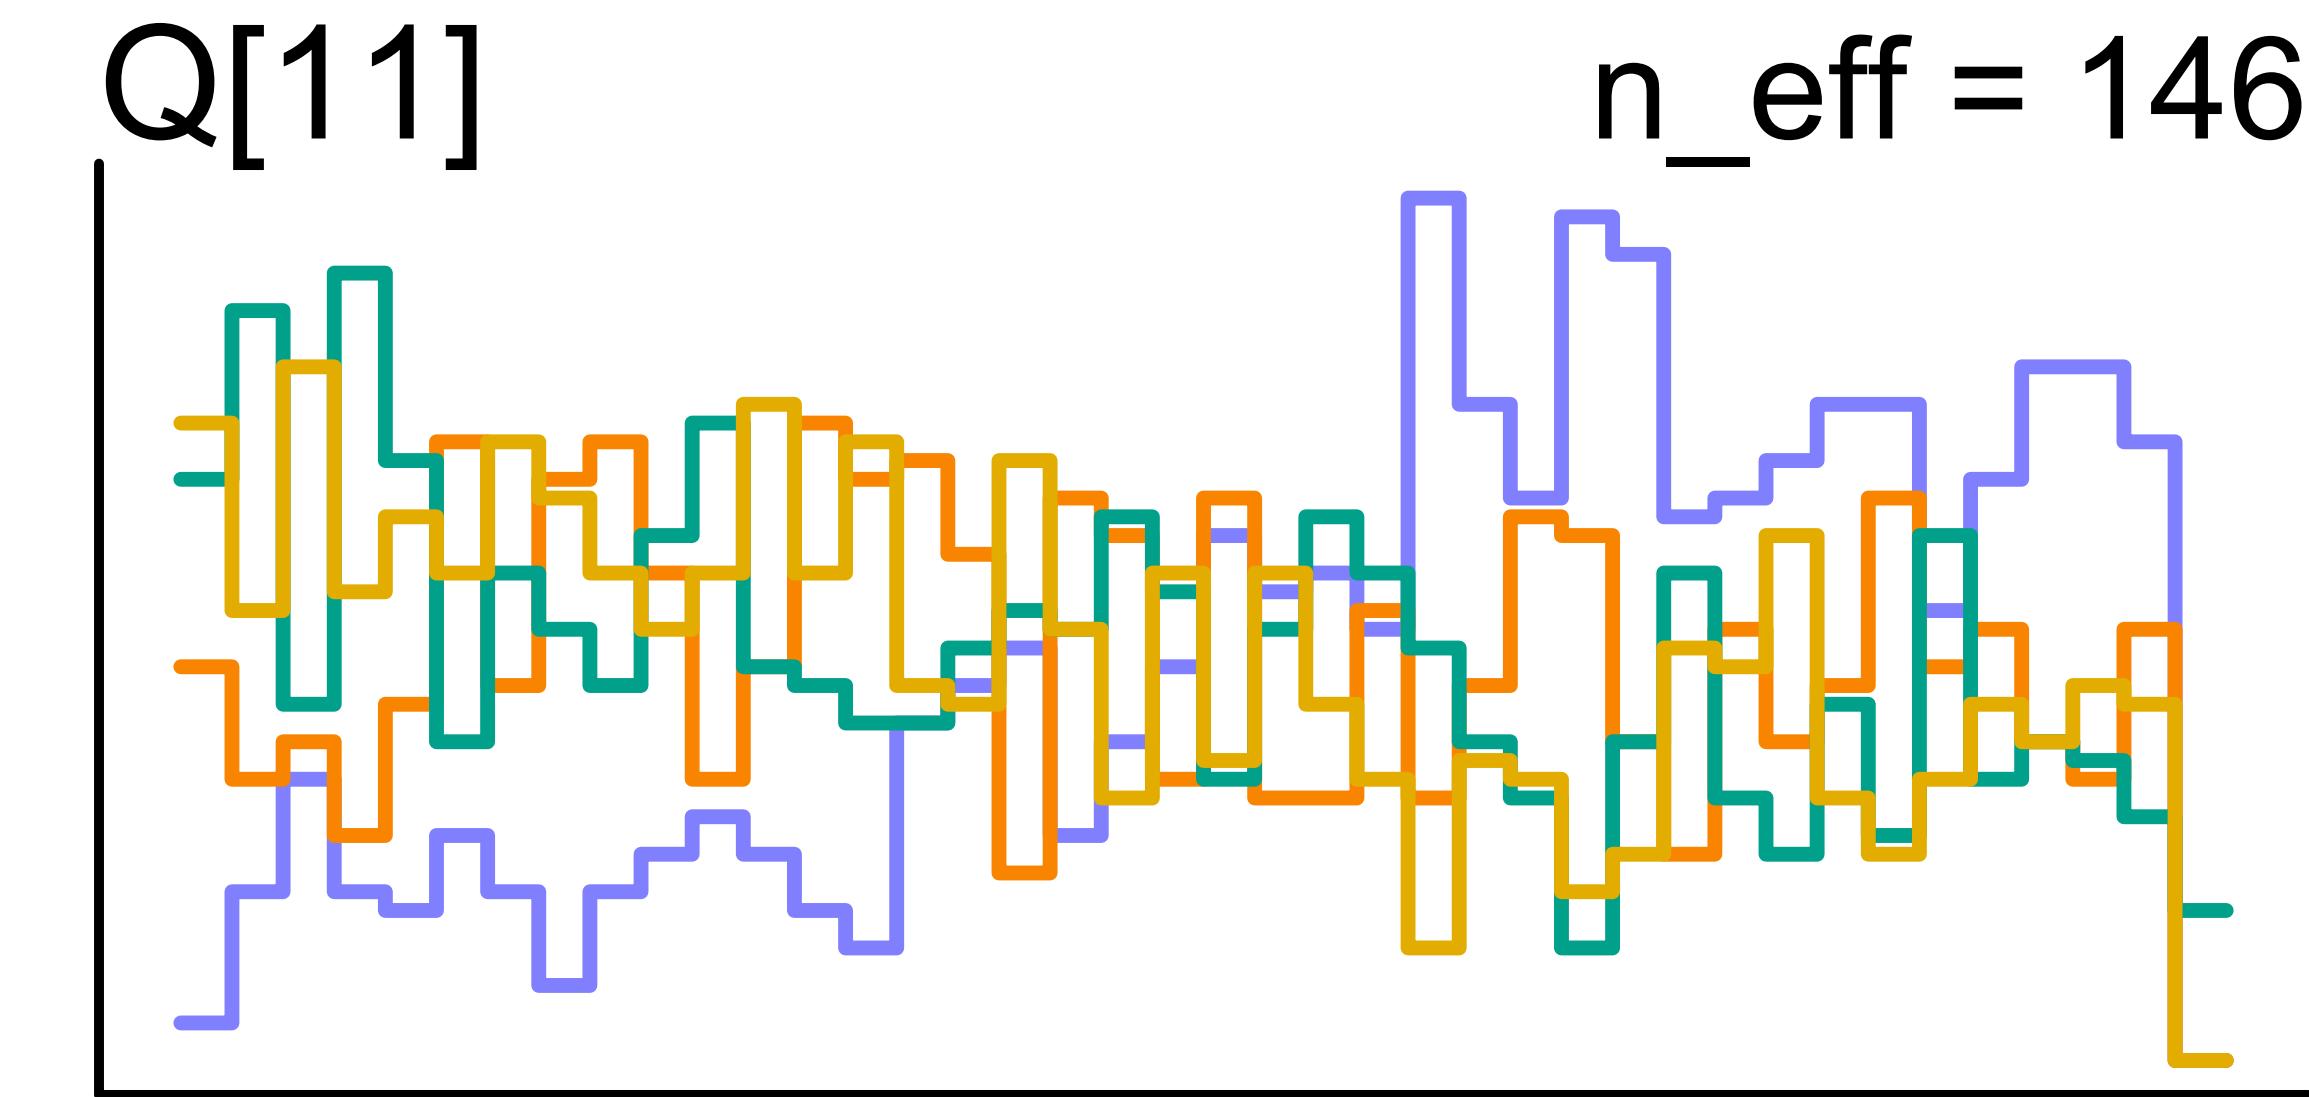
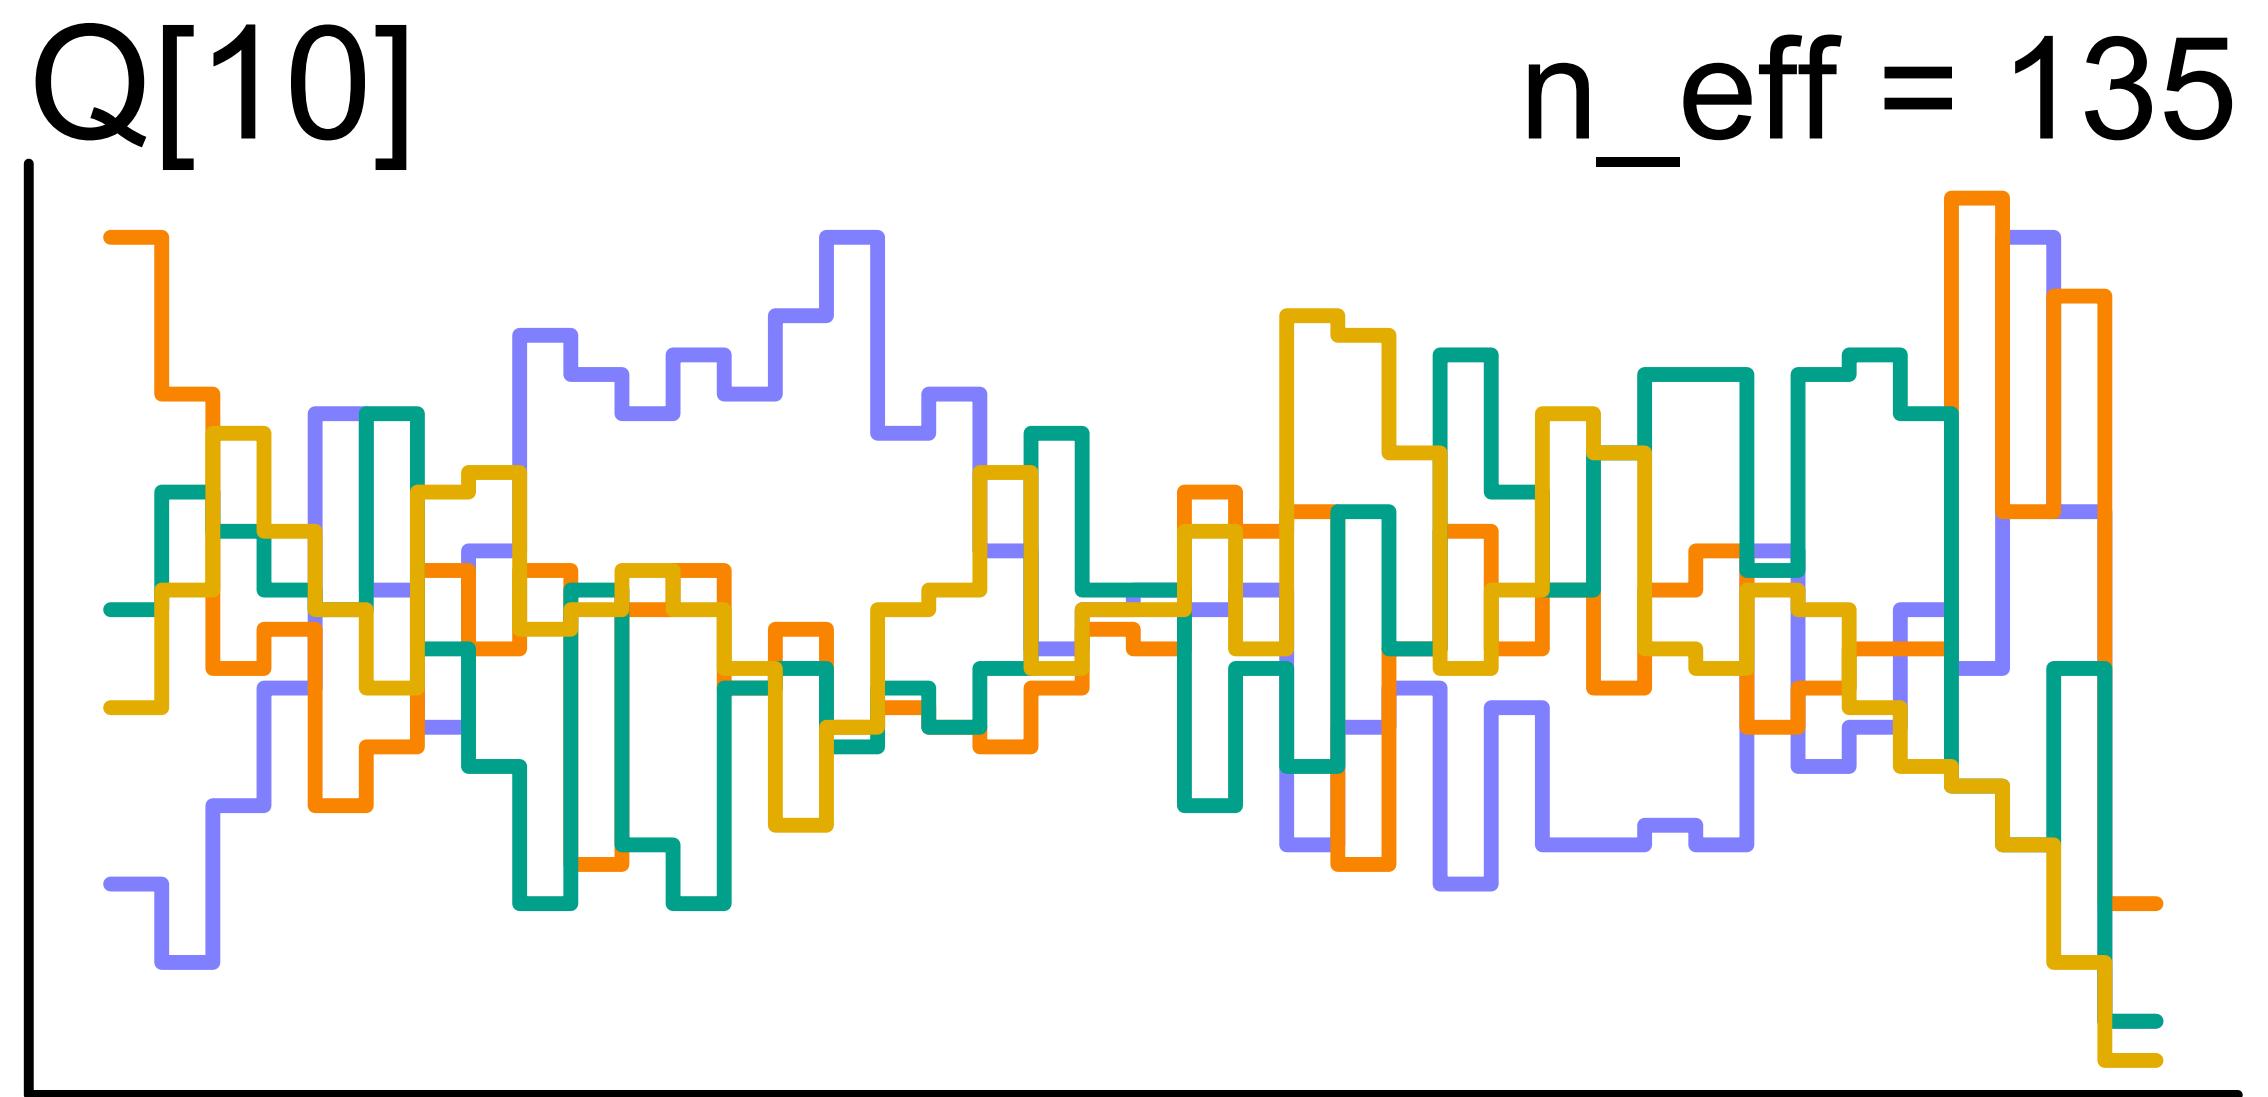
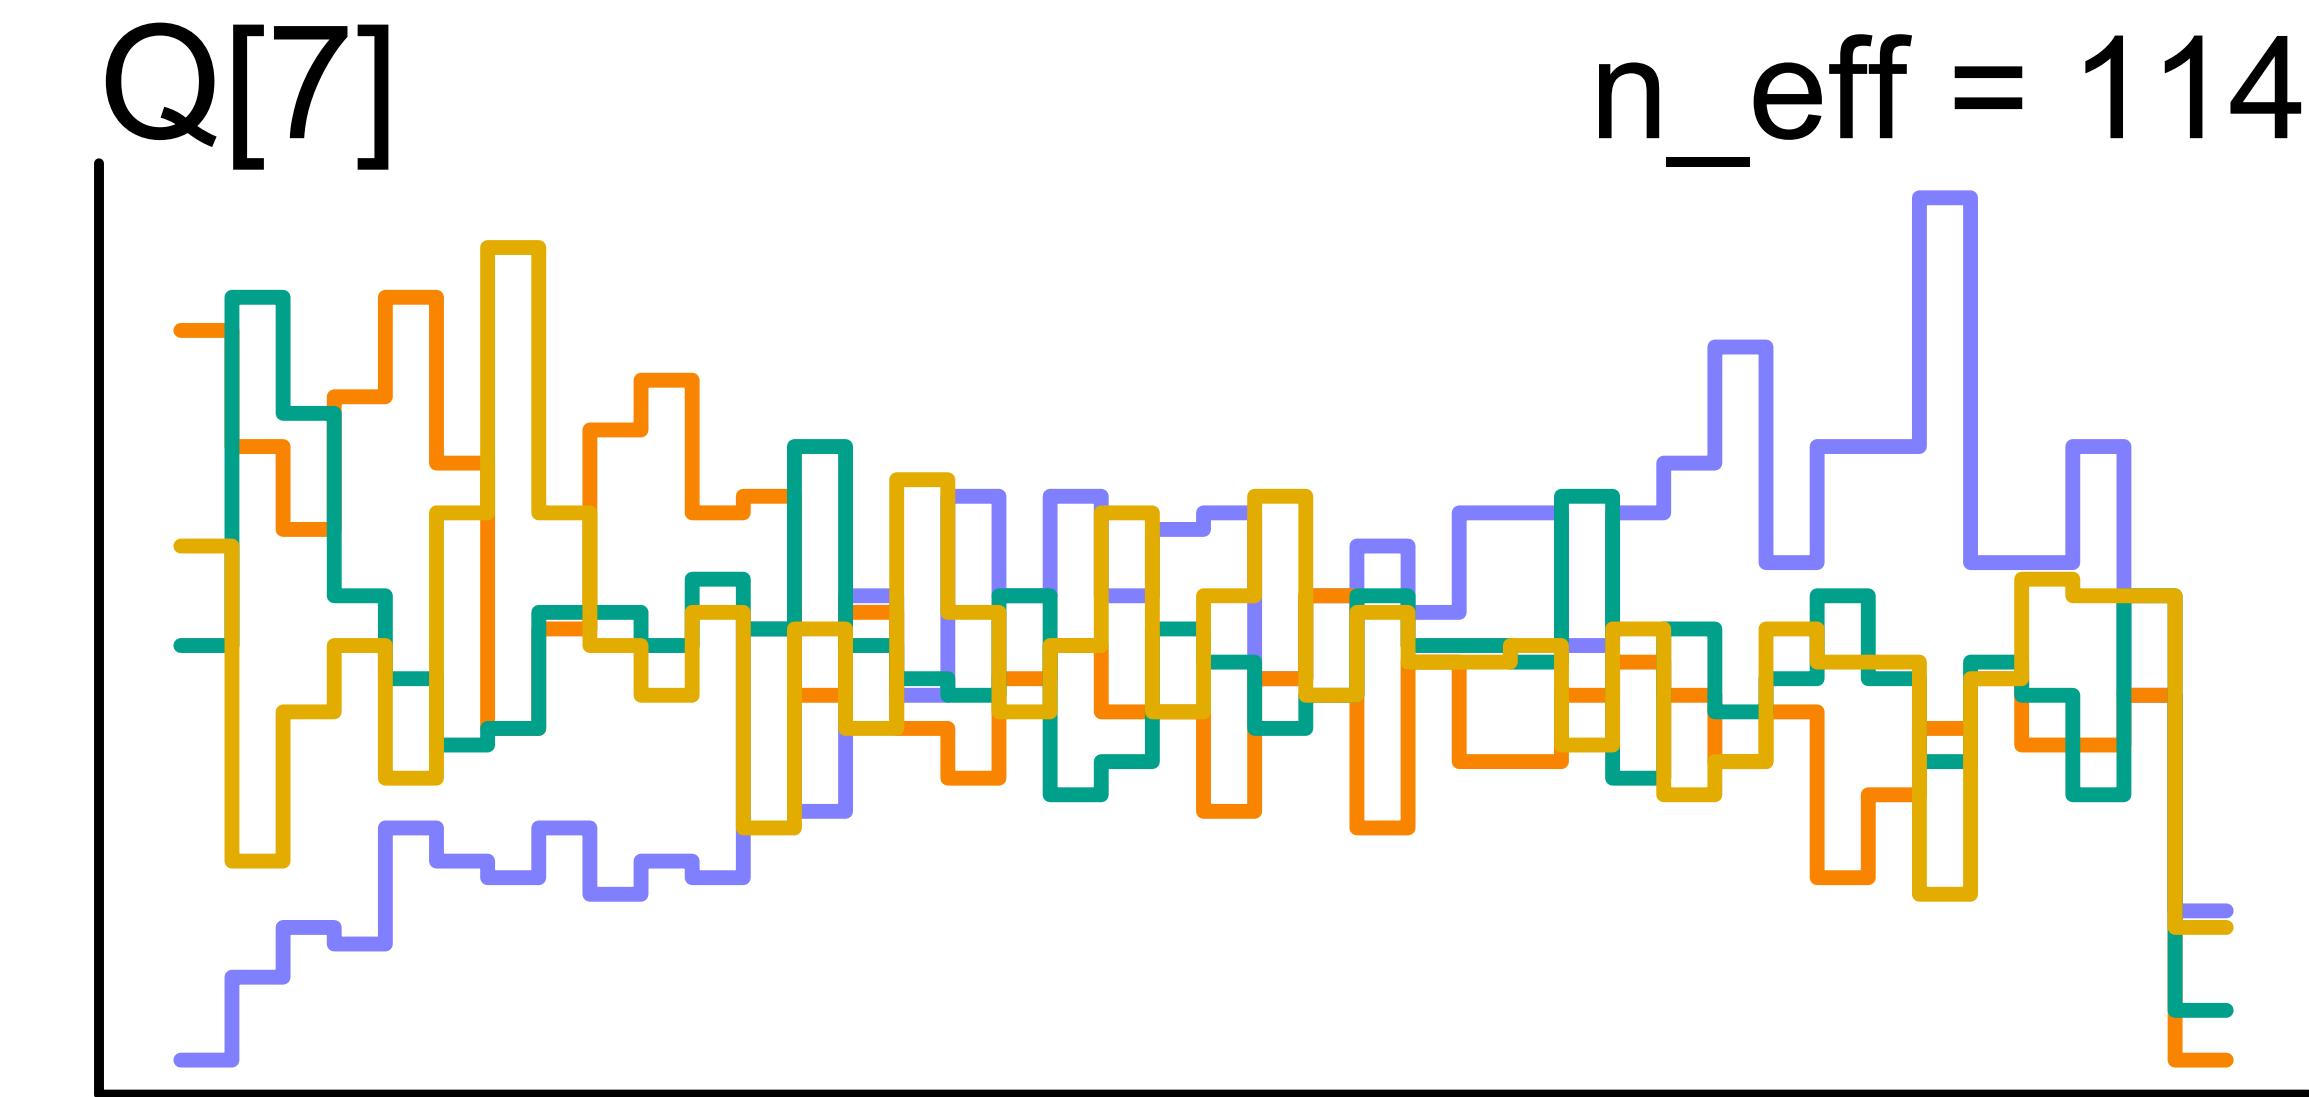
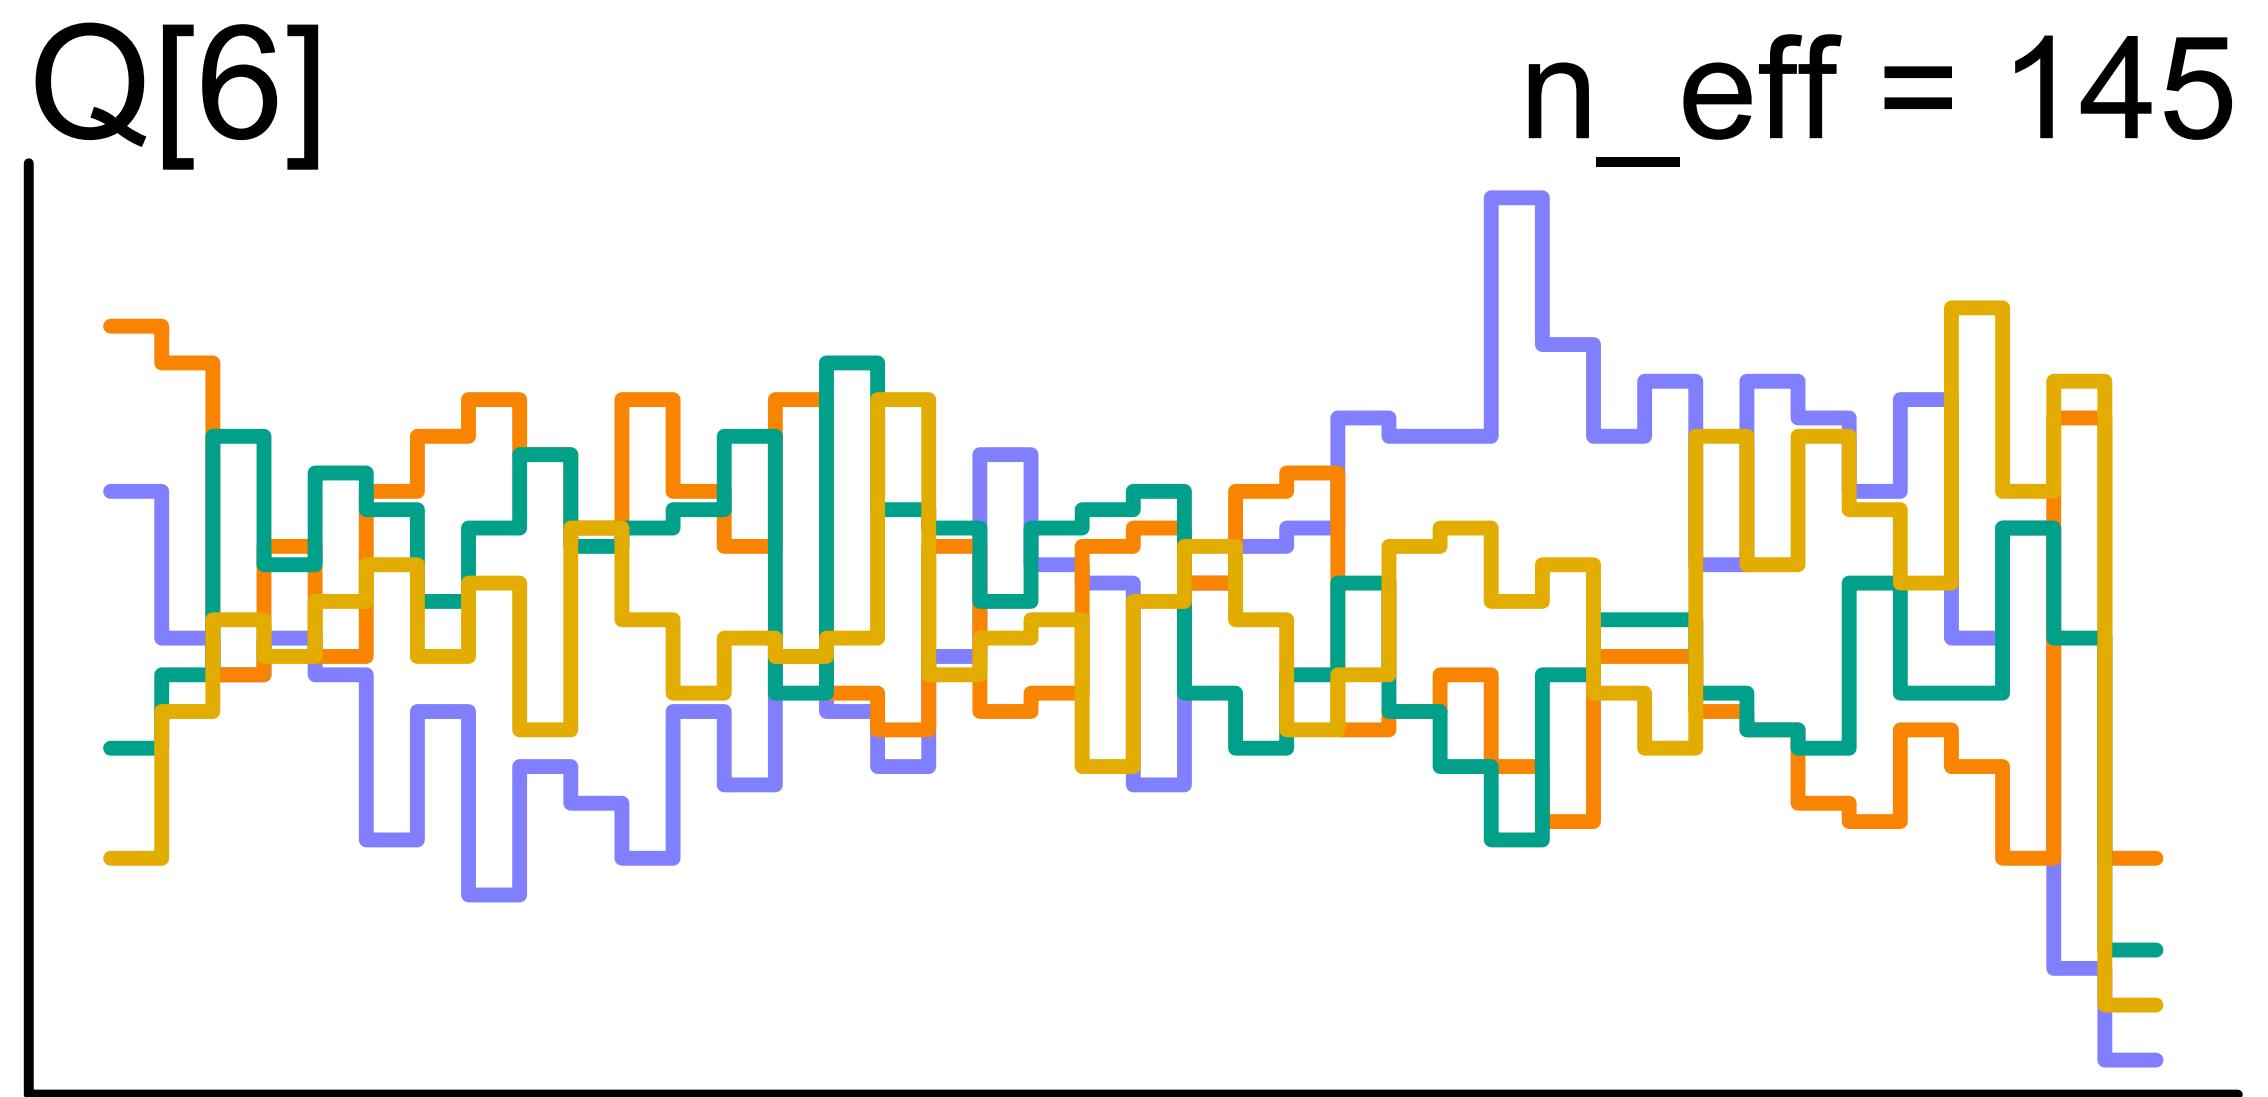


Rank orders of samples. No chain should tend to be below/above others.









R-hat

When chains converge:

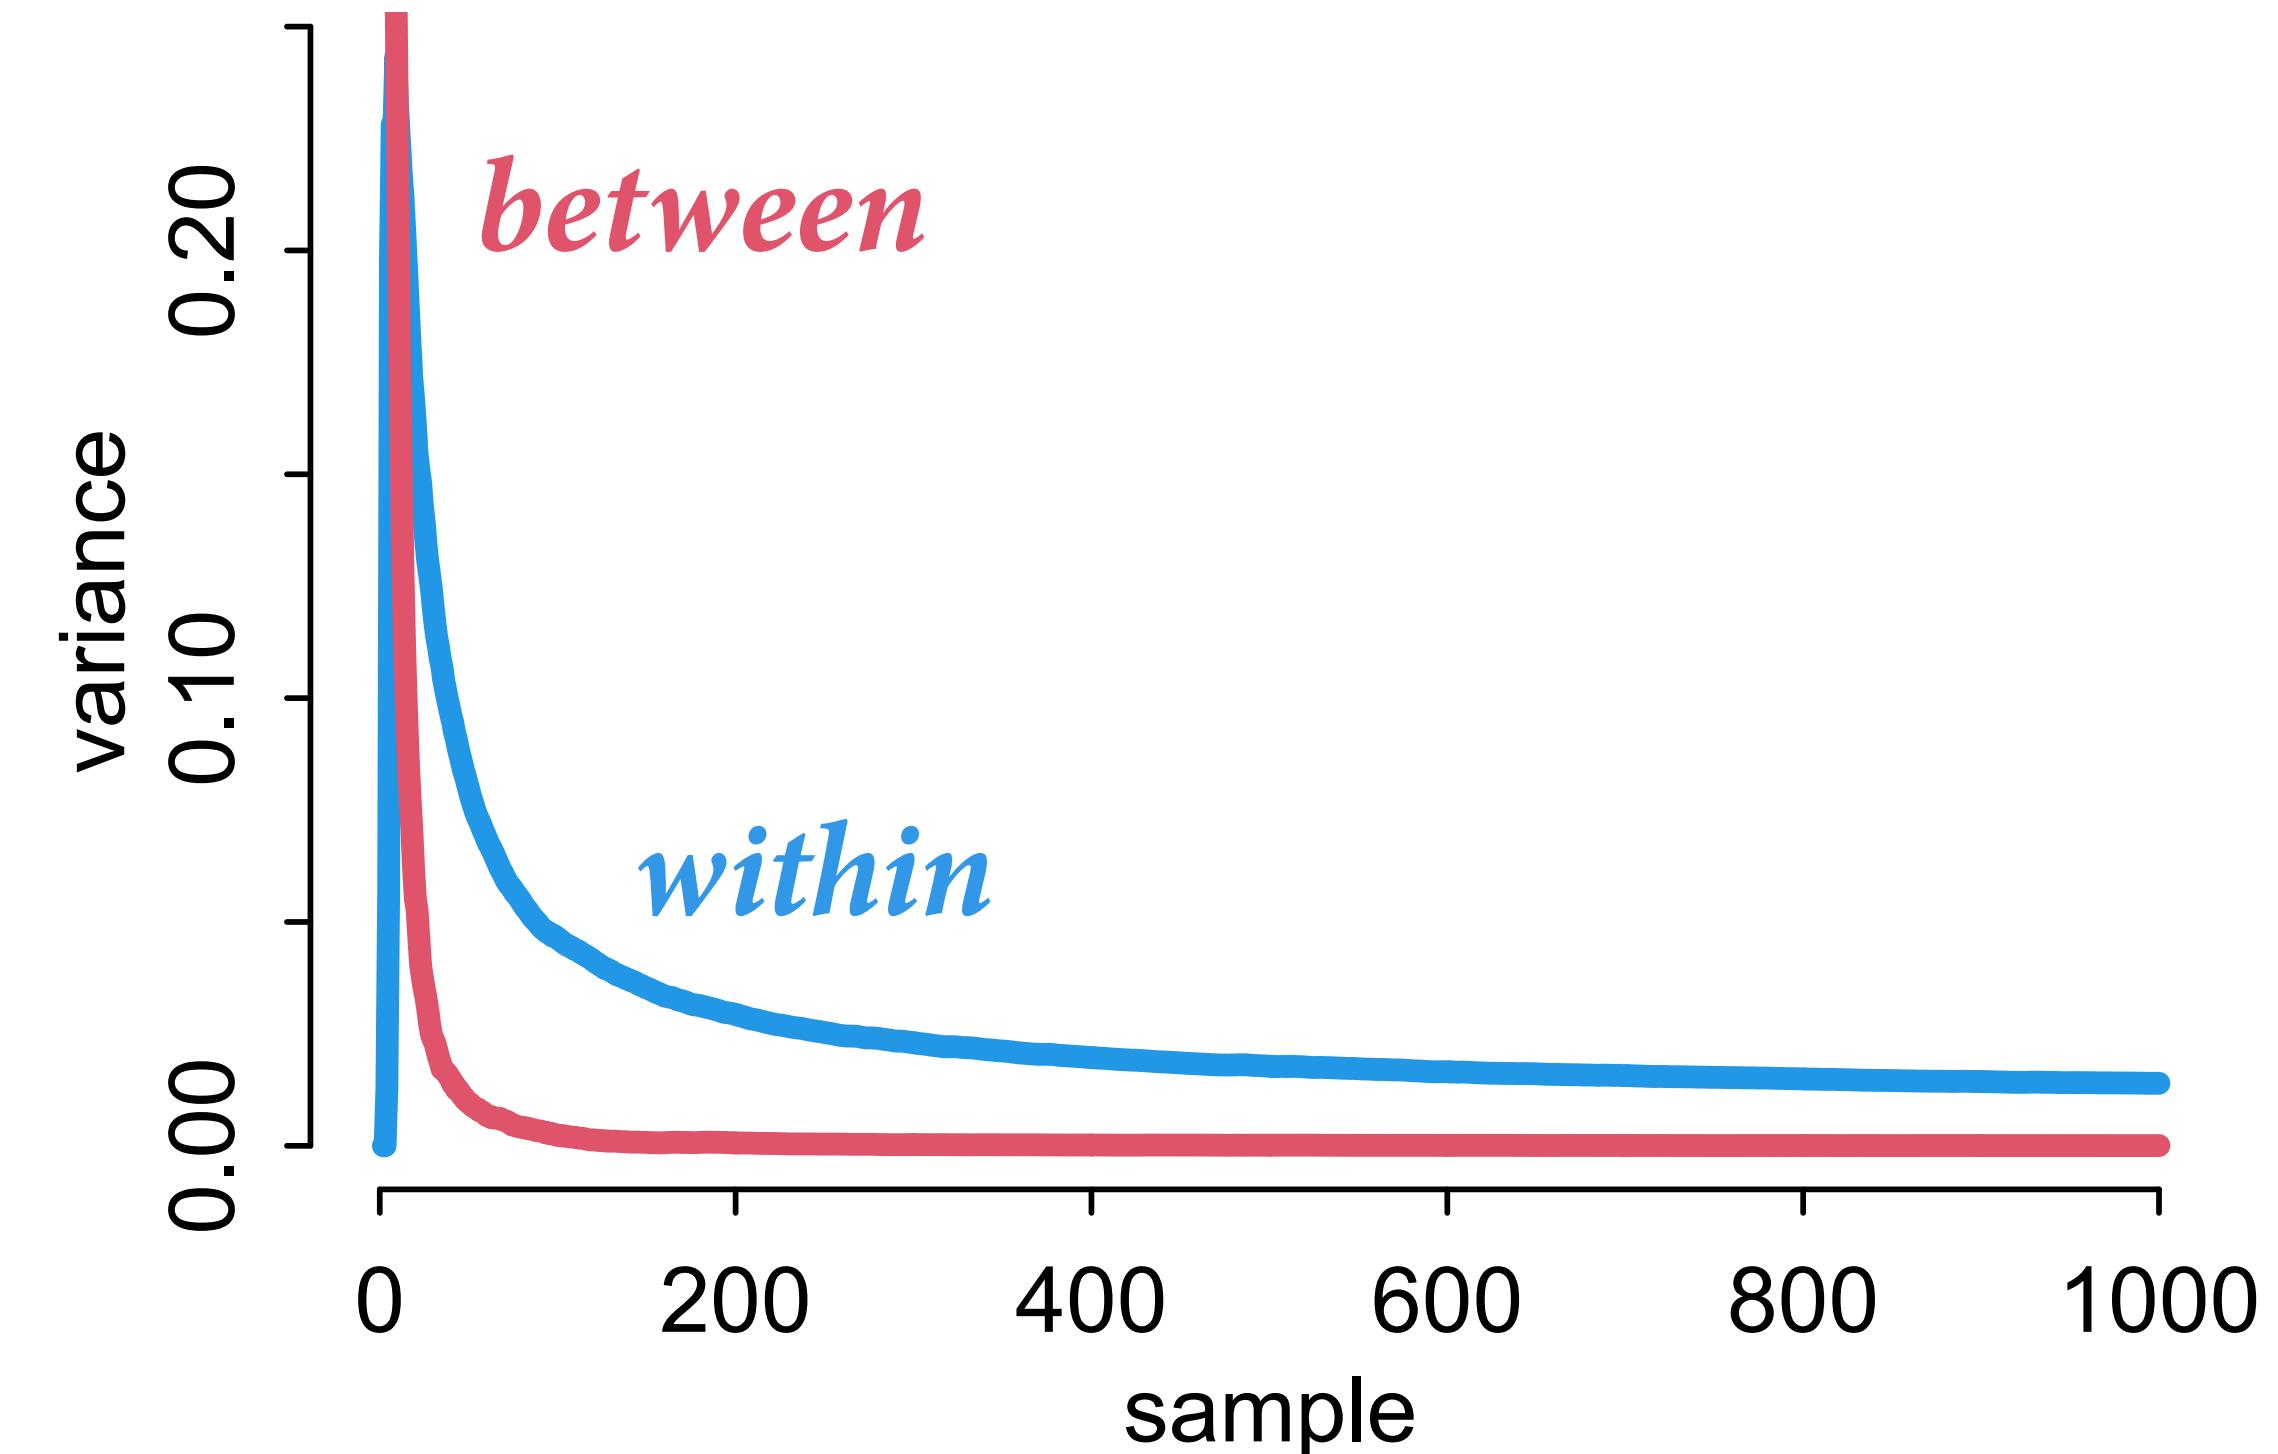
- (1) Start and end of each chain explores same region
- (2) Independent chains explore same region

R-hat is a ratio of variances:

As total variance shrinks to average variance within chains, R-hat approaches 1

NO GUARANTEES; NOT A TEST

	mean	sd	5.5%	94.5%	n_eff	Rhat4
Q[1]	0.14	0.30	-0.34	0.64	2962	1
Q[2]	0.11	0.32	-0.40	0.61	3033	1
Q[3]	0.27	0.31	-0.21	0.75	2608	1



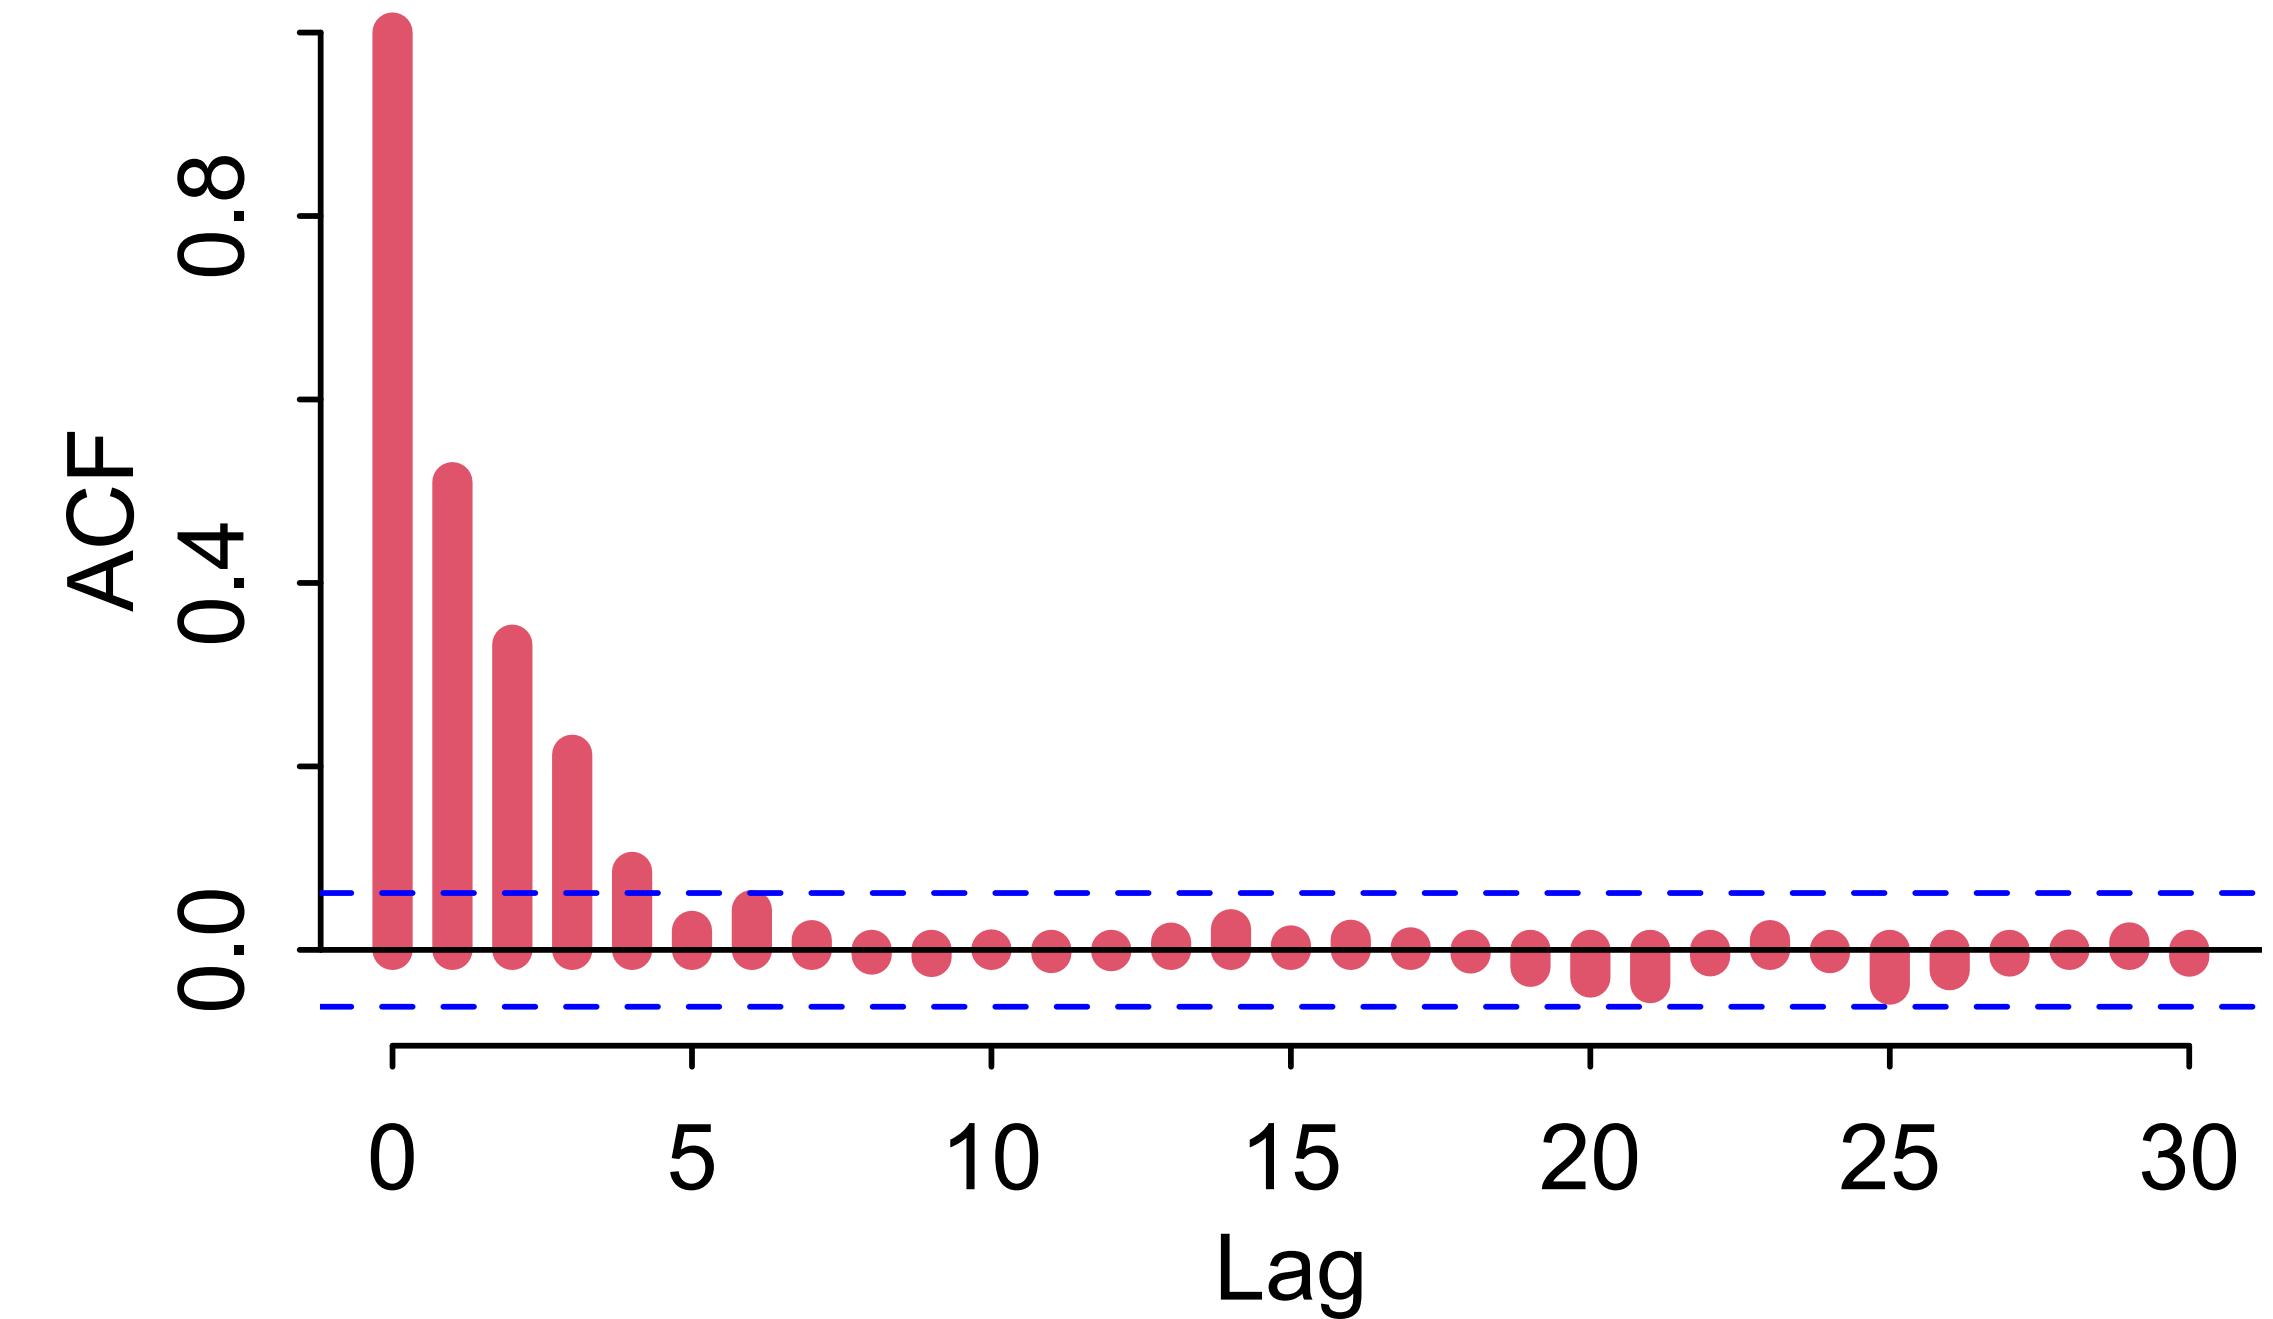
n_eff

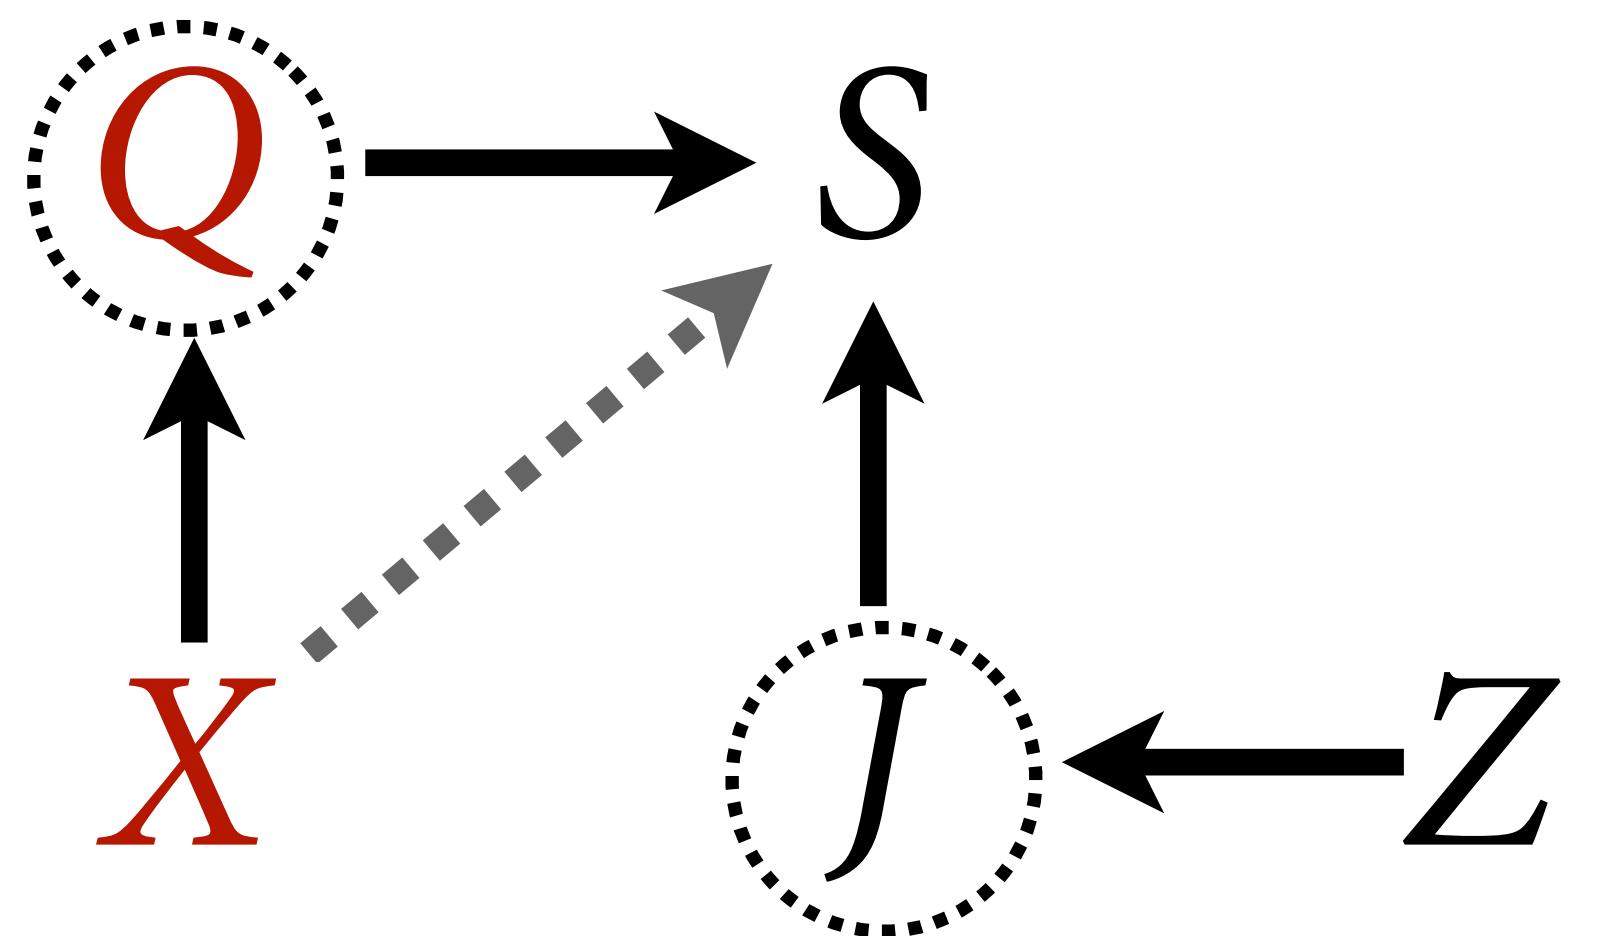
Estimate of number of **effective samples**

“How long would the chain be, if each sample was independent of the one before it?”

When samples are **autocorrelated**, you have fewer *effective* samples

	mean	sd	5.5%	94.5%	n_eff	Rhat4
Q[1]	0.14	0.30	-0.34	0.64	2962	1
Q[2]	0.11	0.32	-0.40	0.61	3033	1
Q[3]	0.27	0.31	-0.21	0.75	2608	1





```

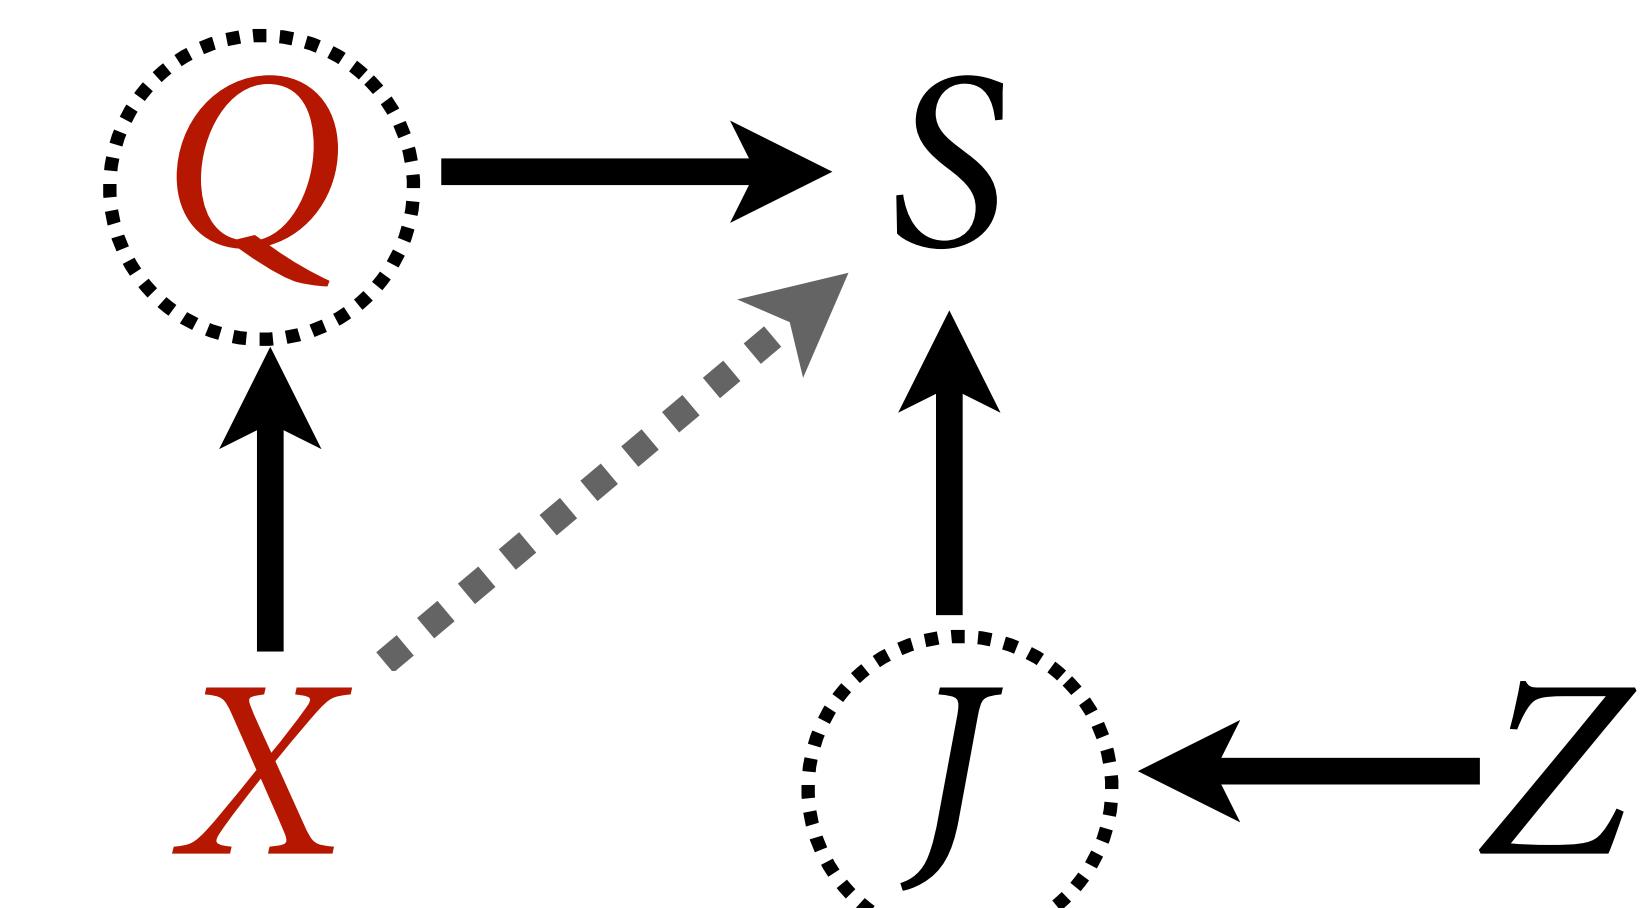
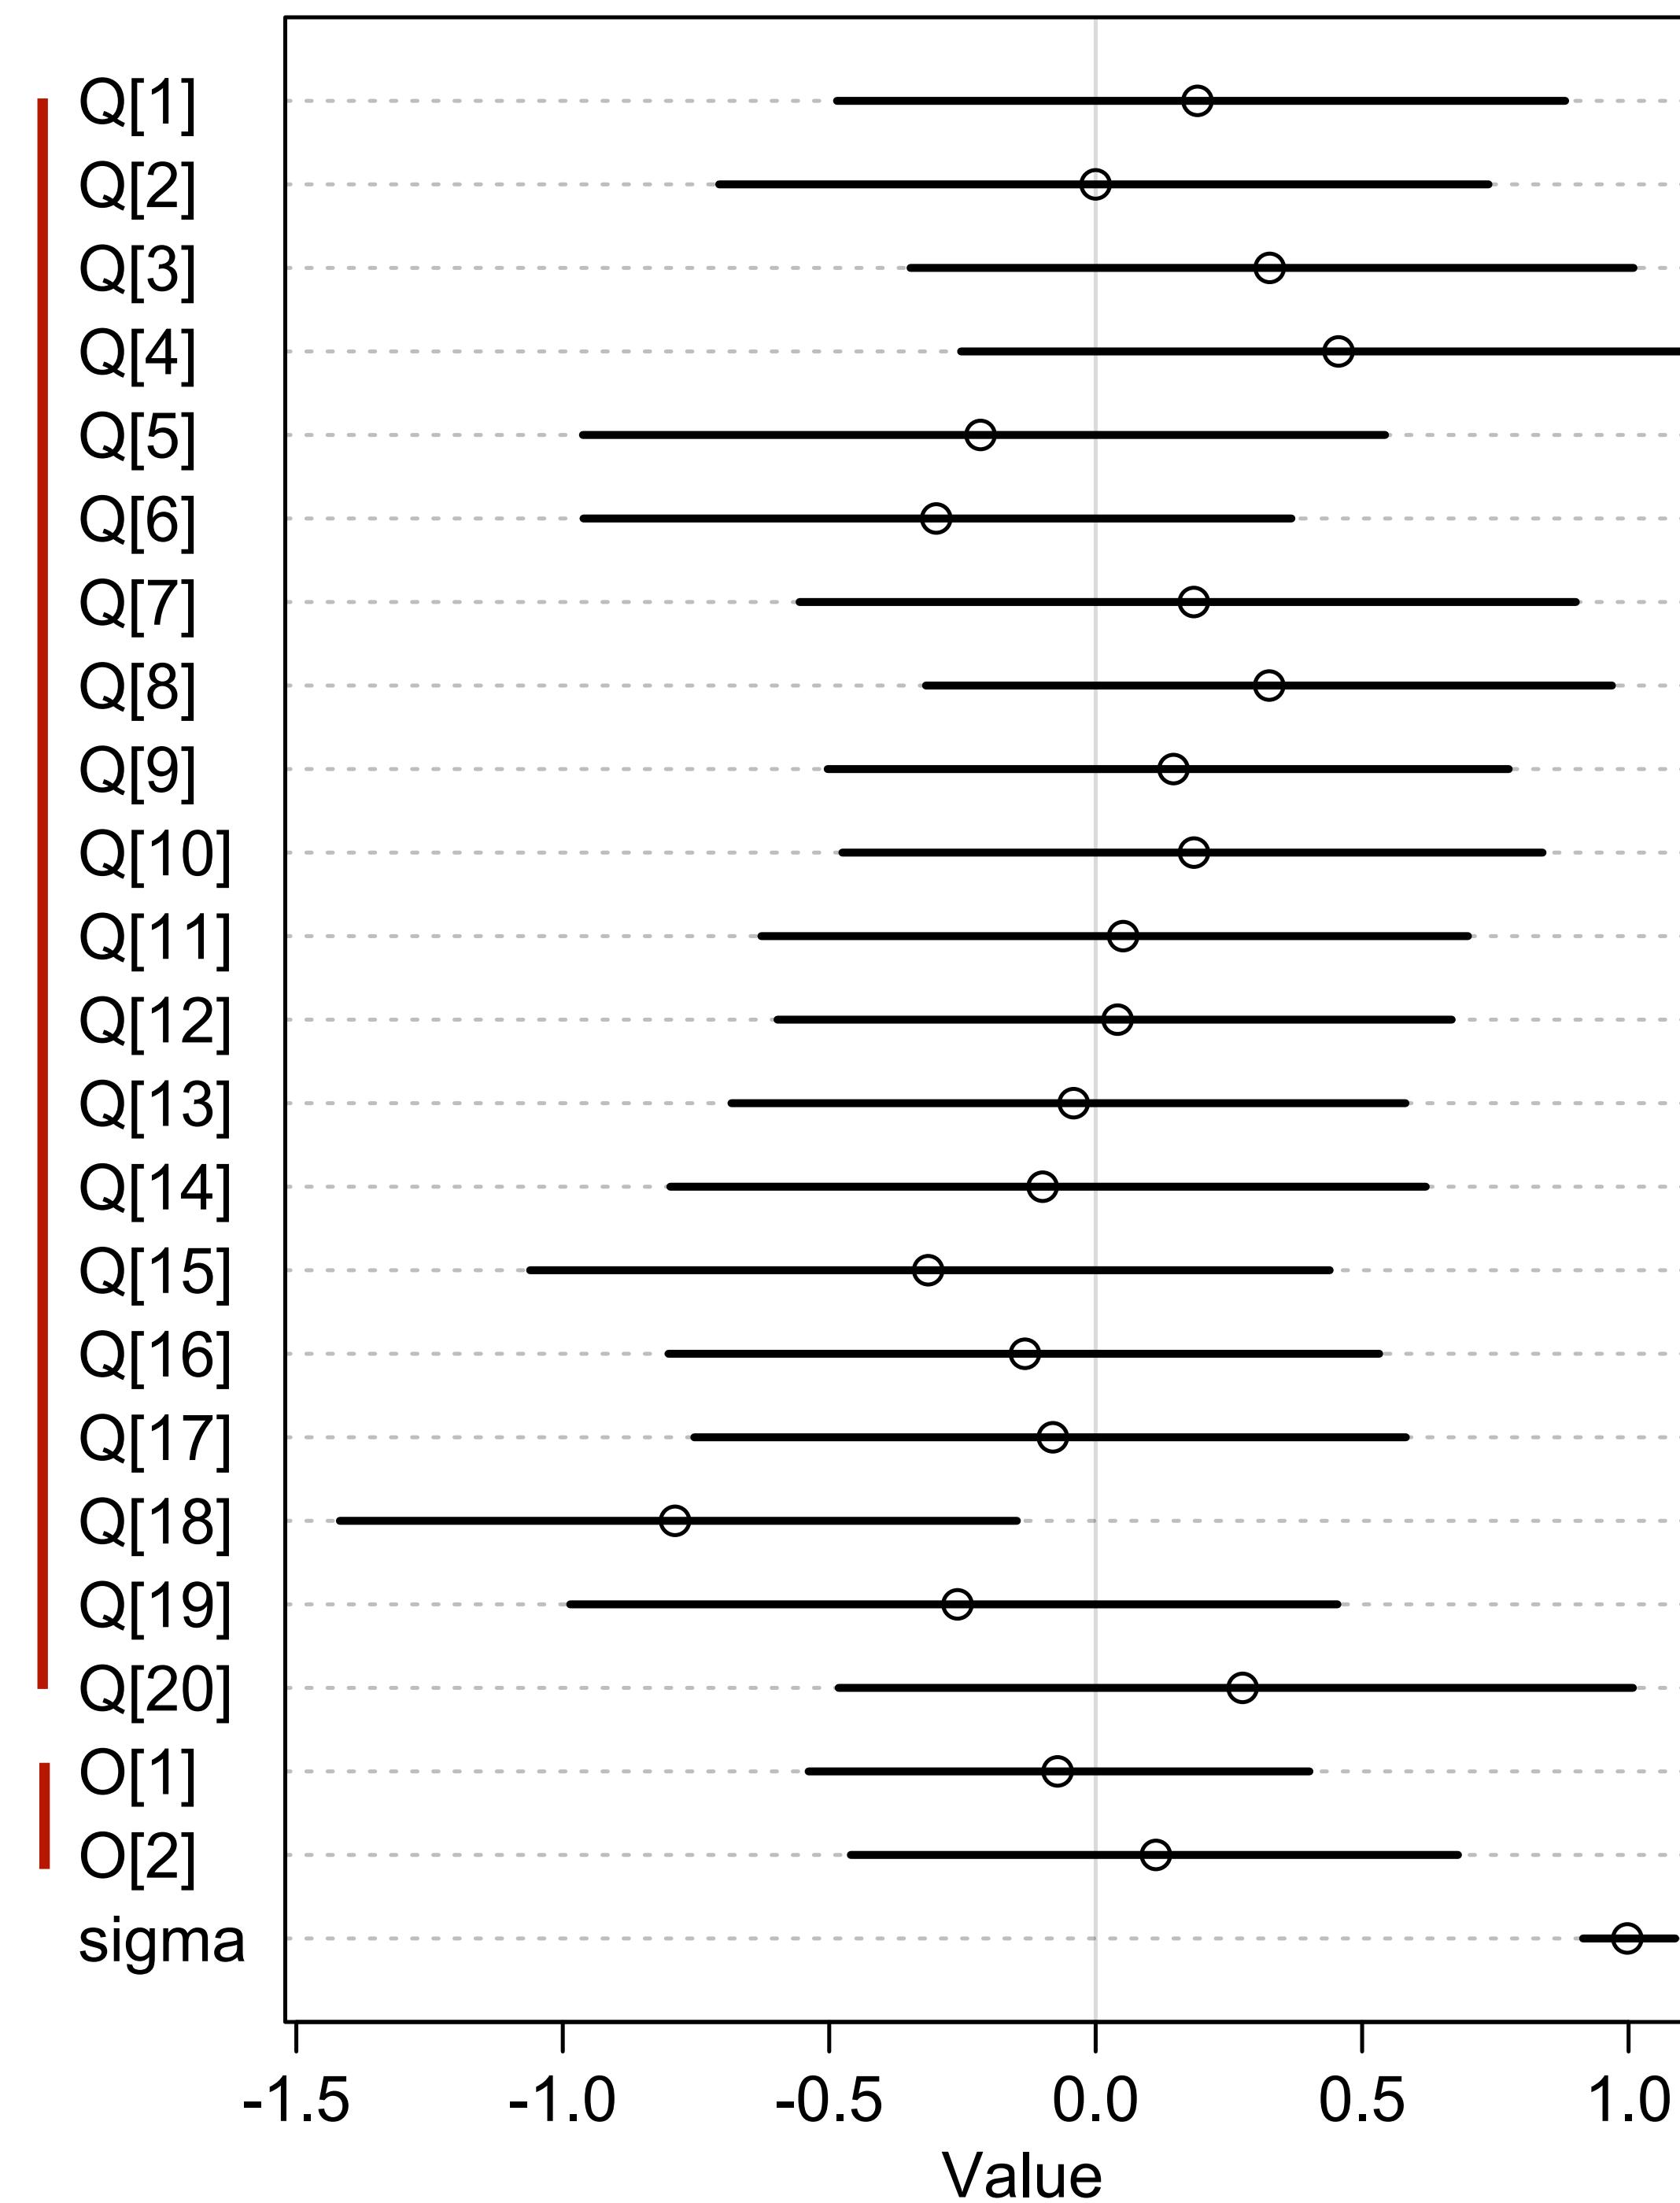
mQ0 <- ulam(
  alist(
    S ~ dnorm(mu,sigma),
    mu <- Q[W] + O[X],
    Q[W] ~ dnorm(0,1),
    O[X] ~ dnorm(0,1),
    sigma ~ dexp(1)
  ) , data=dat , chains=4 , cores=4 )
plot(precis(mQ0,2))

```

$$\begin{aligned}
S_i &\sim \text{Normal}(\mu_i, \sigma) \\
\mu_i &= Q_{W[i]} + O_{X[i]} \\
Q_j &\sim \text{Normal}(0, 1) \\
O_j &\sim \text{Normal}(0, 1) \\
\sigma &\sim \text{Exponential}(1)
\end{aligned}$$

wines

wines



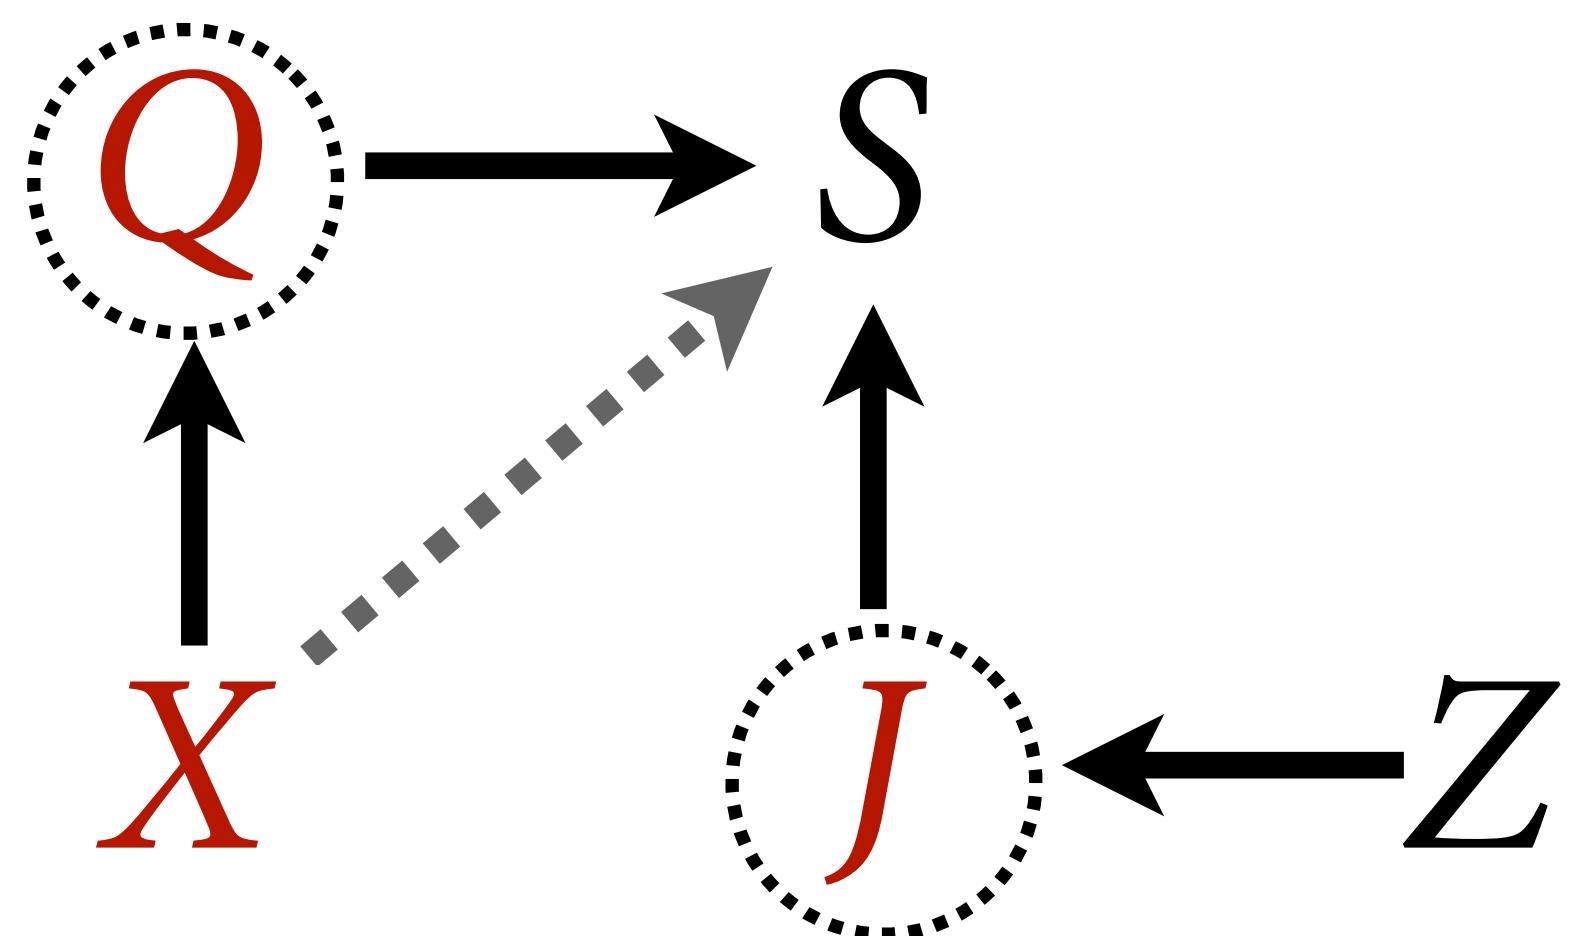
$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = Q_{W[i]} + O_{X[i]}$$

$$Q_j \sim \text{Normal}(0, 1)$$

$$O_j \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$



```
mQOJ <- ulam(
  alist(
    S ~ dnorm(mu,sigma),
    mu <- (Q[W] + O[X] - H[J])*D[J],
    Q[W] ~ dnorm(0,1),
    O[X] ~ dnorm(0,1),
    H[J] ~ dnorm(0,1),
    D[J] ~ dexp(1),
    sigma ~ dexp(1)
  ) , data=dat , chains=4 )
plot(precis(mQOJ,2))
```

$$\begin{aligned}
 S_i &\sim \text{Normal}(\mu_i, \sigma) \\
 \mu_i &= (Q_{W[i]} + O_{X[i]} - H_{J[i]})D_{J[i]} \\
 Q_j &\sim \text{Normal}(0, 1) \\
 O_j &\sim \text{Normal}(0, 1) \\
 H_j &\sim \text{Normal}(0, 1) \\
 D_j &\sim \text{Normal}(0, 1) \\
 \sigma &\sim \text{Exponential}(1)
 \end{aligned}$$

$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]})D_{J[i]}$$

expected score



$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]})D_{J[i]}$$

expected score

wine quality

$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]})D_{J[i]}$$

expected score

wine quality

origin

$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]}) D_{J[i]}$$

expected score

wine quality

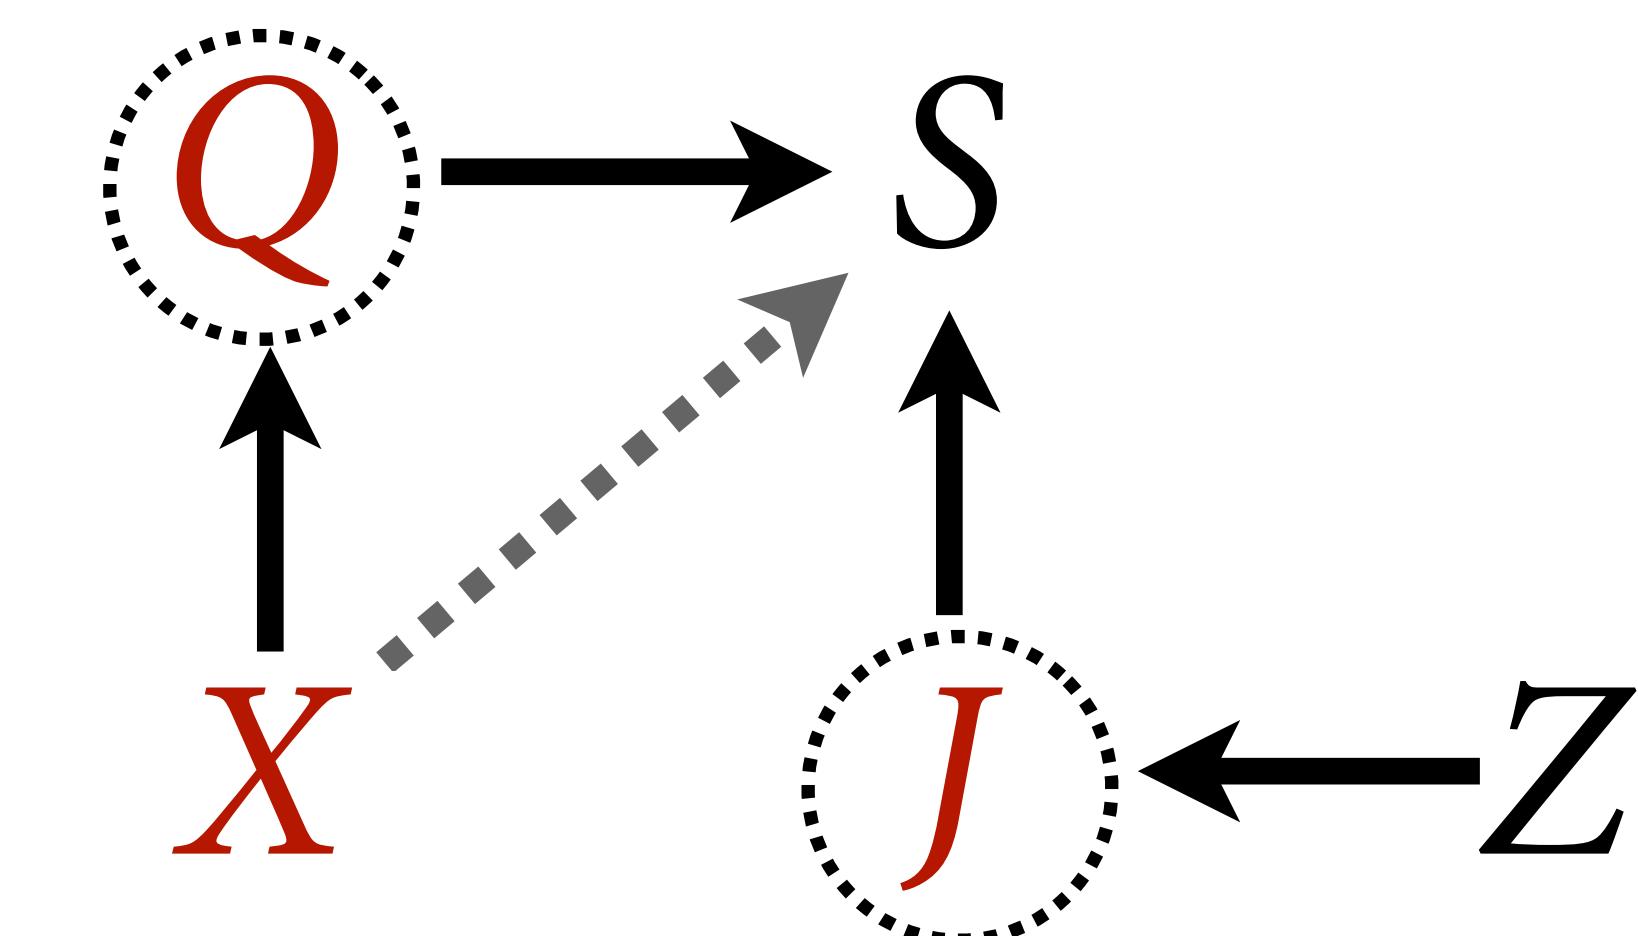
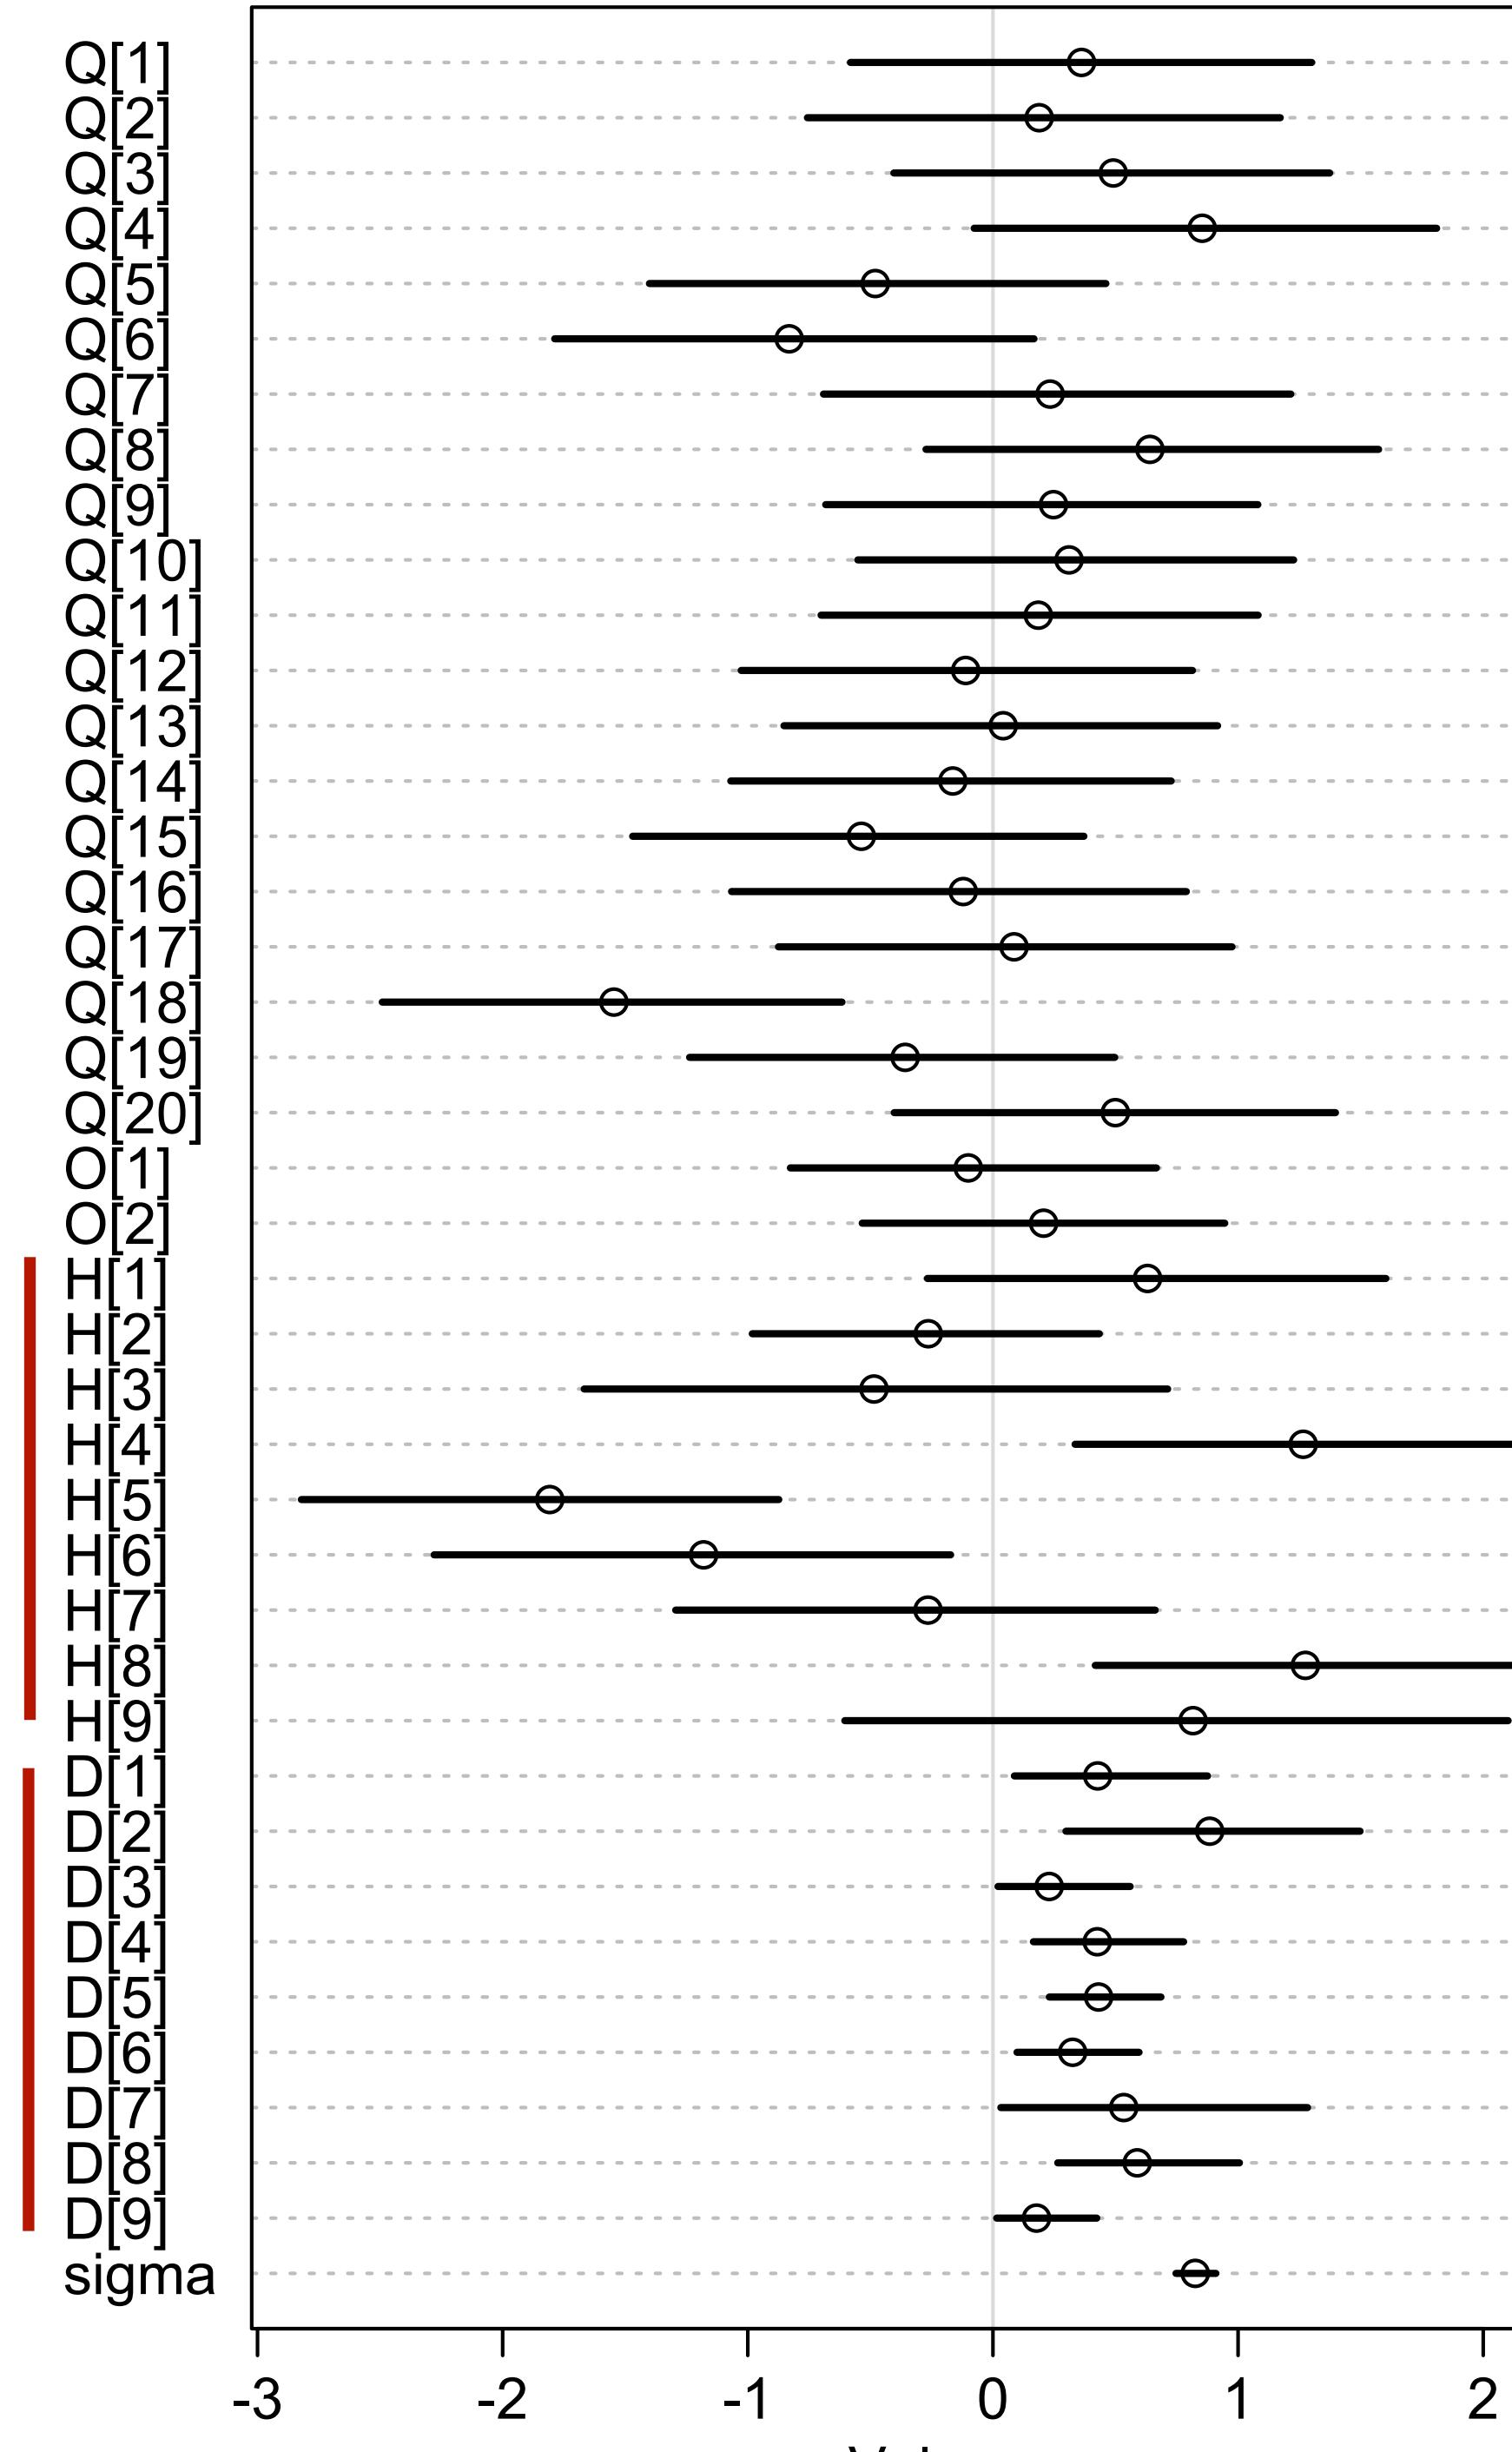
origin

*judge
harshness*

*judge
discrimination*

harshness

discrimination



$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]})D_{J[i]}$$

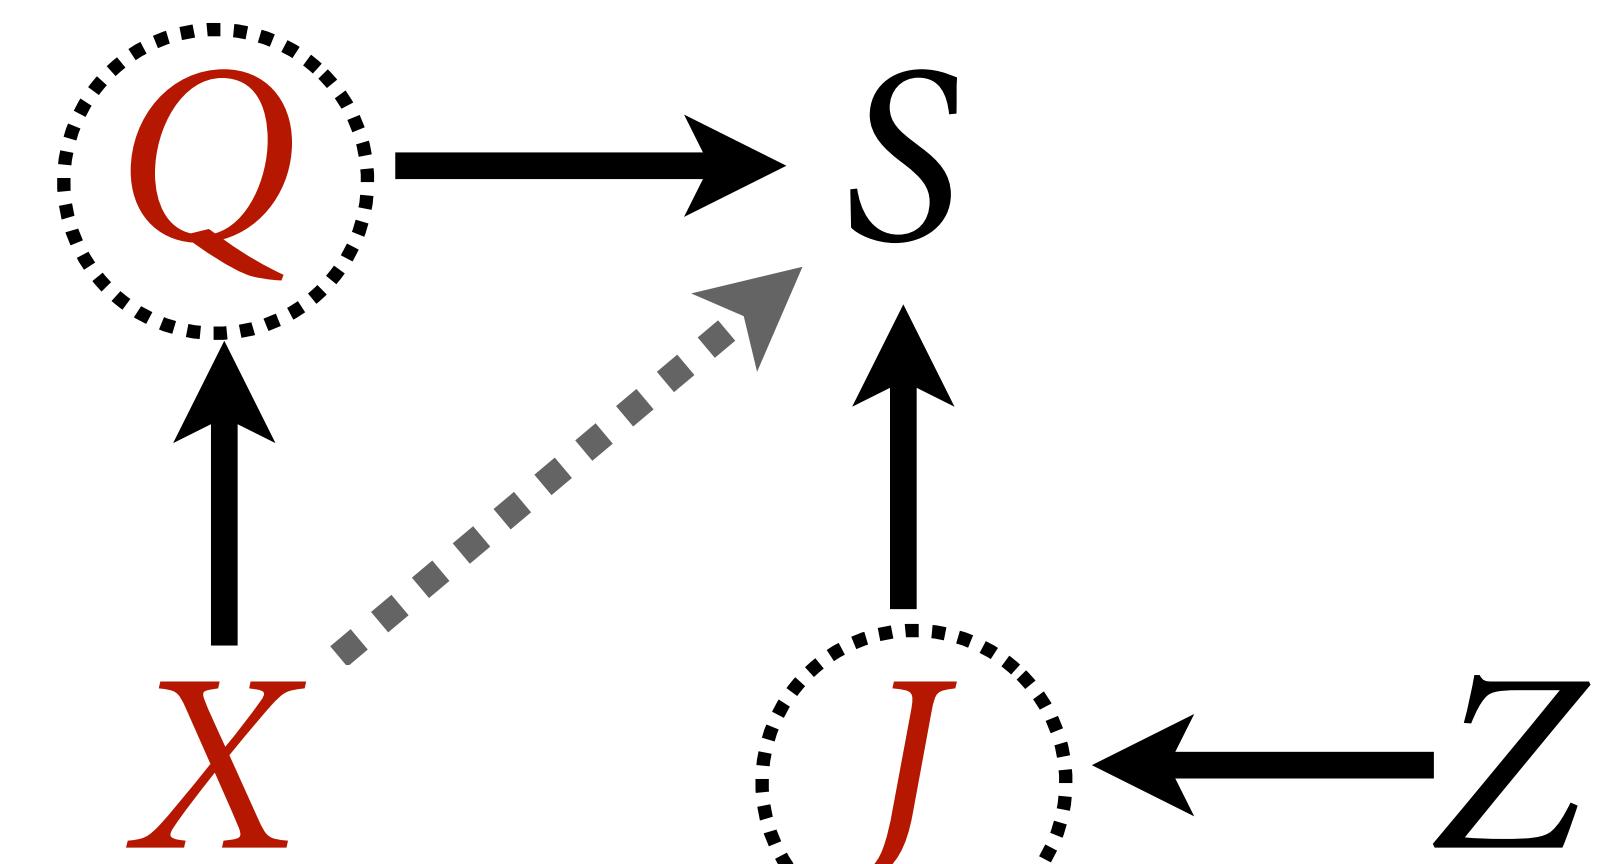
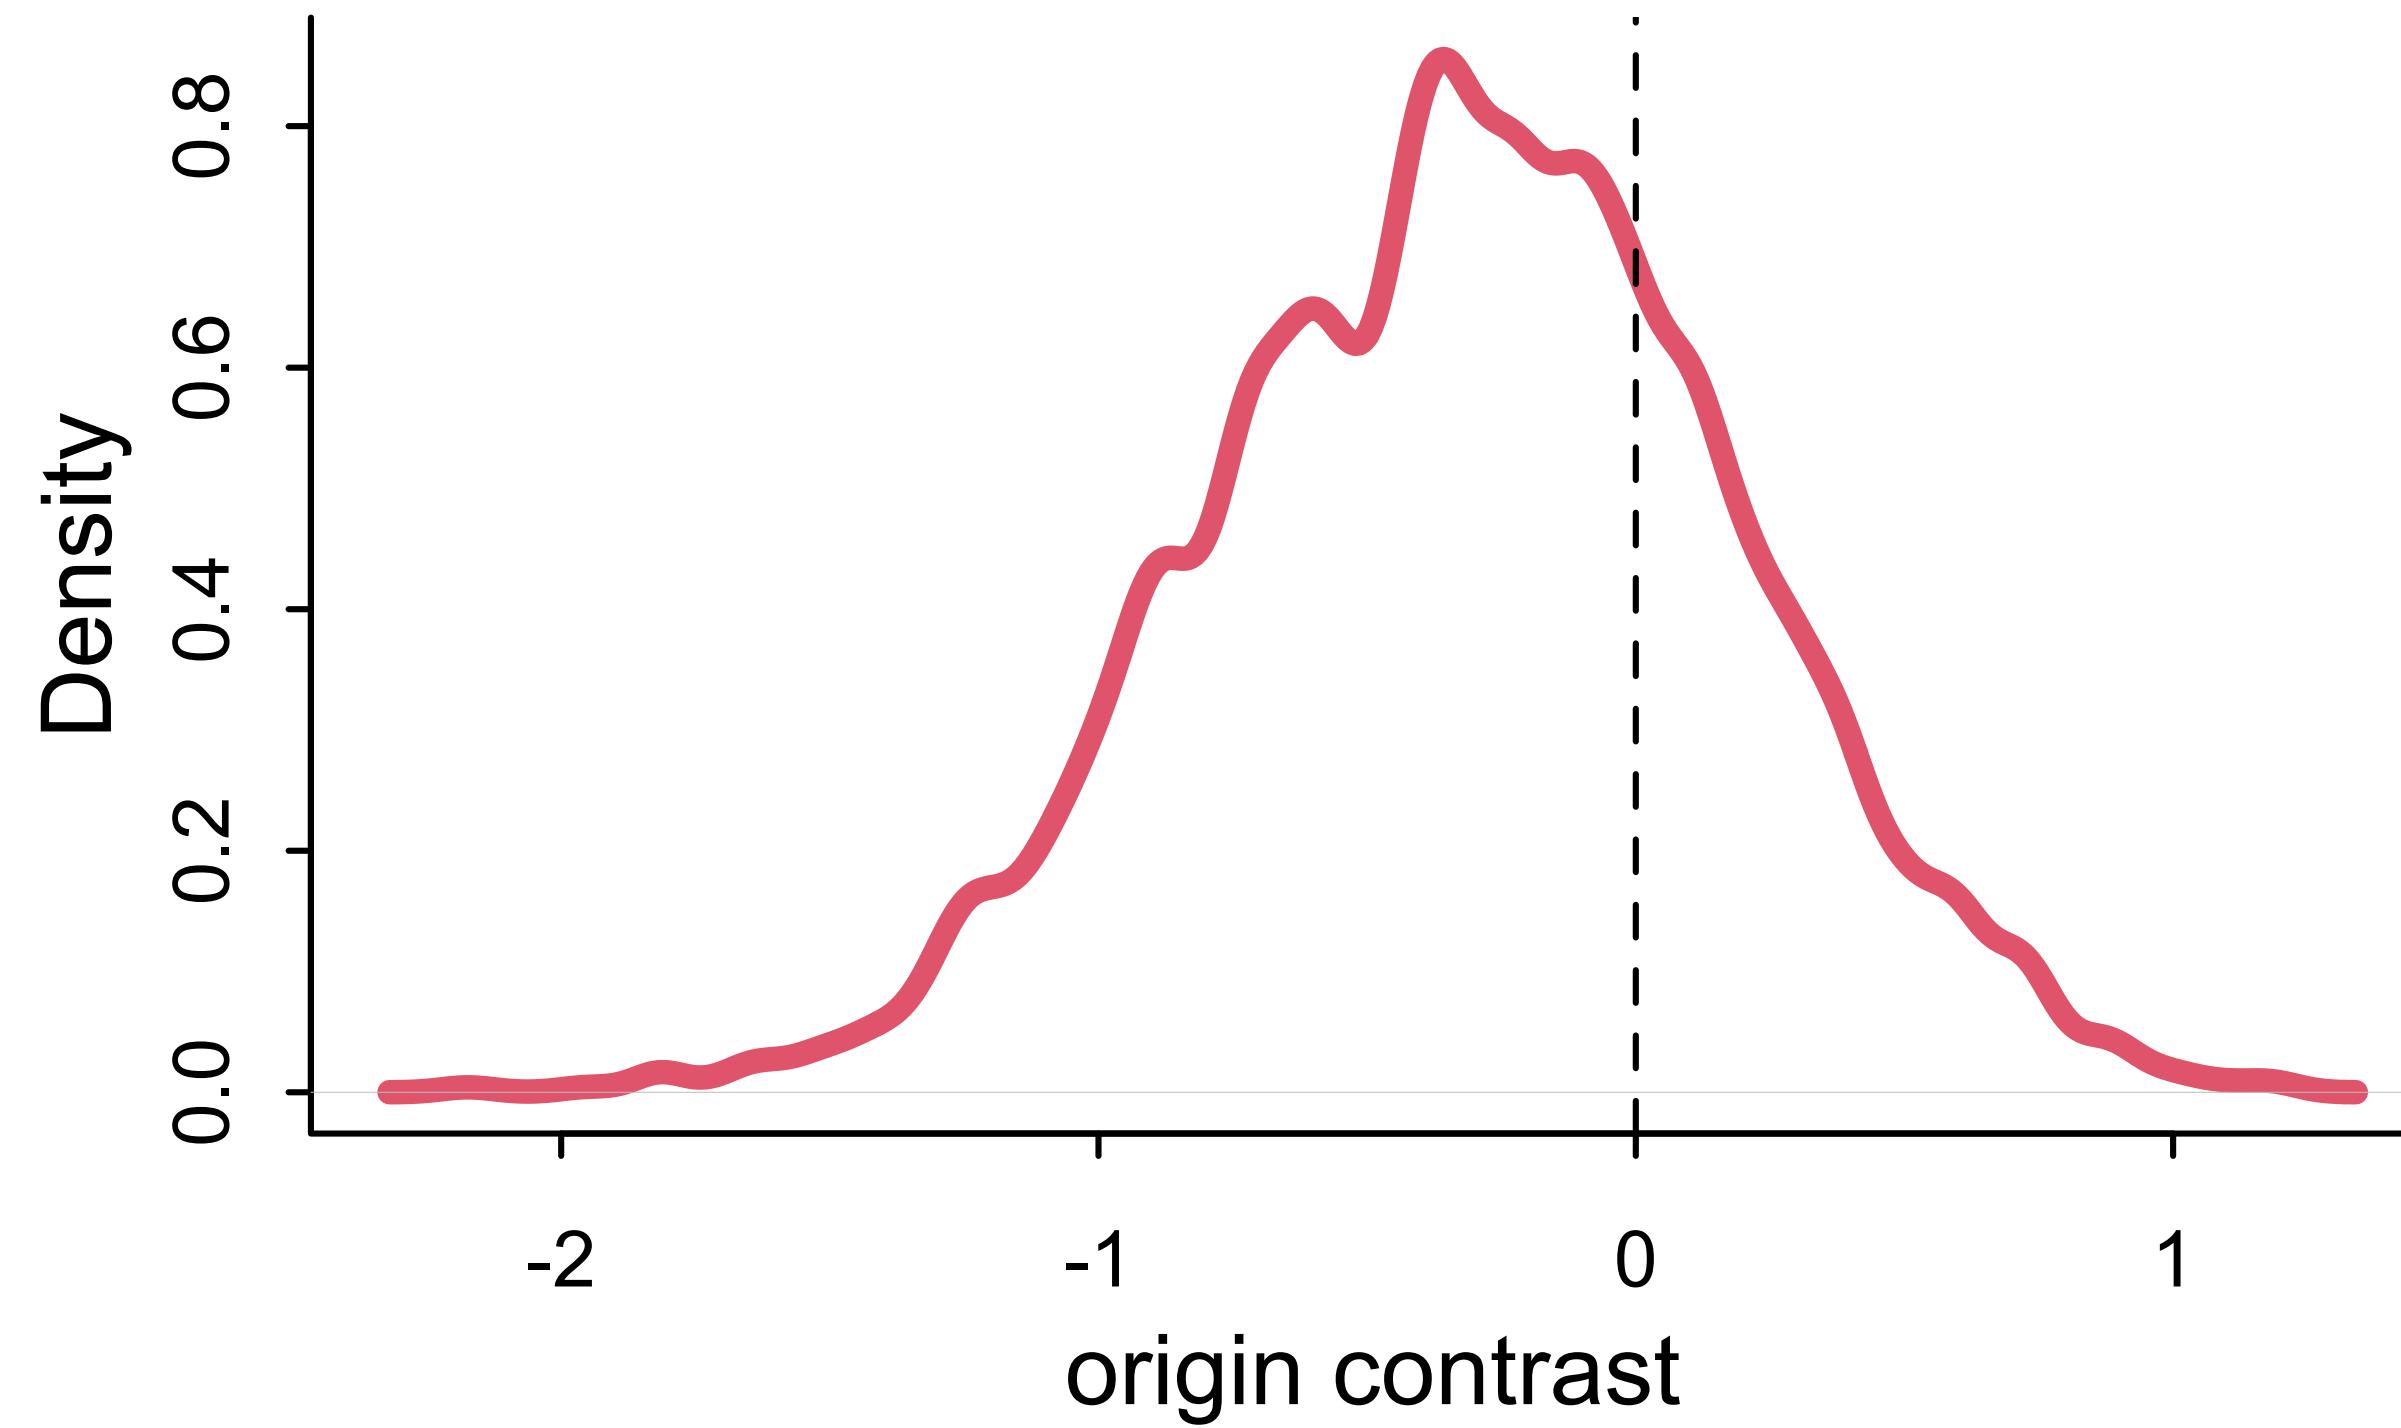
$$Q_j \sim \text{Normal}(0, 1)$$

$$O_j \sim \text{Normal}(0, 1)$$

$$H_j \sim \text{Normal}(0, 1)$$

$$D_j \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$



$$S_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = (Q_{W[i]} + O_{X[i]} - H_{J[i]}) D_{J[i]}$$

$$Q_j \sim \text{Normal}(0, 1)$$

$$O_j \sim \text{Normal}(0, 1)$$

$$H_j \sim \text{Normal}(0, 1)$$

$$D_j \sim \text{Normal}(0, 1)$$

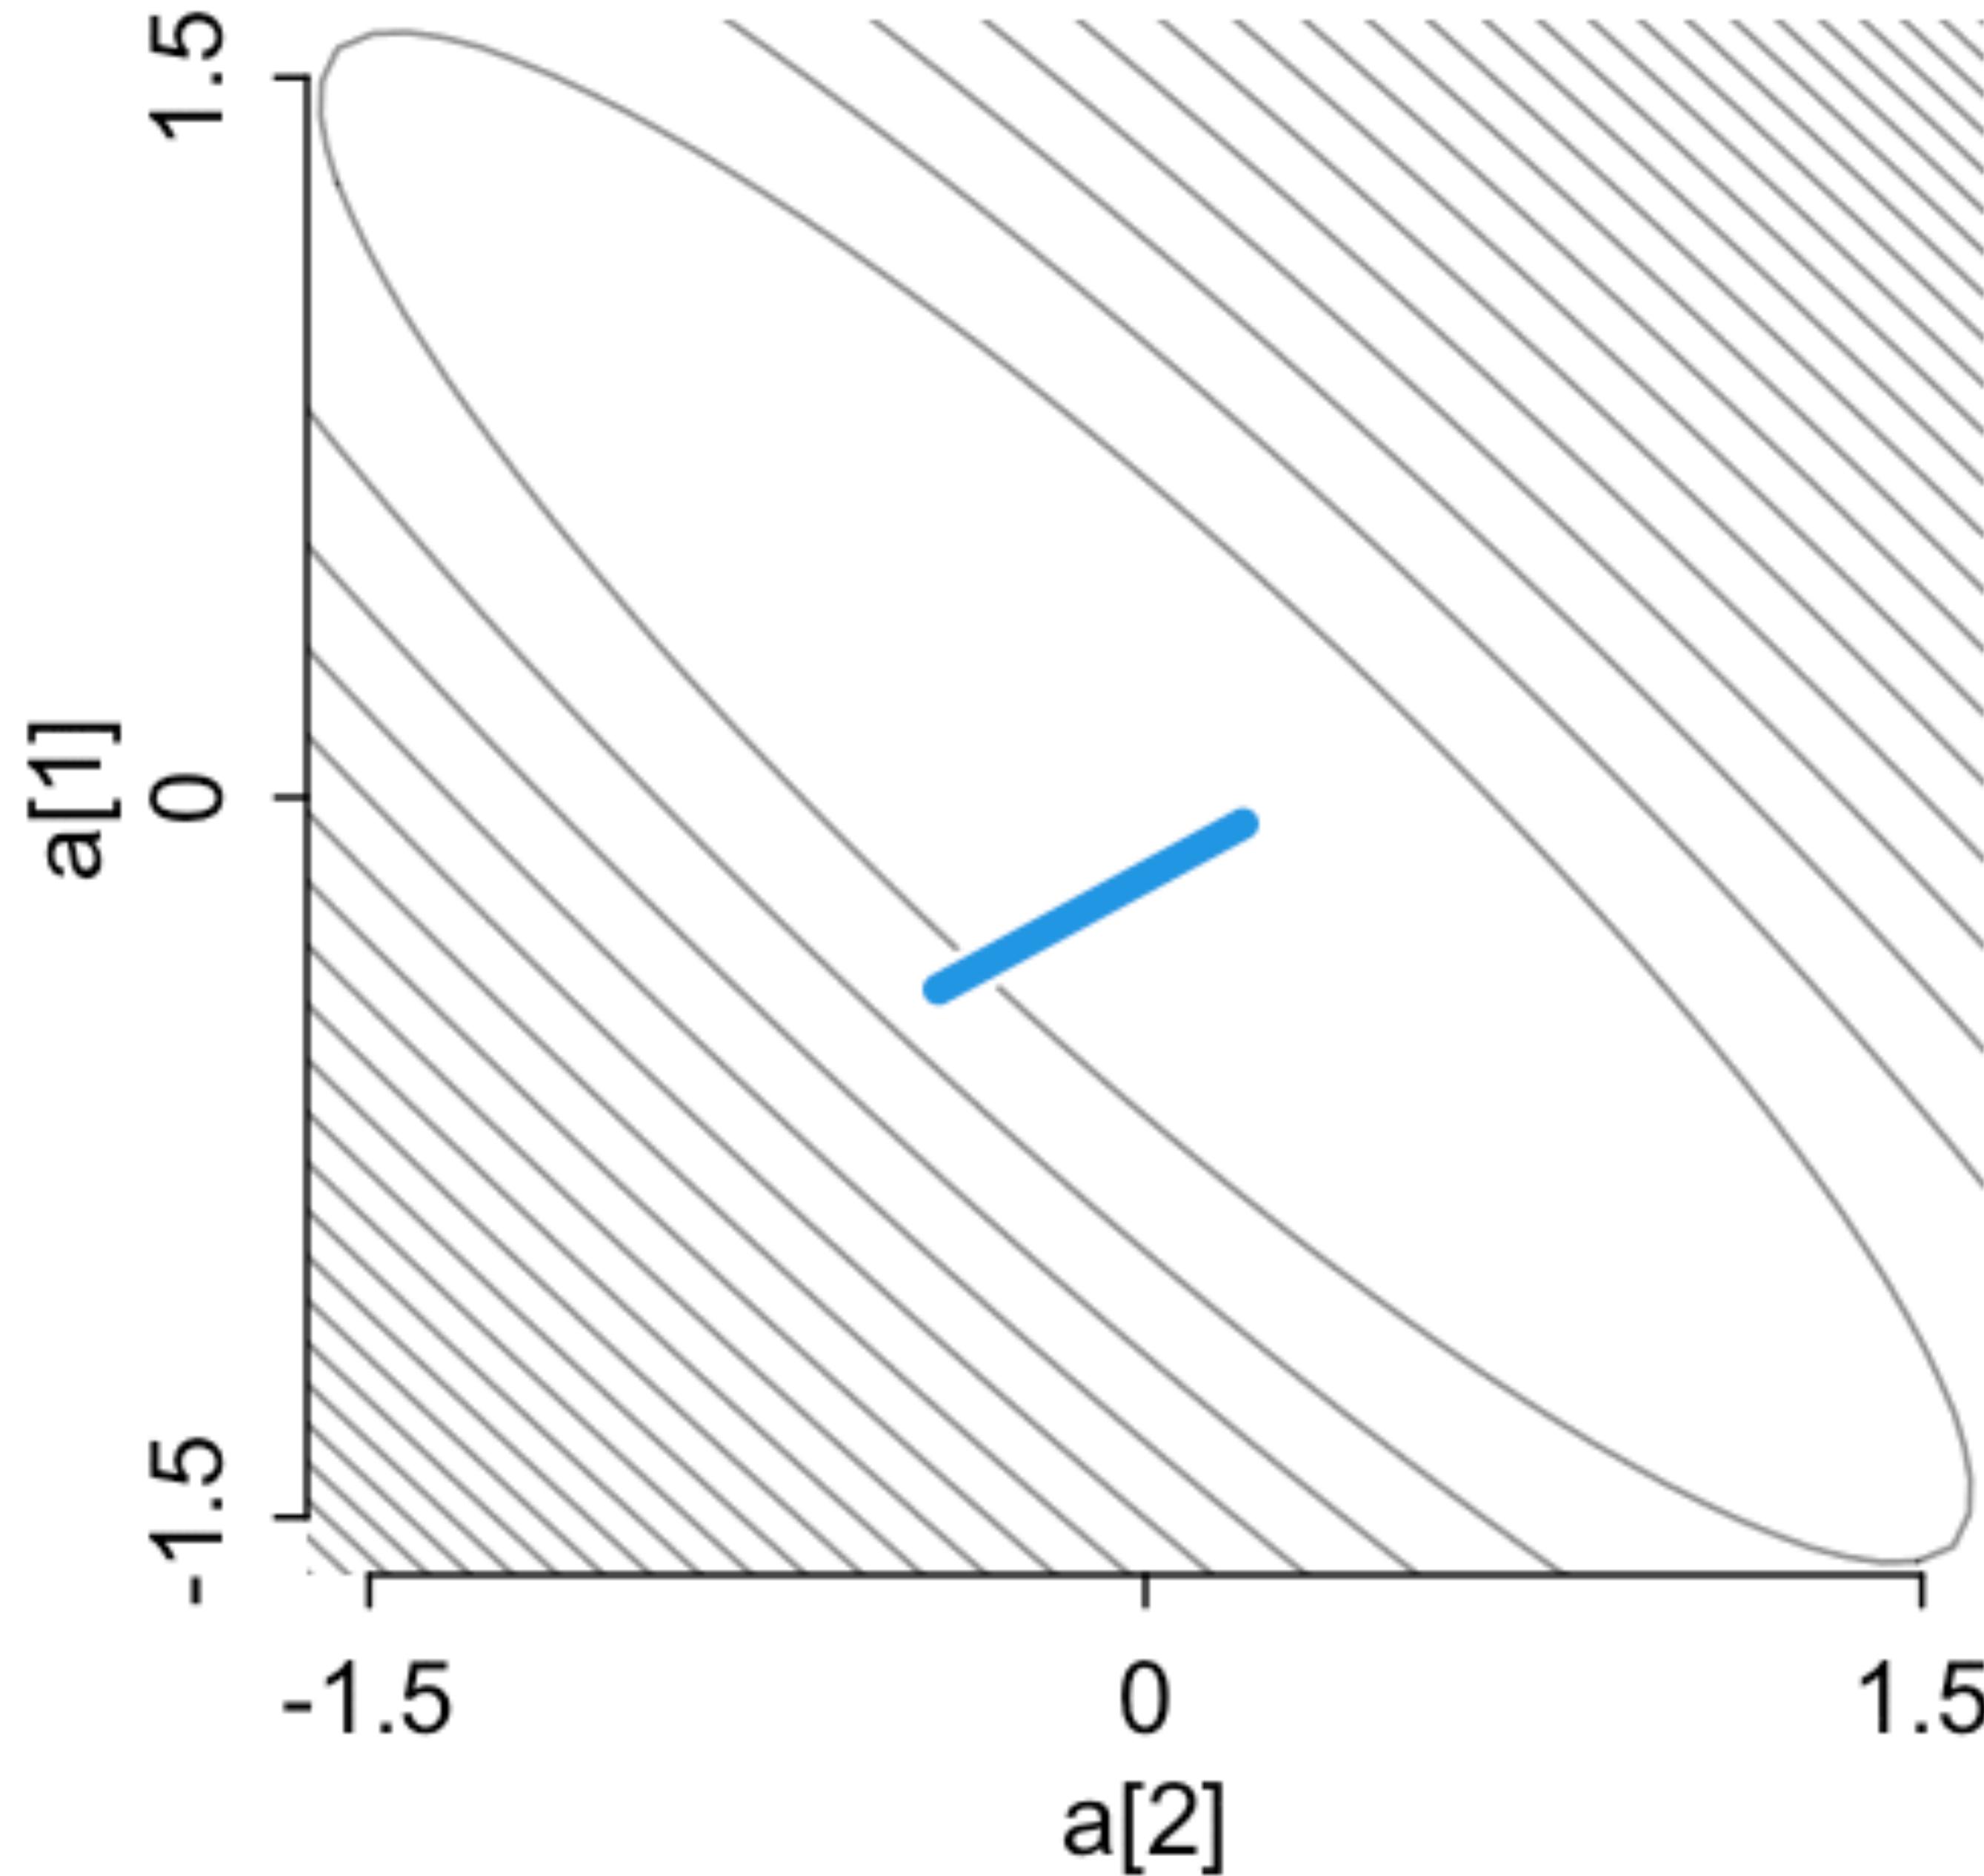
$$\sigma \sim \text{Exponential}(1)$$

The Folk Theorem of Statistical Computing

“When you have computational problems, often there’s a problem with your model.”

Andrew Gelman
Spider-Man of Bayesian data analysis





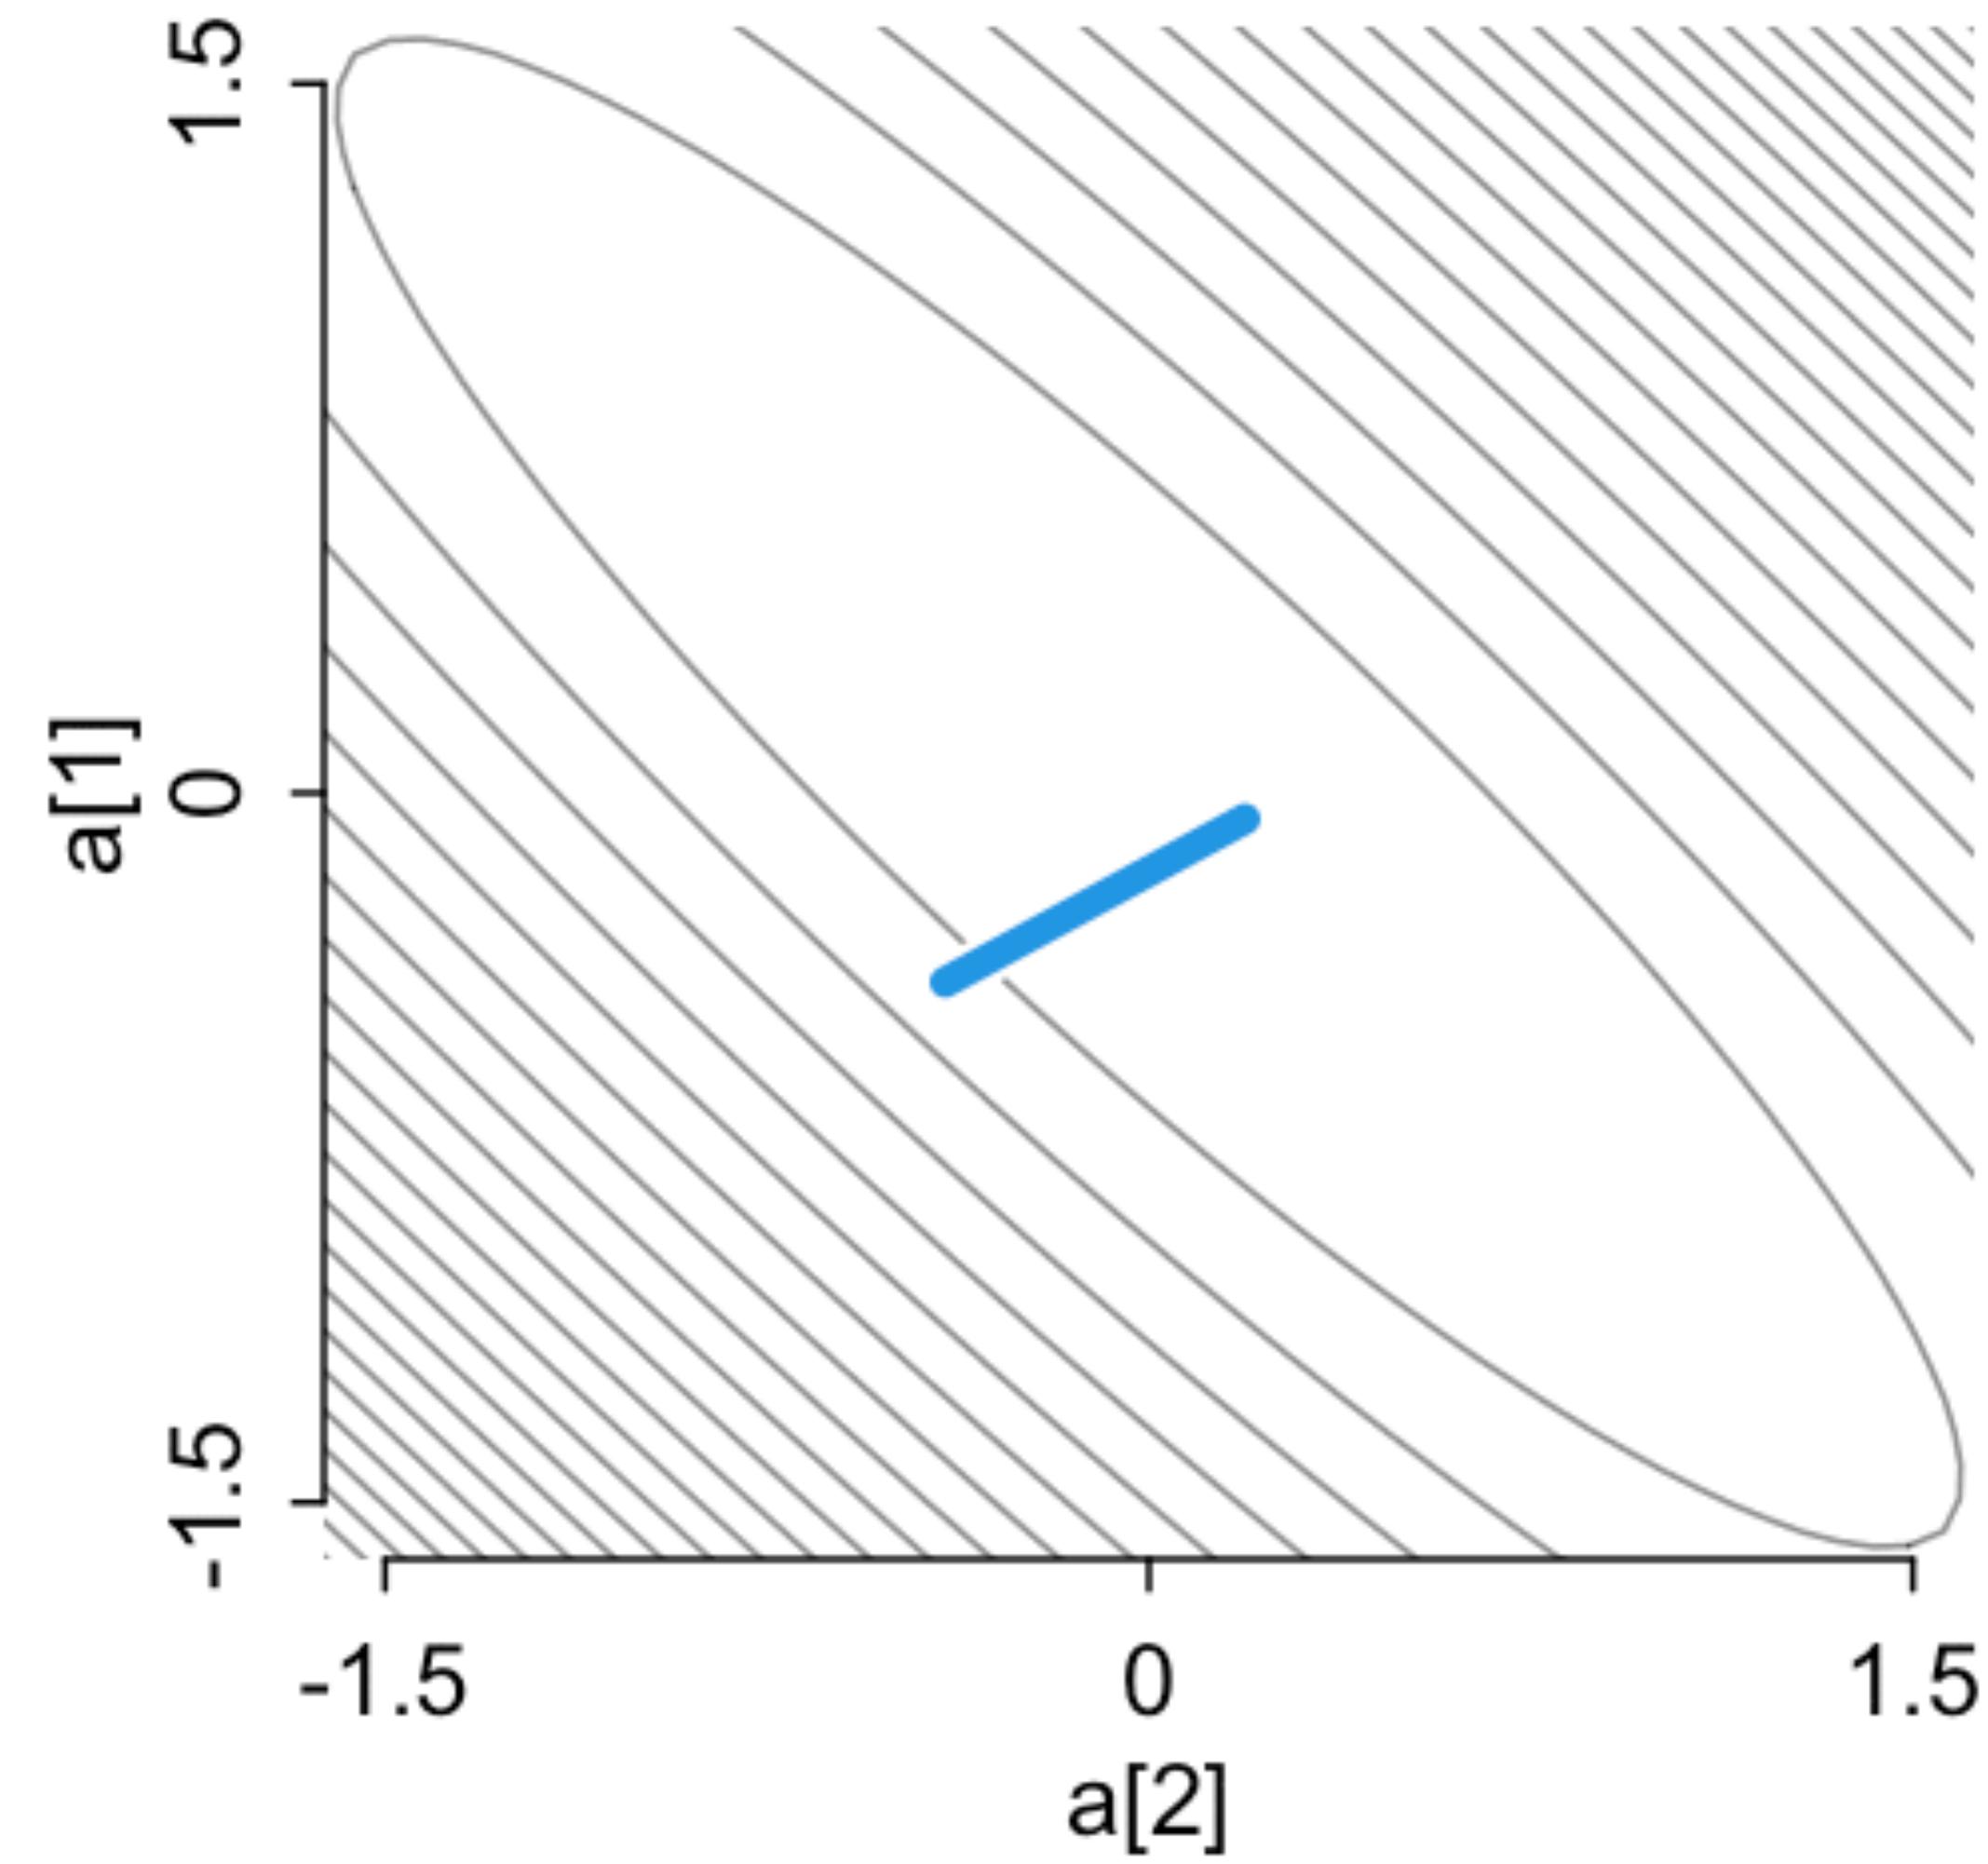
Divergent transitions

Divergent transition: A kind of rejected proposal

Simulation *diverges* from true path

Many DTs: poor exploration & possible bias

Will discuss again in later lecture



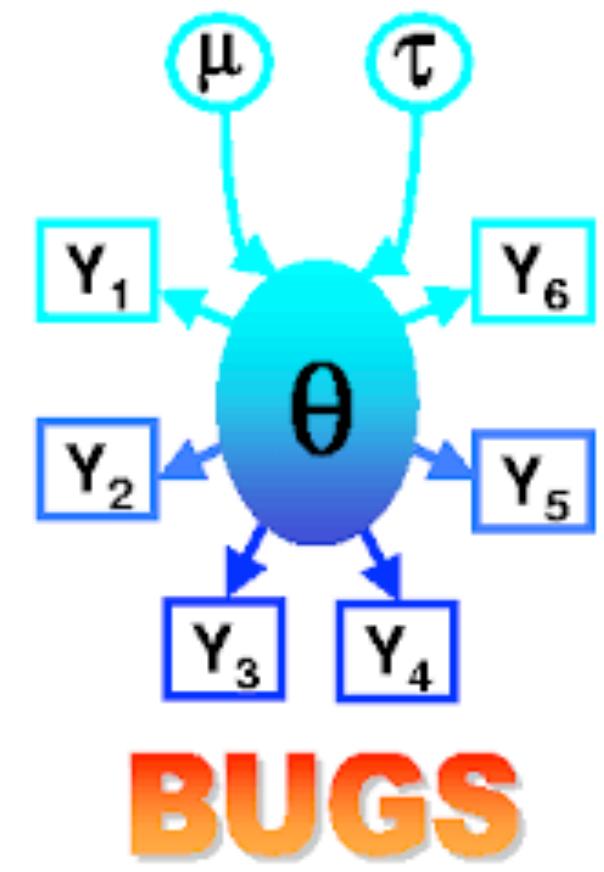
El Pueblo Unido

Desktop MCMC has been a revolution in scientific computing

Custom scientific modeling

High-dimension

Propagate measurement error



Sir David Spiegelhalter

Course Schedule

Week 1	Bayesian inference	Chapters 1, 2, 3
Week 2	Linear models & Causal Inference	Chapter 4
Week 3	Causes, Confounds & Colliders	Chapters 5 & 6
Week 4	Overfitting / MCMC	Chapters 7, 8, 9
Week 5	Generalized Linear Models	Chapters 10, 11
Week 6	Integers & Other Monsters	Chapters 11 & 12
Week 7	Multilevel models I	Chapter 13
Week 8	Multilevel models II	Chapter 14
Week 9	Measurement & Missingness	Chapter 15
Week 10	Generalized Linear Madness	Chapter 16

https://github.com/rmcelreath/stat_rethinking_2023

